



CMS Status and Needs

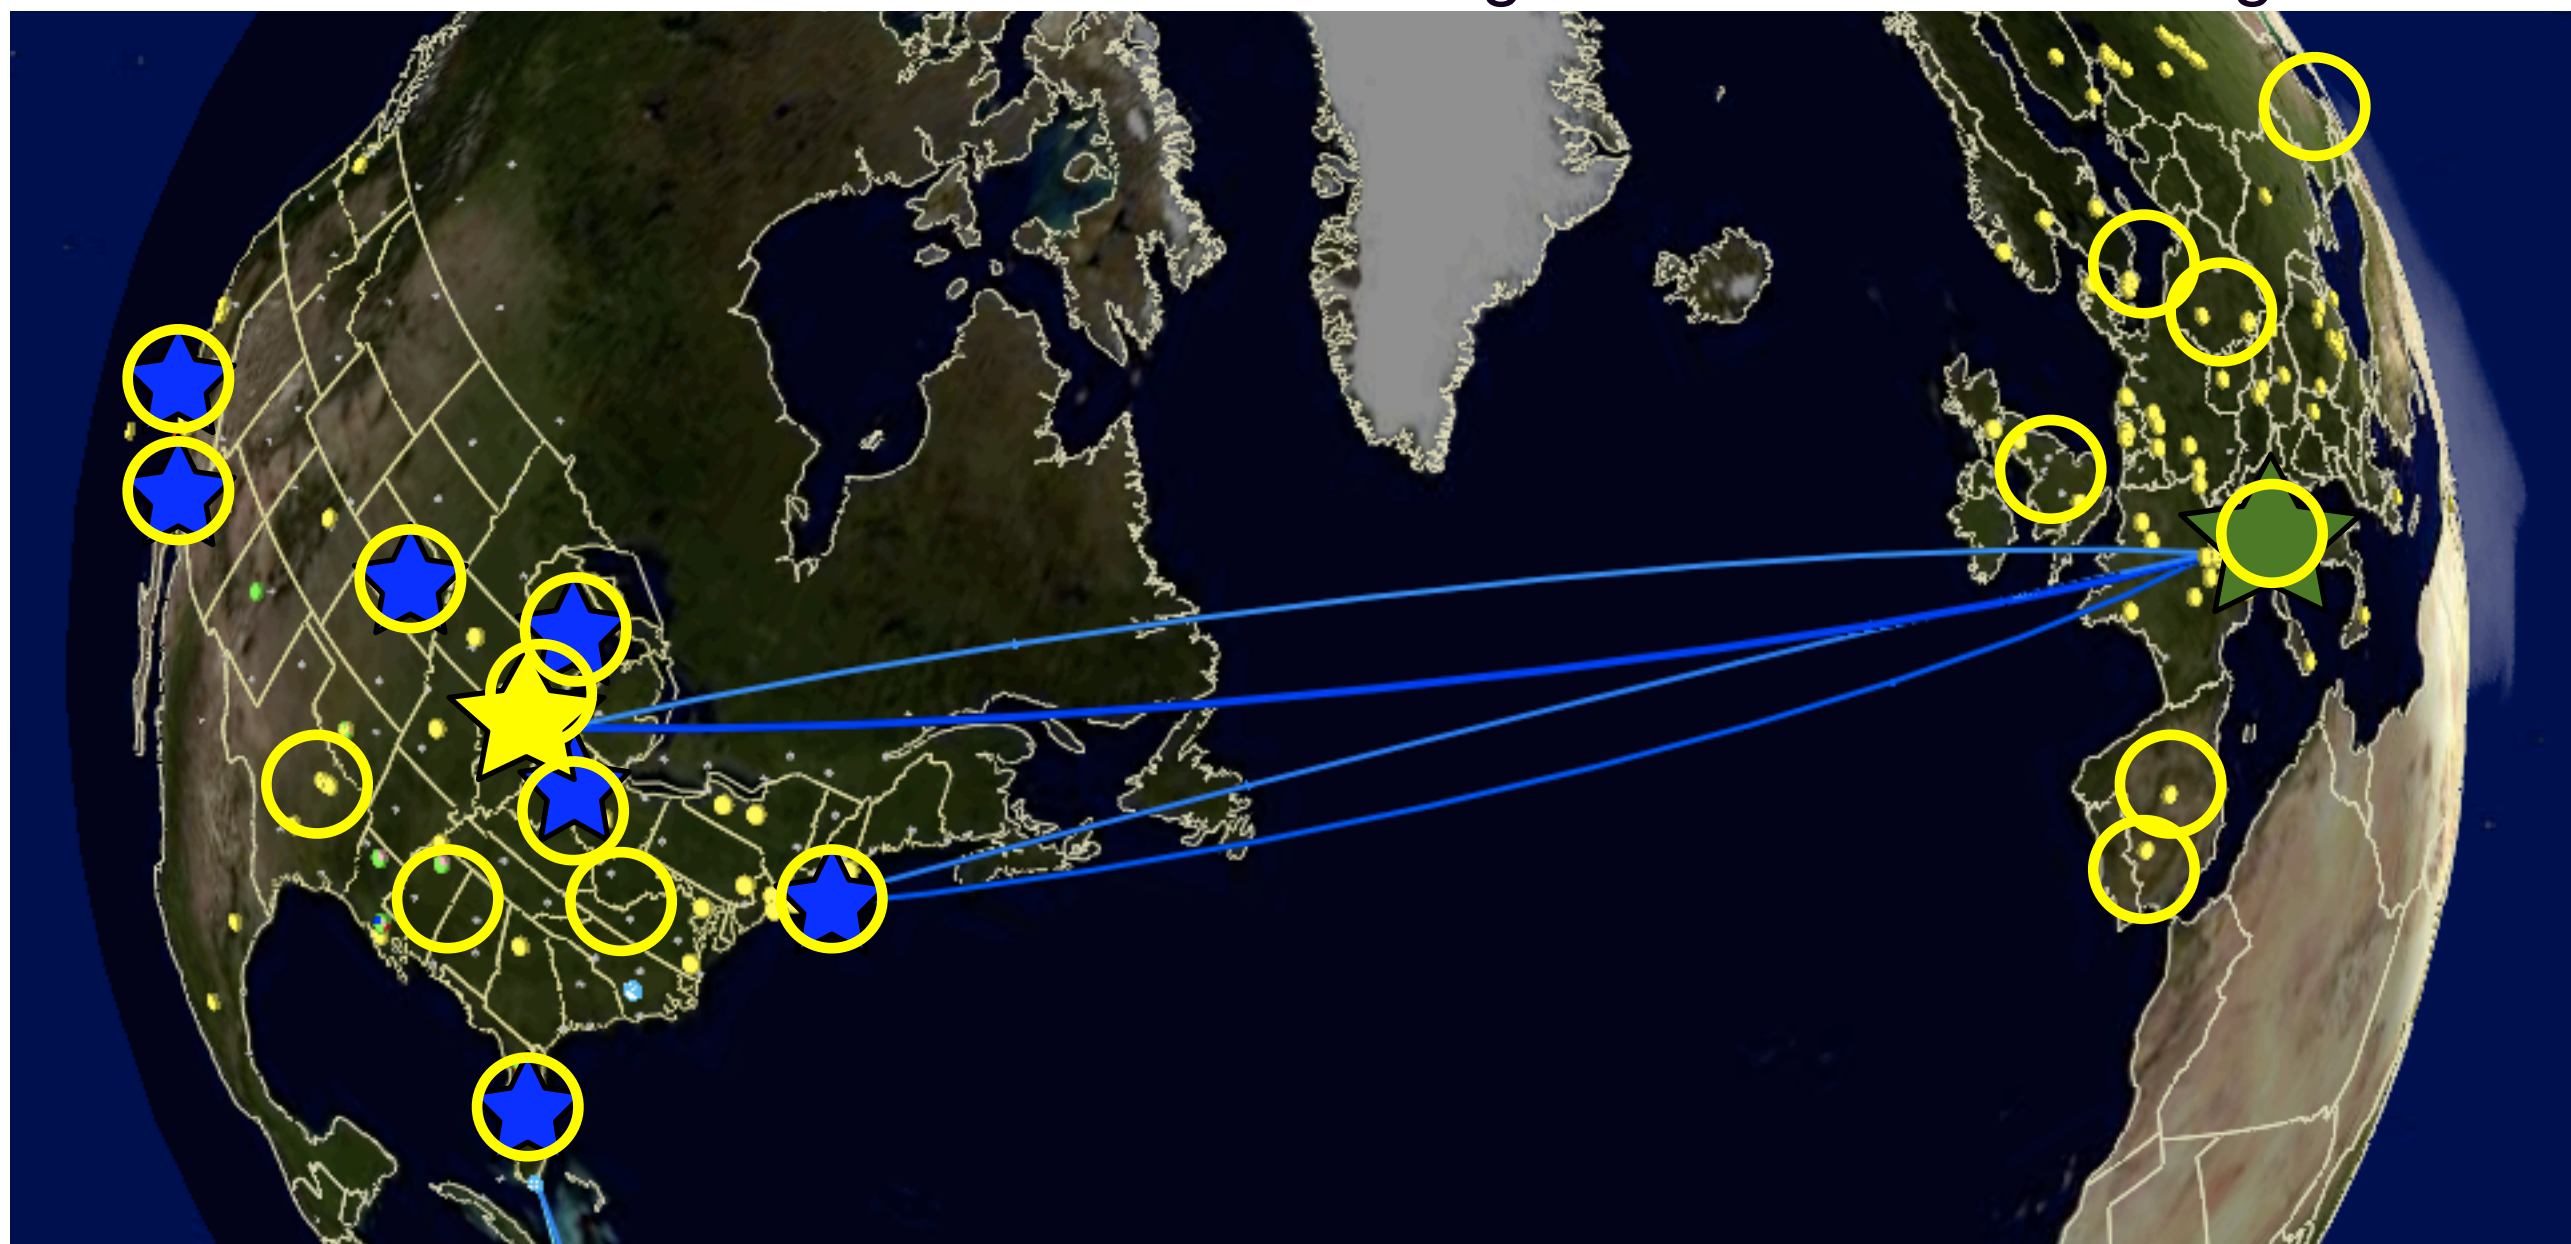
Open Science Grid Council Meeting

L.A.T.Bauerdick/Fermilab



Worldwide LHC Computing Grid

- ★ OSG connected to Grids in Europe and Asia
- ★ U.S. CMS: Tier-1 at Fermilab, 7 Tier-2 centers, ~9 Tier-3 centers
- ★ Grid connectivity to access world-wide distributed data at 7 Tier-1 centers
- ★ Put to real-world tests with Data Challenges and first data taking






First Beam-induced Events in CMS

$\sim 2 \cdot 10^9$ protons on collimator ~ 150 m upstream of CMS


Sept. 10, 2008
Beam Splash





Massive Data Taking of Cosmic Events

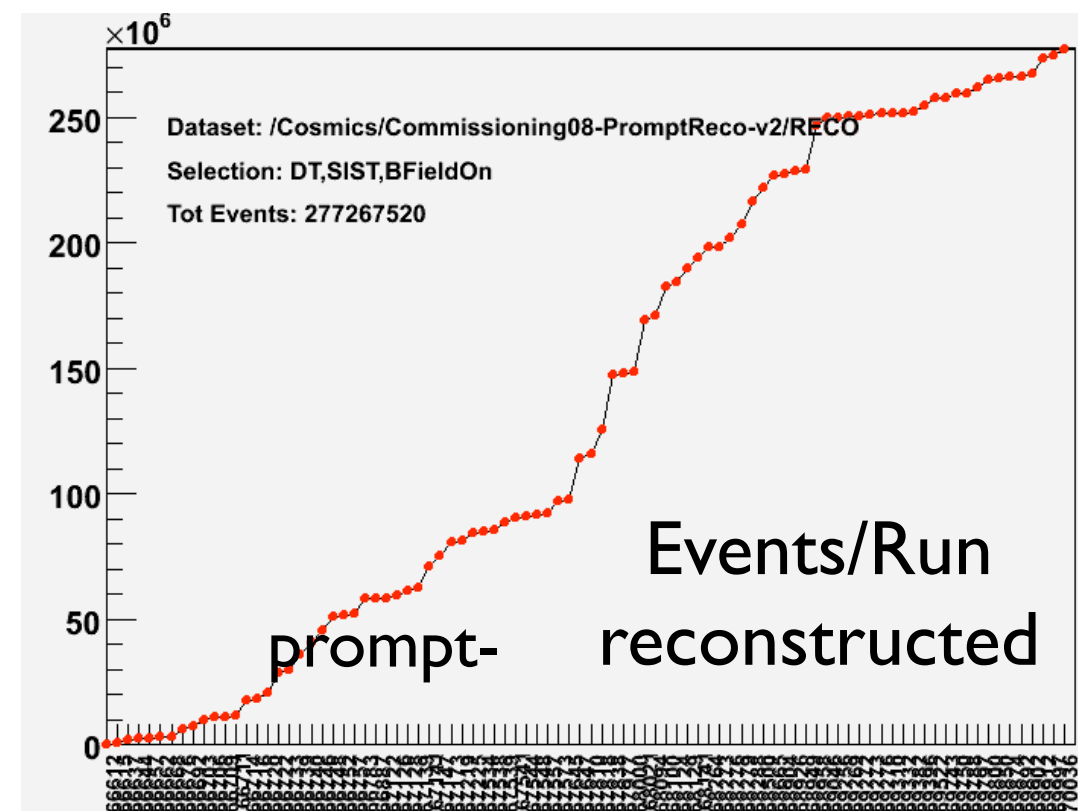

Run 66748, Event 8900172, LS 160, Orbit 167345832, BX 2011





Initial Running Experience


- ◆ Throughout 2008, S&C systems getting successively ready for data
 - ★ data/computing challenges, global runs, first beam, cosmics run w/ field
- ◆ S&C services behave rather well under data taking conditions
 - ★ initial beam data taking fully successful
 - ★ chain of data and work flows from HLT farm through Tier-0/1/2/3 works
 - ★ continuous Cosmic data taking at 300Hz for 4 weeks, 355M events
 - ◆ 277M events available with full detector and magnetic field, reprocessed at Tier-1s





Data Serving on the Grid for Physics Analysis

- ◆ Fermilab well established as the major CMS data center, serving CMS processing needs and supporting U.S. data analysis
- ★ serving data to Tier-2, Tier-3 and LPC-CAF physics users





U.S. Computing Infrastructure: Tier-2 Centers

- ♦ U.S. Tier-2s extensively used for simulation and analysis processing.
 - ★ all the US Tier-2 sites reached the nominal capacity by this summer
 - ★ FY09 Tier-2 ramp to double storage and increase CPU by 50%
- ♦ Tier-2 program great success thanks to very engaged sites
 - ★ Nominal Resources per site:

US-CMS Tier-2 Summer 2008	CPU T2	1MSI2k	Tier-2 Processing Nodes
	Disk T2	200TB	dCache (200MB/s IO)
	Network	2.5-10Gb/s	WAN Networking
	People	2FTE	Supported Tier-2 Operations

- ♦ Organized sites in CMS to support physics analysis services for users
 - ★ associated sites with specific physics/detector performance groups
 - ★ allocated disk space for groups, and disk space for individual users



Using Tier-2s for CMS Physics

- ◆ Tier-1 to Tier-2 data transfers improved dramatically over last year
 - ★ Improvements in full mesh of Tier-1 to Tier-2
 - ★ Have successful demonstrations of 200MB/s - 400MB/s to Tier-2 sites
 - ★ High transfer rates enable Tier-2 storage as dynamic physics cache
- ◆ Simulation at Tier-2s is a well established operation
 - ★ Reached design rates for production
 - ★ Good spread across grids OSG and EGEE
- ◆ Analysis moving away from CERN/Tier-1 down to Tier-2 sites
 - ★ Analysis use and adoption of distributed computing tools is expanding
 - ★ Rates of job submissions have steadily increased
 - ★ Some work left to do to automate and harden the infrastructure
- ◆ Results of “Tier-2 round table” with physics customers:
 - ★ feedback from each of the physics groups and individual users solicited
 - ★ overall very positive feedback, individual issues are being dealt with
 - ★ assessed list of services that would make user experience better



University Computing: Access to U.S. CMS Facilities and CMS Data

- ◆ Number of Tier-3s is getting >> than number of Tier-2s
 - ★ Tier-3s require particular support, from CMS, OSG, community
 - ◆ S&C providing support as part of Grid services, with OSG
 - ◆ OSG has well-defined model how to integrate new Tier-3s
 - ◆ USCMS support CMS-specific services, like software installs
 - ◆ USCMS has hired dedicated Tier-3 support personnel
 - ◆ had started the effort with 25% of an FTE during the fall
- ◆ Strategy for including University computing facility relies on OSG
 - ★ large interest at Universities to bring their resources online to the Grid
 - ★ OSG is providing what sites need to succeed on the Grid
 - ★ setting up OSG enables sites to get access to CMS data and resources at Tier-2 and Tier-1
- ◆ OSG enabler to include University computing ad campus grids into CMS infrastructure, and to allow University scientist access to the large samples of CMS data

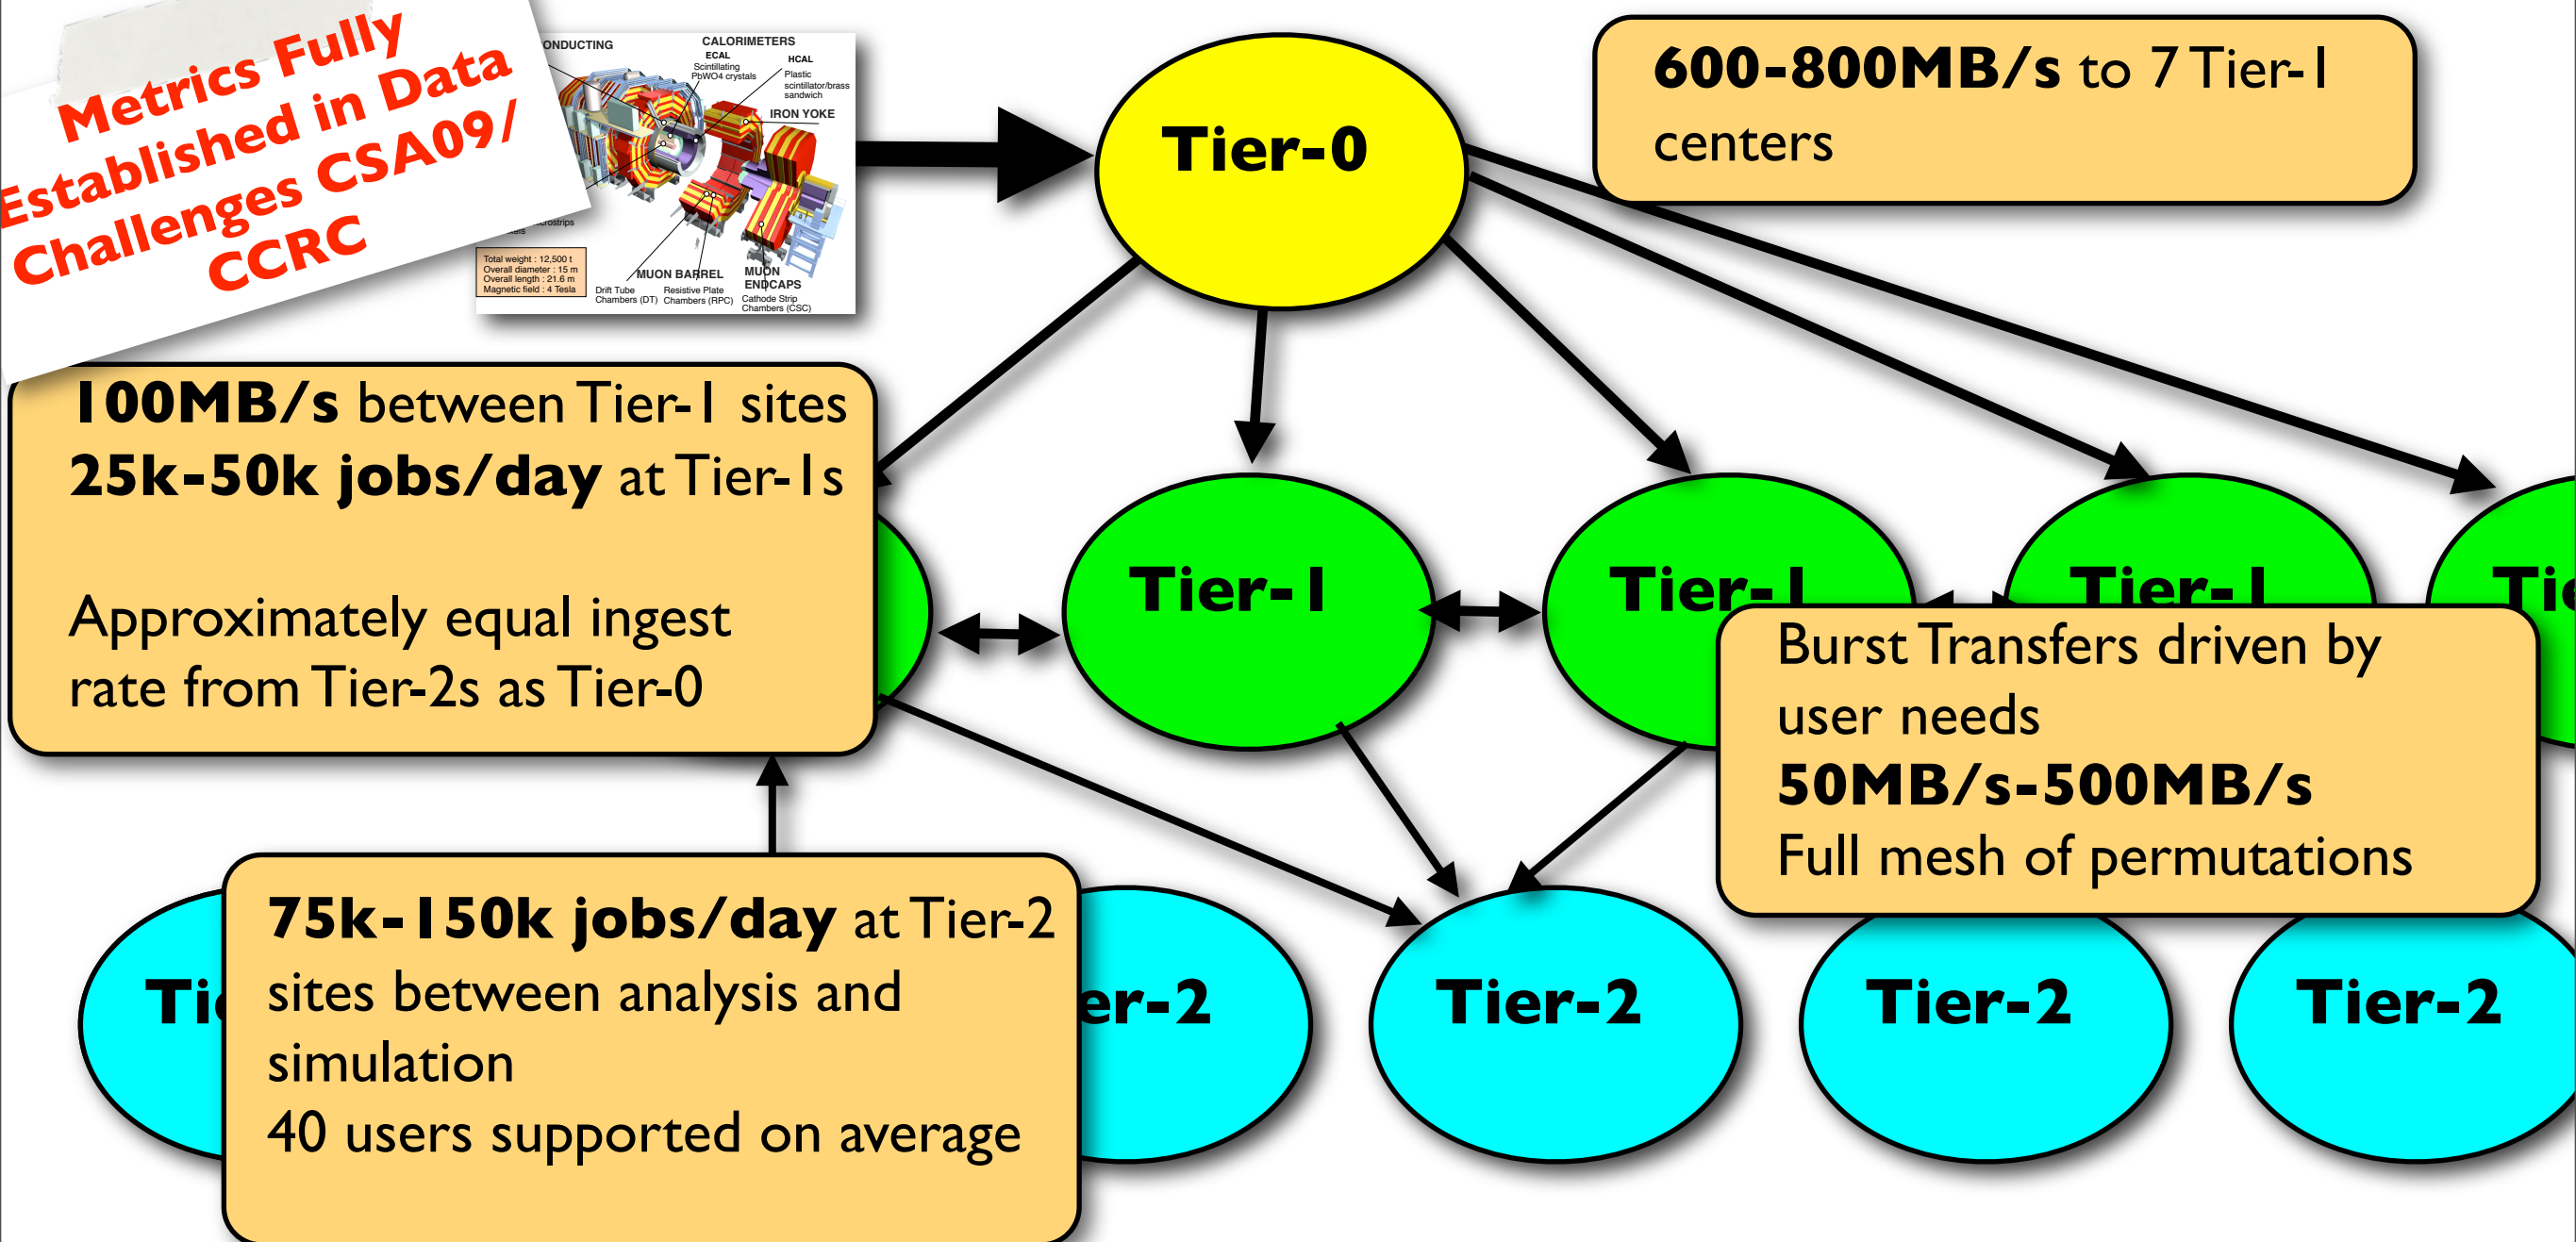



CMS Computing Metrics

◆ Reminder of CMS Computing Model:

- ★ user submit jobs, workload management system submits to where data is
- ★ data movements are triggered by operators, physics organizers, users

**Metrics Fully
Established in Data
Challenges CSA09/
CCRC**





Getting Data to Tier-2s for Physics Analysis

Transfers Fermilab → Tier-2 sites As tested in CCRC08

◆ Goals reached

★ Tier-1 → Tier-2

50-500MB/s in bursts

◆ Grid transfers work in general

★ Still working to get every individual link up to speed

★ great progress due to site and link commissioning work over this year

◆ Broad utilization of analysis tools by user


★ Nearly 800 individuals submitted CRAB jobs in 2008

★ Have reached nearly 100k jobs by 80-100 individual users per day

◆ reaching 50k jobs for individuals as demonstrated in computing challenge

★ Still working to improve the overall user experience on grid

◆ but doing ok according to user feedback





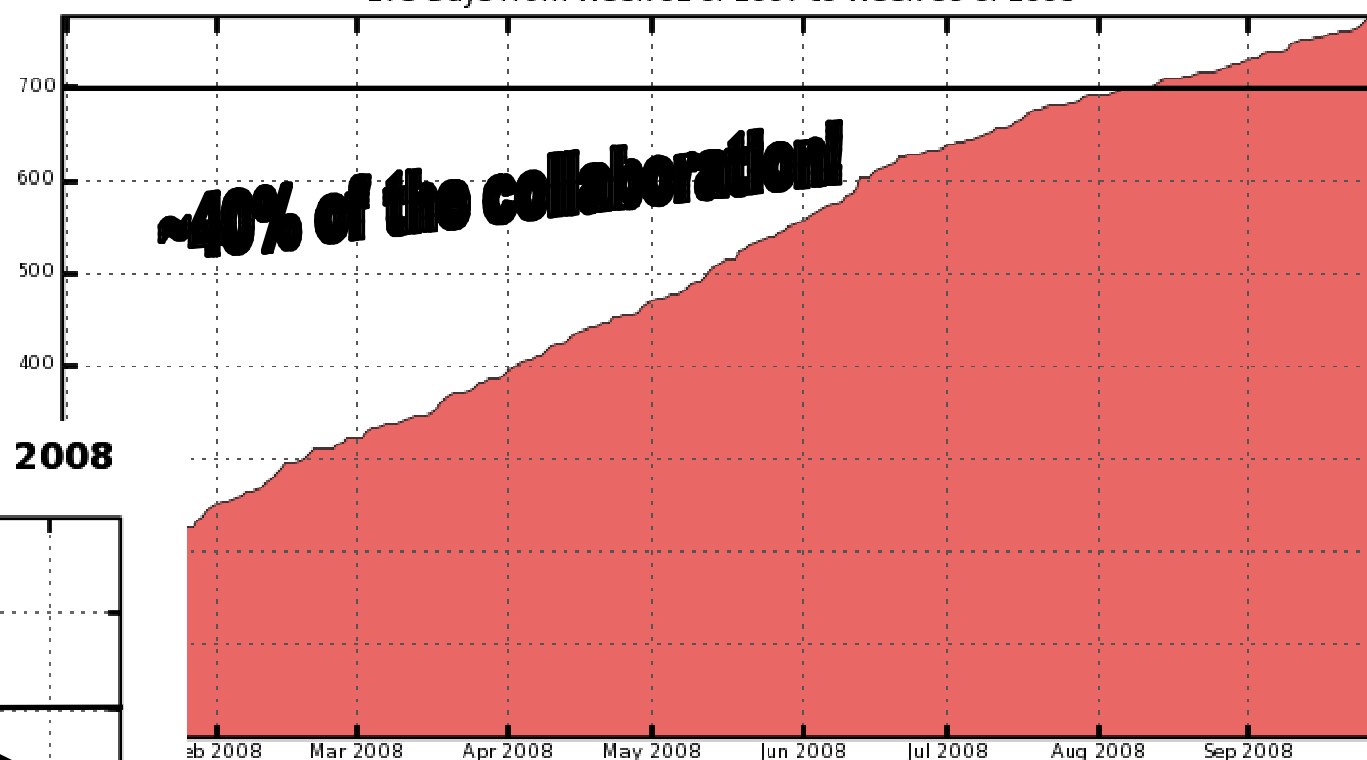
Distributed Data Analysis

- ★ CMS Tool CRAB, with the glide-in WMS
- ◆ Good adoption by the collaboration: 40% of collaboration using it!

Crab distinct users from the beginning of 2008


273 Days from Week 52 of 2007 to Week 39 of 2008

700 people



Analysis jobs terminated per day from the beginning of 2008


259 Days from Week 52 of 2007 to Week 37 of 2008



60k jobs



CMS Largest User VO within OSG




CMS did a massive production of simulated events and a reconstruction processing of cosmics data over the months of Nov 2008 - January 2009



2009 Computing Challenges

- ◆ Start of Data Taking in October, and cosmic data taking in ~July
 - ★ then a looong run for ~44 weeks
- ◆ in Spring 2009: end-to-end exercise to verify that luminosity, conditions and data quality information are consistently handled
 - ★ putting it all together, running through the standard workflows at all tiers



- ◆ Moving towards Analysis Operations
 - ★ task force started to assess effort and tools required —> needs to OSG
 - ★ transition to analysis operations in spring —> load on OSG



Specific Needs: Operations of Services

- ★ Security context. This includes operational, policy, software, validation.
- ★ BDII at the GOC.
- ★ GOC ticketing system.
- ★ Accounting and RSV, and related reporting to WLCG
- ★ Software cache and software lifecycle management via the VDT, including the OSG cache, as well as deployment testing, incl. vtb and itb.
- ★ Variety of validation, including VOMS and GIP.
- ★ Registration and administrative services
- ★ Support, including OSG-storage (office hour, ticketing), and CE deployment.

CMS identified the BDII as special among these because it is a single point of failure which can make the entirety of OSG disappear from view.

Requested scalability testing of BDII, resulting in a clear understanding of operational parameters, now and in the future. (ongoing)



Specific Needs: Accounting and Reporting

- ◆ Accounting, reliability, and availability reporting
 - ★ OSG is responsible for aggregating the reporting for CMS, and present CMS towards WLCG for problem resolution in WLCG accounting.
 - ★ OSG prepare a monthly report for the funding agencies to offset any misrepresentations by WLCG.
 - ◆ OSG to engage in WLCG, and be pro-active about changes in WLCG requirements.
 - ★ E.g. transition to SI2006, space accounting, deployed capacity accounting.
- ★ *All goals are being met by OSG.*



Tier-2 Reliability and Accounting

★ Availability and Reliability is being published since the end of the summer

Federation	Site	CPU Count	Reli-ability	Avail-ability	Reliability History		
					Oct-08	Nov-08	Dec-08
T2_US_Caltech (USA, Caltech CMS T2)							
	cit_cms_t2	N/A	97 %	97 %	68 %	85 %	91 %
T2_US_Florida (USA, Florida CMS T2)							
	uflorida-hpc	N/A	96 %	96 %	94 %	92 %	100 %
	uflorida-ihepa	N/A	100 %	100 %	99 %	100 %	100 %
	uflorida-pg	N/A	100 %	100 %	99 %	98 %	100 %
T2_US_MIT (USA, MIT CMS T2)							
	mit_cms	N/A	77 %	77 %	87 %	93 %	90 %
T2_US_Nebraska (USA, Nebraska CMS T2)							
	nebraska	N/A	99 %	99 %	92 %	97 %	92 %
T2_US_Purdue (USA, Purdue CMS T2)							
	purdue-lear	N/A	N/A	N/A	N/A	N/A	N/A
	purdue-rcac	N/A	97 %	97 %	98 %	93 %	95 %
	purdue-steele	N/A	99 %	99 %	86 %	96 %	97 %
T2_US_UCSD (USA, UC San Diego CMS T2)							
	ucsdt2	N/A	88 %	88 %	100 %	98 %	78 %
T2_US_Wisconsin (USA, U. Wisconsin CMS T2)							
	GLOW	N/A	96 %	89 %	100 %	100 %	100 %



Specific Needs: Site Operations Improvements

- ◆ Ability to announce downtimes of both CE & SE.
- ◆ Ability to ban single FQAN at a site.
- ◆ Integrate all existing monitoring functionalities of CE & SE into one site-wide “monitoring portal”.
- ◆ Ability to have multiple CEs per cluster viewed as one as in EGEE.
- ◆ Ability to reject new jobs while in drain-off.

Adequate progress made, but none of these is completed.

(completion = deployed and in operations at CMS sites)

Additional operational goals:

Both are met.

Reinstallation of CE in 4-8 hours

Reinstallation of SE in 6-12 hours



WMS Issues

- ◆ OSG support during Production roll-out of glideinWMS. (ongoing)
- ◆ Help from “scalability, reliability, usability” area in OSG with condor testing. (done)
- ◆ glideinWMS not to submit glideins to sites that are off in BDII. (incomplete)



Storage Issues

- ◆ A lot of progress, some open issues, and: do we need a road-map?
 - ★ LCG client tools (done)
 - ★ Achieve 2Hz SRM request rate (done)
 - ◆ CMS realized change in needs: 20-100Hz
 - ★ Additional ops tools for dCache ops.
 - ◆ Locally configurable space usage monitor (scheduled for next release)
 - ★ Procedures for operations problems to feed back into development cycle. (scheduled for end of year 3)
 - ★ Successfully deploy BestMan at 5-10 OSG sites for production use with PhEDEx. (partially done)
 - ★ Mutually compatible gLiteUI <-> OSG client.
 - ◆ (partially done via lcg client tools deployment as part of OSG client)
- ◆ future storage systems for Tier-2, Tier-3?
 - ★ need to chart out a way forward
 - ★ require more manageability, reliability, usability
 - ★ the Nebraska Hadoop “coup” shows what is feasible



Summary/Conclusions

- ◆ After several stress tests and data challenges, numerous large scale simulation exercises, a (all too) short experience with beam-induced data, and extensive cosmic runs, the S&C systems and operations teams are ready for sustained data taking, processing and analysis
 - ★ Work remains in the areas of usability, automation, reliability
- ◆ Software and Computing Systems useable for a large # physicists
 - ★ Software is functional, remaining issues are being addressed
 - ★ We have significant computing resources deployed
 - ◆ that are being used for detector commissioning and physics studies
 - ★ CMS OSG Tier-2 and Tier-3 sites are working well and are being used
 - ◆ University sites started to plug into the system
 - ◆ —> system tests show that the system works for many users
- ◆ We are ready to operate software and computing 24/7
 - ★ Still, sizable development, integration, commissioning ahead
- ◆ OSG has been fully engaged and is major contributor to CMS