

Security in OSG

Rob Quick <rquick@iu.edu>
OSG Operations Coordinator
Manager High Throughput Computing Group
Indiana University

Slides By Igor Sfiligoi

Logistical reminder

It is OK to ask questions

During the lecture

During the demos


During the exercises

During the breaks

If I don't know the answer,
I will find someone who likely does

Reminder – Single sign-on

The user should use the same
mechanism to submit jobs to any site
And there are 100s of them in OSG



Passwords a non-starter

We all know username/password is the preferred authentication mechanism

Almost everybody use it!

But not a good solution for distributed systems

Uses a **shared secret** between
the user and the service provider

And secrets stay secret only if few entities know it

Sharing passwords between sites a bad idea!

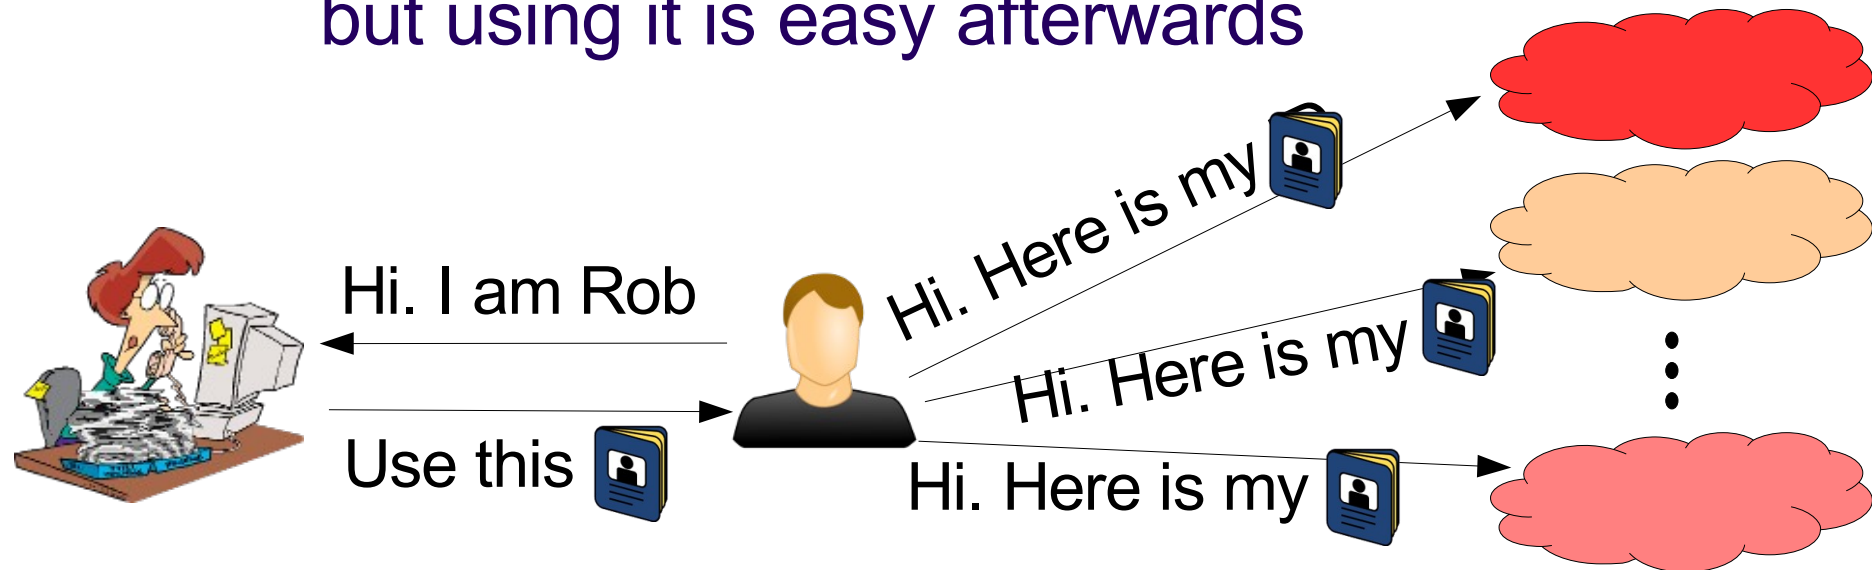
Adding an intermediary

A better approach is to introduce a highly trusted intermediary

Have been used in real life for ages

e.g. States as issuers of IDs/Passports

Getting the ID can be a lengthy process,
but using it is easy afterwards



Adding an intermediary

A better approach is to introduce a highly trusted intermediary

Have been used in the past for e.g.

e.g. State

Getting

but

Chain of trust.

You are trusted because
the site trusts the issuer,
and the issuer trusted you.



Hi. I am

Use this



Hi. Here is my



...

Technical implementations

Many technical solutions

x.509 PKI

Kerberos

OpenID

many more...

All based on the same basic principle

Each has strengths and weaknesses

OSG standardized on x.509



Will not argue
if it is the best
one.

x.509 PKI

Based on public key cryptography

A user has a (private,public) key pair

One signs, the other verifies

The highly trusted entity is called a
Certification Authority (CA)


The user is given a **certificate**

Cert. has user name in it

Cert. also contains the (priv, pub) key pair

Cert. has a limited lifetime

Cert. is signed by the CA private key

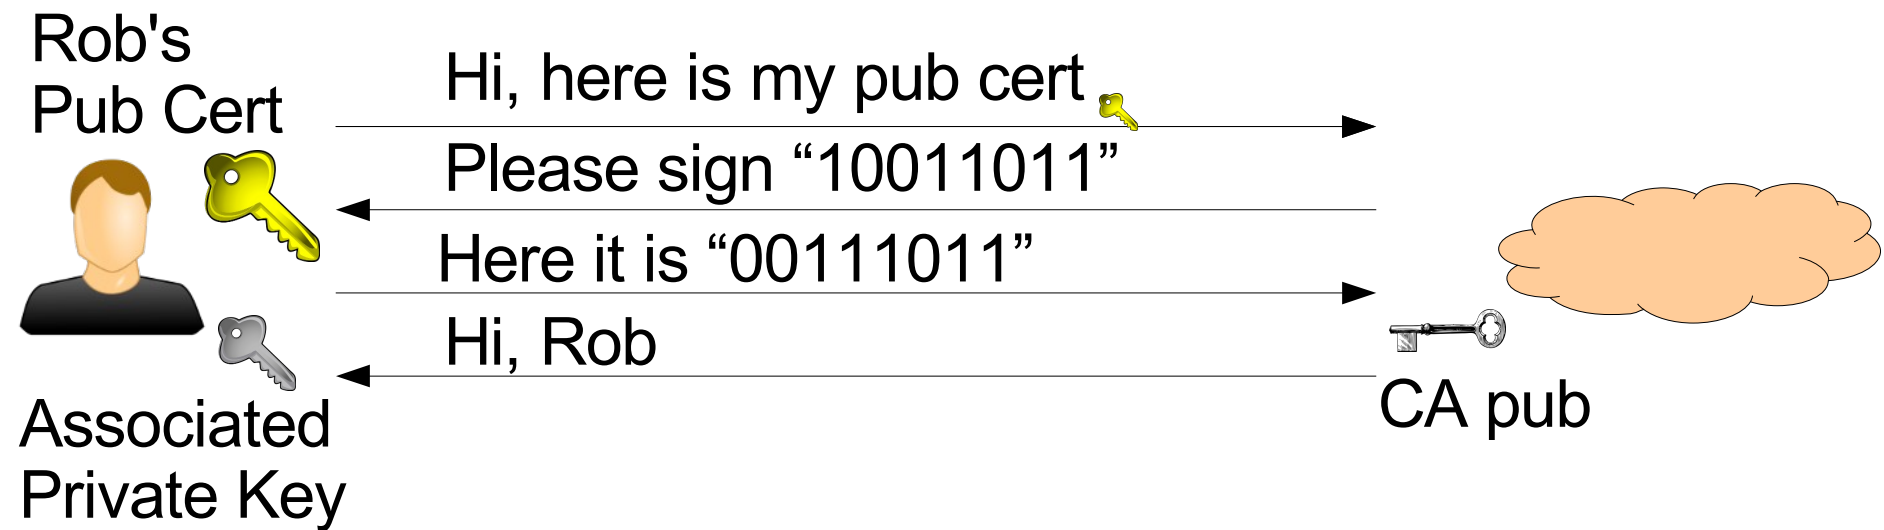


You
should
have one
by now.

x.509 authentication

Sites have CA public key pre-installed

User authenticates by signing a site provided string and providing the public part of the cert



Mutual authentication

The OSG clients also require servers to authenticate

Same principle as before

The site's server owns a x.509 certificate

User client must have the CA pre-installed

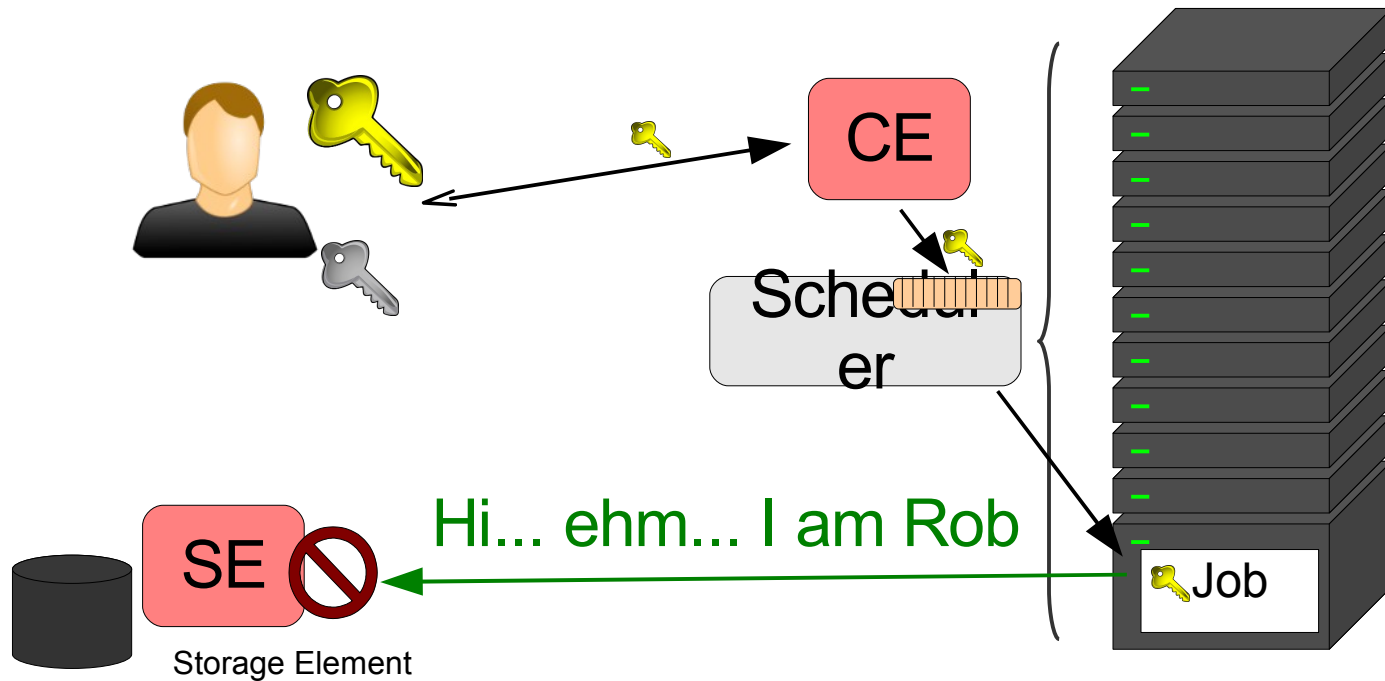
So we have mutual authentication



Impersonation

Sometimes your jobs need to impersonate you

For example to access remote data



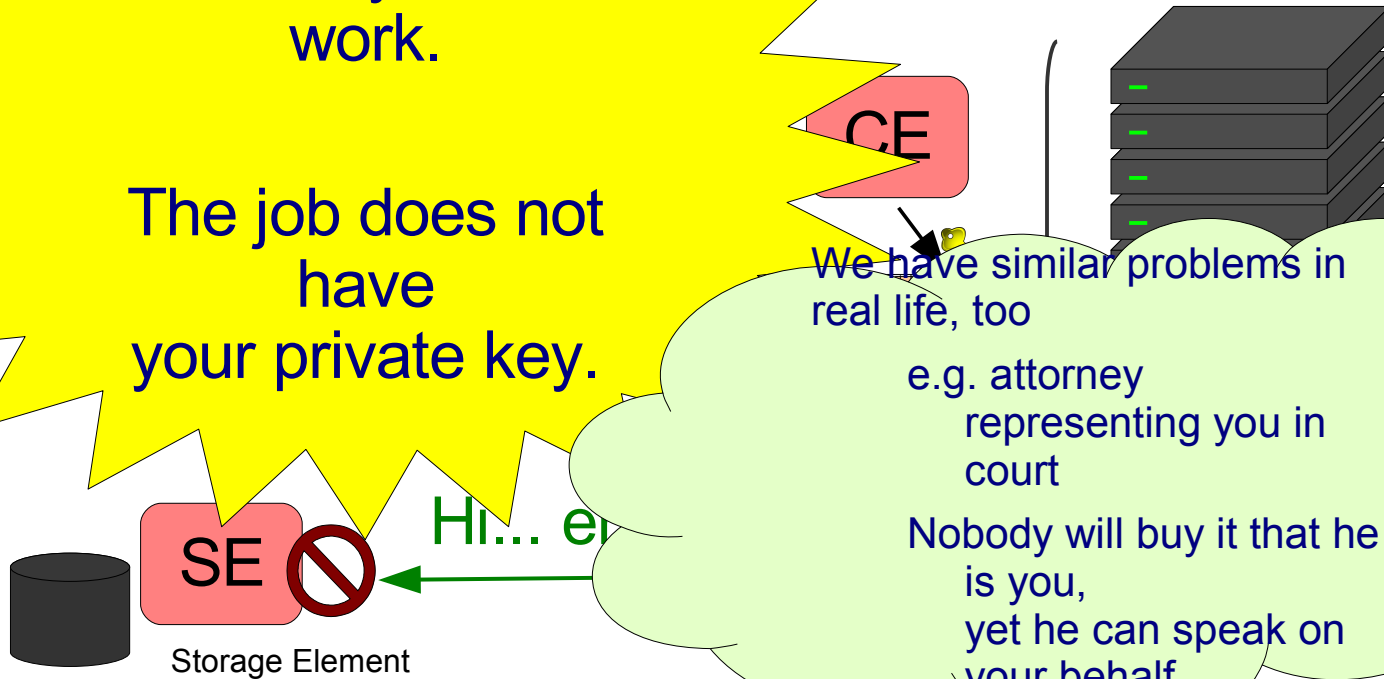


Impersonation

Sometimes your jobs need to impersonate you
For example, to access remote data

Obviously will not work.

The job does not have your private key.



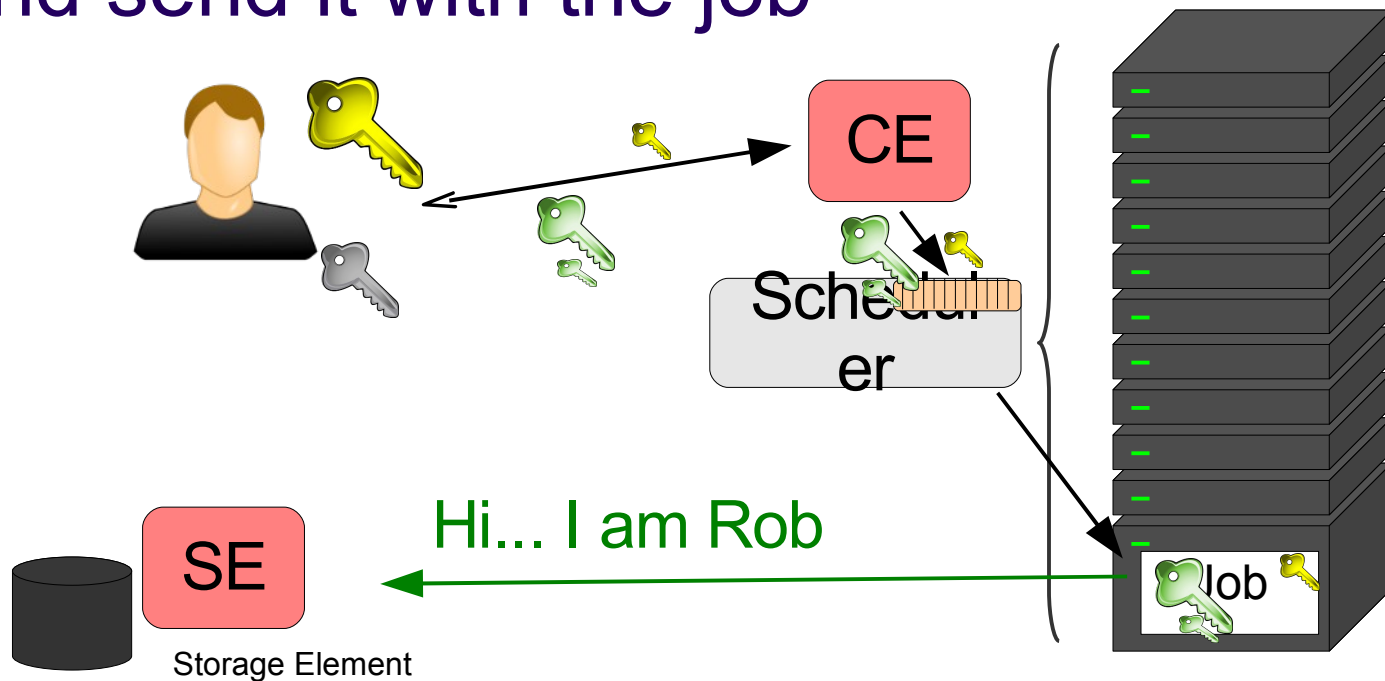
Proxy delegation

The **job** is indeed **not you**

Create a proxy certificate for the job

Add another level of trust delegation

And send it with the job





Proxy delegation

The **job** is indeed **not you**

Create a proxy certificate for the job

Add another delegation

And send

YES!

You are sending
the
proxy's private
key
to the WN.



Risk mitigation

Proxy delegation is risky

Your proxy could be stolen

In OSG, we mitigate by limiting lifetime

At most few hours recommended

After the proxy expires, the proxy is useless

Can be annoying

Must keep renewing, if long running job!



But we don't have
anything better.



Risk mitigation

Proxy delegation is risky

If using glideinWMS, Condor will automatically create a short lived proxy and keep re-delegating it, for long running job!

But we do anything

Completely transparent to you.

x.509 in Overlay systems

x.509 is typically used in Overlay systems as well

For glideinWMS,
all communication between processes
is mutually authenticated using
x.509 (proxy) certificates


Authentication vs. Authorization

Just because you can authenticate yourself, it does not mean you are authorized, too

e.g. your passport tells who you are,
but does not allow you to drive a car

x.509 PKI only covers **authentication**

Tells the site who you are



Need a different
mechanism
for authorization

Per-user authorization not an option


The naive approach is using a list

Since we do not want let just anyone in!

However, the problem is scale

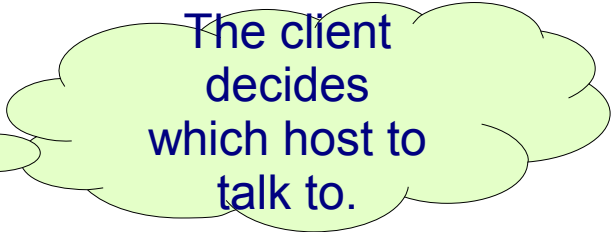
OSG has ~10,000 users!

Sites do not want to decide on
a user-by-user basis!



Server authorization is
easy.

Just require host name
in the certificate name;
CA will enforce this.



The client
decides
which host to
talk to.

Adding roles

Sites want to operate on higher level concepts

Some kind of attribute

Like in real life

Think about passport vs driver's license

Both tell a cop who you are

(and to 1st approx. are issued by the same entity)

But the driver's license tells him

you are allowed to use a car, too

“Class C”

Need for an attribute authority

Users can have many roles

But don't want to have multiple certs

e.g. I may be running HEP jobs or School jobs

So the attributes cannot come from the CA

And you would not just trust the user

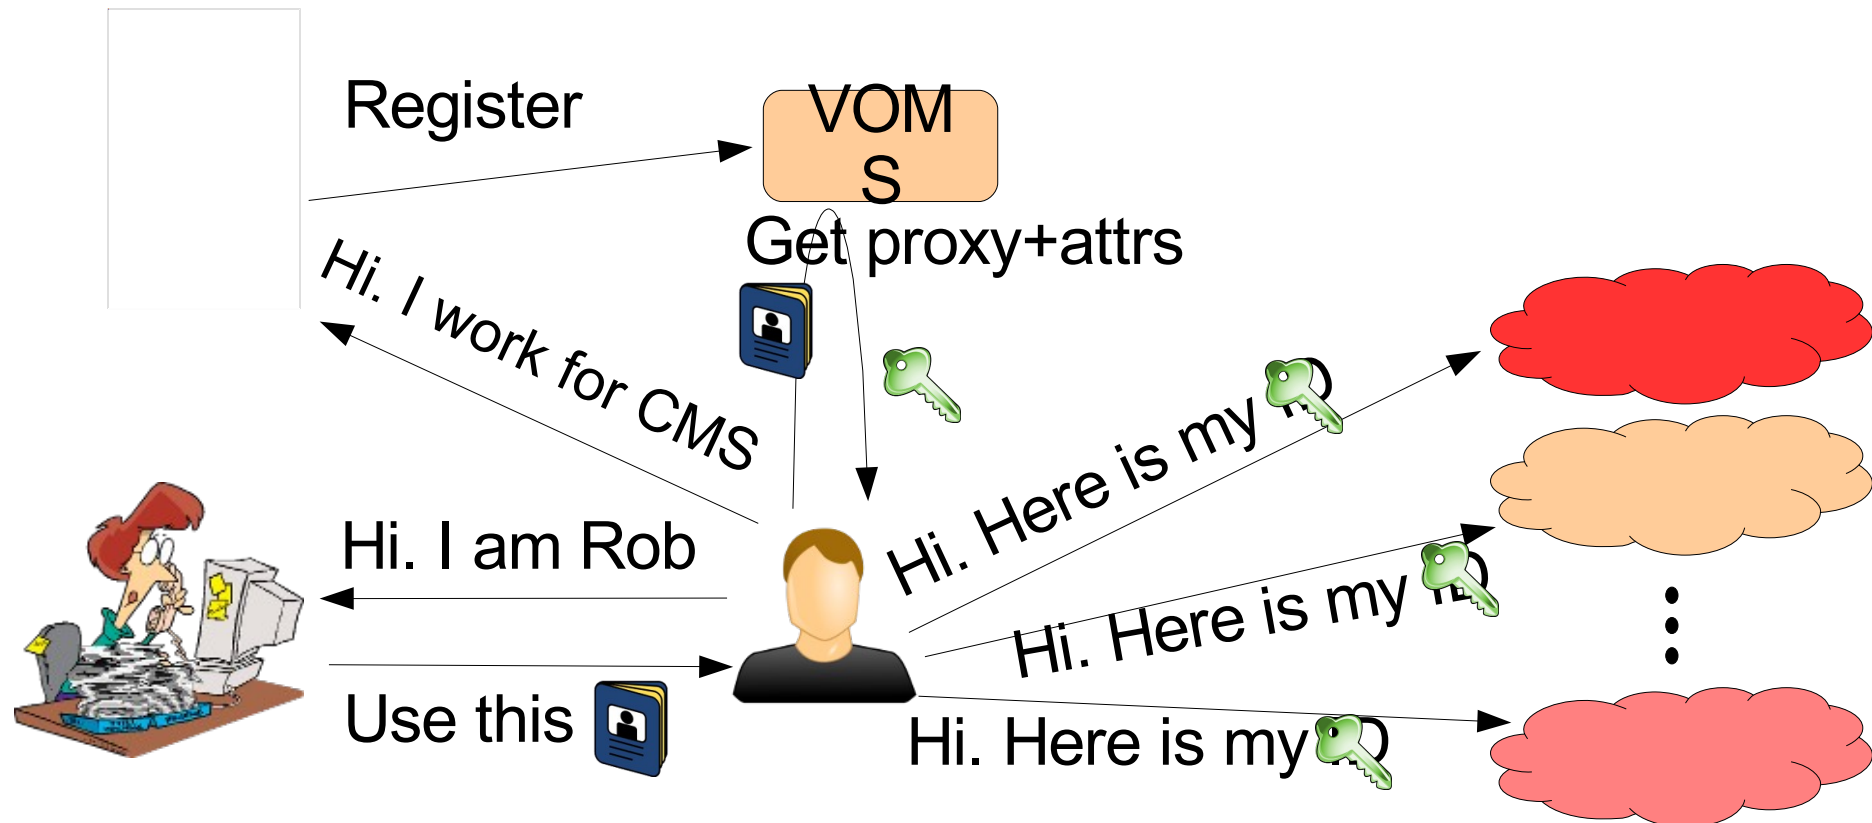
In OSG, we use VOMS

Virtual Organization Management System

OSG expects well organized VOs (e.g. CMS)

VO and VOMS

VO decides who is worthy of an attribute
Site decides based on that attribute



VO and VOMS

VO decides who is worthy of an attribute

Site decides based on that attribute

Register


Hi

Site still knows who
you are
and can act on your
name alone.

Hi. I a

But they will do it only
for exceptional cases.

Use this



More security considerations

There is much more than authentication and authorization to security

But we don't have the time to cover everything

Just briefly

Sharing of resources

Privacy

Acceptable conduct

Sharing of resources

Modern CPUs are many-core, so

Very likely your job will be
sharing the node with other jobs

Sites will map your Grid name into UID

Hopefully unique... be sure to ask

Standard *NIX protections

Act accordingly

e.g. no file should be world writable

Privacy

By default, no privacy in OSG

Assume all your files are publicly readable

Apart from your proxy

If you need privacy, you will have to take explicit measures

Both during network transfers, and

For files on disk

x.509 can be used for encryption

But remember, proxy has new keys

Acceptable conduct

Each OSG user is bound by its AUP
(Acceptable User Policy)

And sites are allowed to have additional
rules in place

In a nutshell

Use only for the declared science purpose

Do not overload the system

Do not attempt to circumvent security



Else,
you will be
banned!



Questions?

Questions? Comments?

Feel free to ask me questions later:

Rob Quick <rquick@iu.edu>