



An Introduction to High-Throughput Computing

Monday morning, 9:15am

Alain Roy <roy@cs.wisc.edu>
OSG Software Coordinator
University of Wisconsin-Madison


Who Am I?

- With Condor since 2001
- Open Science Grid Software Coordinator
- Taught at five previous summer schools
- I have the two cutest kids in the whole world:



Overview of day

- Lectures alternating with exercises
 - Emphasis on lots of exercises
 - Hopefully overcome PowerPoint fatigue



"I need someone well versed in the art of torture—do you know PowerPoint?"

Overview of day

- I've done most of these exercises in a grid school before, but not with this format:
 - It's okay to move ahead on exercises if you have time
 - It's okay to take longer on them if you need to
 - If you move along quickly, try the "On Your Own" sections and "Challenges"

Most important!

- Please ask me questions!
 - During the lectures
 - During the exercises
 - During the breaks
 - During the meals
 - Over dinner
 - The rest of the week

Before we start

- Sometime today, do the exercise on getting a certificate.
 - It is **required** for all exercises Tuesday-Thursday.
 - It will be easiest if you do it today.

Goals for this session

- Understand basics of high-throughput computing
- Understand the basics of Condor
- Run a basic Condor job





What is high-throughput computing? (HTC)

- An approach to distributed computing that focuses on long-term throughput, not instantaneous computing power.
 - We don't care about operations per second
 - We care about operations per year
- Implications:
 - Focus on reliability
 - Use all available resources (not just high-end supercomputers)

Good uses of HTC

- “I need as many simulation results as possible before my deadline...”
- “I have lots of small-ish independent tasks that can be run indepdently”

What's not HTC?

- The need for “real-time” results:
 - (Ignore fact “real-time” is hard to define.)
 - HTC is less worried about latency (delay to answer) than total throughput
- The need to maximize FLOPS
 - “I must use the latest supercomputer, because I need the fastest computer/network/storage/...”

An example problem: BLAST

- A scientist has:
 - Question: Does a protein sequence occur in other organisms?
 - Data: lots of protein sequences from various organisms
 - Parameters: how to search the database.
- More throughput means
 - More protein sequences queried
 - Larger/more protein data bases examined
 - More parameter variation
- We'll try out BLAST later today

Why is HTC hard?

- The HTC system has to keep track of:
 - Individual tasks (a.k.a. jobs) & their inputs
 - Computers that are available
- The system has to recover from failures
 - There will be failures! Distributed computers means more chances for failures.
- You have to share computers
 - Sharing can be within an organization, or between orgs
 - So you have to worry about security.
 - And you have to worry about policies on how you share.
- If you use a lot of computers, you have to deal variety:
 - Different kinds of computers (arch, OS, speed, etc..)
 - Different kinds of storage (access methodology, size, speed, etc...)
 - Different networks interacting (network problems are hard to debug!)

Let's take one step at a time

Small

Local



Large

Distributed

- Can you run one job on one computer?
- Can you run one job on another local computer?
- Can you run 10 jobs on a set of local computers?
- Can you run 1 job on a remote computer?
- Can you run 10 jobs at a remote site?
- Can you run a mix of jobs here and remotely?
- This is the (rough) progress we'll take in the school.

Discussion

- For 5 minutes, talk to a neighbor: If you want to run one job in a local cluster of computers:
 - What do you (the user) need to provide so a single job can be run
 - What does the system need to provide so your single job can be run?
 - Think of this as a set of processes: what needs happen when the job is given? A “process” could be a computer process, or just an abstract task.



Alain's answer:

What does the user provide?

- A “headless job”.
 - Not interactive (no GUI): how do you interact with 1000 simultaneous jobs?
- A set of input files
- A set of output files.
- A set of parameters (command-line arguments).
- Requirements:
 - Ex: My job requires at least 2GB of RAM.
 - Ex: My job requires Linux.
- Control:
 - Ex: Send me email when the job is done.
 - Ex: Job 2 is more important than Job 1.
 - Ex: Kill my job if it's run for more than 6 hours.

Alain's answer:

What does the system provide?

- Methods to:
 - Submit/Cancel job.
 - Check on state of job.
 - Check on state of avail. computers.
- Processes to:
 - Reliably track set of submitted jobs.
 - Reliably track set of available computers.
 - Decide which job runs on where.
 - Advertise a single computer.
 - Start up a single job.

Surprise!

Condor does this (and more)

- Methods to:
 - Submit/Cancel job. `condor_submit/condor_rm`
 - Check on state of job. `condor_q`
 - Check on state of avail. computers. `condor_status`
- Processes to:
 - Reliably track set of submitted jobs. `schedd`
 - Reliably track set of avail. computers. `collector`
 - Decide which job runs on where. `negotiator`
 - Advertise a single computer `startd`
 - Start up a single job `starter`

A brief introduction to Condor



And Google Takes Computers...

I need a Mac!

$$E = mc^2$$


ers

Desktop Computers

$$\begin{aligned} &= 1\text{kg} \times (3 \times 10^8 \text{ ms}^{-1})^2 \\ &= 1\text{kg} \times (3 \times 10^8 \text{ ms}^{-1}) \times (3 \times 10^8 \text{ ms}^{-1}) \\ &= 1\text{kg} \times (9 \times 10^{16} \text{ m}^2 \text{ s}^{-2}) \\ &= 1 \times (9 \times 10^{16}) \text{ kg m}^2 \text{ s}^{-2} \\ &= 9 \times 10^{16} \text{ J} \end{aligned}$$

Match

need a Linux box
with 2GB RAM!



Quick Terminology

- **Cluster**: A dedicated set of computers not for interactive use
- **Pool**: A collection of computers used by Condor
 - May be dedicated
 - May be interactive

Matchmaking

- Matchmaking is fundamental to Condor
- Matchmaking is two-way
 - Job describes what it requires:
I need Linux && 8 GB of RAM
 - Machine describes what it requires:
I will only run jobs from the Physics department
- Matchmaking allows preferences
 - I **need** Linux, and I **prefer** machines with more memory but will run on any machine you provide me

Why Two-way Matching?

- Condor conceptually divides people into three groups:
 - Job submitters
 - Machine owners
 - Pool (cluster) administrator
- } May or may not be the same people
- All three of these groups have preferences

ClassAds

- ClassAds state facts
 - My job's executable is analysis.exe
 - My machine's load average is 5.6
- ClassAds state preferences
 - I require a computer with Linux



ClassAds

- ClassAds are:
 - semi-structured
 - user-extensible
 - schema-free
 - Attribute = Expression

Example:


```
MyType           = "Job" ← String
TargetType       = "Machine"
ClusterId        = 1377 ← Number
Owner            = "roy"
Cmd              = "analysis.exe"
Requirements     =
    (Arch == "INTEL") ← Boolean
    && (OpSys == "LINUX")
    && (Disk >= DiskUsage)
    && ((Memory * 1024) >= ImageSize)
...
```


Schema-free ClassAds

- Condor imposes some schema
 - Owner is a string, ClusterID is a number...
- But users can extend it however they like, for jobs or machines
 - `AnalysisJobType = "simulation"`
 - `HasJava_1_4 = TRUE`
 - `ShoeLength = 7`
- Matchmaking can use these attributes
 - `Requirements = OpSys == "LINUX"`
`&& HasJava_1_4 == TRUE`

Submitting jobs

- Users submit jobs from a computer
 - Jobs described as ClassAds
 - Each submission computer has a queue
 - Queues are **not** centralized
 - Submission computer watches over queue
 - Can have multiple submission computers
 - Submission handled by condor_schedd





Open Science Grid

Advertising computers

- Machine owners describe computers
 - Configuration file extends ClassAd
 - ClassAd has dynamic features
 - Load Average
 - Free Memory
 - ...
 - ClassAds are sent to Matchmaker



ClassAd

Type = "Machine"

Requirements = "..."

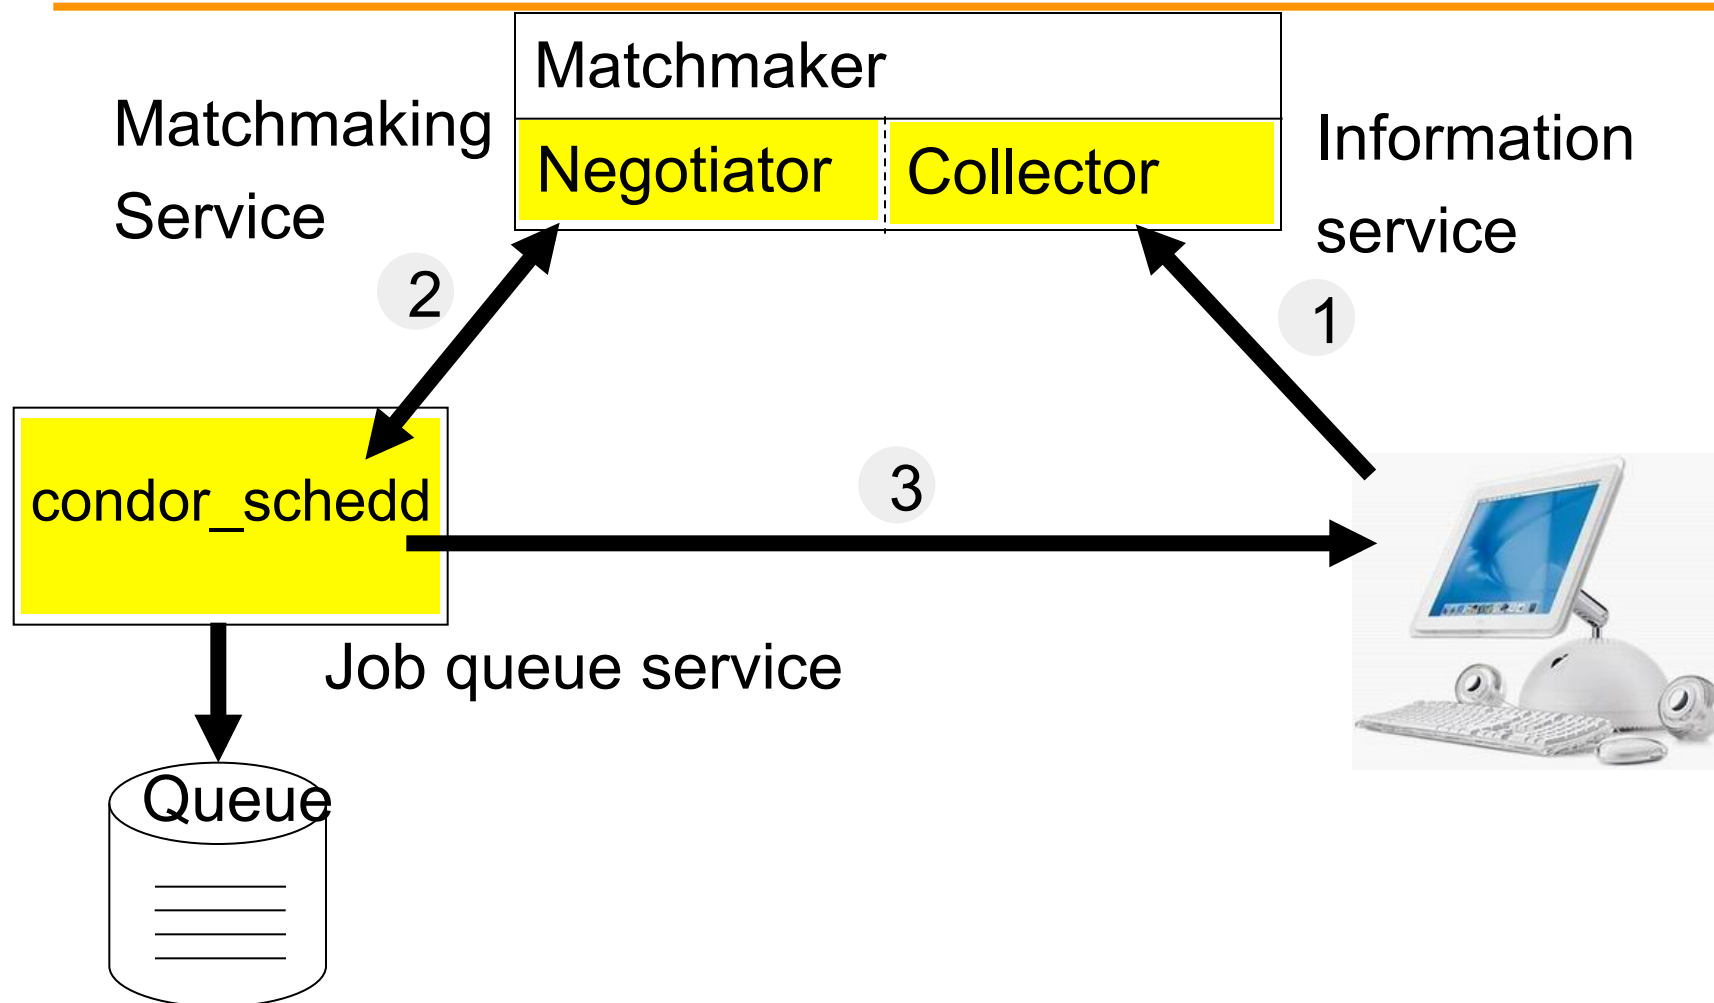
Matchmaker
(Collector)



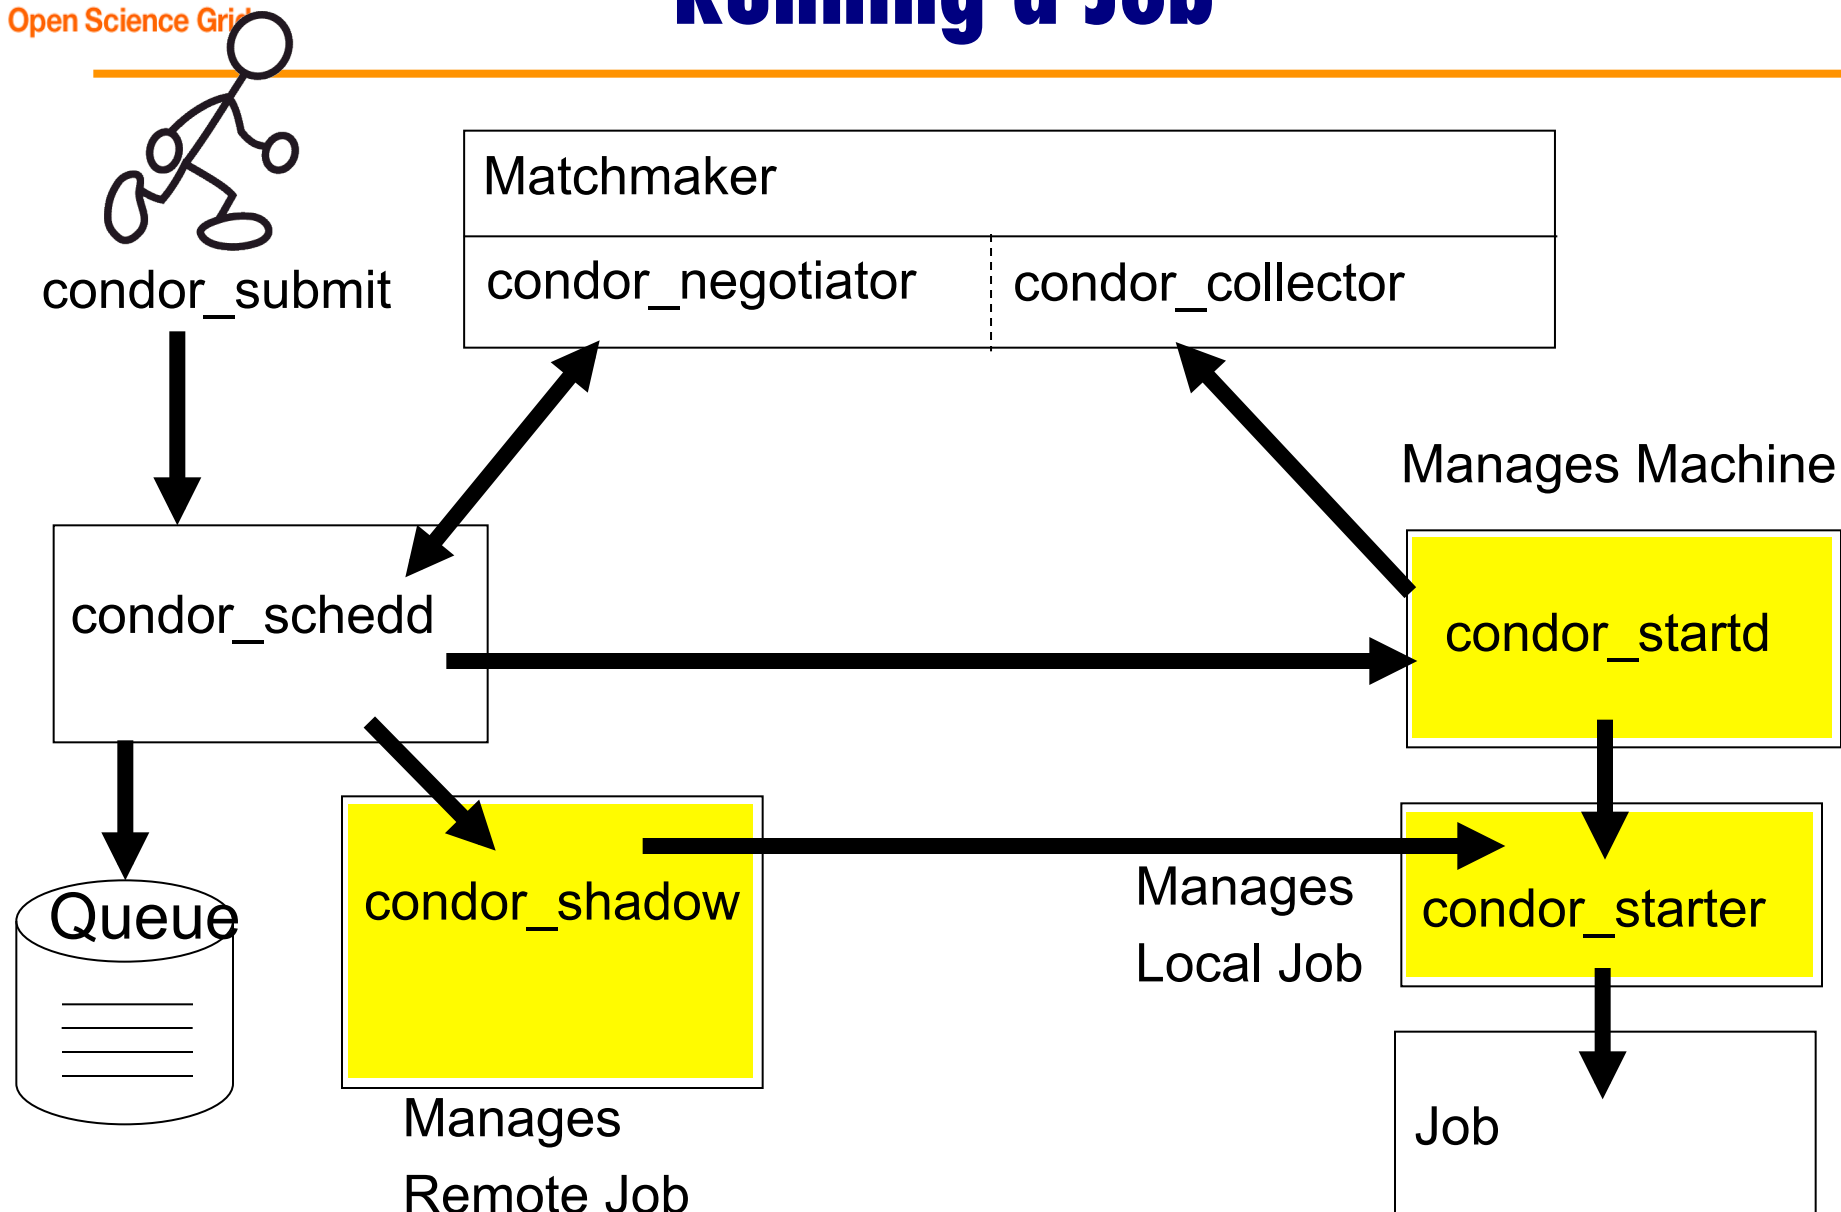
Matchmaking

- Negotiator collects list of computers
- Negotiator contacts each schedd
 - What jobs do you have to run?
- Negotiator compares each job to each computer
 - Evaluate requirements of job & machine
 - Evaluate in context of both ClassAds
 - If both evaluate to true, there is a match
- Upon match, schedd contacts execution computer

Matchmaking diagram



Running a Job



Condor processes

- Master: Takes care of other processes
- Collector: Stores ClassAds
- Negotiator: Performs matchmaking
- Schedd: Manages job queue
- Shadow: Manages job (submit side)
- Startd: Manages computer
- Starter: Manages job (execution side)

If you forget most of these remember two (for other lectures)

- Master: Takes care of other processes
- Collector: Stores ClassAds
- Negotiator: Performs matchmaking

Schedd: Manages job queue

- Shadow: Manages job (submit side)


Startd: Manages computer

- Starter: Manages job (execution side)


Some notes

- One negotiator/collector per pool
- Can have many schedds (submitters)
- Can have many startds (computers)
- A machine can have any combination of:
 - Just a startd (typical for a dedicated cluster)
 - schedd + startd (perhaps a desktop)
 - Personal Condor: everything

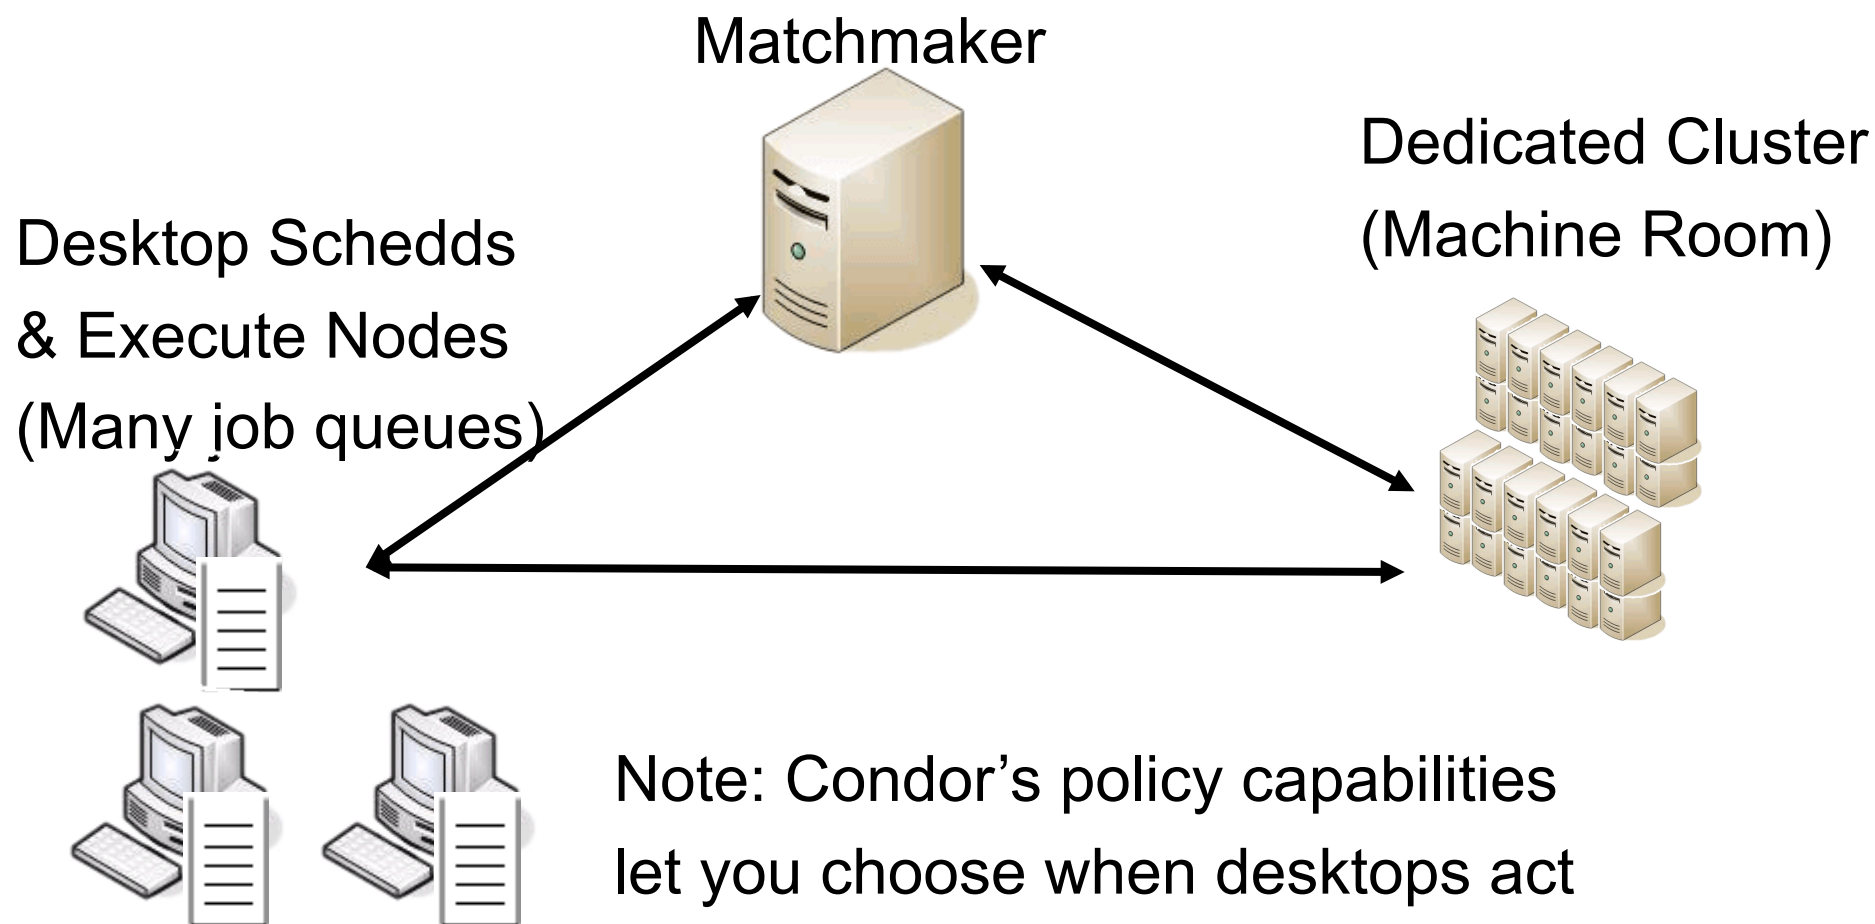
Example Pool 1



Example Pool 1a



Example Pool 2



Note: Condor's policy capabilities let you choose when desktops act as execute nodes.



Our Condor Pools

- One submit computer
 - One schedd/queue for everyone
 - `vdt-itb.cs.wisc.edu`
- One local set of dedicated Condor execute nodes:
 - About 14 computers, nearly 60 available “batch slots”
- Remote resources at UNL
 - For Tuesday, not today.



Open Science Grid

Our Condor Pool

Name	OpSys	Arch	State	Activity	LoadAv	Mem	ActvtyTime
slot1@miniosg-c01.	LINUX	INTEL	Unclaimed	Idle	0.000	2030	0+00:00:04
slot2@miniosg-c01.	LINUX	INTEL	Unclaimed	Idle	0.000	2030	0+00:00:05
slot3@miniosg-c01.	LINUX	INTEL	Unclaimed	Idle	0.000	2030	0+00:29:46
slot4@miniosg-c01.	LINUX	INTEL	Unclaimed	Idle	0.000	2030	0+00:29:47
slot5@miniosg-c01.	LINUX	INTEL	Unclaimed	Idle	0.020	2030	0+00:07:45
slot6@miniosg-c01.	LINUX	INTEL	Unclaimed	Idle	0.000	2030	13+01:43:11
slot7@miniosg-c01.	LINUX	INTEL	Unclaimed	Idle	0.000	2030	13+01:43:12
slot8@miniosg-c01.	LINUX	INTEL	Unclaimed	Idle	0.000	2030	13+01:43:05
slot1@miniosg-c02.	LINUX	INTEL	Unclaimed	Idle	0.000	2030	0+05:35:22
slot2@miniosg-c02.	LINUX	INTEL	Unclaimed	Idle	0.000	2030	0+05:45:29
slot3@miniosg-c02.	LINUX	INTEL	Unclaimed	Idle	0.000	2030	0+01:35:06
...							

Total Owner Claimed Unclaimed Matched Preempting Backfill

INTEL/LINUX	59	0	59	0	0	0
Total	59	0	59	0	0	0

That was a whirlwind tour!

- Let's get some hands-on experience with Condor, to solidify this knowledge.
- Goal: Check out our installation, run some basic jobs.



Questions?

- Questions? Comments?
 - Feel free to ask me questions later:
Alain Roy <roy@cs.wisc.edu>
- Upcoming sessions
 - 9:45-10:30
 - Hands-on exercises
 - 10:30 – 10:45
 - Break
 - 10:45 – 12:15
 - More!