
Grid Compute Resources and Job Management



Open Science Grid

New Mexico Grid School – April 9, 2009

Marco Mambelli – University of Chicago

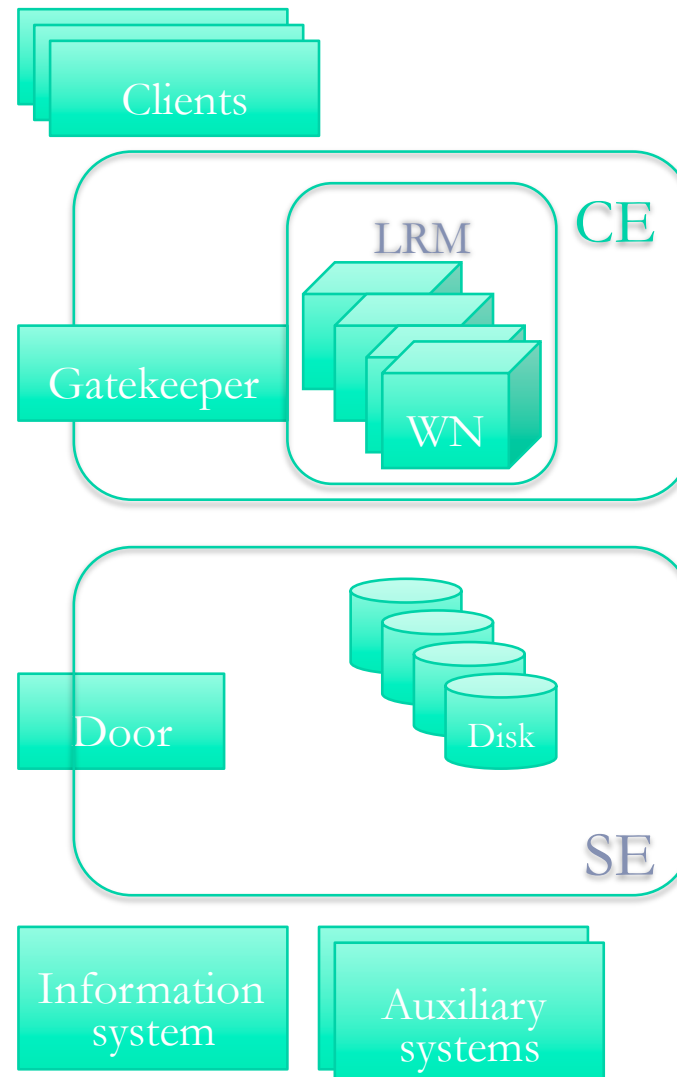
marco@hep.uchicago.edu

Outline

- Grid abstraction
- Clusters and Local Resource Managers
- GRAM, the Grid protocol
- Adding some Security
- Condor
- Something more with Condor
- DAGman
- OSG abstractions

Grid/OSG components


- Standard installation
- Set of services or resources containing data and providing processing power
- Computing elements (CE)
- Storage elements (SE)
- Information systems
- Clients
- How to interact?
 - validation and testing
 - active use



Grid as Abstraction

- Hides lower level components/fabrics
 - you can use it without knowing what is behind
- Provides leverage
 - do so much with so little
- Limit flexibility
- Suggest usage patterns
- Gain/loss

Do not reinvent the wheel!



Job management layers

- High level user tools
(Panda, Swift, Pegasus, ...)
- OSG abstractions
- DAGman
- Condor
- Adding some Security
- GRAM, the Grid protocol
- Clusters and Local
Resource Managers



The power of leverage



Good use of the Grid Abstraction

- Leaks are likely
 - ❑ all abstraction “lie”
 - ❑ some leaks are desired
- Before you can abstract you must see the details
- Understanding the levels below you understand the abstraction
- To build components
 - ❑ must be fully aware of at least a couple of layers below
 - ❑ to EXCELL you must be VERY familiar with several layers above

Job management layers

- High level user tools
(Panda, Swift, Pegasus, ...)
- OSG abstractions
- DAGman
- Condor
- Adding some Security
- GRAM, the Grid protocol
- Clusters and Local
Resource Managers



Local Resource Managers (LRM)

- Compute resources have a **local resource manager** (LRM) that controls:
 - ❑ Who is allowed to run jobs
 - ❑ How jobs run on a specific resource
- *Example policy:*
 - ❑ Each cluster node can run one job.
 - ❑ If there are more jobs, then they must wait in a queue
- LRMs allow nodes in a cluster can be **reserved** for a specific person
- *Examples:* PBS, LSF, Condor

Job management layers

- High level user tools
(Panda, Swift, Pegasus, ...)
- OSG abstractions
- DAGman
- Condor
- Adding some Security
- GRAM, the Grid protocol
- Clusters and Local
Resource Managers




GRAM

Globus Resource Allocation Manager

- **GRAM** = provides a standardised interface to submit jobs to LRMs.
- Clients submit a job request to GRAM
- GRAM translates into something a(ny) LRM can understand
 - Same job request can be used for many different kinds of LRM

Job Management on a Grid



Two versions of GRAM

- There are two versions of GRAM
 - GT2
 - Own protocols
 - Older
 - More widely used
 - No longer actively developed
 - GT4
 - Web services
 - Newer
 - New features go into GRAM4

- In this module, will be using GT2

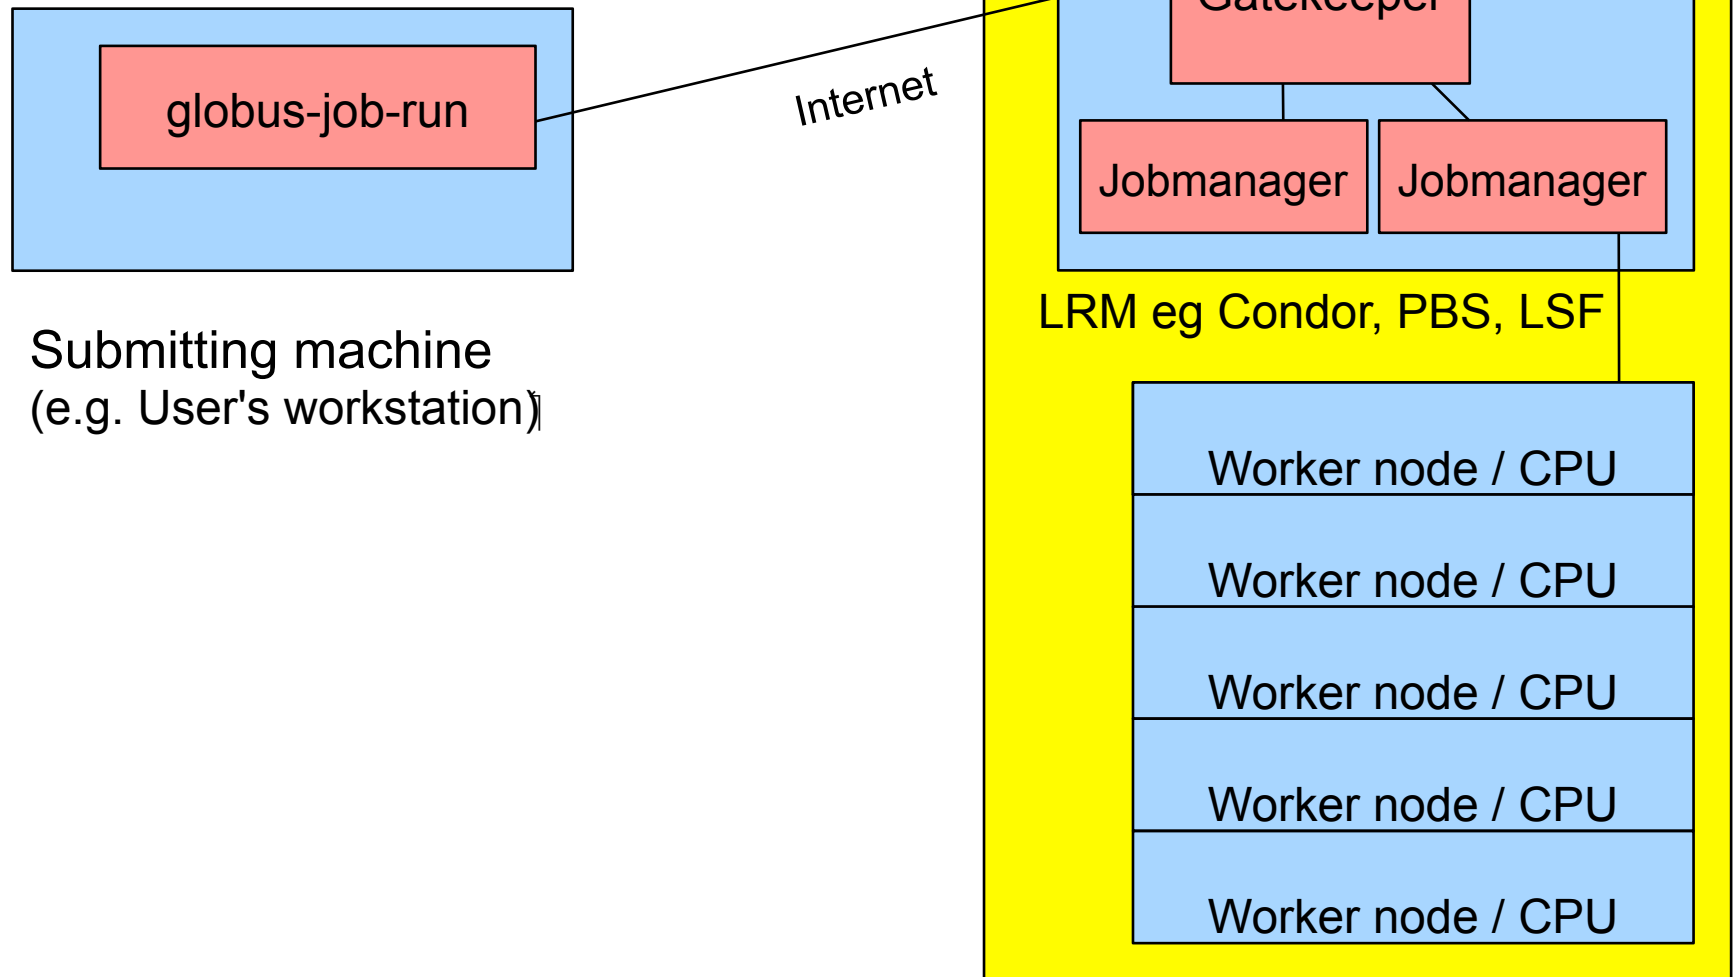
GRAM's abilities

- Given a job specification:
 - ❑ Creates an environment for the job
 - ❑ Stages files to and from the environment
 - ❑ Submits a job to a local resource manager
 - ❑ Monitors a job
 - ❑ Sends notifications of the job state change
 - ❑ Streams a job's stdout/err during execution

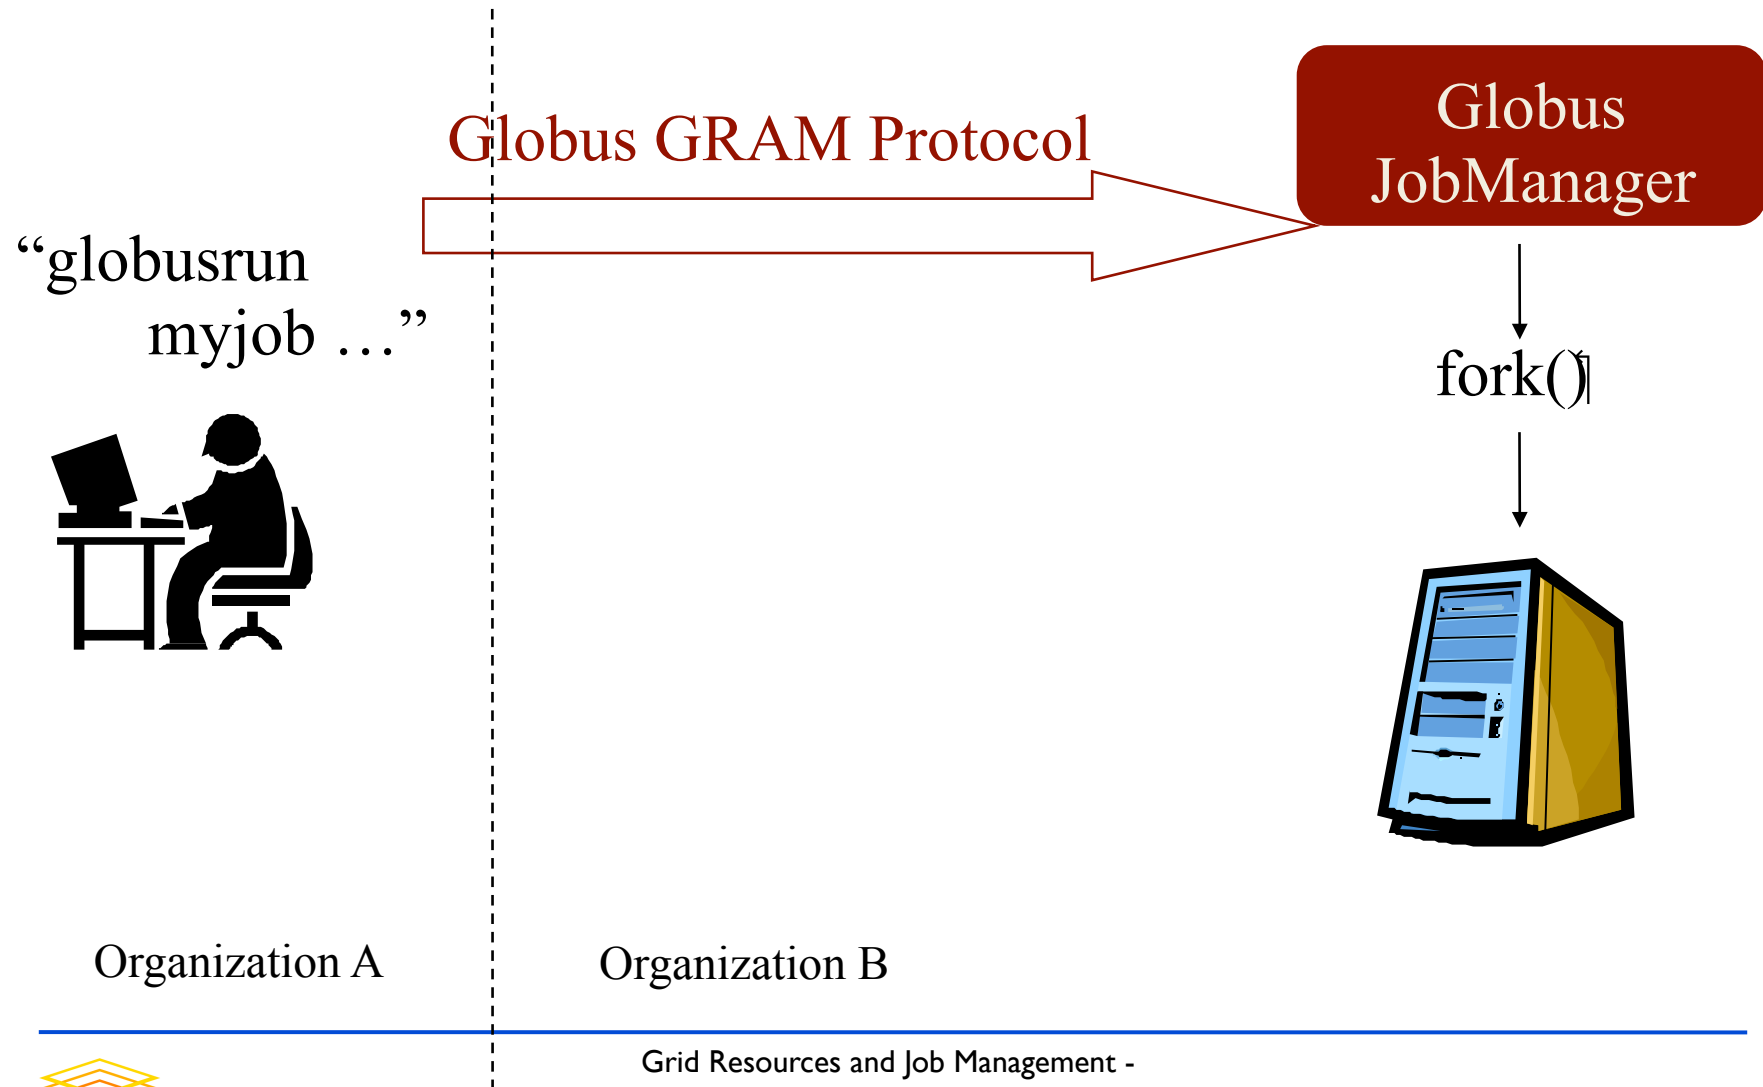
GRAM components

- Clients –
 - ❑ eg. globus-job-submit, globus-run
- Gatekeeper
 - ❑ Server
 - ❑ Accepts job submissions
 - ❑ Handles security
- Jobmanager
 - ❑ Knows how to send a job into the local resource manager
 - ❑ Different job managers for different LRMs

GRAM components



Remote Resource Access: Globus



Job management layers

- High level user tools
(Panda, Swift, Pegasus, ...)
- OSG abstractions
- DAGman
- Condor
- Adding some Security
- GRAM, the Grid protocol
- Clusters and Local
Resource Managers



Security on the Grid

- GRAM is using **GSI** (Grid Security Infrastructure) to **authenticate** users and servers/services using x509 certificates
- An extended certificate identifies the users, their roles and the groups they belong to and it is the basis for the **authorization**
- Mechanisms are in place to guarantee **integrity** of the messages
- Messages can be encrypted to add **privacy**

Certificates and Proxy (certificates)

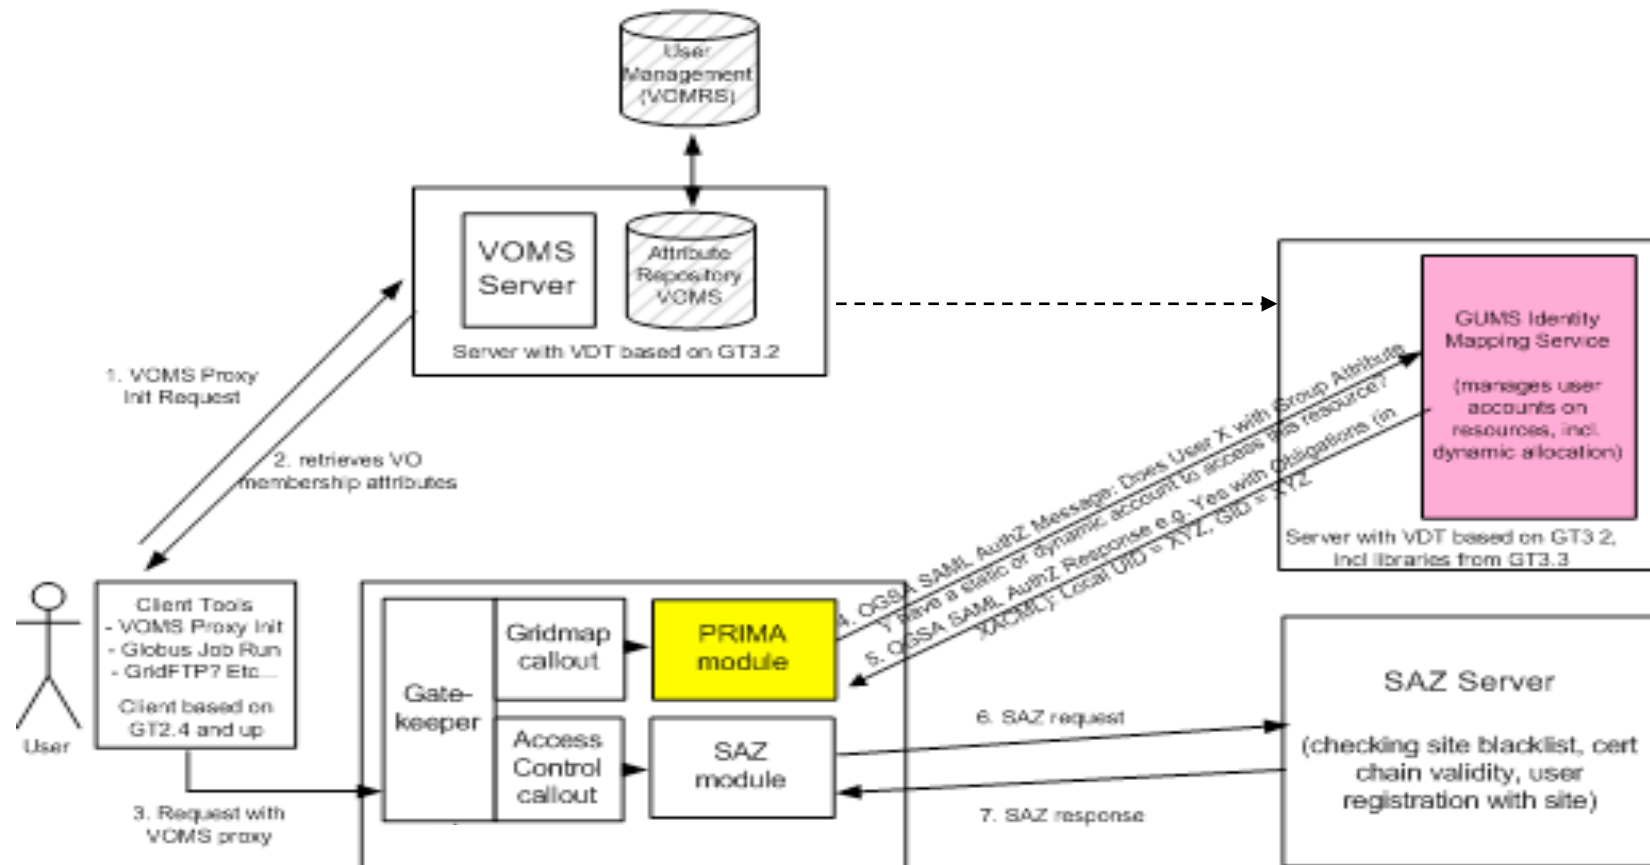
- Certificate content
 - ❑ public part
 - ❑ private key
- Proxy
 - ❑ temporary delegation
- Extended attributes (VOMS)
 - ❑ groups
 - ❑ roles

```
Certificate: Data: Version: 3 (0x2) Serial
Number: 923 (0x39b) Signature Algorithm:
md5WithRSAEncryption Issuer: C=US, O=SDSC, OU=SDSC-
CA, CN=Certificate Authority/UID=certman Validity
Not Before: Jun 22 00:46:02 2006 GMT Not After :
Jul 2 00:46:02 2006 GMT Subject: C=US, O=SDSC,
OU=SDSC, CN=Account Train99/UID=train99 Subject
Public Key Info: Public Key Algorithm:
rsaEncryption RSA Public Key: (2048 bit)
Modulus (2048 bit):
00:af:93:40:80:ce:14:68:d6:6c:67:89:45:0c:3e:
30:98:38:35:c9:bd:b5:08:00:17:4c:e1:fb:38:50:
bd:97:f5:41:92:e7:6e:c4:6f:dc:ad:52:2c:e0:2a:
54:83:79:45:fb:5d:e2:f5:a5:cf:42:94:45:98:22:
d9:5b:81:93:e2:46:5f:e0:7f:71:5f:2d:b0:4a:82:
21:7d:f2:41:f7:b6:33:eb:59:93:f1:71:e3:79:ea:
c0:1b:5e:07:c6:d5:c2:67:41:56:73:d8:1f:a3:fb:
32:4b:f5:96:9f:65:f5:0a:f0:28:d5:90:d6:b0:dc:
4b:29:85:aa:8b:b7:d5:c0:f3:45:28:f9:af:80:7a:
88:40:40:21:60:ea:14:cd:8a:8e:53:40:67:c5:47:
51:bc:95:76:1e:90:b0:ee:ee:41:5a:ec:d4:4c:3c:
ea:eb:2f:f1:55:82:d8:b2:36:d9:92:88:bd:b6:93:
eb:46:69:3b:3a:e2:15:54:82:c0:30:4b:a9:54:3c:
af:52:4e:a5:71:40:a1:58:21:2e:ab:6d:c4:7c:59:
5d:68:b6:95:80:0e:12:91:51:90:0e:38:84:3f:de:
07:99:43:86:a1:0f:70:01:2f:3c:bf:e3:47:b2:16:
67:eb:00:6b:c4:7d:d8:e5:39:77:ac:29:cc:76:94:
2b:d3 Exponent: 65537 (0x10001) X509v3
extensions: X509v3 Basic Constraints:
CA:FALSE Netscape Cert Type: SSL
Client, S/MIME, Object Signing Netscape Comment:
OpenSSL Generated Certificate Netscape CA
Revocation Url: http://www.sdsc.edu/CA/
SDSC_CRL.pem X509v3 Subject Key Identifier:
E1:E3:C9:6E:A6:CF:2C:FC:D3:B7:51:F6:03:66:98:C5:18:71:60:F8
X509v3 Authority Key Identifier:
keyid:BF:A3:87:2C:F6:0D:74:BD:48:6C:0E:27:BF:01:E4:F2:4F:
46:BA:27 DirName:/C=US/O=SDSC/OU=SDSC-CA/
CN=Certificate Authority/UID=certman serial:
00 Signature Algorithm: md5WithRSAEncryption
93:b2:78:07:d9:72:e2:71:d7:66:83:0c:d3:97:0c:9e:24:33:
4e:e3:48:28:9c:44:7e:31:13:70:cc:f8:4a:5d:bc:64:84:3e:
aa:fa:da:86:3f:5e:f8:4a:72:a1:59:57:5a:89:49:5a:2d:c9:
09:5c:a5:69:6e:65:f7:85:8b:07:57:f1:6a:cb:6e:e5:00:17:
```

Virtual Organization Management System

- User management delegated to VO (Virtual Organization)
 - flexible
 - scalable
- Introduction of groups and roles
- Dynamic calls

VOMS Architecture



Submitting a job with GRAM

- **globus-job-run** command

```
$ globus-job-run rookery.uchicago.edu /bin/hostname
```

- Run '/bin/hostname' on the resource rookery.uchicago.edu

- We don't care what LRM is used on 'rookery'. This command works with any LRM.

The client can describe the job with GRAM's Resource Specification Language (RSL)

- Example:

```
& (executable = a.out)
  (directory = /home/nobody )
  (arguments = arg1 "arg 2")
```

- Submit with:

```
globusrun -f spec.rsl -r
rookery.uchicago.edu
```

Use other programs to generate RSL

- RSL job descriptions can become very complicated
- We can use other programs to generate RSL for us
 - Example: Condor-G – next section

Job management layers

- High level user tools
(Panda, Swift, Pegasus, ...)
- OSG abstractions
- DAGman
- Condor
- Adding some Security
- GRAM, the Grid protocol
- Clusters and Local
Resource Managers



Condor

- is a software system that creates an HTC environment
 - ❑ Created at [UW-Madison](#)
- Condor is a specialized workload management system for compute-intensive jobs.
 - ❑ Detects machine availability
 - ❑ Harnesses available resources
 - ❑ Uses remote system calls to send R/W operations over the network
 - ❑ Provides powerful resource management by *matching* resource owners with consumers (broker)

How Condor works

Condor provides:

- a job queueing mechanism
- scheduling policy
- priority scheme
- resource monitoring, and
- resource management.

Users **submit** their serial or parallel jobs to Condor,

Condor places them into a **queue**,

... chooses **when** and **where** to run the jobs based upon a policy,

... carefully **monitors** their progress, and

... ultimately **informs** the user upon completion.

Condor - features

- Checkpoint & migration
- Remote system calls
 - Able to transfer data files and executables across machines
- Job ordering
- *Job requirements and preferences can be specified via powerful expressions*

Condor lets you manage a large number of jobs.

- Specify the jobs in a file and submit them to Condor
- Condor runs them and keeps you notified on their progress
 - ❑ Mechanisms to help you manage huge numbers of jobs (1000's), all the data, etc.
 - ❑ Handles inter-job dependencies (DAGMan)
- Users can set Condor's job priorities
- Condor administrators can set user priorities
- Can do this as:
 - ❑ Local resource manager (LRM) on a compute resource
 - ❑ Grid client submitting to GRAM (as Condor-G)

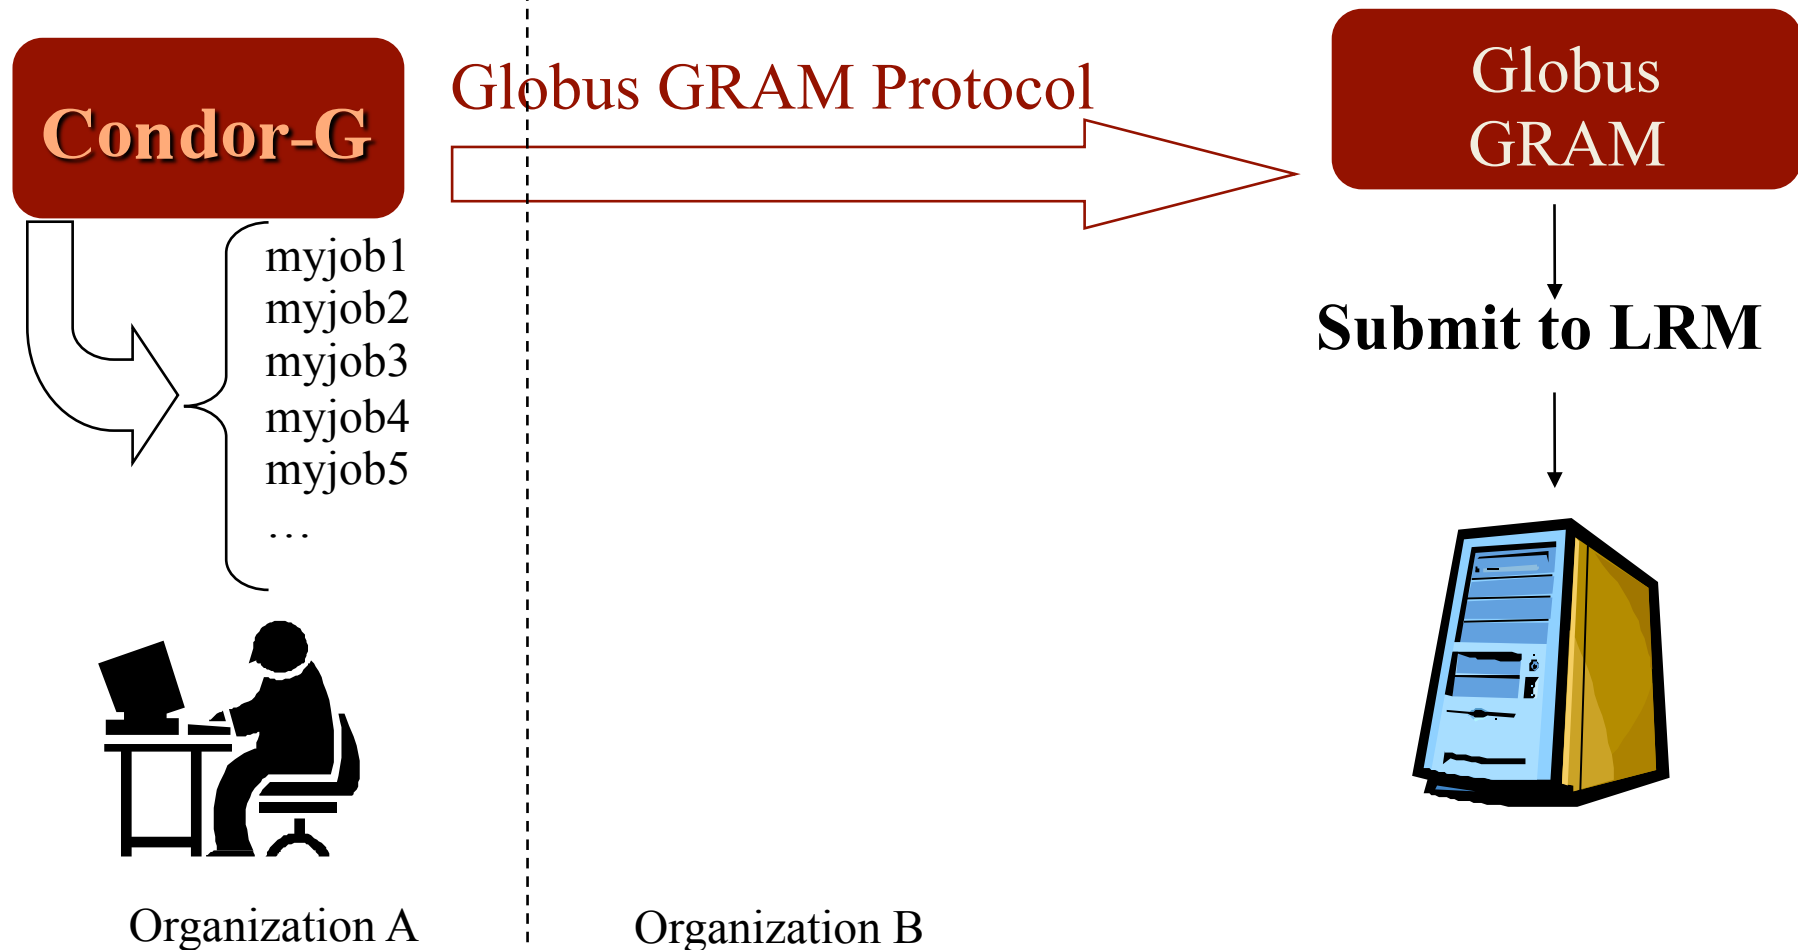
Condor-G

- is the job management part of Condor.
- *Hint*: Install Condor-G to submit to resources accessible through a Globus interface.
- Condor-G does not *create* a grid service.
- It only deals with *using* remote grid services.

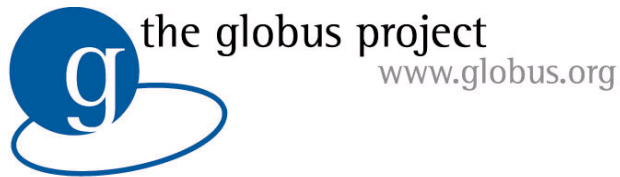
Condor-G ...

- does whatever it takes to run your jobs, even if ...
 - ❑ The gatekeeper is temporarily unavailable
 - ❑ The job manager crashes
 - ❑ Your local machine crashes
 - ❑ The network goes down

Remote Resource Access: Condor-G + Globus + Condor



Condor-G: Access non-Condor Grid resources



- middleware deployed across entire Grid
- remote access to computational resources
- dependable, robust data transfer



- job scheduling across multiple resources
- strong fault tolerance with checkpointing and migration
- layered over Globus as “personal batch system” for the Grid

Four Steps to Run a Job with Condor

- These choices tell Condor
 - ❑ **how**
 - ❑ **when**
 - ❑ **where** to run the job,
 - ❑ and describe exactly **what** you want to run.
- Choose a Universe for your job
- Make your job batch-ready
- Create a *submit description* file
- Run *condor_submit*

I. Choose a Universe

- There are many choices
 - ❑ **Vanilla:** any old job
 - ❑ **Standard:** checkpointing & remote I/O
 - ❑ **Java:** better for Java jobs
 - ❑ **MPI:** Run parallel MPI jobs
 - ❑ Virtual Machine: Run a virtual machine as job
 - ❑ ...
- *For now, we'll just consider vanilla*

2. Make your job batch-ready

- Must be able to run in the background:
 - no interactive input, windows, GUI, etc.
- Condor is designed to run jobs as a batch system, with pre-defined inputs for jobs
- Can still use STDIN, STDOUT, and STDERR (the keyboard and the screen), but files are used for these instead of the actual devices
- Organize data files

3. Create a Submit Description File

- A plain ASCII text file
- Condor does not care about file extensions
- Tells Condor about your job:
 - ❑ Which executable to run and where to find it
 - ❑ Which universe
 - ❑ Location of input, output and error files
 - ❑ Command-line arguments, if any
 - ❑ Environment variables
 - ❑ Any special requirements or preferences

Simple Submit Description File

```
# myjob.submit file
# Simple condor_submit input file
# (Lines beginning with # are comments)
# NOTE: the words on the left side are not
#       case sensitive, but filenames are!
Universe      = vanilla
Executable    = analysis
Log            = my_job.log
Queue
```


4. Run `condor_submit`

- You give *condor_submit* the name of the submit file you have created:

```
condor_submit my_job.submit
```

- *condor_submit* parses the submit file

Another Submit Description File

```
# Example condor_submit input file
# (Lines beginning with # are comments)
# NOTE: the words on the left side are not
#       case sensitive, but filenames are!
Universe      = vanilla
Executable    = /home/wright/condor/my_job.condor
Input         = my_job.stdin
Output        = my_job.stdout
Error         = my_job.stderr
Arguments     = -arg1 -arg2
InitialDir    = /home/wright/condor/run_1
Queue
```

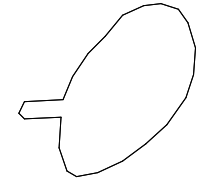
Details

- Lots of options available in the submit file
- Commands to
 - watch the queue,
 - the state of your pool,
 - and lots more
- You'll see much of this in the hands-on exercises.

Other Condor commands

- `condor_q` – show status of job queue
- `condor_status` – show status of compute nodes
- `condor_rm` – remove a job
- `condor_hold` – hold a job temporarily
- `condor_release` – release a job from hold

Submitting more complex jobs



- express dependencies between jobs
⇒ WORKFLOWS
- And also, we would like the workflow to be managed even in the face of failures

Want other Scheduling possibilities?

Use the Scheduler Universe

- In addition to VANILLA, another job universe is the *Scheduler Universe*.
- Scheduler Universe jobs run on the submitting machine and serve as a meta-scheduler.
- **Condor's Scheduler Universe lets you set up and manage job workflows.**
- DAGMan meta-scheduler included
 - DAGMan manages these jobs

Job management layers

- High level user tools
(Panda, Swift, Pegasus, ...)
- OSG abstractions
- **DAGman**
- Condor
- Adding some Security
- GRAM, the Grid protocol
- Clusters and Local
Resource Managers




DAGMan

- Directed Acyclic Graph Manager
- DAGMan allows you to specify the *dependencies* between your Condor jobs, so it can *manage* them automatically for you.
- (e.g., “Don’t run job “B” until job “A” has completed successfully.”)

What is a DAG?

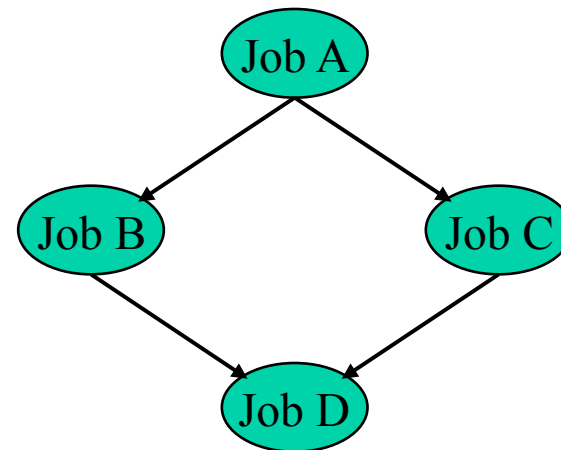
- A DAG is the **data structure** used by DAGMan to represent these dependencies.
- Each job is a **“node”** in the DAG.
- Each node can have any number of “parent” or “children” nodes – as long as there are **no loops**!



Defining a DAG

- A DAG is defined by a *.dag file*, listing each of its nodes and their dependencies:

```
# diamond.dag
Job A a.sub
Job B b.sub
Job C c.sub
Job D d.sub
Parent A Child B C
Parent B C Child D
```



- each node will run the Condor job specified by its accompanying *Condor submit file*

Submitting a DAG

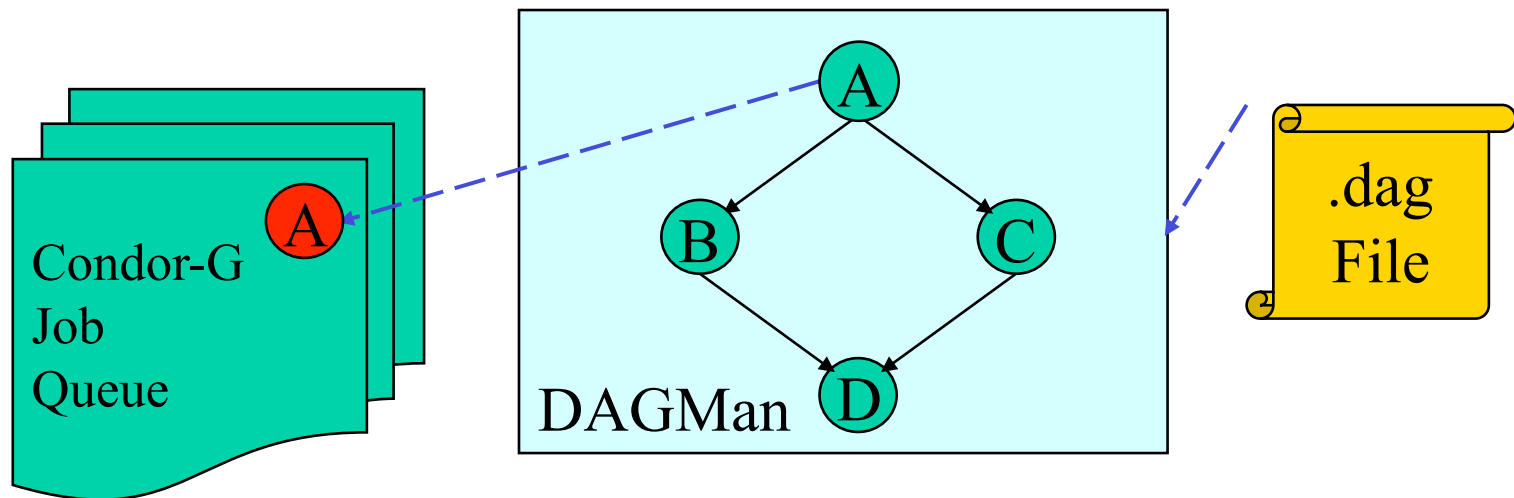
- To start your DAG, just run **condor_submit_dag** with your .dag file, and Condor will start a personal DAGMan daemon which to begin running your jobs:

```
% condor_submit_dag diamond.dag
```

- `condor_submit_dag` submits a Scheduler Universe Job with DAGMan as the executable.
- Thus the DAGMan daemon itself **runs as a Condor job**, so you don't have to baby-sit it.

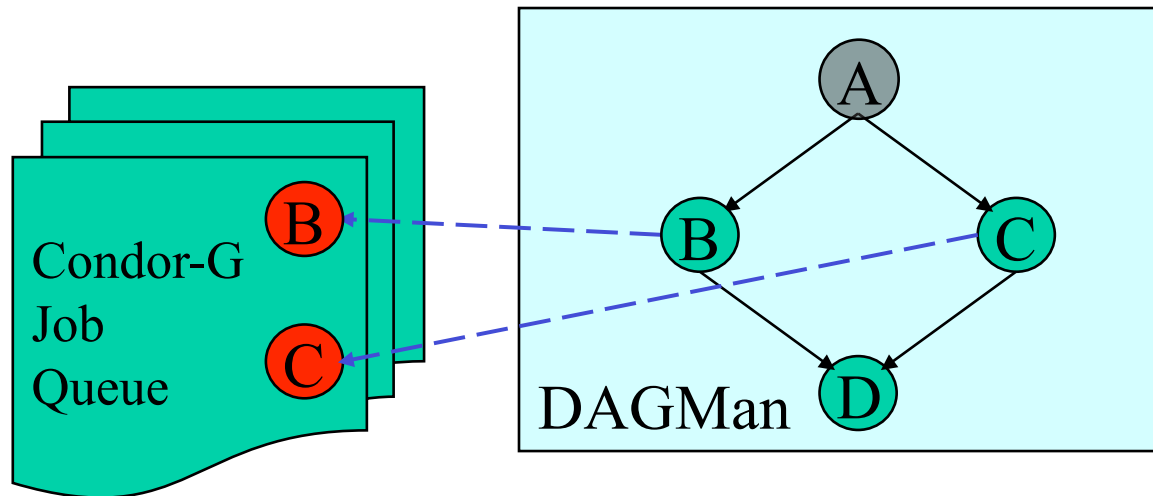
Running a DAG

- DAGMan acts as a “meta-scheduler”, managing the submission of your jobs to Condor-G based on the DAG dependencies.



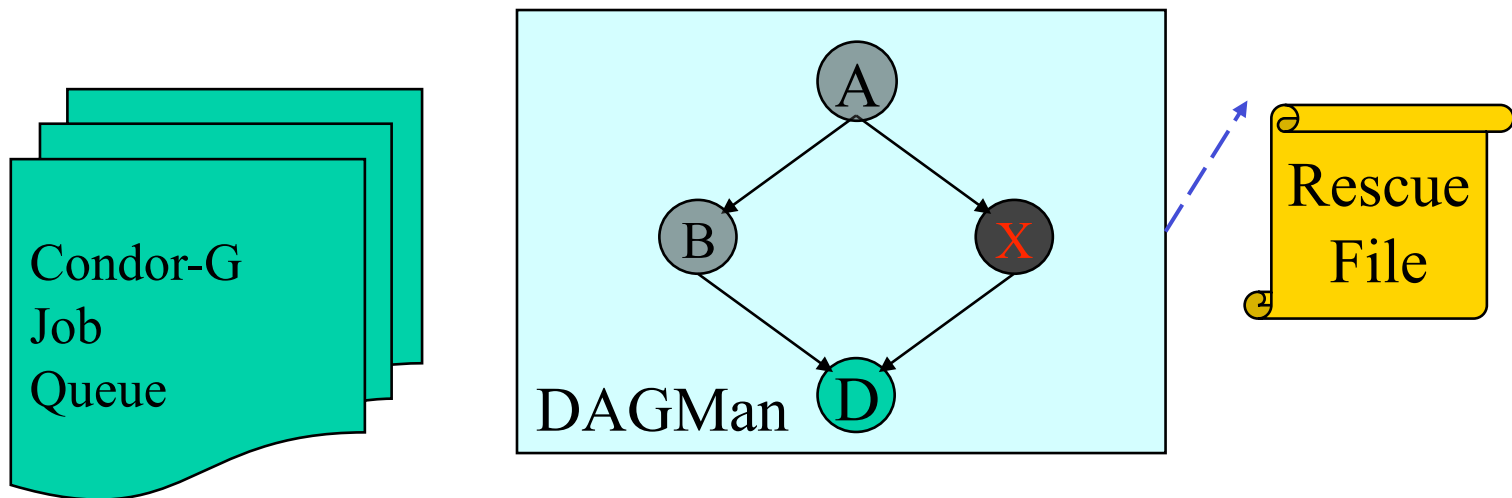
Running a DAG (cont'd)

- DAGMan holds & submits jobs to the Condor-G queue at the appropriate times.



Running a DAG (cont'd)

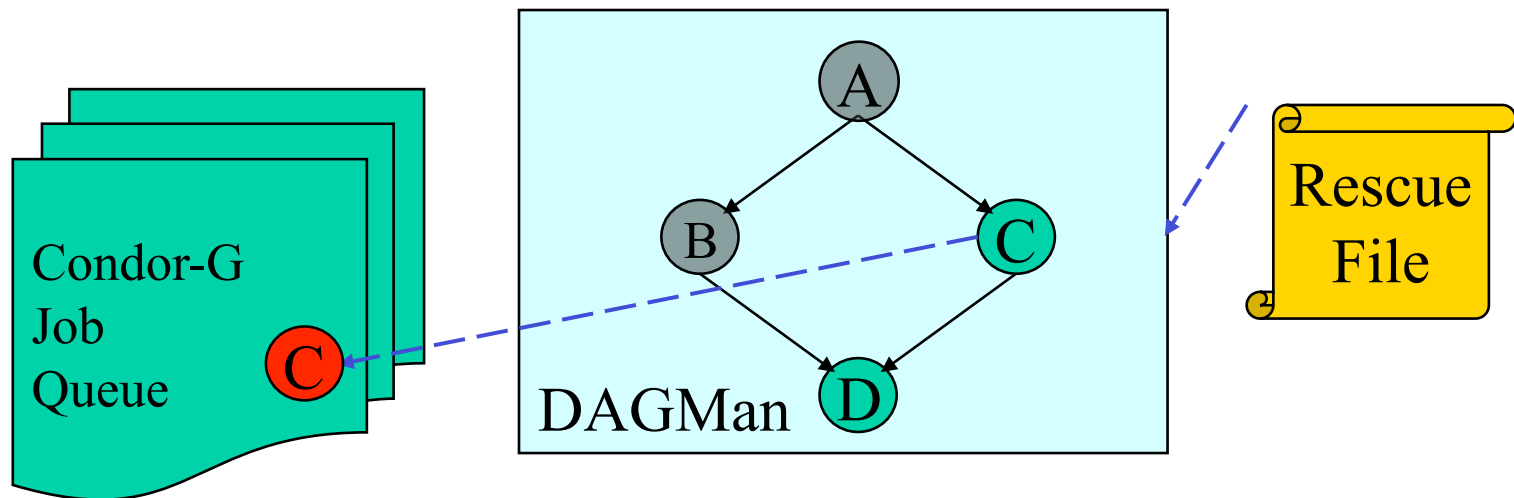
- In case of a job failure, DAGMan continues until it can no longer make progress, and then creates a *“rescue” file* with the current state of the DAG.



Recovering a DAG


-- fault tolerance

- Once the failed job is ready to be re-run, the rescue file can be used to restore the prior state of the DAG.




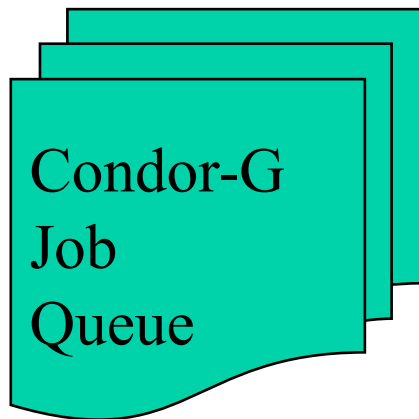
Recovering a DAG (cont'd)

- Once that job completes, DAGMan will continue the DAG as if the failure never happened.



Finishing a DAG

- Once the DAG is complete, the DAGMan job itself is finished, and exits.




We have seen how Condor:

- ... monitors submitted jobs and reports progress
- ... implements your policy on the execution order of the jobs
- ... keeps a log of your job activities

Long jobs: if my jobs run for weeks

...



- What happens to my job when
 - ❑ a machine is shut down
 - ❑ there is a network outage, or
 - ❑ another job with higher priority preempts it?
- Do I lose all of those hours or days of computation time??
- What happens when they get pre-empted?
- *How can I add fault tolerance to my jobs?*

Condor's Standard Universe to the rescue!

- Condor can support various combinations of features/environments in different “Universes”
- Different Universes provide different functionalities to your job:
 - ❑ Vanilla: Run any serial job
 - ❑ Scheduler: Plug in a scheduler
 - ❑ Standard: Support for *transparent process checkpoint and restart*

**provides two important services to your job:
process checkpoint
remote system calls.**

Process Checkpointing

- Condor's process checkpointing mechanism *saves the entire state of a process into a checkpoint file*
 - Memory, CPU, I/O, etc.
- The process can then be *restarted* from the point it left off
- Typically no changes to your job's source code needed—however, your job must be relinked with Condor's Standard Universe support library

Job management layers

- High level user tools
(Panda, Swift, Pegasus, ...)
- OSG abstractions
- DAGman
- Condor
- Adding some Security
- GRAM, the Grid protocol
- Clusters and Local
Resource Managers



OSG & job submissions

- OSG sites present interfaces allowing remotely submitted jobs to be accepted, queued and executed locally.
- OSG supports the Condor-G job submission client which interfaces to either the pre-web service or web services GRAM Globus interface at the executing site.
- Job managers at the backend of the GRAM gatekeeper support job execution by local Condor, LSF, PBS, or SGE batch systems.

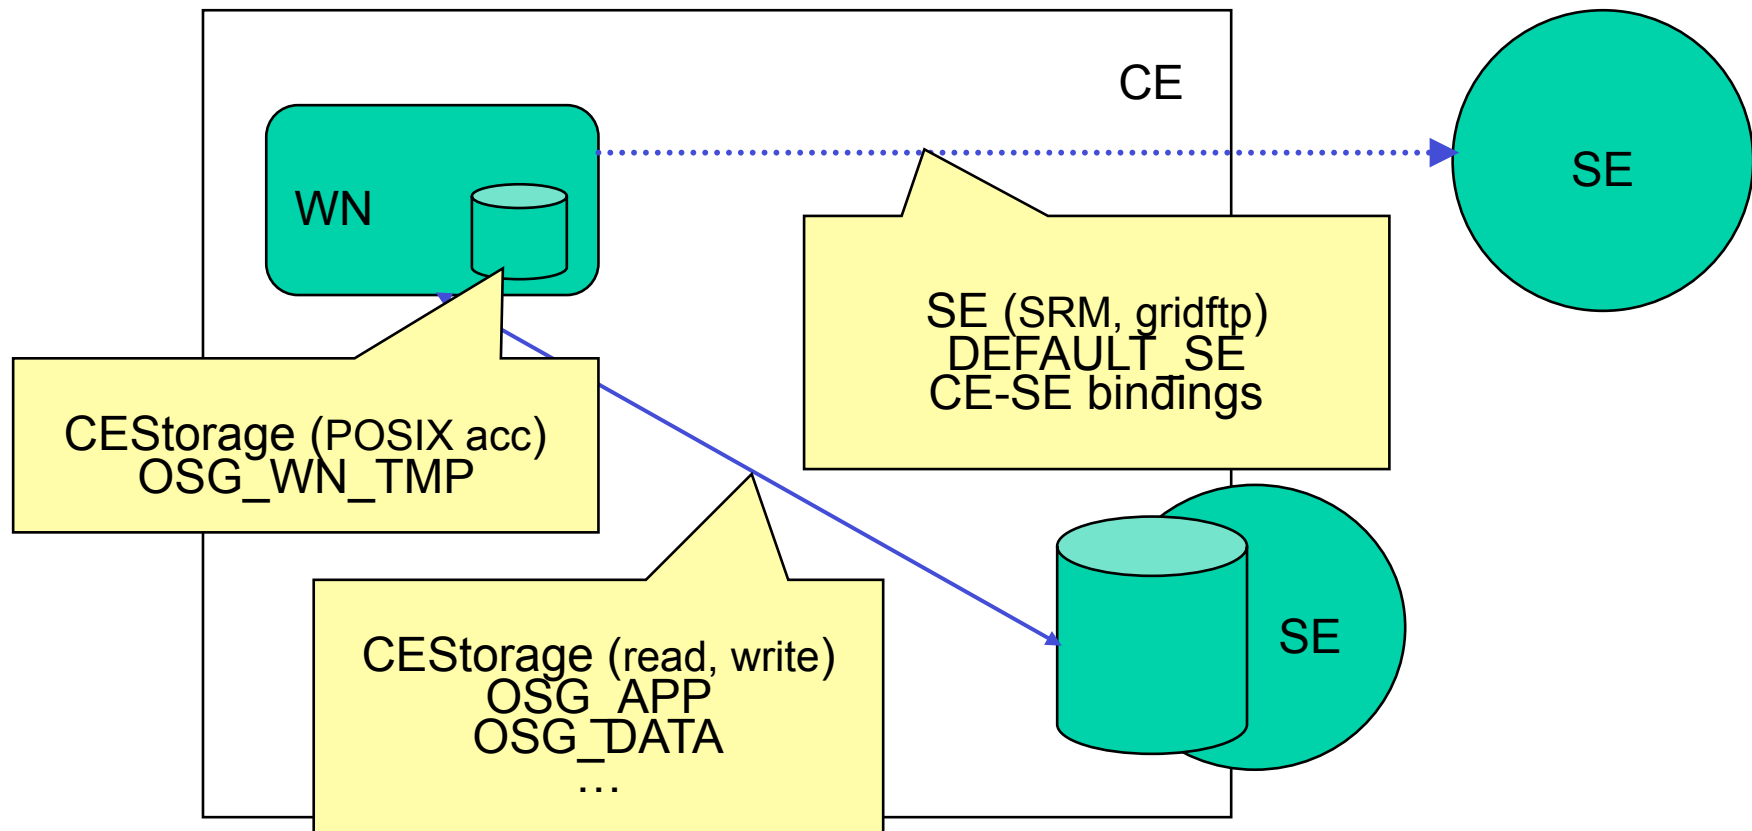
OSG and Environment

- Information systems to discover and describe resources
- Guidelines on how to execute jobs
- Guidelines on how to use available disk spaces
- Jobs find an uniform standard environment
 - OSG_AAA variables

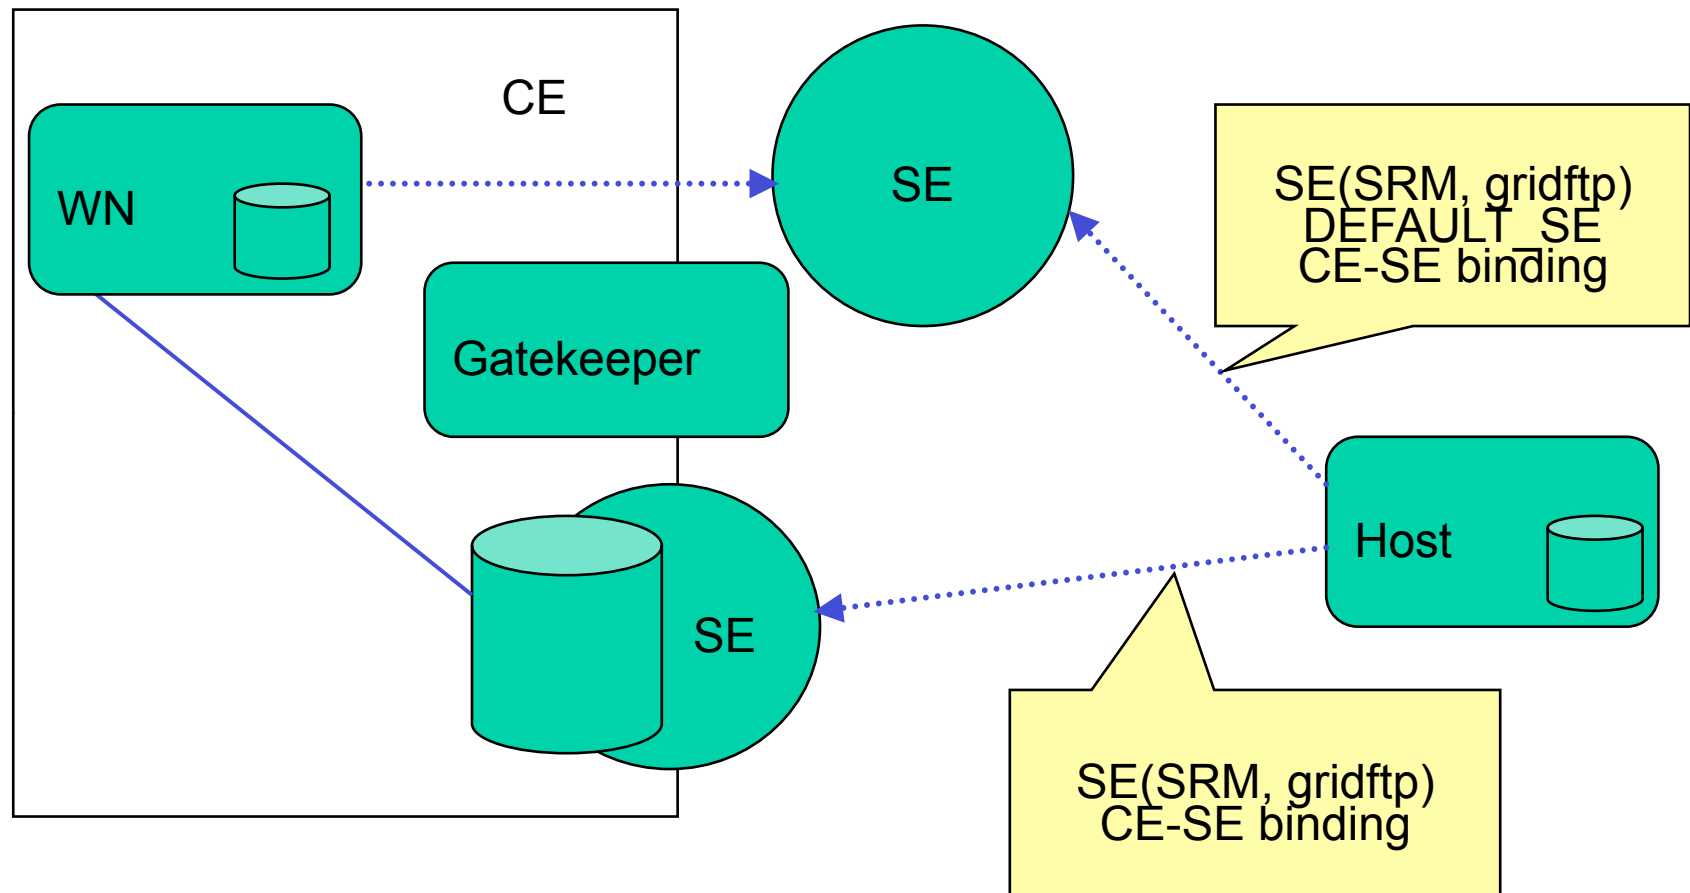
Everything has a place

- OSG_APP
 - ❑ directory to install applications
- OSG_GRID
 - ❑ directory to find grid clients (OSG:WN)
- OSG_WN_TMP
 - ❑ local working directory
- OSG_DATA
 - ❑ shared directory
- OSG_SITE_READ, OSG_SITE_WRITE
 - ❑ optimized for efficient read/write
- DEFAULT_SE
 - ❑ storage element close to this computing element

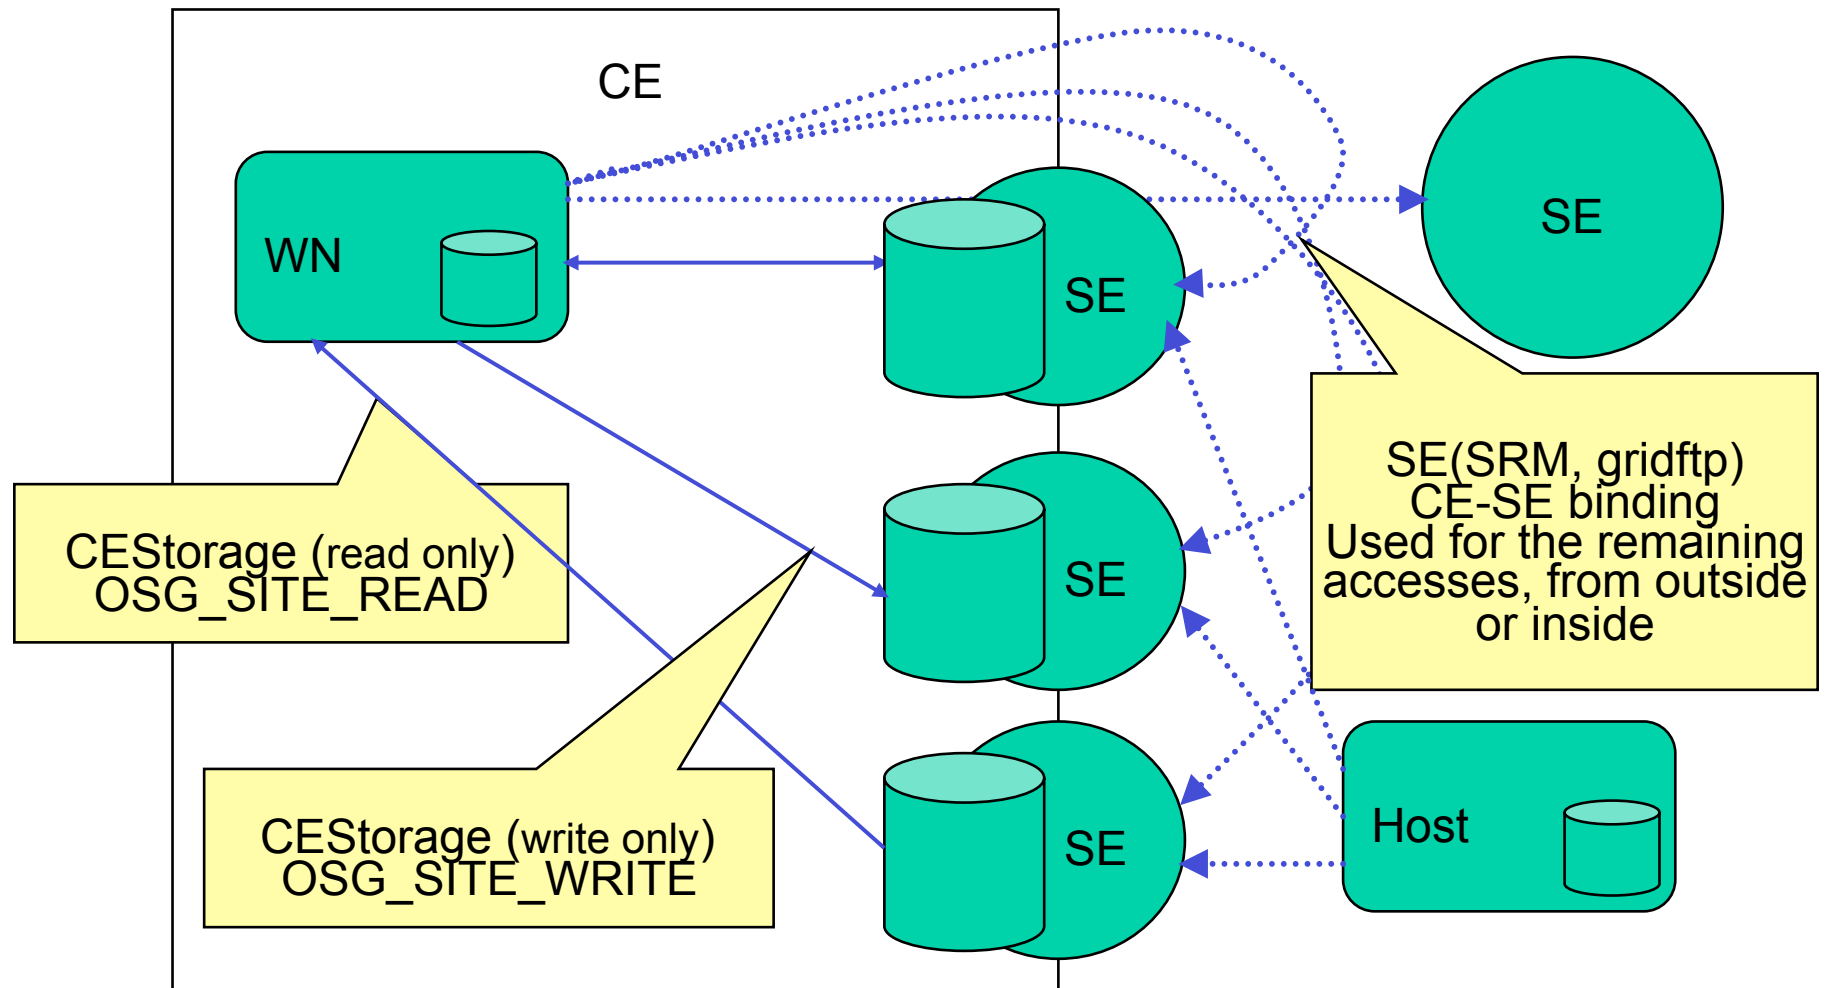
CE View



Outside View



CE View (OSG_WRITE, OSG_READ)



Acknowledgments:

The Globus and Condor parts of this presentation are based on:

Grid Resources and Job Management



Open Science Grid

Jaime Frey and Becky
Gietzel

Condor Project

U. Wisconsin-Madison

?

!

Getting started with the hands-on session

- Pacman is a packet manager that we use in osg to install our software
- OSG-Client is a package that includes a lot of applications to interact with the grid (all the ones that you'll need for today's hands-on sessions)
- The virtual machine and the host grid07 have this software already installed for you but here I'll show you how to start from zero.

What to do

- Get Pacman
- Install and setup pacman
- make osg-client directory
- Install OSG-Client
- Answer installation questions
- Post install setup (installing CA certificates)

Installation

```
wget http://atlas.bu.edu/~youssef/pacman/sample_cache/tarballs/pacman-latest.tar.gz
tar -xzf pacman-latest.tar.gz
cd pacman-3.26
source setup.sh
cd ..
mkdir osg-client
cd osg-client
pacman --get OSG:client
[installation starts]
yall [when asked if you trust the software caches]
[some time later]
y [when asked if you accept licenses]
l ['L' when asked if you want to install locally the certificates]
```

Installation (cont)

[at the end of the installation you need to install the CA certificates. VDT is warning about it. This is done to allow you to choose which CA (Certification Authority) you trust. You do that editing `$VDT_LOCATION/vdt/etc/vdt-update-certs.conf`, e.g. uncomment the line about OSG CA-bundle in the following steps]

```
source setup.sh  
vi $VDT_LOCATION/vdt/etc/vdt-update-certs.conf  
source $VDT_LOCATION/vdt-questions.sh; \  
$VDT_LOCATION/vdt/sbin/vdt-setup-ca-certificates
```

[Now you can continue with the hands-on session:

<http://www.ci.uchicago.edu/osgedu/schools/2009/megs/intro.html>

complete the first 2 session: basic job submission and job submission with Condor]