

---

# Grid Compute Resources and Job Management



**Open Science Grid**

New Mexico Grid School – April 9, 2009

Marco Mambelli – University of Chicago

[marco@hep.uchicago.edu](mailto:marco@hep.uchicago.edu)

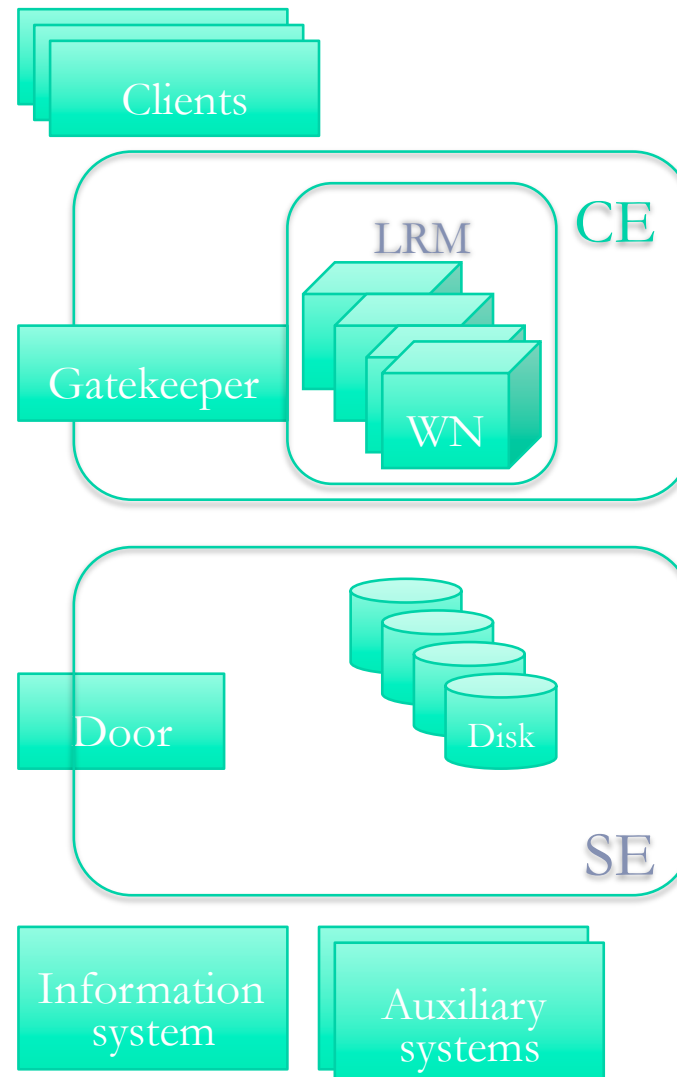
---

# Outline

- Grid abstraction
- Clusters and Local Resource Managers
- GRAM, the Grid protocol
- Adding some Security
- Condor
- Something more with Condor
- DAGman
- OSG abstractions

# Grid/OSG components

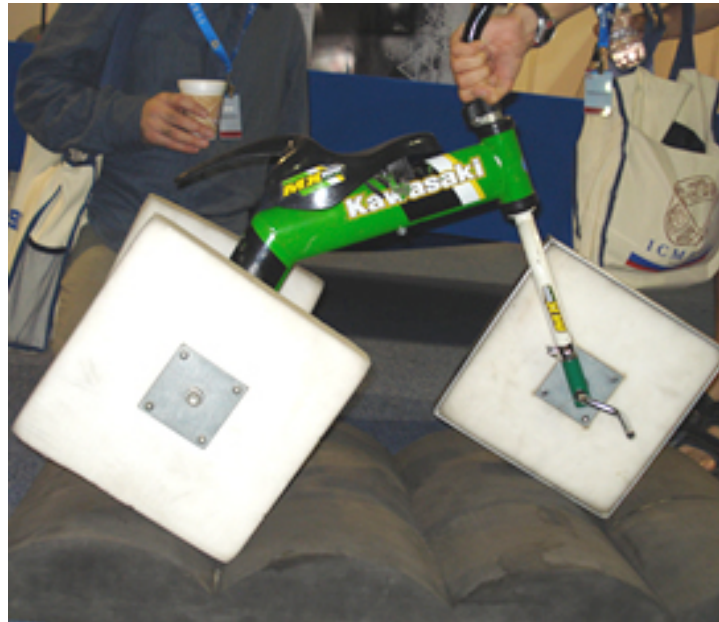
- Standard installation
- Set of services or resources containing data and providing processing power
- Computing elements (CE)
- Storage elements (SE)
- Information systems
- Clients
- How to interact?
  - validation and testing
  - active use



# Grid as Abstraction

- Hides lower level components/fabrics
  - you can use it without knowing what is behind
- Provides leverage
  - do so much with so little
- Limit flexibility
- Suggest usage patterns
- Gain/loss

# Do not reinvent the wheel!



# Job management layers

- High level user tools  
(Panda, Swift, Pegasus, ...)
- OSG abstractions
- DAGman
- Condor
- Adding some Security
- GRAM, the Grid protocol
- Clusters and Local  
Resource Managers



# The power of leverage



# Good use of the Grid Abstraction

- Leaks are likely
  - ❑ all abstraction “lie”
  - ❑ some leaks are desired
- Before you can abstract you must see the details
- Understanding the levels below you understand the abstraction
- To build components
  - ❑ must be fully aware of at least a couple of layers below
  - ❑ to EXCELL you must be VERY familiar with several layers above





# Job management layers

- High level user tools  
(Panda, Swift, Pegasus, ...)
- OSG abstractions
- DAGman
- Condor
- Adding some Security
- GRAM, the Grid protocol
- Clusters and Local  
Resource Managers



# Local Resource Managers (LRM)

- Compute resources have a **local resource manager** (LRM) that controls:
  - ❑ Who is allowed to run jobs
  - ❑ How jobs run on a specific resource
- *Example policy:*
  - ❑ Each cluster node can run one job.
  - ❑ If there are more jobs, then they must wait in a queue
- LRMs allow nodes in a cluster can be **reserved** for a specific person
- *Examples:* PBS, LSF, Condor

# Job management layers

- High level user tools  
(Panda, Swift, Pegasus, ...)
- OSG abstractions
- DAGman
- Condor
- Adding some Security
- GRAM, the Grid protocol
- Clusters and Local  
Resource Managers

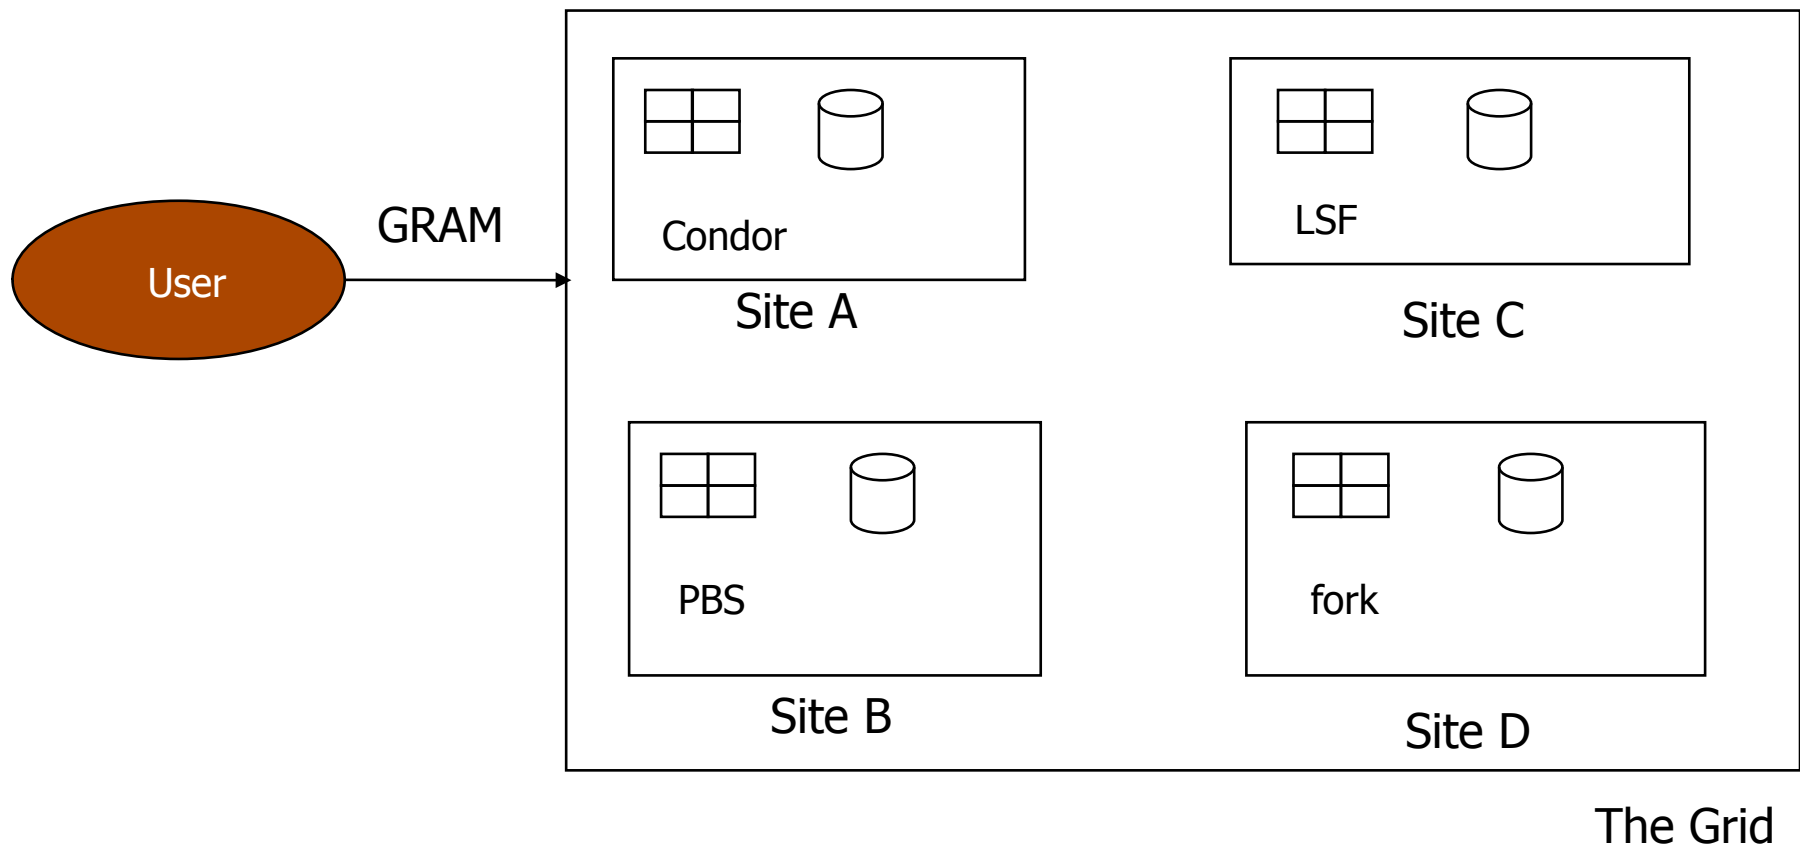


# GRAM

## Globus Resource Allocation Manager

- **GRAM** = provides a standardised interface to submit jobs to LRMs.
- Clients submit a job request to GRAM
- GRAM translates into something a(ny) LRM can understand
  - .... Same job request can be used for many different kinds of LRM

# Job Management on a Grid



# Two versions of GRAM

- There are two versions of GRAM
  - GT2
    - Own protocols
    - Older
    - More widely used
    - No longer actively developed
  - GT4
    - Web services
    - Newer
    - New features go into GRAM4
  
- In this module, will be using GT2

# GRAM's abilities

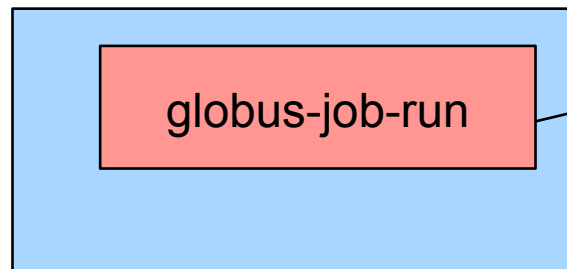
- Given a job specification:
  - ❑ Creates an environment for the job
  - ❑ Stages files to and from the environment
  - ❑ Submits a job to a local resource manager
  - ❑ Monitors a job
  - ❑ Sends notifications of the job state change
  - ❑ Streams a job's stdout/err during execution

# GRAM components

- Clients –
  - ❑ eg. globus-job-submit, globus-run
- Gatekeeper
  - ❑ Server
  - ❑ Accepts job submissions
  - ❑ Handles security
- Jobmanager
  - ❑ Knows how to send a job into the local resource manager
  - ❑ Different job managers for different LRMs

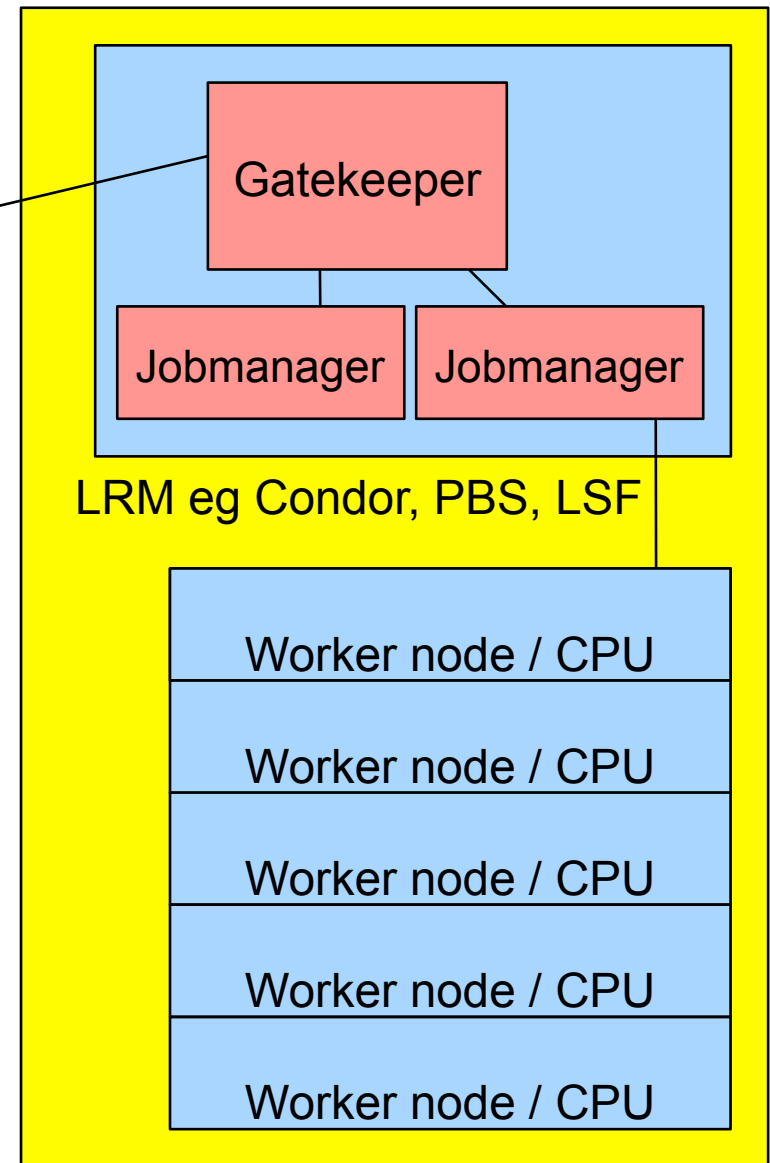


# GRAM components

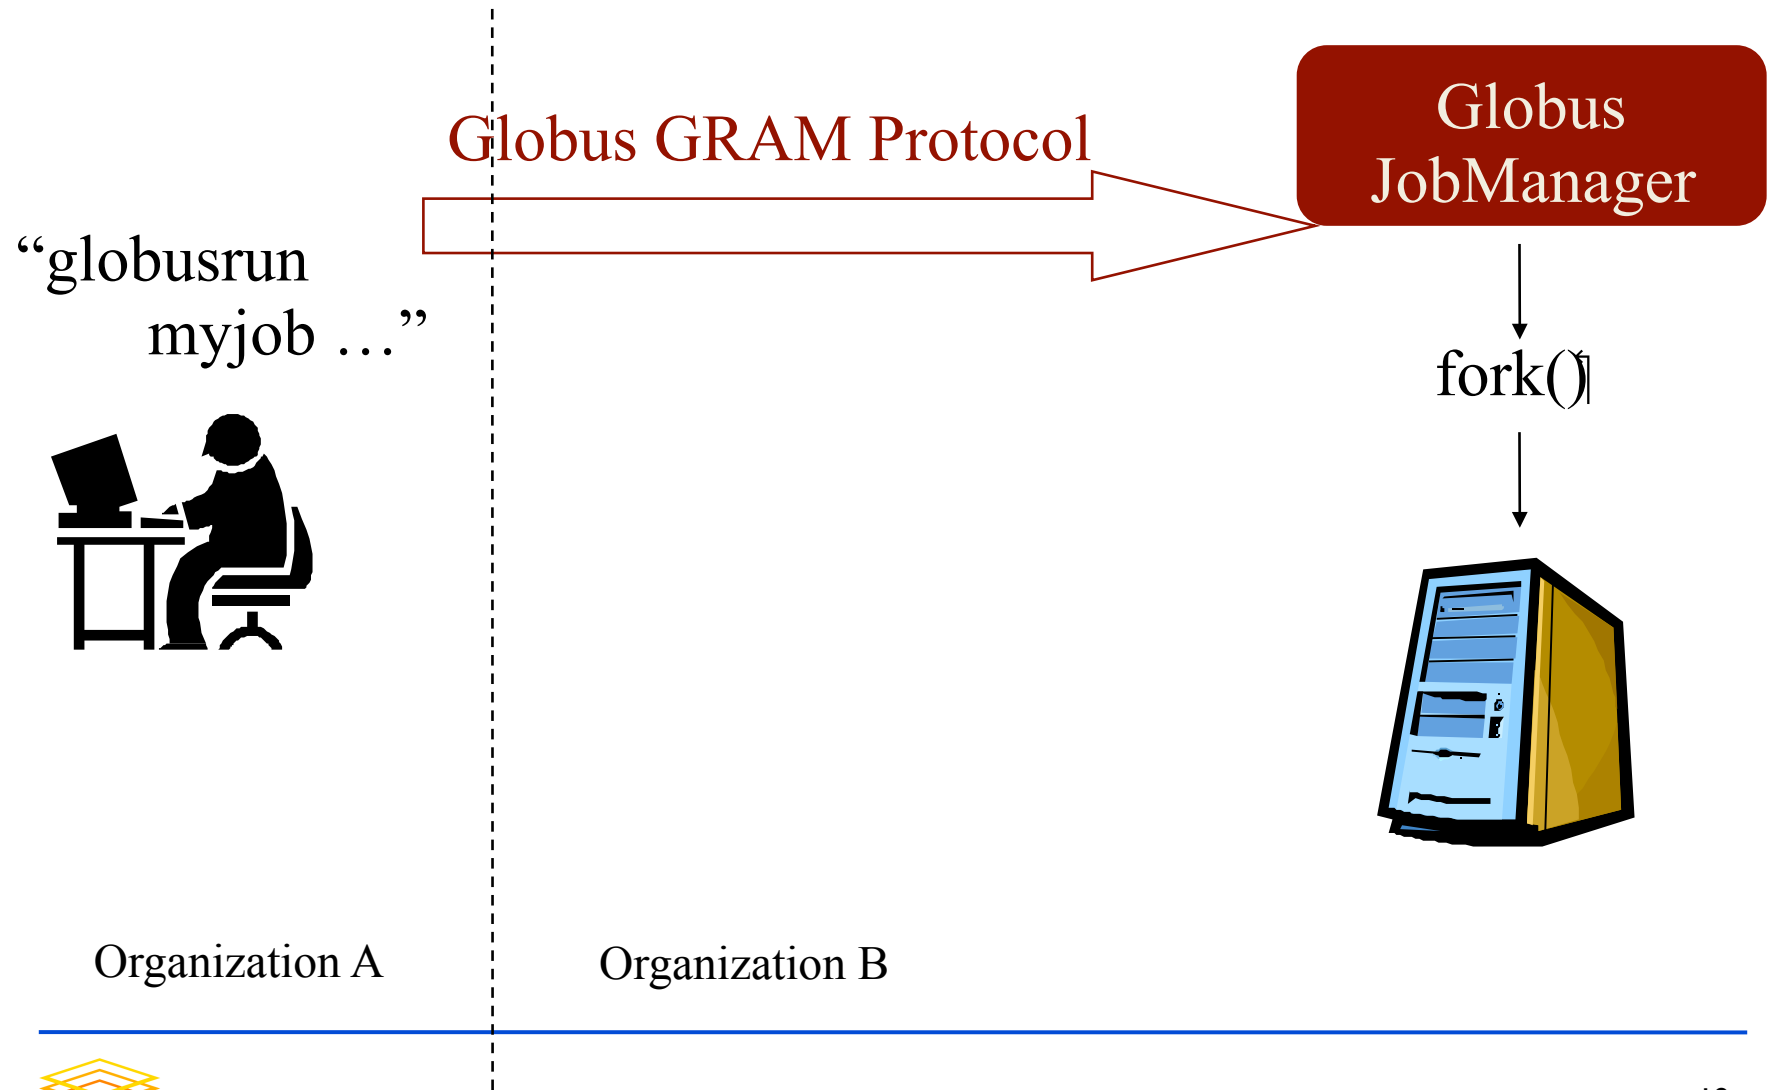


Submitting machine  
(e.g. User's workstation)

Internet



# Remote Resource Access: Globus



# Job management layers

- High level user tools  
(Panda, Swift, Pegasus, ...)
- OSG abstractions
- DAGman
- Condor
- Adding some Security
- GRAM, the Grid protocol
- Clusters and Local  
Resource Managers



# Security on the Grid

- GRAM is using **GSI** (Grid Security Infrastructure) to **authenticate** users and servers/services using x509 certificates
- An extended certificate identifies the users, their roles and the groups they belong to and it is the basis for the **authorization**
- Mechanisms are in place to guarantee **integrity** of the messages
- Messages can be encrypted to add **privacy**

# Certificates and Proxy (certificates)

- Certificate content
  - ❑ public part
  - ❑ private key
- Proxy
  - ❑ temporary delegation
- Extended attributes (VOMS)
  - ❑ groups
  - ❑ roles

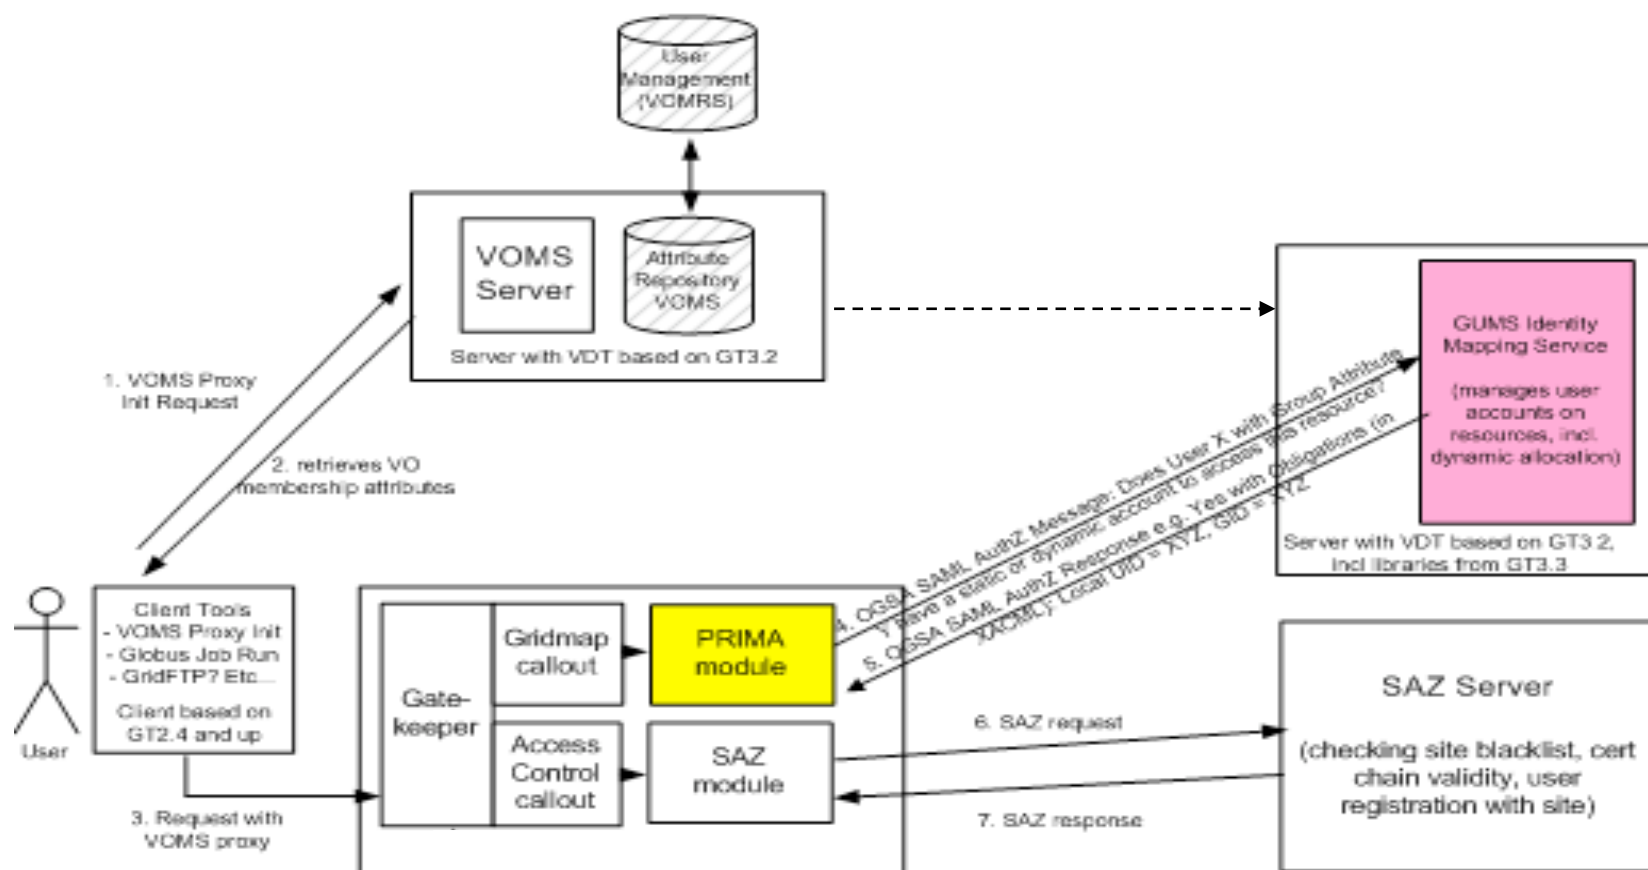
```
Certificate: Data: Version: 3 (0x2) Serial
Number: 923 (0x39b) Signature Algorithm:
md5WithRSAEncryption Issuer: C=US, O=SDSC, OU=SDSC-
CA, CN=Certificate Authority/UID=certman Validity
Not Before: Jun 22 00:46:02 2006 GMT Not After :
Jul 2 00:46:02 2006 GMT Subject: C=US, O=SDSC,
OU=SDSC, CN=Account Train99/UID=train99 Subject
Public Key Info: Public Key Algorithm:
rsaEncryption RSA Public Key: (2048 bit)
Modulus (2048 bit):
00:af:93:40:80:ce:14:68:d6:6c:67:89:45:0c:3e:
30:98:38:35:c9:bd:b5:08:00:17:4c:e1:fb:38:50:
bd:97:f5:41:92:e7:6e:c4:6f:dc:ad:52:2c:e0:2a:
54:83:79:45:fb:5d:e2:f5:a5:cf:42:94:45:98:22:
d9:5b:81:93:e2:46:5f:e0:7f:71:5f:2d:b0:4a:82:
21:7d:f2:41:f7:b6:33:eb:59:93:f1:71:e3:79:ea:
c0:1b:5e:07:c6:d5:c2:67:41:56:73:d8:1f:a3:fb:
32:4b:f5:96:9f:65:f5:0a:f0:28:d5:90:d6:b0:dc:
4b:29:85:aa:8b:b7:d5:c0:f3:45:28:f9:af:80:7a:
88:40:40:21:60:ea:14:cd:8a:8e:53:40:67:c5:47:
51:bc:95:76:1e:90:b0:ee:ee:41:5a:ec:d4:4c:3c:
ea:eb:2f:f1:55:82:d8:b2:36:d9:92:88:bd:b6:93:
eb:46:69:3b:3a:e2:15:54:82:c0:30:4b:a9:54:3c:
af:52:4e:a5:71:40:a1:58:21:2e:ab:6d:c4:7c:59:
5d:68:b6:95:80:0e:12:91:51:90:0e:38:84:3f:de:
07:99:43:86:a1:0f:70:01:2f:3c:bf:e3:47:b2:16:
67:eb:00:6b:c4:7d:d8:e5:39:77:ac:29:cc:76:94:
2b:d3 Exponent: 65537 (0x10001) X509v3
extensions: X509v3 Basic Constraints:
CA:FALSE Netscape Cert Type: SSL
Client, S/MIME, Object Signing Netscape Comment:
OpenSSL Generated Certificate Netscape CA
Revocation Url: http://www.sdsc.edu/CA/
SDSC_CRL.pem X509v3 Subject Key Identifier:
E1:E3:C9:6E:A6:CF:2C:FC:D3:B7:51:F6:03:66:98:C5:18:71:60:F8
X509v3 Authority Key Identifier:
keyid:BF:A3:87:2C:F6:0D:74:BD:48:6C:0E:27:BF:01:E4:F2:4F:
46:BA:27 DirName:/C=US/O=SDSC/OU=SDSC-CA/
CN=Certificate Authority/UID=certman serial:
00 Signature Algorithm: md5WithRSAEncryption
93:b2:78:07:d9:72:e2:71:d7:66:83:0c:d3:97:0c:9e:24:33:
4e:e3:48:28:9c:44:7e:31:13:70:cc:f8:4a:5d:bc:64:84:3e:
aa:fa:da:86:3f:5e:f8:4a:72:a1:59:57:5a:89:49:5a:2d:c9:
09:5c:a5:69:6e:65:f7:85:8b:07:57:f1:6a:cb:6e:e5:00:17:
```



# Virtual Organization Management System

- User management delegated to VO (Virtual Organization)
  - flexible
  - scalable
- Introduction of groups and roles
- Dynamic calls

# VOMS Architecture



# Submitting a job with GRAM

- **globus-job-run** command

```
$ globus-job-run rookery.uchicago.edu /bin/hostname
```

- Run '/bin/hostname' on the resource rookery.uchicago.edu

- We don't care what LRM is used on 'rookery'. This command works with any LRM.



# The client can describe the job with GRAM's Resource Specification Language (RSL)

- Example:

```
& (executable = a.out)
  (directory = /home/nobody )
  (arguments = arg1 "arg 2")
```

- Submit with:

```
globusrun -f spec.rsl -r
rookery.uchicago.edu
```

# Use other programs to generate RSL

- RSL job descriptions can become very complicated
- We can use other programs to generate RSL for us
  - Example: Condor-G – next section

# Job management layers

- High level user tools  
(Panda, Swift, Pegasus, ...)
- OSG abstractions
- DAGman
- Condor
- Adding some Security
- GRAM, the Grid protocol
- Clusters and Local  
Resource Managers



# Condor

- is a software system that creates an HTC environment
  - ❑ Created at [UW-Madison](#)
- Condor is a specialized workload management system for compute-intensive jobs.
  - ❑ Detects machine availability
  - ❑ Harnesses available resources
  - ❑ Uses remote system calls to send R/W operations over the network
  - ❑ Provides powerful resource management by *matching* resource owners with consumers (broker)

# How Condor works

## Condor provides:

- a job queueing mechanism
- scheduling policy
- priority scheme
- resource monitoring, and
- resource management.

Users **submit** their serial or parallel jobs to Condor,

Condor places them into a **queue**,

... chooses **when** and **where** to run the jobs based upon a policy,

... carefully **monitors** their progress, and

... ultimately **informs** the user upon completion.

# Condor - features

- Checkpoint & migration
- Remote system calls
  - Able to transfer data files and executables across machines
- Job ordering
- *Job requirements and preferences can be specified via powerful expressions*

# Condor lets you manage a large number of jobs.

- Specify the jobs in a file and submit them to Condor
- Condor runs them and keeps you notified on their progress
  - ❑ Mechanisms to help you manage huge numbers of jobs (1000's), all the data, etc.
  - ❑ Handles inter-job dependencies (DAGMan)
- Users can set Condor's job priorities
- Condor administrators can set user priorities
- Can do this as:
  - ❑ Local resource manager (LRM) on a compute resource
  - ❑ Grid client submitting to GRAM (as Condor-G)

# Condor-G

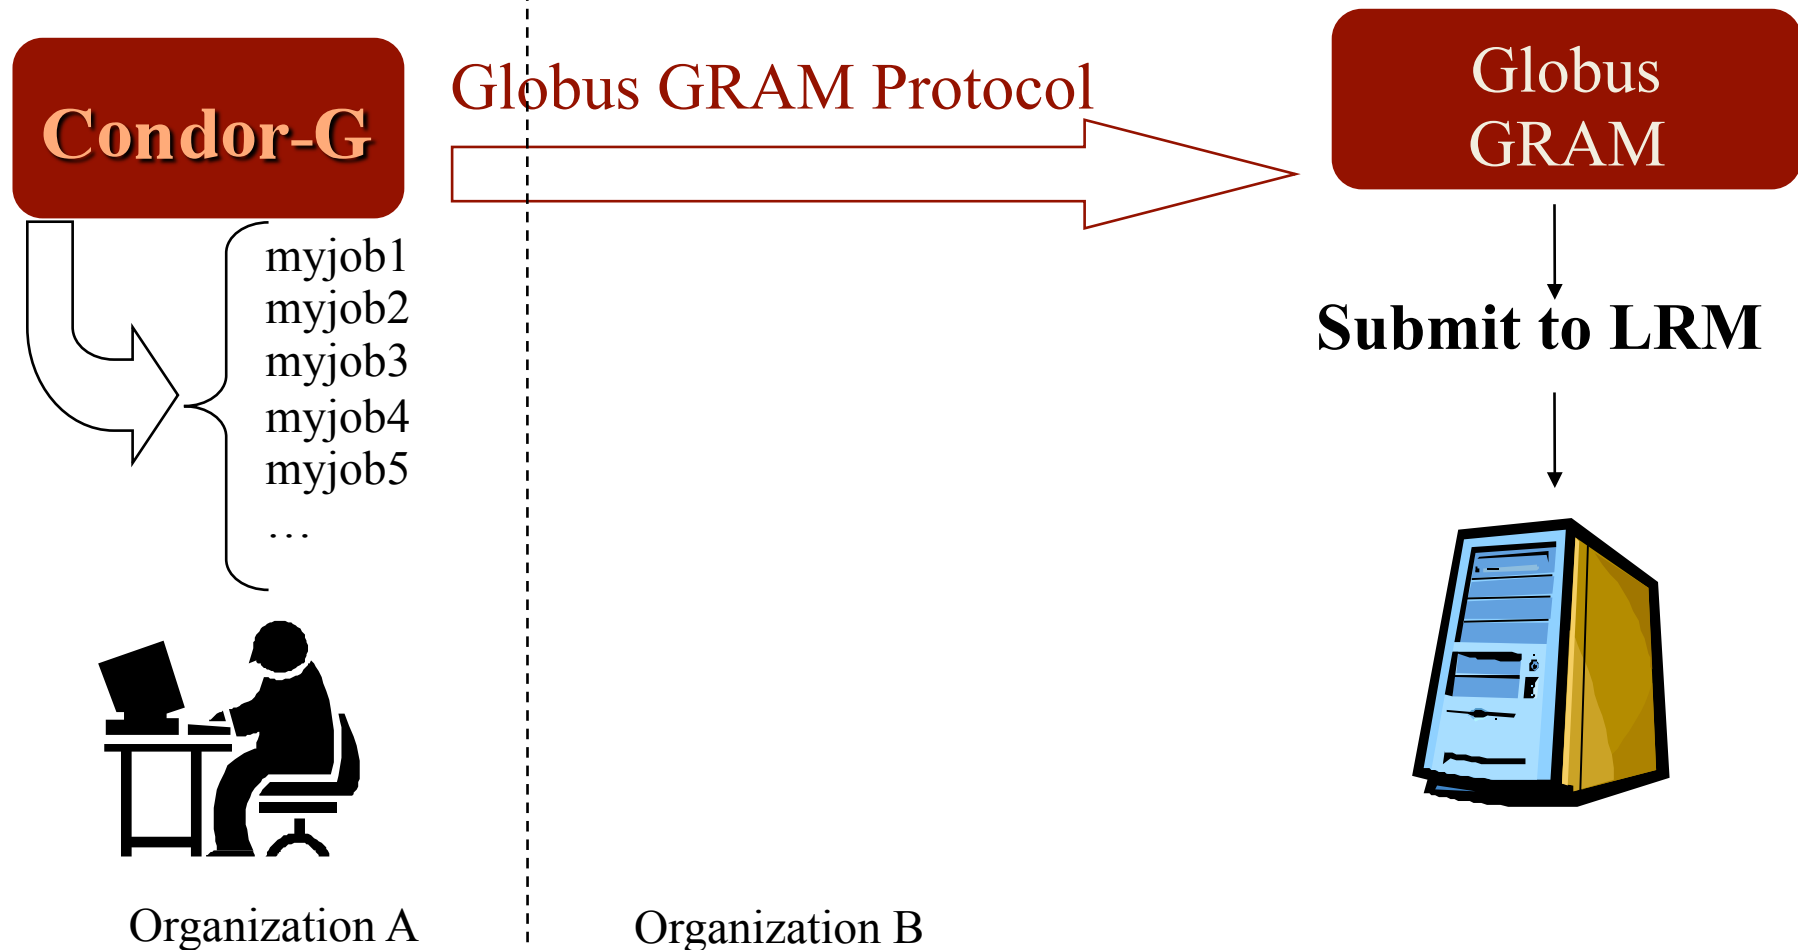
- is the job management part of Condor.
- *Hint:* Install Condor-G to submit to resources accessible through a Globus interface.
- Condor-G does not *create* a grid service.
- It only deals with *using* remote grid services.



# Condor-G ...

- does whatever it takes to run your jobs, even if ...
  - ❑ The gatekeeper is temporarily unavailable
  - ❑ The job manager crashes
  - ❑ Your local machine crashes
  - ❑ The network goes down

# Remote Resource Access: Condor-G + Globus + Condor



# Condor-G: Access non-Condor Grid resources



- middleware deployed across entire Grid
- remote access to computational resources
- dependable, robust data transfer



- job scheduling across multiple resources
- strong fault tolerance with checkpointing and migration
- layered over Globus as “personal batch system” for the Grid

# Four Steps to Run a Job with Condor

- These choices tell Condor
  - ❑ **how**
  - ❑ **when**
  - ❑ **where** to run the job,
  - ❑ and describe exactly **what** you want to run.
- Choose a Universe for your job
- Make your job batch-ready
- Create a *submit description* file
- Run *condor\_submit*

# I. Choose a Universe

- There are many choices
  - ❑ **Vanilla:** any old job
  - ❑ **Standard:** checkpointing & remote I/O
  - ❑ **Java:** better for Java jobs
  - ❑ **MPI:** Run parallel MPI jobs
  - ❑ Virtual Machine: Run a virtual machine as job
  - ❑ ...
- *For now, we'll just consider vanilla*

## 2. Make your job batch-ready

- Must be able to run in the background:
  - no interactive input, windows, GUI, etc.
- Condor is designed to run jobs as a batch system, with pre-defined inputs for jobs
- Can still use STDIN, STDOUT, and STDERR (the keyboard and the screen), but files are used for these instead of the actual devices
- Organize data files

# 3. Create a Submit Description File

- A plain ASCII text file
- Condor does not care about file extensions
- Tells Condor about your job:
  - ❑ Which executable to run and where to find it
  - ❑ Which universe
  - ❑ Location of input, output and error files
  - ❑ Command-line arguments, if any
  - ❑ Environment variables
  - ❑ Any special requirements or preferences

# Simple Submit Description File

```
# myjob.submit file
# Simple condor_submit input file
# (Lines beginning with # are comments)
# NOTE: the words on the left side are not
#       case sensitive, but filenames are!
Universe      = vanilla
Executable    = analysis
Log            = my_job.log
Queue
```



## 4. Run `condor_submit`

- You give *condor\_submit* the name of the submit file you have created:

```
condor_submit my_job.submit
```

- *condor\_submit* parses the submit file

# Another Submit Description File

```
# Example condor_submit input file
# (Lines beginning with # are comments)
# NOTE: the words on the left side are not
#       case sensitive, but filenames are!
Universe      = vanilla
Executable    = /home/wright/condor/my_job.condor
Input         = my_job.stdin
Output        = my_job.stdout
Error         = my_job.stderr
Arguments     = -arg1 -arg2
InitialDir    = /home/wright/condor/run_1
Queue
```

---

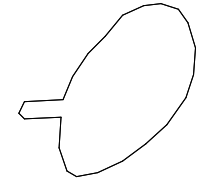
# Details

- Lots of options available in the submit file
- Commands to
  - ❑ watch the queue,
  - ❑ the state of your pool,
  - ❑ and lots more
- You'll see much of this in the hands-on exercises.

# Other Condor commands

- `condor_q` – show status of job queue
- `condor_status` – show status of compute nodes
- `condor_rm` – remove a job
- `condor_hold` – hold a job temporarily
- `condor_release` – release a job from hold

# Submitting more complex jobs



- express dependencies between jobs  
⇒ WORKFLOWS
- And also, we would like the workflow to be managed even in the face of failures

---

# Want other Scheduling possibilities?

## Use the Scheduler Universe

- In addition to VANILLA, another job universe is the *Scheduler Universe*.
- Scheduler Universe jobs run on the submitting machine and serve as a meta-scheduler.
- **Condor's Scheduler Universe lets you set up and manage job workflows.**
- DAGMan meta-scheduler included
  - DAGMan manages these jobs

# Job management layers

- High level user tools  
(Panda, Swift, Pegasus, ...)
- OSG abstractions
- **DAGman**
- Condor
- Adding some Security
- GRAM, the Grid protocol
- Clusters and Local  
Resource Managers



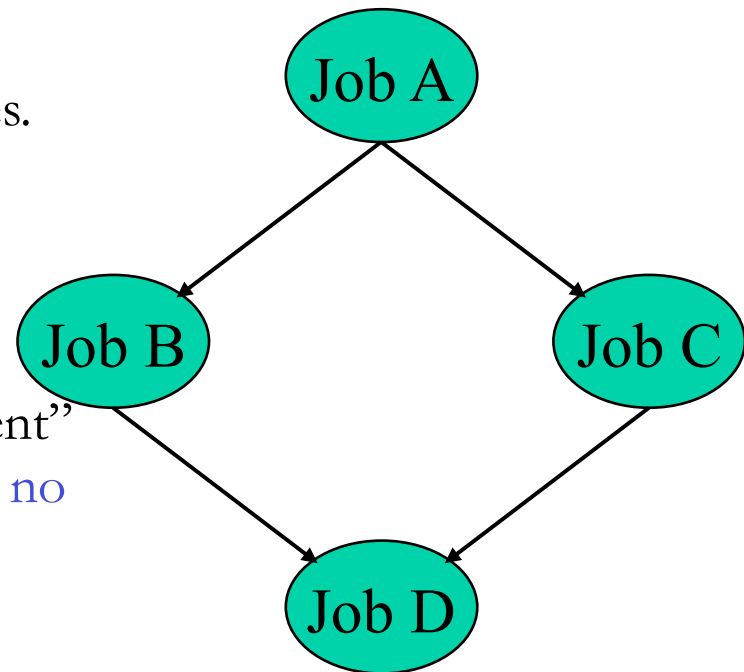
# DAGMan

- **Directed Acyclic Graph Manager**
- DAGMan allows you to specify the *dependencies* between your Condor jobs, so it can *manage* them automatically for you.
- (e.g., “Don’t run job “B” until job “A” has completed successfully.”)



# What is a DAG?

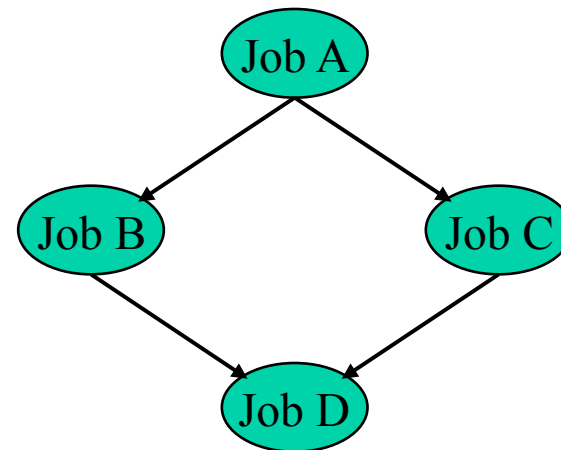
- A DAG is the **data structure** used by DAGMan to represent these dependencies.
- Each job is a **“node”** in the DAG.
- Each node can have any number of “parent” or “children” nodes – as long as there are **no loops**!



# Defining a DAG

- A DAG is defined by a *.dag file*, listing each of its nodes and their dependencies:

```
# diamond.dag
Job A a.sub
Job B b.sub
Job C c.sub
Job D d.sub
Parent A Child B C
Parent B C Child D
```



- each node will run the Condor job specified by its accompanying *Condor submit file*

# Submitting a DAG

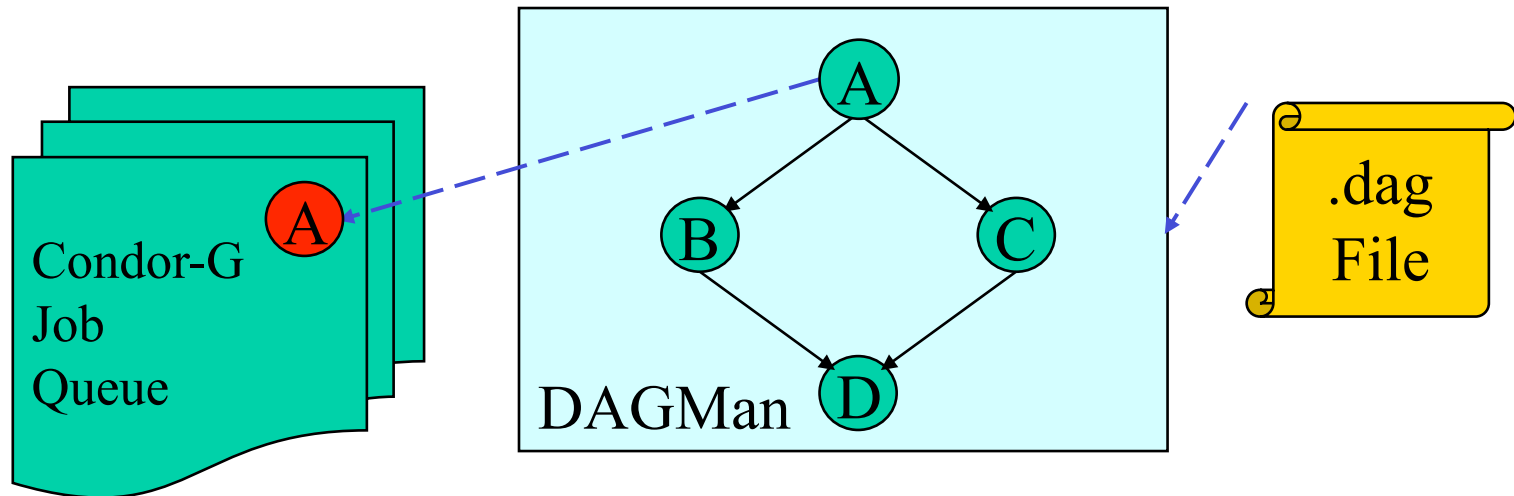
- To start your DAG, just run ***condor\_submit\_dag*** with your .dag file, and Condor will start a personal DAGMan daemon which to begin running your jobs:

```
% condor_submit_dag diamond.dag
```

- `condor_submit_dag` submits a Scheduler Universe Job with DAGMan as the executable.
- Thus the DAGMan daemon itself **runs as a Condor job**, so you don't have to baby-sit it.

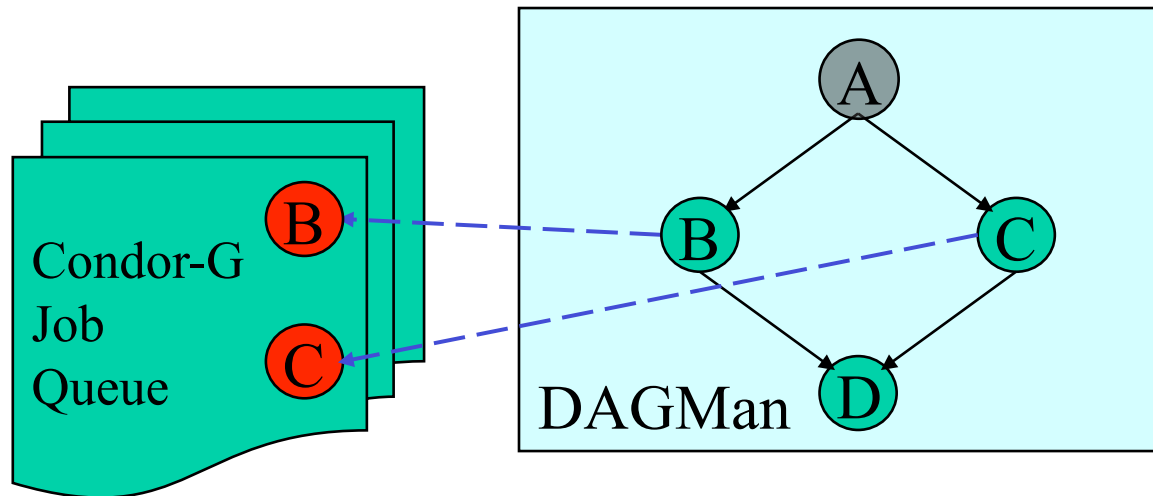
# Running a DAG

- DAGMan acts as a “meta-scheduler”, managing the submission of your jobs to Condor-G based on the DAG dependencies.



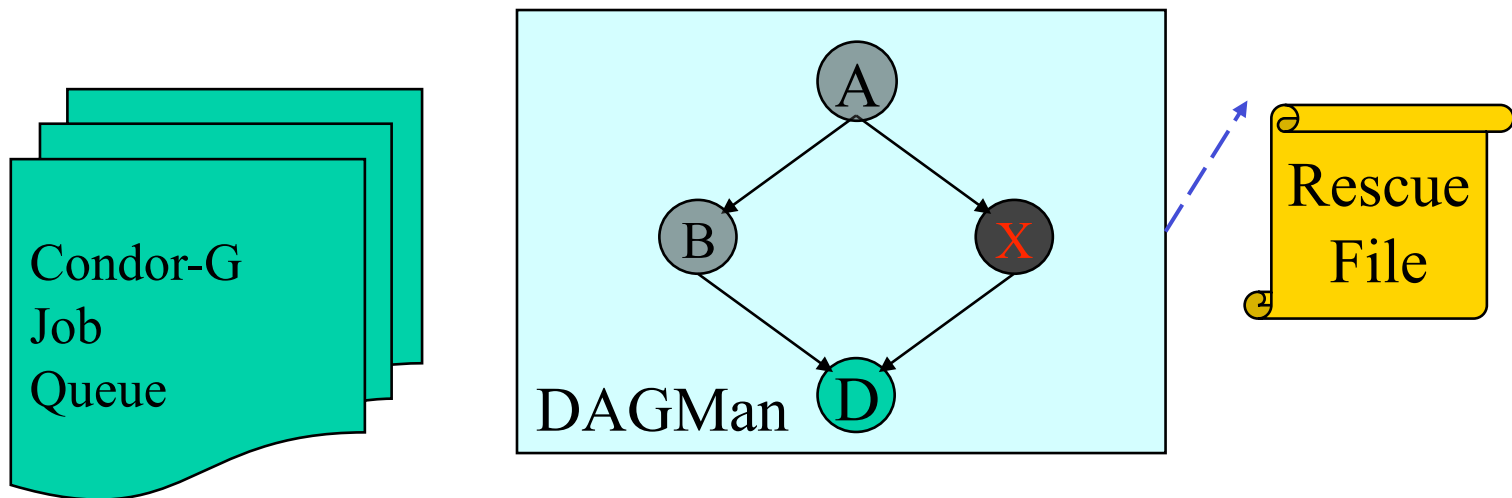
# Running a DAG (cont'd)

- DAGMan holds & submits jobs to the Condor-G queue at the appropriate times.



# Running a DAG (cont'd)

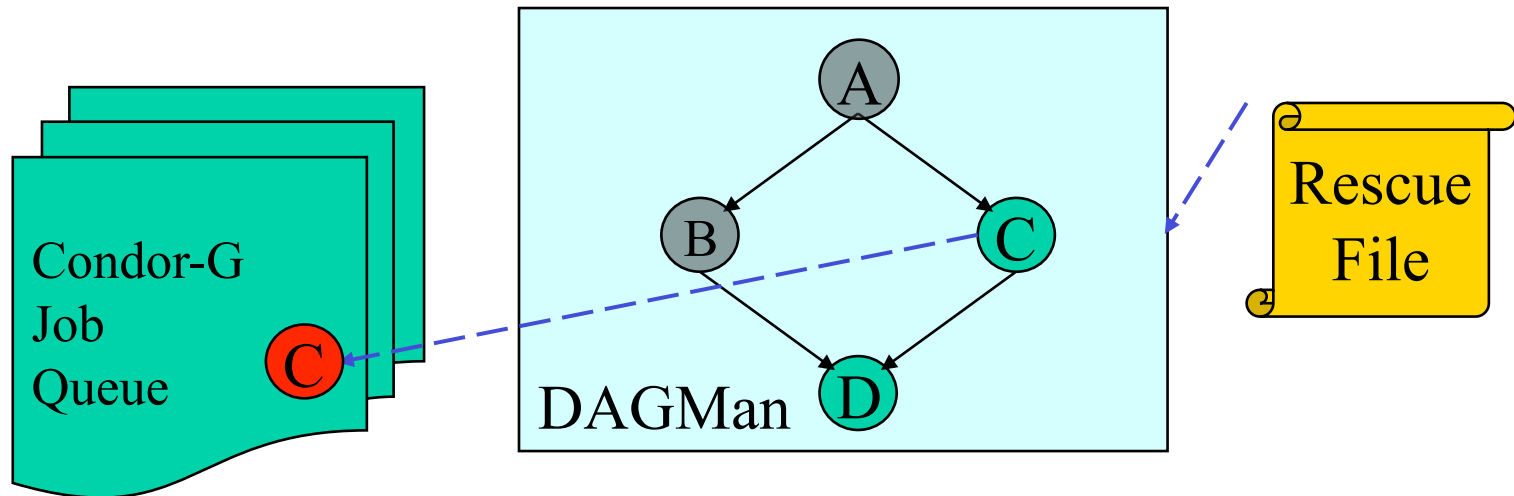
- In case of a job failure, DAGMan continues until it can no longer make progress, and then creates a *“rescue” file* with the current state of the DAG.



# Recovering a DAG

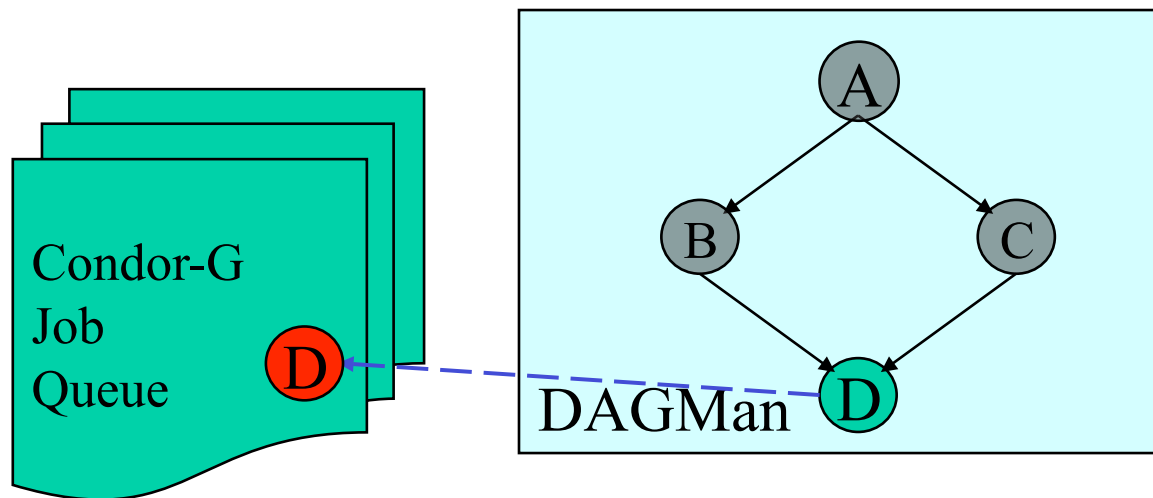
## -- fault tolerance

- Once the failed job is ready to be re-run, the rescue file can be used to restore the prior state of the DAG.



# Recovering a DAG (cont'd)

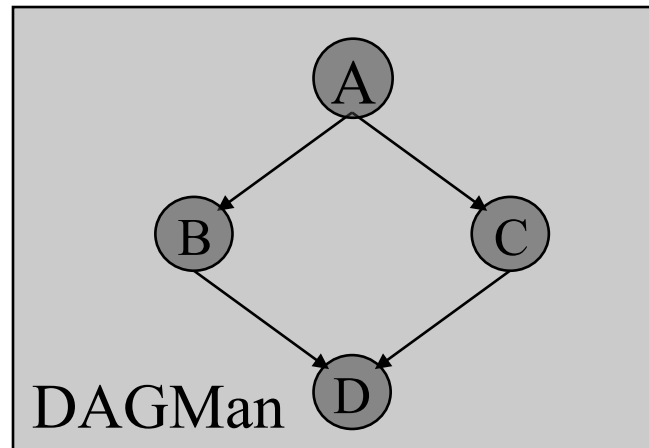
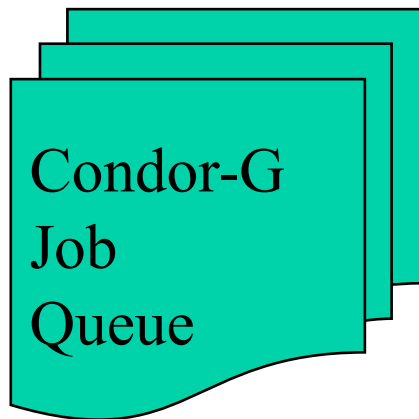
- Once that job completes, DAGMan will continue the DAG as if the failure never happened.





# Finishing a DAG

- Once the DAG is complete, the DAGMan job itself is finished, and exits.

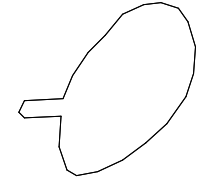


## We have seen how Condor:

- ... monitors submitted jobs and reports progress
- ... implements your policy on the execution order of the jobs
- ... keeps a log of your job activities

# Long jobs: if my jobs run for weeks

...



- What happens to my job when
  - ❑ a machine is shut down
  - ❑ there is a network outage, or
  - ❑ another job with higher priority preempts it?
- Do I lose all of those hours or days of computation time??
- What happens when they get pre-empted?
- *How can I add fault tolerance to my jobs?*

# Condor's Standard Universe to the rescue!

- Condor can support various combinations of features/environments in different “Universes”
- Different Universes provide different functionalities to your job:
  - ❑ Vanilla: Run any serial job
  - ❑ Scheduler: Plug in a scheduler
  - ❑ Standard: Support for *transparent process checkpoint and restart*

**provides two important services to your job:  
process checkpoint  
remote system calls.**

# Process Checkpointing

- Condor's process checkpointing mechanism *saves the entire state of a process into a checkpoint file*
  - Memory, CPU, I/O, etc.
- The process can then be *restarted* from the point it left off
- Typically no changes to your job's source code needed—however, your job must be relinked with Condor's Standard Universe support library

# Job management layers

- High level user tools  
(Panda, Swift, Pegasus, ...)
- OSG abstractions
- DAGman
- Condor
- Adding some Security
- GRAM, the Grid protocol
- Clusters and Local  
Resource Managers



# OSG & job submissions

- OSG sites present interfaces allowing remotely submitted jobs to be accepted, queued and executed locally.
- OSG supports the Condor-G job submission client which interfaces to either the pre-web service or web services GRAM Globus interface at the executing site.
- Job managers at the backend of the GRAM gatekeeper support job execution by local Condor, LSF, PBS, or SGE batch systems.

# OSG and Environment

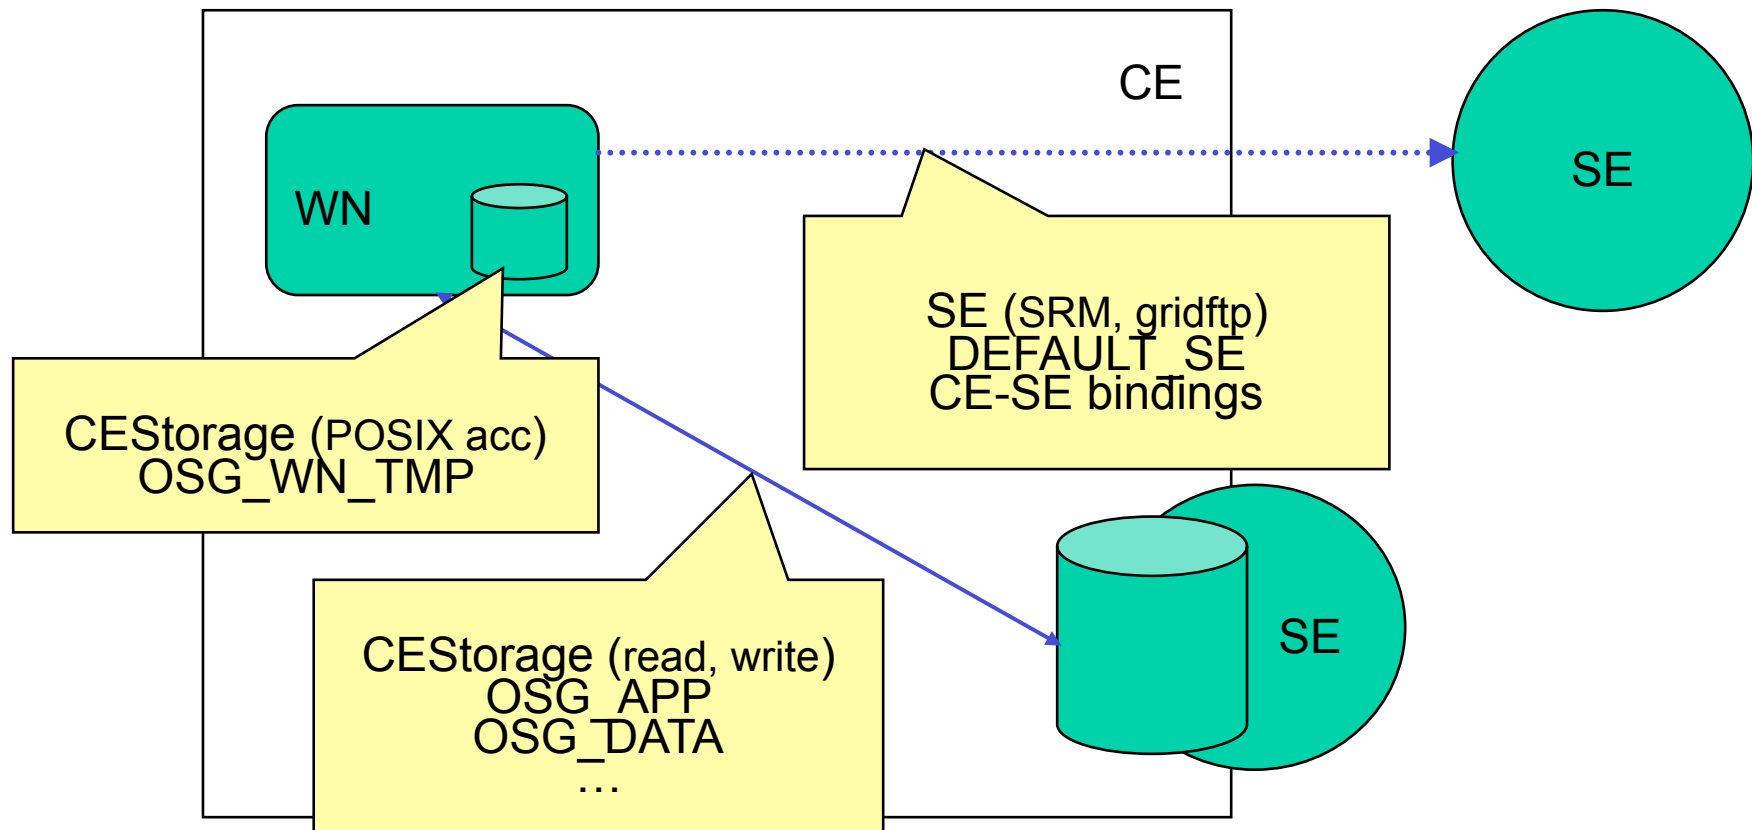
- Information systems to discover and describe resources
- Guidelines on how to execute jobs
- Guidelines on how to use available disk spaces
- Jobs find an uniform standard environment
  - OSG\_AAA variables



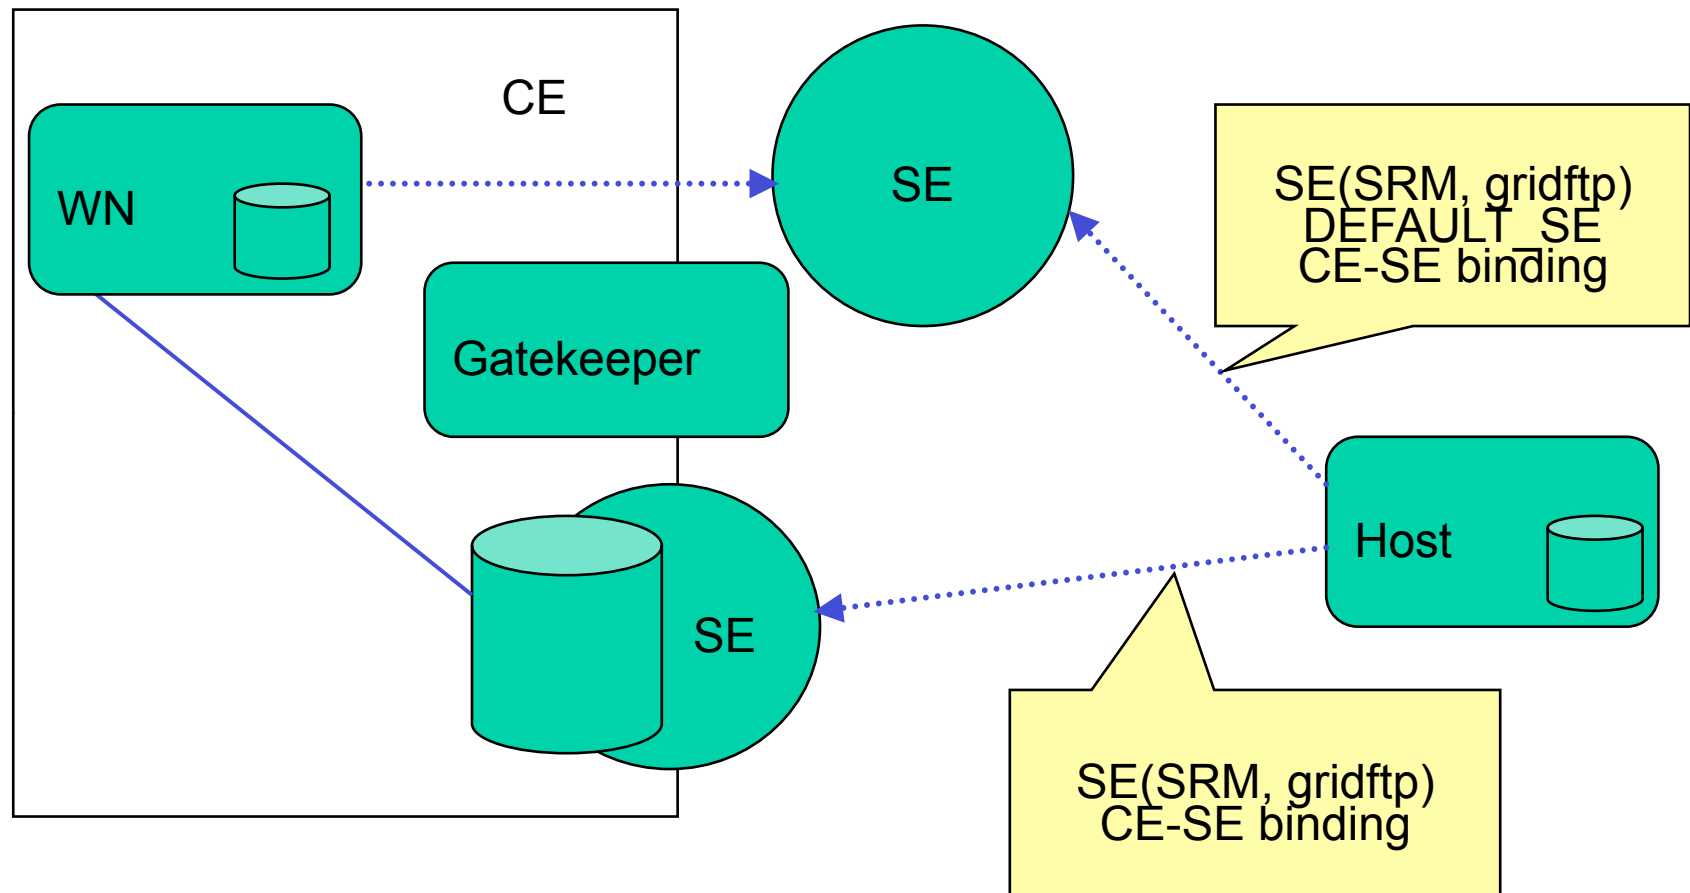
# Everything has a place

- OSG\_APP
  - ❑ directory to install applications
- OSG\_GRID
  - ❑ directory to find grid clients (OSG:WN)
- OSG\_WN\_TMP
  - ❑ local working directory
- OSG\_DATA
  - ❑ shared directory
- OSG\_SITE\_READ, OSG\_SITE\_WRITE
  - ❑ optimized for efficient read/write
- DEFAULT\_SE
  - ❑ storage element close to this computing element

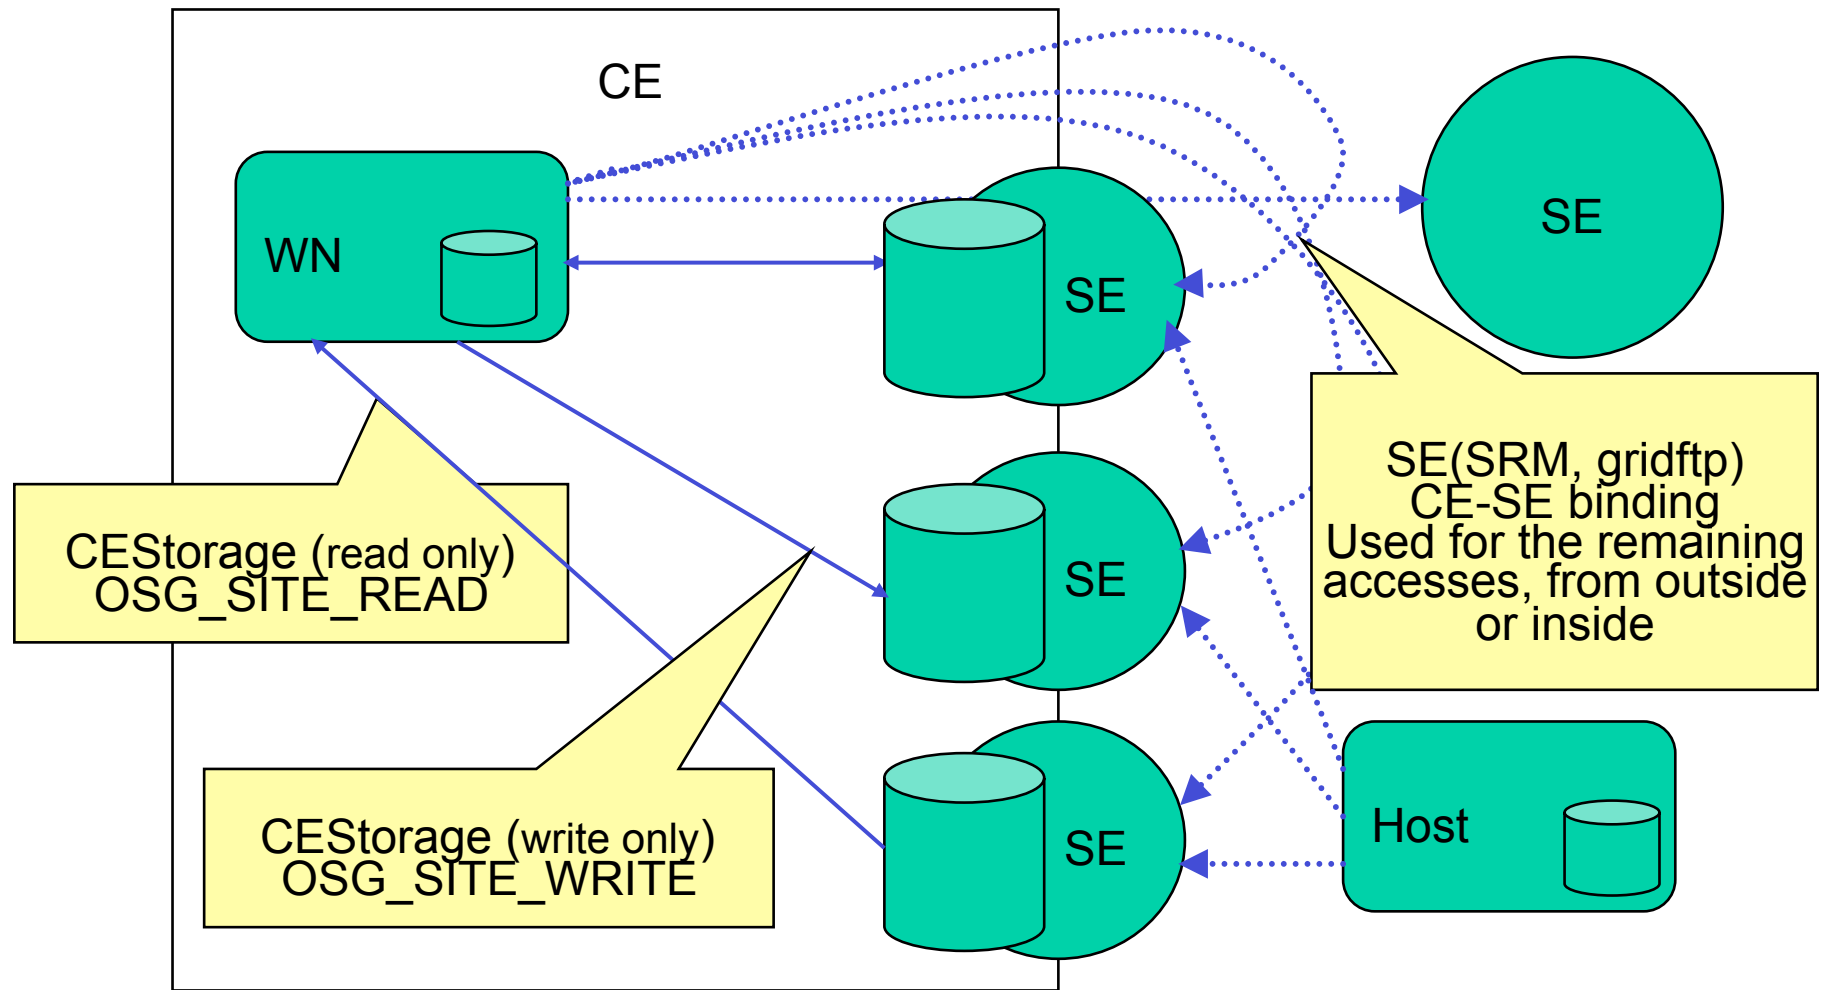
# CE View



# Outside View



# CE View (OSG\_WRITE, OSG\_READ)





## Acknowledgments:

---

This presentation based on:  
Grid Resources and Job Management



**Open Science Grid**

Jaime Frey and Becky  
Gietzel

Condor Project

U. Wisconsin-Madison

---

# Job management layers

- High level user tools  
(Panda, Swift, Pegasus, ...)
- OSG abstractions
- DAGman
- Condor
- Adding some Security
- GRAM, the Grid protocol
- Clusters and Local  
Resource Managers

