

---

# Grid Compute Resources and Job Management

---



July 11-15, 2005

Lecture3: Grid Job Management

1

---

## Job and compute resource management

- This module is about running jobs on remote compute resources

---

July 11-15, 2005

Lecture3: Grid Job Management

2

---

# Job and resource management


- Compute resources have a local resource manager
  - This controls who is allowed to run jobs and how they run, on a resource
- GRAM
  - Helps us run a job on a remote resource
- Condor
  - Manages jobs

---

# Local Resource Managers

- Local Resource Managers (LRMs) – software on a compute resource such a multi-node cluster.
- Control which jobs run, when they run and on which processor they run
- Example policies:
  - Each cluster node can run one job. If there are more jobs, then the other jobs must wait in a queue
  - Reservations – maybe some nodes in cluster reserved for a specific person
- eg. PBS, LSF, Condor

# Job Management on a Grid



## GRAM

- *Globus Resource Allocation Manager*
- Provides a standardised interface to submit jobs to different types of LRM
- Clients submit a job request to GRAM
- GRAM translates into something the LRM can understand
- Same job request can be used for many different kinds of LRM

---

# GRAM

- Given a job specification:
    - Create an environment for a job
    - Stage files to and from the environment
    - Submit a job to a local resource manager
    - Monitor a job
    - Send notifications of the job state change
    - Stream a job's stdout/err during execution
- 

---

## Two versions of GRAM

- There are two versions of GRAM
    - GRAM2
      - Own protocols
      - Older
      - More widely used
      - No longer actively developed
    - GRAM4
      - Web services
      - Newer
      - New features go into GRAM4
  
  - In this module, will be using GRAM2
-

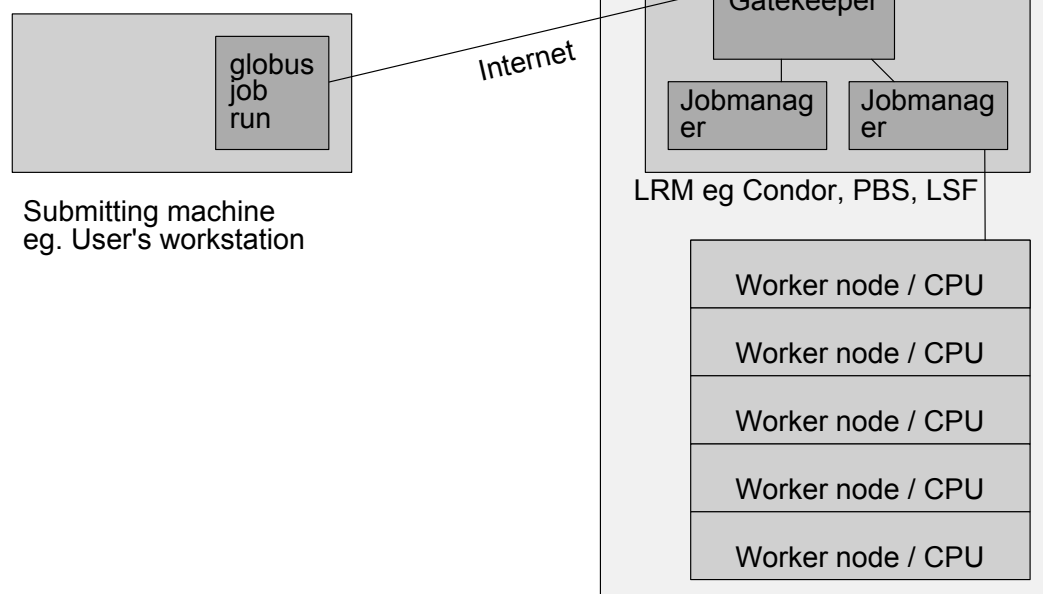
---

## GRAM components

- Clients – eg. Globus-job-submit, globusrun
  - Gatekeeper
    - Server
    - Accepts job submissions
    - Handles security
  - Jobmanager
    - Knows how to send a job into the local resource manager
    - Different job managers for different LRMs
- 

---

## GRAM components



---

## Submitting a job with GRAM

- Globus-job-run command
  - `globus-job-run rookery.uchicago.edu /bin/hostname rook11`
- Run '/bin/hostname' on the resource `rookery.uchicago.edu`
- We don't care what LRM is used on 'rookery'. This command works with any LRM.

---

The client can describe the job with GRAM's Resource Specification Language (RSL)

- Example:

```
&(executable = a.out)
(directory = /home/nobody )
(arguments = arg1 "arg 2")
```
  - Submit with:

```
globusrun -f spec.rsl -r
rookery.uchicago.edu
```
-

---

## Use other programs to generate RSL

- RSL job descriptions can become very complicated
  - We can use other programs to generate RSL for us
  - Example: Condor-G – next section
- 

---

## Condor

- Globus-job-run submits jobs, but...
    - No job tracking: what happens when something goes wrong?
  - Condor:
    - Many features, but in this module:
    - Condor-G for reliable job management
-

---

# Condor can manage a large number of jobs

- Managing a large number of jobs
  - You specify the jobs in a file and submit them to Condor, which runs them all and keeps you notified on their progress
  - Mechanisms to help you manage huge numbers of jobs (1000's), all the data, etc.
  - Condor can handle inter-job dependencies (DAGMan)
  - Condor users can set job priorities
  - Condor administrators can set user priorities
- Can do this as:
  - a local resource manager on a compute resource
  - a grid client submitting to GRAM (Condor-G)

---

# Condor can manage compute resource

- Dedicated Resources
  - Compute Clusters
- Non-dedicated Resources
  - Desktop workstations in offices and labs
    - Often idle 70% of time
- Condor acts as a Local Resource Manager





---

## ... and Condor Can Manage Grid jobs

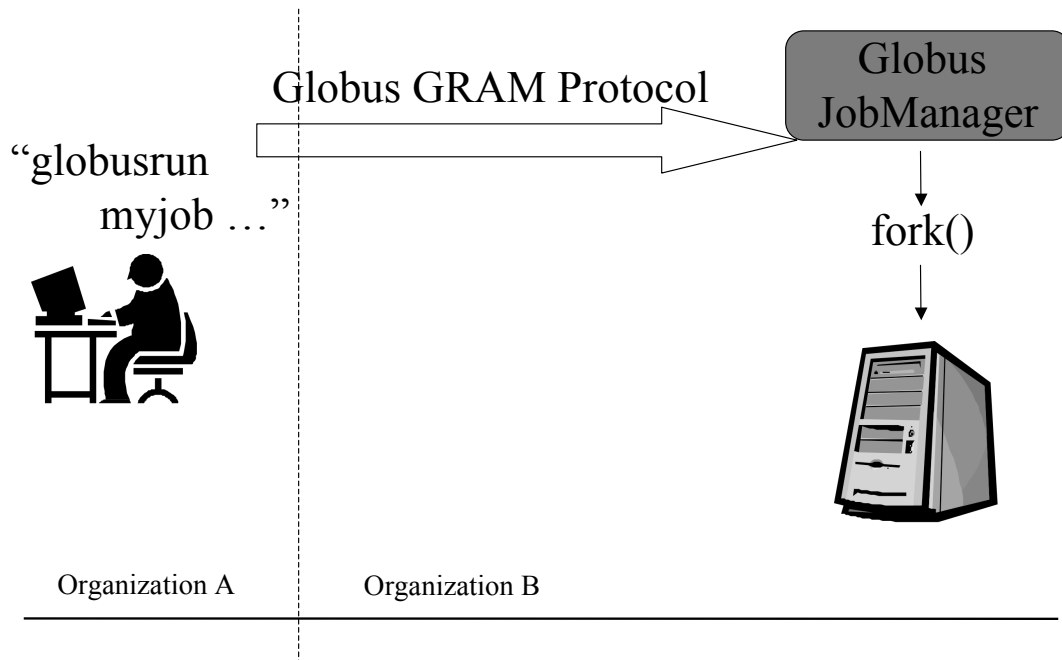
- Condor-G is a specialization of Condor. It is also known as the “Grid universe”.
- Condor-G can submit jobs to Globus resources, just like globus-job-run.
- Condor-G benefits from Condor features, like a job queue.

---

## Some Grid Challenges

- Condor-G does whatever it takes to run your jobs, even if ...
  - The gatekeeper is temporarily unavailable
  - The job manager crashes
  - Your local machine crashes
  - The network goes down

## Remote Resource Access: Globus

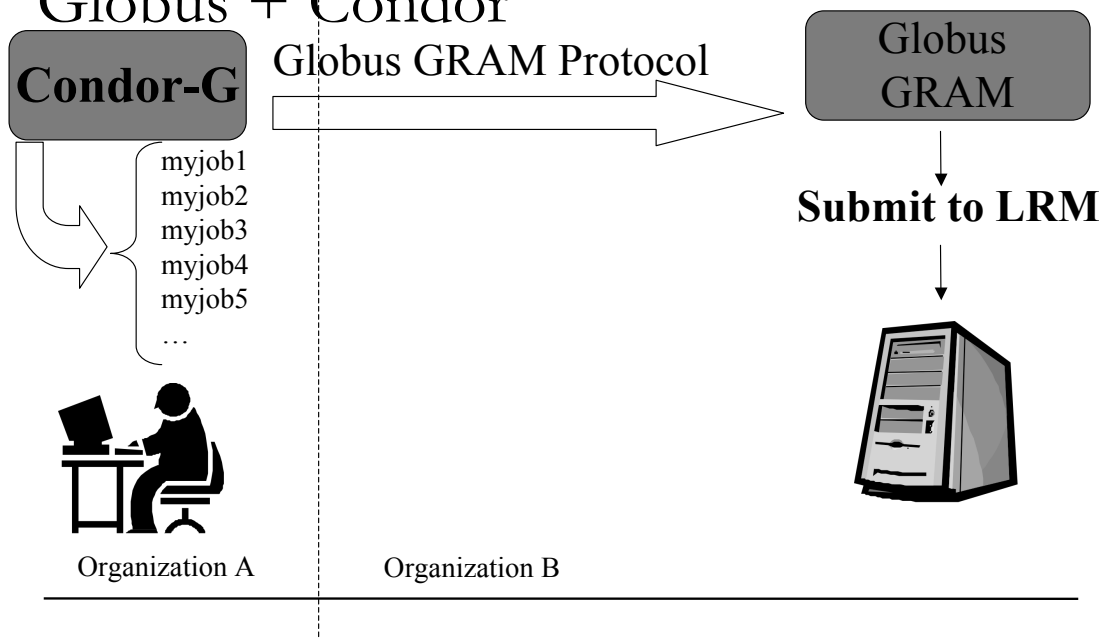


July 11-15, 2005

Lecture3: Grid Job Management

19

## Remote Resource Access: Condor-G + Globus + Condor



July 11-15, 2005

Lecture3: Grid Job Management

20

---

## Example Application ...

Simulate the behavior of  $F(x,y,z)$  for 20 values of  $x$ , 10 values of  $y$  and 3 values of  $z$  ( $20*10*3 = 600$  combinations)

- ❑  $F$  takes on the average 3 hours to compute on a “typical” workstation (total = 1800 hours)
- ❑  $F$  requires a “moderate” (128MB) amount of memory
- ❑  $F$  performs “moderate” I/O -  $(x,y,z)$  is 5 MB and  $F(x,y,z)$  is 50 MB
  
- ❑ 600 jobs

---

## Creating a Submit Description File

- A plain ASCII text file
- Tells Condor about your job:
  - ❑ Which executable, universe, input, output and error files to use, command-line arguments, environment variables, any special requirements or preferences (more on this later)
- Can describe many jobs at once (a “cluster”) each with different input, arguments, output, etc.

---

# Simple Submit Description File

```
# Simple condor_submit input file
# (Lines beginning with # are comments)
# NOTE: the words on the left side are not
#       case sensitive, but filenames are!
Universe    = vanilla
Executable = my_job
Queue
```

```
$ condor_submit myjob.sub
```

---

---

## Other Condor commands

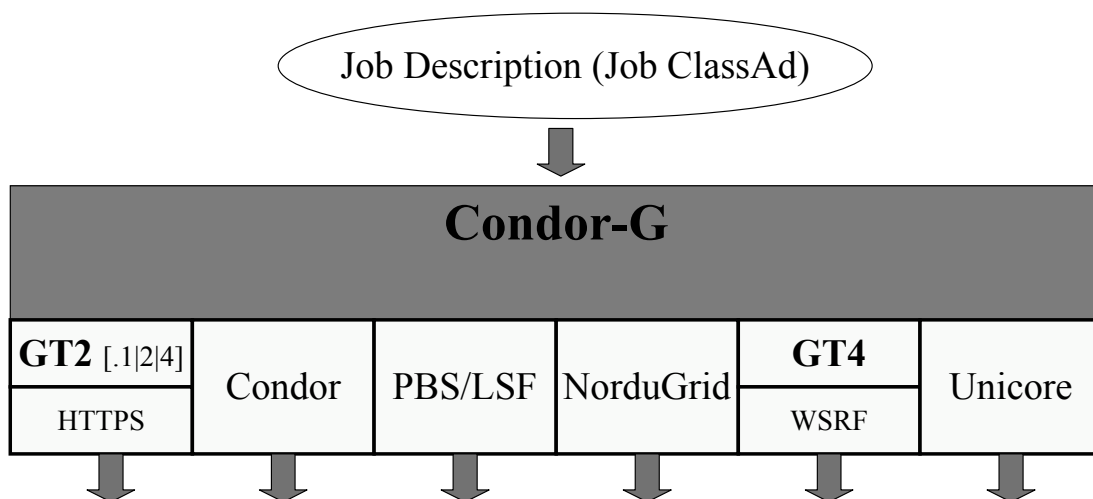
- `condor_q` – show status of job queue
- `condor_status` – show status of compute nodes
- `condor_rm` – remove a job
- `condor_hold` – hold a job temporarily
- `condor_release` – release a job from hold

# Condor-G: Access non-Condor Grid resources



- middleware deployed across entire Grid
- remote access to computational resources
- dependable, robust data transfer
- job scheduling across multiple resources
- strong fault tolerance with checkpointing and migration
- layered over Globus as “personal batch system” for the Grid

## Condor-G



---

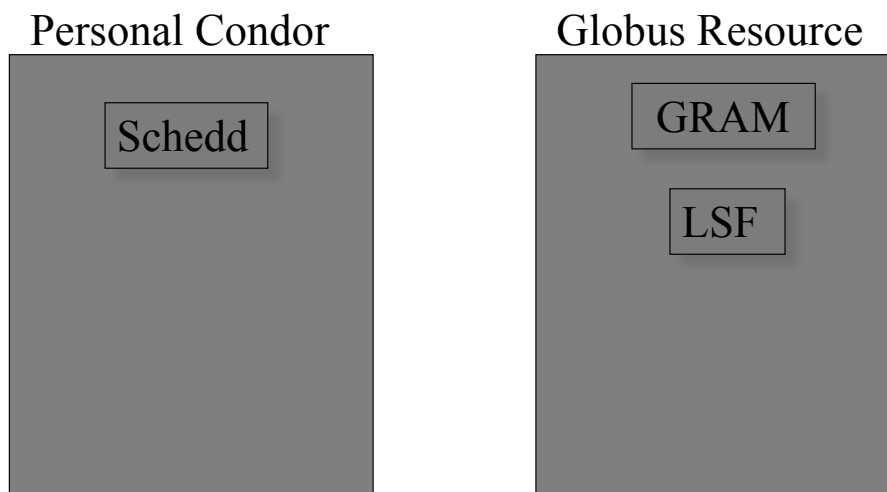
## Submitting a GRAM Job

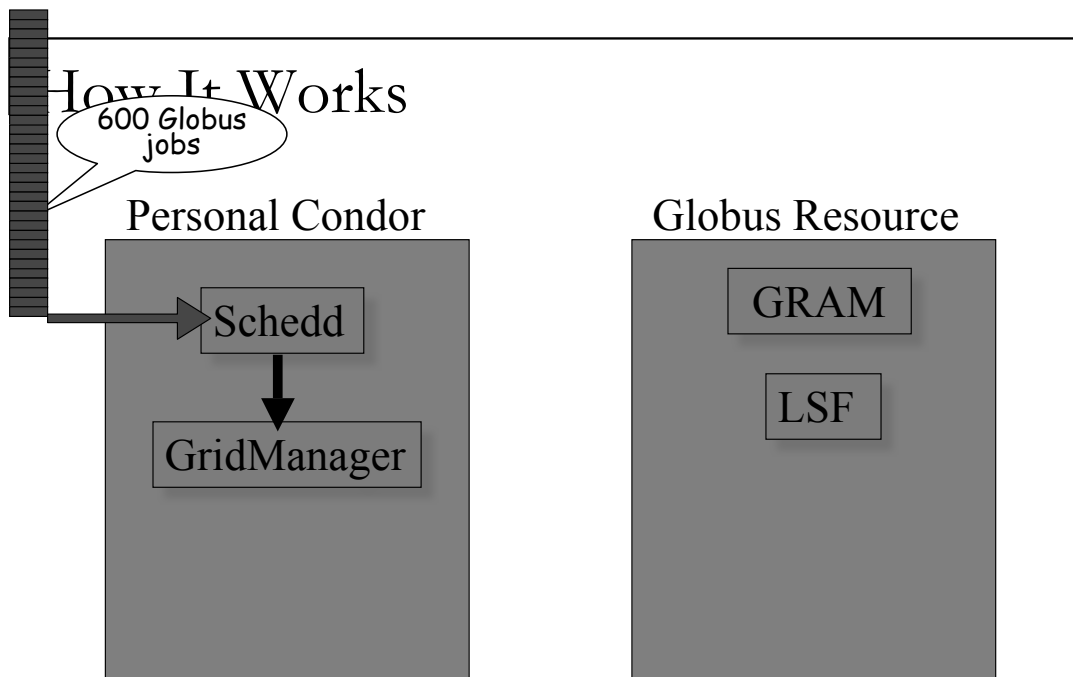
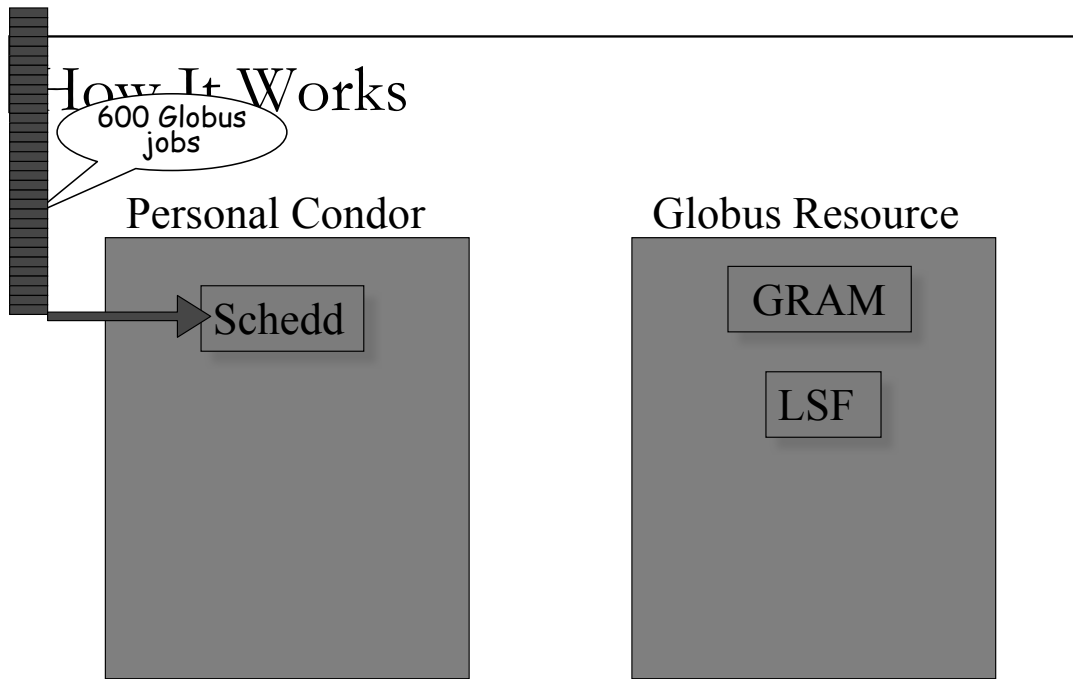
- In submit description file, specify:
  - Universe = grid
  - Grid\_Resource = gt2 <gatekeeper host>
    - 'gt2' means GRAM2
  - Optional: Location of file containing your X509 proxy

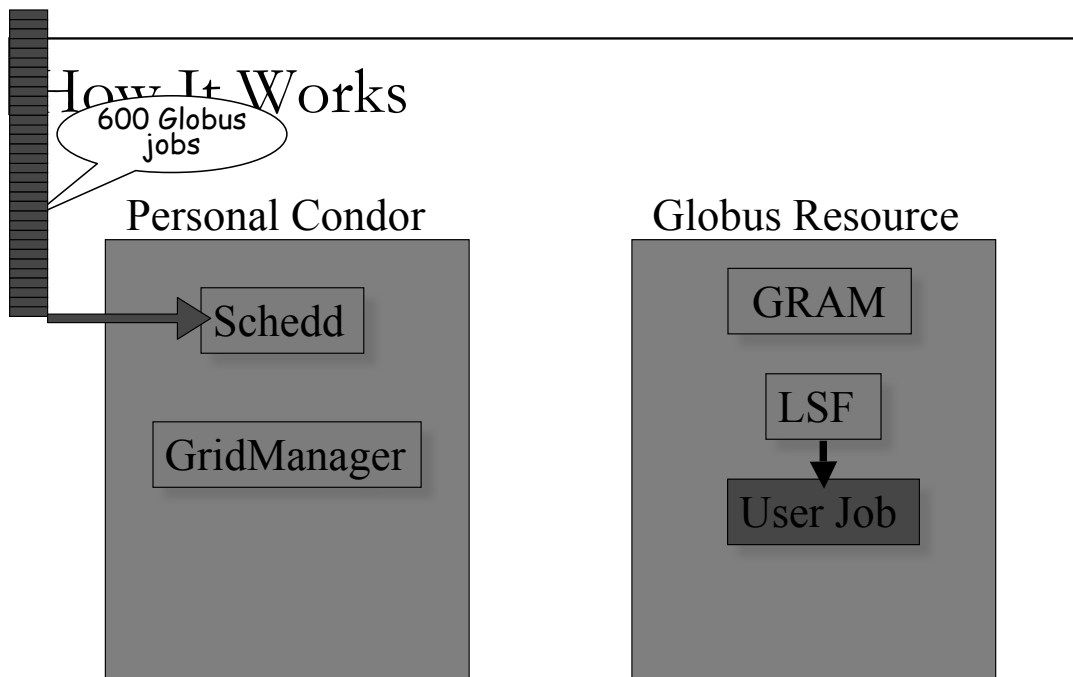
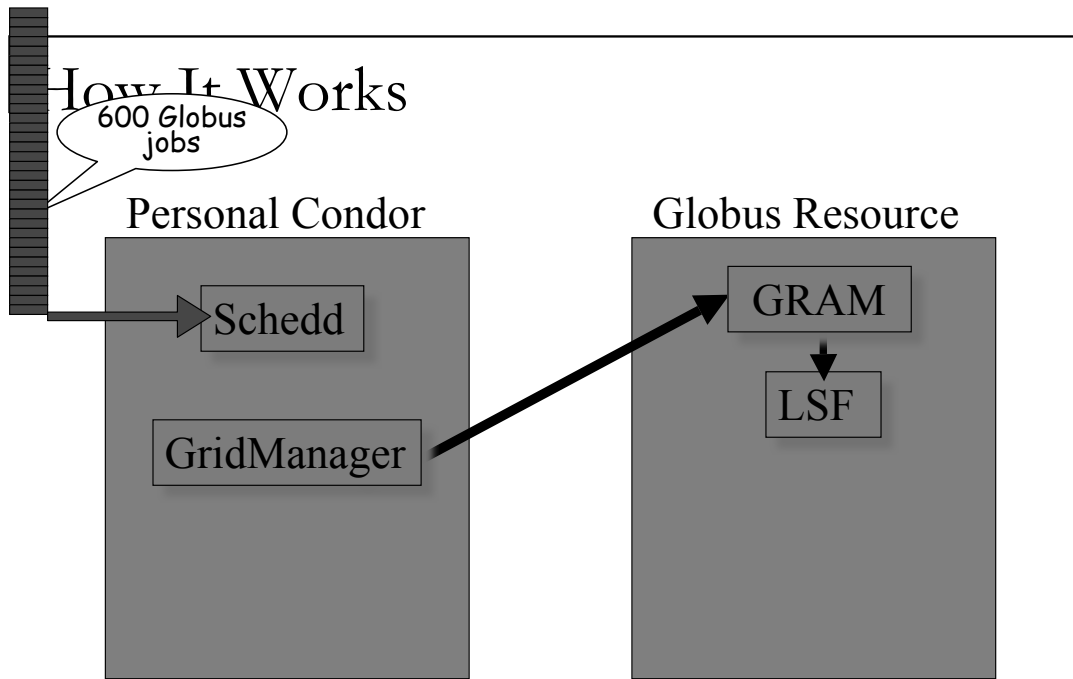
```
universe      = grid
grid_resource = gt2 beak.cs.wisc.edu/jobmanager-pbs
executable   = progname
queue
```

---

## How It Works









---

# Grid Universe Concerns

- What about Fault Tolerance?
    - Local Crashes
      - What if the submit machine goes down?
    - Network Outages
      - What if the connection to the remote Globus jobmanager is lost?
    - Remote Crashes
      - What if the remote Globus jobmanager crashes?
      - What if the remote machine goes down?
  - Condor-G's persistent job queue lets it recover from all of these failures
  - If a JobManager fails to respond...
- 

---

## This presentation based on: Grid Resources and Job Management

---

Jaime Frey  
Condor Project,  
University of Wisconsin-Madison  
[jfrey@cs.wisc.edu](mailto:jfrey@cs.wisc.edu)



Open Science Grid



Grid Summer Workshop  
June 26-30, 2006