

# Grid Data Management

---



Open Science Grid

---

# Data Management

- Distributed community of users need to access and analyze large amounts of data



- Requirement arises in both simulation and experimental science

# Data Management

- Huge raw volume of data
  - ❑ Measured in terabytes, petabytes, and further ...
  - ❑ Data sets can be partitioned as small number of large files or large number of small files
  - ❑ Store it long term in appropriate places (e.g., tape silos)
  - ❑ Move input to where your job is running
  - ❑ Move output data from where your job ran to where you need it (eg. your workstation, long term storage)

# Data Management on the Grid

- How to move data/files to where I want?
  - GridFTP
- Data sets replicated for reliability and faster access
- Files have logical names
- Service that maps logical file names to physical locations
  - Replica Location Service (RLS)
  - Where are the files I want?

# GridFTP

- High performance, secure, and reliable data transfer protocol based on the standard FTP
  - ❑ <http://www.ogf.org/documents/GFD.20.pdf>
  - ❑ Multiple implementations exist, we'll focus on Globus GridFTP
- Globus GridFTP Features include
  - ❑ Strong authentication, encryption via Globus GSI
  - ❑ Multiple transport protocols - TCP, UDT
  - ❑ Parallel transport streams for faster transfer
  - ❑ Cluster-to-cluster or striped data movement
  - ❑ Multicasting and overlay routing
  - ❑ Support for reliable and restartable transfers

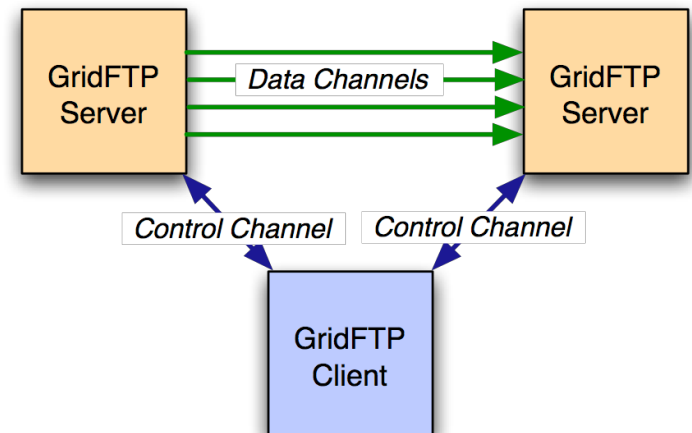

# Basic Definitions

## ■ Control Channel

- ❑ TCP link over which **commands** and **responses** flow
- ❑ Low bandwidth; encrypted and integrity protected by default

## ■ Data Channel

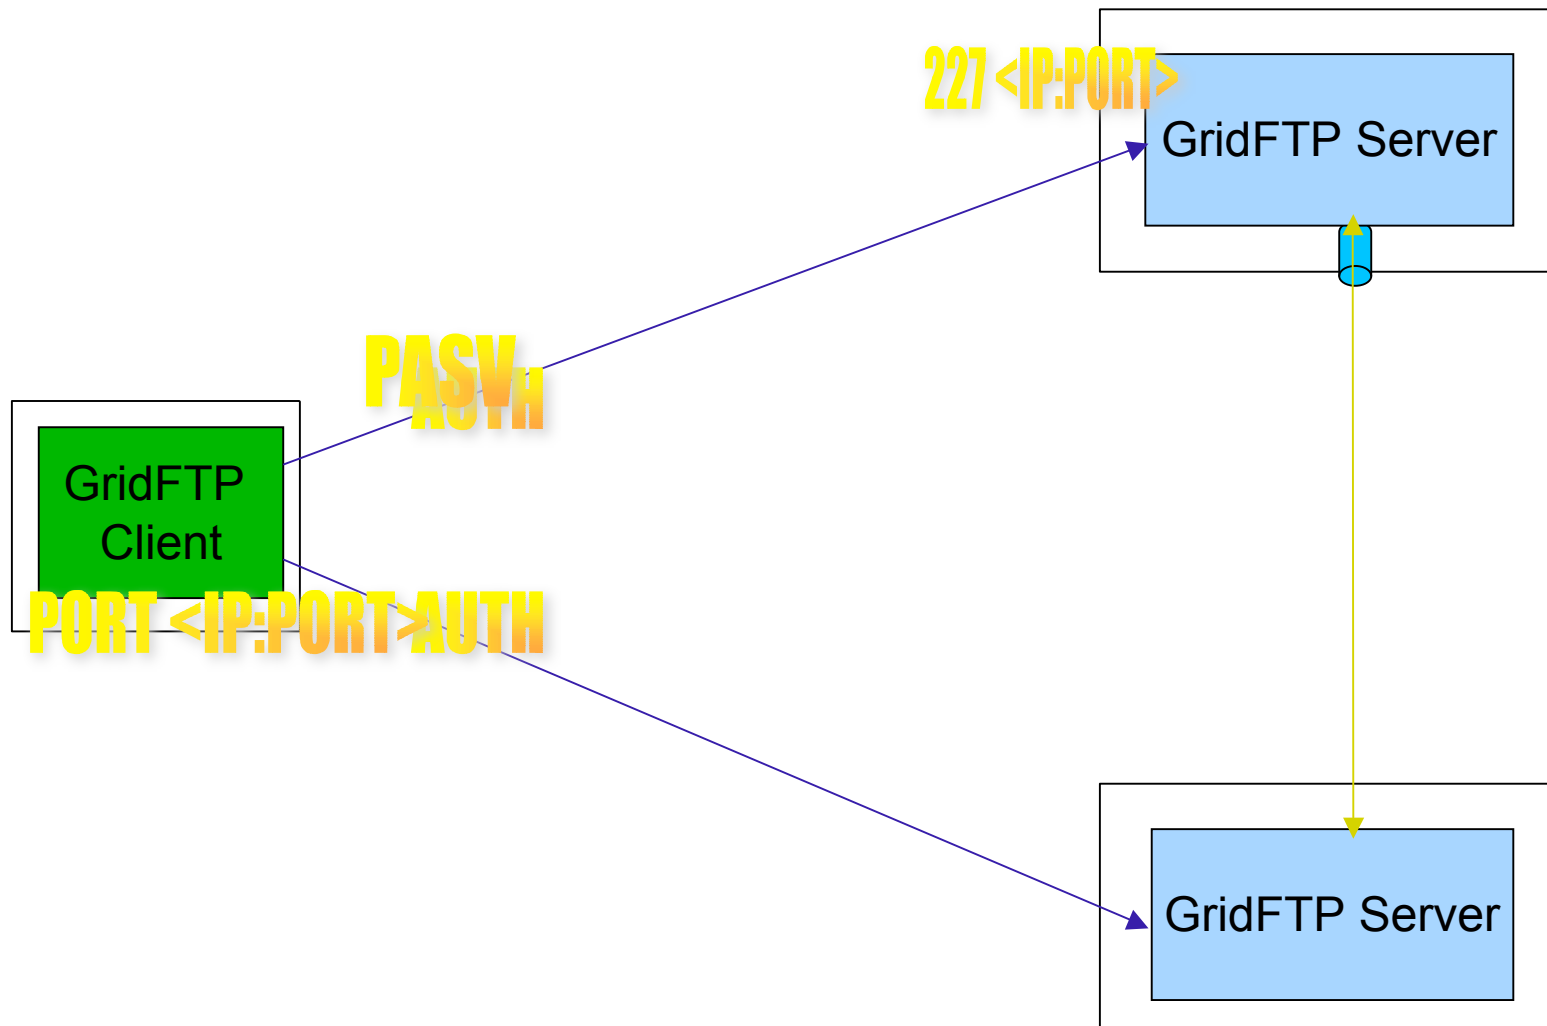
- ❑ Communication link(s) over which the actual **data** of interest flows
- ❑ High Bandwidth; authenticated by default; encryption and integrity protection optional



# Control Channel Establishment

- Server listens on a well-known port (2811)
- Client form a TCP Connection to server
- Authentication
  - ❑ Anonymous
  - ❑ Clear text USER <username>/PASS <pw>
  - ❑ Base 64 encoded GSI handshake

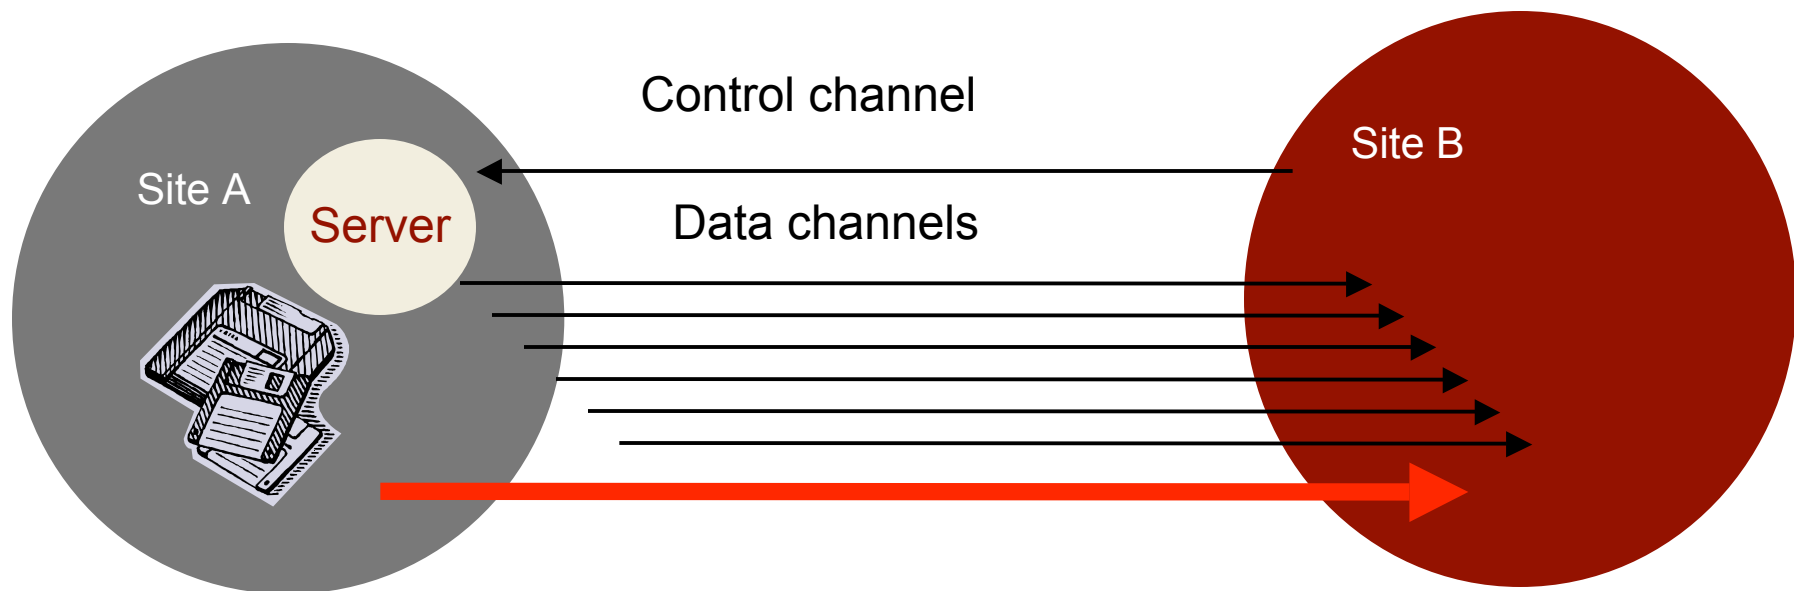
# Data Channel Establishment







# Going fast – parallel streams

- Use several data channels
- TCP - default transport protocol used by GridFTP
- TCP has limitations on high bandwidth wide area networks




# Parallel Streams



# Cluster-to-cluster data transfer

GridFTP can do coordinated data transfer utilizing multiple computer nodes at source and destination



# GridFTP usage

- **globus-url-copy - commonly used GridFTP client**
  - **Usage: globus-url-copy [options] srcurl dsturl**
- Conventions on URL formats:
  - **file:///home/YOURLOGIN/dataex/largefile**
    - a file called **largefile** on the local file system, in directory **/home/YOURLOGIN/dataex/**
  - **gsiftp://osg-edu.cs.wisc.edu/scratch/YOURLOGIN/**
    - a directory accessible via gsiftp on the host called **osg-edu.cs.wisc.edu** in directory **/scratch/YOURLOGIN**.

# GridFTP transfers using globus-url-copy

- globus-url-copy

`file:///home/YOURLOGIN/dataex/myfile`

`gsiftp://osg-edu.cs.wisc.edu/nfs/osgedu/YOURLOGIN/ex1`

- globus-url-copy

`gsiftp://osg-edu.cs.wisc.edu/nfs/osgedu/YOURLOGIN/ex2`

`gsiftp://tp-osg.ci.uchicago.edu/YOURLOGIN/ex3`

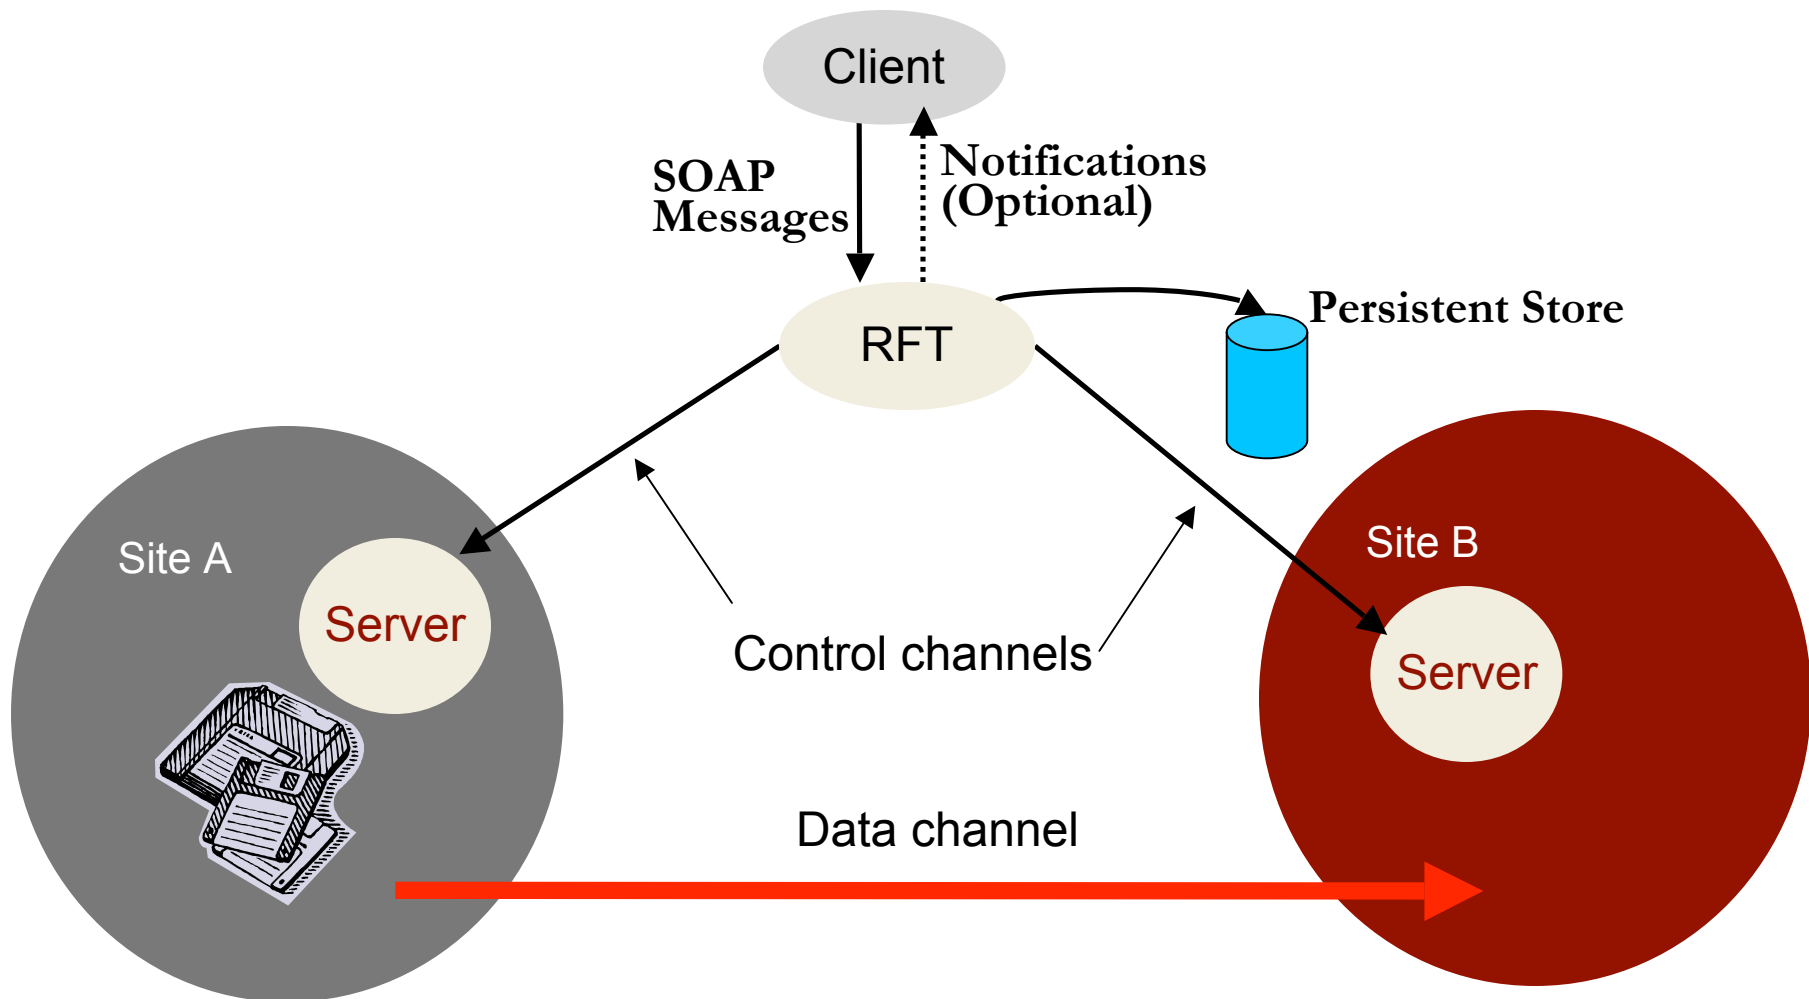
# Handling failures

- GridFTP server sends restart and performance markers periodically
  - ❑ Restart markers are helpful if there is any failure
  - ❑ No need to transfer the entire file again
  - ❑ Use restart markers and transfer only the missing pieces
- GridFTP supports partial file transfers
  - ❑ Globus-url-copy has a retry option
  - ❑ Recover from transient server and network failures
  - ❑ What if the client (globus-url-copy) fails in the middle of a transfer?

# RFT = Reliable file transfer

- GridFTP client that provides more reliability and fault tolerance for file transfers
  - Part of the Globus Toolkit
- RFT acts as a client to GridFTP, providing management of a large number of transfer jobs (same as Condor to GRAM)
- RFT can
  - keep track of the state of each job
  - run several transfers at once
  - deal with connection failure, network failure, failure of any of the servers involved.

# RFT





# RFT example

- Use the rft command with a .xfr file
- `cp /soft/globus-4.0.3-r1/share/globus_wsrf_rft_client/transfer.xfr rft.xfr`
- Edit rft.xfr to match your needs
- `rft -h terminable.ci.uchicago.edu -f ./rft.xfr`

# RLS - Replica Location Service

## ■ RLS

- ❑ Component of the data grid architecture (Globus component)
- ❑ It provides access to mapping information from logical names to physical names of items
- ❑ Its goal is to reduce access latency, improve data locality, improve robustness, scalability and performance for distributed applications

## ■ Logical Filenames (LFN)

- ❑ Names a file with interesting data in it
- ❑ Doesn't refer to location (which host, or where in a host)

## ■ Physical Filenames (PFN)

- ❑ Refers to a file on some filesystem somewhere
- ❑ Often use `gsiftp://` URLs to specify

## ■ Two RLS catalogs:

- ❑ Local Replica Catalog (LRC) and
- ❑ Replica Location Index (RLI)

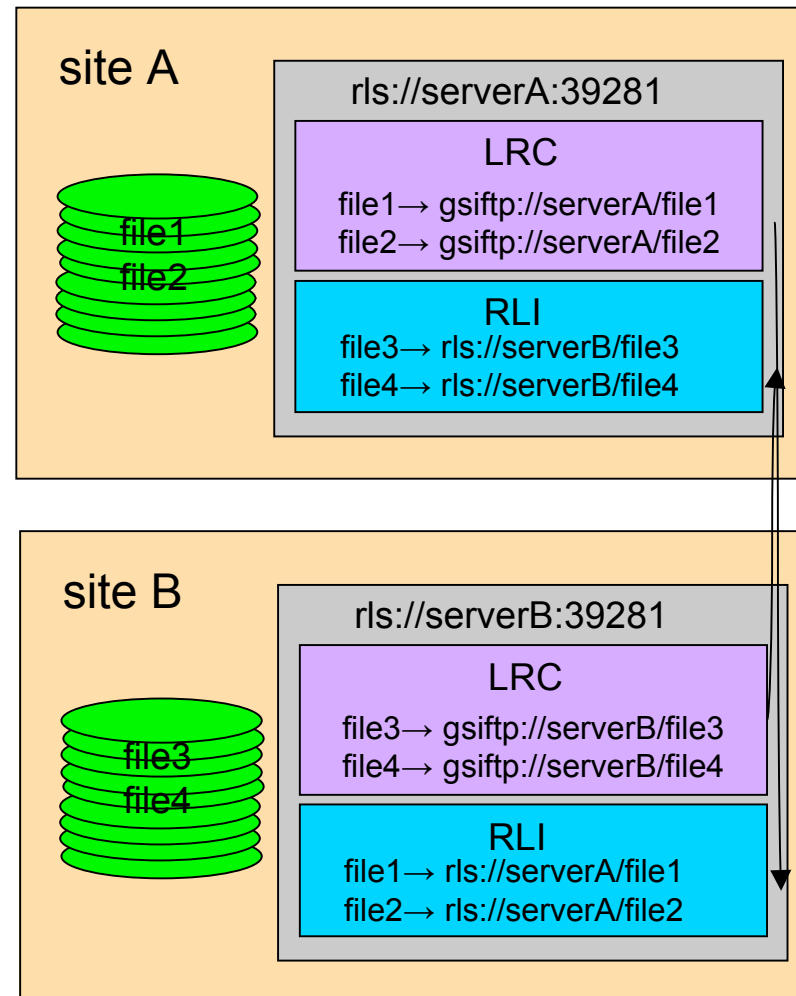
# Local Replica Catalog (LRC)

- Stores mappings from LFNs to PFNs.
- Interaction:
  - ❑ *Q*: Where can I get filename 'experiment\_result\_1'?
  - ❑ *A*: You can get it from  
gsiftp://gridlab1.ci.uchicago.edu/home/benc/r.txt
- Undesirable to have one of these for whole grid
  - ❑ Lots of data
  - ❑ Single point of failure
- Index LRCs

# Replica Location Index (RLI)

- Stores mappings from LFNs to LRCs.
- Interaction:
  - ❑ *Q*: Who can tell me about filename 'experiment\_result\_1'.
  - ❑ *A*: You can get more info from the LRC at gridlab1
  - ❑ (Then go to ask that LRC for more info)
- Failure of one RLI or LRC doesn't break everything
- RLI stores reduced set of information, so can cope with many more mappings

# Globus RLS



# Globus RLS

## ■ Quick Review

- ❑ LFN → logical filename (think of as simple filename)
- ❑ PFN → physical filename (think of as a URL)
- ❑ LRC → your local catalog of maps from LFNs to PFNs
  - H-R-792845521-16.gwf → gsiftp://dataserver.phys.uwm.edu/LIGO/H-R-792845521-16.gwf
- ❑ RLI → your local catalog of maps from LFNs to LRCs
  - H-R-792845521-16.gwf → LRCs at MIT, PSU, Caltech, and UW-M
- ❑ LRCs inform RLIs about mappings known

## ■ Can query for files is a 2-step process: find files on your Grid by

- ❑ querying RLI(s) to get LRC(s)
- ❑ then query LRC(s) to get URL(s)

# Globus RLS: Server Perspective

- Mappings LFNs → PFNs kept in database
- **Mappings LFNs → LRCs** stored in 1 of 2 ways
- **Table in database**
  - full, complete listing from LRCs that update your RLI
  - requires each LRC to send your RLI full, complete list
    - as number of LFNs in catalog grows, this becomes substantial
    - $10^8$  filenames at 64 bytes per filename ~ 6 GB
- **In memory in a special hash called Bloom filter**
  - $10^8$  filenames stored in as little as 256 MB
    - easy for LRC to create Bloom filter and send over network to RLIs
  - can cause RLI to lie when asked if knows about a LFN
    - only false-positives
    - acceptable in many contexts
  - Wild carding not possible with Bloom Filters

# Bloom filters

- LRC-to-RLI flow can happen in two ways:
  - LRC sends list of all its LFNs (but not PFNs) to the RLI. RLI stores whole list.
    - Answer accurately: “Yes I know” / “No I don’t know”
    - Expensive to move and store large list
  - Bloom filters
    - LRC generates a Bloom filter of all of its LFNs
    - Bloom filter is a bitmap that is much smaller than whole list of LFNs
    - Answers less accurately: “Maybe I know” / “No I don’t know”. Might end up querying LRCs unnecessarily (but we won’t ever get wrong answers)
    - can’t do a wildcard search



# RLS command line tools

- **globus-rls-admin**
  - administrative tasks
    - ping server
    - connect RLIs and LRCs together
- **globus-rls-cli**
  - end user tasks
    - query LRC and RLI
    - add mappings to LRC

# Globus RLS: Client Perspective

Two ways for clients to interact with RLS Server

- **globus-rls-cli** simple command-line tool
  - ❑ query
  - ❑ create new mappings
- “roll your own” client by coding against API
  - ❑ Java
  - ❑ C
  - ❑ Python

# Globus-rls-cli

Simple query to LRC to find a PFN for LFN

- Note more than one PFN may be returned

```
$ globus-rls-cli query lrc lfn some-file.jpg rls://dataserver:39281
```

```
some-file.jpg : file://localhost/netdata/s001/S1/R/H/714023808-714029599/some-file.jpg
```

```
some-file.jpg : file://medusa-slave001.medusa.phys.uwm.edu/data/S1/R/H/714023808-714029599/some-file.jpg
```

```
some-file.jpg : gsiftp://dataserver.phys.uwm.edu:15000/data/gsiftp_root/cluster_storage/data/s001/S1/R/H/714023808-714029599/some-file.jpg
```

- Server and client sane if LFN not found

```
$ globus-rls-cli query lrc lfn foo rls://dataserver
```

```
LFN doesn't exist: foo
```

```
$ echo $?
```

```
1
```

# Globus-rls-cli

## Wildcard searches of LRC supported

- ❑ probably a good idea to quote LFN wildcard expression

```
$ globus-rls-cli query wildcard lrc lfn H-R-7140242*-16.gwf
rls://dataserver:39281
H-R-714024208-16.gwf:
gsiftp://dataserver.phys.uwm.edu:15000/data/gsiftp_root/cluster_storage/data/s001/S1/R/H/714023808-714029599/H-R-714024208-16.gwf
H-R-714024224-16.gwf:
gsiftp://dataserver.phys.uwm.edu:15000/data/gsiftp_root/cluster_storage/data/s001/S1/R/H/714023808-714029599/H-R-714024224-16.gwf
```

# Globus-rls-cli

Bulk queries also supported

- obtain PFNs for more than one LFN at a time

```
$ globus-rls-cli bulk query lrc lfn H-R-714024224-16.gwf  
H-R-714024320-16.gwf rls://dataserver
```

```
H-R-714024320-16.gwf:
```

```
gsiftp://dataserver.phys.uwm.edu:15000/data/gsiftp_root/  
cluster_storage/data/s001/S1/R/H/714023808-714029599/H-  
R-714024320-16.gwf
```

```
H-R-714024224-16.gwf:
```

```
gsiftp://dataserver.phys.uwm.edu:15000/data/gsiftp_root/  
cluster_storage/data/s001/S1/R/H/714023808-714029599/H-  
R-714024224-16.gwf
```

# Globus-rls-cli

Simple query to RLI to locate a LFN -> LRC mapping

❑ then query that LRC for the PFN

```
$ globus-rls-cli query rli lfn example-file.gwf  
rls://dataserver
```

```
example-file.gwf: rls://ldas-cit.ligo.caltech.edu:39281
```

```
$ globus-rls-cli query lrc lfn example-file.gwf rls://ldas-  
cit.ligo.caltech.edu:39281
```

```
example-file: gsiftp://ldas-  
cit.ligo.caltech.edu:15000/archive/S1/L0/LHO/H-R-7140/H-R-  
714024224-16.gwf
```

# Globus-rls-cli

- Bulk queries to RLI also supported

```
$ globus-rls-cli bulk query rli lfn H-R-714024224-16.gwf H-R-714024320-16.gwf rls://dataserver
H-R-714024320-16.gwf: rls://ldas-cit.ligo.caltech.edu:39281
H-R-714024224-16.gwf: rls://ldas-cit.ligo.caltech.edu:39281
```

- Wildcard queries to RLI may not be supported!

- no wildcards when using Bloom filter updates

```
$ globus-rls-cli query wildcard rli lfn "H-R-7140242*-16.gwf"
rls://dataserver
```

Operation is unsupported: Wildcard searches with Bloom filters

# Globus-rls-cli

## Create new LFN → PFN mappings

- ❑ use **create** to create 1<sup>st</sup> mapping for a LFN

```
$ globus-rls-cli create file1 gsiftp://dataserver/file1  
rls://dataserver
```

- ❑ use **add** to add more mappings for a LFN

```
$ globus-rls-cli add file1 file://dataserver/file1  
rls://dataserver
```

- ❑ use **delete** to remove a mapping for a LFN

- when last mapping is deleted for a LFN the LFN is also deleted
- cannot have LFN in LRC without a mapping

```
$ globus-rls-cli delete file1 file://file1 rls://dataserver
```



# Globus-rls-cli

LRC can also store attributes about LFN and PFNs

- ❑ size of LFN in bytes?
- ❑ md5 checksum for a LFN?
- ❑ ranking for a PFN or URL?
- ❑ extensible...you choose attributes to create and add
- ❑ can search catalog on the attributes
- ❑ attributes limited to
  - strings
  - integers
  - floating point (double)
  - date/time

# Globus-rls-cli

- Create attribute first then add values for LFNs

```
$ globus-rls-cli attribute define md5checksum lfn string  
rls://dataserver
```

```
$ globus-rls-cli attribute add file1 md5checksum lfn  
string 42947c86b8a08f067b178d56a77b2650 rls://dataserver
```

- Then query on the attribute

```
$ globus-rls-cli attribute query file1 md5checksum lfn  
rls://dataserver  
md5checksum: string: 42947c86b8a08f067b178d56a77b2650
```

# Storage and Grid

- Storage resources are shared on the Grid
  - ❑ Layered storage on Grid systems (memory, disk and tape)
  - ❑ Need to make sure that the file is not swapped off of the disk to tape while the file is being transferred
  - ❑ Data receiving end can run out of space in the middle of a transfer
  - ❑ Reserve space on the destination end

# Storage and Grid

- ❑ Unlike compute/network resources, storage resources are not available when jobs are done
- ❑ Release resource usage when done, unreleased resource need to be garbage collected
- ❑ Need to enforce quotas
- ❑ Need to ensure fairness of space allocation and scheduling

# What is SRM?

- Storage Resource Managers (SRMs) are middleware components
  - whose function is to provide
    - dynamic space allocation
    - file management on shared storage resources on the Grid
  - Different **implementations** for underlying storage systems are based on the same SRM **specification**

# Managing spaces

- Negotiation
  - Client asks for space: `Guaranteed_C`, `MaxDesired`
  - SRM return: `Guaranteed_S ≤ Guaranteed_C`,  
best effort `≤ MaxDesired`
- Types of spaces
  - Access Latency (Online, Nearline)
  - Retention Policy (Replica, Output, Custodial)
  - Subject to limits per client (SRM or VO policies)
- Lifetime
  - Negotiated: `Lifetime_C` requested
  - SRM return: `Lifetime_S ≤ Lifetime_C`
- Reference handle
  - SRM returns space reference handle (space token)
- Updating space
  - Resize for more space or release unused space
  - Extend or shorten the lifetime of a space

# File Management

- Assignment of files to spaces
  - ❑ Files can be assigned to any space, provided that their lifetime is shorter than the remaining lifetime of the space
  - ❑ Files can be put into an SRM without explicit reservation
  - ❑ Default spaces are not visible to client
- Files already in the SRM can be moved to other spaces
  - ❑ By `srmChangeSpaceForFiles`
- Files already in the SRM can be pinned in spaces
  - ❑ By requesting specific files (`srmPrepareToGet`)
  - ❑ By pre-loading them into online space (`srmBringOnline`)
- Releasing files from space by a user
  - ❑ Release all files that user brought into the space whose lifetime has not expired

# Transfer protocol negotiation

## ■ Negotiation

- ❑ Client provides an ordered list of preferred transfer protocols
- ❑ SRM returns first protocol from the list it supports
- ❑ Example
  - Client provided protocols list: bbftp, gridftp, ftp
  - SRM returns: gridftp

## ■ Advantages

- ❑ Easy to introduce new protocols
- ❑ User controls which transfer protocol to use

## ■ How it is returned?

- ❑ Transfer URL (TURL)
- ❑ Example: bbftp://dm.slac.edu//temp/run11/File678.txt



# Site URL and Transfer URL

- Provide: Site URL (SURL)
  - URL known externally – e.g. in Replica Catalogs
  - e.g. `srm://ibm.cnaf.infn.it:8444/dteam/test.10193`
- Get back: Transfer URL (TURL)
  - Path can be different from SURL – SRM internal mapping
  - Protocol chosen by SRM based on request protocol preference
  - e.g. `gsiftp://ibm139.cnaf.infn.it:2811//gpfs/sto1/dteam/test.10193`
- One SURL can have many TURLs
  - Files can be replicated in multiple storage components
  - Files may be in near-line and/or on-line storage
  - In a light-weight SRM (a single file system on disk)
    - SURL may be the same as TURL except protocol

# Directory management

- Usual unix semantics
  - `srmLs`, `srmMkdir`, `srmMv`, `srmRm`, `srmRmdir`
- A single directory for all spaces
  - No directories for each file type
  - File assignment to spaces is virtual

# OSG & Data management

- OSG relies on GridFTP protocol for the raw transport of the data using Globus GridFTP in all cases except where interfaces to storage management systems (rather than file systems) dictate individual implementations.
- OSG supports the SRM interface to storage resources to enable management of space and data transfers to prevent unexpected errors due to running out of space, to prevent overload of the GridFTP services, and to provide capabilities for pre-staging, pinning and retention of the data files. OSG currently provides reference implementations of two storage systems the [\(BeStMan\)](#) and [dCache](#)

# Credits

Bill Allcock [allcock@mcs.anl.gov](mailto:allcock@mcs.anl.gov)

Ben Clifford [benc@ci.uchicago.edu](mailto:benc@ci.uchicago.edu)

Scott Koranda [skoranda@uwm.edu](mailto:skoranda@uwm.edu)

Alex Sim [asim@lbl.gov](mailto:asim@lbl.gov)



Open Science Grid



Open Science Grid