



Building a Campus Grid



Mats Rynge – rynge@renci.org

Renaissance Computing Institute
University of North Carolina, Chapel Hill






Outline

- Recipe
- Condor Daemons
- ClassAds
- Configuration Files
- Configuration Examples
 - Creating a central manager
 - Joining a dedicated processing machine
 - Joining a interactive workstation
- Troubleshooting
- Security




Why a Campus Grid?

- Existing, under-utilized, resources all over campus!
- Enable departments and researchers on campus to share their *existing* computers
- The same departments and researchers will have access to a **larger pool** of resources when they need to do large computations
- Community driven project
 - Enabled by Central IT
- Corner stone to join national cyber infrastructure



What is a campus Condor pool? (cont)





Recipe

- Central IT will provide central services (collector, negotiator, one submit node, ...)
- Campus departments / researches can join their existing compute resources to the pool, and share with the community
- Resource owners have full control over the policy for their resources



Recipe - Required Tools

- Condor
- A central IT champion
 - Configure and maintain the central services
 - Outreach to departments and users (remember, this is about distributed resources!)
 - Documentation and support





Condor Daemons

- You only have to run the daemons for the services you need to provide
- **DAEMON_LIST** is a comma separated list of daemons to start
 - **DAEMON_LIST=MASTER, SCHEDD, STARTD**



Condor Daemons

- **condor_master** - controls everything else
- **condor_startd** - executing jobs
- **condor_schedd** - submitting jobs
- **condor_collector** - Collects system information; only on Central Manager
- **condor_negotiator** - Assigns jobs to machines; only on Central Manager




condor_master

- Provides access to many remote administration commands:
 - `condor_reconfig`, `condor_restart`, `condor_off`, `condor_on`, etc.
- Default server for many other commands:
 - `condor_config_val`, etc.



condor_startd

- Represents a machine willing to run jobs to the Condor pool
- Run on any machine you want to run jobs on
- Enforces the wishes of the machine owner (the owner's "policy")




condor_schedd

- Represents jobs to the Condor pool
- Maintains persistent queue of jobs
 - Queue is not strictly first-in-first-out (priority based)
 - Each machine running **condor_schedd** maintains its own independent queue
- Run on any machine you want to submit jobs from



condor_collector

- Collects information from all other Condor daemons in the pool
- Each daemon sends a periodic update called a ClassAd to the collector
 - Old ClassAds removed after a time out
- Services queries for information:
 - Queries from other Condor daemons
 - Queries from users (**condor_status**)



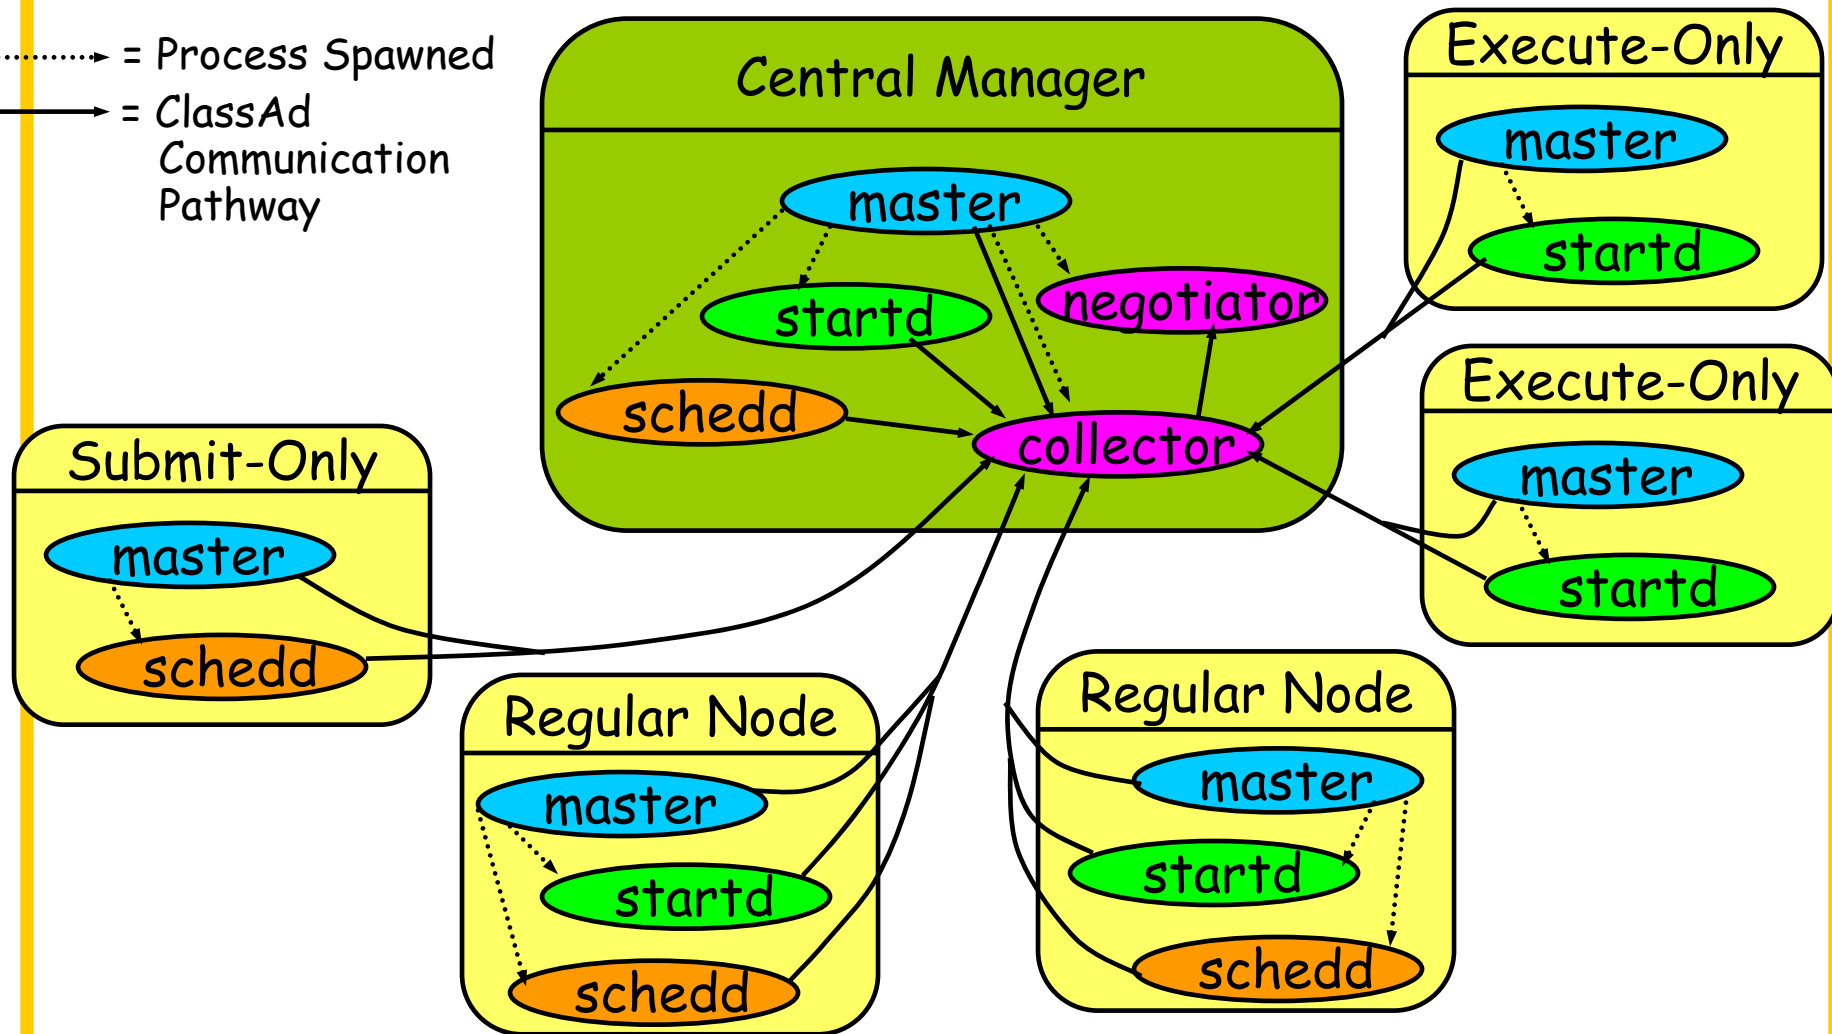
condor_negotiator

- Performs matchmaking in Condor
 - Pulls list of available machines and job queues from **condor_collector**
 - Matches jobs with available machines
 - Both the job and the machine must satisfy each other's requirements (2-way matching)
- Handles ranks and priorities




Typical Condor Pool


.....> = Process Spawned
——> = ClassAd Communication Pathway






Open Science Grid Job Startup





ClassAds – the basic building block of Condor

- Set of key-value pairs
- Values can be expressions
- Can be matched against each other
 - Requirements and Rank
 - MY.name – Looks for “name” in local ClassAd
 - TARGET.name – Looks for “name” in the other ClassAd
 - Name – Looks for “name” in the local ClassAd, then the other ClassAd




ClassAd Expressions

- Some configuration file macros specify expressions for the Machine's ClassAd
 - Notably **START**, **RANK**, **SUSPEND**, **CONTINUE**, **PREEMPT**, **KILL**
- Can contain a mixture of macros and ClassAd references
- Notable: **UNDEFINED**, **ERROR**



ClassAd Expressions

- `+`, `-`, `*`, `/`, `<`, `<=`, `>`, `>=`, `==`, `!=`, `&&`, and `||` all work as expected
- `TRUE==1` and `FALSE==0` (guaranteed)



ClassAd Expressions: UNDEFINED and ERROR

- Special values
- Passed through most operators
 - Anything == UNDEFINED is UNDEFINED
- && and || eliminate if possible.
 - UNDEFINED && FALSE is FALSE
 - UNDEFINED && TRUE is UNDEFINED



Open Science Grid ClassAd Expressions:


=?= and !==

- **=?=** and **!==** are similar to **==** and **!=**
- **=?=** tests if operands have the same type and the same value.
 - **10 == UNDEFINED -> UNDEFINED**
 - **UNDEFINED == UNDEFINED -> UNDEFINED**
 - **10 ==?= UNDEFINED -> FALSE**
 - **UNDEFINED ==?= UNDEFINED -> TRUE**
- **!==** inverts **=?=**



Configuration Files

- Multiple files concatenated
 - Later definitions overwrite previous ones
- Order of files:
 - Global configuration file (only required file)
 - Local and shared configuration files




Global Configuration File

- Found either in file pointed to with the **CONDOR_CONFIG** environment variable, **/etc/condor/condor_config**, or **~condor/condor_config**
- All settings can be in this file
- “Global” on assumption it’s shared between machines. NFS, automated copies, etc.




Local Configuration File

- **LOCAL_CONFIG_FILE** macro
- Machine-specific settings
 - local policy settings for a given owner
 - different daemons to run (for example, on the Central Manager!)



Local Configuration File

- Can be on local disk of each machine
`/var/adm/condor/condor_config.local`
- Can be in a shared directory
 - Use `$(HOSTNAME)` which expands to the machine's name
`/shared/condor/condor_config.$(HOSTNAME)`
`/shared/condor/hosts/$(HOSTNAME) /
condor_config.local`



Policy

- Allows machine owners to specify job priorities, restrict access, and implement other local policies




Policy Expressions

- Specified in **condor_config**
 - Ends up startd/machine ClassAd
- Policy evaluates both a machine ClassAd and a job ClassAd together
 - Policy can reference items in either ClassAd (See manual for list)




Machine (Startd) Policy Expression Summary

- **START** – When is this machine willing to start a job
- **RANK** - Job preferences
- **SUSPEND** - When to suspend a job
- **CONTINUE** - When to continue a suspended job
- **PREEMPT** – When to nicely stop running a job
- **KILL** - When to immediately kill a preempting job



RANK

- Indicates which jobs a machine prefers
 - Jobs can also specify a rank
- Floating point number
 - Larger numbers are higher ranked
 - Typically evaluate attributes in the Job ClassAd
 - Typically use + instead of &&



RANK

- Often used to give priority to owner of a particular group of machines
- Claimed machines still advertise looking for higher ranked job to preempt the current job



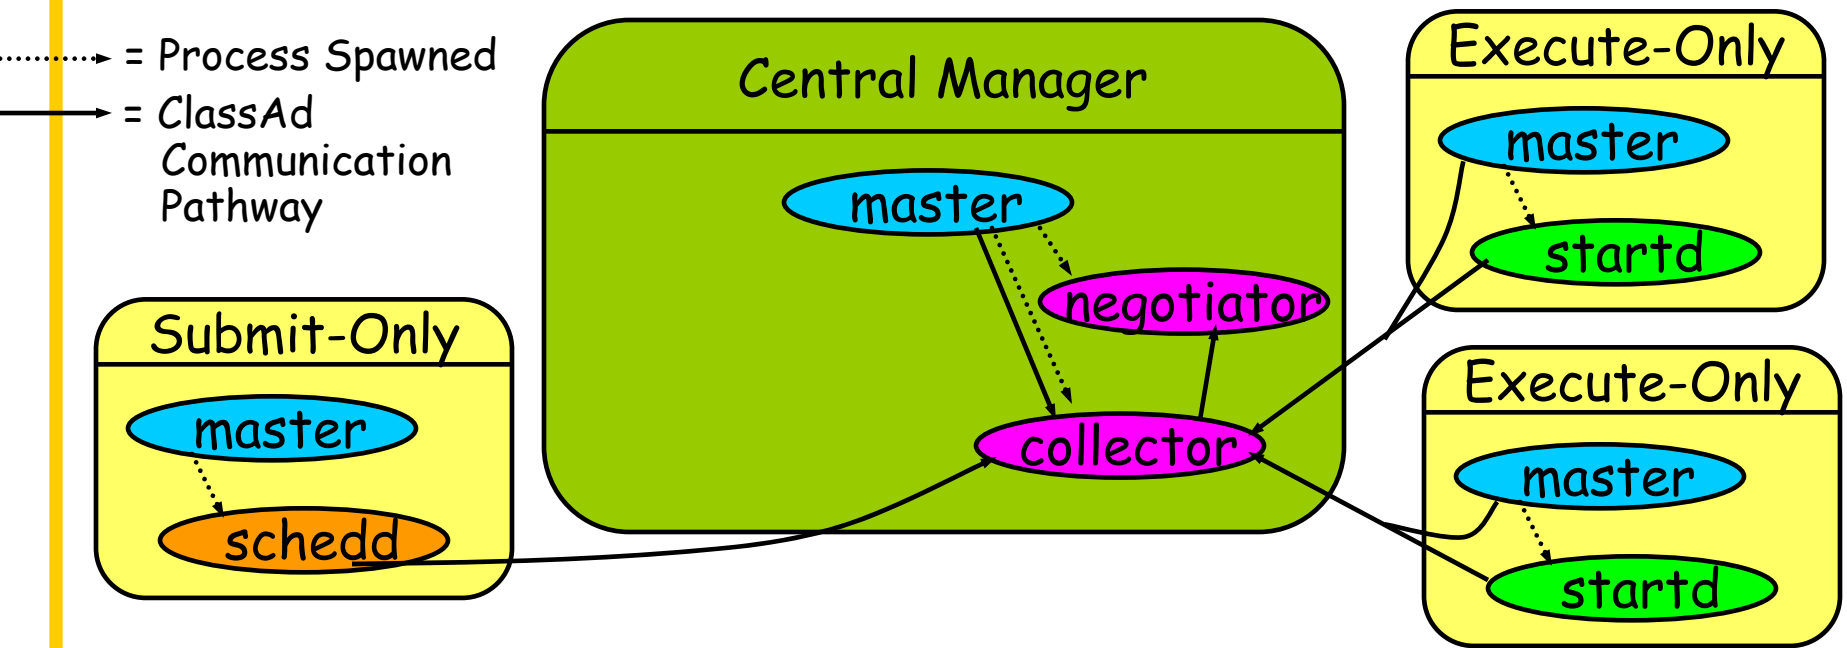
Configuration Case #1

Central/department IT setting up a
central manager



Needed Daemons

DAEMON_LIST = MASTER, COLLECTOR, NEGOTIATOR





Security


```
HOST_ALLOW_READ = *.unc.edu, 152.54.0.0/20
```

```
HOST_ALLOW_WRITE = *.unc.edu, 152.54.0.0/20
```

```
HOSTDENY_WRITE = *.wireless.unc.edu
```

Limit the number of allowed submit nodes (more on the security later)

```
HOSTALLOW_ADVERTISE_SCHEDD =  
    tarheelgrid.isis.unc.edu, fire.renci.org, ...
```

Configuration Case #2

Department or lab wanting to join
their existing **dedicated** processing
machine to campus pool



Case #2 - Configuration

`CONDOR_HOST = cm1.isis.unc.edu`

`DAEMON_LIST = MASTER, STARTD`

`HOST_ALLOW_READ = *.unc.edu, 152.54.0.0/20`

`HOST_ALLOW_WRITE = *.unc.edu, 152.54.0.0/20`

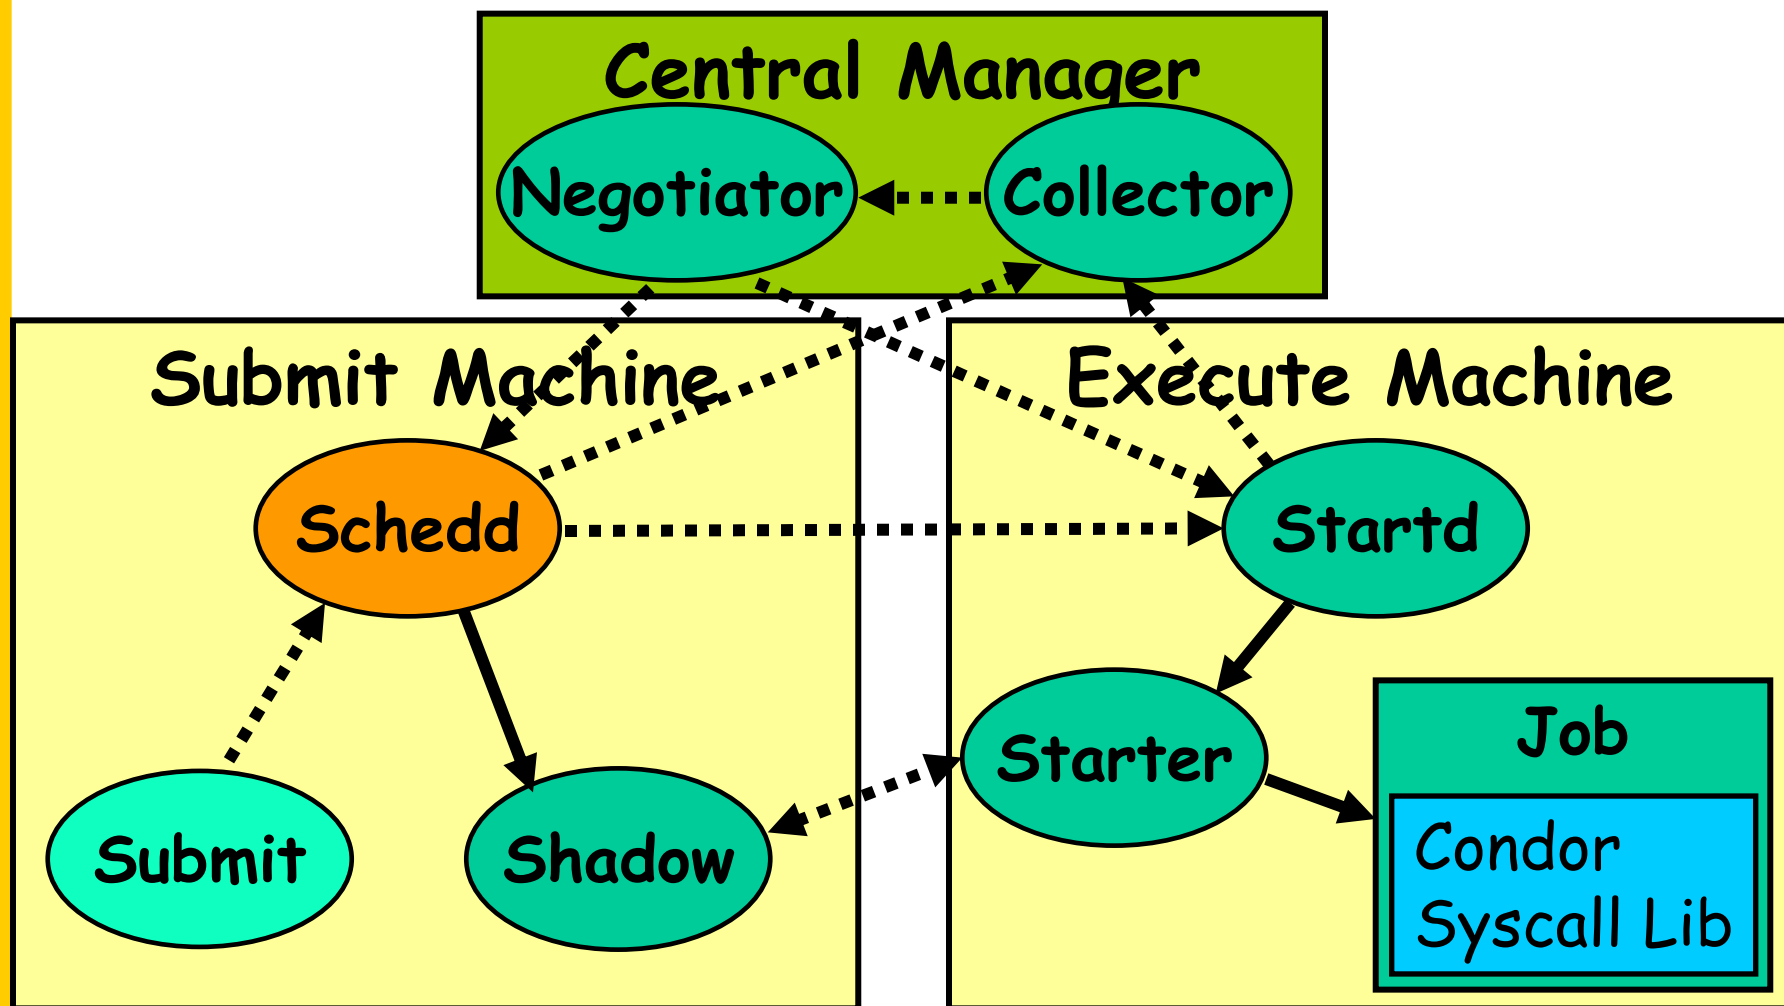
`HIGHPORT = 12000`

`LOWPORT = 10000`

Note that the firewall needs to be open, for the specific port range, against all hosts specified in
`HOST_ALLOW_READ/HOST_ALLOW_WRITE`



Open Science Grid Job Startup





Case #2 – Policy

Prefer Chemistry jobs

START = True


RANK = *Department != UNDEFINED &&
Department == "Chemistry"*

SUSPEND = False

CONTINUE = True


PREEMPT = False

KILL = False



Submit file with Custom Attribute

- Prefix an entry with “+” to add to job ClassAd
`Executable = charm-run`
`Universe = standard`
`+Department = "Chemistry"`
`queue`



Use Case #3

Department or lab wanting to join a
interactive workstation



Case #3 - Configuration

`CONDOR_HOST = cm1.isis.unc.edu`

`DAEMON_LIST = MASTER, STARTD`


`HOST_ALLOW_READ = *.unc.edu, 152.54.0.0/20`

`HOST_ALLOW_WRITE = *.unc.edu, 152.54.0.0/20`

`HIGHPORT = 12000`

`LOWPORT = 10000`

Note that the firewall needs to be open, for the specific port range, against all hosts specified in
`HOST_ALLOW_READ/HOST_ALLOW_WRITE`



Defining Idle

- One possible definition:
 - No keyboard or mouse activity for 5 minutes
 - Load average below 0.3



Desktops should...

- **START** jobs when the machine becomes idle
- **SUSPEND** jobs as soon as activity is detected
- **PREEMPT** jobs if the activity continues for 5 minutes or more
- **KILL** jobs if they take more than 5 minutes to preempt



Useful Attributes

- **LoadAvg**
 - Current load average
- **CondorLoadAvg**
 - Current load average generated by Condor
- **KeyboardIdle**
 - Seconds since last keyboard or mouse activity



Useful Attributes

- **CurrentTime**
 - Current time, in Unix epoch time (seconds since midnight Jan 1, 1970)
- **EnteredCurrentActivity**
 - When did Condor enter the current activity, in Unix epoch time



Macros in Configuration Files

```
NonCondorLoadAvg = (LoadAvg - CondorLoadAvg)
```

```
BgndLoad = 0.3
```


```
CPU_Busy = ($(NonCondorLoadAvg) >= $(BgndLoad))
```

```
CPU_Idle = ($(NonCondorLoadAvg) < $(BgndLoad))
```

```
KeyboardBusy = (KeyboardIdle < 10)
```

```
MachineBusy = ($(CPU_Busy) || $(KeyboardBusy))
```

```
ActivityTimer = \  
    (CurrentTime - EnteredCurrentActivity)
```



Desktop Machine Policy


START = \$(CPU_Idle) && KeyboardIdle > 300

SUSPEND = \$(MachineBusy)


CONTINUE = \$(CPU_Idle) && KeyboardIdle > 120

PREEMPT = (Activity == "Suspended") && \
\$(ActivityTimer) > 300

KILL = \$(ActivityTimer) > 300



Command Line Tools





condor_config_val

- Find current configuration values

```
% condor_config_val MASTER_LOG  
/var/condor/logs/MasterLog  
% cd `condor_config_val LOG`
```



condor_config_val -v


- Can identify source

```
% condor_config_val -v CONDOR_HOST
```

```
CONDOR_HOST: condor.cs.wisc.edu
```

```
Defined in
```

```
`/etc/condor_config.hosts', line 6
```

condor_config_val -config

- What configuration files are being used?

```
% condor_config_val -config
```

Config source:

```
/var/home/condor/condor_config
```

Local config sources:

```
/unsup/condor/etc/condor_config.hosts
```

```
/unsup/condor/etc/condor_config.global
```


```
/unsup/condor/etc/condor_config.policy
```

```
/unsup/condor-test/etc/hosts/puffin.local
```



Querying daemons `condor_status`


- Queries the collector for information about daemons in your pool
- Defaults to finding `condor_startds`
- `condor_status -schedd` summarizes all job queues
- `condor_status -master` returns list of all `condor_masters`



condor_status

- **-long** displays the full ClassAd
- Optionally specify a machine name to limit results to a single host

```
condor_status -l  
node4.cs.wisc.edu
```



condor_status -constraint

- Only return ClassAds that match an expression you specify
- Show me idle machines with 1GB or more memory

```
-condor_status -constraint  
'Memory >= 1024 && Activity ==  
"Idle"'
```



`condor_status -format`

- Controls format of output
- Useful for writing scripts
- Uses C printf style formats
 - One field per argument



condor_status -format


- Census of systems in your pool:

```
% condor_status -format '%s'  
Arch -format '%s\n' OpSys | sort  
| uniq -c  
797 INTEL LINUX  
118 INTEL WINNT50  
108 SUN4u SOLARIS28  
6 SUN4x SOLARIS28
```



Examining Queues `condor_q`

- View the job queue
- The “**-long**” option is useful to see the entire ClassAd for a given job
- supports **-constraint** and **-format**
- Can view job queues on remote machines with the “**-name**” option



condor_q -better-analyze

- (Heavily truncated output)

The Requirements expression for your job is:

```
( ( target.Arch == "SUN4u" ) && ( target.OpSys ==  
  "WINNT50" ) && [snip]
```

Condition	Machines	Suggestion
1 (target.Disk > 100000000)	0	MODIFY TO 14223201
2 (target.Memory > 10000)	0	MODIFY TO 2047
3 (target.Arch == "SUN4u")	106	
4 (target.OpSys == "WINNT50")	110	MOD TO "SOLARIS28"

Conflicts: conditions: 3, 4



Debugging Jobs: `condor_q`

- Examine the job with `condor_q`
 - especially `-long` and `-analyze`
 - Compare with `condor_status -long` for a machine you expected to match



Debugging Jobs: User Log

- Examine the job's user log

- Can find with:

```
condor_q -format '%s\n' UserLog 17.0
```

- Set with “log” in the submit file

- Contains the life history of the job
- Often contains details on problems



Debugging Jobs: ShadowLog

- Examine **ShadowLog** on the submit machine
 - Note any machines the job tried to execute on
 - There is often an “ERROR” entry that can give a good indication of what failed



Debugging Jobs: Matching Problems

- No **ShadowLog** entries? Possible problem matching the job.
 - Examine **ScheddLog** on the submit machine
 - Examine **NegotiatorLog** on the central manager



Debugging Jobs: Local Problems

- **ShadowLog** entries suggest an error but aren't specific?
 - Examine **StartLog** and **StarterLog** on the execute machine



Debugging Jobs: Reading Log Files

- Condor logs will note the job ID each entry is for
 - Useful if multiple jobs are being processed simultaneously
 - grepping for the job ID will make it easy to find relevant entries




Security





Security Threats

- Condor System
 - Authentication / Authorization
 - Next couple of slides
- Using Condor as a vehicle for attacks
 - Hard to prevent
 - Local exploits (privilege escalation)
 - Condor jobs are not sandboxed
 - Another example: Distributed DoS




Security

- Authentication
 - Users
 - Done on the submit machine, using OS authentication mechanisms
 - Machines
 - IP based
 - Machines on the wireless subnet are not allowed to join the pool
 - Other subnets to block?



Security - Users

- Based on the local UID
- Example: cdpoon@tarheelgrid.isis.unc.edu
- Note that these are not Campus wide identifiers. The components of the identifier is local username @ hostname of the submit node.
- Another example: engage@belhaven-1.renci.org
- These are getting logged during the negotiation cycle on the central manager
- Authorization can be done at the central manager, or in the policy on the execution machine



Security – Machines

- Anybody on campus can join a machine for job execution
- Only a set of machines will be allowed for job submits – this will ensure we have an audit trail
 - The allowed set is configured on the central manager



Security Links

- Condor documentation on security:
http://www.cs.wisc.edu/condor/manual/v7.2/3_6Security.html
- Building a secure Condor pool in an open academic environment
<http://www.allhands.org.uk/2005/proceedings/papers/435.pdf>