

---

# Grid Compute Resources and Job Management



Open Science Grid

---

# Grid middleware - “glues” all pieces together

- Offers services that couple **users** with **remote resources** through **resource brokers**
- Remote process management
- Co-allocation of resources
- Storage access
- Information
- Security
- QoS

# Terms:

- Globus
- GRAM
- Condor
- Condor-G

# Local Resource Managers (LRM)

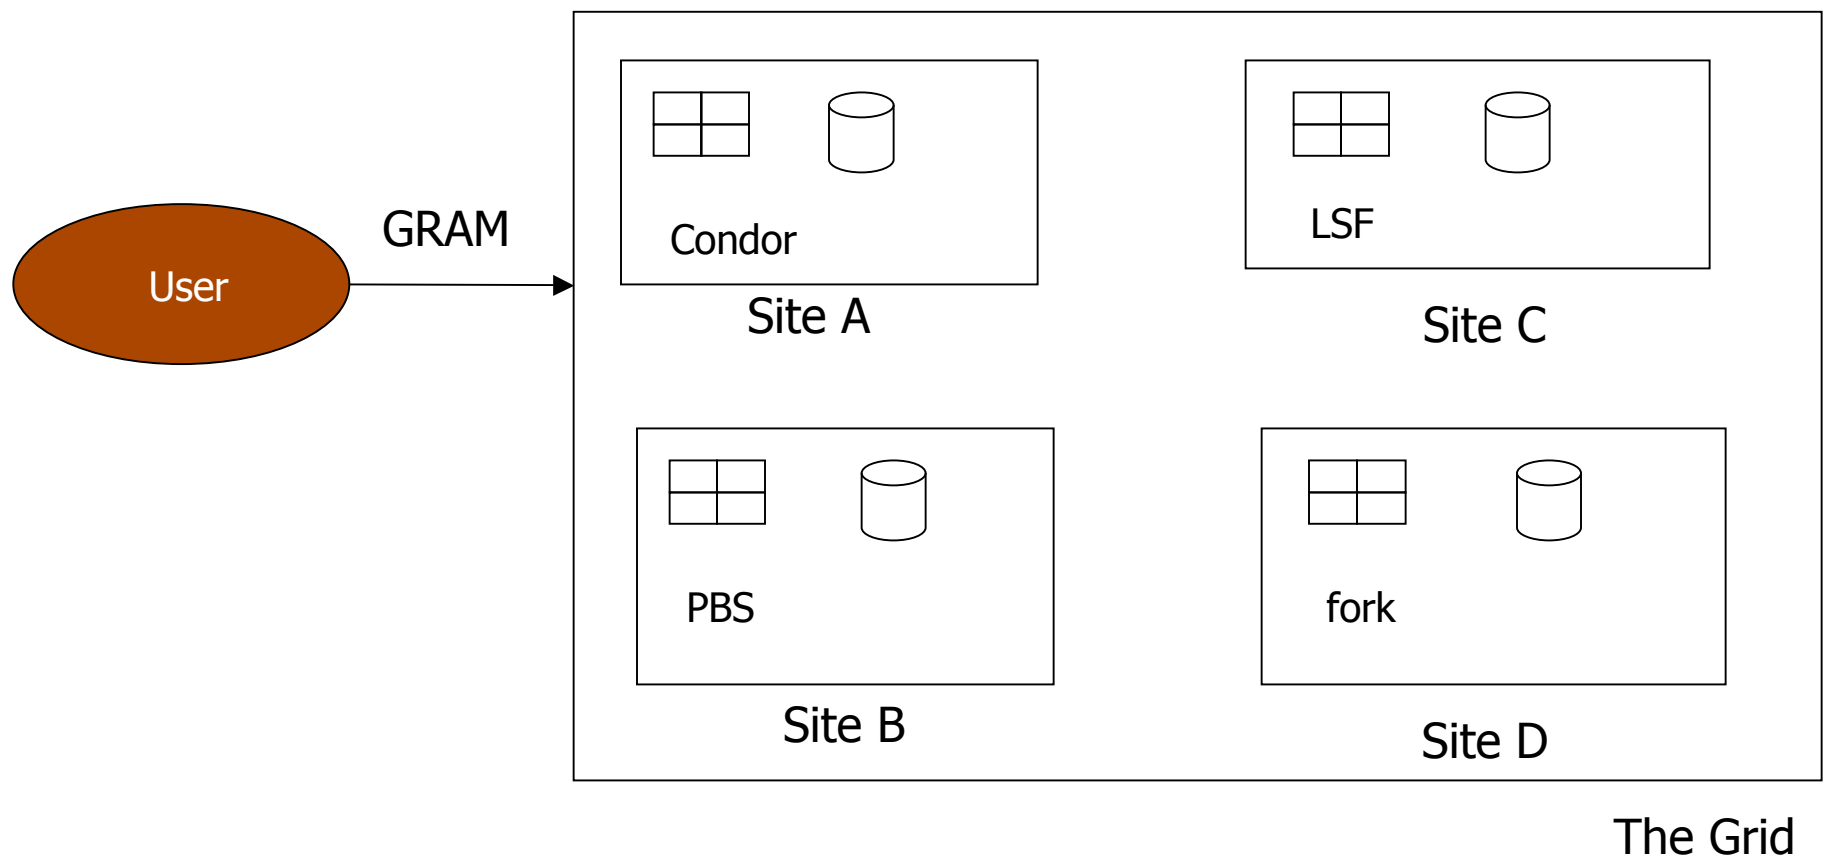
- Compute resources have a **local resource manager** (LRM) that controls:
  - ❑ Who is allowed to run jobs
  - ❑ How jobs run on a specific resource
  - ❑ Specifies the order and location of jobs
- *Example policy:*
  - ❑ Each cluster node can run one job.
  - ❑ If there are more jobs, then they must wait in a queue
- *Examples:* PBS, LSF, Condor

# GRAM

## Globus Resource Allocation Manager

- **GRAM** = provides a standardised interface to submit jobs to LRMs.
- Clients submit a job request to GRAM
- GRAM translates into something a(ny) LRM can understand
  - .... Same job request can be used for many different kinds of LRM


# Job Management on a Grid



# GRAM's abilities

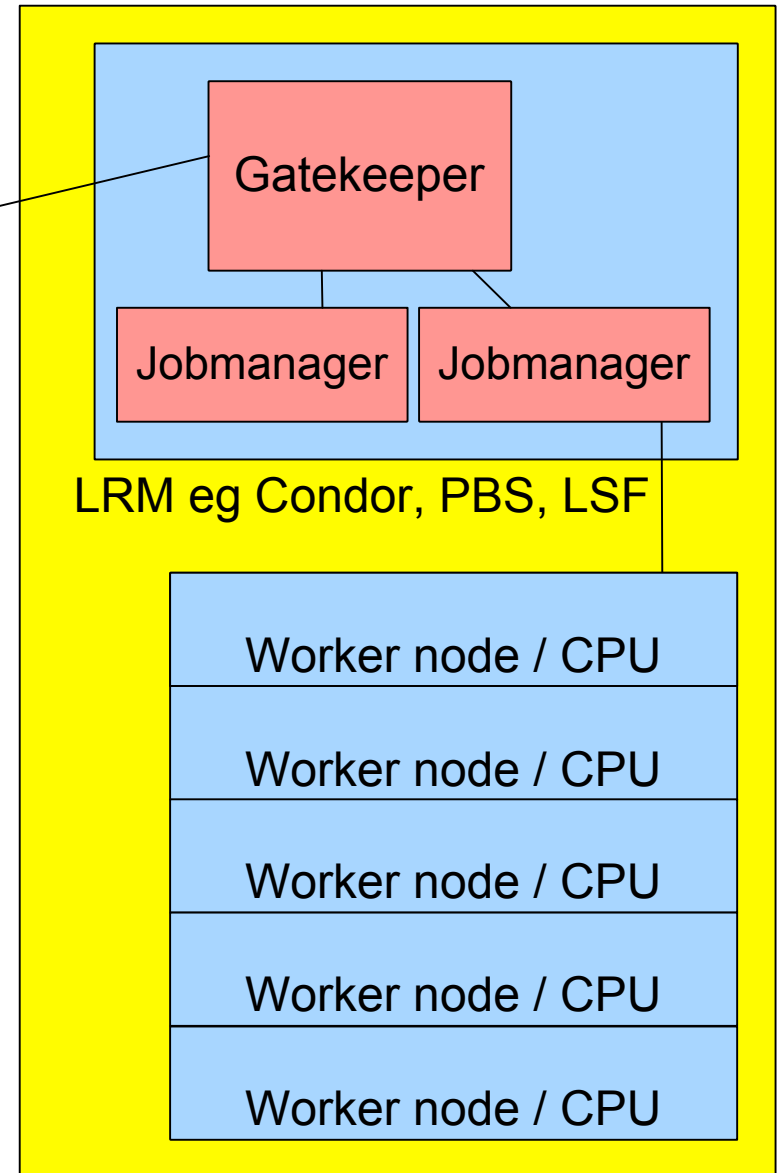
- Given a job specification:
  - ❑ Creates an environment for the job
  - ❑ Stages files to and from the environment
  - ❑ Submits a job to a local resource manager
  - ❑ Monitors a job
  - ❑ Sends notifications of the job state change
  - ❑ Streams a job's stdout/err during execution

# GRAM components



Submitting machine  
(e.g. User's workstation)

Internet





# Submitting a job with GRAM

- **globus-job-run** command

```
$ globus-job-run grid07.uchicago.edu /bin/hostname
```

- Run '/bin/hostname' on the resource grid07.uchicago.edu

- We don't care what LRM is used on 'grid07'.
- This command works with any LRM.

# Condor

- Condor is a specialized workload management system for compute-intensive jobs.
- is a software system that creates an HTC environment (created at [UW-Madison](#))
  - ❑ Detects machine availability
  - ❑ Harnesses available resources
  - ❑ Provides powerful resource management by *matching* resource owners with consumers (broker)

# How Condor works

## Condor provides:

- a job queueing mechanism
- scheduling policy
- priority scheme
- resource monitoring, and
- resource management.

Users **submit** their serial or parallel jobs to Condor,  
Condor places them into a **queue**,  
... chooses **when** and **where** to run the jobs based upon a policy,  
... carefully **monitors** their progress, and  
... ultimately **informs** the user upon completion.

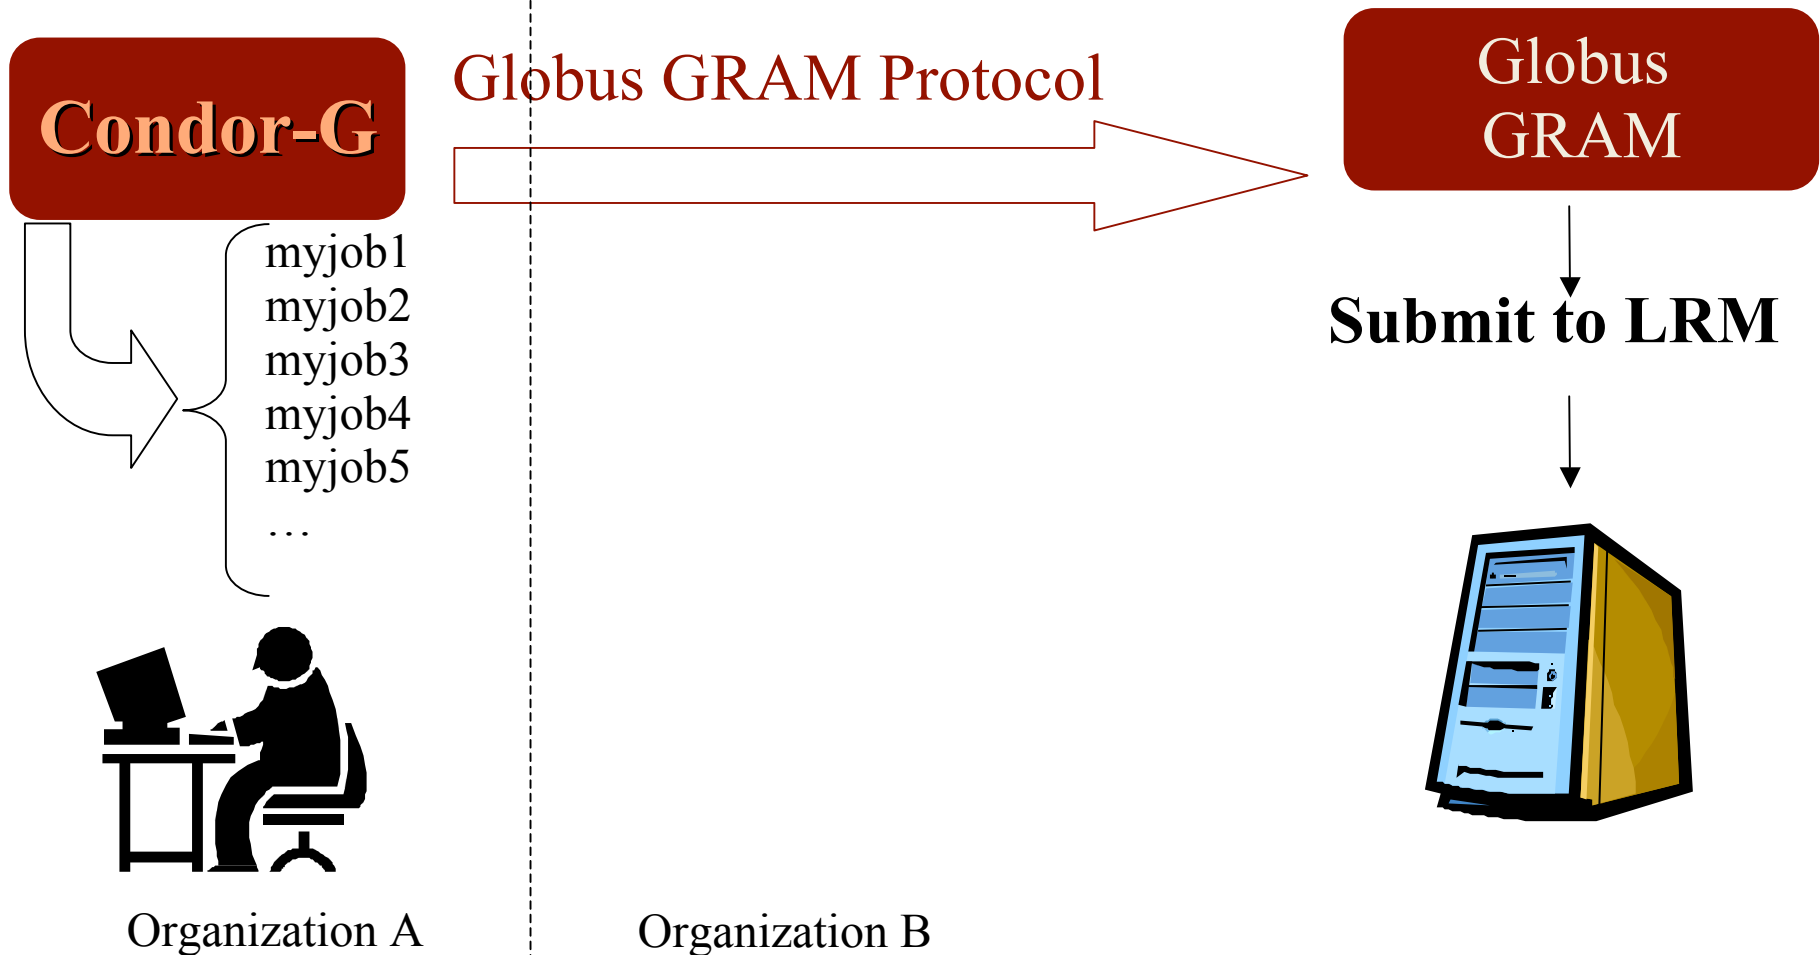
# Condor lets you manage a large number of jobs.

- Specify the jobs in a file and submit them to Condor
- Condor runs them and keeps you notified on their progress
  - ❑ Mechanisms to help you manage huge numbers of jobs (1000's), all the data, etc.
  - ❑ Handles inter-job dependencies (DAGMan)
- Users can set Condor's job priorities
- Condor administrators can set user priorities
- Can do this as:
  - ❑ Local resource manager (LRM) on a compute resource
  - ❑ Grid client submitting to GRAM (as Condor-G)

# Condor-G

- is the job management part of Condor.
  - *Hint:* Install Condor-G to submit to resources accessible through a Globus interface.
- Condor-G does not *create* a grid service.
- It only deals with *using* remote grid services.

# Remote Resource Access: Condor-G + Globus + Condor



# Four Steps to Run a Job with Condor

- These choices tell Condor
  - ❑ **how**
  - ❑ **when**
  - ❑ **where** to run the job,
  - ❑ and describe exactly **what** you want to run.
- Choose a Universe for your job
- Make your job batch-ready
- Create a *submit description* file
- Run *condor\_submit*

# Simple Submit Description File

```
# myjob.submit file
# Simple condor_submit input file
# (Lines beginning with # are comments)
# NOTE: the words on the left side are not
#       case sensitive, but filenames are!
```

```
Universe    = grid
Executable  = analysis
Log         = my_job.log
Queue
```



# Run condor\_submit

- You give *condor\_submit* the name of the submit file you have created:

```
condor_submit my_job.submit
```

- *condor\_submit* parses the submit file

# Another Submit Description File

```
# Example condor_submit input file
```

```
Universe      = grid
Executable    = /home/wright/condor/my_job.condor
Input         = my_job.stdin
Output        = my_job.stdout
Error         = my_job.stderr
Arguments     = -arg1 -arg2
InitialDir    = /home/wright/condor/run_1
Queue
```

# Other Condor commands

- `condor_q` – show status of job queue
- `condor_status` – show status of compute nodes
- `condor_rm` – remove a job
- `condor_hold` – hold a job temporarily
- `condor_release` – release a job from hold

# Submitting more complex jobs

- express dependencies between jobs  
⇒ WORKFLOWS
- And also, we would like the workflow to be managed even in the face of failures

# Want other Scheduling possibilities?

## Use the Scheduler Universe


- In addition to VANILLA, another job universe is the *Scheduler Universe*.
- Scheduler Universe jobs run on the submitting machine and serve as a meta-scheduler.
- **Condor's Scheduler Universe lets you set up and manage job workflows.**
- DAGMan meta-scheduler included
  - DAGMan manages these jobs

# DAGMan

- **Directed Acyclic Graph Manager**
- DAGMan allows you to specify the *dependencies* between your Condor jobs, so it can *manage* them automatically for you.
- (e.g., “Don’t run job “B” until job “A” has completed successfully.”)

# What is a DAG?


- A DAG is the **data structure** used by DAGMan to represent these dependencies.
- Each job is a **“node”** in the DAG.
- Each node can have any number of “parent” or “children” nodes – as long as there are **no loops**!



# Defining a DAG

- A DAG is defined by a *.dag file*, listing each of its nodes and their dependencies:

```
# diamond.dag
Job A a.sub
Job B b.sub
Job C c.sub
Job D d.sub
Parent A Child B C
Parent B C Child D
```



- each node will run the Condor job specified by its accompanying *Condor submit file*



# Submitting a DAG

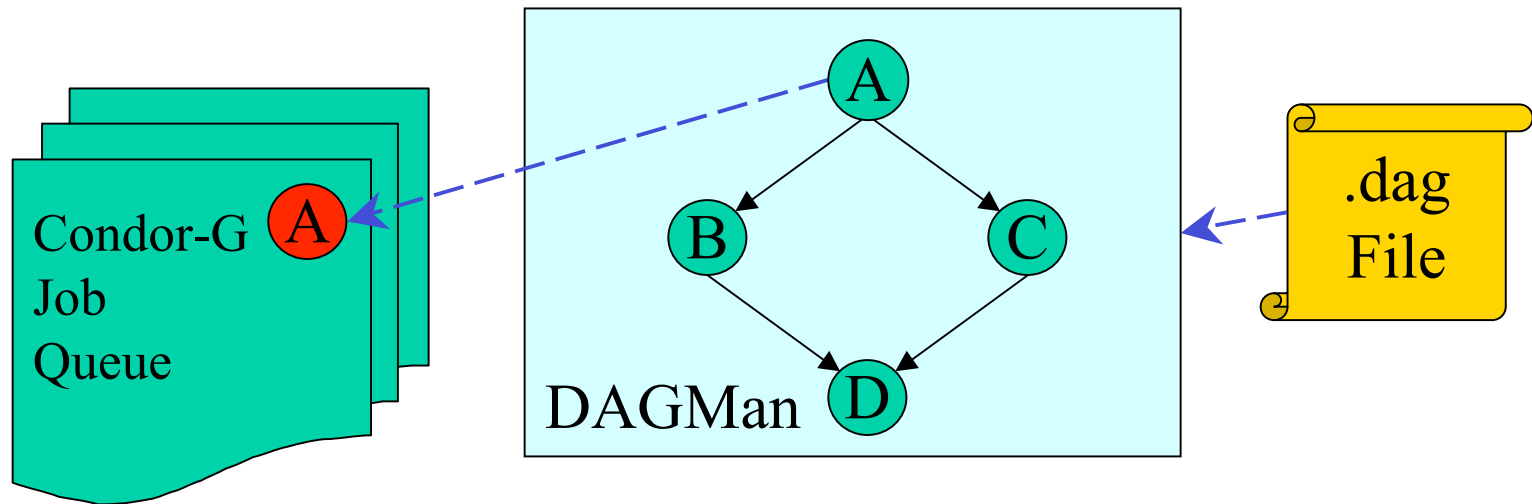
- To start your DAG, just run **condor\_submit\_dag** with your .dag file, and Condor will start a personal DAGMan daemon which to begin running your jobs:

```
% condor_submit_dag diamond.dag
```

- `condor_submit_dag` submits a Scheduler Universe Job with DAGMan as the executable.
- Thus the DAGMan daemon itself runs as a Condor job, so you don't have to baby-sit it.


# Running a DAG

- DAGMan acts as a “meta-scheduler”, managing the submission of your jobs to Condor-G based on the DAG dependencies.



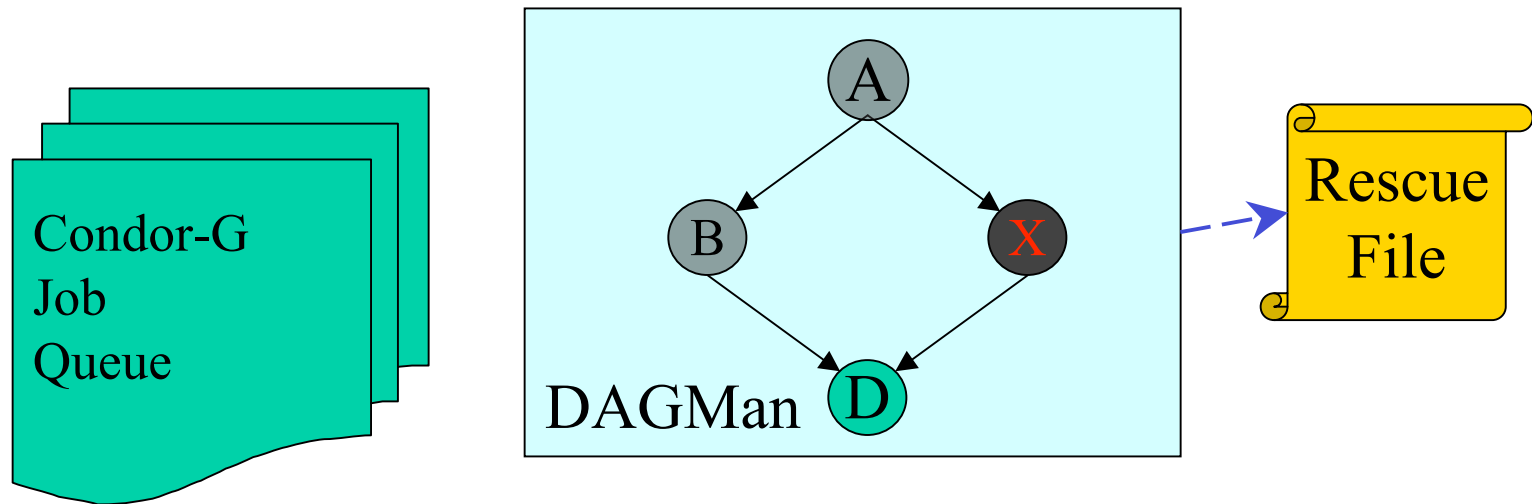
# Running a DAG (cont'd)

- DAGMan holds & submits jobs to the Condor-G queue at the appropriate times.



# Running a DAG (cont'd)

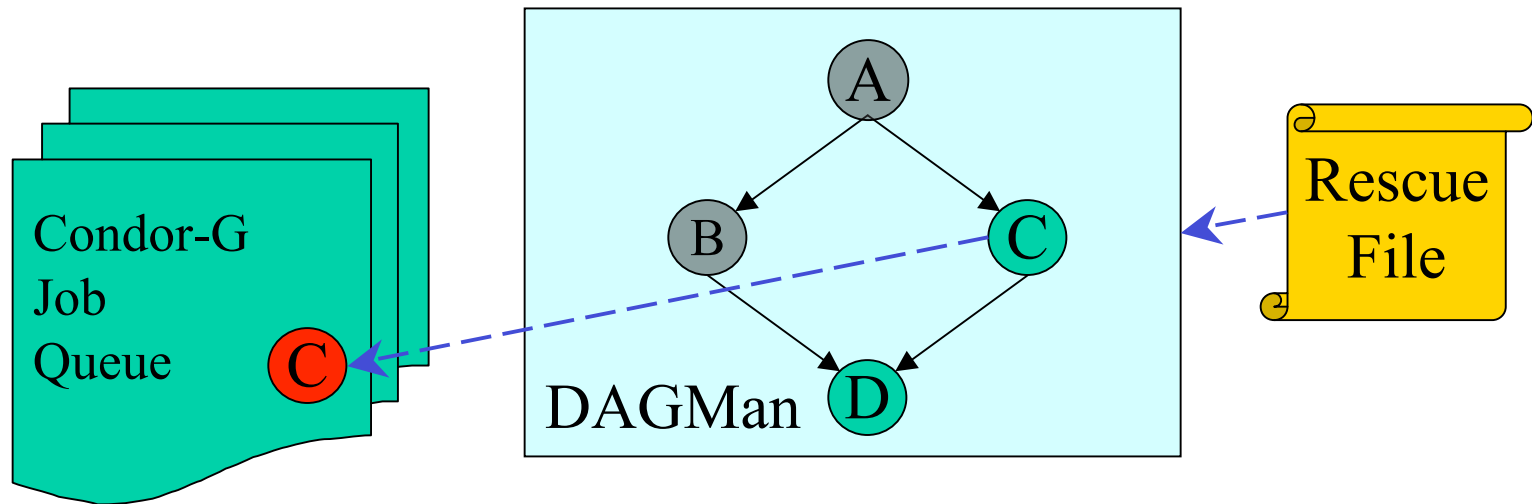
- In case of a job failure, DAGMan continues until it can no longer make progress, and then creates a *“rescue” file* with the current state of the DAG.



# Recovering a DAG


## -- fault tolerance

- Once the failed job is ready to be re-run, the rescue file can be used to restore the prior state of the DAG.





# Recovering a DAG (cont'd)

- Once that job completes, DAGMan will continue the DAG as if the failure never happened.



# Finishing a DAG

- Once the DAG is complete, the DAGMan job itself is finished, and exits.



## We have seen how Condor:

- ... monitors submitted jobs and reports progress
- ... implements your policy on the execution order of the jobs
- ... keeps a log of your job activities



.... Now go to the Lab part ....

## Acknowledgments:

---

This presentation based on:  
Grid Resources and Job Management



**Open Science Grid**

Jaime Frey and Becky  
Gietzel

Condor Project

U. Wisconsin-Madison

---