

2011 OSG Summer School

An introduction to
Distributed High-Throughput Computing
with emphasis on
Grid computing


by Igor Sfiligoi
University of California San Diego

Me and DHTC

- Working with distributed computing since 1996
- Working with Grids since 2005
- Leader of the OSG glidein factory ops since 2009
- Deeply involved in glidein development and deployments
- Mostly worked with HEP_(KLOE, CDF, CMS)

High Throughput Computing

- Alain yesterday introduced you to HTC
 - The concept of getting as many CPU cycles as possible over the long run
 - Based on batch job processing
 - No interactive access to resources



HTC in words

- As our esteemed Miron would put it

HTC is about extending the compute power of my own machine.

I **could** run my work on my own machine, but then it would take a very large number of calendar days/months/years to complete.

To finish the computation in a reasonable time, I have to expand the capacity of my own machine by obtaining and using temporary resources.

DHTC


Introduction to **Distributed High Throughput Computing**

What is DHTC

- So what is **Distributed** HTC???
- HTC is always distributed, right?
- What we mean here is
MASSIVELY distributed
 - i.e. **more than you can afford**
to host and operate **in one place**

Anatomy of DHTC

- So DHTC is about computing on **more than one HTC system**



Why is it different?

- Not a single system anymore
 - Most likely does not have a shared file system
 - Different clusters likely operated by different people
 - You may have a different account
 - Different clusters may use different technologies (e.g. Condor vs PBS)
- Usually harder to use than regular HTC
 - Just as using a HTC cluster is harder than using a single desktop

Why DHTC

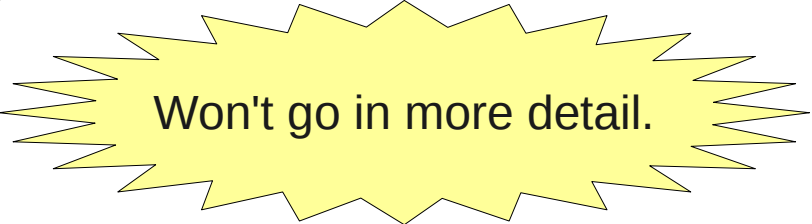
- Many reasons:
 - Practical
(a site has a limit to how much HW can host)
 - Political
(you only get money for HW if it is hosted at X)
 - Economical
(hosting and operating HW myself is too expensive)
(someone else can offer you hosted HW for less)
 - Opportunistic
(owners of site X have temporarily no jobs, might as well allow others to use them (for free or for pay))

DHTC in real life

- Campus-wide scheduling
 - e.g. U.Wisc Condor flocking
- Scientific Grids
 - e.g. OSG, Teragrid, EGI
- Hosted servers
 - e.g. Rackspace
- Cloud computing
 - e.g. Amazon EC2

Campus-wide scheduling

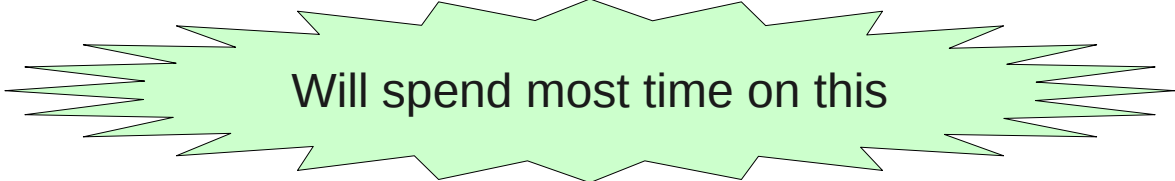
- Connects HTC clusters ran by different departments
- High trust between them
 - Sysadmins can talk to each other
 - Trust may be imposed from above (e.g. CIO)
- May use the same technology, which makes life much easier
 - Condor flocking an excellent example



Won't go in more detail.

Scientific Grids


- Widely distributed
 - OSG is US wide
- Moderate trust
 - Too many participants
 - OSG – $O(100)$ sites and $O(1k)$ users
- Many technologies
 - Joining sites may have existing infrastructure



Will spend most time on this

Hosted computing

- Commercial offering
- Lease a server for \$\$/month
- Similar to buying and hosting your own HW
 - Install whatever you want on them
 - But not in your LAN
- May be a good solution if you need a boost in capacity for a few months



Won't go in more detail.

Cloud computing

- A mix between hosted computing and Grids
- Job-based like a Grid
 - But “jobs” are Virtual Machines, not just processes
- You get your own machines like in hosted cmp.
 - They just happen to be VMs
 - You install whatever you want in them
 - There is an economic factor (although there is a push for scientific clouds as well)

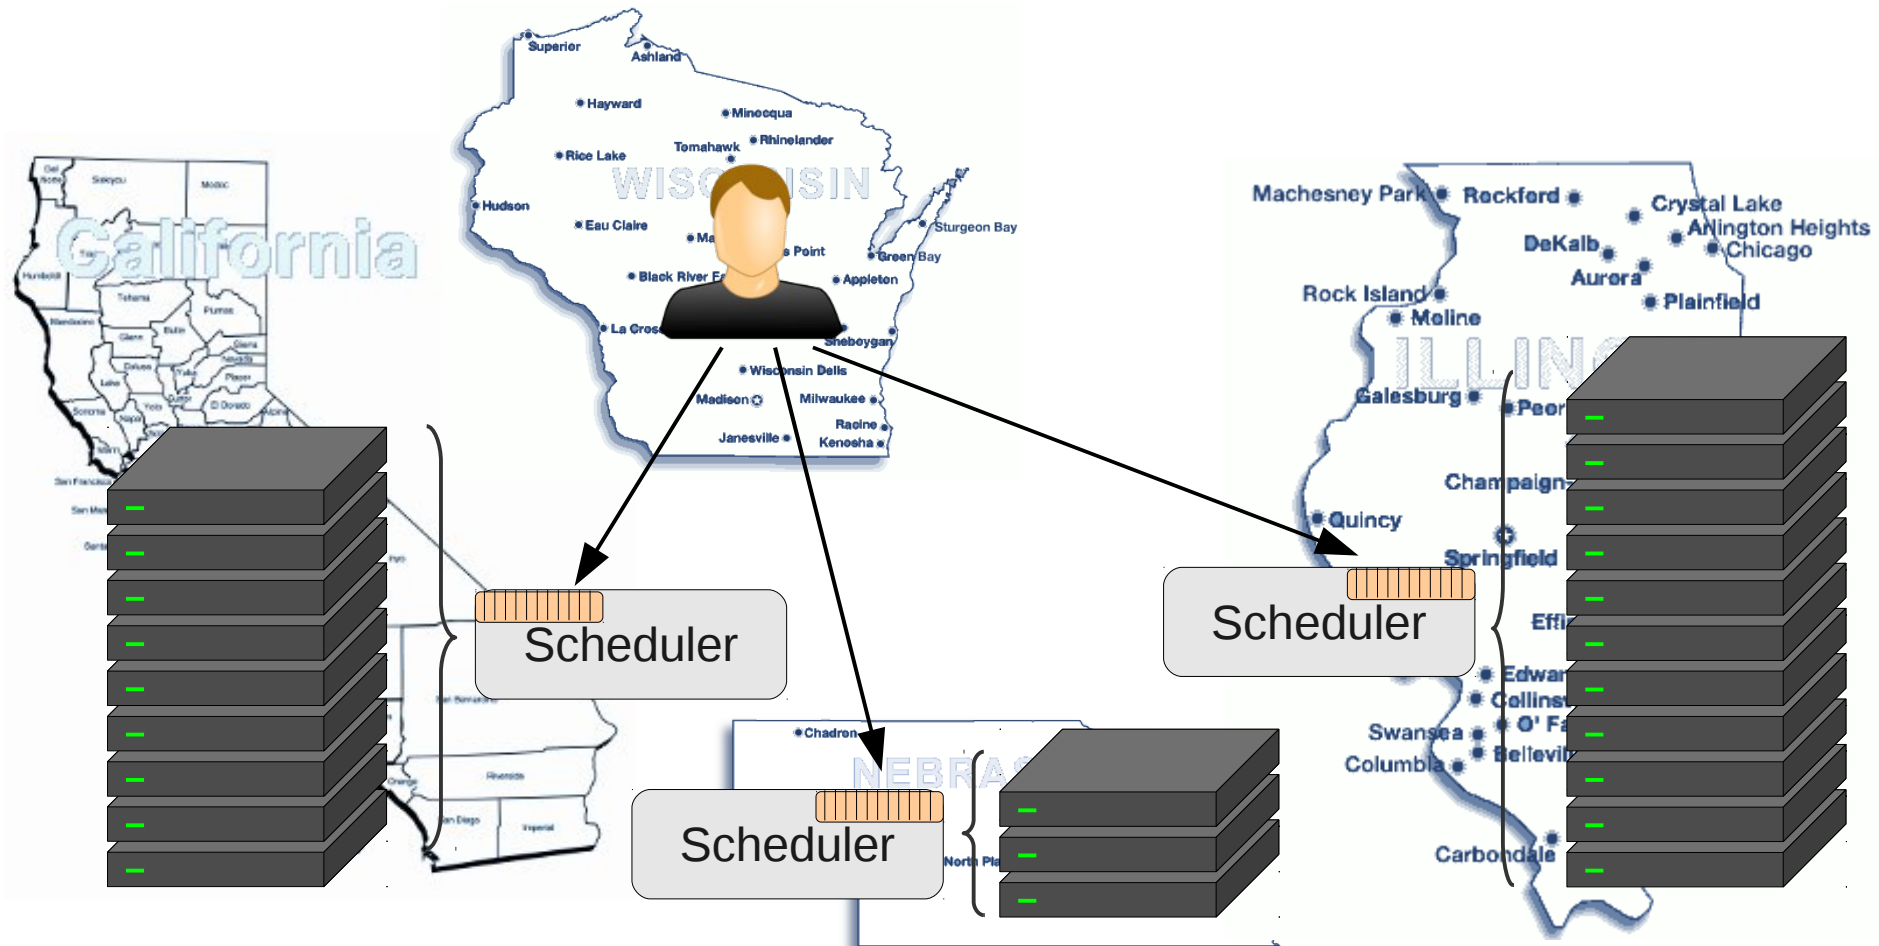
Won't go in more detail.

DHTC

Grid computing

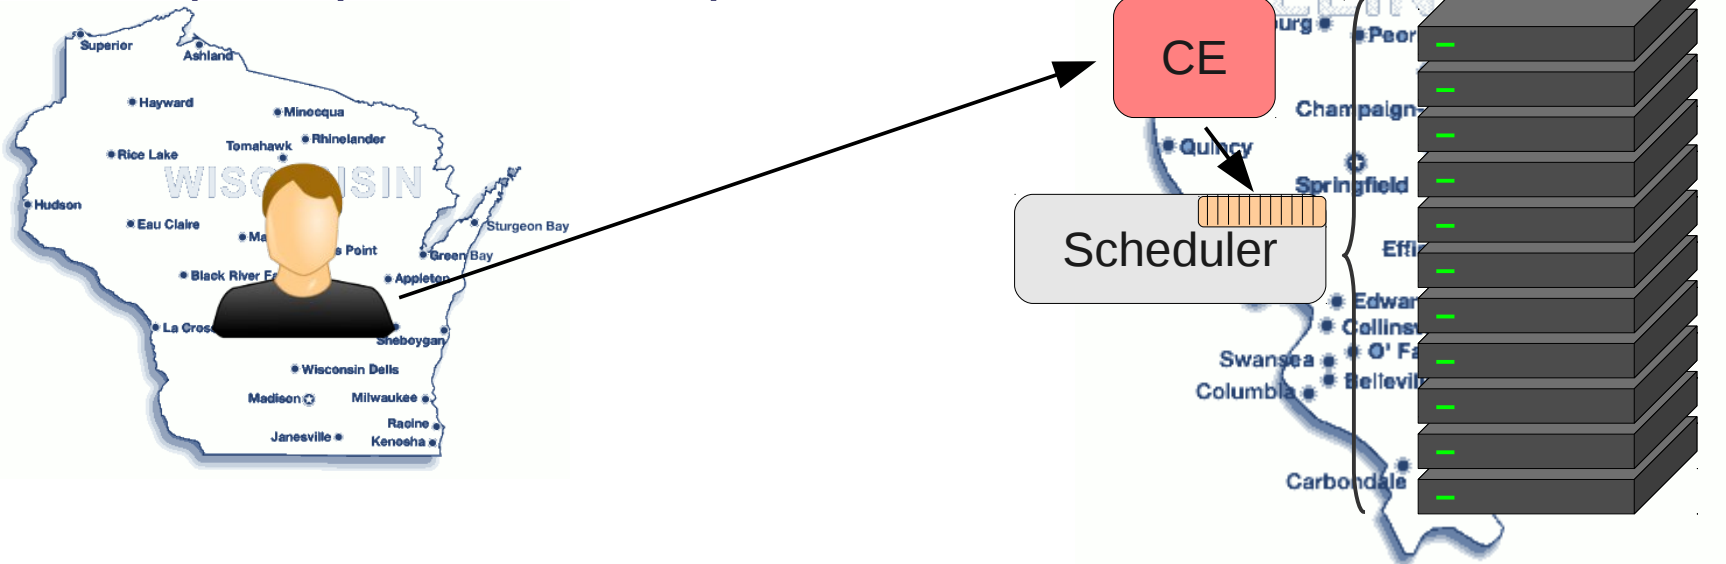
Grid computing

- Think of it as DHTC over WAN



Remote access

- Major difference compared to HTC is that we need remote access
 - A gatekeeper
 - Often called a CE (Compute Element)




Many options

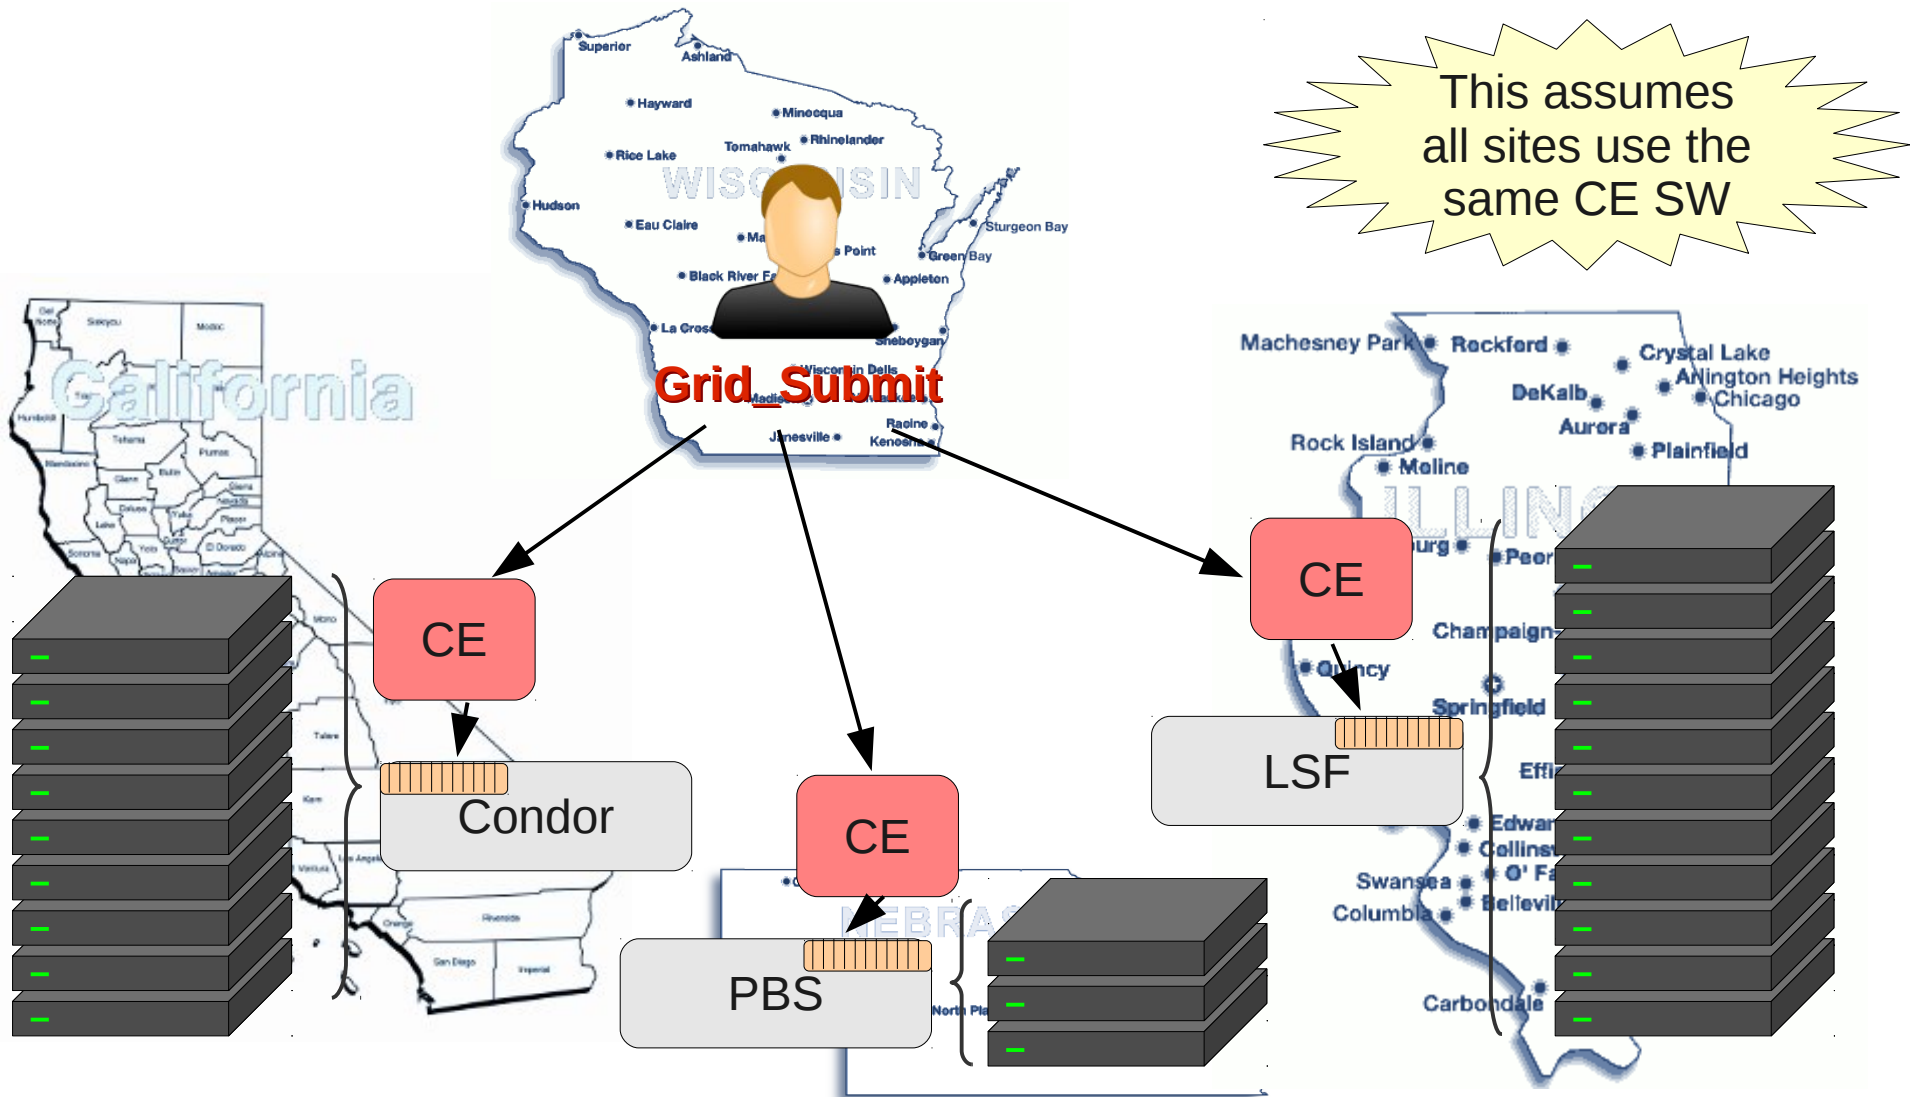
- Batch system native
 - e.g. Condor-C
 - SSH
 - Then locally `condor_submit`, `bsub`, ...
 - Grid gatekeepers
 - Globus
 - CREAM
 - ARC
- } Will hide site details

Grid gatekeepers

- Each site potentially uses a different local HTC system
- A Grid gatekeeper abstracts the API
 - Same remote calls for site-local Condor, PBS, ...
- Makes life easier for users
 - No need to know site details
- Used by most major Grid communities
 - e.g. OSG, Teragrid, EGI



Grid CE = Abstraction layer



Condor as a Grid client

- Condor can be used as a “universal client”
- Can submit jobs to remote HTC systems
 - Using the “Grid” universe
 - Supports most CE APIs
- Same job commands as for Condor HTC
 - `condor_submit`, `condor_q`, etc.
- But very different under the hood!

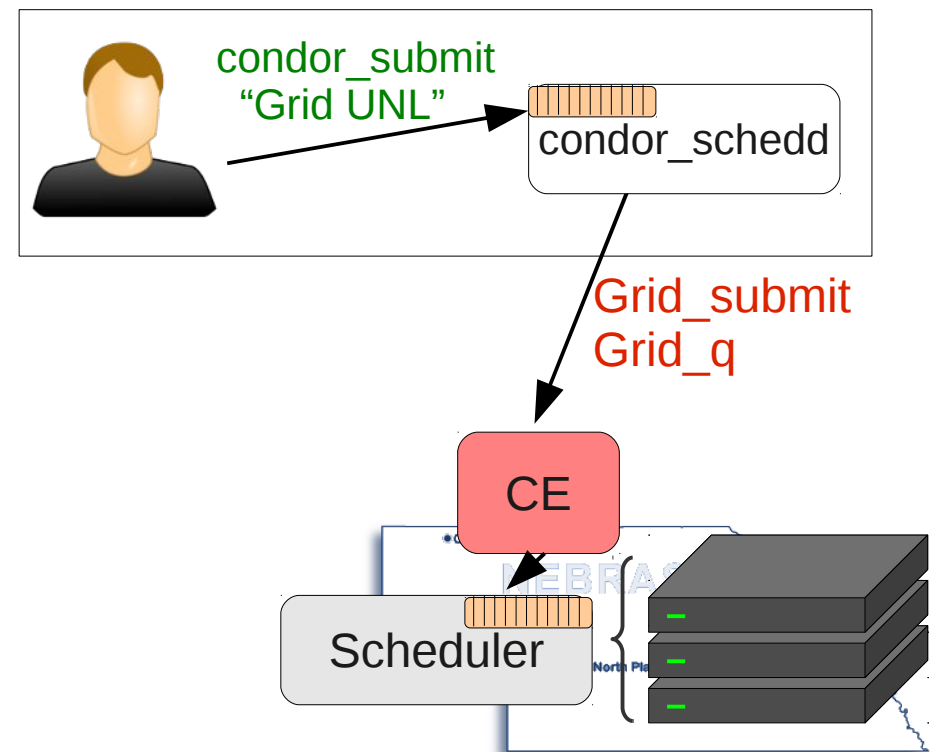
Known as
Condor-G

See next slide

Condor-G

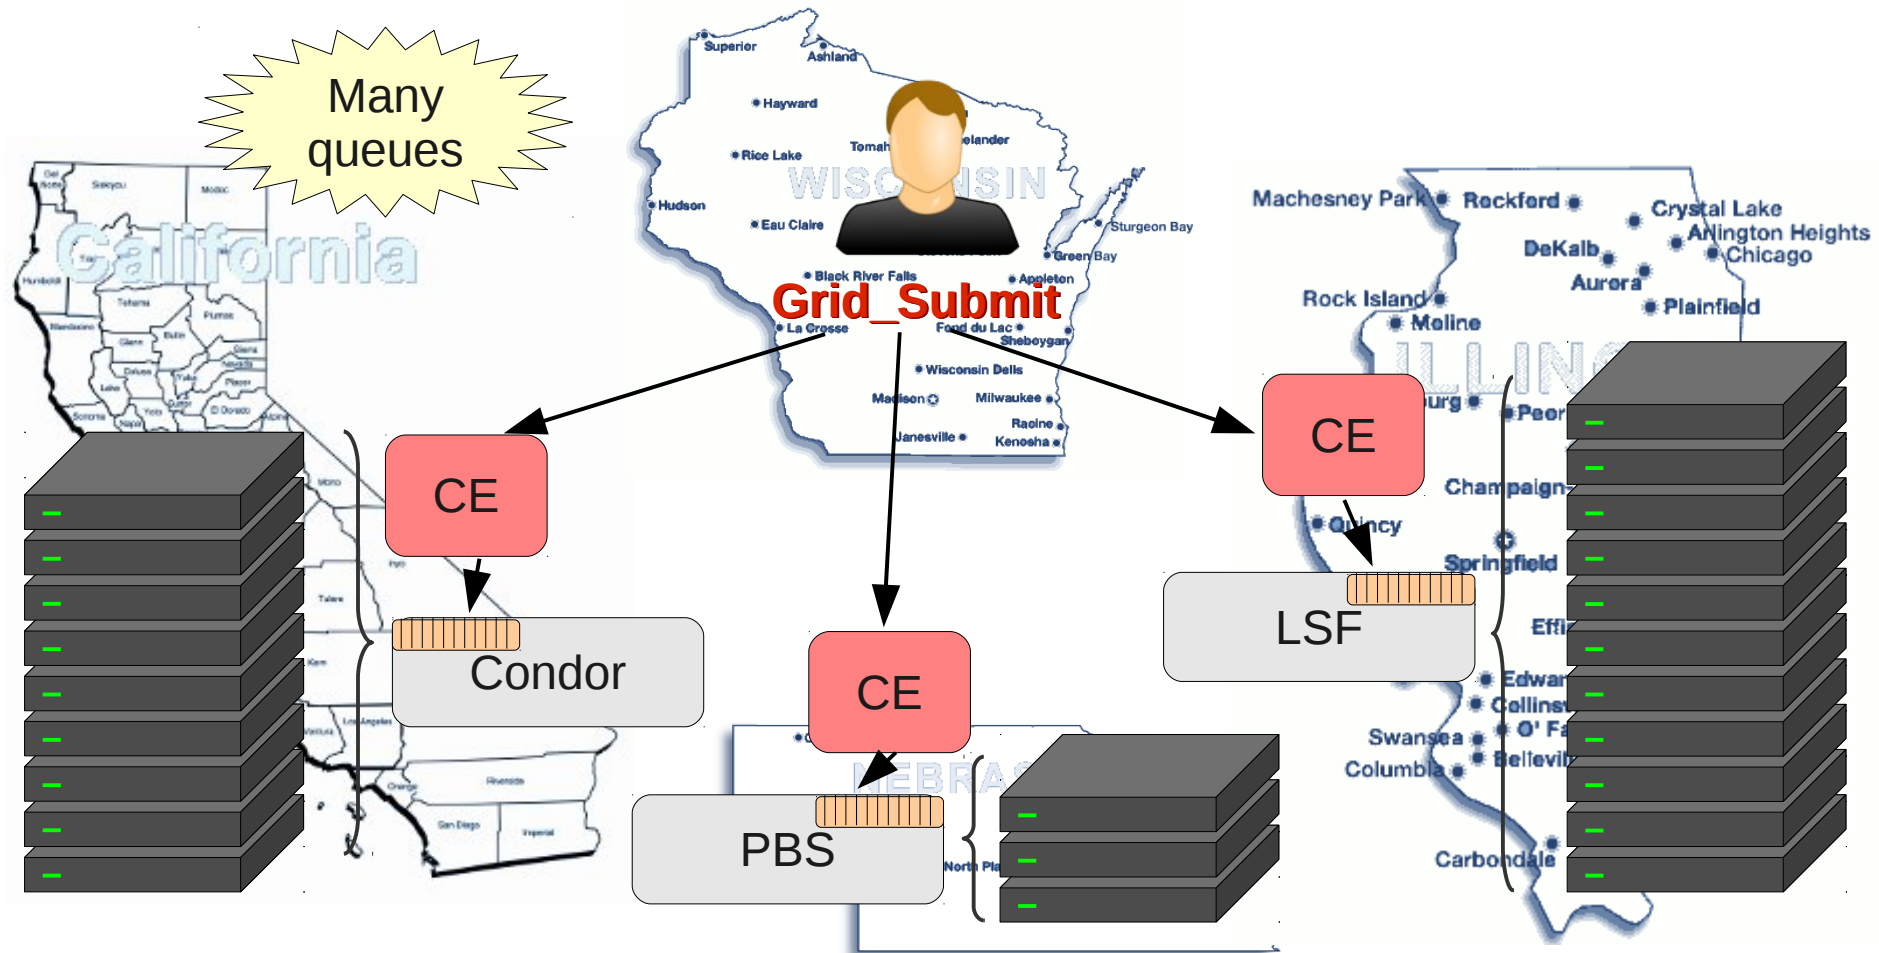
- Condor-G does not manage remote resources
 - It just forwards and monitors jobs to remote HTC sys
 - No condor_status
- No matchmaking
 - User explicitly specifies API and site to use

Limited resource selection available with external support




Job partitioning problem

- Pure Grid computing requires job partitioning



Job partitioning = Hard problem

- Job partitioning is a hard problem
- Especially in the Grid
 - Many different technologies (e.g. Condor vs PBS)
 - Owned by many different admins
 - With an abstracted API in between
- Some automation highly desirable
 - gliteWMS
 - OSG MM



No really good solutions

Get your hands dirty

- This is all the theory you need to know for now
- Exercise time
- Feel free to ask question