

How to get the needed computing

Tuesday morning, 11:15am

Igor Sfiligoi <isfiligoi@ucsd.edu>

Leader of the OSG Glidein Factory Operations

University of California San Diego




Logistical reminder

- It is OK to ask questions
 - During the lecture
 - During the demos
 - During the exercises
 - During the breaks
- If I don't know the answer,
I will find someone who likely does

Reminder - (D)HTC

- (D)HTC is about getting as many CPU cycles as possible over the long run
 - No matter where they are located



But where do you get the CPUs?

- You likely have a desktop/laptop
- But that's normally not enough
 - Or you wouldn't be here!



So....

How do I get access
to the needed computing?

Getting the needed CPUs

You either:

- Buy it
 - Get a computing grant
 - Use leftover CPU cycles

Getting the needed CPUs

You either:

- Buy it
 - Your own hardware (e.g. a cluster)
 - Contribute hardware to a common pool
 - Lease machines on monthly/yearly basis
 - Rent machines on a per hour basis
- Get a computing grant
- Use leftover CPU cycles
 - On friends' hardware
 - At your home institution
 - On a large-scale scientific infrastructure

Getting the needed CPUs

You either:

- Buy it
 - Your own hardware (e.g. a cluster)
 - Contribute hardware to a common pool
 - Lease machines on monthly/yearly basis
 - Rent machines on a per hour basis
- Get a computing grant
- Use leftover CPU cycles
 - On friends' hardware
 - At your home institution
 - On a large-scale scientific infrastructure

Which way to go?

- Buying is more reliable
 - You can plan on how much computing you will be able to do
 - But not always an option
(We all have a limited budget)
- Grants are similar in nature
 - But may be harder to get than money!
- Opportunistic use can give you vastly more
 - But you will have to be flexible
 - And there is no guarantee you will get what you hope for

Don't assume flat usage

- While HTC paradigm postulates maximizing results out of available resources
- Most people really have spiky compute needs, e.g.
 - You come up with a great idea... need a gazillion CPUs **now** to verify it! Then nothing for a month or more while you look for the next great one.
 - Everybody wants to run a last computing pass just before that important conference

Impact on planning

- You likely cannot afford to buy enough dedicated resources to cover the peaks
- You should **always** plan on using leased resources
 - At least for part of the year
 - i.e. make sure you choose tools that allow you to do DHTC

Buying dedicated hardware

- Buying your own hardware is the most straightforward approach
- But you must also budget for
 - Floor space
 - Electricity
 - Cooling
 - Personnel costs

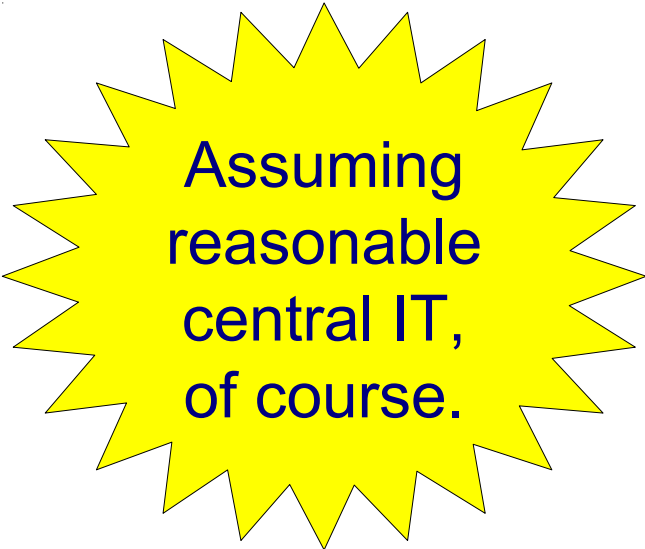
} Often more than the HW itself
- And install your own HTC system
 - Which requires significant expertise

Server hosting

- You can lease hardware from commercial entities
 - Typically 3-12 months leases
 - Popular vendor: Rackspace
- Great for shorter projects
 - Likely gives you lower Total Cost of Ownership
 - May cost less than buying the HW itself
 - And you save on the infrastructure costs
 - But still requires you to operate an HTC system

Contributing to a common pool

- If you have an existing HTC system at your institution/campus, it **may be cheaper** to contribute to it
 - Economies of scale
 - Better expertise
- You will have to trade some flexibility
 - But often it is well worth it



Assuming
reasonable
central IT,
of course.

Contributing to a common pool

- If you have a system

The HTC system
may also be located
at a different location,
if they allow remote access.

Usually called
Grid computing

Computing Grants

- Some of the US large-scale computing infrastructures are grant based
 - e.g. XSEDE
- You make a proposal, and if you make a good case, you can get a substantial CPU allocation on their HTC system
- Initial effort comparable to buying HW
 - A lot of paperwork needed
 - And long lead times

Cloud computing

- Similar to server hosting
 - But you pay by the hour
- Most famous is Amazon EC2
(but not the only one)
- Great for spike leveling
 - Can get a lot of resources on short notice, if you have the needed money
 - But can be quite expensive

Cloud computing

- Similar to server hosting

You may hear talking
about “Scientific Clouds”.

That's really just another name for
VM-based HTC systems

(so **not** what I call Cloud here)

Going for opportunistic resources

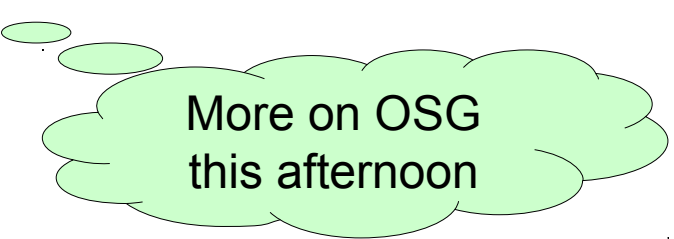
- All the methods described so far give you dedicated resources
 - i.e. you can count on them being there when you need them
(at least, after the initial setup period)
 - But that may not be enough for your needs
 - And you cannot afford more
- Opportunistic resources may come to the rescue
 - Just remember there are **no guarantees** here

Opportunistic resources

- The opportunistic resources are essentially machines that are **currently** not needed by the owner
 - e.g. owners are in the low part of a spike
- Reasons why they may allow you to use them (instead of turning them off)
 - Connections – e.g. Friend's desktop
 - Politics – e.g. Funding agency requirement
 - Money – e.g. Amazon's spot instances

Free HTC resources

- You should first look close by
 - The HTC system at your home institution may frequently have spare capacity
 - They are likely happy to share
- Once that is not enough, remote HTC clusters may provide substantial additional free resources as well
 - e.g. Sites on OSG



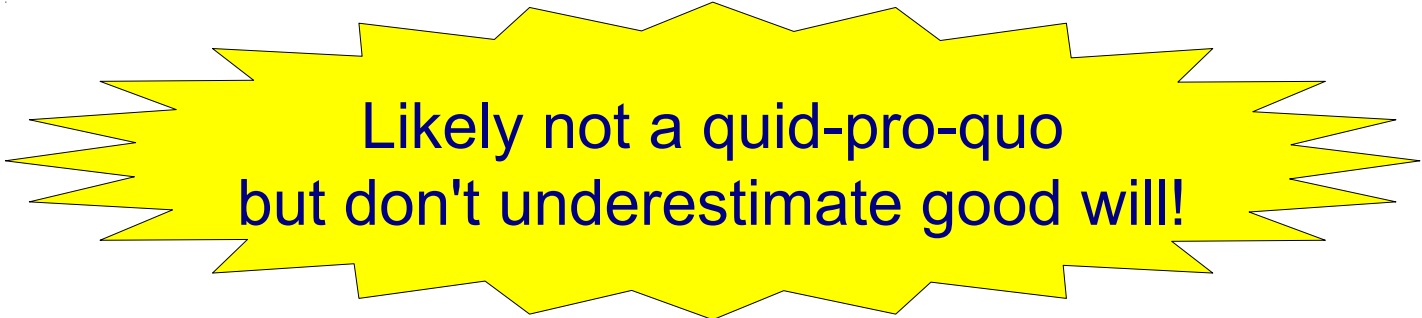
More on OSG
this afternoon

Beggars can't be choosers

- Remember, when you scout for free resources, you have very few rights
 - You are effectively “a beggar”
- You will have to adapt
 - Don't expect to get it your way
 - Even though some sysadmins may be willing to help
 - The more flexible you are, the more free resources you will be able to use

Consider contributing back

- If you buy hardware, consider contributing back
 - i.e. slightly overprovision your cluster and give unused CPU cycles to others
- A modest amount of gifted CPU can buy you a lot of good will
 - For when you need to level your own spikes




Likely not a quid-pro-quo
but don't underestimate good will!

Opportunistic Cloud Resources

- Here opportunistic means **cheap(er)**
 - Not free
- Basically one bids (real money) for resources
 - The highest bidder gets the resources
 - Until someone else bids higher
 - e.g. Amazon EC2 Spot Instances
- Can get a lot for relatively little money
 - But not always

Example

- For Amazon EC2 m1.xlarge
 - Full price: \$0.5/hour
 - Min spot price: \$0.05/hour
 - Many periods when you could get that price



Putting everything together

- You will get most work done if you mix owned, leased and opportunistic resources
 - i.e. DHTC
- Plan on using an overlay system early on
 - So you don't have to use N different tools to use N different clusters
 - Searching for/Learning about new tools when you are close to a deadline is no fun

Questions?

- Questions? Comments?
 - Feel free to ask me questions later:
Igor Sfiligoi <isfiligoi@ucsd.edu>
- Upcoming sessions
 - Now – noon
 - Hands-on and demo
 - noon – 12:40pm
 - Tour
 - 12:40pm – 1:30pm
 - Lunch



Grids and Clouds

