

# More HTCondor

2013 OSG User School, Monday, Lecture 2

**Tim Cartwright**  
`cat@cs.wisc.edu`

University of Wisconsin–Madison  
OSG Software Team Manager  
OSG Education Coordinator



# Questions so far?

# Goals For This Session

---

- Understand the mechanisms of HTCondor (and HTC in general) a bit more deeply
- Use a few more HTCondor features
- Run more (and more complex) jobs at once



# HTCondor in Depth

# Why Is HTC Difficult?

- System must track jobs, machines, policy, ...
- System must recover gracefully from failures
- Try to use all available resources, all the time
- Lots of variety in users, machines, networks, ...
- Sharing is hard (e.g., policy, security)
- More about the principles of HTC on Thursday

# Main Parts of HTCondor

---

# Main Parts of HTCondor

## Function

Track waiting/running jobs

Track available machines

Match jobs and machines

Manage one machine

Manage one job (on submitter)

Manage one job (on machine)

# Main Parts of HTCondor

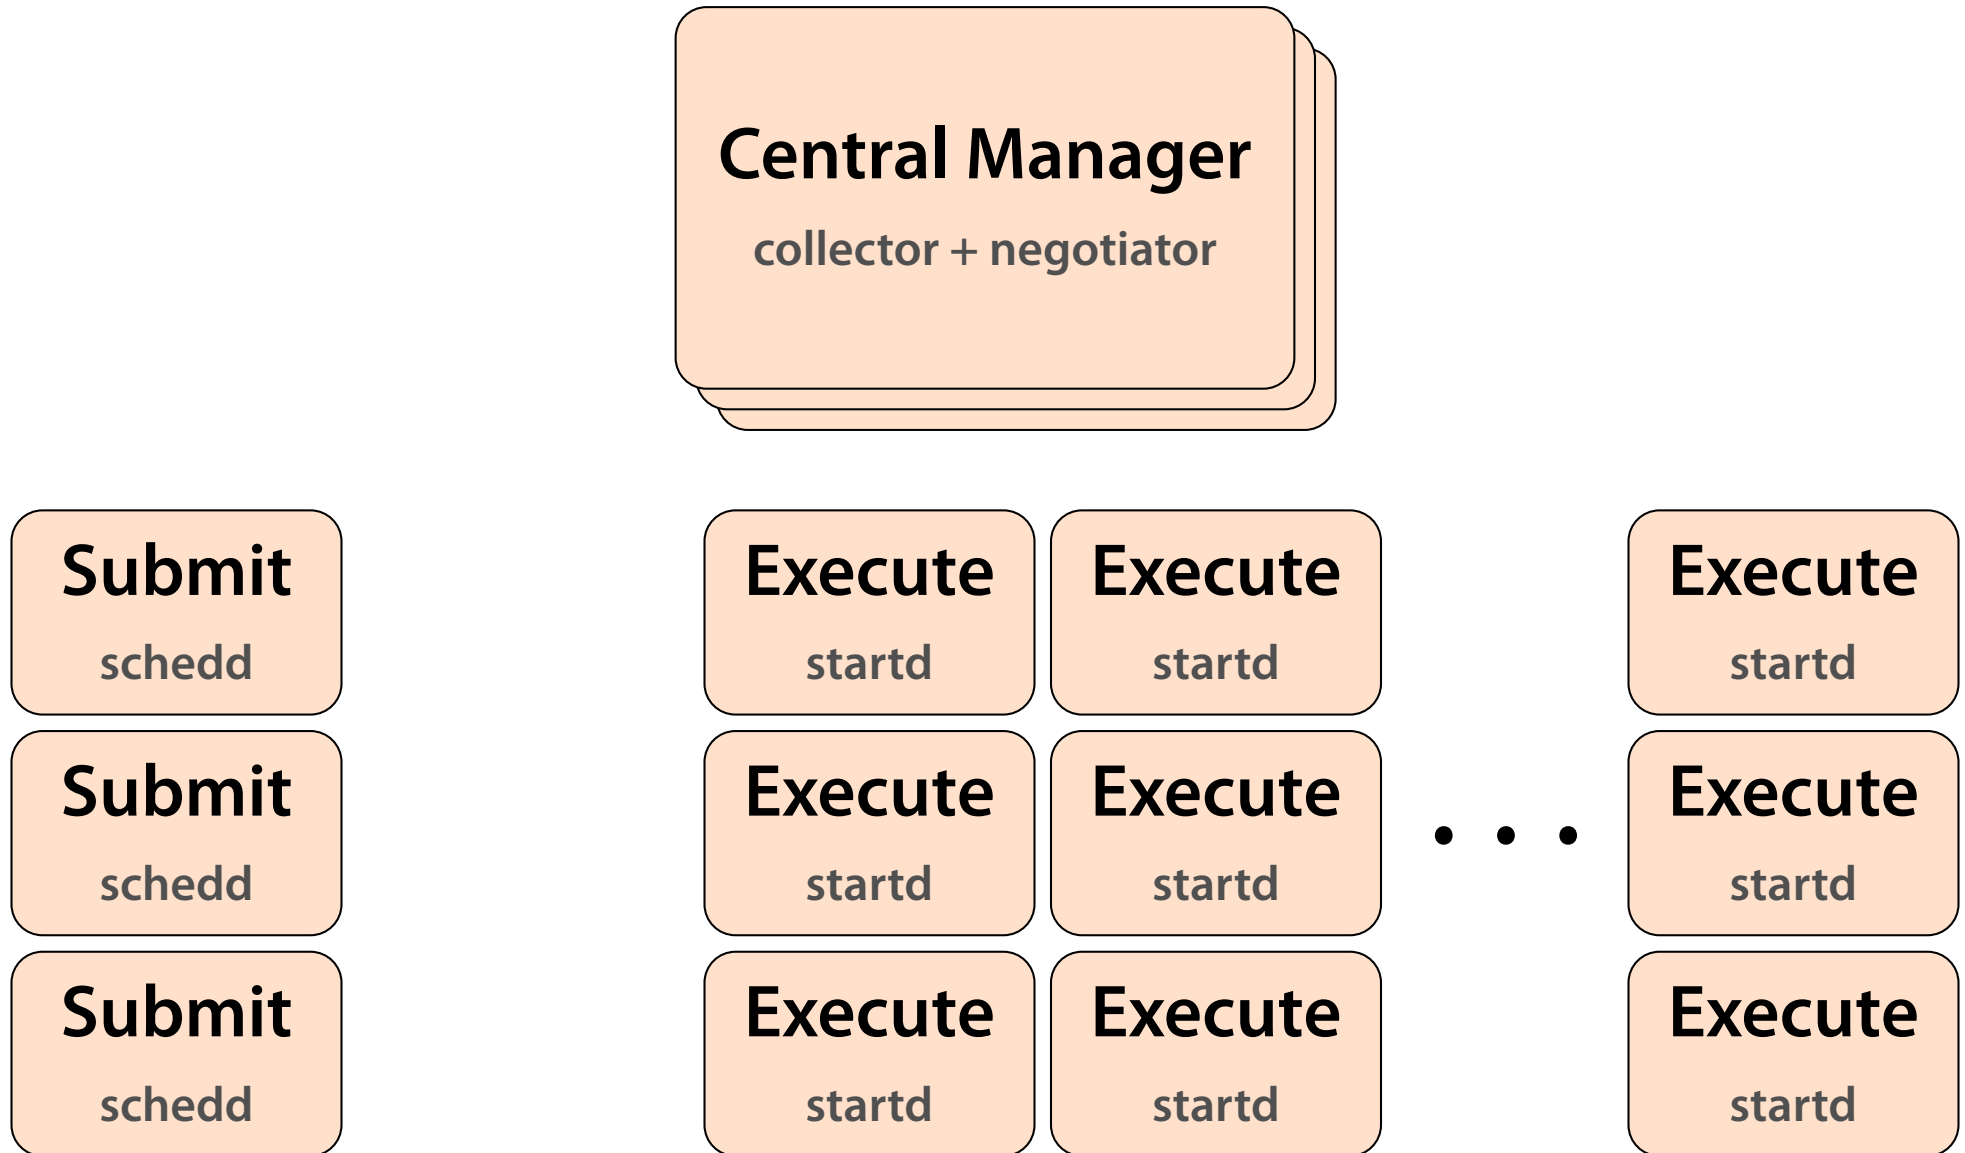
Function	HTCondor Name
Track waiting/running jobs	schedd ("sked-dee")
Track available machines	collector
Match jobs and machines	negotiator
Manage one machine	startd ("start-dee")
Manage one job (on submitter)	shadow
Manage one job (on machine)	starter



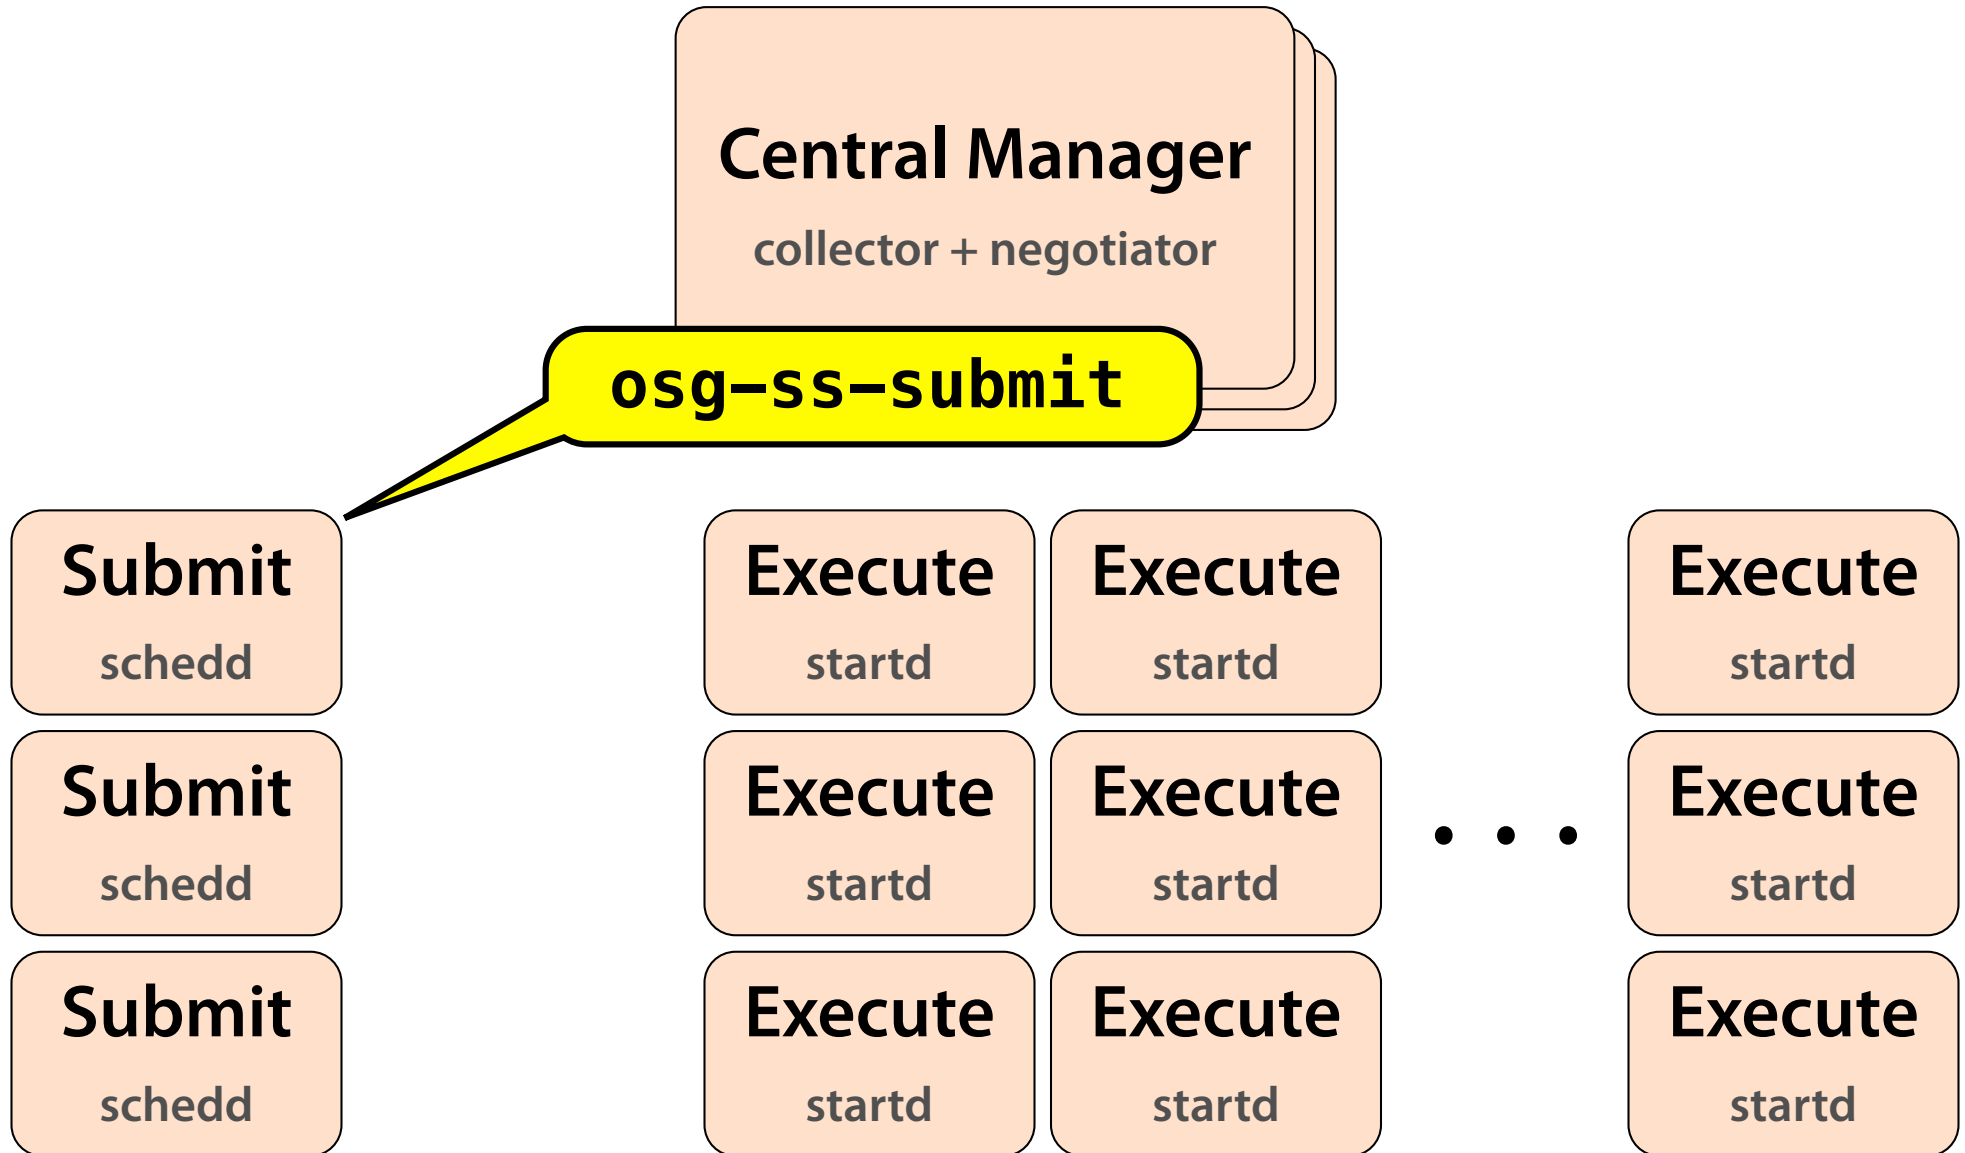
# Main Parts of HTCondor

Function	HTCondor Name	#
Track waiting/running jobs	schedd ("sked-dee")	1+
Track available machines	collector	1
Match jobs and machines	negotiator	1
Manage one machine	startd ("start-dee")	per machine
Manage one job (on submitter)	shadow	per job running
Manage one job (on machine)	starter	per job running

# Typical Architecture



# Typical Architecture



# Typical Architecture

`cm.chtc.wisc.edu`

**Central Manager**  
collector + negotiator

**Submit**

schedd

**Submit**

schedd

**Submit**

schedd

**Execute**

startd

**Execute**

startd

**Execute**

startd

**Execute**

startd

**Execute**

startd

**Execute**

startd

...

**Execute**

startd

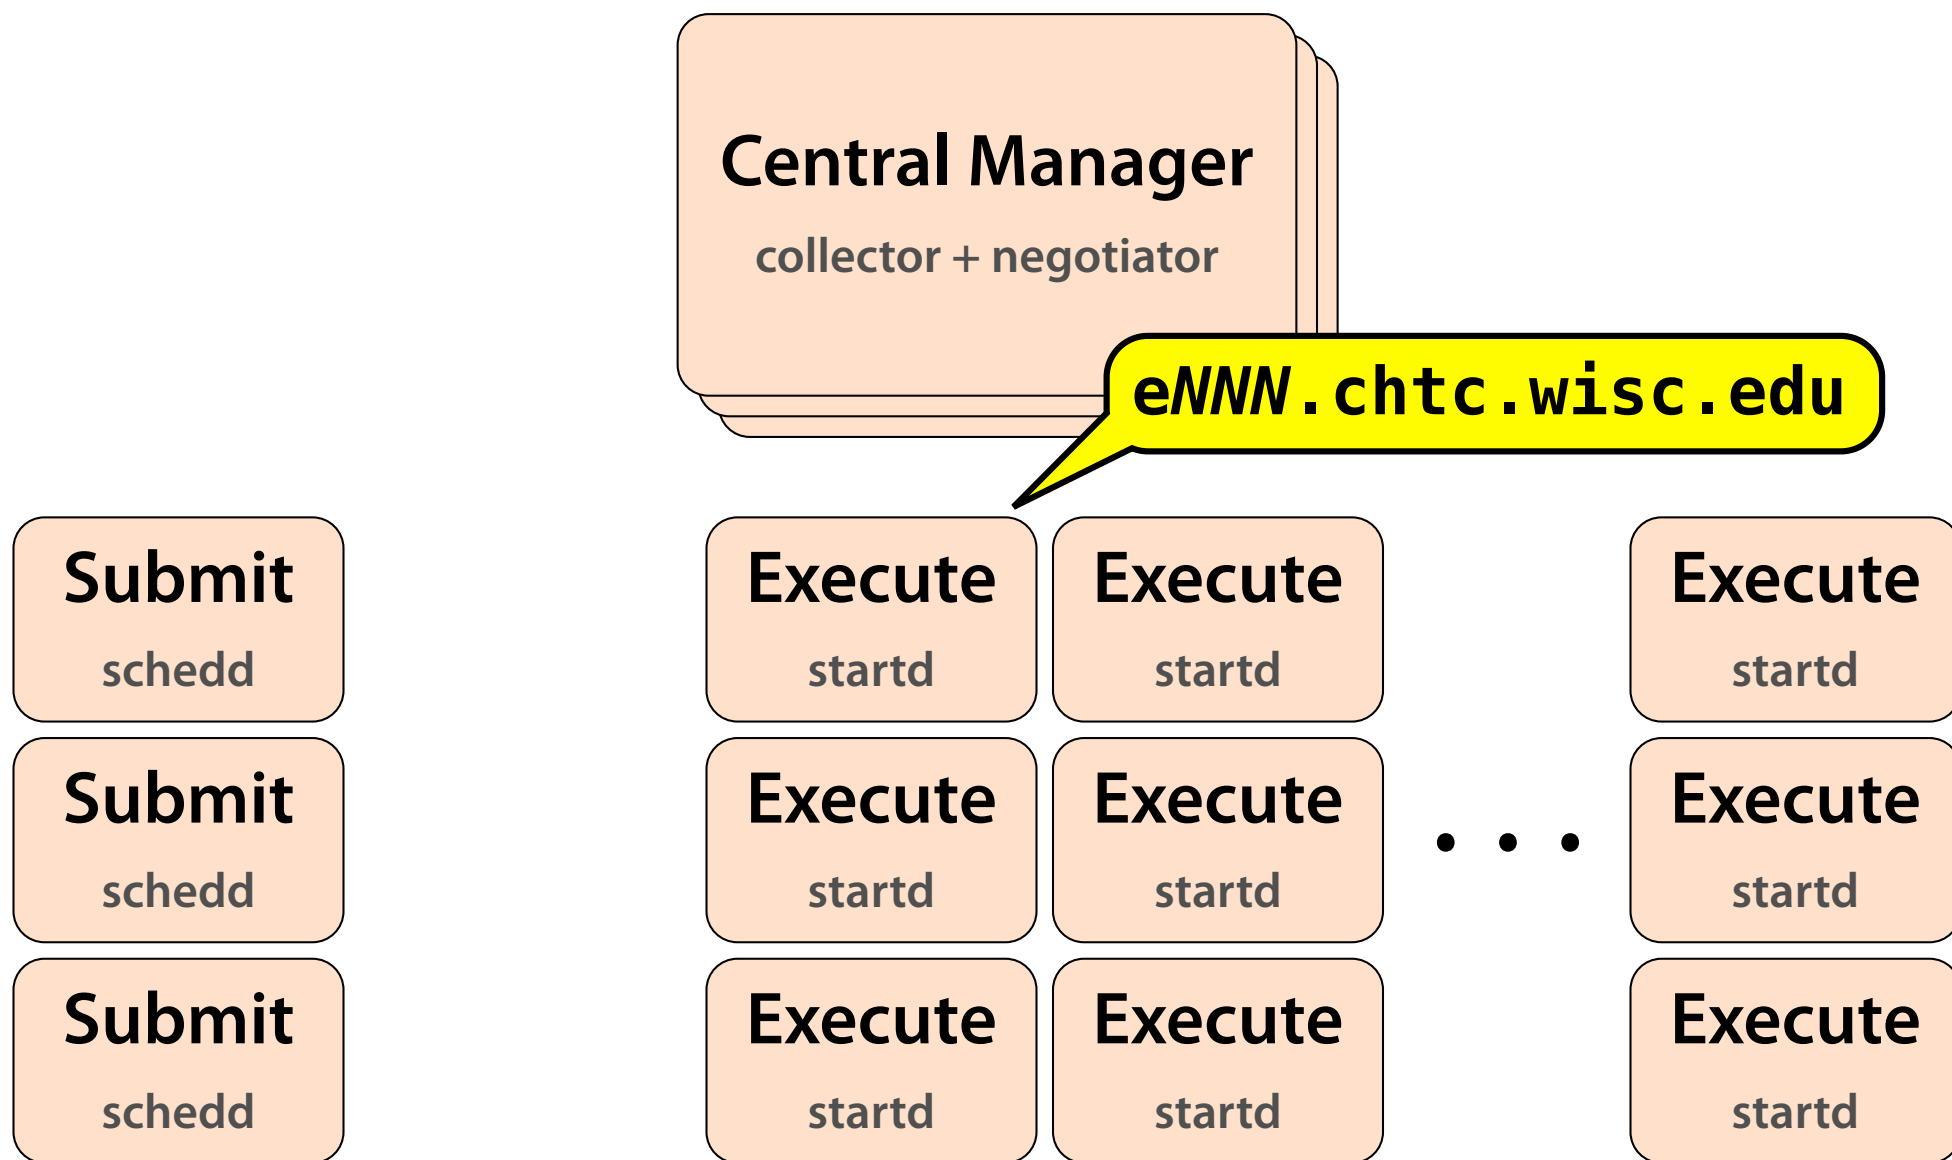
**Execute**

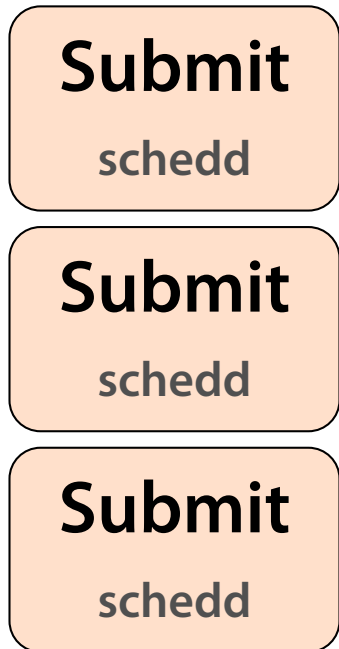
startd

**Execute**

startd

# Typical Architecture





# The Life of an HTCondor Job

## *Central Manager*

negotiator


collector

schedd

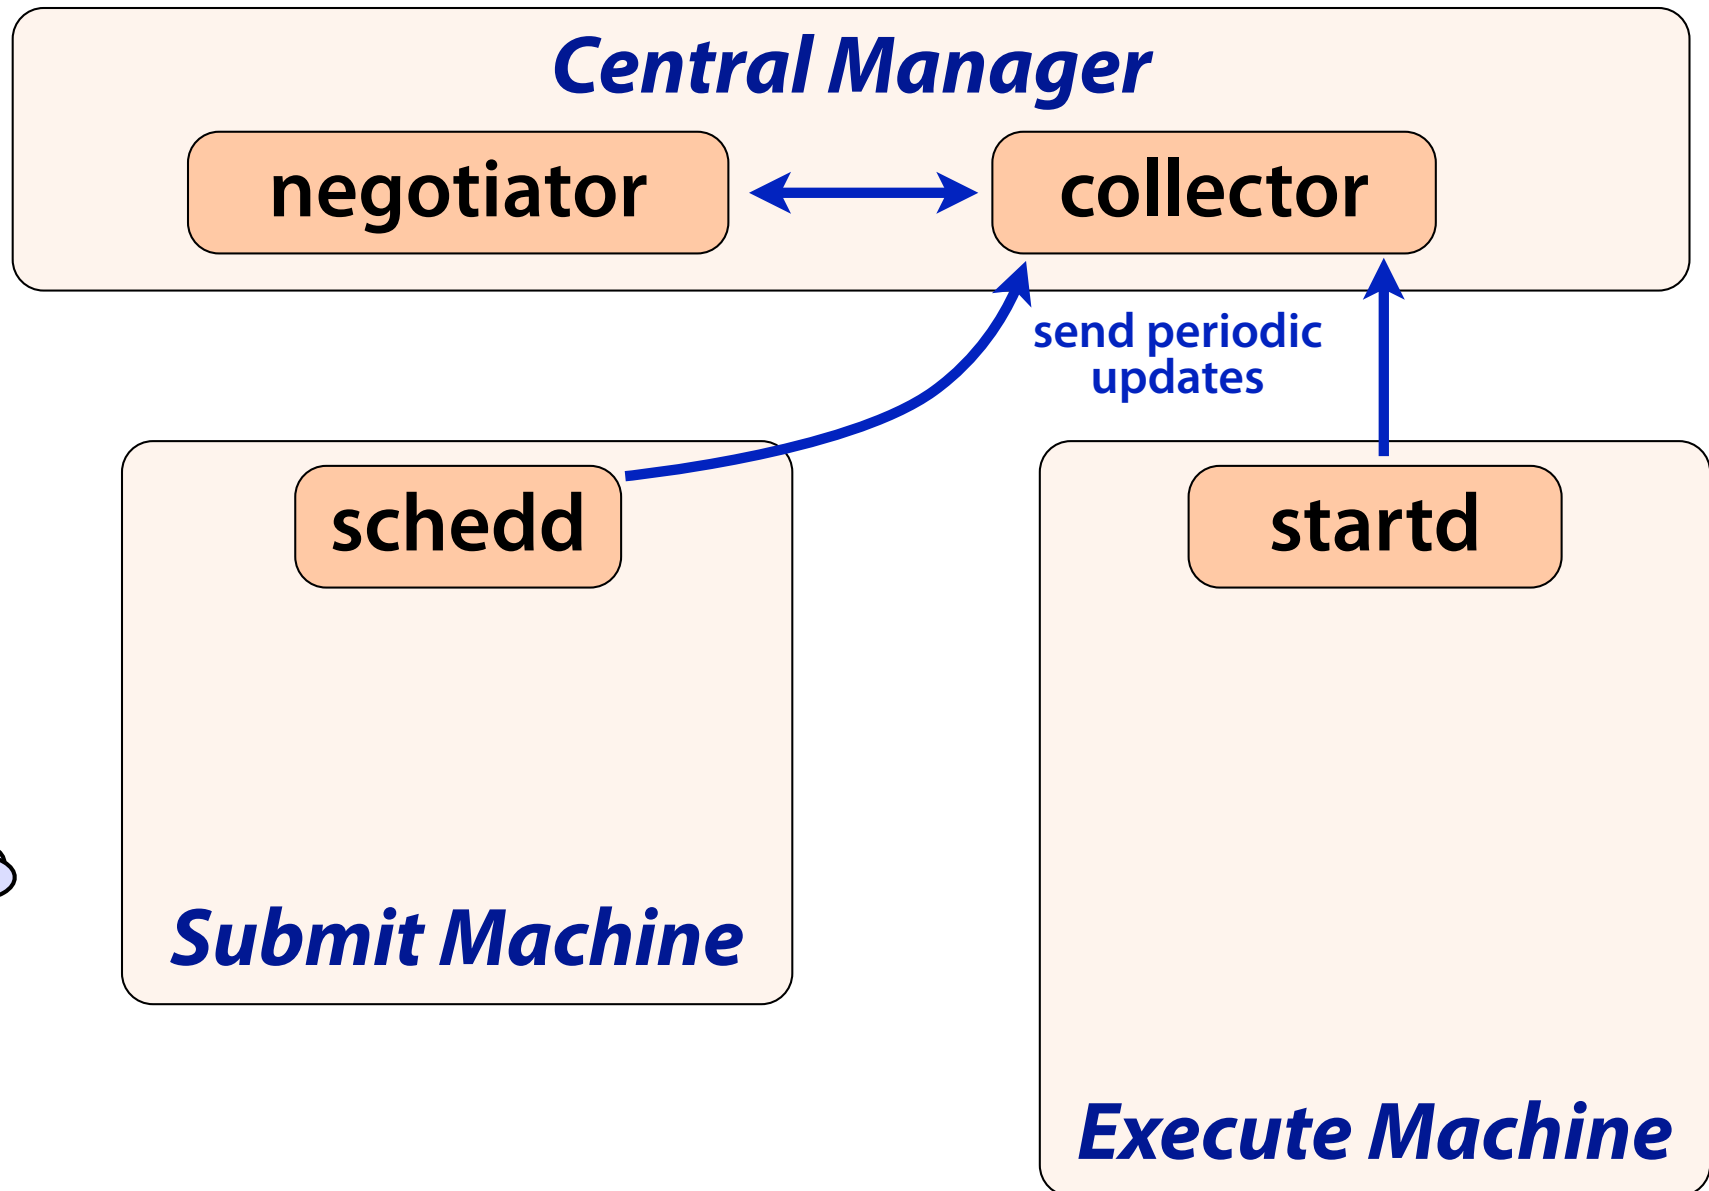

*Submit Machine*

startd

*Execute Machine*

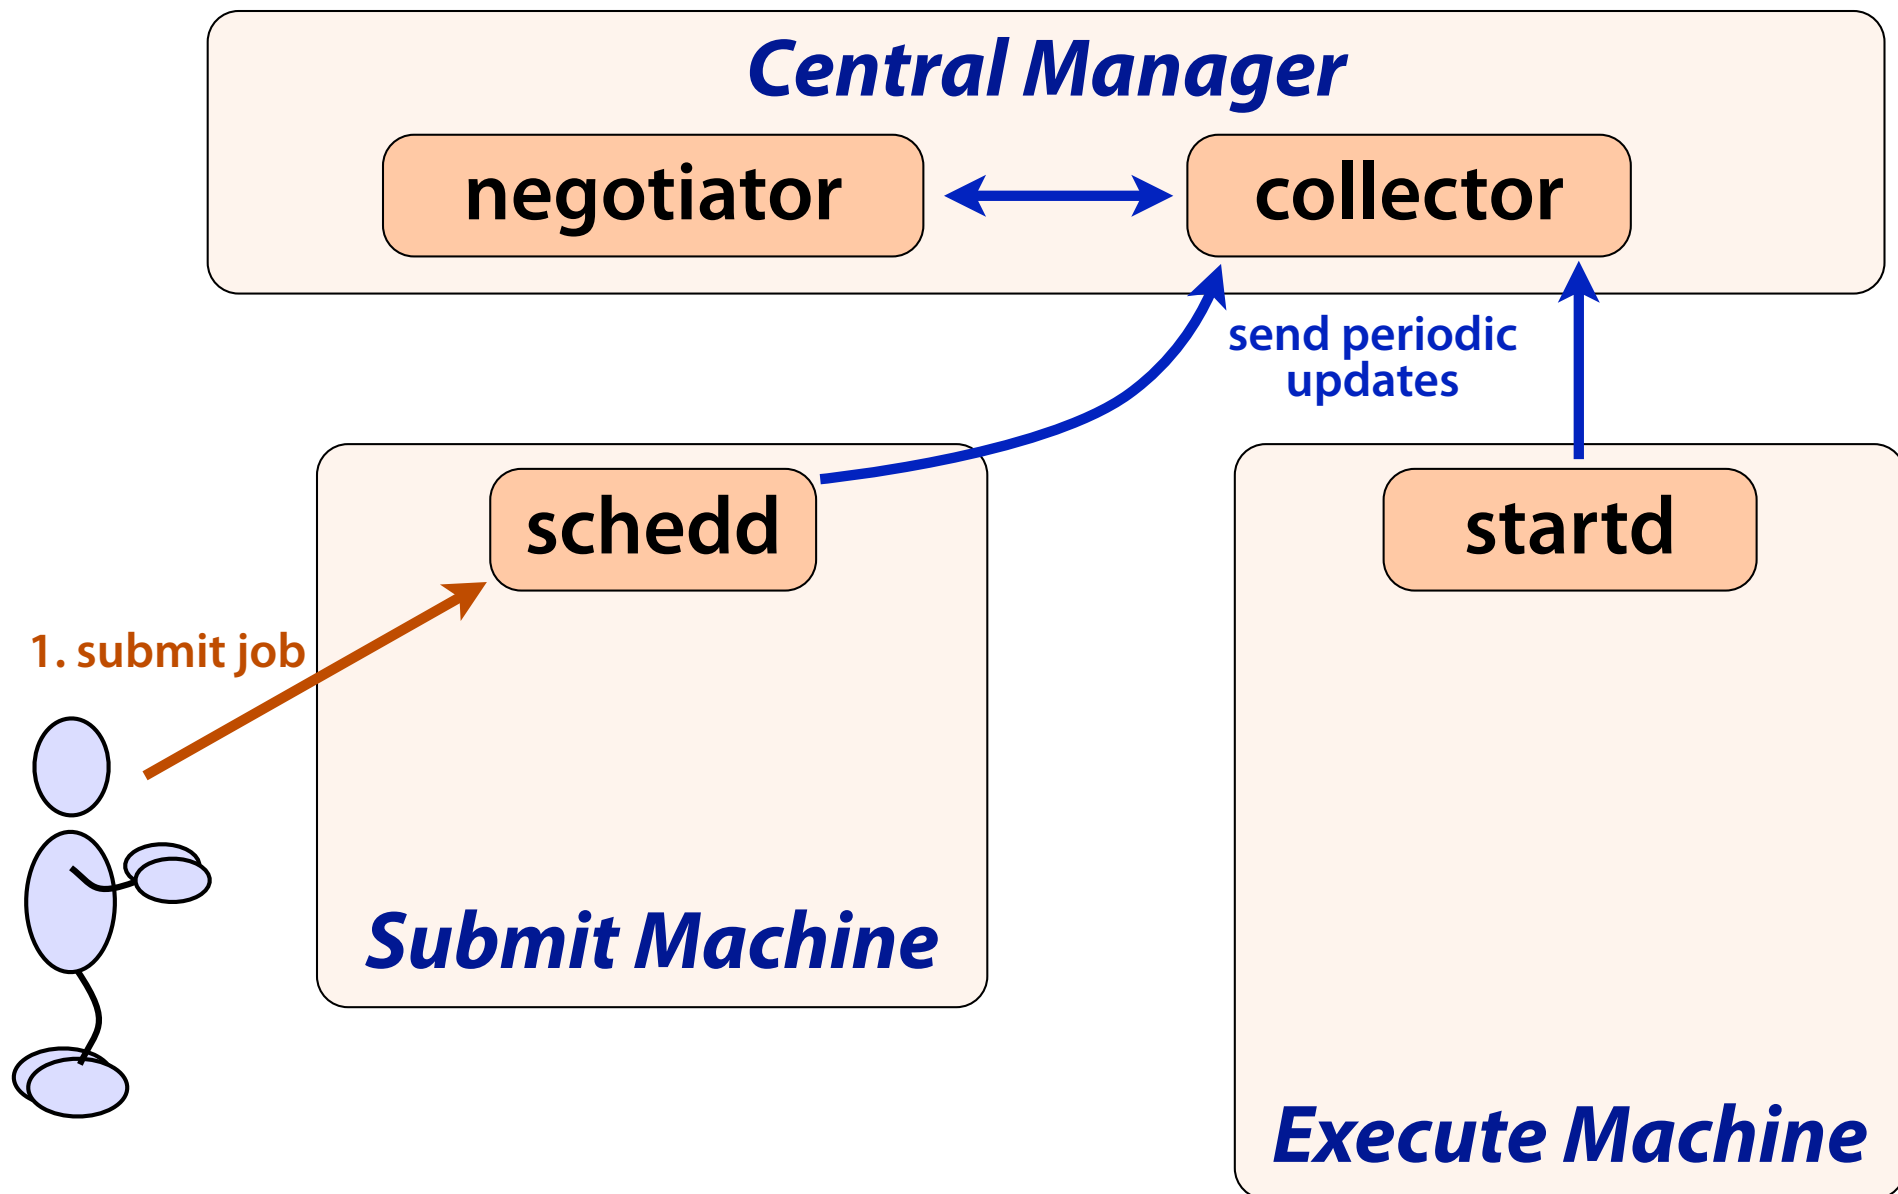


# The Life of an HTCondor Job

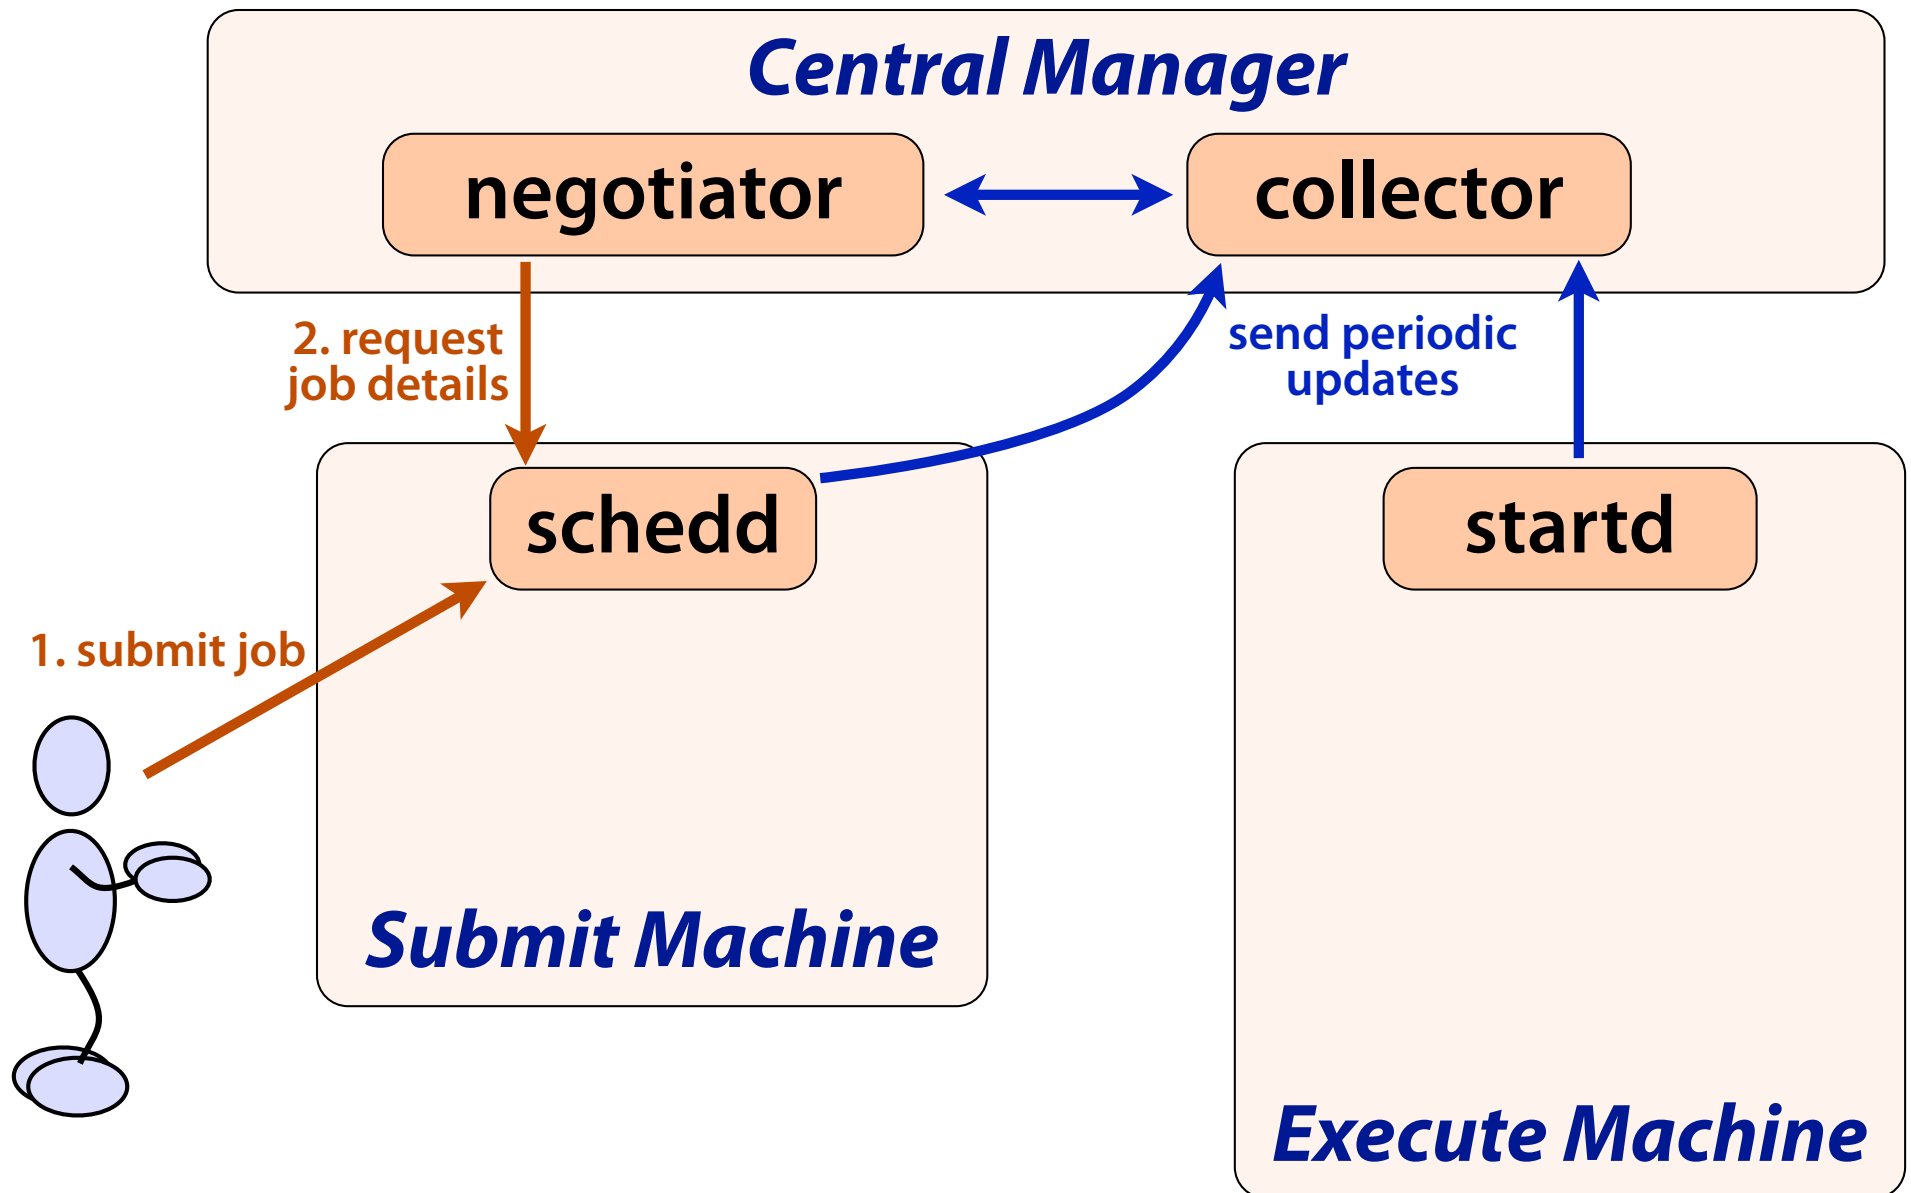




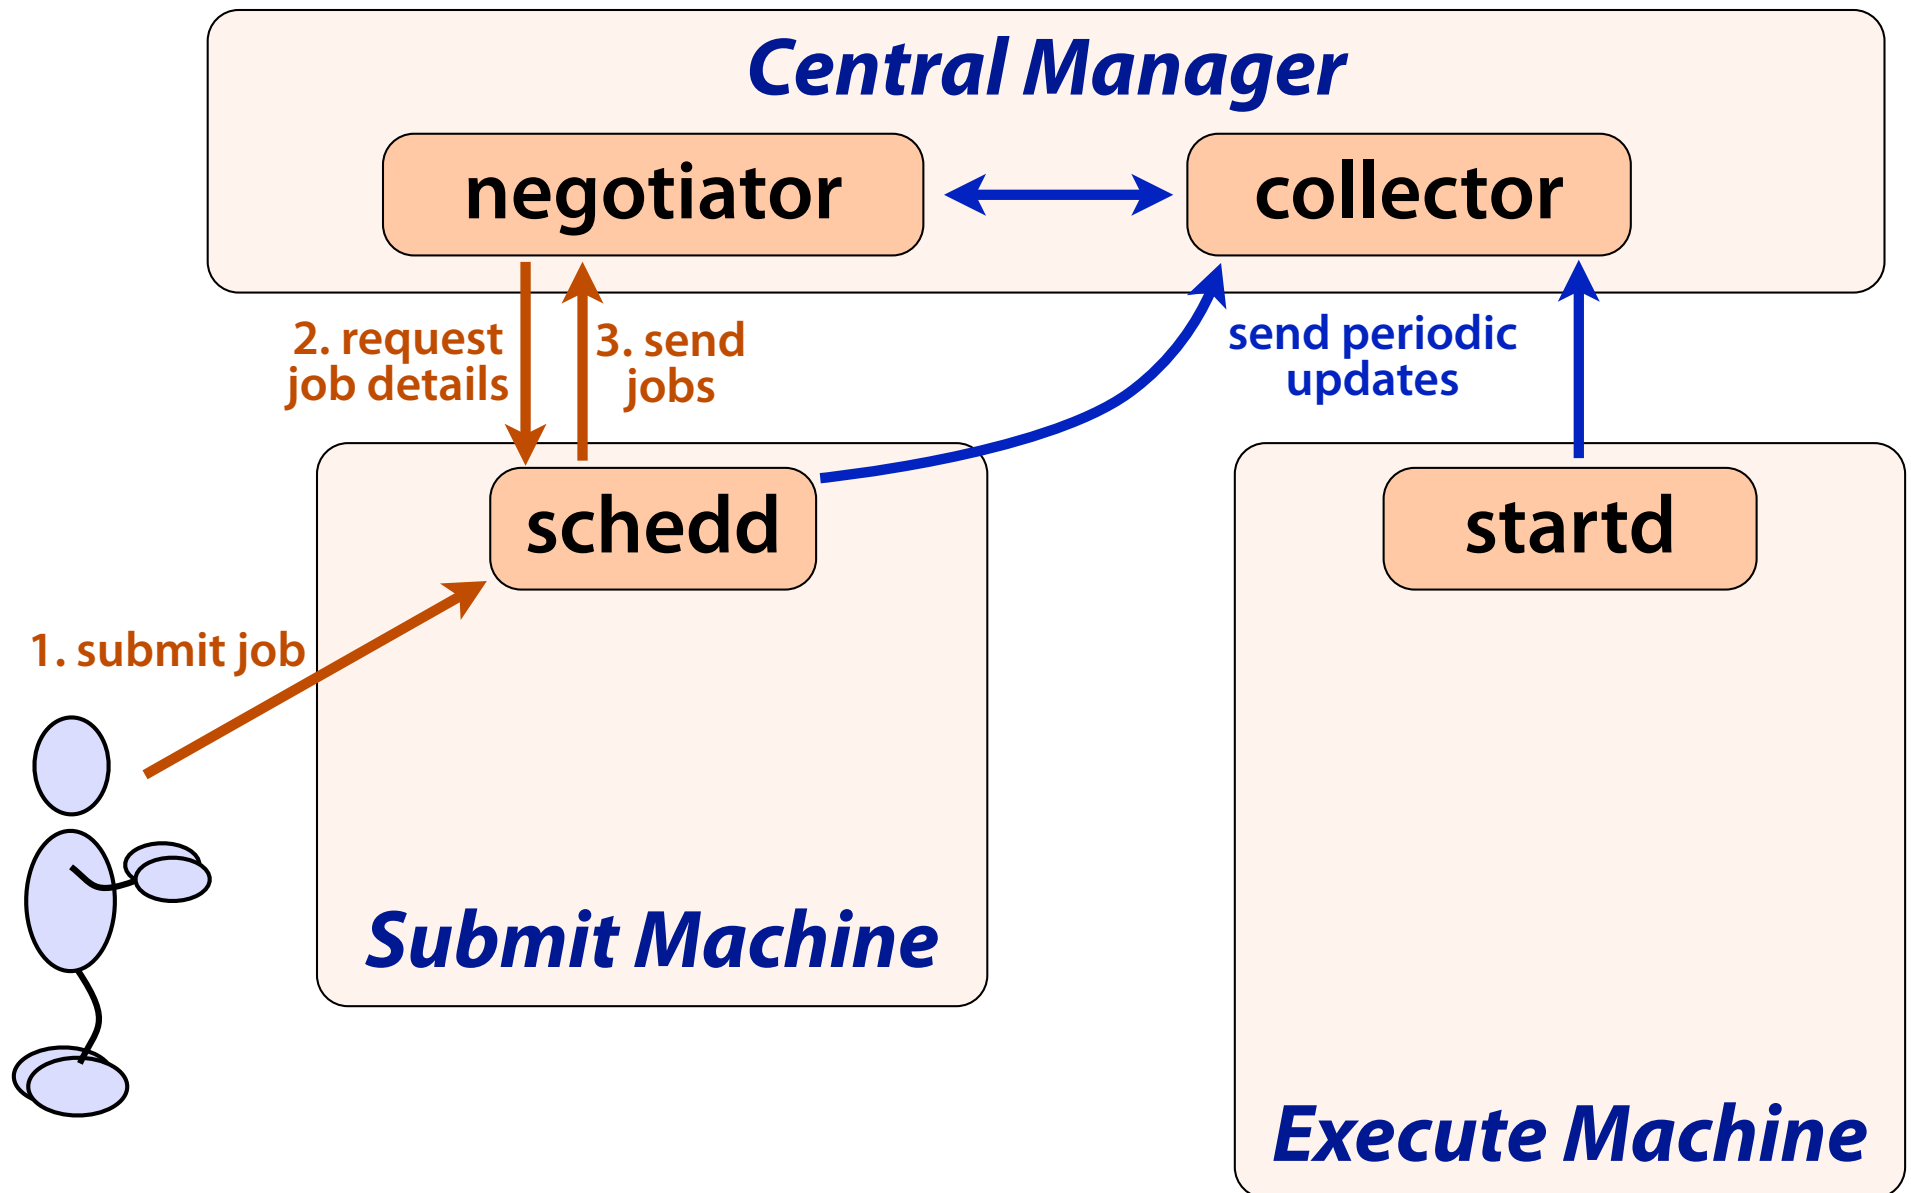
# The Life of an HTCondor Job



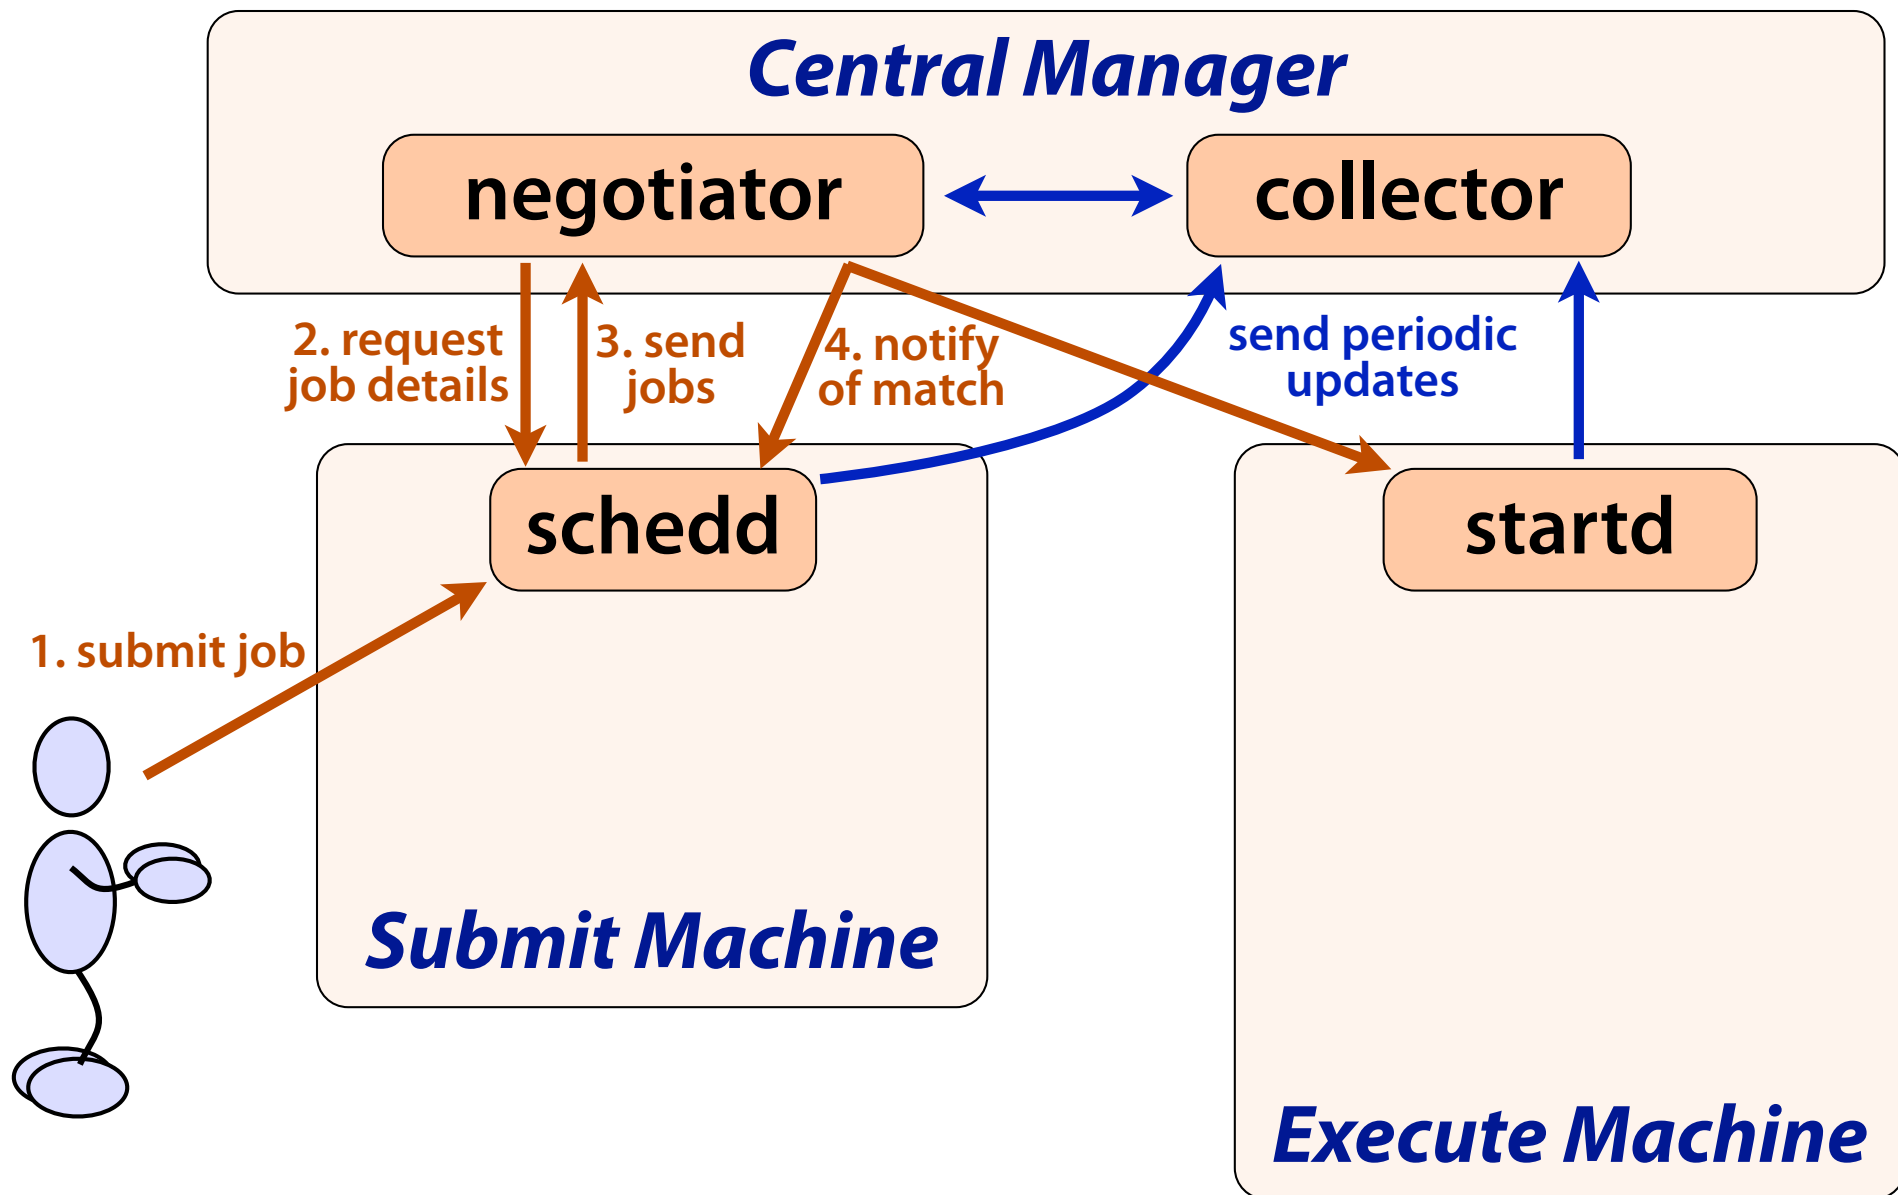
# The Life of an HTCondor Job



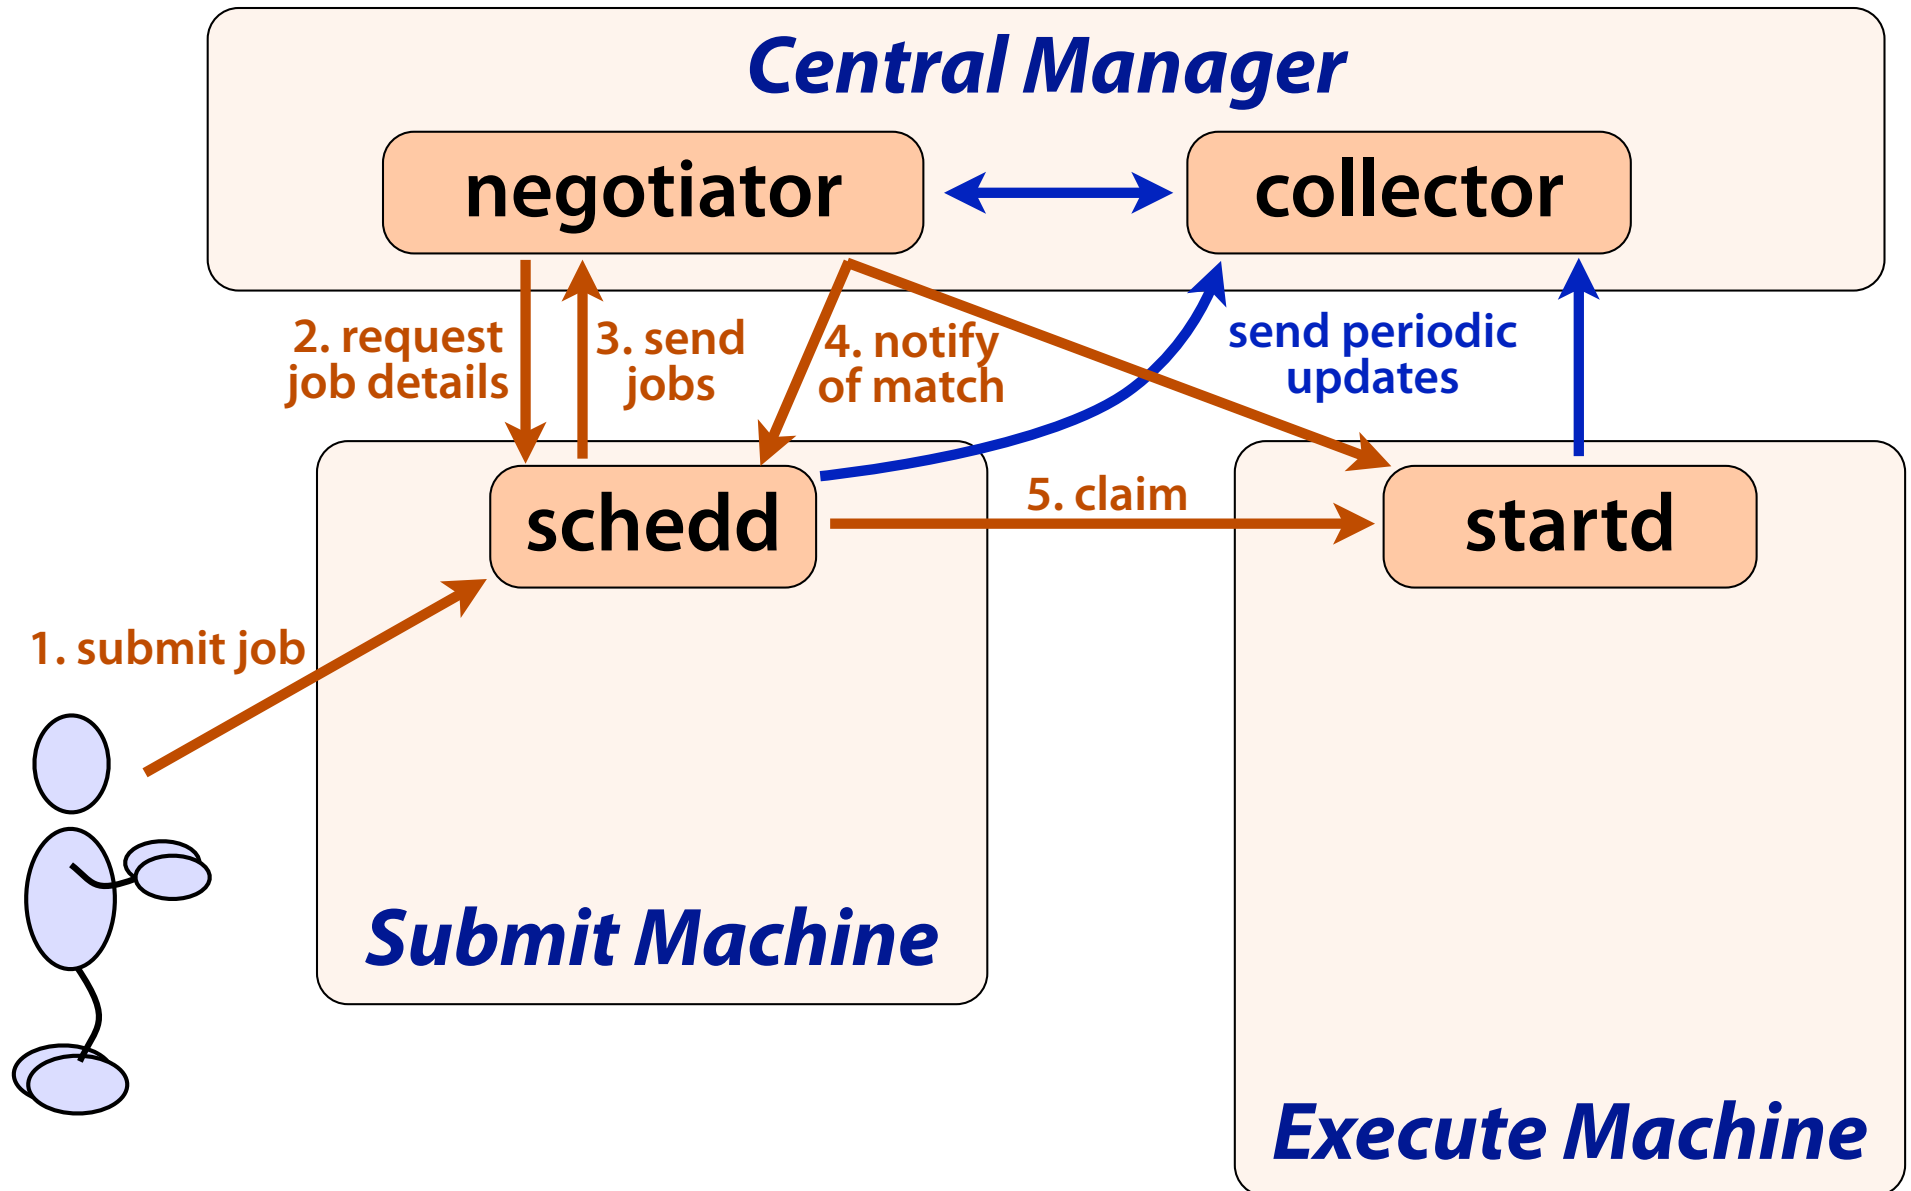
# The Life of an HTCondor Job



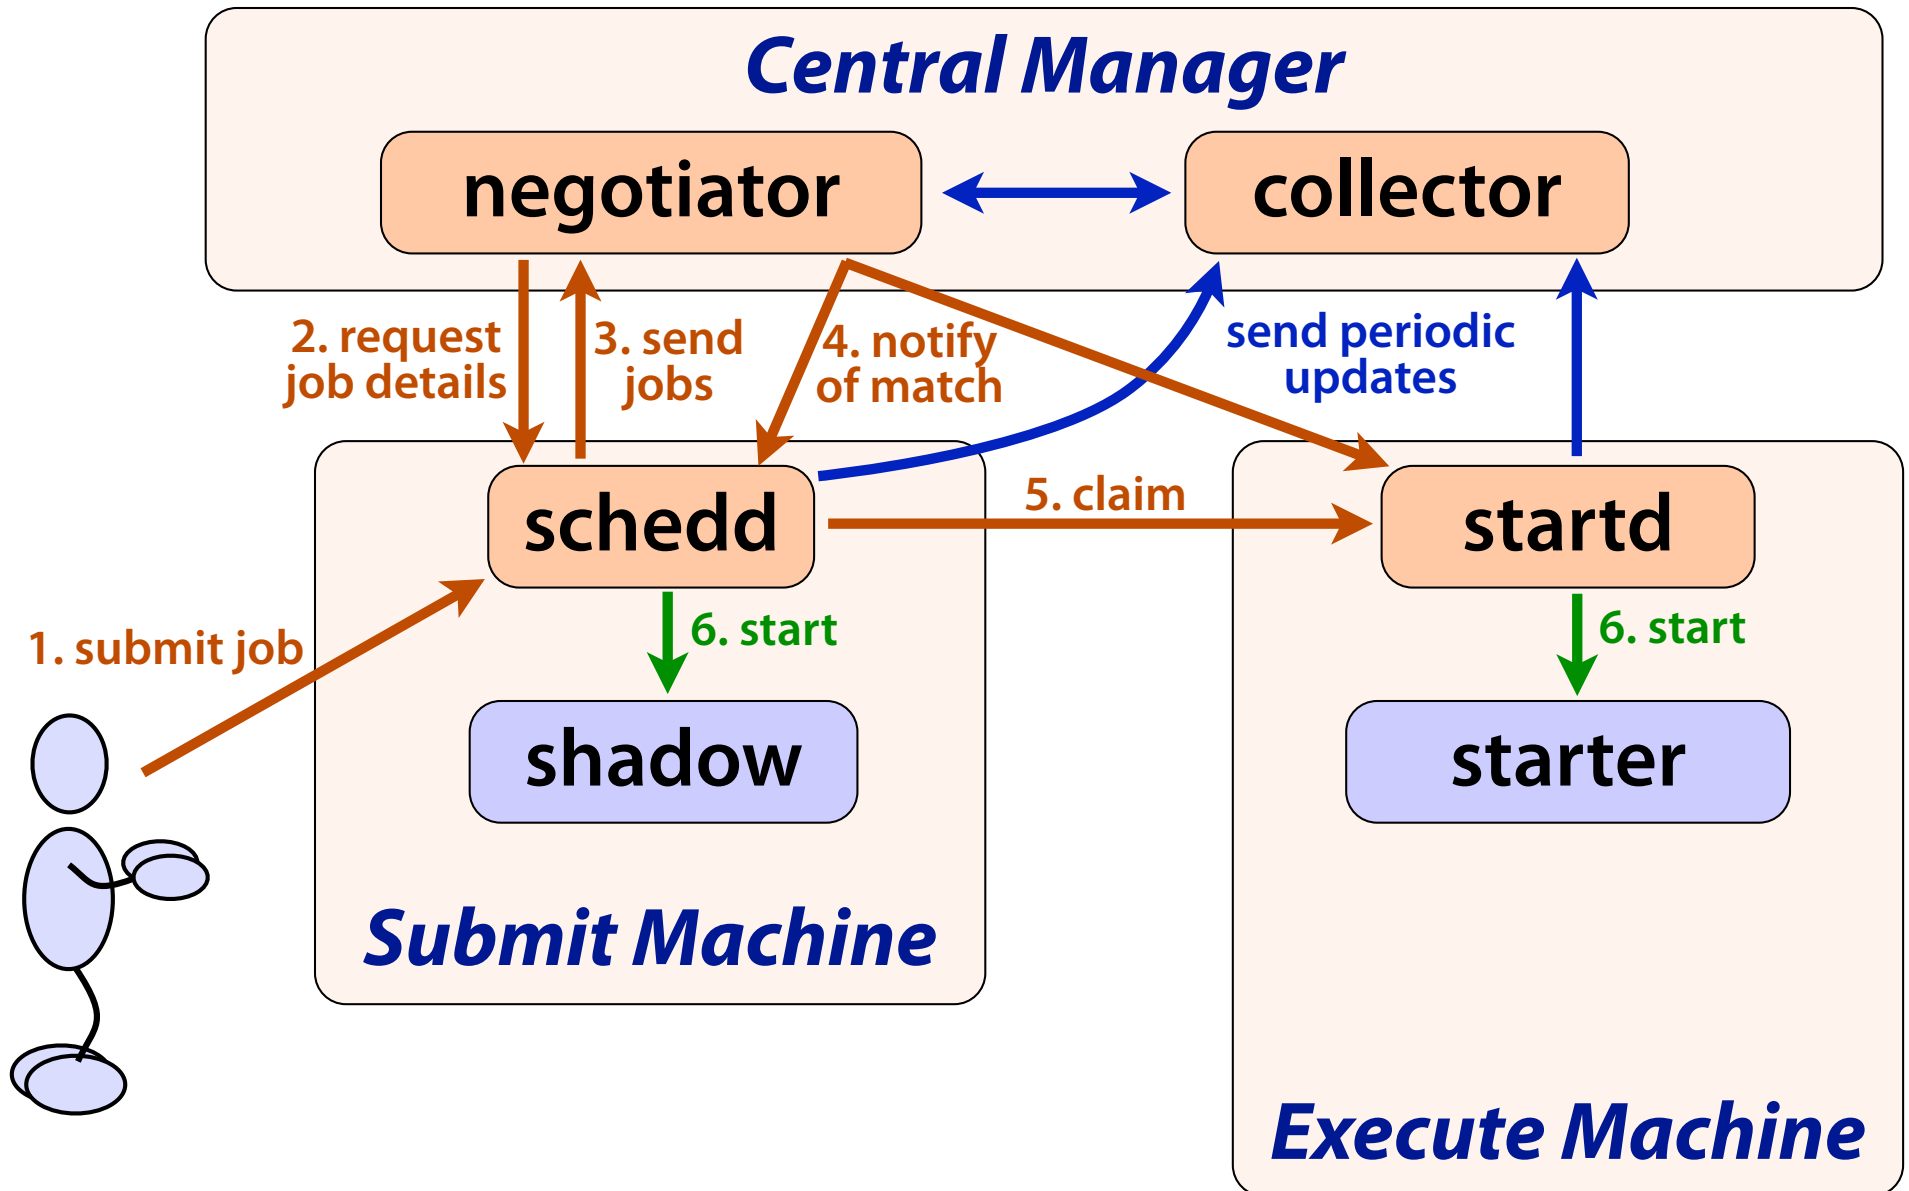
# The Life of an HTCondor Job



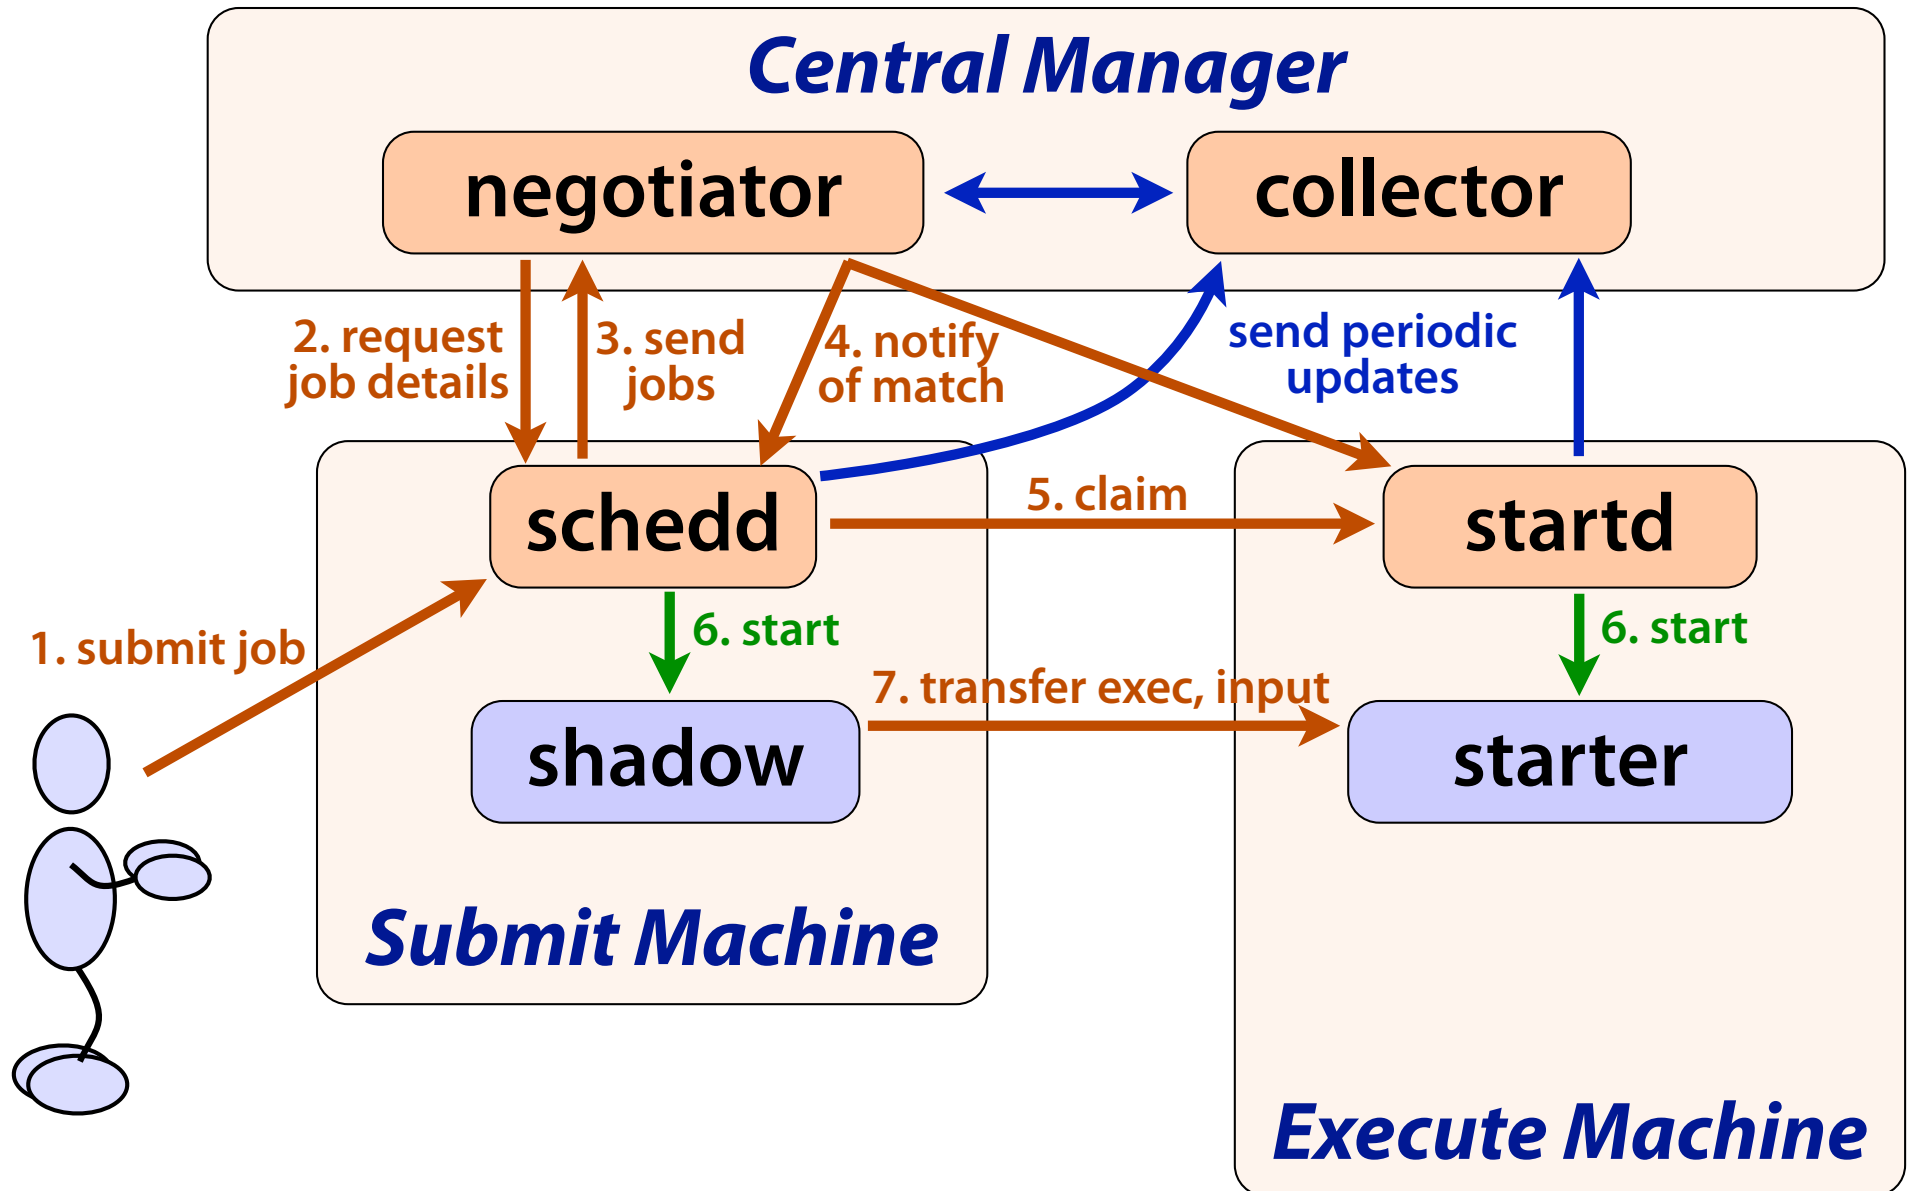
# The Life of an HTCondor Job



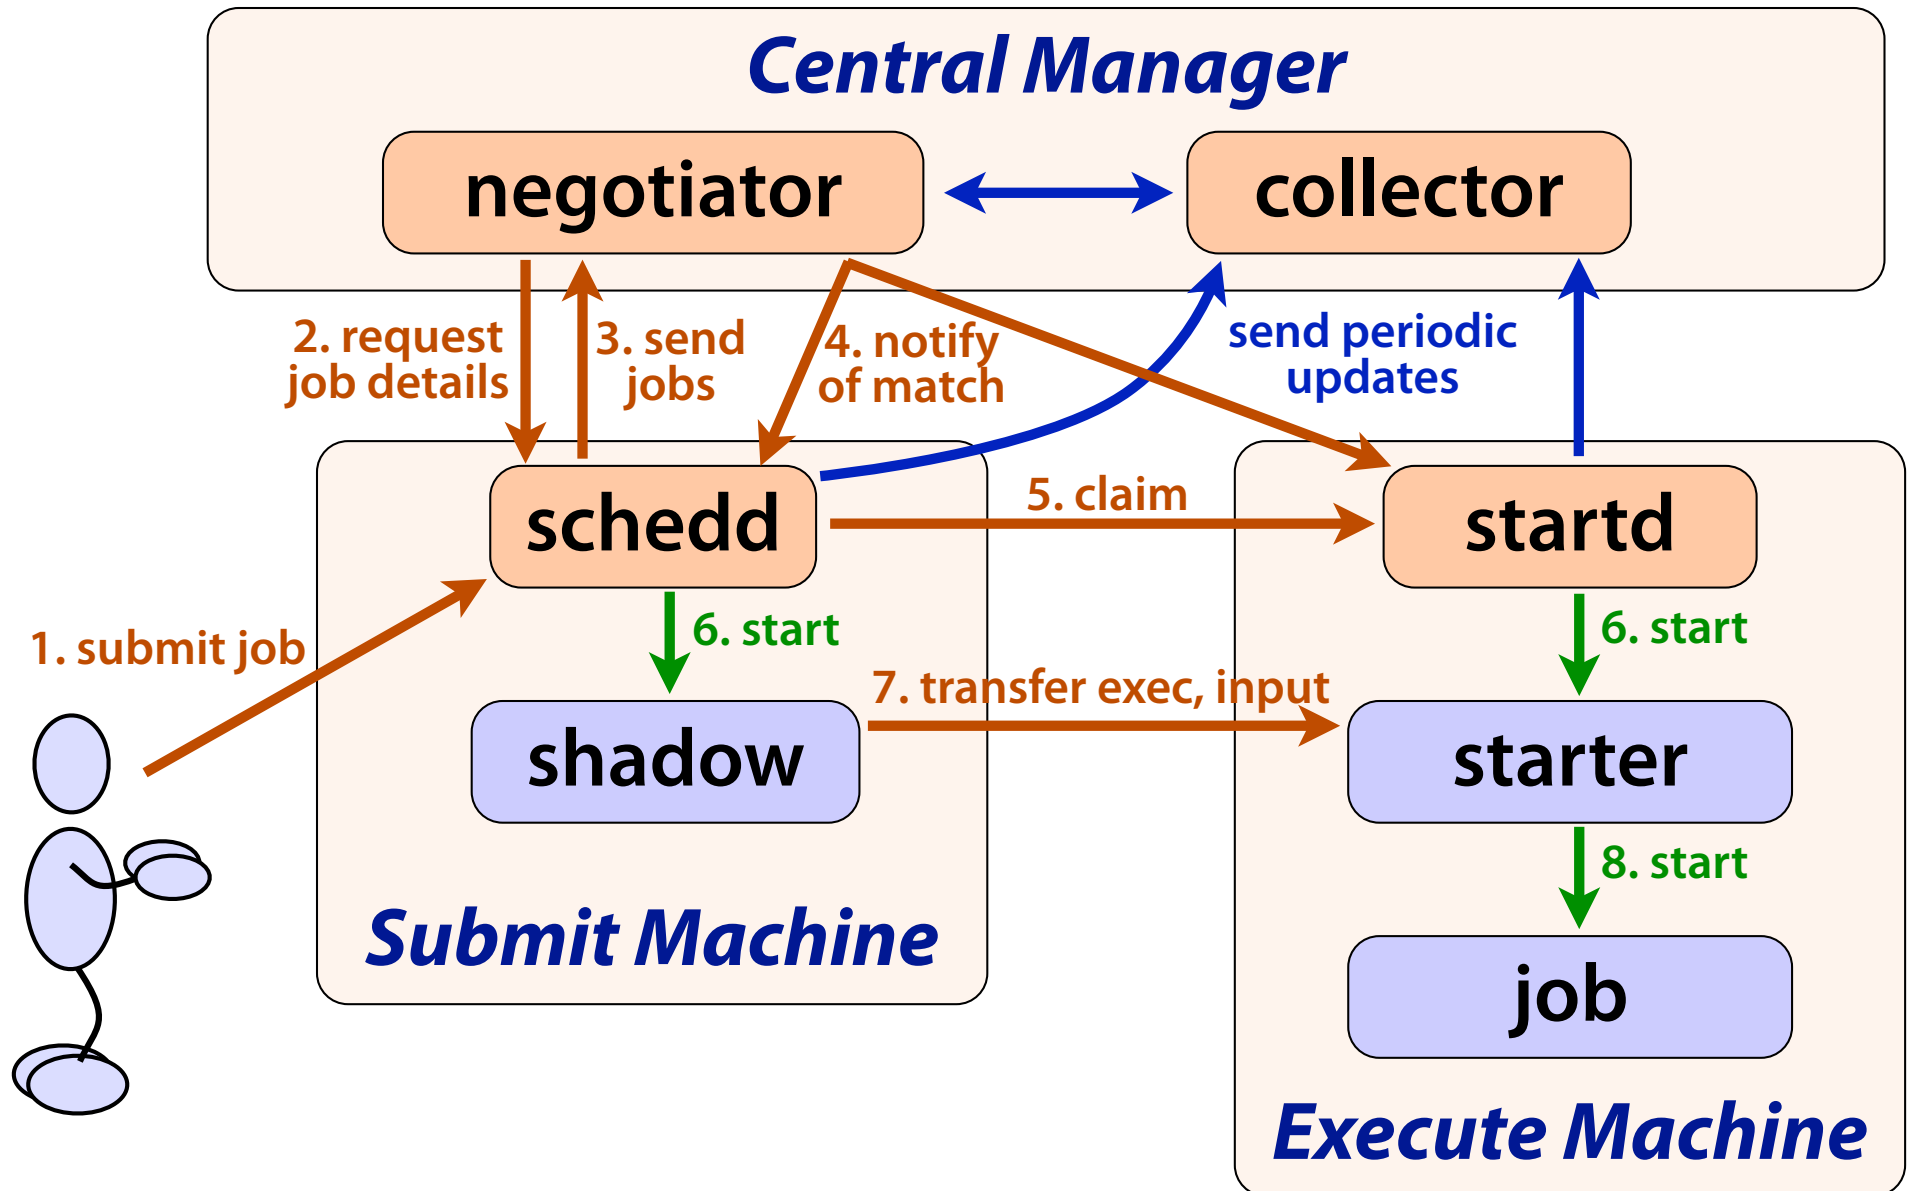
# The Life of an HTCondor Job



# The Life of an HTCondor Job

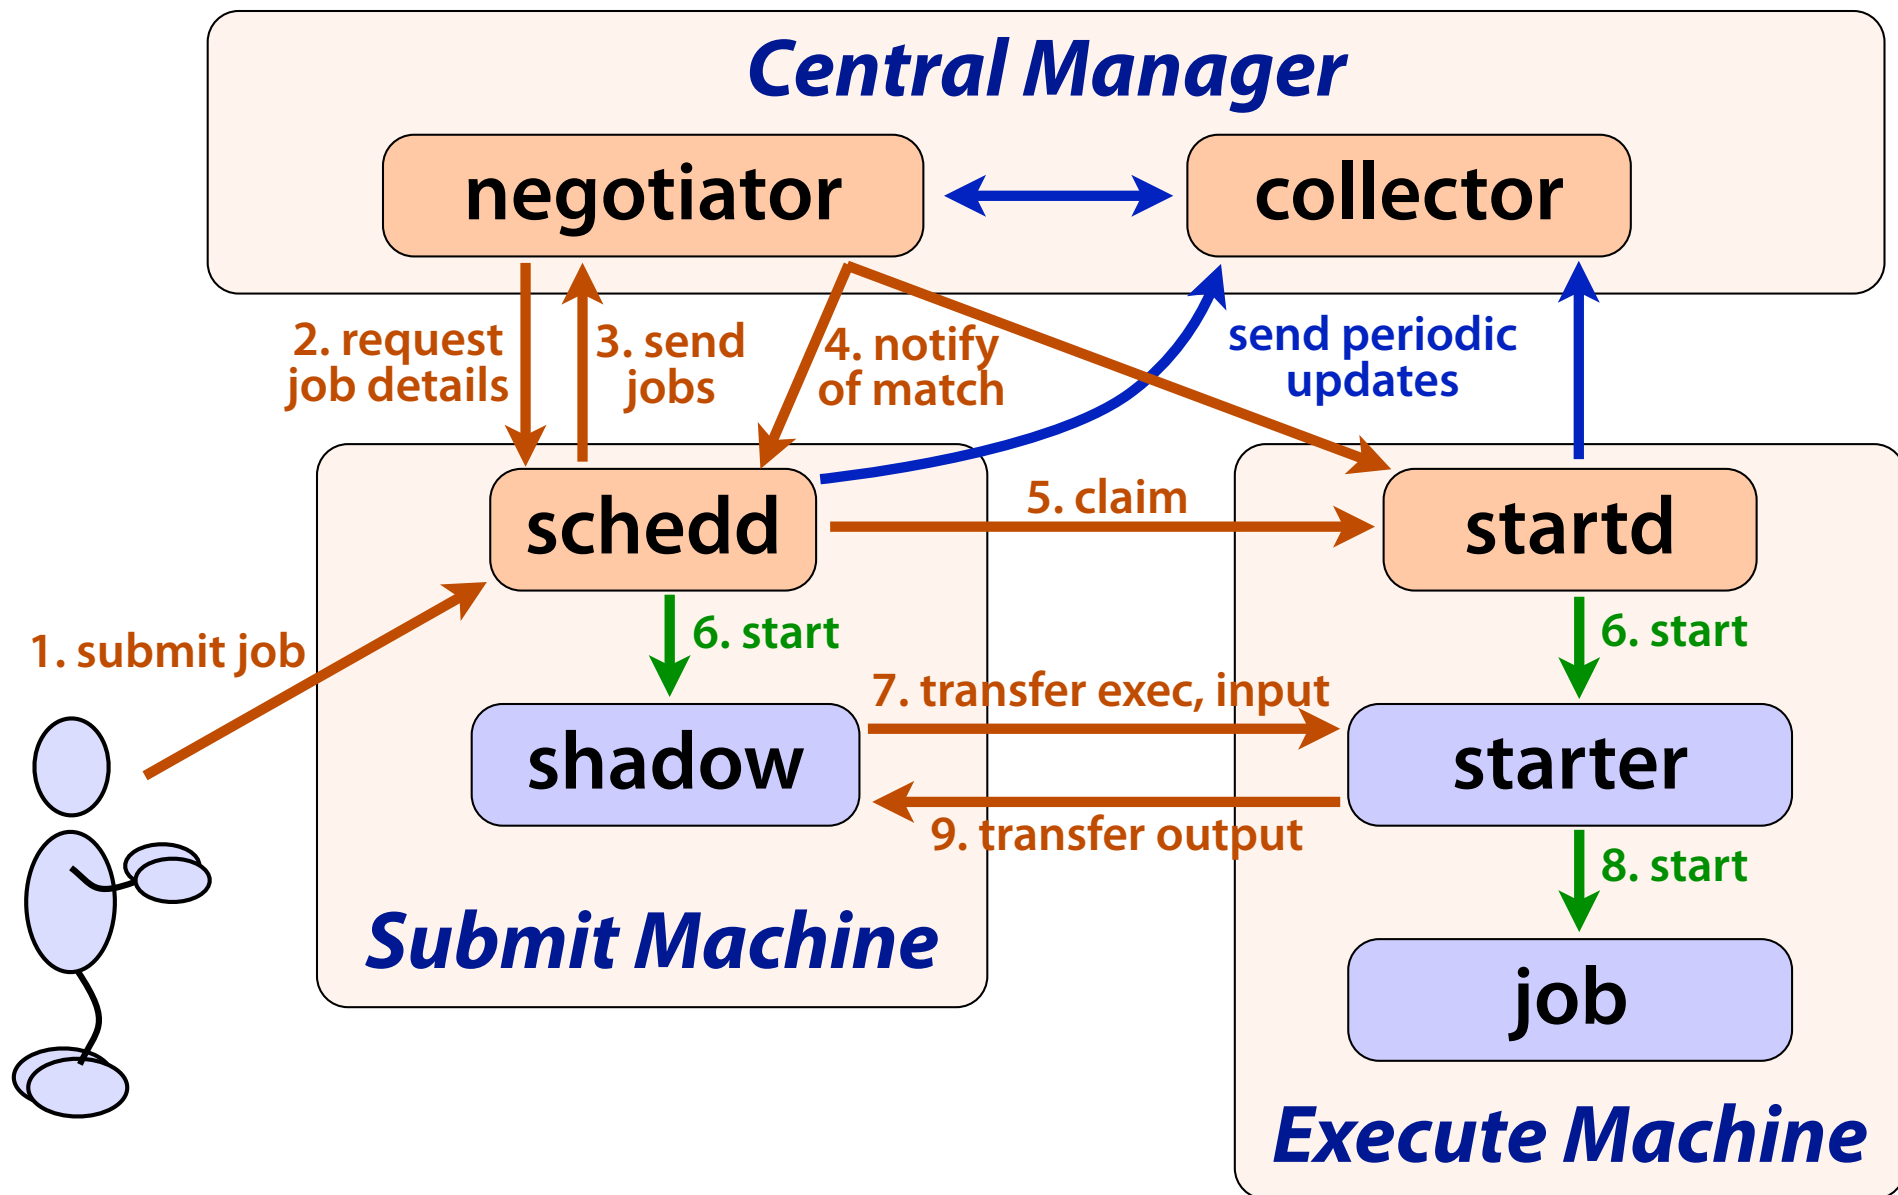


# The Life of an HTCondor Job





# The Life of an HTCondor Job



# Matchmaking Algorithm (sort of)

A. Gather lists of machines and waiting jobs

B. For each user:

1. Compute maximum # of slots to allocate to user  
(the user's "fair share", a % of whole pool)
2. For each job (up to user's maximum slots):
  - a. Find all machines that are acceptable  
(i.e., machine *and* job requirements are met)
  - b. If there are no acceptable machines, skip to next job
  - c. Sort acceptable machines by job preferences
  - d. Pick the best one
  - e. Record match of job and slot

# ClassAds

- In HTCondor, information about machines and jobs (and more) are represented by ClassAds
- You do not write ClassAds (much), but reading them may help understanding and debugging
- ClassAds can represent persistent facts, current state, preferences, requirements, ...
- HTCondor uses a core of predefined attributes, but users can add other, new attributes, which can be used for matchmaking, reporting, etc.

# Sample ClassAd Attributes

```
MyType = "Job"
TargetType = "Machine"
ClusterId = 14
Owner = "cat"
Cmd = ".../test-job.py"
Requirements = (Arch == "X86_64") &&
               (OpSys == "LINUX") &&
               ...
Rank = 0.0
In = "/dev/null"
UserLog = ".../test-job.log"
Out = "test-job.out"
Err = "test-job.err"
NiceUser = false
ShoeSize = 10
```


# Sample ClassAd Attributes

string

```
MyType = "Job"
TargetType = "Machine"
ClusterId = 14
Owner = "cat"
Cmd = "/.../test-job.py"
Requirements = (Arch == "X86_64") &&
               (OpSys == "LINUX") &&
               ...
Rank = 0.0
In = "/dev/null"
UserLog = "/.../test-job.log"
Out = "test-job.out"
Err = "test-job.err"
NiceUser = false
ShoeSize = 10
```

# Sample ClassAd Attributes

```
MyType = "Job"
TargetType = "Machine"
ClusterId = 14
Owner = "cat"
Cmd = "/.../test-job.py"
Requirements = (Arch == "X86_64") &&
               (OpSys == "LINUX") &&
               ...
Rank = 0.0
In = "/dev/null"
UserLog = "/.../test-job.log"
Out = "test-job.out"
Err = "test-job.err"
NiceUser = false
ShoeSize = 10
```



number


# Sample ClassAd Attributes

```
MyType = "Job"
TargetType = "Machine"
ClusterId = 14
Owner = "cat"
Cmd = ".../test-job.py"
Requirements = (Arch == "X86_64") &&
               (OpSys == "LINUX") &&
               ...
Rank = 0.0
In = "/dev/null"
UserLog = ".../test-job.log"
Out = "test-job.out"
Err = "test-job.err"
NiceUser = false
ShoeSize = 10
```

operations/  
expressions

# Sample ClassAd Attributes

```
MyType = "Job"
TargetType = "Machine"
ClusterId = 14
Owner = "cat"
Cmd = ".../test-job.py"
Requirements = (Arch == "X86_64") &&
               (OpSys == "LINUX") &&
               ...
Rank = 0.0
In = "/dev/null"
UserLog = ".../test-job.log"
Out = "test-job.out"
Err = "test-job.err"
NiceUser = false
ShoeSize = 10
```





# Sample ClassAd Attributes

```
MyType = "Job"
TargetType = "Machine"
ClusterId = 14
Owner = "cat"
Cmd = ".../test-job.py"
Requirements = (Arch == "X86_64") &&
               (OpSys == "LINUX") &&
               ...
Rank = 0.0
In = "/dev/null"
UserLog = ".../test-job.log"
Out = "test-job.out"
Err = "test-job.err"
NiceUser = false
ShoeSize = 10
```

arbitrary

# HTCondor Universes

- Different combinations of configurations and features are bundled as *universes*:

<b>vanilla</b>	A “normal” job; default, fine for today
<b>standard</b>	Supports checkpointing and remote I/O
<b>java</b>	Special support for Java programs
<b>parallel</b>	Supports parallel jobs (such as MPI)
<b>grid</b>	Submits to remote system (more tomorrow)
... and more!	

# HTCondor Priorities

- **Job priority**
  - ▶ Set per job by the user (owner)
  - ▶ Relative to that user's other jobs
  - ▶ Set in submit file or change later with **condor\_prio**
  - ▶ Higher number means run sooner
- **User priority**
  - ▶ Computed based on past usage
  - ▶ Determines user's "fair share" percentage of slots
  - ▶ Lower number means run sooner (0.5 is minimum)
- **Preemption**
  - ▶ Low priority jobs stopped for high priority ones (stopped jobs go back into the regular queue)
  - ▶ Governed by fair-share algorithm and pool policy
  - ▶ Not enabled on all pools



# HTCondor Commands

# List Jobs: `condor_q`

- Select jobs: by user (e.g., you), cluster, job ID
- Format output as you like
- View full ClassAd(s), typically 80–90 attributes (most useful when limited to a single job ID)
- Ask HTCondor why a job is not running
  - ▶ May not explain everything, but can help
  - ▶ Remember: Negotiation happens periodically
- Explore **`condor_q`** options in next exercises

# List Slots: `condor_status`

- Select slots: available, host, specific slot
- Select slots by ClassAd expression  
E.g., slots with SL 6 (OS) and  $\geq 10$  GB memory
- Format output as you like
- View full ClassAd(s), typically 120–250 attributes  
(most useful when limited to a single slot)
- Explore **`condor_status`** options in exercises



# Submit Files

# Resource Requests

```
request_cpus = ClassAdExpression  
request_disk = ClassAdExpression  
request_memory = ClassAdExpression
```

- Ask for minimum resources of execute machine
- May be dynamically allocated (very advanced!)
- *Check job log for actual usage!!!*

```
request_disk = 20000000    # in KB by default  
request_disk = 2GB        # KB, MB, GB, TB  
  
request_memory = 2000      # in MB by default  
request_memory = 2GB       # KB, MB, GB, TB
```



# File Access in HTCondor

- **Option 1: Shared filesystem**

- ▶ Easy to use (jobs just access files)
- ▶ But, must exist and be ready handle load

```
should_transfer_files = NO
```

- **Option 2: HTCondor transfers files for you**

- ▶ Must name all input files (except executable)
- ▶ May name output files; defaults to all new/changed

```
should_transfer_files = YES  
when_to_transfer_output = ON_EXIT  
transfer_input_files = a.txt, b.tgz
```

# Email Notifications

**notification = Always | Complete | Error | Never**

- When to send email
  - ▶ **Always**: job checkpoints or completes
  - ▶ **Complete**: job completes (default)
  - ▶ **Error**: job completes with error
  - ▶ **Never**: do not send email

**notify\_user = email**

- Where to send email
- Defaults to ***user@submit-machine***

# Requirements and Rank

**requirements = *ClassAdExpression***

- Expression must evaluate to **true** to match slot
- HTCondor adds defaults! Check ClassAds ...
- See HTCondor Manual (esp. 2.5.2 & 4.1) for more

**rank = *ClassAdExpression***

- Ranks matching slots in order by preference
- Must evaluate to a FP number, higher is better
  - ▶ False becomes 0.0, True becomes 1.0
  - ▶ Undefined or error values become 0.0
- Writing rank expressions is an art form

# Arbitrary Attributes

**+*AttributeName* = *value***

- Adds arbitrary attribute(s) to job's ClassAd
- Useful in (at least) 2 cases:
  - ▶ Affect matchmaking with special attributes
  - ▶ Report on jobs with specific attribute value
- Experiment with reporting during exercises!

# Many Jobs Per Submit File, Pt. 1

- Can use **queue** statement many times
- Make changes between **queue** statements
  - ▶ Change **arguments**, **log**, **output**, input files, ...
  - ▶ Whatever is not explicitly changed remains the same

```
executable = test.py
```

```
log          = test.log
```

```
output      = test-1.out  
arguments   = "test-input.txt 42"  
queue
```

```
output      = test-2.out  
arguments   = "test-input.txt 43"  
queue
```

# Many Jobs Per Submit File, Pt. 1

- Can use **queue** statement many times
- Make changes between **queue** statements
  - ▶ Change **arguments**, **log**, **output**, input files, ...
  - ▶ Whatever is not explicitly changed remains the same

```
executable = test.py
```

```
log = test.log
```

```
output = test-1.out  
arguments = "test-input.txt 42"  
queue
```

```
output = test-2.out  
arguments = "test-input.txt 43"  
queue
```

**log = test.log (still)**

# Many Jobs Per Submit File, Pt. 2

## queue $N$

- Submits  $N$  copies of the job
  - ▶ One cluster number for all copies, just as before
  - ▶ Process numbers go from 0 to  $(N-1)$
- What good is having  $N$  copies of the same job?
  - ▶ Randomized processes (e.g., Monte Carlo)
  - ▶ Job fetches work description from somewhere else
  - ▶ But what about overwriting output files, etc.?
- Wouldn't it be nice to have different files and/or arguments automatically applied to each job?

# Separating Files by Run

```
output = program.out.$(Cluster).$(Process)
```

- Can use these variables anywhere in submit file
  - ▶ Often used in **output**, **error**, and **log** files
- Maybe use **\$(Process)** in arguments?
  - ▶ Can't perform math on values; code must accept as is

...

```
output = test.$(Cluster)_$(Process).out  
log      = test.$(Cluster)_$(Process).log
```

```
arguments = "test-input.txt $(Process)"  
queue 10
```



# Separating Directories by Run

```
initialdir = path
```

- Use *path* (instead of submit dir.) to locate files
  - ▶ E.g.: **output**, **error**, **log**, **transfer\_input\_files**
  - ▶ *Not executable*; it is still relative to submit directory
- Use **\$(Process)** to separate all I/O by job ID

```
initialdir = run-$(Process)
transfer_input_files = input-$(Process).txt
output = test.$(Cluster)-$(Process).out
log     = test.$(Cluster)-$(Process).log

arguments = "input-$(Process).txt $(Process)"
queue 10
```



# Your Turn!

# Exercises!

- Ask questions!
- Lots of instructors around
- Reminder: Get your X.509 certificate today!
- Coming next:

Now – 12:15      Hands-on exercises

12:15–1:15      Lunch

1:15–5:30      Afternoon sessions with Zach