

Principles for HTC

2013 OSG User School, Thursday, Lecture 3

Tim Cartwright
cat@cs.wisc.edu

University of Wisconsin–Madison
OSG Software Team Manager
OSG Education Coordinator



The Hope of HTC

Reminder: What is HTC?

“... the use of *many computing resources* over *long periods of time* to accomplish a computational task” (Wikipedia, retr. 25 June 2013)

- Try to use all resources, reliably, all the time
- Maximize operations per *year*
- Use any computers, even old or slow ones

What System Is This?

- ***Mystery System X:***
 - ▶ Provides a *lot* of computing
 - ▶ Has high availability and reliability
 - ▶ Degrades gracefully
 - ▶ Spreads the workload automatically
 - ▶ Grows (and shrinks) easily when needed
 - ▶ Responds well to temporary overloads
 - ▶ Adapts easily to new uses
- HTCondor? OSG? Amazon EC2? Other Clouds?

Actually...

Actually...

- Those were all *promised* features!

Actually...

- Those were all *promised* features!
- ... of distributed data processing systems

Actually...

- Those were all *promised* features!
- ... of distributed data processing systems
- ... from the 1970s!!!

(Adapted from: Enslow, P. H., Jr. (1978). What is a “distributed” data processing system? *Computer*, 11(1), 13–21. doi:10.1109/C-M.1978.217901)

Actually...

- Those were all *promised* features!
- ... of distributed data processing systems
- ... from the 1970s!!!

(Adapted from: Enslow, P. H., Jr. (1978). What is a “distributed” data processing system? *Computer*, 11(1), 13–21. doi:10.1109/C-M.1978.217901)

- Sound like promises of today: HTC, grid, cloud



The *Hype* of HTC?

Miron Livny



- Founder and leader of Condor Project since mid-1980s
- PI and Technical Director of OSG
- Coined term “high throughput computing”
- Has principled approach to HTC

Miron's Reminder

What has been
is what will be,
and what has been done
is what will be done,
and *there is nothing new
under the sun.*

— Ecclesiastes 1:9 (ESV)

Attributed to Koheleth, who was Ecclesiastes or its author, traditionally believed to be Solomon, son of David, king in Jerusalem, ~950 BCE



Ecclesiastes, (קְהֵלֶת, Kohelet, “son of David, and king in Jerusalem,” alias Solomon, wood engraving, Gustave Doré (1832–1883)

Nothing New Under the Sun

- New terms, new hype

Distributed data processing, cluster computing, grid computing, cloud computing, virtualization, peer-to-peer, client-server, cyberinfrastructure

- But, the underlying problems are the same
- Principles to address the problems are the same


But, But ... Aren't Clouds New?

- Amazon EC2 makes it easy to create virtual machines on the fly. That's new, right?
- **HTCondor:** Allocate a computer + Start a job
- **Amazon EC2:** Allocate a VM + Start it
- Conceptually, a virtual machine is just another kind of job: You start it, you stop it, you have some control over it

What *Is* New About Clouds?

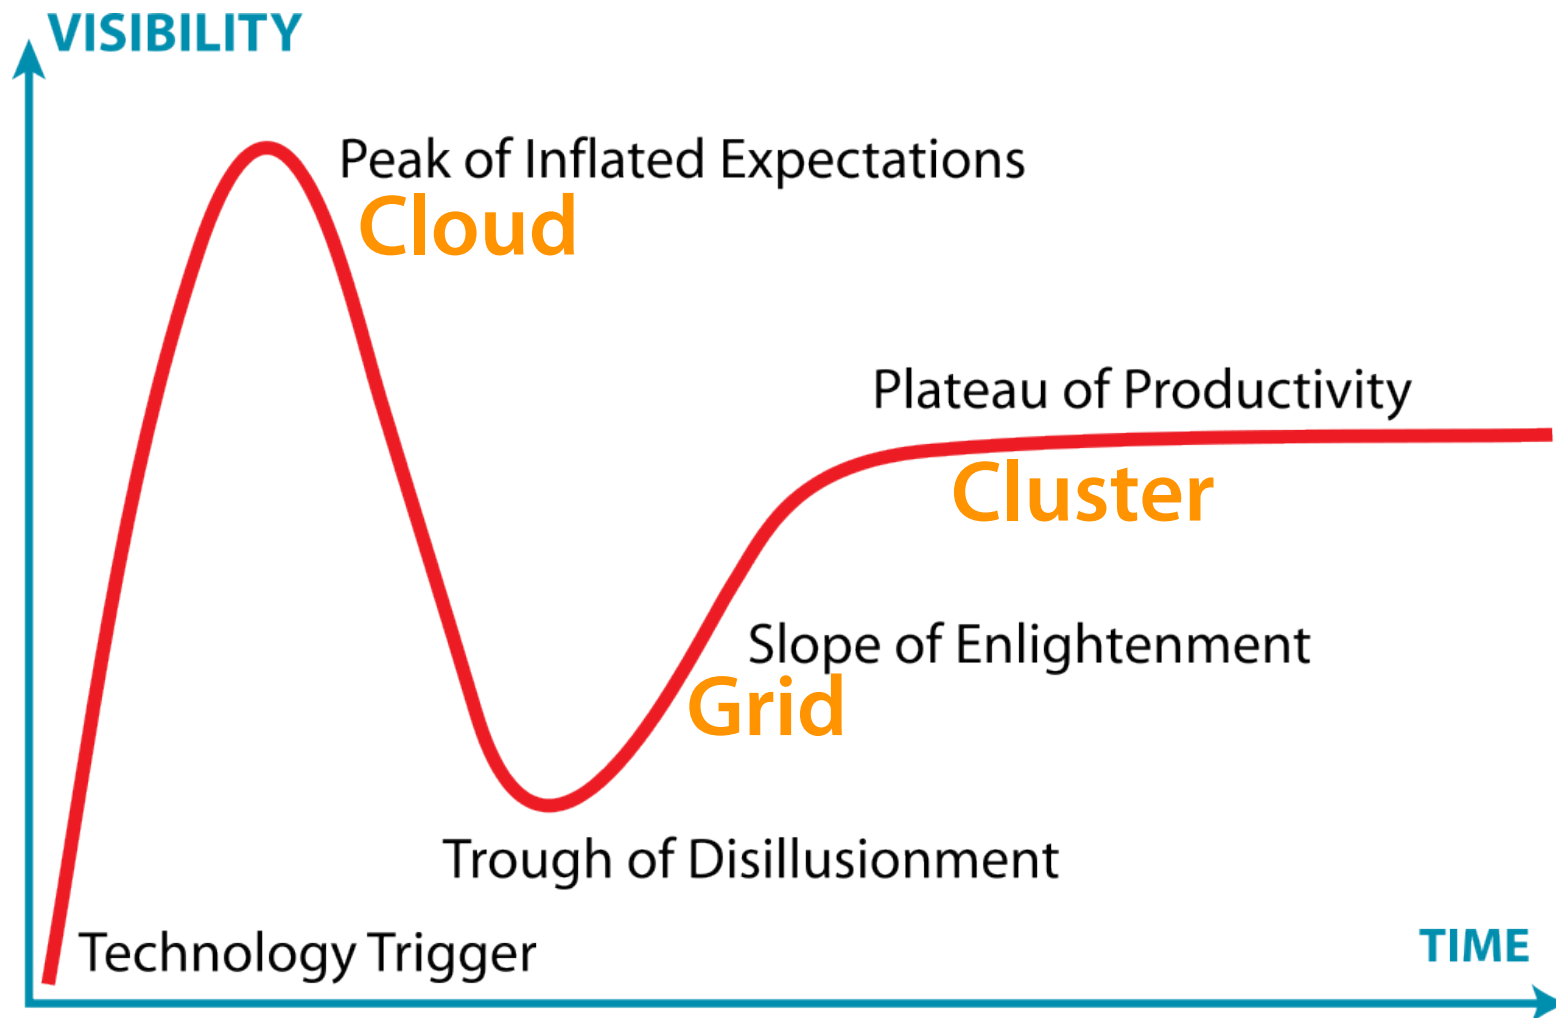
- **OSG/XSEDE:** User cost for existing resources: **\$0**
 - ▶ OSG: Submit jobs, wait for available slots
 - ▶ XSEDE: Submit proposal, get free allocation
- **Amazon EC2:** Bring a credit card, **≥\$0.05/hour**
 - ▶ Whether you use it or not
- Really, the difference is between gov't-funded and commercial resources, ***not*** technology

Gartner Hype Cycle



<http://www.gartner.com/technology/research/methodologies/hype-cycle.jsp>

Gartner Hype Cycle



<http://www.gartner.com/technology/research/methodologies/hype-cycle.jsp>

Instead

- Know the principles
- Understand how the principles apply
- And you will be able to use the next technology in the next hype cycle to help your science
- So let's look understand the principles



Distributed Computing

Distributed Computing???

“You know you have a distributed system when the crash of a computer you’ve never heard of stops you from getting any work done.”

— Leslie Lamport

Definition of Distributed Computing

1. Multiplicity of resources
2. Component interconnection
3. Unity of control
4. System transparency
5. Component autonomy

Enslow, P. H., Jr., & Saponas, T. G. (**1981**). *Distributed and decentralized control in fully distributed processing systems: A survey of applicable models* (GIT-ICS-81/02). Georgia Institute of Technology.

1. Multiplicity of Resources

- Use many general purpose components
 - ▶ Physical or logical
 - ▶ Do not need to be the same (homogenous)
 - ▶ But should have same/similar capabilities
 - ▶ I.e., want to assign tasks anywhere
- Improves reliability and throughput

2. Component Interconnection

- Thus, physically and/or logically distributed
- Connected via network (these days)
 - ▶ Use a two-way cooperative protocol
 - ▶ Requests can be refused based on local knowledge
- Example: HTTP



3. Unity of Control

- ***Not*** centralized control
 - ▶ Creates single point of failure
 - ▶ Avoid critical paths, critical components
- Unify around a common goal
 - ▶ Via shared policy and, maybe, software
 - ▶ Multiple centers of control
- Improves overall reliability and availability

4. System Transparency

- Users should feel like the whole system is one giant virtual machine
- Should not need to know:
 - ▶ Which physical component will perform task
 - ▶ Where the component is located
- May be the hardest part to achieve

5. Component Autonomy

- ***Autonomy*** (act locally)
 - ▶ Resources make their own decisions
 - ▶ Accept or refuse requests
 - ▶ Based on local policy
- ***Cooperative*** (think globally)
 - ▶ I.e., common policies and goals

Laptops: Distributed System?

Multiplicity?	Sort of: multiple cores
Interconnection?	No: not physically distributed
Unity of control?	Yes
Transparency?	Yes: <i>is</i> one system
Autonomous?	No: parts under central control

So, your laptop alone is not a distributed system

glideinWMS: Distributed System?

Multiplicity?	Yes: many diverse resources
Interconnection?	Yes: distributed and connected
Unity of control?	Yes: one system and goal
Transparency?	Yes: appears to be one pool
Autonomous?	Yes: sites have local control

So, glideinWMS is a distributed system [pheew!]



Principles for HTC

4 Principles for DHTC

- Enslow's definitions are ~30 years old
- From them, OSG has derived guiding principles
- Copied from the recent OSG funding proposal

DHTC Principles

- ***Resource Diversity***

The system must be flexible enough to accept many types of resources, software, and services

- ***Dependability***

Throughput must be tolerant to faults: There will always be services or resources that are not available

- ***Autonomy***

You must allow resource providers from different organizations to share; you must allow them to preserve local autonomy to set policies and select technologies

- ***Mutual Trust***

You must support complex trust relationships across organizations and software tools



got principles?

glideinWMS: Resource Diversity?

- Can use different kinds of Compute Elements:
GRAM, CREAM, NorduGrid
- Compute Elements work with different clusters:
HTCondor, Torque/PBS, LSF, SGE, (SLURM soon)
- Supports multiple platforms:
Intel 32- & 64-bit × RHEL, CentOS, SL 5 & 6

glideinWMS: Dependability?

- Uses HTCondor:
 - ▶ Reliable queue to manage jobs
 - ▶ Resources appear and disappear
 - ▶ Pool (often) relies on one Central Manager...
- Accesses multiple sites:
 - ▶ Individual sites come and go

glideinWMS: Autonomy?

- *Sites*

- ▶ Choose to accept VOs and users
- ▶ Can limit how many pilot jobs run
- ▶ Can limit how long pilot jobs run
- ▶ Can prefer local users over glide-ins

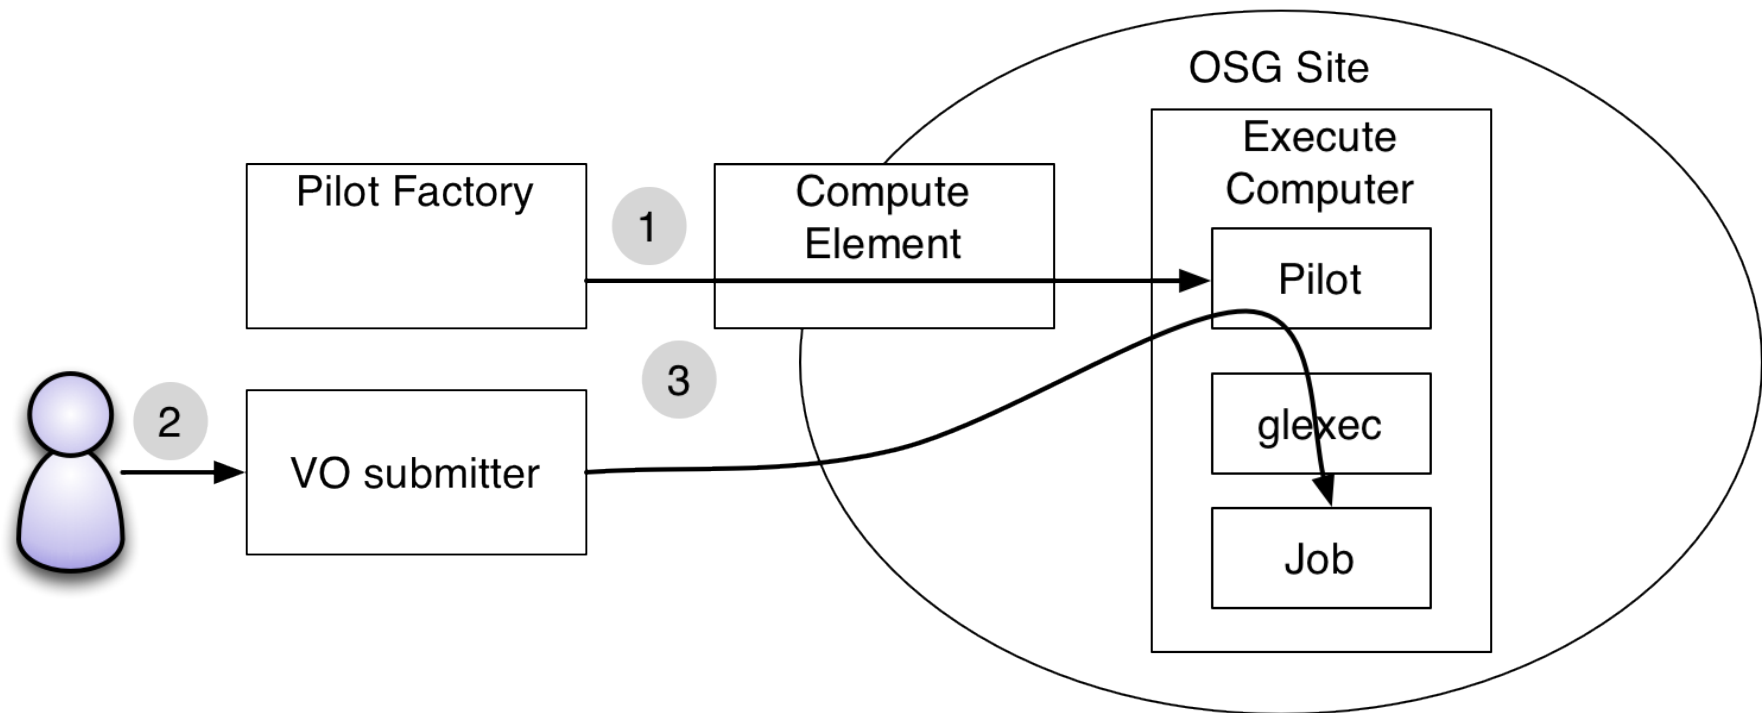
- *Glide-ins*

- ▶ Factory decides how many pilots to submit
- ▶ Factory controls policy on running pilots
- ▶ Pool has policy for how users share resources

glideinWMS: Mutual Trust?

- Pilot jobs are authenticated and authorized
 - ▶ Based on X.509 certificates
- gLExec
 - ▶ Some sites are uncomfortable that *pilot* user is not *job* user
 - ▶ gLExec is voluntary way to authorize and/or log the *job* users

Digression: gLExec



Principles and Scale

- The principles apply:
 - ▶ For a cluster
 - ▶ For a campus
 - ▶ For a grid (like OSG) or cloud (EC2)
 - ▶ For ... ???



Principles and You

How should you apply the principles?

- ***Resource Diversity***

Can your jobs run on diverse resources? (Location, platform, configuration, storage, etc.) If you get more, can you use them?

- ***Dependability***

Are your jobs robust? Can they recover from (or at least report) errors? Do you use methods to check results, retry nodes, etc.?

- ***Autonomy***

Be prepared to deal with the consequences of local decisions; jobs may be rejected, removed, throttled; resources come & go

- ***Mutual Trust***

Who do you trust with your data? Do you believe your results ("trust but verify")? Must deal with security systems...

Be Skeptical!

- When someone wants to sell you a new way of doing distributed computing, ask:
 - ▶ Can it handle diverse resources?
 - ▶ Is it dependable? (Oh yeah, how so?)
 - ▶ Are components autonomous?
 - ▶ Does it manage trust relationships?
 - ▶ Does it improve on the ability to provide the fundamental principles, or is it merely a shiny new version of what we already have?

There is nothing new under the sun!



Questions?

- Discussion can continue into break
- Coming next:

Now – 2:45	Break
2:45–4:20	HTC Showcase
4:20–4:45	Break + Showcase Q&A
4:45–5:15	Wrap up