

# **Workflows: from Development to *Automated* Production**

**Thursday morning, 10:00 am**

Lauren Michael <lmichael@wisc.edu>

Research Computing Facilitator

University of Wisconsin - Madison

# 'Engineering' a Good Workflow

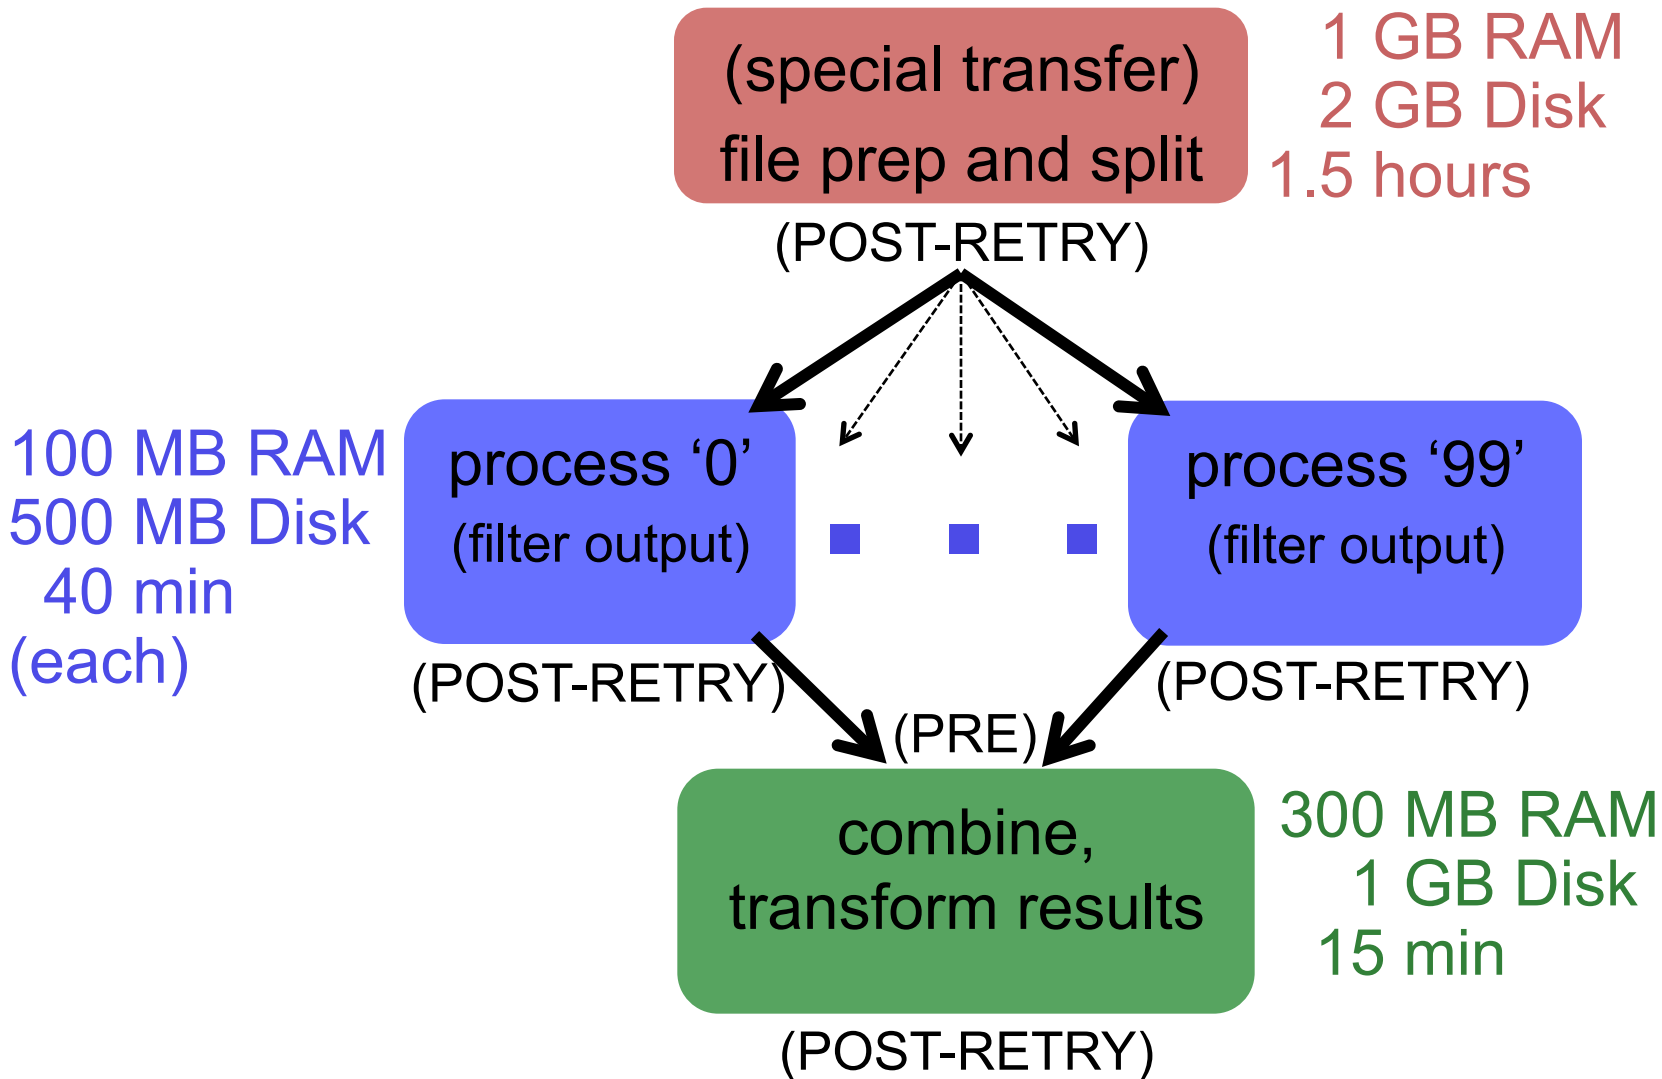
---

1. Draw out the *general* workflow
2. Define details (test 'pieces' with HTCondor jobs)
  - divide or consolidate 'pieces'
  - off-load file transfers and consider file transfer times
  - identify steps to be automated or checked
3. Build it piece-by-piece; test and optimize
4. Scale-up: data and computing resources
5. What more can you automate or error-check?

(And remember to document)



# From This . . .





# End Up with This

**DATA**



# Scaling and Optimization

---

- Your ‘small’ DAG runs (once)! Now what?
  - Need to make it run *full-scale*
  - Need to make it run *everywhere, everytime*
  - Need to make it run *unattended*
  - Need to make it run *when someone else tries*

# Scaling Up – Things to Think About

---

- More jobs:
  - 50-MB files may be fine for 10 or 100 jobs, but not for 1000 jobs. Why?
  - Steps to identify rare errors
- Larger files:
  - more execute RAM and disk space
  - potentially more transfer and compute time
- Scale-up gradually

# Data Scaling Solutions (larger files)

---

- File manipulations
  - split into pieces, when possible (HTC!)
  - filter files to only essential data
  - compression/decompression
- Listen to Lincoln's methods (Yesterday):
  - Sandbox
  - Caching
  - Pre-staging
  - Storage Element (SE) horsepower

# Make It Run Everywhere

---

- What does an OSG machine have?
  - Assume the worst: nothing
- Bring as much as possible with you:
  - Won't that slow me down?
- Bring:
  - Executable
  - Environment
  - Parameters (using parameter files)
  - Random numbers (generate before-submission)



# Scaling Out to OSG: Rules of Thumb

---

- CPU (single-threaded)
  - Best jobs run between **5 min** and **2 hours**
    - Upper limit somewhat soft
- Disk
  - Keep scratch working space < 20 GB
  - Intermediate needs (/tmp?)
  - Submit disk: think total and I/O transfer
- **Batch or Break Up Jobs**
- Use squid caching where appropriate

# The expanding onion

---

- Laptop (1 machine)
  - You control everything!
- Local cluster (1000 cores)
  - You can ask an admin nicely
- Campus (5000 cores)
  - It better be important/generalizable
- OSG (50,000 cores)
  - Good luck finding the pool admins

# Bringing It With You: MATLAB

---

- What's the problem with MATLAB?
  - license limitations
- What's the solution?
  - “compiling”
- Similar measures for other interpreter languages (R, Python, etc)

# How to bring MATLAB along

---

1) Purchase & install MATLAB, which comes with a “compiler”

2) Run compiler as follows: (online guides)

```
$ mcc -m -R -singleCompThread -R -nodisplay -R -nojvm -nocache foo.m
```

3) This creates run\_foo.sh (et. al.)

4) Create tarball of the runtime

```
$ cd /usr/local/mathworks-R2014b
```

```
$ tar cvzf ~/matlab.tgz ../mathworks-R2014b
```

## 5) Edit the run\_foo.sh

```
tar xzf matlab.tgz
mkdir cache
chmod 0777 cache
export MCR_CACHE_ROOT=`pwd`/cache
```

## Make a submit file:

```
universe = vanilla
executable = run_foo.sh
arguments = ./mathworks-R2014b
should_transfer_files = yes
when_to_transfer_output = on_exit
transfer_input_files = matlab.tgz, foo
queue
```

# Make It Work Everytime

---

- What could possibly go wrong?
  - Eviction
  - Non-existent dependencies
  - File corruption
  - Performance surprises
    - Network
    - Disk
    - ...
  - *Maybe* even a bug in your code

# Self-Checkpointing (for long jobs and shish-kebabs)

---

## 1. Changes to your code

- Save information about progress to a new file, at least every 60 minutes
- At the beginning of code:
  - If progress file exists, start from where the program (or script) left off
  - Otherwise, start from the beginning

## 2. Change to submit file:

`when_to_transfer_output = ON_EXIT_OR_EVICT`

# Error Checks Are Essential

---

- If you don't check, it will happen...
- Check expected file existence (transfer or creation), and repeat with a finite loop
  - better yet, check *rough* file size too
- Advanced:
  - Error-check with wrapper-RETRY combo (RETRY for *specific* error codes from wrapper)



# What to do if a check fails

---

- Understand something about failure
- Use DAG “RETRY”, when useful
- Let the rescue dag continue...

# Performance Surprises

---

One bad node can ruin your whole day

- “Black Hole” machines
  - GLIDEIN whitelist/blacklist if a site is somehow ‘bad’. (But talk to the GOC first!)
- *REALLY* slow machines
  - Use `periodic_hold` / `periodic_release`

# Make It Work Unattended

---

- Remember the ultimate goal:  
Automation! Time savings!
- Need to automate:
  - Data collection?
  - Data cleansing
  - Submission (condor cron)
  - Analysis and verification
  - LaTeX and paper submission ☺

# Make *Science* Work Unattended?



Well, maybe not, but a scientist can dream ...

# Make It Run(-able) for Someone Else

---

- If others can't reproduce your work, it isn't real science!
  - Work hard to make this happen.
  - It's *their* throughput, too.

Only ~10% of published cancer research  
is reproducible!

(Yet another argument for automation)

# Documentation at Multiple Levels

---

- In job files: comment lines
  - submit files, wrapper scripts, executables
- In README files
  - describe file purposes
  - define overall workflow, justifications
- In a Document
  - draw the workflow!

# Make It Run Faster? Maybe.

---

Throughput, throughput, throughput

- Resource reductions (match more slots!)
- Wall-time reductions
  - if significant *per workflow*
  - Why not *per job*?

Think in orders of magnitude:

- Say you have 1000 hour-long jobs that are matched at a rate of 1 per minute ...

*Waste the computer's time, not yours.*

# Maybe Not Worth It

---

## Maybe Not Worth It:

- Rewriting your code in a different language
- Targeting “faster” machines
- Targeting machines that will run longer jobs
- Others?

*Waste the computer's time, not yours.*



# Testing, Testing, 1-2-3 ...

---


- ALWAYS test a subset after making changes
  - How big of a change needs retesting?
- Scale up gradually
- Avoid making problems for others (and for yourself)



# If this were a test...

- 20 points for finishing at all
- 10 points for the right answer
- 1 point for every error check
- 1 point per documentation line

*Out of 100 points? 200 points?*



# Questions?

---

- Feel free to contact me:
  - [lmichael@wisc.edu](mailto:lmichael@wisc.edu)
- Now: Break
  - 10:30-10:45am
- Next:
  - **10:45am-12:15pm: Exercises 6.2, 6.3**
  - 12:15pm: Lunch
  - 1:15-2:30pm: Principles of High-Throughput Computing