

# Introduction to HTC

2014 OSG User School, Monday, Lecture 1

Greg Thain

University of Wisconsin–  
Madison

Center For High Throughput  
Computing



# Welcome!



# Why Are We Here?

# *Transform* Your Research With Computing



# Overview

# Overview of Week

---

---

## Monday

- High Throughput Computing locally
- Miscellaneous
  - Survey
  - UW reimbursement form
  - Certificate

---

## Tuesday

- Distributed High Throughput Computing
- Security
- Tour of Wisconsin Institutes for Discovery

---

## Wednesday

- Distributed storage
- Practical issues with DHTC

---

## Thursday

- From science to production
  - Principles of HTC
  - HTC Showcase
  - Next steps
-

# Overview of a Day

---

- Short introductory lectures
- Lots of hands-on exercises
- Some demos, interactive sessions, etc.
- Optional evening sessions
  - ▶ Monday – Wednesday, 7–9 p.m.
  - ▶ Union South (check TITU)
  - School staff on hand
  - ▶

# Keys to Success

---

- Work hard
- Ask questions!
  - ... during lectures
  - ... during exercises
  - ... during breaks
  - ... during meals
  - ... in person is best, email is OK
- If we do not know an answer, we will try to find the person who does





# Ready?

One Thing

Computing is Cheap!

# Goals For This Session


---

- Understand the basics of High Throughput Computing
- Understand a few things about HTCondor, which is one kind of HTC system
- Use basic HTCondor commands
- Run a job locally!

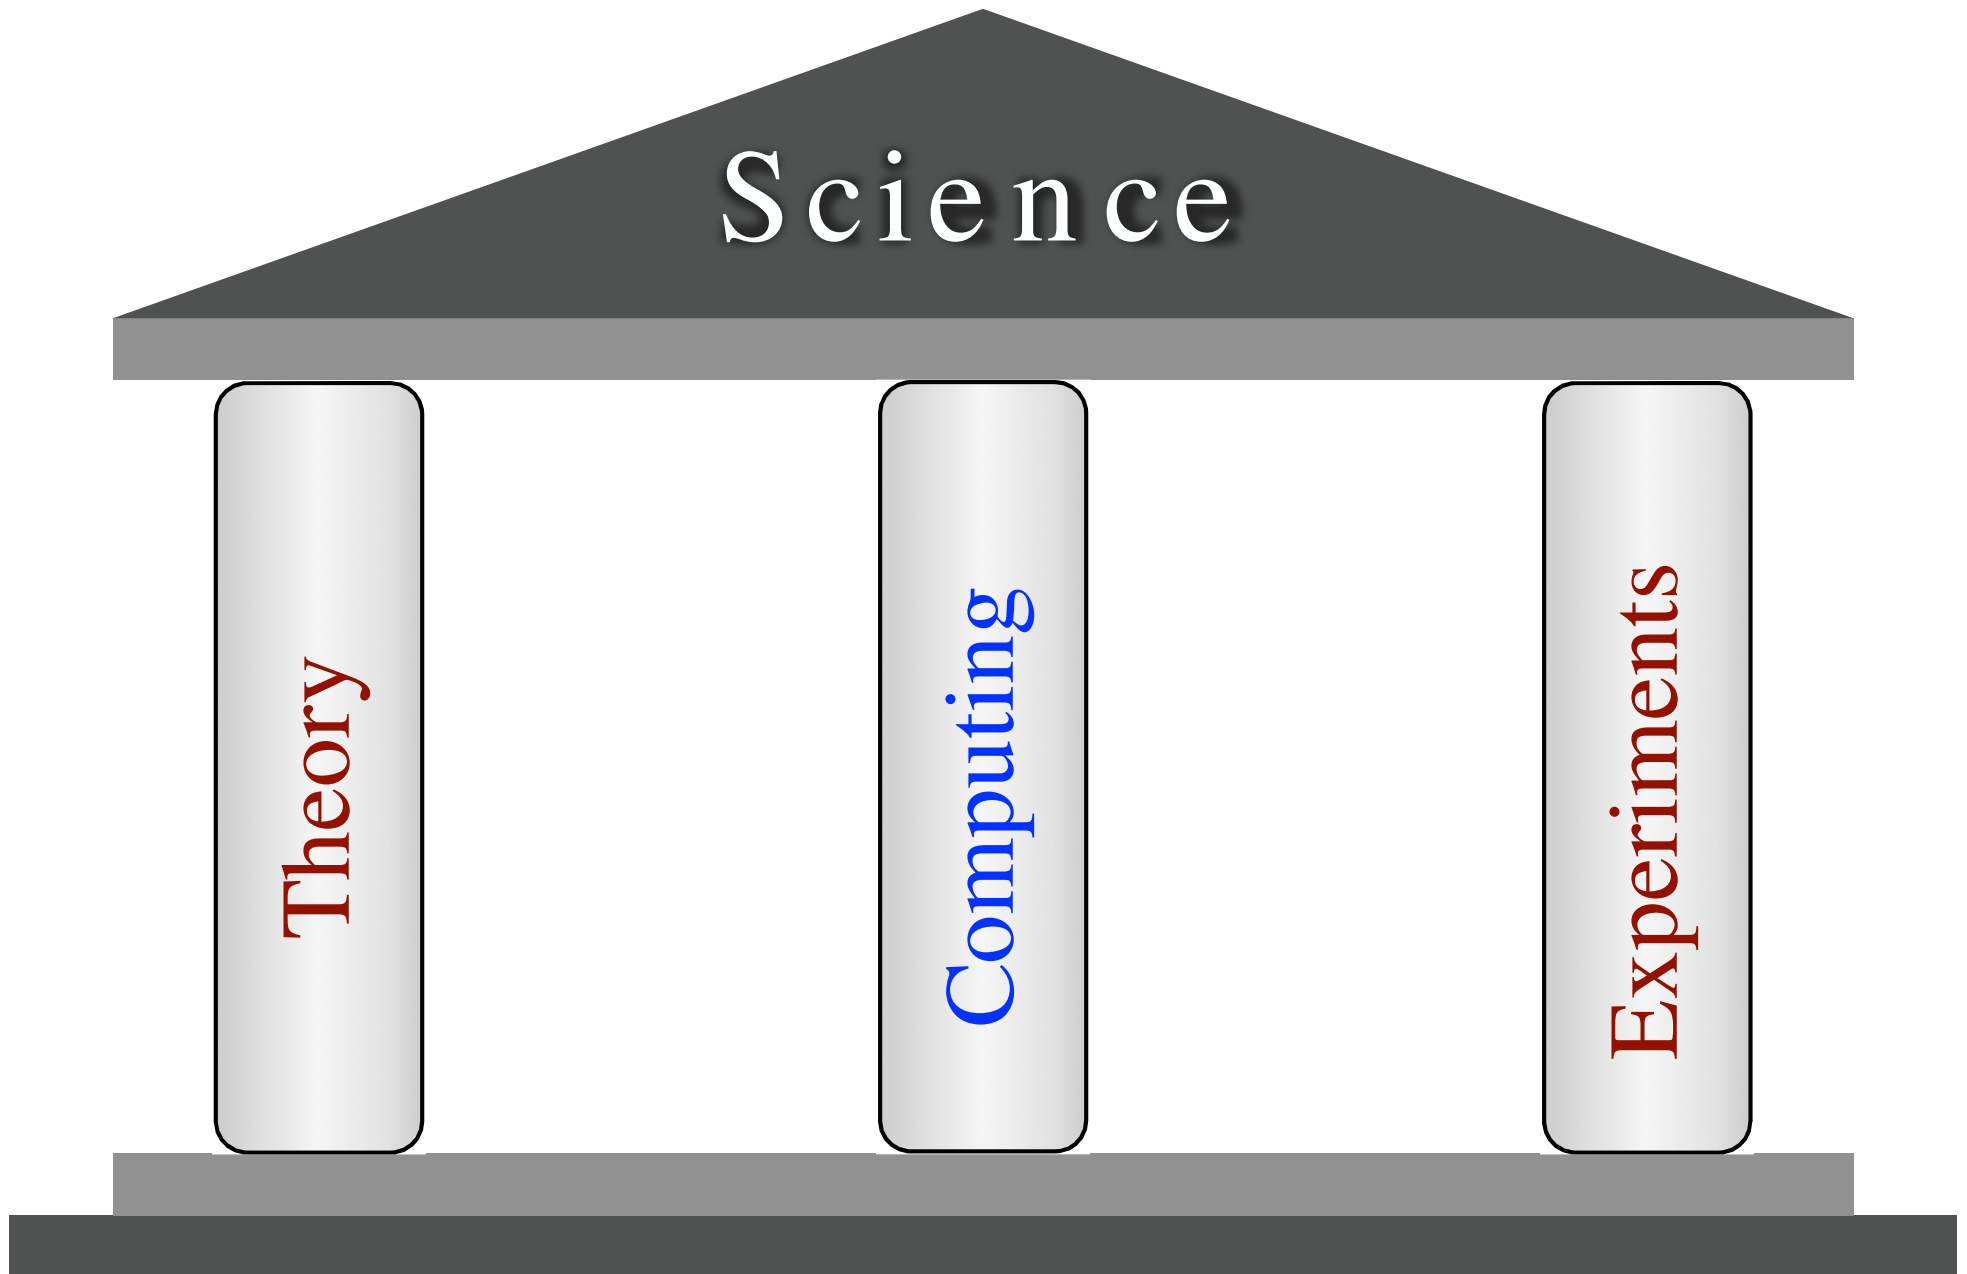


# Why HTC?

# Computing in Science



# Computing in Science



# Example Challenge

---

- You have a program to run (simulation, Monte Carlo, data analysis, image analysis, stats, ...)
- Each run takes about 1 hour
- You want to run the program  $8 \times 12 \times 100$  times
- 9600 hours  $\approx$  1.1 years ... running nonstop!
- Conference is next week



# Distributed Computing

- Use many computers to perform 1 computation
- Example:
  - ▶ 2 computers  $\Rightarrow$  4,800 hours  $\approx$  1/2 year
  - ▶ 8 computers  $\Rightarrow$  1,200 hours  $\approx$  2 months
  - ▶ 100 computers  $\Rightarrow$  96 hours = 4 days
  - ▶ 9,600 computers  $\Rightarrow$  1 hour! (but ...)

These computers are no faster than your laptop!

# Performance vs. Throughput

- High *Performance* Computing (HPC)
  - Focus on biggest, fastest systems (supercomputers)
  - Maximize operations per *second*
  - Often requires special code
  - Often must request and wait for access
- High *Throughput* Computing (HTC)
  - Focus on using all resources, reliably, all the time
  - Maximize operations per *year*
  - Use any kind of computer, even old, slow ones
  - Must break task into separate, independent parts
  - Access varies by availability, usage, etc.

# HPC vs HTC: An Analogy






# HPC vs HTC: An Analogy



- Our local HTC systems
- Recent CPU hours:
  - ~ 280,000 / day
  - ~ 8.3 million / month
  - ~ 78 million / year



# Open Science Grid

- HTC scaled way up
  - ▶ Over 110 sites
  - ▶ Mostly U.S.
  - ▶ Some others
  - ▶ Past year:
    - ~170 million jobs
    - ~770 million CPUhours
    - ~372 petabytes transferred
- Can submit jobs locally, move to OSG
- <http://www.opensciencegrid.org/>



# Other Distributed Computing

- Other systems to manage a local cluster:
  - PBS/Torque
  - LSF
  - Sun Grid Engine/Oracle Grid Engine
  - SLURM
- Other wide-area systems:
  - European Grid Infrastructure
  - Other national and regional grids
  - Commercial cloud systems used to augment grids
- HPC
  - Various supercomputers (e.g., TOP500 list)
  - XSEDE



# HTCondor



# HTCondor History and Status

- History
  - ▶ Started in 1988 as a “cycle scavenger”
  - ▶ Protected interests of users and machine owners
- Today
  - ▶ Expanded to become CHTC team: 20+ full-time staff
  - ▶ Current production release: HTCondor 8.0.0
  - ▶ HTCondor software: ~700,000 lines of C/C++ code
- Miron Livny
  - ▶ Professor, UW–Madison CompSci
  - ▶ Director, CHTC
  - ▶ Dir. of Core Comp. Tech.,
  - ▶ WID/MIR Tech. Director & PI, OSG



# HTCondor Functions

---

- Users
  - Define jobs, their requirements, and preferences
  - Submit and cancel jobs
  - Check on the state of a job
  - Check on the state of the machines
- Administrators
  - Configure and control the HTCondor system
  - Declare policies on machine use, pool use, etc.
- Internally
  - Match jobs to machines (enforcing all policies)
  - Track and manage machines
  - Track and run jobs

# Terminology: *Job*

- ***Job***: A computer program *or* one run of it
- *Not* interactive, *no* GUI (e.g., not Word or email)  
(How could you interact with 1,000 programs running at once?)
  1. Input: command-line arguments and/or files
  2. Run: do stuff
  3. Output: standard output & error and/or files
- Scheduling
  - User decides when to *submit* job to be run
  - System decides when to run job, based on policy

# Terminology: *Machine*, *Slot*

- *Machine*
  - A *machine* is a physical computer (typically)
  - May have multiple *processors* (computer chips)
  - One processor may have multiple *cores* (CPUs)
- HTCondor: *Slot*
  - One assignable unit of a machine (i.e., 1 job per slot)
  - Most often, corresponds to one core
  - Thus, typical machines today have 4–40 slots
- Advanced HTCondor feature: Can get 1 slot with many cores on 1 machine, for MPI(-like) jobs

# Terminology: Matchmaking

Two-way process of finding a slot for a job

Jobs have requirements and preferences

*E.g.: I need Red Hat Linux 6 and 100 GB of disk space, I and*

*prefer to get as much memory as possible*

Machines have requirements and preferences I

*E.g.: I run jobs only from users in the Comp. Sci. dept., and prefer to run ones that ask for a lot of memory*

*Important jobs may replace less important ones*

Thus: Not as simple as waiting in a line!



# Running a Job

# Viewing Slots

## condor\_status

- With no arguments, lists *all* slots currently in pool
- Summary info is printed at the end of the list
- For more info: exercises, **-h**, manual, next lecture

<a href="#">slot6@opt-a001.cht</a>	LINUX	X86_64	Claimed	Busy	1.000	1024	0+19:09:32
<a href="#">slot7@opt-a001.cht</a>	LINUX	X86_64	Claimed	Busy	1.000	1024	0+19:09:31
<a href="#">slot8@opt-a001.cht</a>	LINUX	X86_64	Unclaimed	Idle	1.000	1024	0+17:37:54
<a href="#">slot9@opt-a001.cht</a>	LINUX	X86_64	Claimed	Busy	1.000	1024	0+19:09:32
<a href="#">slot10@opt-a002.ch</a>	LINUX	X86_64	Unclaimed	Idle	0.000	1024	0+17:55:15
<a href="#">slot11@opt-a002.ch</a>	LINUX	X86_64	Unclaimed	Idle	0.000	1024	0+17:55:16

	Total	Owner	Claimed	Unclaimed	Matched	Preempting	Backfill
INTEL/WINNT51	2	0	0	2	0	0	0
INTEL/WINNT61	52	2	0	50	0	0	0
X86_64/LINUX	2086	544	1258	284	0	0	0
Total	2140	546	1258	336	0	0	0

# Viewing Jobs

## condor\_q

- With no args, lists all jobs waiting or running here
- For more info: exercises, **-h**, manual, next lecture

-- Submitter: osg-ss-submit.chtc.wisc.edu : <...> : ...

ID	OWNER	SUBMITTED	RUN_TIME	ST	PRI	SIZE	CMD
6.0	cat	11/12 09:30	0+00:00:00	I	0	0.0	explore.py
6.1	cat	11/12 09:30	0+00:00:00	I	0	0.0	explore.py
6.2	cat	11/12 09:30	0+00:00:00	I	0	0.0	explore.py
6.3	cat	11/12 09:30	0+00:00:00	I	0	0.0	explore.py
6.4	cat	11/12 09:30	0+00:00:00	I	0	0.0	explore.py

5 jobs; 5 idle, 0 running, 0 held



# Basic Submit File

```
executable = word_freq.py
universe = vanilla
arguments = "words.txt 1000"

output = word_freq.out
error = word_freq.err
log = word_freq.log

should_transfer_files = YES
when_to_transfer_output = ON_EXIT
transfer_input_files = words.txt

queue
```

# Basic Submit File

**executable** = word\_freq.py

**universe** = vanilla

**arguments** = "words.txt 1000" Command-line arguments to pass

**output** = word\_freq.out to executable when run; surround with

**error** = word\_freq.err

**log** = word\_freq.log double quotes *[opt]*

**should\_transfer\_files** = YES

**when\_to\_transfer\_output** = ON\_EXIT

**transfer\_input\_files** = words.txt

**queue**

# Basic Submit File

```
executable = word_freq.py  
universe = vanilla  
arguments = "words.txt 1000"
```

```
output = word_freq.out  
error = word_freq.err  
log = word_freq.log
```

Local files that will receive the contents of standard output and error from the run *[opt]*

```
should_transfer_files = YES  
when_to_transfer_output = ON_EXIT  
transfer_input_files = words.txt  
queue
```

# Basic Submit File

```
executable = word_freq.py  
universe = vanilla  
arguments = "words.txt 1000"
```

```
output = word_freq.out  
error = word_freq.err  
log = word_freq.log
```

```
should_transfer_files = YES
```

```
when_to_transfer_output = ON_EXIT
```

```
transfer_input_files = words.txt
```

```
queue
```

Condor's log file  
from running the  
job; very helpful,  
do not omit!

# Basic Submit File

```
executable = word_freq.py  
universe = vanilla  
arguments = "words.txt 1000"
```

```
output = word_freq.out  
error = word_freq.err  
log = word_freq.log
```

```
should_transfer_files = YES  
when_to_transfer_output = ON_EXIT  
transfer_input_files = words.txt  
queue
```

separated list of input files to transfer to machine [opt]

# Basic Submit File

```
executable = word_freq.py  
universe = vanilla  
arguments = "words.txt 1000"  
  
output = word_freq.out  
error = word_freq.err  
log = word_freq.log  
  
should_transfer_files = YES  
when_to_transfer_output = ON_EXIT  
transfer_input_files = words.txt  
  
queue
```

# Submit a Job

## **condor\_submit** *submit-file*

- Submits job to local submit machine
- Use condor\_q to track

**Submitting job(s).**

**1 job(s) submitted to cluster *NNN*.**

Each condor\_submit creates one Cluster

A *job ID* is written as *cluster.process*(e.g., **8.0**)

We will see how to make multiple processes later

# Remove a Job

```
condor_rm cluster [...]  
condor_rm cluster.process [...]
```

- Removes one or more jobs from the queue
- Identify jobs by whole cluster or single job ID
- Only you (*or admin*) can remove your jobs

Cluster **NNN** has been marked for removal.





# Your Turn!

# Thoughts on Exercises

---

- Copy-and-paste is quick, but you may learn more by typing out commands yourself
- Experiment!
  - Try your own variations on the exercises
  - If you have time, try to apply to your own work
- If you do not finish, that's OK — you can make up work later or during evenings, if you like

● If you finish early, try any extra challenges or optional sections, or move ahead to the next section if you are brave

# Sometime Today ...

---

- Sometime today, do the exercise on getting an X.509 personal certificate
- It is **not** required today
- It **will be** required tomorrow afternoon
- It is best to start the process early

# Exercises!

---

- Ask questions!
- Lots of instructors around
- Coming next:

Now – 10:30      Hands-on exercises

10:30–10:45      Break

10:45–11:15      Lecture

11:15–12:15      Hands-on exercises