

Intermediate HTCondor: Workflows

Monday pm

Greg Thain
Center For High Throughput Computing
University of Wisconsin-Madison

Before we begin...

- Any questions on the lectures or exercises up to this point?



Quick Review: 1

1 Job

```
Universe = vanilla  
Executable = runme.sh  
Arguments = 1 2 true
```

```
Output    = out  
Error     = err  
Log       = log
```

```
queue
```

Quick Review: 2

**Many
Jobs**

```
Universe = vanilla
Executable = runme.sh
Arguments = 1 2 true

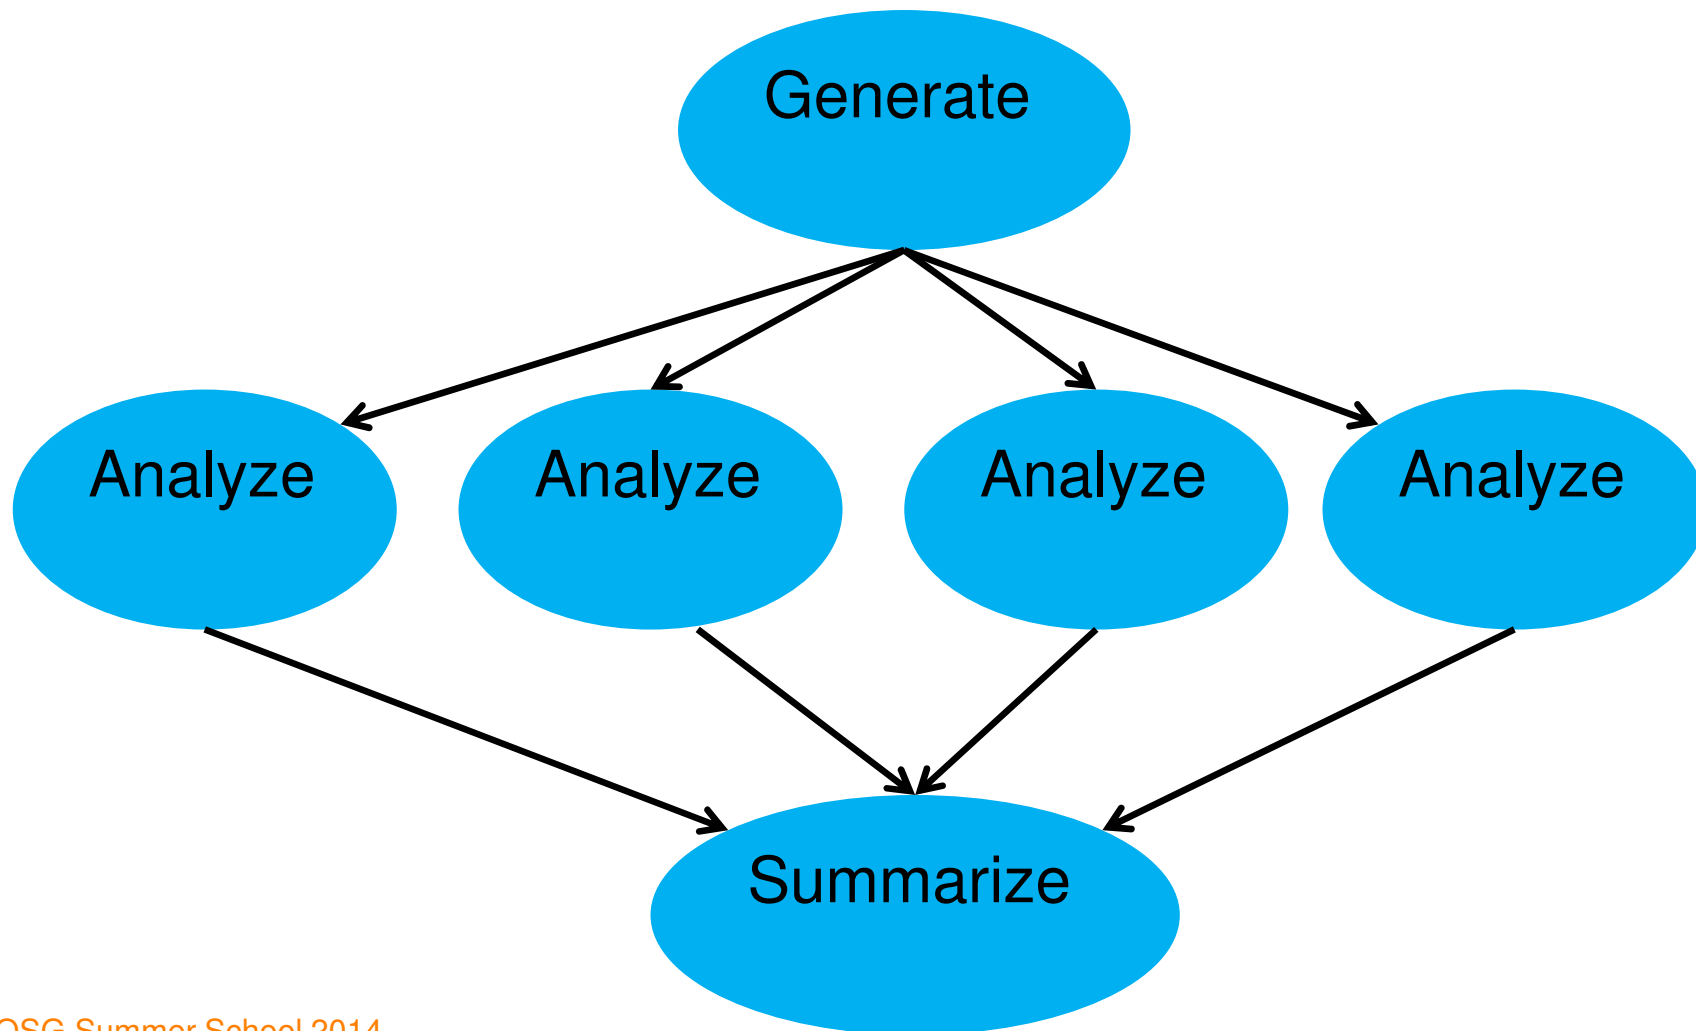
Output    = out.$(PROCESS)
Error     = err.$(PROCESS)
Log       = log.$(PROCESS)

Queue 10000
```

Workflows

- Often, you don't have independent tasks!
- Common example:
 - You want to analyze a set of images
 1. You need to generate N images (once)
 2. You need to analyze all N images
 - One job per image
 3. You need to summarize all results (once)

Do you really want to do this manually?

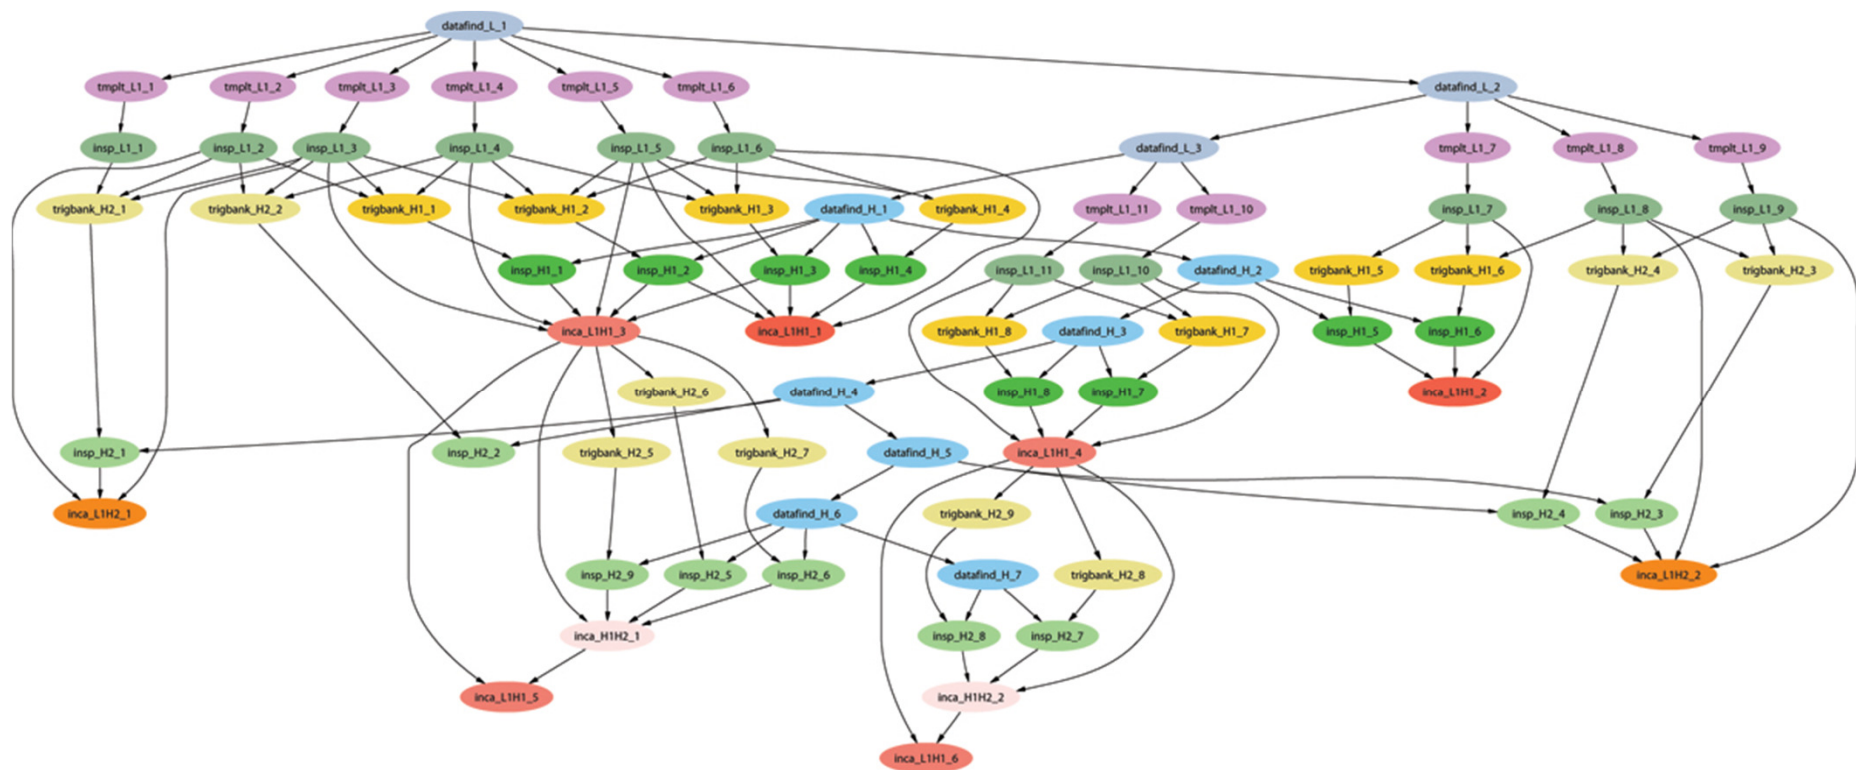


Workflows: The HTC definition

Workflow:

A graph of jobs to run: one or more jobs must **succeed** before one or more others can start running

Example of a LIGO Inspiral DAG





DAGMan

- DAGMan: HTCondor's workflow manager
Directed Acyclic Graph (DAG)
Manager (Man)
- Allows you to specify the dependencies between your HTCondor jobs
- Manages the jobs and their dependencies
- That is, it manages a workflow of HTCondor jobs

What is a DAG?

- A DAG is the structure used by DAGMan to represent these dependencies.
- Each job is a node in the DAG.
- Each node can have any number of “parent” or “children” nodes – as long as there are no loops!




Defining a DAG

- A DAG is defined by a *.dag file*, listing each of its nodes and their dependencies. For example:

```
# Comments are good
Job A a.sub
Job B b.sub
Job C c.sub
Job D d.sub

Parent A Child B C
Parent B C Child D
```



DAG Files....

- This complete DAG has five files

One DAG File:

```
Job A a.sub
Job B b.sub
Job C c.sub
Job D d.sub

Parent A Child B C
Parent B C Child D
```

Four Submit Files:

```
Universe = Vanilla
Executable = analysis...

Universe = ...
```

Submitting a DAG

- To start your DAG, just run `condor_submit_dag` with your .dag file, and HTCondor will start a DAGMan process to manage your jobs:

```
% condor_submit_dag diamond.dag
```

- `condor_submit_dag` submits a Scheduler Universe job with DAGMan as the executable
- Thus the DAGMan daemon itself runs as an HTCondor job, so you don't have to baby-sit it

DAGMan is a HTCondor job


- DAGMan itself is a condor job with a job id, so

```
% condor_rm job_id_of_dagman  
% condor_hold job_id_of_dagman  
% condor_q -dag # is magic
```

- DAGMan submits jobs, one cluster per node
- Don't confuse dagman as job with jobs of dagman


Running a DAG

- DAGMan acts as a job scheduler, managing the submission of your jobs to HTCondor based on the DAG dependencies





Running a DAG (cont'd)

- DAGMan submits jobs to HTCondor at the appropriate times
- For example, after A finishes, it submits B & C



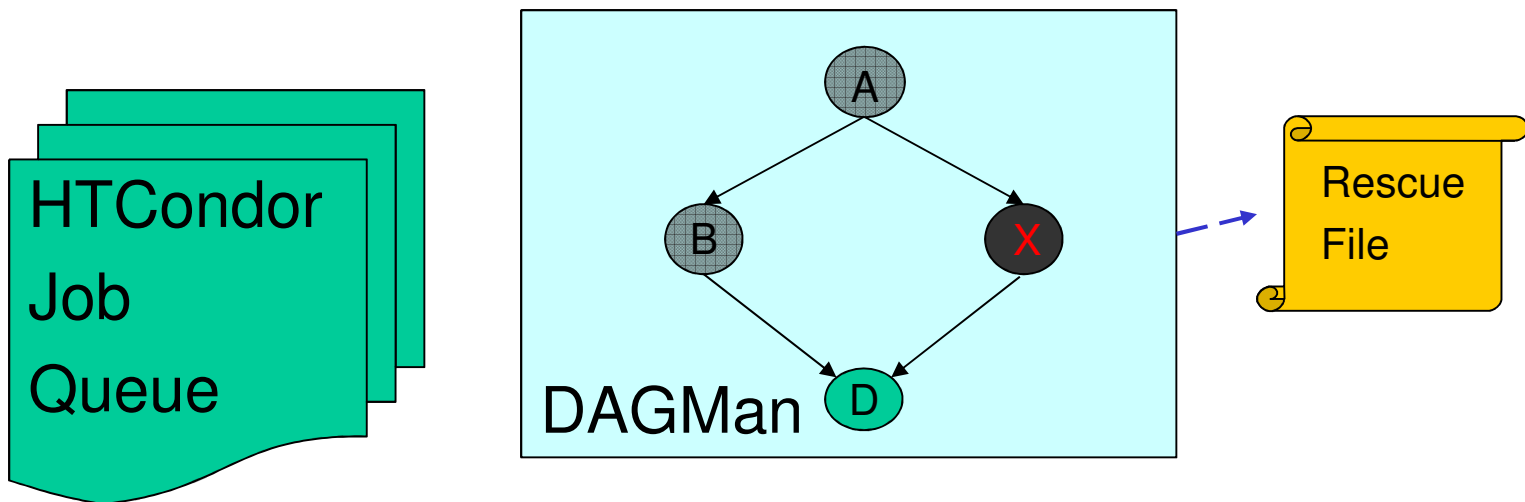
Finishing a DAG

- Once the DAG is complete, the DAGMan job itself is finished, and exits



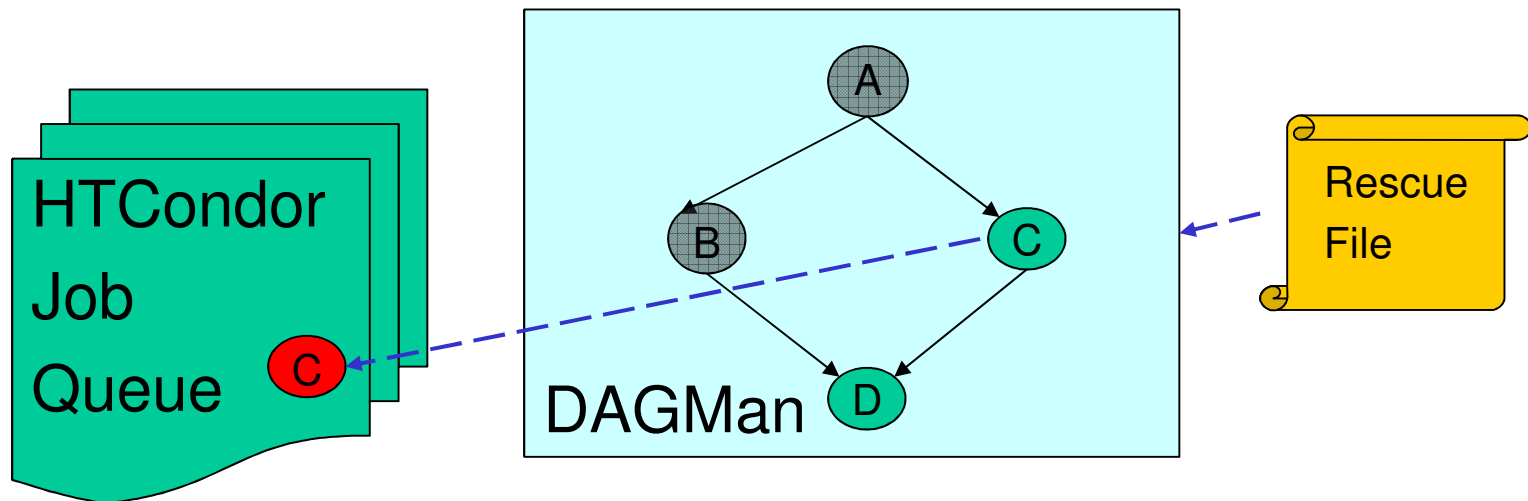
Successes and Failures

- A job *fails* if it exits with a non-zero exit code
- In case of a job failure, DAGMan runs other jobs until it can no longer make progress, and then creates a “*rescue*” file with the current state of the DAG




Recovering a DAG

- Once the failed job is ready to be re-run, the rescue file can be used to restore the prior state of the DAG
 - Another example of reliability for HTC!



Recovering a DAG (cont'd)

- Once that job completes, DAGMan will continue the DAG as if the failure never happened



DAGMan & Fancy Features

- DAGMan doesn't have a lot of “fancy features”
 - No loops
 - Not much assistance in writing very large DAGs (script it yourself)
- Focus is on solid core
 - Add the features people need in order to run large DAGs well
 - People build systems on top of DAGMan

Related Software

Pegasus: <http://pegasus.isi.edu/>

- Writes DAGs based on abstract description
- Runs DAG on appropriate resource (HTCondor, OSG, EC2...)
- Locates data, coordinates execution
- Uses DAGMan, works with large workflows

Makeflow: <http://nd.edu/~ccl/software/makeflow/>

- User writes *make* file, not DAG
- Works with HTCondor, SGE, Work Queue...
- Handles data transfers to remote systems
- Does not use DAGMan

DAGMan: Reliability

- For each job, HTCondor generates a log file
- DAGMan reads this log to see what has happened
- If DAGMan dies (crash, power failure, etc...)
 - HTCondor will restart DAGMan
 - DAGMan re-reads log file
 - DAGMan knows everything it needs to know
 - Principle: DAGMan can recover state from files and without relying on a service (HTCondor queue, database...)
- Recall: HTC requires reliability!

Let's try it out!

- Exercises with DAGMan.



Questions?

- Questions? Comments?
- Feel free to ask me questions later:


Bonus Material...



Use of HTCondor by the LIGO Scientific Collaboration

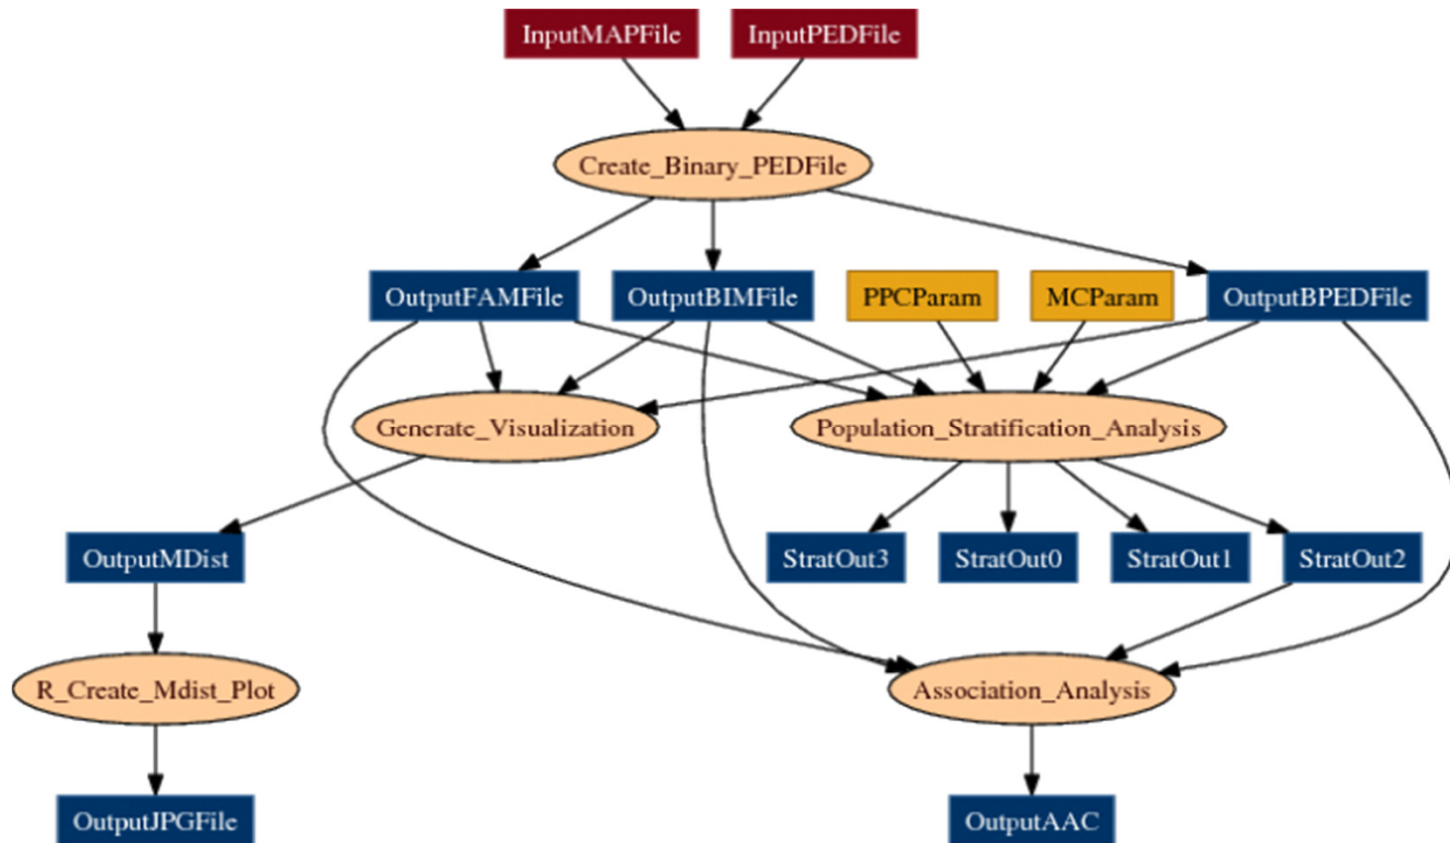


- HTCondor handles 10's of millions of jobs per year running on the LDG, and up to 500k jobs per DAG.
- HTCondor standard universe check pointing widely used, saving us from having to manage this.
- At Caltech, 30 million jobs processed using 22.8 million CPU hrs. on 1324 CPUs in last 30 months.
- For example, to search 1 yr. of data for GWs from the inspiral of binary neutron star and black hole systems takes ~2 million jobs, and months to run on several thousand ~2.6 GHz nodes.



(Statement from 2010—"last 30 months" isn't from now. Also, I think they do up to 1 million jobs per DAG now.)

Example workflow: Bioinformatics



From Mason, Sanders, State (Yale)

http://pegasus.isi.edu/applications/association_test

Example workflow: Astronomy

From Berriman & Good (JPAC)

<http://pegasus.isi.edu/applications/galactic-plane>

Montage Galactic Plane Workflow

