

Introduction to HTC

2015 OSG User School, Monday, Lecture 1

Greg Thain

University of Wisconsin—
Madison

Center For High Throughput
Computing



Welcome!



Why Are We Here?

Transform Your Research With Computing



Overview

Overview of Week

Monday

- High Throughput Computing locally
- Miscellaneous
 - Survey
 - UW reimbursement form
 -

Tuesday

- Distributed High Throughput Computing
- Security
- Tour of Wisconsin Institutes for Discovery

Wednesday

- Distributed storage
- Practical issues with DHTC

Thursday

- From science to production
 - Principles of HTC
 - HTC Showcase
 - Next steps
-

Overview of a Day

- Short introductory lectures
- Lots of hands-on exercises
- Some demos, interactive sessions, etc.
- Optional evening sessions
 - Monday – Wednesday, 7–9 p.m.
 - Union South (check TITU)
 - School staff on hand
 -

Keys to Success

- Work hard
- Ask questions!
 - ... during lectures
 - ... during exercises
 - ... during breaks
 - ... during meals
 - ... in person is best, email is OK
- If we do not know an answer, we will try to find the person who does



Ready?

One Thing

Computing is Cheap!


Goals For This Session

- Understand the basics of High Throughput Computing
- Understand a few things about HTCondor, which is one kind of HTC system
- Use basic HTCondor commands
- Run a job locally!

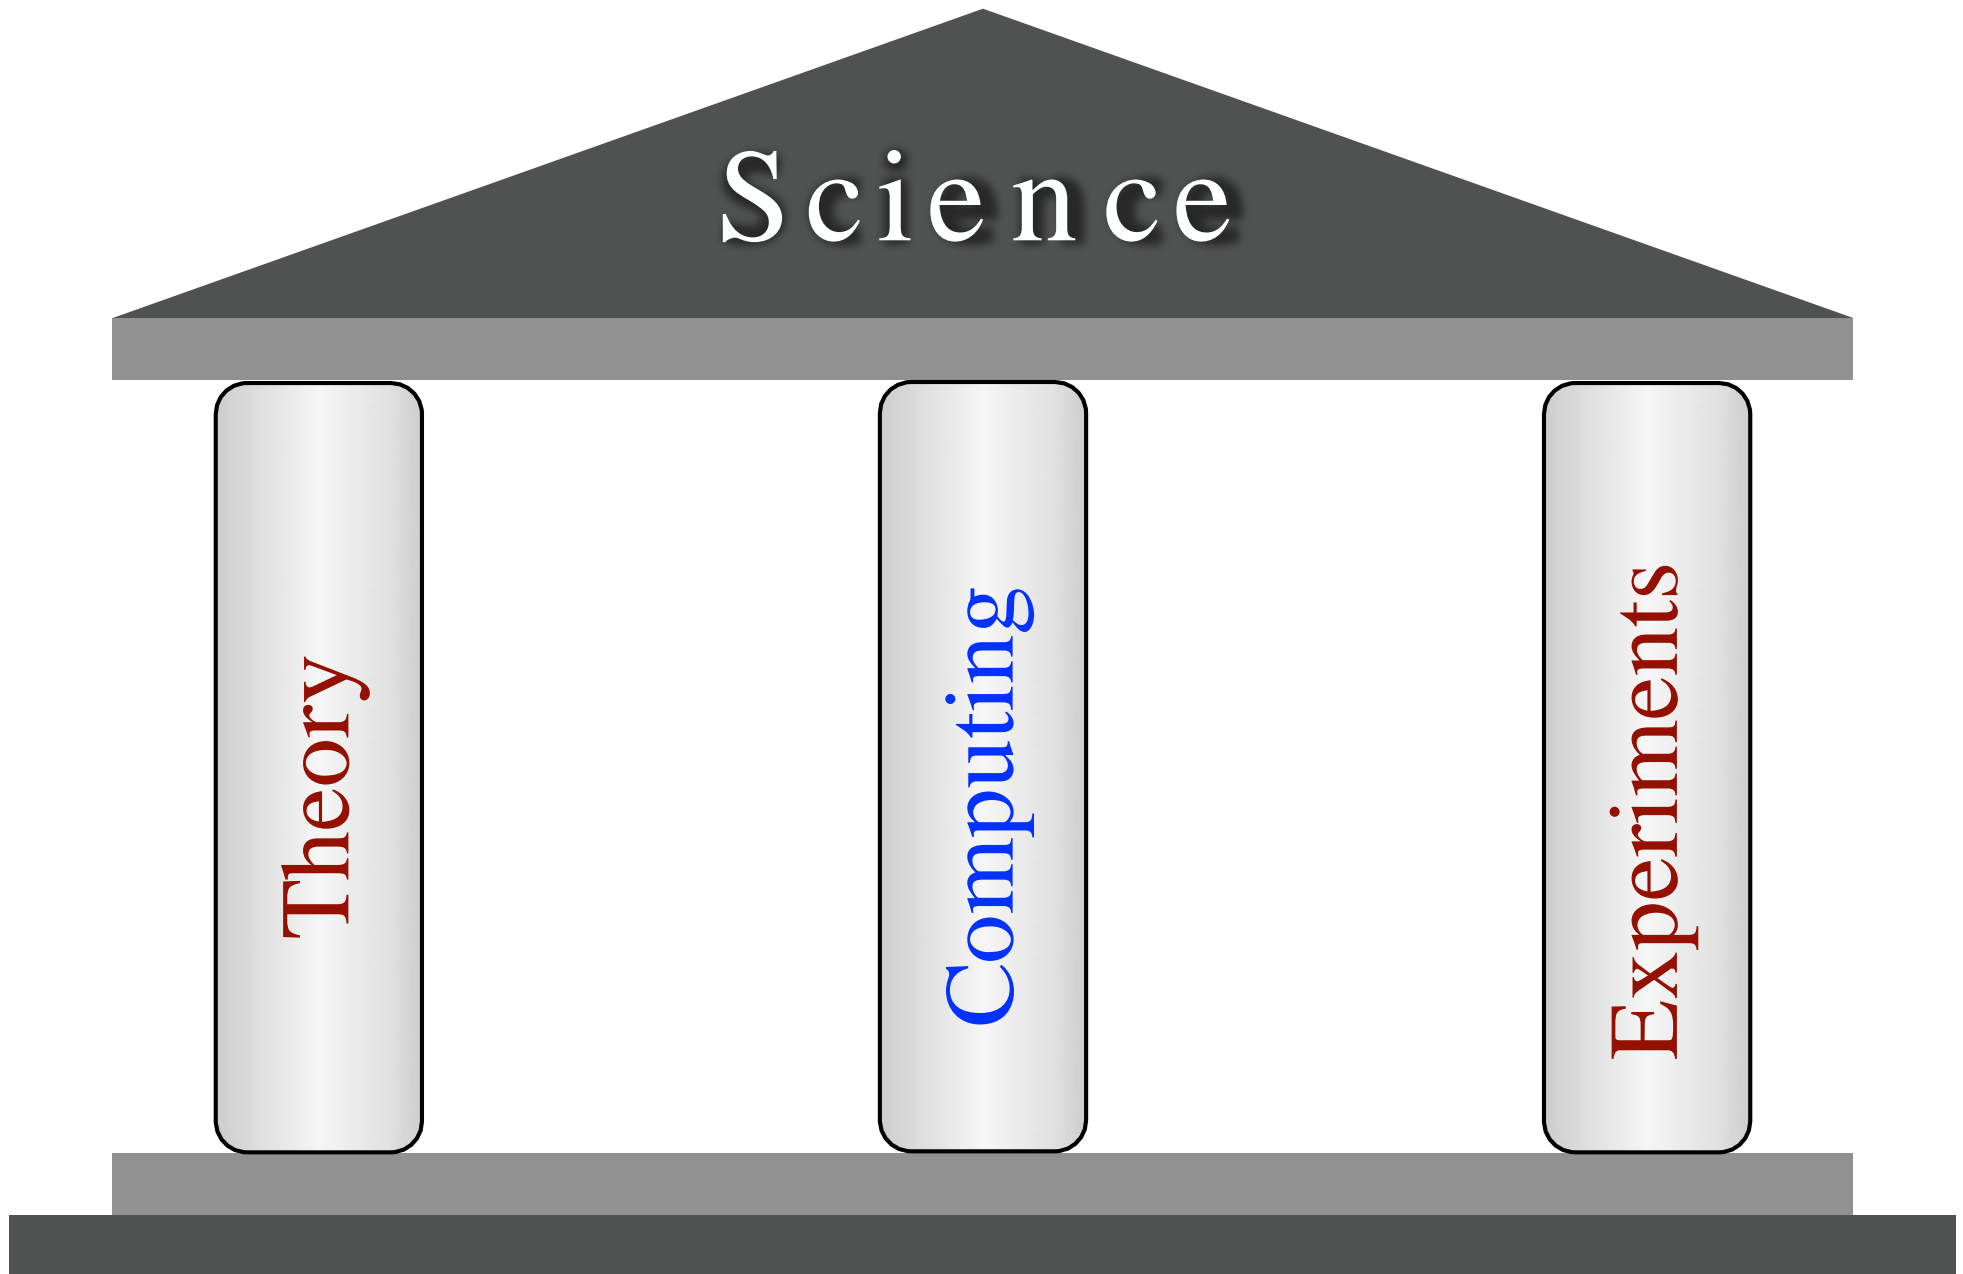


Why HTC?

Computing in Science



Computing in Science



Example Challenge

- You have a program to run (simulation, Monte Carlo, data analysis, image analysis, stats, ...)
- Each run takes about 1 hour
- You want to run the program $8 \times 12 \times 100$ times
- 9600 hours \approx 1.1 years ... running nonstop!
- Conference is next week

Distributed Computing

- Use many computers to perform 1 computation
- Example:
 - ▶ 2 computers \Rightarrow 4,800 hours \approx 1/2 year
 - ▶ 8 computers \Rightarrow 1,200 hours \approx 2 months
 - ▶ 100 computers \Rightarrow 96 hours = 4 days
 - ▶ 9,600 computers \Rightarrow 1 hour! (but ...)

These computers are no faster than your laptop!

Performance vs. Throughput

- High *Performance* Computing (HPC)
 - Focus on biggest, fastest systems (supercomputers)
 - Maximize operations per *second*
 - Often requires special code
 - Often must request and wait for access
- High *Throughput* Computing (HTC)
 - Focus on using all resources, reliably, all the time
 - Maximize operations per *year*
 - Use any kind of computer, even old, slow ones
 - Must break task into separate, independent parts
 - Access varies by availability, usage, etc.


HPC vs HTC: An Analogy



HPC vs HTC: An Analogy



- Our local HTC systems
- Recent CPU hours:
 - ~ 280,000 / day
 - ~ 8.3 million / month
 - ~ 78 million / year



Open Science Grid

- HTC scaled way up
 - Over 110 sites
 - Mostly U.S.
 - Some others
 - Past year:
 - ~170 million jobs
 - ~770 million CPUhours
 - ~372 petabytes transferred
- Can submit jobs locally, move to OSG
- <http://www.opensciencegrid.org/>



Other Distributed Computing

- Other systems to manage a local cluster:
 - PBS/Torque
 - LSF
 - Sun Grid Engine/Oracle Grid Engine
 - SLURM
- Other wide-area systems:
 - European Grid Infrastructure
 - Other national and regional grids
 - Commercial cloud systems used to augment grids
- HPC
 - Various supercomputers (e.g., TOP500 list)
 - XSEDE



HTCondor

HTCondor History and Status

- History
 - Started in 1988 as a “cycle scavenger”
 - Protected interests of users and machine owners
- Today
 - Expanded to become CHTC team: 20+ full-time staff
 - Current production release: HTCondor 8.4.0
 - HTCondor software: ~700,000 lines of C/C++ code
- Miron Livny
 - Professor, UW–Madison CompSci
 - Director, CHTC
 - Dir. of Core Comp. Tech.,
 - WID/MIR Tech. Director & PI, OSG



HTCondor Functions

- Users
 - Define jobs, their requirements, and preferences
 - Submit and cancel jobs
 - Check on the state of a job
 - Check on the state of the machines
- Administrators
 - Configure and control the HTCondor system
 - Declare policies on machine use, pool use, etc.
- Internally
 - Match jobs to machines (enforcing all policies)
 - Track and manage machines
 - Track and run jobs

Terminology: *Job*

- ***Job***: A computer program *or* one run of it
- *Not* interactive, *no* GUI (e.g., not Word or email)
(How could you interact with 1,000 programs running at once?)
 1. Input: command-line arguments and/or files
 2. Run: do stuff
 3. Output: standard output & error and/or files
- Scheduling
 - User decides when to *submit* job to be run
 - System decides when to run job, based on policy

Terminology: *Machine*, *Slot*

- *Machine*
 - A *machine* is a physical computer (typically)
 - May have multiple *processors* (computer chips)
 - One processor may have multiple *cores* (CPUs)
- HTCondor: *Slot*
 - One assignable unit of a machine (i.e., 1 job per slot)
 - Most often, corresponds to one core
 - Thus, typical machines today have 4–40 slots
- Advanced HTCondor feature: Can get 1 slot with many cores on 1 machine, for MPI(-like) jobs

Terminology: Matchmaking

Two-way process of finding a slot for a job

Jobs have requirements and preferences

E.g.: I need Red Hat Linux 6 and 100 GB of disk space, I and

prefer to get as much memory as possible

Machines have requirements and preferences I

E.g.: I run jobs only from users in the Comp. Sci. dept., and prefer to run ones that ask for a lot of memory

Important jobs may replace less important ones

Thus: Not as simple as waiting in a line!



Running a Job

Viewing Slots

condor_status

- With no arguments, lists *all* slots currently in pool
- Summary info is printed at the end of the list
- For more info: exercises, **-h**, manual, next lecture

```
slot6@opt-a001.cht LINUX      X86_64 Claimed   Busy      1.000  1024  0+19:09:32
slot7@opt-a001.cht LINUX      X86_64 Claimed   Busy      1.000  1024  0+19:09:31
slot8@opt-a001.cht LINUX      X86_64 Unclaimed Idle       1.000  1024  0+17:37:54
slot9@opt-a001.cht LINUX      X86_64 Claimed   Busy      1.000  1024  0+19:09:32
slot10@opt-a002.ch LINUX      X86_64 Unclaimed Idle       0.000  1024  0+17:55:15
slot11@opt-a002.ch LINUX      X86_64 Unclaimed Idle       0.000  1024  0+17:55:16
```

	Total	Owner	Claimed	Unclaimed	Matched	Preempting	Backfill
NTEL/WINNT51	2	0	0	2	0	0	0
INTEL/WINNT61	52	2	0	50	0	0	0
X86_64/LINUX	2086	544	1258	284	0	0	0
Total	2140	546	1258	336			

Viewing Jobs

condor_q

- With no args, lists all jobs waiting or running here
- For more info: exercises, **-h**, manual, next lecture

-- Submitter: osg-ss-submit.chtc.wisc.edu : <...> : ...

ID	OWNER	SUBMITTED	RUN_TIME	ST	PRI	SIZE	CMD
6.0	cat	11/12 09:30	0+00:00:00	I	0	0.0	explore.py
6.1	cat	11/12 09:30	0+00:00:00	I	0	0.0	explore.py
6.2	cat	11/12 09:30	0+00:00:00	I	0	0.0	explore.py
6.3	cat	11/12 09:30	0+00:00:00	I	0	0.0	explore.py
6.4	cat	11/12 09:30	0+00:00:00	I	0	0.0	explore.py

5 jobs; 5 idle, 0 running, 0 held

Basic Submit File

```
executable = word_freq.py
universe = vanilla
arguments = "words.txt 1000"

output = word_freq.out
error = word_freq.err
log = word_freq.log

should_transfer_files = YES
when_to_transfer_output = ON_EXIT
transfer_input_files = words.txt

queue
```

Submit a Job

condor_submit *submit-file*

- Submits job to local submit machine
- Use condor_q to track

Submitting job(s).

1 job(s) submitted to cluster *NNN*.

Each condor_submit creates one Cluster

A *job ID* is written as **cluster**.**process**(e.g., **8.0**)

We will see how to make multiple processes later

Remove a Job

```
condor_rm cluster [...]  
condor_rm cluster.process [...]
```

- Removes one or more jobs from the queue
- Identify jobs by whole cluster or single job ID
- Only you (*or admin*) can remove your jobs

Cluster **NNN** has been marked for removal.



Your Turn!

Thoughts on Exercises

- Copy-and-paste is quick, but you may learn more by typing out commands yourself
- Experiment!
 - Try your own variations on the exercises
 - If you have time, try to apply to your own work
- If you do not finish, that's OK — you can make up work later or during evenings, if you like

• If you finish early, try any extra challenges or optional sections, or move ahead to the next section if you are brave

Exercises!

- Ask questions!
- Lots of instructors around
- Coming next:

Now – 10:30 Hands-on exercises

10:30–10:45 Break

10:45–11:15 Lecture

11:15–12:15 Hands-on exercises