

# Data Considerations

Thursday AM, Lecture 2

Lauren Michael

CHTC, UW-Madison

# Overview – Data Handling

---

- Review of HTCondor Data Handling
- What is ‘Large’ Data?
- Data Management Tips
- Next talks: Dealing with Large Data


# Overview – Data Handling

---

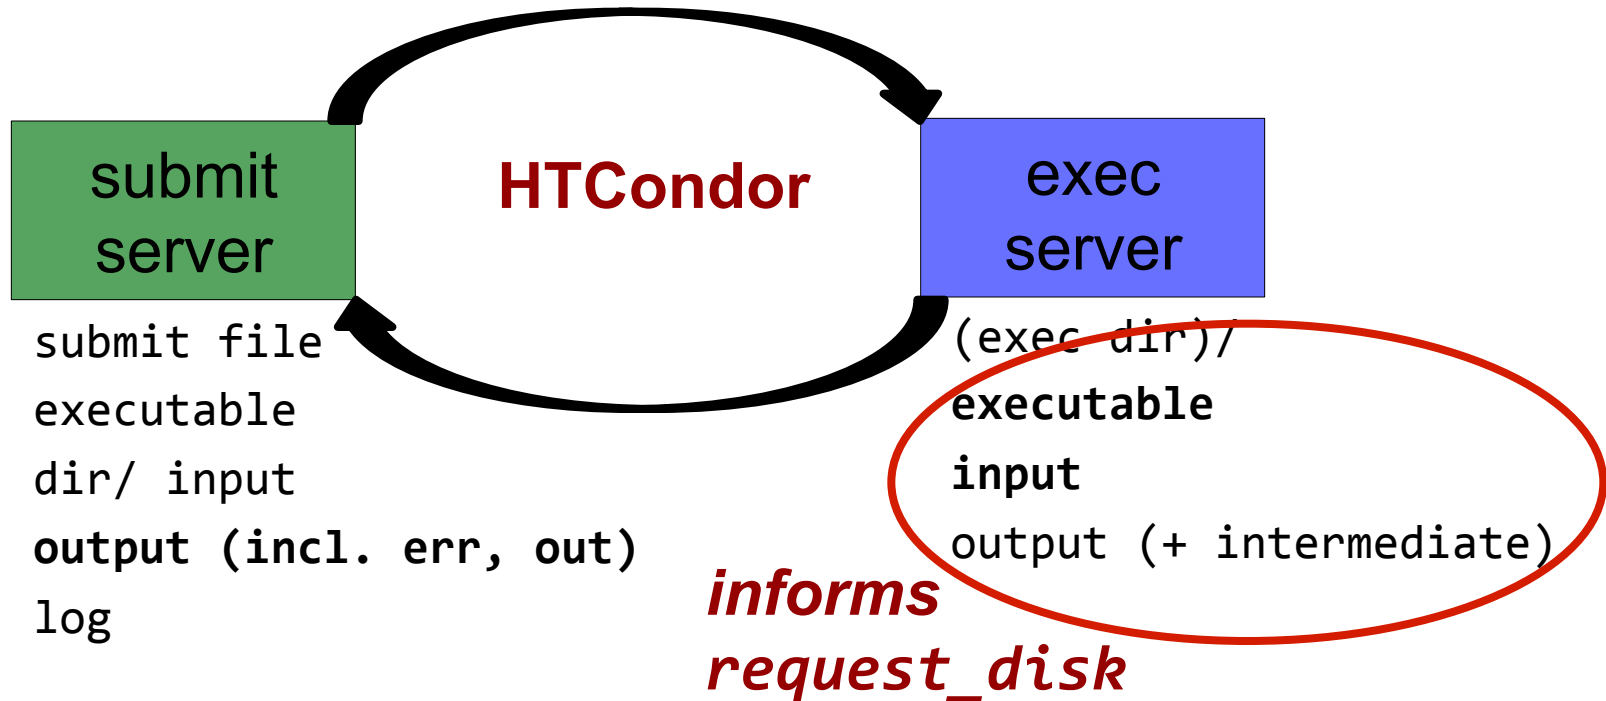
- **Review of HTCondor Data Handling**
- What is ‘Large’ Data?
- Data Management Tips
- Next talks: Dealing with Large Data




# Review: HTCondor Data Handling



# Server-specific data needs



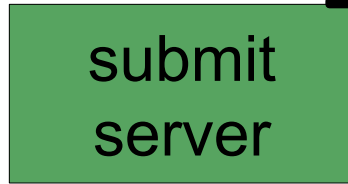
# Network needs: the submit server bottleneck






# Network needs: the submit server bottleneck

*Input transfers  
for many jobs  
will coincide*



submit file  
executable  
dir/ input  
output

**HTCondor**

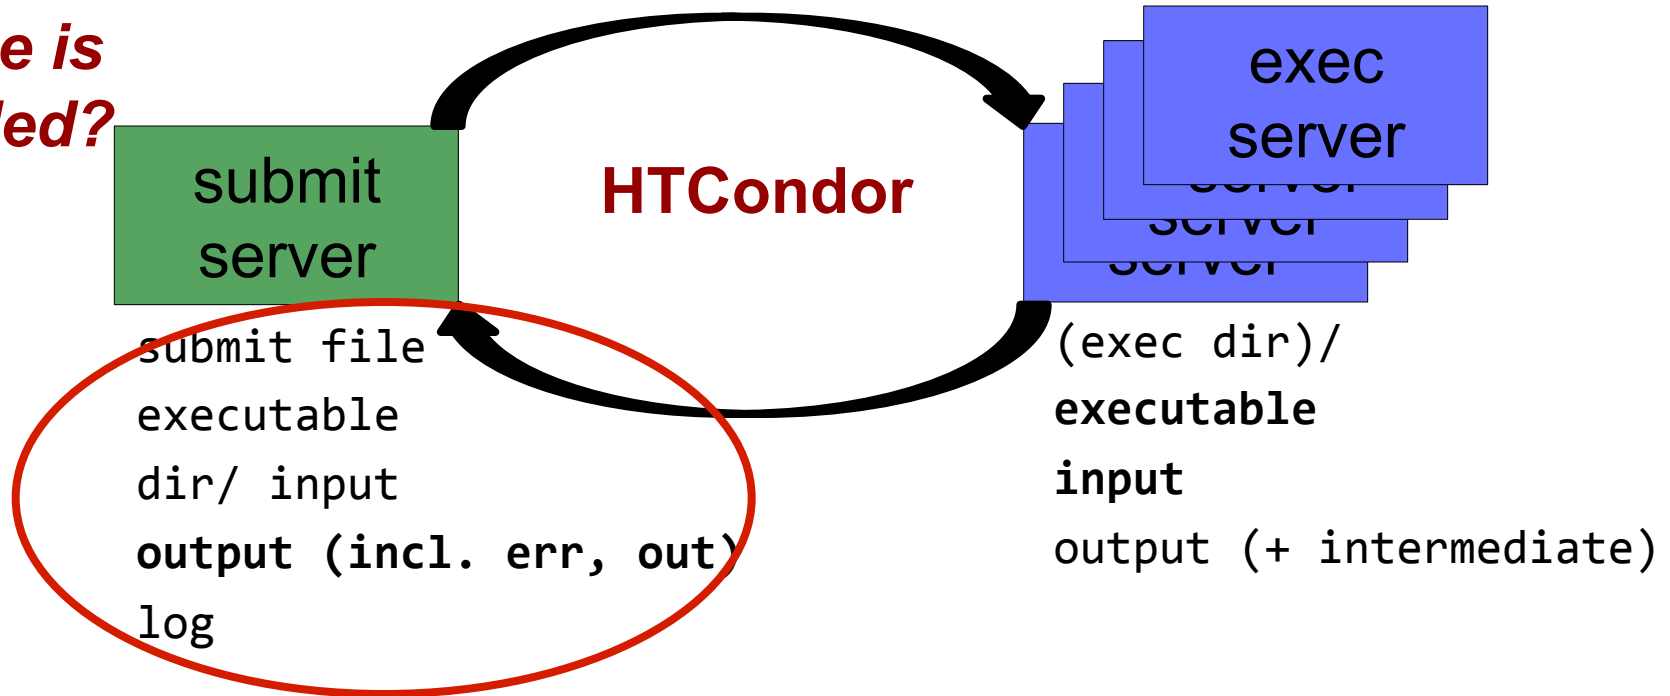


*Output  
transfers are  
staggered*

(exec dir)/  
executable  
input  
output

# Server-specific data needs

*What total disk space is needed?*





# Overview – Data Handling

---

- Review of HTCondor Data Handling
- **What is ‘Large’ Data?**
- Data Management Tips
- Next talks: Dealing with Large Data

# What is ~~big~~ large data?

---

- For researchers “big data” is relative
  - What is ‘big’ for you? Why?

# What is ~~big~~ large data?

---

- For researchers “big data” is relative
  - What is ‘big’ for you? Why?
- Volume, velocity, variety!
  - think: a million 1-KB files, versus one 1-GB file


# HTCondor transfer limits

*Input transfers  
for many jobs  
will coincide*



submit file  
executable  
dir/ input  
output

**HTCondor**



*Output  
transfers are  
staggered*

(exec dir)/  
executable  
input  
output

# Per-job transfer limits

*Input transfers  
for many jobs  
will coincide*

submit  
server

submit file  
executable  
dir/ input  
output

HTCondor

exec  
server  
server  
server  
server

(exec dir)/  
executable

input  
output

*Output  
transfers are  
staggered*

# 'Large' input: The collaborator analogy



# 'Large' input: The collaborator analogy

amount	method of delivery
words	
tiny – 10MB	
10MB – GBs	
TBs	

# 'Large' input: The collaborator analogy

amount	method of delivery
words	email body
tiny – 10MB	
10MB – GBs	
TBs	



# 'Large' input: The collaborator analogy

amount	method of delivery
words	email body
tiny – 10MB	email attachment (managed transfer)
10MB – GBs	
TBs	


# 'Large' input: The collaborator analogy

amount	method of delivery
words	email body
tiny – 10MB	email attachment (managed transfer)
10MB – GBs	download from Google Drive, Drop/Box, other web-accessible server
TBs	

# 'Large' input: The collaborator analogy


amount	method of delivery
words	email body
tiny – 10MB	email attachment (managed transfer)
10MB – GBs	download from Google Drive, Drop/Box, other web-accessible server
TBs	ship an external drive (local copy needed)

# Large input in HTC and OSG



file size	method of delivery
words	within executable or arguments?
tiny – 10MB per file	HTCondor file transfer (up to 1GB total per-job)
10MB – 1GB, shared	download from web proxy (network-accessible server)
1GB - 10GB, unique or shared	StashCache (regional replication)
10 GB - TBs	shared file system (local copy, local execute servers)

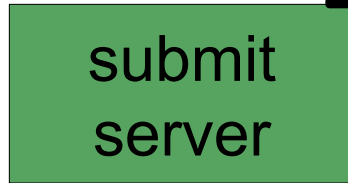
# Large input in HTC and OSG



file size	method of delivery
words	within executable or arguments?
<b>tiny – 10MB per file</b>	<b>HTCondor file transfer (up to 1GB total per-job)</b>
10MB – 1GB, shared	download from web proxy (network-accessible server)
1GB - 10GB, unique or shared	StashCache (regional replication)
10 GB - TBs	shared file system (local copy, local execute servers)


# HTCondor transfer limits

*Input transfers  
for many jobs  
will coincide*



submit file  
executable  
dir/ input  
output


**HTCondor**




*Output  
transfers are  
staggered*

(exec dir)/  
executable  
input  
output

# Per-job transfer limits



# Per-job transfer limits





# Output for HTC and OSG



file size	method of delivery
words	<del>within executable or arguments?</del>
tiny – <b><u>1GB</u></b>	HTCondor file transfer (up to 1 GB total per-job)
1GB+	shared file system (local copy, local execute servers)

# Output for HTC and OSG



amount	method of delivery
words	<del>within executable or arguments?</del>
tiny – <b><u>1GB</u></b>	HTCondor file transfer (up to 1 GB total per-job)
1GB+	shared file system (local copy, local execute servers)

- Why are there fewer options?

# A OSG-wide File System?

---

## WHY NOT?

- Technology barriers:
  - need local copies everywhere
    - too many users, too much data
  - need unrealistic network
  - security: need global user accounts
- Social barriers:
  - local sites unwilling/able to do the above

# Overview – Data Handling

---

- Review of HTCondor Data Handling
- What is ‘Large’ Data?
- **Data Management Tips**
- Next talks: Dealing with Large Data

# Data Management Tips


---

- **Know your job needs**
- Know your *batch* needs
- Reduce per-job data needs
- Test, then scale up gradually



# Review: HTCondor Data Handling

***What total disk space is needed?***



# In-Job Data Needs

- “Input” includes:
  - data *and* software
    - executable
    - transfer\_input\_files
  - files downloaded into the job (not transferred by HTCondor)
- “Output” includes:
  - your output, output, error
  - intermediate data (sub-directories and deleted files)
- **Total informs request\_disk, but request a little extra for first tests!**

exec  
server

(exec dir)/  
executable  
input  
output  
(intermediate)

# More on memory and disk requests

---

- Memory might be informed by input file sizes, but this really depends on your software
- A good place to start for tests?
  - If your files are smaller than 1 GB:  
    `request_memory = 1GB`  
    `request_disk = 1GB`
- Be prepared to increase or decrease for a batch!



# Per-job Submit Node Data

- Submit file, log file, dag file
  - only exist on the submit node
  - *usually* fairly small (KBs at most)
- All files transferred *by HTCondor*
  - executable, transfer\_input\_files
- Any files copied back by HTCondor
  - stdout, stderr
  - default: any files remaining at end of job (not subdirectories)

submit  
server

submit file  
executable  
dir/ input  
output (incl. err, out)  
log

# Data Management Tips

---

- Know your job needs
- **Know your *batch* needs**
- Reduce per-job data needs
- Test, then scale up gradually

# In-job Batch Needs

- Per-job values remain the same
  - Same request\_memory and request\_disk as these are per-job values
  - Go by your tests!
  - Some software will vary considerably (e.g. our blastx example, today)

exec  
server

# Submit Node Batch Needs

submit  
server


- Mostly multiplicative!
  - total can get BIG, quickly
- What will *NOT* multiply?
  - “shared” data: software and common input used by every job
  - versus “unique” input for each job, which multiplies

# Data Management Tips

---

- Know your job needs
- Know your *batch* needs
- **Reduce per-job data needs**
- Test, then scale up gradually

# Per-job transfer limits



# Reducing data needs

---

- *An HTC best practice!*
- split large input for better throughput *and* less per-job data
- eliminate unnecessary data
- file compression and consolidation

# Split Up Your Input!

- Think throughput!
  - Can the data be processed as smaller pieces?
  - (Exercises 2.3 and 3.1)



- Rather than “shared” input ...
  - If each job only needs *some* of the data from a big file
  - (e.g. parameter sweeps)



# Eliminate Unnecessary Data!

---

- Test for files on your local computer that aren't actually necessary.
  - maybe: strip down complex software.
- Remove excess data from input.
- Transfer input carefully. Remember:  
`transfer_input_files = file1,dir,dir_contents/`

# Compress and Combine!

1. Input: prior to job submission, decompress in-job
2. Output: as part of the job
3. Prior to transfer between your computer and the submit server.

```
tar czf newfile.tar.gz <contents>
tar xzf newfile.tar.gz
```

- Also use to combine many files, especially for “shared” data and software


```
transfer_input_files = file1,many_files.tar.gz,dir.tar.gz
```

# Data Management Tips

---

- Know your job needs
- Know your *batch* needs
- Reduce per-job data needs
- **Test, then scale up gradually**

# Per-job transfer limits



**Number of jobs matters!**

# Test, and Scale Up *Gradually*

---

- Capabilities of submit servers and local networks may vary!
- Just because it isn't a problem for **10 jobs** doesn't mean you won't kill the submit server or network performance when you submit **100 or 1000 jobs**.

# Overview – Data Handling

---

- Review of HTCondor Data Handling
- What is ‘Large’ Data?
- Data Management Tips
- **Next talks: Dealing with Large Data**

# Exercises

---

- 2.1 Understanding a job's data needs
  - revisit the “blast” example from yesterday
- 2.2 Using data compression to reduce transfer load
  - compress/combine the database files
- 2.3 Splitting input (prep for large run in 3.1)
  - split large blast input

# Questions?

---

- Feel free to contact me:
  - [lmichael@wisc.edu](mailto:lmichael@wisc.edu)
- Next: Exercises 2.1-2.3
- Later: Handling large input data