

# Workflows: from Development to Automated Production

Friday afternoon, 1:15pm

Christina Koch [ckoch5@wisc.edu](mailto:ckoch5@wisc.edu)  
Research Computing Facilitators  
University of Wisconsin - Madison

# OR, GETTING THE MOST OUT OF WORKFLOWS, PART 2

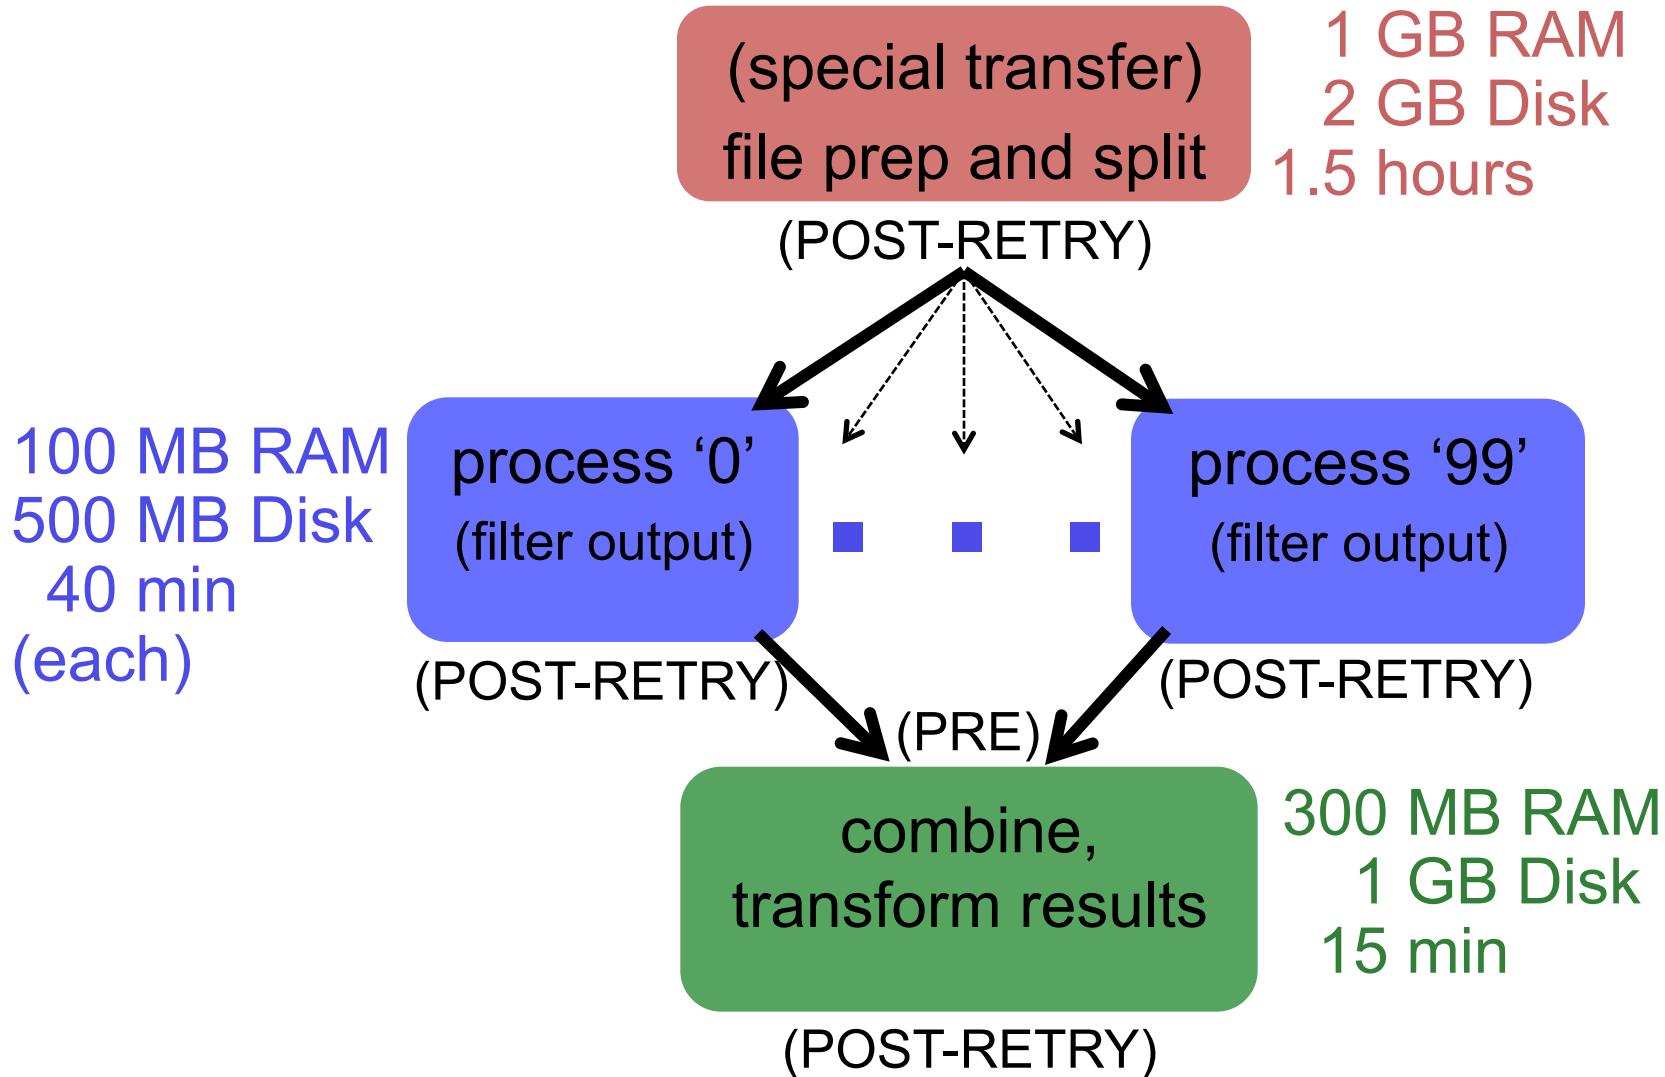
# Building a Good Workflow

---

1. Draw out the *general* workflow
2. Define details (test ‘pieces’ with HTCondor jobs)
  - divide or consolidate ‘pieces’
  - determine resource requirements
  - identify steps to be automated or checked
3. **Build it modularly; test and optimize**
4. Scale-up gradually
5. Make it work consistently
6. What more can you automate or error-check?

(And remember to document!)

# To Get Here ...



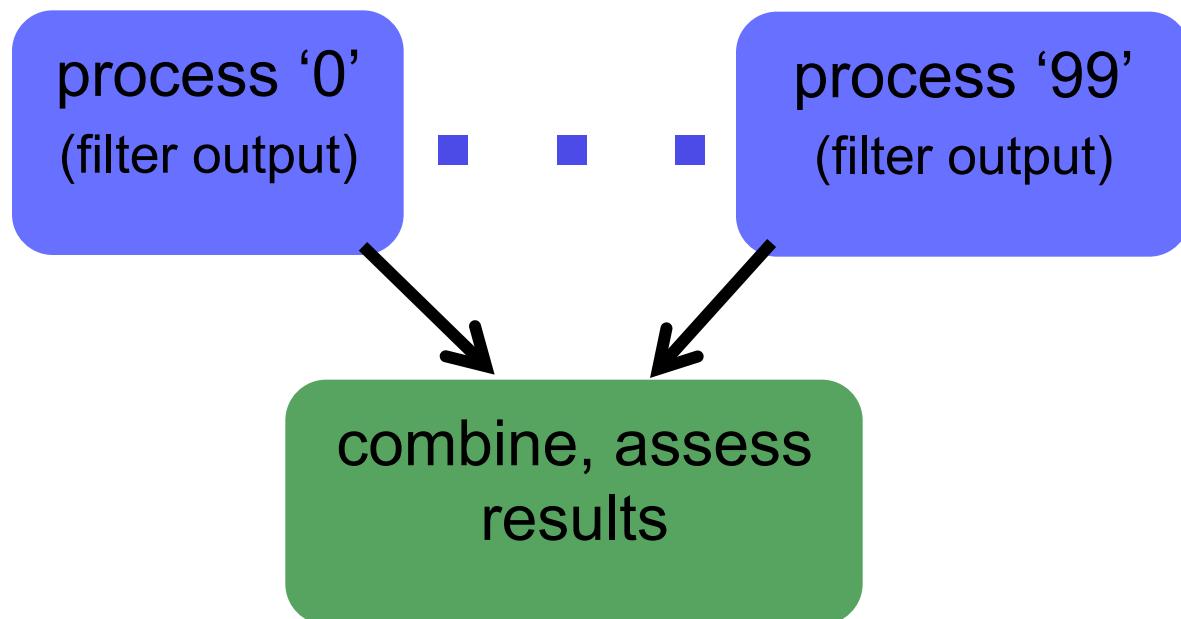
# start with the HTC “step” in the DAG...

process '0'  
(filter output)

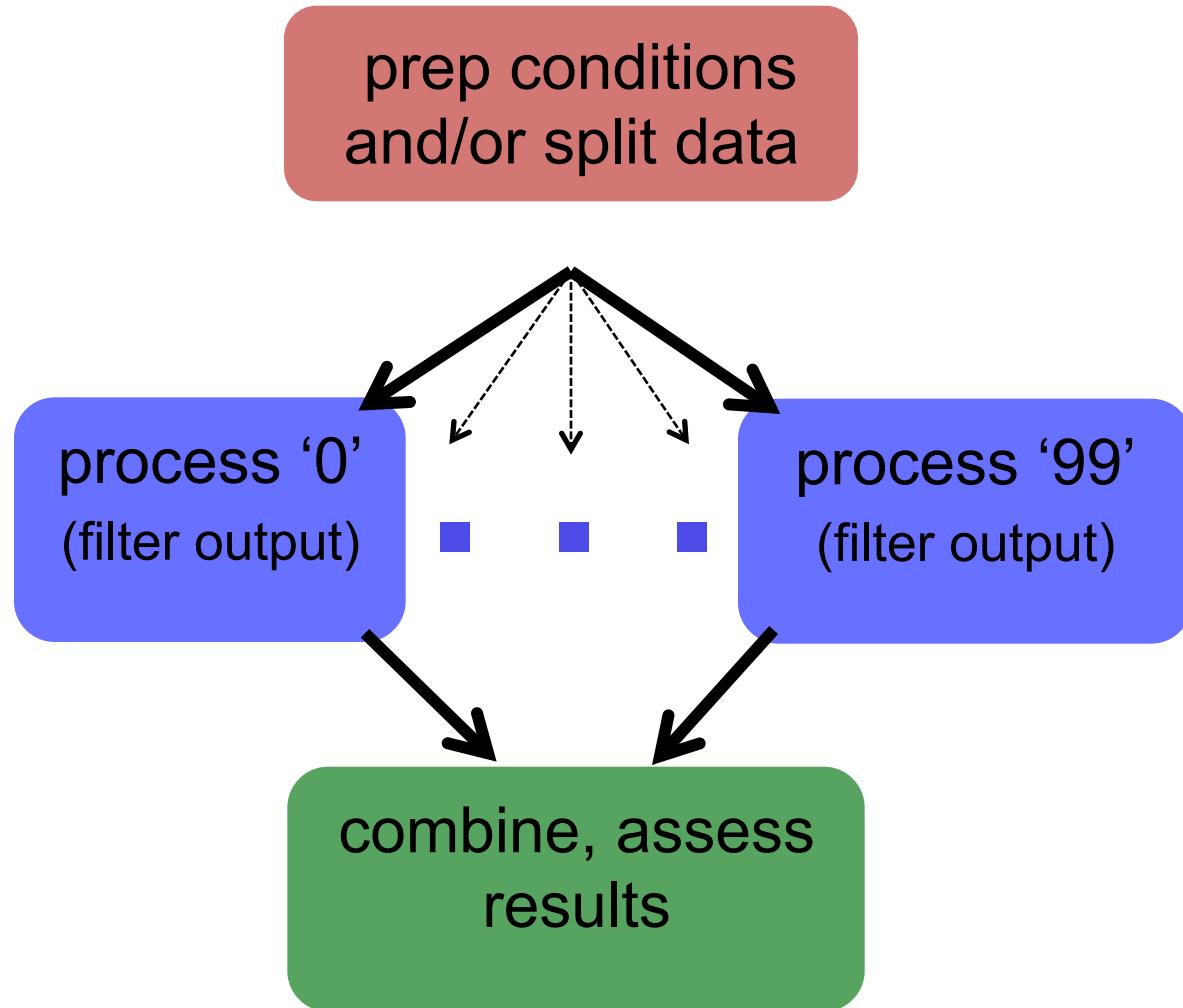


process '99'  
(filter output)

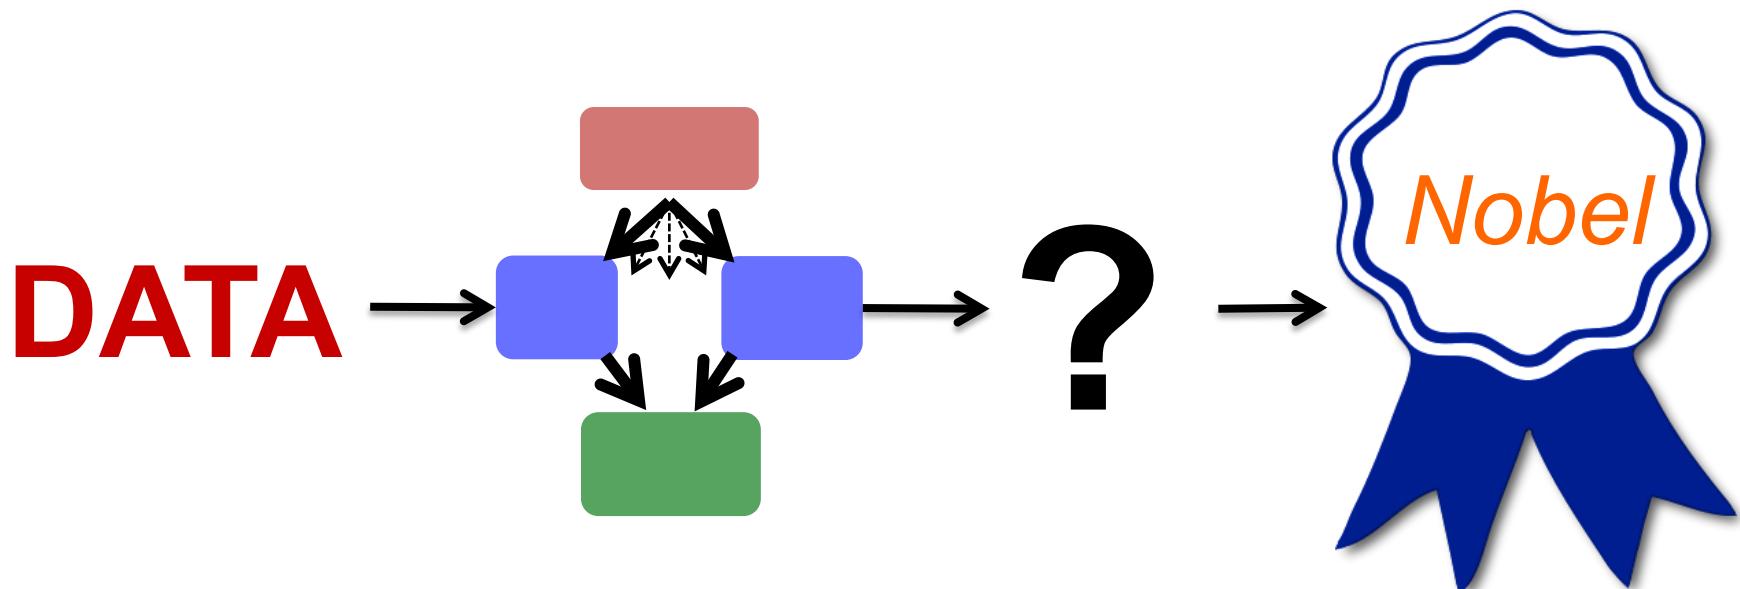
# ... then add in another step ...



# ... and another step ...



# and End Up with This



# Building a Good Workflow

---

1. Draw out the *general* workflow
2. Define details (test ‘pieces’ with HTCondor jobs)
  - divide or consolidate ‘pieces’
  - determine resource requirements
  - identify steps to be automated or checked
3. Build it modularly; test and optimize
4. **Scale-up gradually**
5. Make it work consistently
6. What more can you automate or error-check?

(And remember to document!)

# Scaling and Optimization

- Your ‘small’ DAG runs (once)! Now what?
  - Need to make it run *full-scale*
  - Need to make it run *everywhere, everytime*
  - Need to make it run *unattended*
  - Need to make it run *when someone else tries*



to the  
moon!



# Scaling Up: OSG Rules of Thumb

---

- CPU (single-threaded)
  - Best jobs run between **10 min** and **10 hrs**  
(Upper limit somewhat soft)
- Data (disk and network)
  - Keep scratch working space < 20 GB
  - Intermediate needs (/tmp?)
  - Use alternative data transfer appropriately
- Memory
  - Closer to 1 GB than 8 GB

# Testing, Testing, 1-2-3 ...

---

- ALWAYS test a subset after making changes
  - How big of a change needs retesting?
- Scale up gradually
- Avoid making problems for others (and for yourself)

# Scaling Up - Things to Think About

---

- More jobs:
  - 100-MB per input files may be fine for 10 or 100 jobs, but not for 1000 jobs. Why?
  - most submit queues will falter beyond ~10,000 total jobs
- Larger files:
  - more disk space, perhaps more memory
  - potentially more transfer and compute time

**Be kind to your submit and execute nodes  
and to fellow users!**

# Solutions for More Jobs

---

- Use a DAG to throttle the number of idle or queued jobs (“max-idle” and/or “DAGMAN CONFIG”)
- Add more resiliency measures
  - “RETRY” (works per-submit file)
  - “SCRIPT POST” (use \$RETURN, check output)
- Use SPLICE, VAR, and DIR for modularity/organization

# Solutions for Larger Files

---

- File manipulations
  - split input files to **send minimal data** with each job
  - **filter** input *and* output files to transfer only essential data
  - use compression/decompression
- Follow file delivery methods from yesterday for files that are still “large”

# Self-Checkpointing (solution for long jobs and “shish-kebabs”)

## 1. Changes to your code

- Periodically save information about progress to a new file (every hour?)
- At the beginning of script:
  - If progress file exists, read it and start from where the program (or script) left off
  - Otherwise, start from the beginning

## 2. Change to submit file:

```
when_to_transfer_output = ON_EXIT_OR_EVICT
```

# Building a Good Workflow

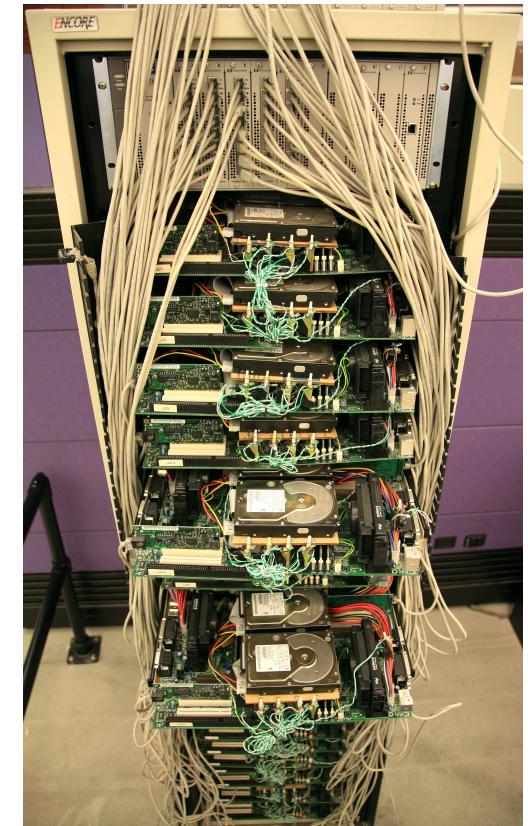
---

1. Draw out the *general* workflow
2. Define details (test ‘pieces’ with HTCondor jobs)
  - divide or consolidate ‘pieces’
  - determine resource requirements
  - identify steps to be automated or checked
3. Build it modularly; test and optimize
4. Scale-up gradually
5. **Make it work consistently**
6. What more can you automate or error-check?

(And remember to document!)

# Make It Run Everywhere

- What does an OSG machine have?
  - assume the worst
- Bring as much as possible with you:
  - won't that slow me down?
- Bring:
  - executable
  - likely, more of the “environment”



# The expanding onion

---

- Laptop (1 machine)
  - You control everything!
- Local cluster (1000 cores)
  - You can ask an admin nicely
- Campus (5000 cores)
  - It better be important/generalizable
- OSG (50,000 cores)
  - Good luck finding the pool admins\_

# Make It Work Everytime

- What could possibly go wrong?
  - Eviction
  - Non-existent dependencies
  - File corruption
  - Performance surprises
    - Network
    - Disk
    - ...
  - *Maybe even a bug in your code*



# Performance Surprises

---

One bad node can ruin your whole day

- “**Black Hole**” machines
  - Depending on the error, email OSG!
- ***REALLY* slow machines**
  - use `periodic_hold` / `periodic_release`

# Error Checks Are Essential

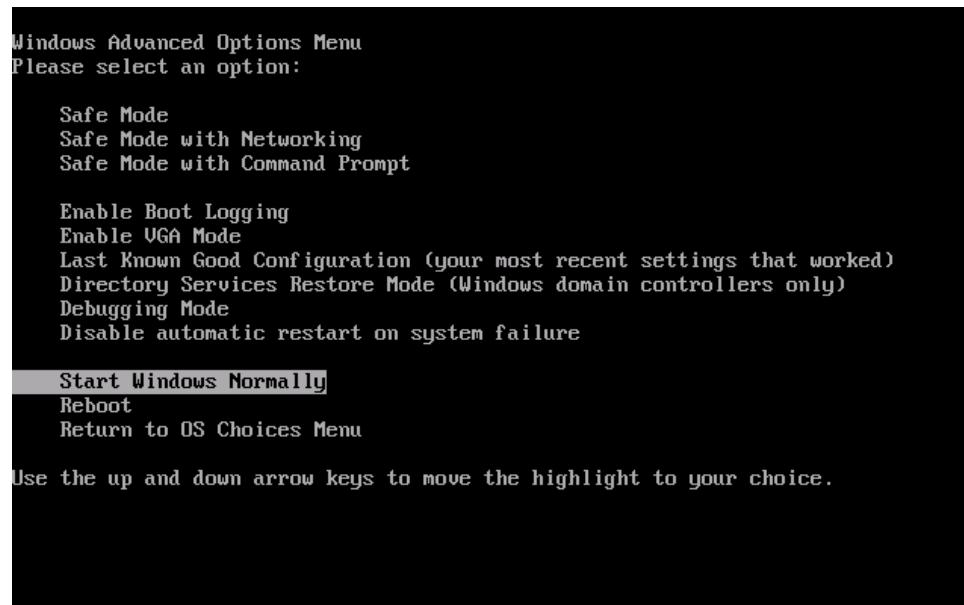
---

**If you don't check, it will happen...**

- Check expected file existence, and repeat with a finite loop or number of retries
  - better yet, check *rough* file size too
- Advanced:
  - RETRY for *specific* error codes from wrapper
  - “periodic\_release” for specific hold reasons

# What to do if a check fails

- Understand something about failure
- Use DAG “RETRY”, when useful
- Let the rescue dag continue...



# Building a Good Workflow

---

1. Draw out the *general* workflow
2. Define details (test ‘pieces’ with HTCondor jobs)
  - divide or consolidate ‘pieces’
  - determine resource requirements
  - identify steps to be automated or checked
3. Build it modularly; test and optimize
4. Scale-up gradually
5. Make it work consistently
6. **What more can you automate or error-check?**

(And remember to document!)

# Automate All The Things

---

- Well, not really, but kind of ...
- Really: What is the minimal number of manual steps necessary?  
even 1 might be too many; zero is perfect!
- Consider what you get out of automation
  - time savings (including less ‘babysitting’ time)
  - reliability and reproducibility

# Automation Trade-offs

HOW LONG CAN YOU WORK ON MAKING A ROUTINE TASK MORE  
EFFICIENT BEFORE YOU'RE SPENDING MORE TIME THAN YOU SAVE?  
(ACROSS FIVE YEARS)

		HOW OFTEN YOU DO THE TASK					
		50/DAY	5/DAY	DAILY	WEEKLY	MONTHLY	YEARLY
HOW MUCH TIME YOU SHAVE OFF	1 SECOND	1 DAY	2 HOURS	30 MINUTES	4 MINUTES	1 MINUTE	5 SECONDS
	5 SECONDS	5 DAYS	12 HOURS	2 HOURS	21 MINUTES	5 MINUTES	25 SECONDS
	30 SECONDS	4 WEEKS	3 DAYS	12 HOURS	2 HOURS	30 MINUTES	2 MINUTES
	1 MINUTE	8 WEEKS	6 DAYS	1 DAY	4 HOURS	1 HOUR	5 MINUTES
	5 MINUTES	9 MONTHS	4 WEEKS	6 DAYS	21 HOURS	5 HOURS	25 MINUTES
	30 MINUTES	6 MONTHS	5 WEEKS	5 DAYS	1 DAY	2 HOURS	
	1 HOUR	10 MONTHS	2 MONTHS	10 DAYS	2 DAYS	5 HOURS	
	6 HOURS			2 MONTHS	2 WEEKS	1 DAY	
	1 DAY				8 WEEKS	5 DAYS	

# Make It Work Unattended

- Remember the ultimate goal:  
**Automation! Time savings!**
- Things to automate:
  - Data collection?
  - Data preparation and staging
  - Submission (condor cron)
  - Analysis and verification
  - LaTeX and paper submission ☺



# Make It Run(-able) for Someone Else

---

- If others can't reproduce your work, it's *NOT* real science!
  - Work hard to make this happen.
  - It's *their* throughput, too.

Only ~10% of published cancer research  
is reproducible!

# Building a Good Workflow

---

1. Draw out the *general* workflow
2. Define details (test ‘pieces’ with HTCondor jobs)
  - divide or consolidate ‘pieces’
  - determine resource requirements
  - identify steps to be automated or checked
3. Build it modularly; test and optimize
4. Scale-up gradually
5. Make it work consistently
6. What more can you automate or error-check?

**(And remember to document!)**

# Documentation at Multiple Levels

---

- In job files: comment lines
  - submit files, wrapper scripts, executables
- In README files
  - describe file purposes
  - define overall workflow, justifications
- In a document!
  - draw the workflow, explain the big picture



Open Science Grid

---

# PARTING THOUGHTS

# Make It Run Faster? Maybe.

---

Throughput, throughput, throughput

- Resource reductions (match more slots!)
- Wall-time reductions
  - if significant *per workflow*
  - Why not *per job?*

Think in orders of magnitude:

- Say you have 1000 hour-long jobs that are matched at a rate of 100 per hour ...

*Waste the computer's time, not yours.*

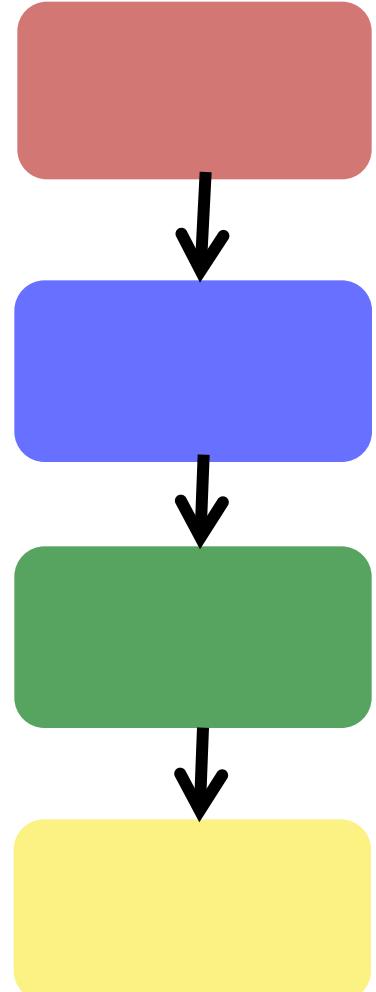
# Maybe Not Worth The Time

---

- Rewriting your code in a different, faster language
- Targeting “faster” machines
- Targeting machines that run jobs longer
- Others?

# Golden Rules for DAGs

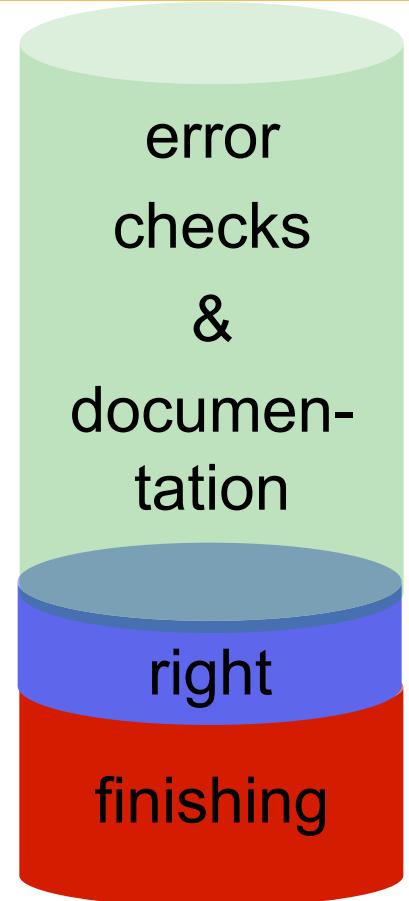
- Beware of lo-on<sup>g</sup> shish kebabs!
  - (self-checkpointing)
- Use PRE and POST script generously
- RETRY is your friend
- DIR and VAR for organization
- DAGs of DAGs are good
  - SPLICE
  - SUB\_DAG\_EXTERNAL



# If HTC workflows were a test...

- 20 points for finishing at all
- 10 points for the right answer
- 1 point for every error check
- 1 point per documentation line

*Out of 100 points? 200 points?*



# Questions?

---

- Now: Exercises 2.1 (2.2 Bonus)
- Next:
  - HTC Showcase!