

Troubleshooting Your Jobs

Brian Lin

OSG Software Team

University of Wisconsin - Madison

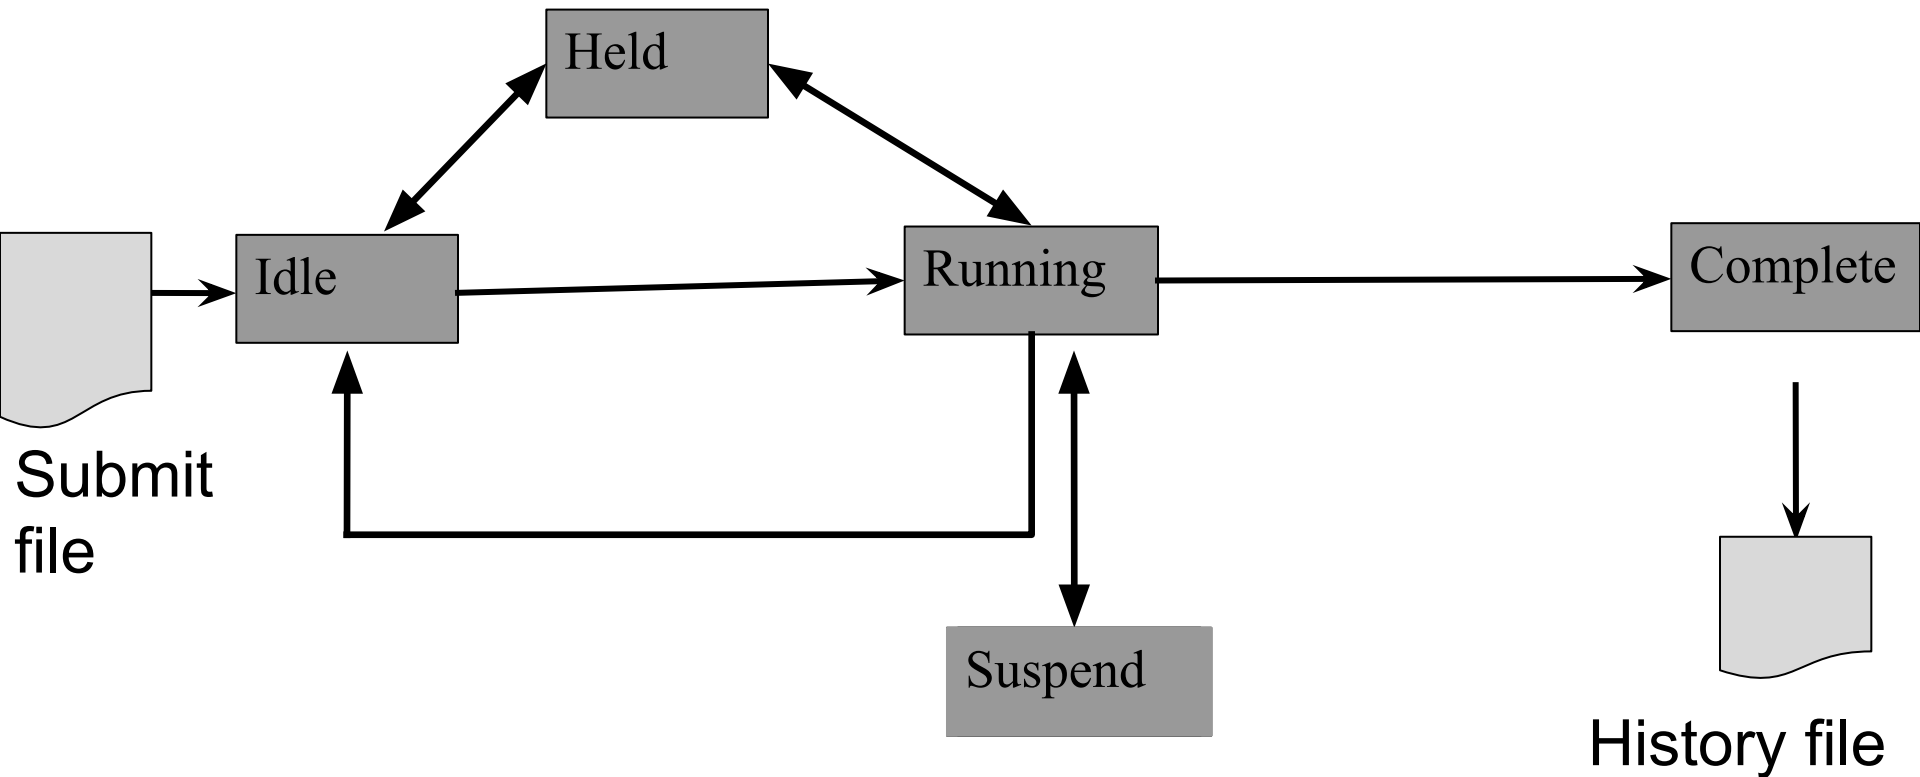
The grid is your oyster!



**The grid is
your oyster!**
...if you can get it to work.

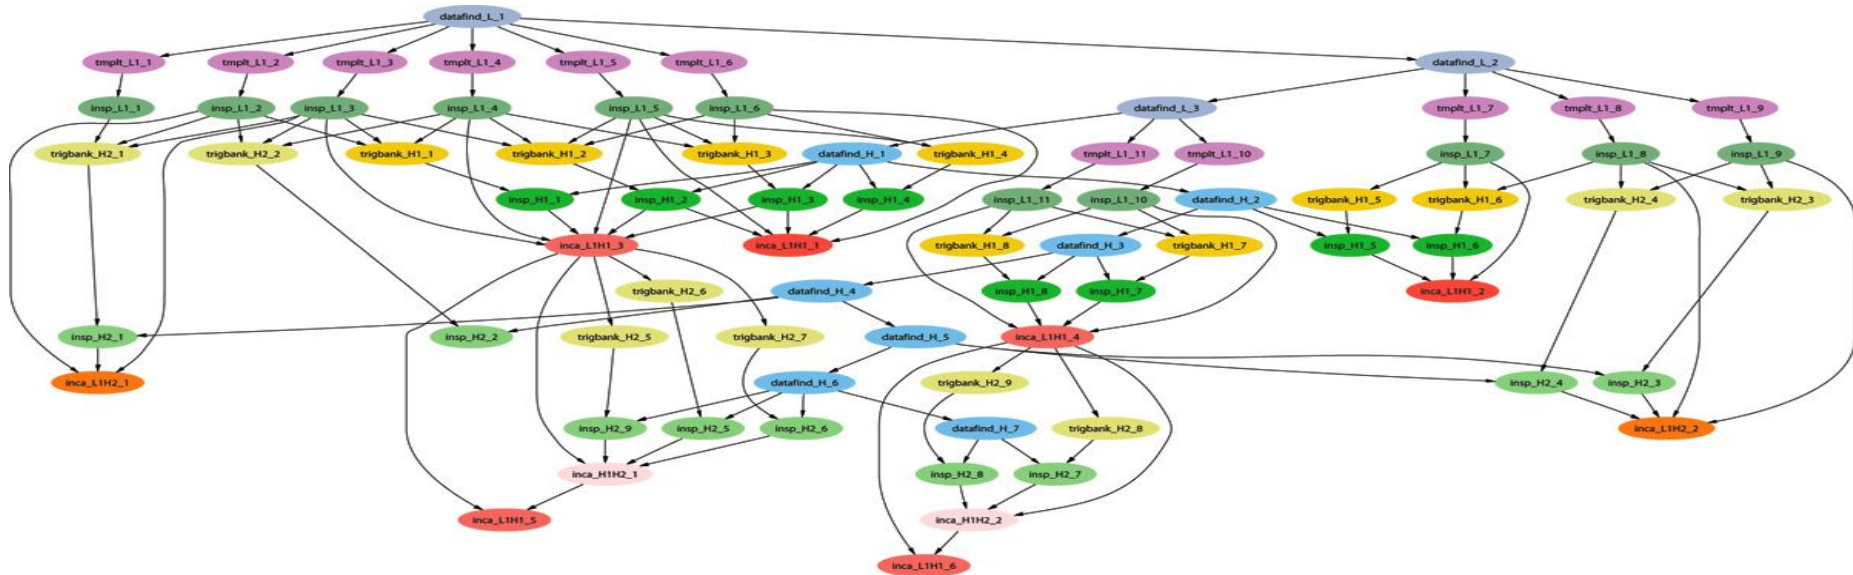


General Troubleshooting



General Troubleshooting

LIGO Example of a LIGO Inspiral DAG



Keys to Troubleshooting

- Identify the problem and break it into sub-problems
- Know where to find more information:
 - condor_q and its options
 - Output, error, log, and rescue files
 - Google
 - Local support
 - htcondor-users@cs.wisc.edu

Common Issues

Why can't I submit my job?

```
$ condor_submit sleep.sh
```

Submitting job(s)

ERROR: on Line 2 of submit file:

ERROR: Failed to parse command file (line 2).

Why can't I submit my job?

```
$ condor_submit sleep.sh
```

Submitting job(s)

ERROR: on Line 2 of submit file:

ERROR: Failed to parse command file (line 2).

Huh? You've tried to submit something that wasn't your submit file.

Solution: Submit your .sub file!

Why can't I submit my job?

\$ condor_submit sleep.sub

Submitting job(s)

- ERROR: Can't open "/cloud/login/blin/school/inptu_data" with flags 00 (No such file or directory)
- Submitting job(s)No 'executable' parameter was provided
- ERROR: I don't know about the 'vanila' universe.
- ERROR: Executable file /bin/slep does not exist

Why can't I submit my job?

\$ condor_submit sleep.sub

Submitting job(s)

- ERROR: Can't open "/cloud/login/blin/school/inptu_data" with flags 00 (No such file or directory)
- Submitting job(s)No 'executable' parameter was provided
- ERROR: I don't know about the 'vanila' universe.
- ERROR: Executable file /bin/slep does not exist

Huh? There are typos in your submit file.

Solution: Fix your typos! Condor can only catch a select few of them.

Why can't I submit my DAG?

Errors appears in *dagman.out files instead of STDOUT or STDERR:

07/25/16 17:14:42 From submit: Submitting job(s)

07/25/16 17:14:42 From submit: **ERROR: Invalid log file:**
"/home/blin/sleep/sleep.log" (No such file or directory)

07/25/16 17:14:42 failed while reading from pipe.

07/25/16 17:14:42 Read so far: Submitting job(s)**ERROR: Invalid log file:**
"/home/blin/sleep/sleep.log" (No such file or directory)

07/25/16 17:14:42 **ERROR: submit attempt failed**

What are my jobs up to?

\$ condor_q -help status

JobStatus codes:

1 I IDLE

2 R RUNNING

3 X REMOVED

4 C COMPLETED

5 H HELD

6 > TRANSFERRING_OUTPUT

7 S SUSPENDED

Why are my jobs idle?

```
$ condor_q -better 29486
```

```
[...]
```

The Requirements expression for job 29486.000 reduces to these conditions:

Slots

Step	Matched	Condition
------	---------	-----------

-----	-----	-----
-------	-------	-------

[1]	8348	Target.OpSysMajorVer == 6
-----	------	---------------------------

[7]	10841	TARGET.Disk >= RequestDisk
-----	-------	----------------------------

[9]	1	TARGET.Memory >= RequestMemory
-----	---	--------------------------------

[11]	10852	TARGET.HasFileTransfer
------	-------	------------------------

Why are my jobs still running?

```
$ condor_q -nobatch
```

```
-- Schedd: learn.chtc.wisc.edu : <128.104.100.43:9618?...
```

ID	OWNER	SUBMITTED	RUN_TIME	ST	PRI	SIZE	CMD
14665.0	blin	7/25 18:19	0+23:00:03	R	0	0.3	sleep.sh

Solution: Use ``condor_ssh_to_job <job ID>`` to open an SSH session to the worker node running your job.

Why are my jobs held?

```
$ condor_q -held
```

```
-- Schedd: learn.chtc.wisc.edu : <128.104.100.43:9618?... @ 07/18/17 15:13:42
```

ID	OWNER	HELD_SINCE	HOLD_REASON
19.0	blin	7/14 15:07	Error from fermicloud113.fnal.gov: Failed to execute '/cloud/login/blin/school/sleep.sh': (errno=13: 'Permission denied')
14662.0	blin	7/25 18:05	Error from slot1_12@e163.chtc.wisc.edu: Failed to execute '/var/lib/condor/execute/slot1/dir_3090825/condor_exec.exe': (errno=8: 'Exec format error')

Why are my jobs held?

```
$ condor_q -held
```

```
-- Schedd: learn.chtc.wisc.edu : <128.104.100.43:9618?... @ 07/18/17 15:13:42
```

ID	OWNER	HELD_SINCE	HOLD_REASON
19.0	blin	7/14 15:07	Error from fermicloud113.fnal.gov: Failed to execute '/cloud/login/blin/school/sleep.sh': (errno=13: 'Permission denied')

14662.0	blin	7/25 18:05	Error from slot1_12@e163.chtc.wisc.edu: Failed to execute '/var/lib/condor/execute/slot1/dir_3090825/condor_exec.exe': (errno=8: 'Exec format error')
---------	------	------------	---

Huh? Condor couldn't run your executable.

Solution: Set the executable bit on your executable (`chmod +x <filename>`) and/or add the missing shebang line at the top of the executable, e.g. `#!/bin/bash`.

Why are my jobs held?

\$ condor_q -held

-- Schedd: learn.chtc.wisc.edu : <128.104.100.43:9618?... @ 07/18/17 15:13:42

ID	OWNER	HELD_SINCE	HOLD_REASON
29494.0	blin	7/16 17:31	Failed to initialize user log to /home/blin/foo/bar/test-000.log or /home/blin/foo/./test.dag.nodes.log

Why are my jobs held?

```
$ condor_q -held
```

```
-- Schedd: learn.chtc.wisc.edu : <128.104.100.43:9618?... @ 07/18/17 15:13:42
```

ID	OWNER	HELD_SINCE	HOLD_REASON
29494.0	blin	7/16 17:31	Failed to initialize user log to /home/blin/foo/bar/test-000.log or /home/blin/foo/./test.dag.nodes.log

Huh? Condor couldn't read job or DAG log files

Solution: Ensure existence/proper write permissions for each folder in the specified path or choose a new location for your files!

Why are my jobs held?

```
$ condor_q -held
```

```
-- Schedd: learn.chtc.wisc.edu : <128.104.100.43:9618?... @ 07/18/17 15:13:42
```

ID	OWNER	HELD_SINCE	HOLD_REASON
19.0	blin	7/14 15:07	Error from fermicloud113.fnal.gov: Failed to execute '/cloud/login/blin/school/sleep.sh': invalid interpreter (/bin/bash) specified on first line of script (errno=2: 'No such file or directory')

Why are my jobs held?

```
$ condor_q -held
```

```
-- Schedd: learn.chtc.wisc.edu : <128.104.100.43:9618?... @ 07/18/17 15:13:42
```

ID	OWNER	HELD_SINCE	HOLD_REASON
19.0	blin	7/14 15:07	Error from fermicloud113.fnal.gov: Failed to execute '/cloud/login/blin/school/sleep.sh': invalid interpreter (/bin/bash) specified on first line of script (errno=2: 'No such file or directory')

Huh? There are carriage returns (^M) in your executable.

Solution: Use `dos2unix` or `vi -b <filename>` to see and delete the carriage returns (use 'x' to delete).

Why are my jobs held?

\$ condor_q -held

-- Schedd: learn.chtc.wisc.edu : <128.104.100.43:9618?... @ 07/18/17 15:13:42

ID	OWNER	HELD_SINCE	HOLD_REASON
19.0	blin	7/14 15:07	Error from slot1_1@e026.chtc.wisc.edu: Job has gone over memory limit of 1 megabytes. Peak usage: 1 megabytes.

Why are my jobs held?

```
$ condor_q -held
```

```
-- Schedd: learn.chtc.wisc.edu : <128.104.100.43:9618?... @ 07/18/17 15:13:42
```

ID	OWNER	HELD_SINCE	HOLD_REASON
19.0	blin	7/14 15:07	Error from slot1_1@e026.chtc.wisc.edu: Job has gone over memory limit of 1 megabytes. Peak usage: 1 megabytes.

Huh? You've used too many resources.

Solution: Request more!

Why are my jobs held?

Edit your jobs on the fly with ``condor_qedit`` and ``condor_release``:

1. `$ condor_qedit <job ID> <resource> <value>`
 - a. `$ condor_qedit <job ID> RequestMemory <mem in MB>`
 - b. `$ condor_qedit -const 'JobStatus =?= 5' RequestDisk <disk in KiB>`
 - c. `$ condor_qedit -const 'Owner =?= "blin"' RequestCpus <CPUs>`
2. `$ condor_release <job ID>`

Or remove your jobs, fix the submit file, and resubmit:

1. `$ condor_rm <job ID>`
2. Add `request_disk`, `request_mem`, or `request_cpus` to your submit file
3. `$ condor_submit <submit file>`


Why are my jobs held?

\$ condor_q -held

-- Schedd: learn.chtc.wisc.edu : <128.104.100.43:9618?... @ 07/18/17 15:13:42

24.0 blin 7/14 16:12 Error from fermicloud113.fnal.gov: STARTER at 128.104.100.52 failed to send file(s) to <128.104.100.43:64130>: error reading from /var/lib/condor/execute/dir_24823/bar: (errno 2) No such file or directory; SHADOW failed to receive file(s) from <128.104.100.52:10507>

Why are my jobs held?



Why are my jobs held?

```
$ condor_q -held
```

```
-- Schedd: learn.chtc.wisc.edu : <128.104.100.43:9618?... @ 07/18/17 15:13:42
```

```
24.0 blin 7/14 16:12 Error from fermicloud113.fnal.gov: STARTER at  
128.104.100.52 failed to send file(s) to <128.104.100.43:64130>: error reading from  
/var/lib/condor/execute/dir_24823/bar: (errno 2) No such file or directory; SHADOW  
failed to receive file(s) from <128.104.100.52:10507>
```

Huh? Your job did not create the files that you specified in *transfer_output_files*.

Solution: Check for typos in *transfer_output_files*, job runtime, add debugging information to your code.

My jobs completed but...

The output is wrong:

1. Check *.log files for return codes or unexpected behavior: short runtimes, using too much or too few resources
2. Check *.err and *.out for error messages.
3. Submit an interactive job: ``condor_submit -i <submit file>`` and run the executable manually.
 - a. If it succeeds, does your submit file have the correct args? If yes, try adding `'GET_ENV=True'` to your submit file.
 - b. If it fails, there is an issue with your code or your invocation!

Troubleshooting DAGs

1. Check `*.rescue*` files (which DAG nodes failed)
2. Check `*.dagman.out` (errors with job submission)
3. Check `*.nodes.log` (return codes, PRE/POST script failures).
4. If PRE/POST scripts failed, run them manually to see where they failed.

Troubleshooting Exercise



petshop bones

@direlog

Follow



'programming' is a delightful pastime in which you debug for hours only to discover the real problem is: you can't read, and you can't count

6:56 PM - 16 Jul 2017

877 Retweets 1,735 Likes



12



877



1.7K



Source: @direlog, <https://twitter.com/direlog/status/886721271102410752>

Thanks!

Questions?