CMS Connect

Planning

Goal

Explore creating a distributed job service for USCMS, similar to ATLAS Connect and OSG Connect

http://connect.usatlas.org/

http://osgconnect.net/

- “Virtual cluster” to augment local resources
- Ease burden of site admins (resource targets)

Service Elements

- Identity & group management service
- Auto account provisioning to login host
- Flocking service (BOSCO-based)
- Data service (for quasi-transient job data)
- Job monitoring service
- Accounting service
- Web page for navigation
- Service documentation page

Identity and Group Management

- Leverage the Nexus service provided by Globus
- Use campus identity
- Use CI-Logon, InCommon federation
- Users authorized in standard way (acceptable use agreement, personal verification) but no X509 or VOMS needed

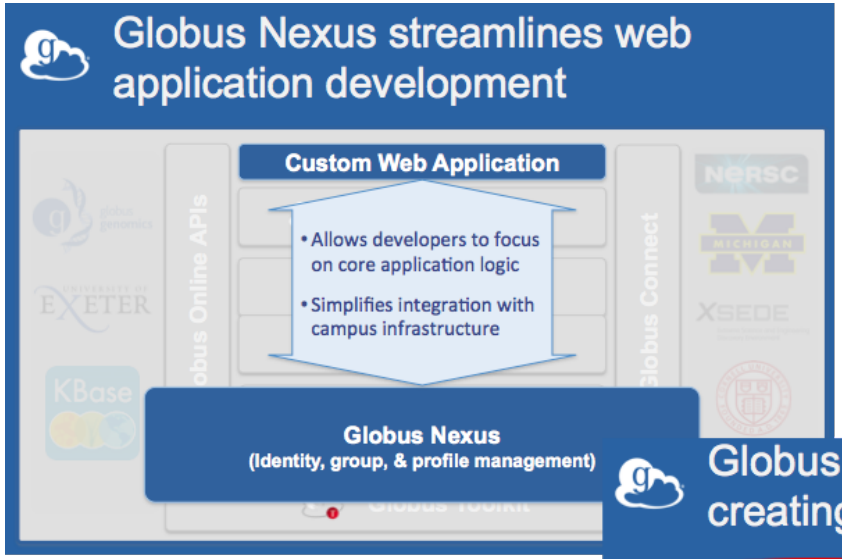

Globus Nexus makes it easy for individuals, teams, and institutions to create web applications for the science community

It provides a flexible, powerful Platform-as-a-Service to which developers can outsource their identity, group, and profile management needs.

CMS Connect would leverage all these components.

CI-Logon plus home university ID provider used to login.

CMS Operations staff control authorization.

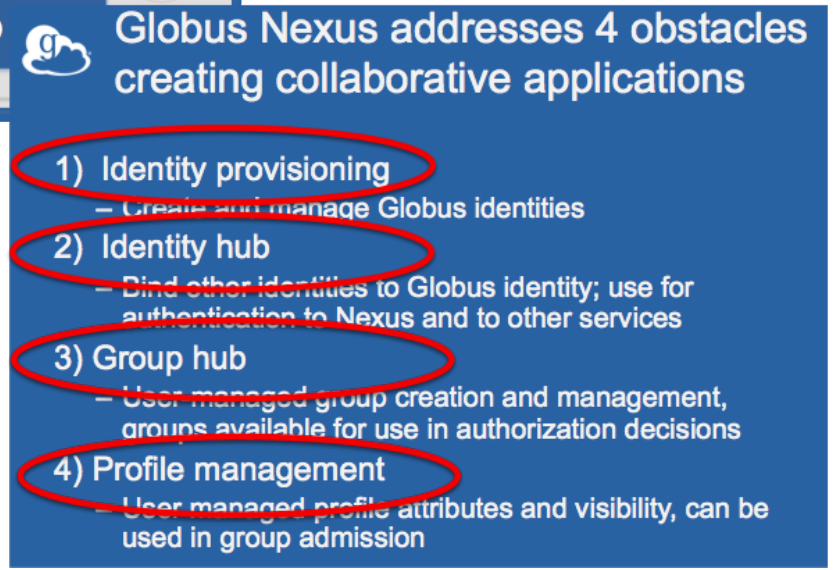

Auto provisioning of user accounts

- A cron runs on the login hosts and checks for newly authorized members in the CMS Connect group.
- Creates Unix account, and users can quickly login and start work (no admin needed) (~10 mins)
- Has `/cvmfs/cms.cern.ch/`, OASIS mounted plus other CMS compatibility libs as needed.

Flocking service

- `cms-connect` user account created on target clusters (runs Glideins)
- Only ssh-based protocol to reach clusters
- Resource targets can advertise attributes back to the schedd as usual
- Usual ClassAds and Condor submit scripts

Resource targets

- Any cluster running Condor, PBS, SLURM, or SGE
- User account created
- Best if local squid service is available
- Can also submit to any GlideinWMS VO front-end (also OSG VO front-end for opportunistic cycles)

CVMFS

There are a number of options is the target cluster does not have CVMFS. We are exploring these in ATLAS + CCTools group at Notre Dame

Storage service

- Provide a quasi-transient storage service for job input/output data
- Implemented with Ceph object store
- POSIX access provided to login host
- Globus Server for managed xfers
- Xrootd, and http service endpoints
- Would provision ~ 20 TB

Other services

- Web skin for navigation to services
- Cycle Server for job monitoring
- Gratia for accounting
- Confluence wiki for service documentation (CMS user documentation found elsewhere)

How to do it

- Meeting to discuss services; capabilities and limitations
- Identify one or more resource targets, setup technical phone call
- Meeting to discuss web branding, information, domain name, email support lists, etc.

Possible timeline

1. Meeting next week or soon thereafter
2. Two-three weeks for UC team to deploy services, including Globus-based backend
3. Week of testing with first couple of resource targets. New targets can be added usually quickly, with a brief technical phone call
4. First users in August