



OSG Area Coordinators

Network Monitoring Update: **August 16, 2017**

Shawn McKee

Status of Issues from Last Time

- Rebalancing of VM resources successful
 - Small tweaks to fix some remaining issues
- Successfully migrated mesh-configuration from MyOSG/OIM to the new standalone instance (<http://meshconfig.grid.iu.edu>)
 - URL redirection in place for toolkit instances using the old URL
 - Docs updated to point to new location
- The ETF monitoring instance has been updated and is working well

OSG Network Services Status

- Since the last meeting the OSG network services have been running well (mostly)
 - Some issues with RSV stopping mysteriously
 - Edgar couldn't find a root cause
 - Restart of rsv fixes the issue
- No significant other issues discovered and services are generally running well
 - Would like to maintain current status of 'non-production' until RSV issues are resolved and new data flow is in place and tested

perfSONAR Update Released

- Yesterday the perfSONAR developers release perfSONAR v4.0.1
 - This is a minor feature and bugfix release.
 - Sites with auto-updates enabled will automatically install the release sometime in the next few days
 - For a list of changes see <https://www.perfsonar.net/release-notes/version-4-0-1/>
- As of mid-day today about 15% of our instances had upgraded.
 - We expect ~80% to be upgraded in about 48 hours

Updated Documentation

- Our existing documentation at <https://twiki.opensciencegrid.org/bin/view/Documentation/NetworkingInOSG> has been revised
 - Streamlined
 - Updated for v4.0
- The perfSONAR developers provide excellent documentation on installation and configuration of perfSONAR
 - We pointed to that and provided details on the few specific things we need to worry about for OSG/WLCG installations
 - Firewalls configured allowing our access to the data
 - Configuration for auto-mesh and information about MCA
 - Information about OSG network services
 - Information about OSG and WLCG support
- Two reviewers of the docs suggested some updates...

Documentation Additions Planned

- Create a page discussing **http** and **https** requirements
 - The pscheduler now **requires** port 443 to be open between perfSONAR instances that test to one another
- Add some sections to cover important deployment options
 - Deployment models (VMS, containers, etc) and their pros/cons
 - Multiple NICs and their pros/cons
 - Requirements for bandwidth vs latency perfSONAR boxes
- Describe psetf (Experiments Testing Framework for perfSONAR; see https://psetf.grid.iu.edu/etf/check_mk/)
- Update the **FAQ** for new tests available in perfSONAR v4.x
- Augment troubleshooting information to enhance pS troubleshooting web page
- Create page on **installation/upgrade** to CentOS7 so users know they don't need to migrate their data
- Refactor docs to provided a better detailed a step-by-step guide

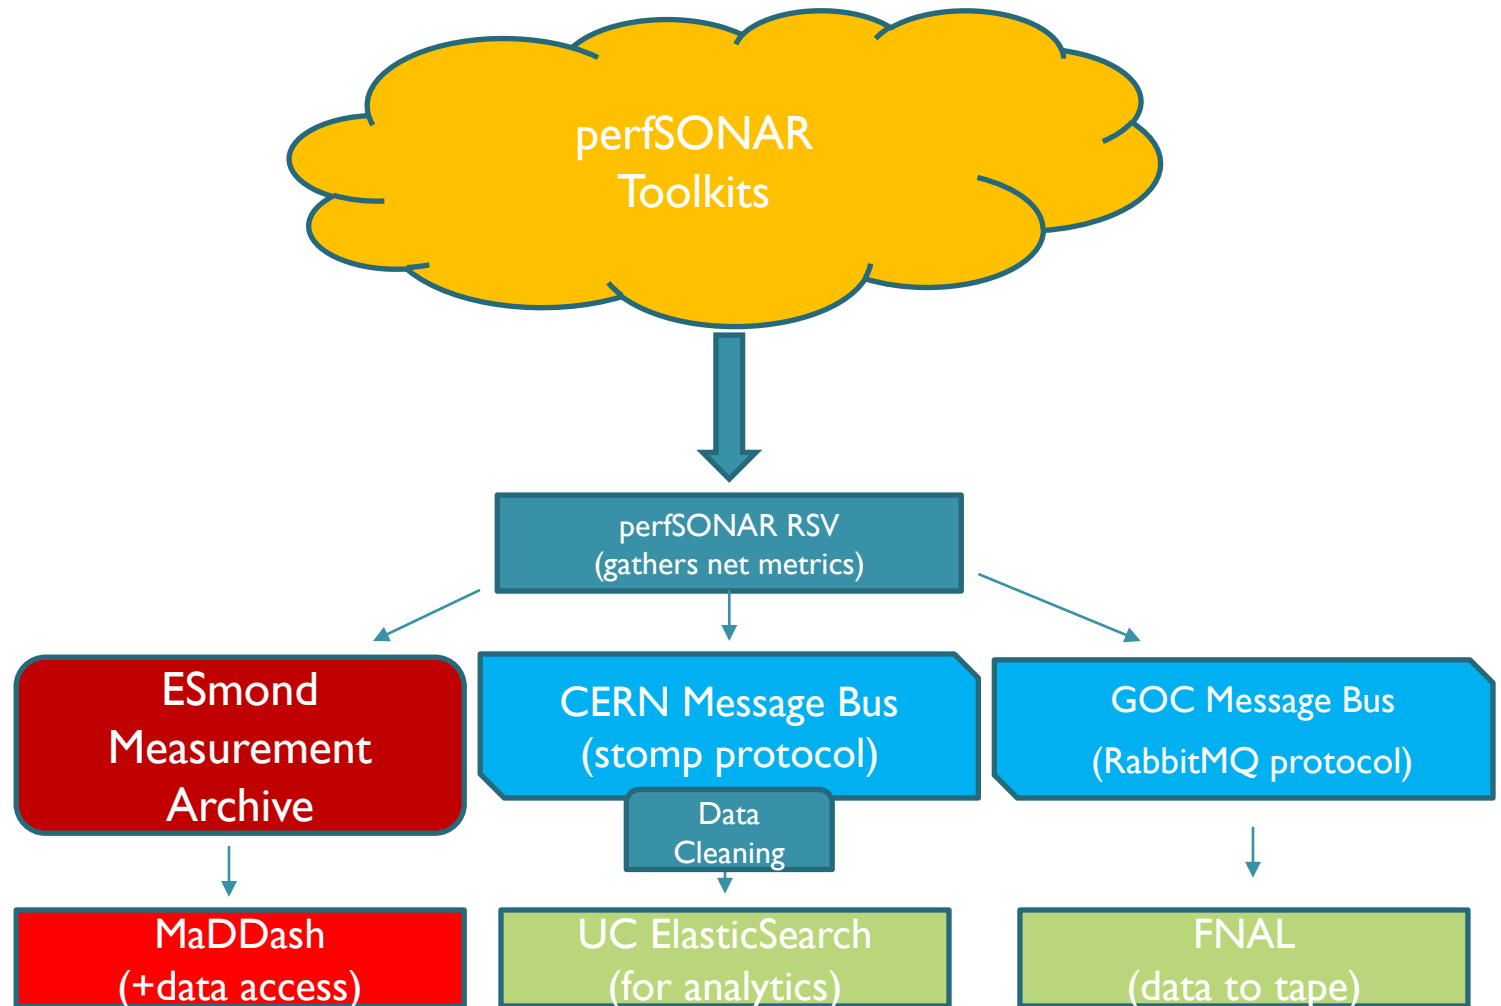
Recent GGUS Network Tickets

- Contacted by **CERN** about **Pakistan** transfer issues
 - Transfer rates are $\sim < 2$ Mbps
 - **perfSONAR** results showed problem was some kind of network shaper/limiter
- **Lancaster** – **SARA** (bad link on a LAG)
- **Dubna** – **Moscow** (Bad fiber identified)
- Transfer to **Oxford** (dCache issue)
- **Brazil** (CBPF) to elsewhere (MTU issue)

Action Items from the Retreat

- During the retreat we identified a few action items
 - We wanted to start gathering the HTCondor Schedd network information
 - We wanted to change the network data pipeline to ensure the critical services for OSG networking were within OSG control
 - We will decrease our reliance on Esmond, keeping the current instance for only a few weeks of data (supports MaDDash)

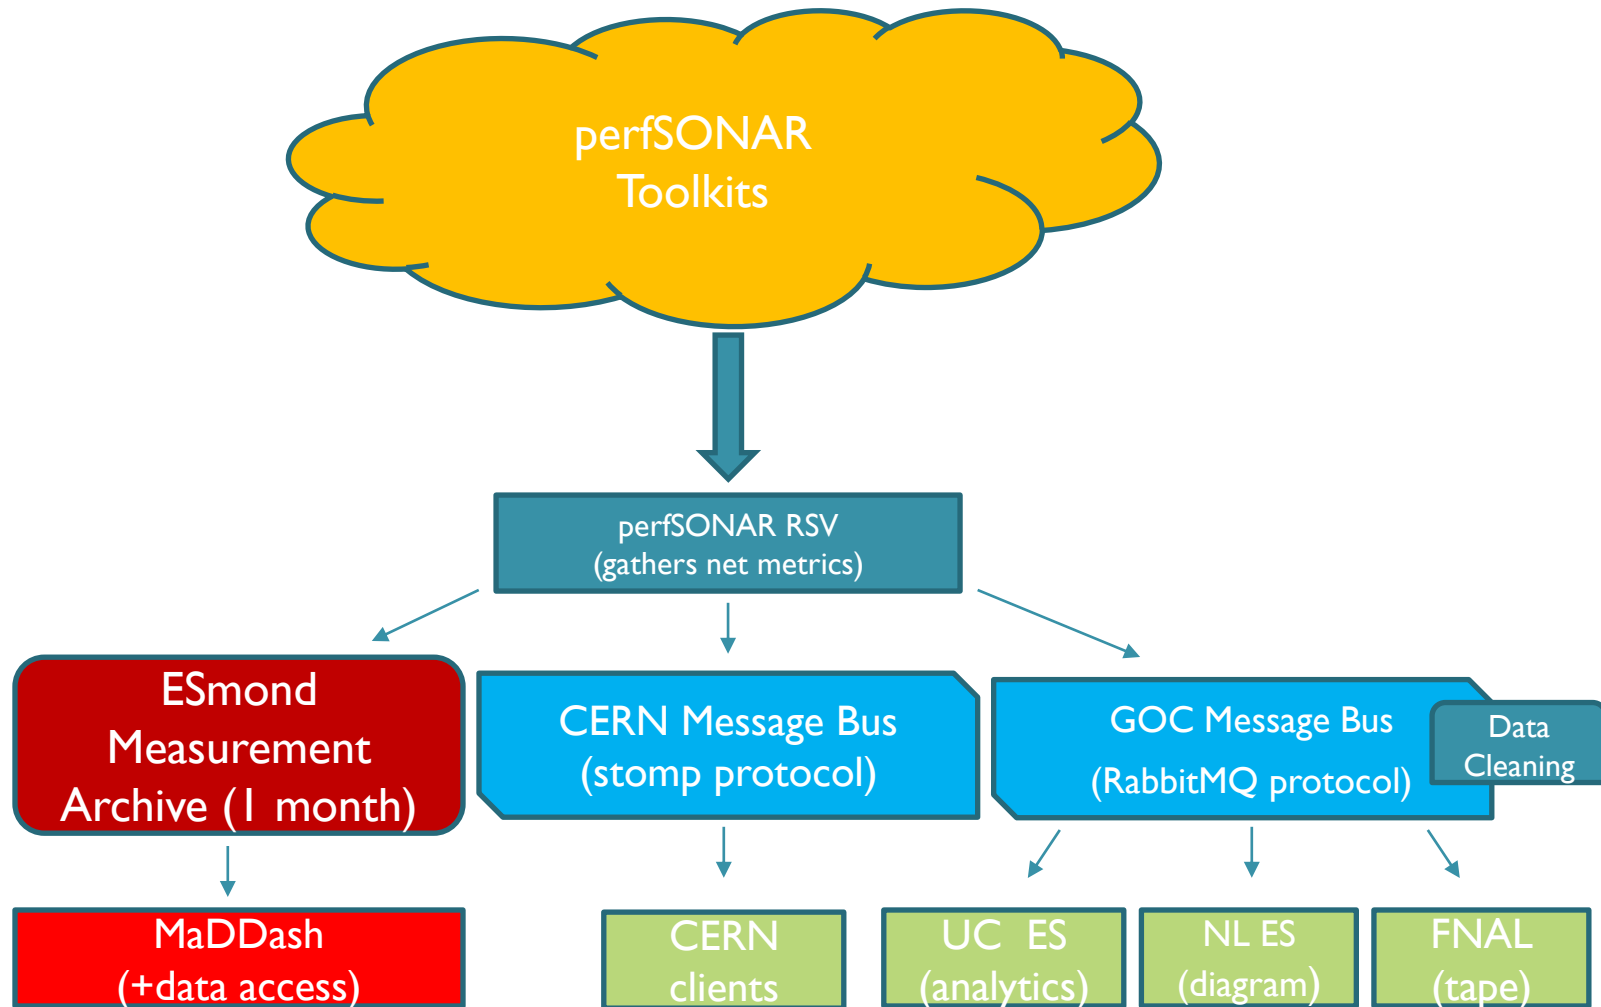
Current Net Metrics Flow



Planned Pipeline Changes

- The data pipeline will be modified to have use the GOC RabbitMQ message bus.
 - Raw JSON data will go to tape at FNAL for long term storage
 - Network metrics will be sent to ElasticSearch at Nebraska for use in non-interactive exploration and display of data, e.g., 1 year graphs of quantities of interested updated daily
 - Network metrics currently be read from CERN's bus and sent to UC will be reconfigured to use the OSG message bus as a source
 - Data cleaning and transformation is currently done by Ilija using a VM at CERN. This will be migrated to the US and will use the GOC bus
 - The **STOMP** protocol will be enabled on the GOC bus
 - We will continue to send data to CERN for their use
 - We will reconfigure ESmond on OSG to only host a few weeks of data to support our MaDDash instance (to quickly view net status)

Planned Net Metrics Flow



Results of Changes

- While it may some more complex, the pipeline changes result in simplifications in use and operations
 - All critical data stays within OSG controlled resources allowing easier debugging and management
 - One primary bus setup at GOC that we need to worry about
 - Data endpoints have very specific tasks allowing simpler configuration and better resource use
 - The Esmond instance can be very much scaled down. **No need to address data lifecycle management via Esmond!**
- Enables a future improvement: perfSONAR 4.1 should support central configuration to send data from our toolkit instances directly to a RabbitMQ bus like the one at GOC

Status of Retreat Changes

- Lincoln Bryant / UC tried to reconfigure HTCondor per instructions to cause it to emit the network data to a log file.
 - Didn't work: "I tried turning on the recommended knobs but never saw any network logging."
- Data pipeline reconfiguration has not happened yet...Ilija, Shawn and others too busy with various meetings since the retreat
- **Esmond use in OSG still working OK. Next steps will be to setup 8 week limit on the data retained.**

New 'https' issue found

- Since the release of perfSONAR 4.0 we have seen issues in measurement coverage (enabled by our new dashboards)
 - We have identified an issue: perfSONAR 4.0 now requires port 443 (https) be open between instances.
 - We previously told sites to open either 80 or 443
 - A test from site 'A' to site 'B' initiated by site 'B' requires that site 'A' upload its results to site 'B' via https (pscheduler) and some sites block 443
 - We will discuss this at the next perfSONAR developers meeting tomorrow

Publishing New Dashboards?

- We (Marian Babik and I) are almost ready to publish some new Kibana dashboards which provide useful task-specific information for various use-cases:
 - Site problems
 - Link problems
 - perfSONAR infrastructure problems
 - R&E network visibility
- **Question to OSG: are we ready to do this?**
 - Concern is that the UC analytics platform really needs to be “production” level once we point users at these new dashboards
 - Discussion? Additional concerns? Suggestions?

Involving an Undergraduate in OSG Networking

- I am putting together an UROP proposal (Undergraduate Research Opportunity Program) to have an suitable undergraduate work with me on OSG network related tasks
 - Would be for Fall 2017 (no cost to OSG)
 - If they work out I can offer to hire them hourly continuing forward
 - Proposal is due August 21st

Concerns

- Issue with https needs resolving: what should we tell our sites?
- Need to implement data pipeline changes but will require help from outside OSG networking
 - Main issues are location of VM for data cleaning/augmentation and setup of GOC message bus to support “transparent” change for VM to receive data
- Still on the list: recruit additional non-WLCG sites using our developing capabilities as a “carrot” (new dashboards and docs will help)

Questions or Comments?

Thanks!

URLs for Reference

- OSG Network Datastore Documents
 - Operations https://docs.google.com/document/d/11144BSO-88M0cLMMjKcKMIE-Q5s2IX-w3IYl-0Pn_08/edit#
 - SLA <https://twiki.grid.iu.edu/bin/view/Operations/PSServiceLevelAgreement>
 - Data lifecycle https://docs.google.com/document/d/1mJlkf43nZf6gvKoNtiTOc0g0MYDv_wSfSm7YdiMs3Lo/edit#
- Current OSG network documentation <https://www.opensciencegrid.org/bin/view/Documentation/NetworkingInOSG>
- OSG networking year-5 goals and milestones: <https://docs.google.com/document/d/1FzmXZinO4Pb8NAfd5SWUzaAFYOL23dt66hQsDmaP-VWl/edit>
- perfSONAR adoption tracking: http://grid-monitoring.cern.ch/perfsonar_coverage.txt
- Deployment documentation for both OSG and WLCG hosted in OSG (migrated from CERN) <https://twiki.opensciencegrid.org/bin/view/Documentation/DeployperfSONAR>
- ATLAS Analytics:
 - Packet-loss: <http://tiny.cc/PktLossNoUnknown> (6 month view)
 - perfSONAR dashboard: <http://tiny.cc/pSDash>
 - perfSONAR link details: <http://tiny.cc/pSLink>
- Mesh-config in OSG <https://meshconfig.grid.iu.edu/>
- Pre-Production Meshconfig <https://meshconfig-itb.grid.iu.edu/>
- perfSONAR homepage: <http://www.perfsonar.net/>