

Technology Investigations Update

Brian Bockelman

Area Coordinator's Meeting - September 20

Highlights

- GRÅCC: Cleanups and more data sources.
- Starting to execute on plan for small, concrete improvements in monitoring.
- Ongoing Stash-Bonanza at Nebraska.

GRÅCC


- perfSonar data is now archived @ FNAL via GRACC. Appears on the weekly backup summary.
 - Waiting on STOMP endpoint (currently in ITB; waiting for Marian Babik @ CERN to test) for sending network data into GRACC ElasticSearch proper.
- ElasticSearch database backups are mostly in place (this is separate from tape archive of raw records).
- Puppetization and improved check_mk coverage is ongoing.
 - Currently working on nginx / Let's Encrypt.
- Two rounds of “cleaning the crap” over the last month. Cleaned up bad records for VOs and site names.
- Found and re-implemented another piece of hidden Gratia logic. Unfortunately, affected WLCG reports.
 - Working to document cases where workarounds for bug probes were hardcoded into Gratia collector but never fixed in the probe.
- Aiming to improve monitoring around GRACC-to-APEL uploads so we catch these issues first.

Monitoring Improvements

- **GlideInFactory logs into Elasticsearch:**
 - Done in ITB; working to roll out to production.
 - Each log line is augmented with various pilot-related attributes to help support staff search for records.
- **XRootD and GridFTP publishing to GRACC.** No progress to report - focus is on network data right now.
- **GOC historical service status archiving:** Prototype converter done, waiting on GOC to install FileBeats and test.

OSG Storage Progress

- The osgstorage.org CVMFS repositories continue to be used heavily by LIGO and others.
 - In addition to detector data, LIGO is now using CVMFS to distribute non-public results to the rest of the collaboration
- Monitoring improvements: data moved via HTTP in StashCache infrastructure now shows up in monitoring!



Stash-Bonanza

- Today & yesterday, Nebraska is hosting a hackathon for doing a write-enabled stashcp *for OSG Connect*.
Basic idea:
 - OSG-Connect submit host issues a security token.
 - HTCondor brings token along with job, renewing it as necessary.
 - When stashcp is used to stageout, token is sent along with write request.
 - Stash origin server (Xrootd server using HTTPS protocol) is augmented with the ability to authorize requests based on attached token & write into CephFS as the correct user ID.
- Even better: original idea was to have an “OSG-Connect” private CA and use GSI proxies. However: why start a new project based on X509?
 - Instead, this is the inaugural use of SciTokens' OAuth2-inspired token format.
- We have end-to-end working!
Components we had to touch:
 - Requires two new plugins to Xrootd.
 - One small patch to Xrootd.
 - Three patches to bugfix HTCondor.
 - stashcp changes.
 - “credmon” plugin for HTCondor.
 - Finish up initial release of SciTokens python library.
- Development! Unit tests! Deployment! Packaging!

Containers / GPUs

- Starting to look at how to best include GPUs in GRACC.
- OSG User Support documentation includes GPU information:
 - How to run TensorFlow on OSG: <https://support.opensciencegrid.org/support/solutions/articles/12000028940-tensorflow>
 - How to allocate GPUs on OSG: <https://support.opensciencegrid.org/support/solutions/articles/5000653025-gpu-jobs>
 - Relies on Singularity to manage complexity of library stack of modern ML environments.
- Initial successes running IceCube (via GLOW) on GPUs at Syracuse!
 - At time of writing, 34 jobs running.
- OSG and GLOW VOs should be able to access GPUs at UCSD, Nebraska, and Syracuse.
- Heterogeneity remains a problem - starting to organize thinking on how we do this for RHEL6/7 in order to better tackle GPUs.

Retirements

- HTCondor-CE 3.0.0 should have the **traceability improvements** needed to meet FNAL requirements to switch on singularity.
 - Any updates from FNAL? Gone quiet since release... assuming all is good.
- **GUMS retirement**
 - Nothing to report. Some T3s are showing interest. Need to make a push internal to CMS.
 - What is dead remains dead: Nebraska has been adjusting to the post-GUMS lifestyle.
- **VOMS-Admin retirement:**
 - OSG still needs to help glideinWMS users to sign pilot credentials directly. GLOW remains initial target.
 - This is still a waiting game: need sufficient number of sites to upgrade to OSG 3.4.

Globus Toolkit -> Closed Source

- Reminder: **OSG will support our stakeholder's usage of Globus.** Open-ended for as long as it is needed.
- Another round of meetings hosted by XSEDE (included XSEDE, PRACE, EGI, OSG, Globus).
- Starting to understand each stakeholder's needs — but no one appears to be “cutting metal” yet.

Concerns

- **Lack of progress for Globus Toolkit migrations!**
Need to get the ball rolling on this.
- **Blueprints:** Need to get these back on track. At least made a list of desired topics at the OSG Staff Retreat.
- **Effort:** Will continue to need time from Dave Dykstra and Kevin Retzke @ FNAL.