

glideinWMS Training @ IU

Condor overview

by Igor Sfiligoi (UCSD)

Acknowledgement

- These slides are heavily based on the presentation Todd Tannenbaum gave at CERN in Feb 2011

<https://indico.cern.ch/conferenceTimeTable.py?confId=124982#20110214.detailed>

Outline

- What is Condor
- Condor principles
- Condor daemons
- Condor protocol overview

What is Condor

What is Condor

- Condor is a Workload Management System
 - i.e. a batch system
- Strong points
 - Fault tolerant
 - Robust feature set
 - Flexible
- Development team dedicated to working closely w/ scientific community as priority #1

How can Condor be used

- Managing local processes (local)
- Managing local cluster (~vanilla) ← Only vanilla in this talk
- Connecting clusters (flocking)
- Handling resource overlays (glideins)
- Swiss-knife for accessing other WMS (Condor-G)
 - e.g. Grid, Cloud, pbs, etc.

(Vanilla) Condor principles

(Vanilla) Condor principles

- Two parts of the equation
 - Jobs
 - Machines/Resources
- Jobs
 - Condor's quanta of work
 - Like a UNIX process
 - Can be an element of a workflow
- Machines
 - Represent available resources
 - Mostly CPU, but indirectly memory and disk as well

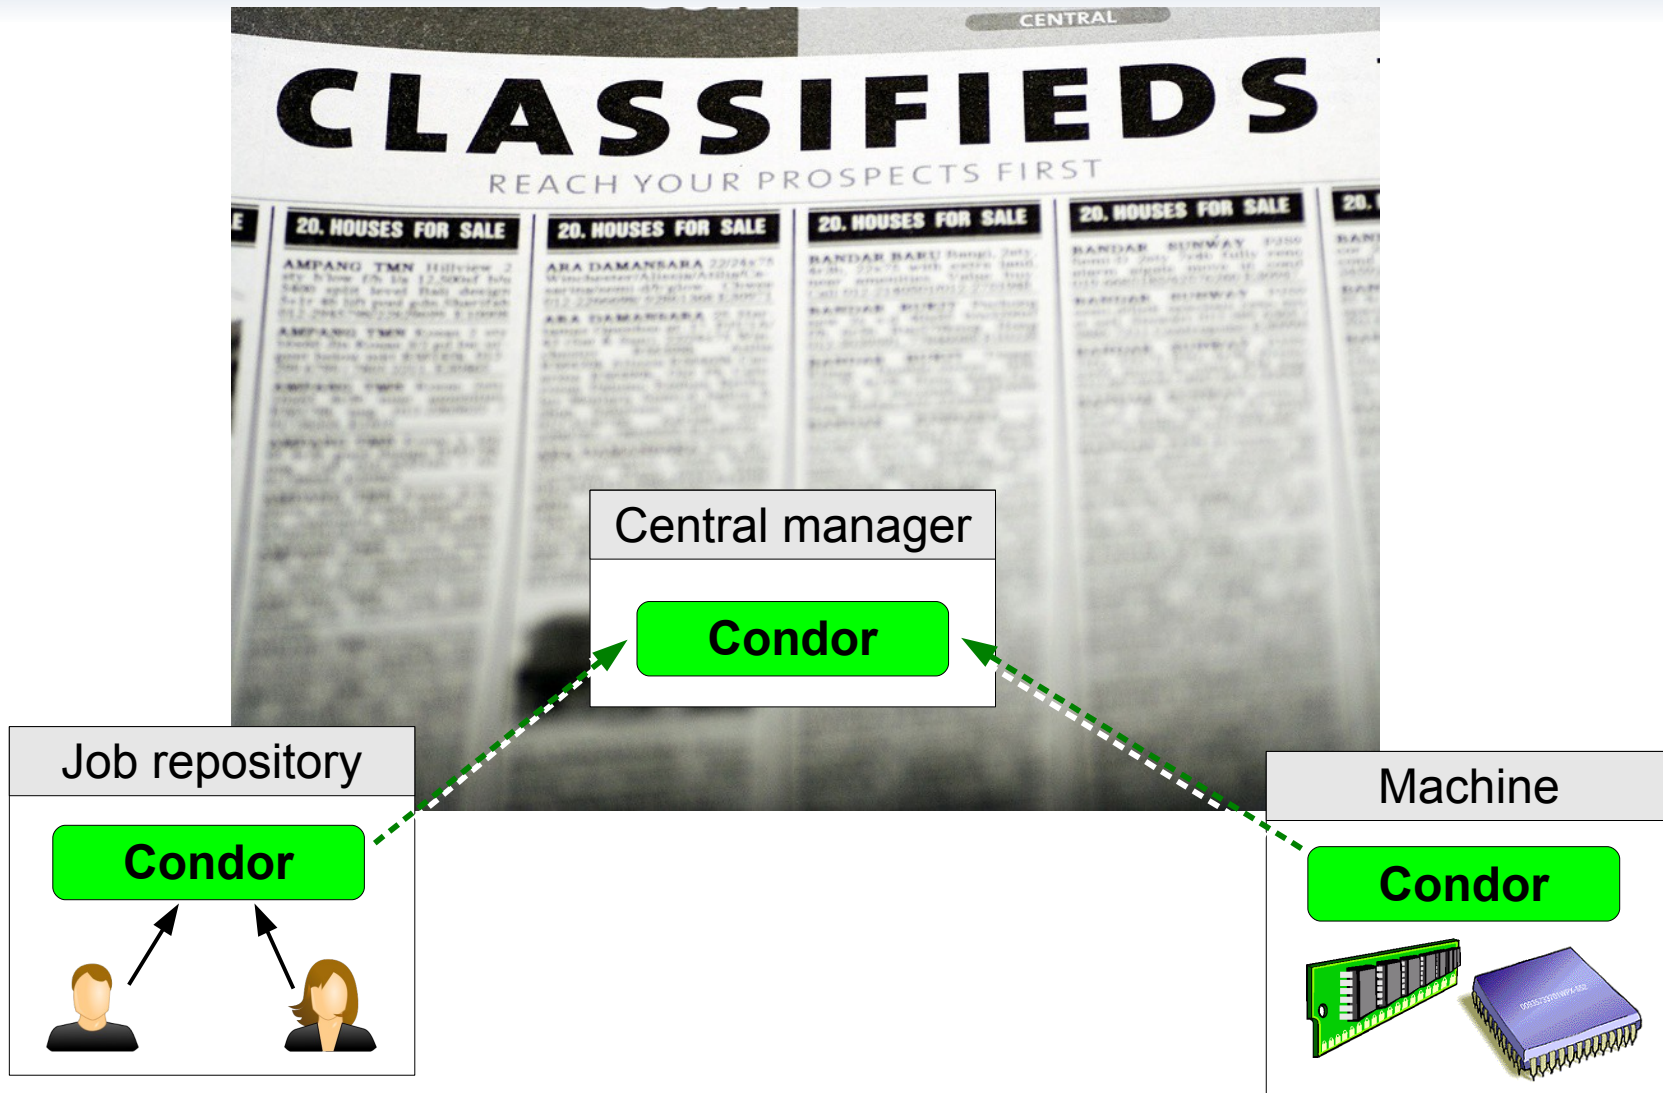
Jobs Have Wants & Needs

- Jobs state their requirements and preferences:
 - Requirements:
 - I **require** a Linux/x86 platform
 - Preferences ("Rank"):
 - I **prefer** a machine owned by CMS
- Jobs describe themselves via attributes:
 - Standard, i.e. defined by Condor:
 - I am owned by Albert
 - Custom, i.e. specified by the user (or the administrator):
 - I am a Monte Carlo job
 - I will be done within 12h

Machines Do Too!

- Machine requirements and preferences:
 - Requirements:
 - I **require** that jobs declare a runtime shorter than 18h
 - Preferences ("Rank"):
 - I **prefer** Monte Carlo jobs
- Machine attributes:
 - Standard, i.e. defined by Condor:
 - I am a Linux node
 - I control 2GB of memory
 - Custom, i.e. specified by the administrator:
 - I have been paid with CMS money

Condor brings them together



Condor ClassAds



What are Condor ClassAds?

- ClassAds is a language for objects (jobs and machines) to
 - Express attributes about themselves
 - Express what they require/desire in a match (similar to personal classified ads)
- Structure
 - Set of attribute name/value pairs
 - Value : Literals (string, bool, int, float) or an expression

Example ClassAd

```
MyType = "Machine"  
TargetType = "Job"  
Name = "glidein_999@cabinet-2-2-1.t2.ucsd.edu"  
Machine = "cabinet-2-2-1.t2.ucsd.edu"  
StartDlpAddr = "<169.228.131.179:56787>"  
State = "Claimed"  
Activity = "Busy"  
Cpus = 1  
Memory = 36170  
Disk = 231463800  
OpSys = "LINUX"  
Arch = "X86_64"  
Requirements = JOB_Is_ITB != true  
Rank = 1  
KFlops = 972989  
Mips = 3499  
HasFileTransfer = true  
IS_GLIDEIN = true  
GLIDEIN_SEs = "bsrm-1.t2.ucsd.edu"  
DaemonStartTime = 1324784426
```

ClassAd Expressions

- Similar look to C : operators, references, functions
- **Operators**: +, -, *, /, <, <=, >, >=, ==, !=, &&, and || all work as expected
 - Type checking ops: ==?, !=
- **Functions**: if/then/else, string manipulation, list operations, dates, randomization, ...
- **References**: to other attributes in the same ad, or attributes in an ad that is a candidate for a match
- True==1 and False==0 (guaranteed)
 - e.g. (3 == (2+True)) is identical to True


http://www.cs.wisc.edu/~border/manual/v7.64_1/Condr_s_ClassAd.html#SECTION00512300000000000000

Example Expression

```
ifthenelse(LastVacateTime=?=UNDEFINED,  
    ifthenelse(NormMaxMins!=UNDEFINED,  
        (NormMaxMins*60)<(ToRetire+JobMaxTime-MyCurrentTime),  
        (8*3600)<(ToRetire+JobMaxTime-MyCurrentTime)),  
    ifthenelse(MaxMins!=UNDEFINED,  
        (MaxMins*60)<(ToRetire+JobMaxTime-MyCurrentTime),  
        (16*3600)<(ToRetire+JobMaxTime-MyCurrentTime)))&&  
(ImageSize<(MaxMemMBs*1024))&&  
(stringListMember(GLIDEIN_SEs,DESIRED_SEs,",")=?=True)&&  
(JOB_Is_ITB != TRUE)
```

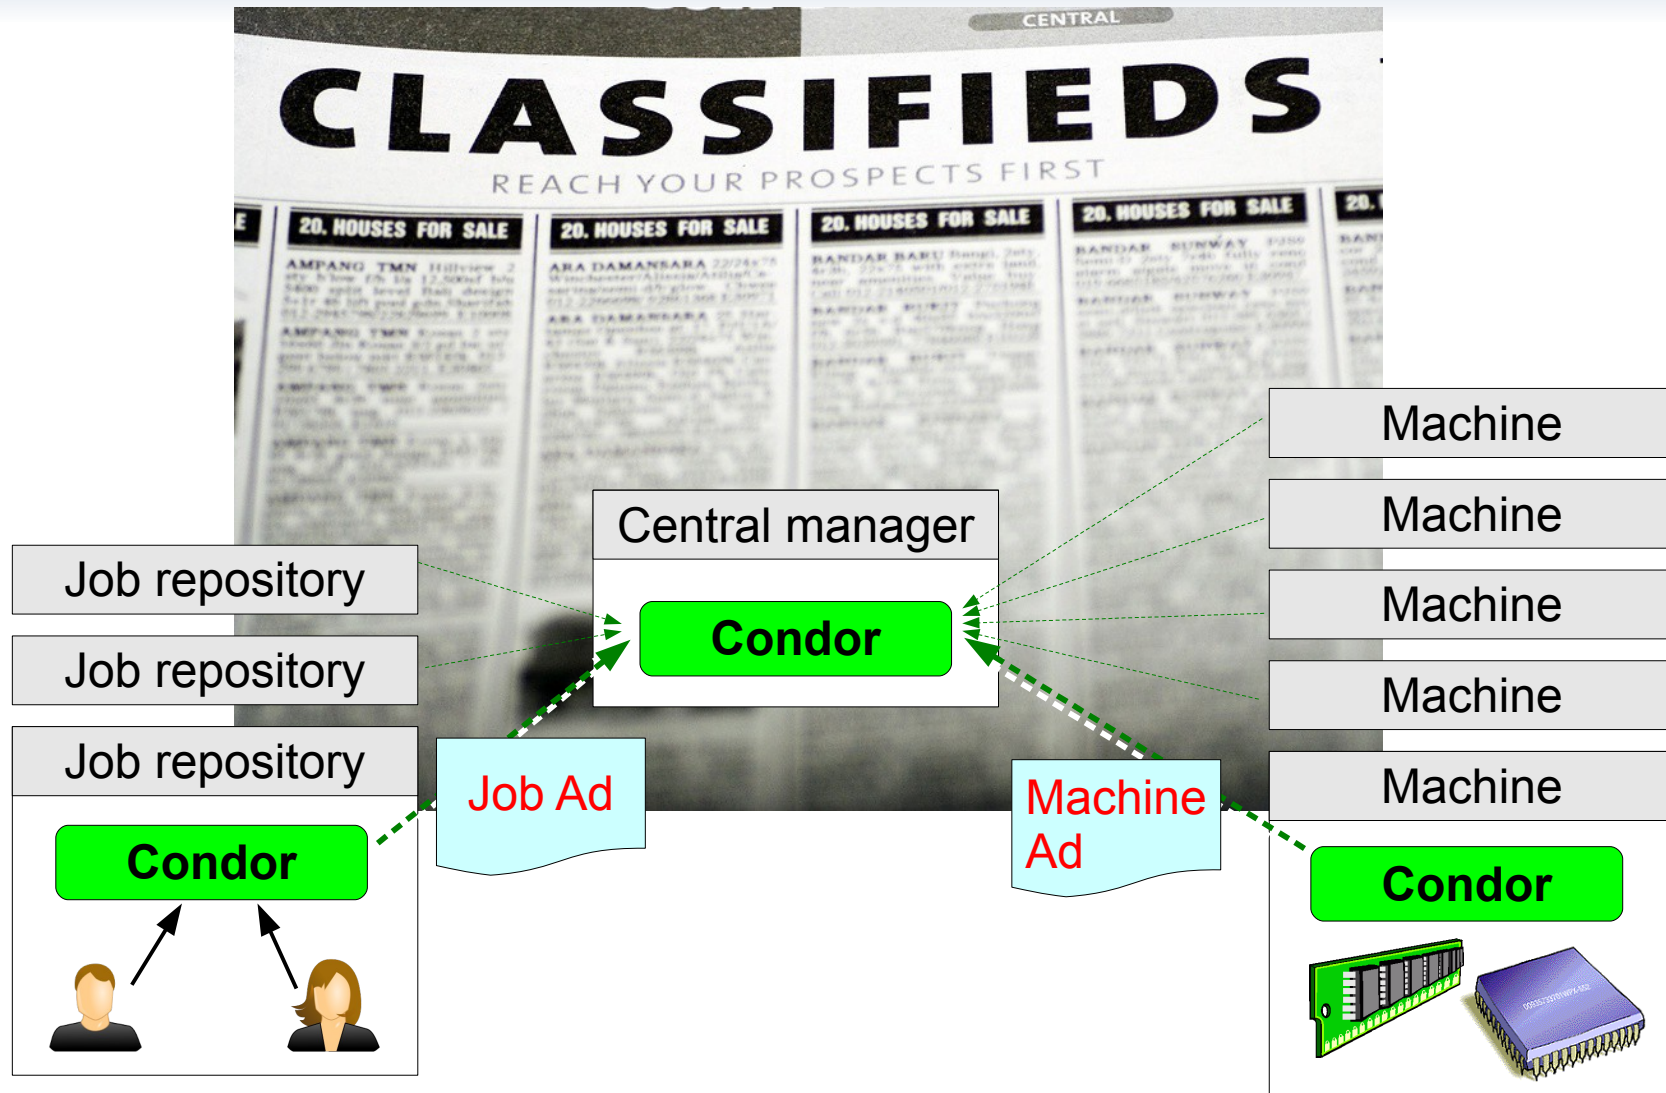

ClassAd Types

- Condor has many types of ClassAds
 - A "**Job Ad**" represents a job to Condor
 - A "**Machine Ad**" represents a computing resource
 - Others types of ads represent instances of other services, users, licenses, etc

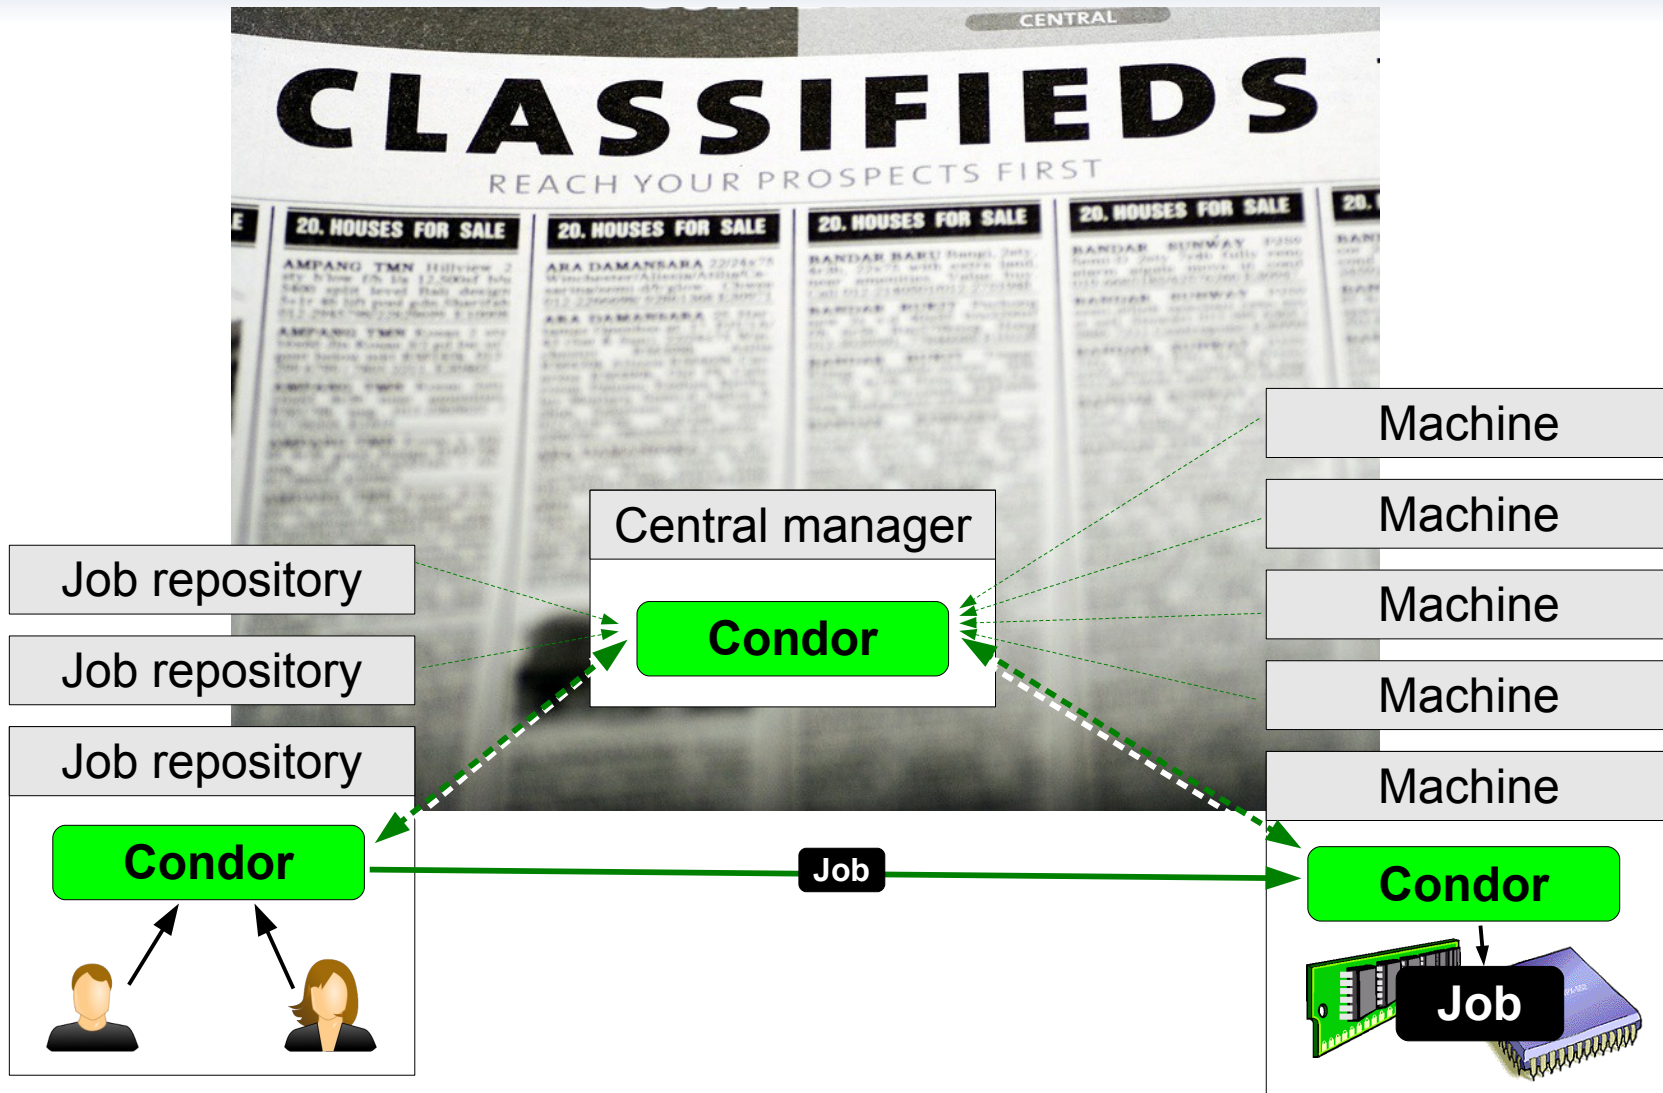


glideinWMS defines some

Central Manager holds them all



Match & start



The Magic of Matchmaking

- Two ads match if **both their Requirements** expressions evaluate to True
 - If more than one match, the match with the **highest Rank is preferred** (float)
- Condor evaluates job ads in the context of a candidate machine ad looking for a match
 - MY.name – Value for attribute “name” in local ClassAd
 - TARGET.name – Value for attribute “name” in match candidate ClassAd
 - Name – Looks for “name” in the local ClassAd, then

Example Fancy Match

Pet Ad

MyType = "Pet"
TargetType = "Buyer"
Requirements =
 DogLover =?= True
Rank = 0
PetType = "Dog"
Color = "Brown"
Price = 75
Breed = "Saint Bernard"
Size = "Very Large"
...

Dog == Resource ~= Machine

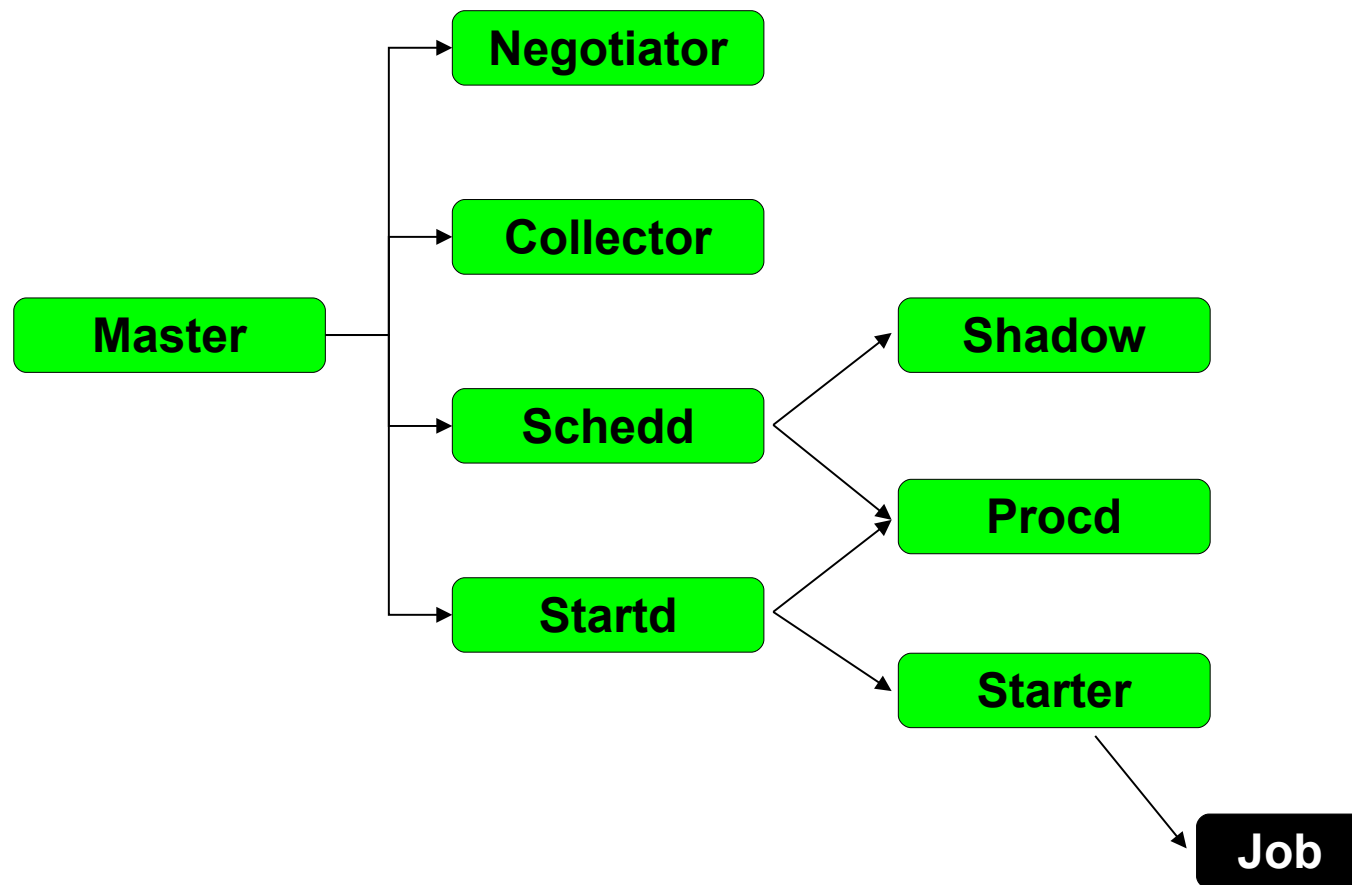
Buyer Ad

MyType = "Buyer"
TargetType = "Pet"
Requirements =
 (PetType == "Dog") &&
 (TARGET.Price <= MY.AcctBalance) &&
 (Size == "Large" || Size == "Very Large")
Rank = (Breed == "Saint Bernard")
AcctBalance = 100
DogLover = True
...

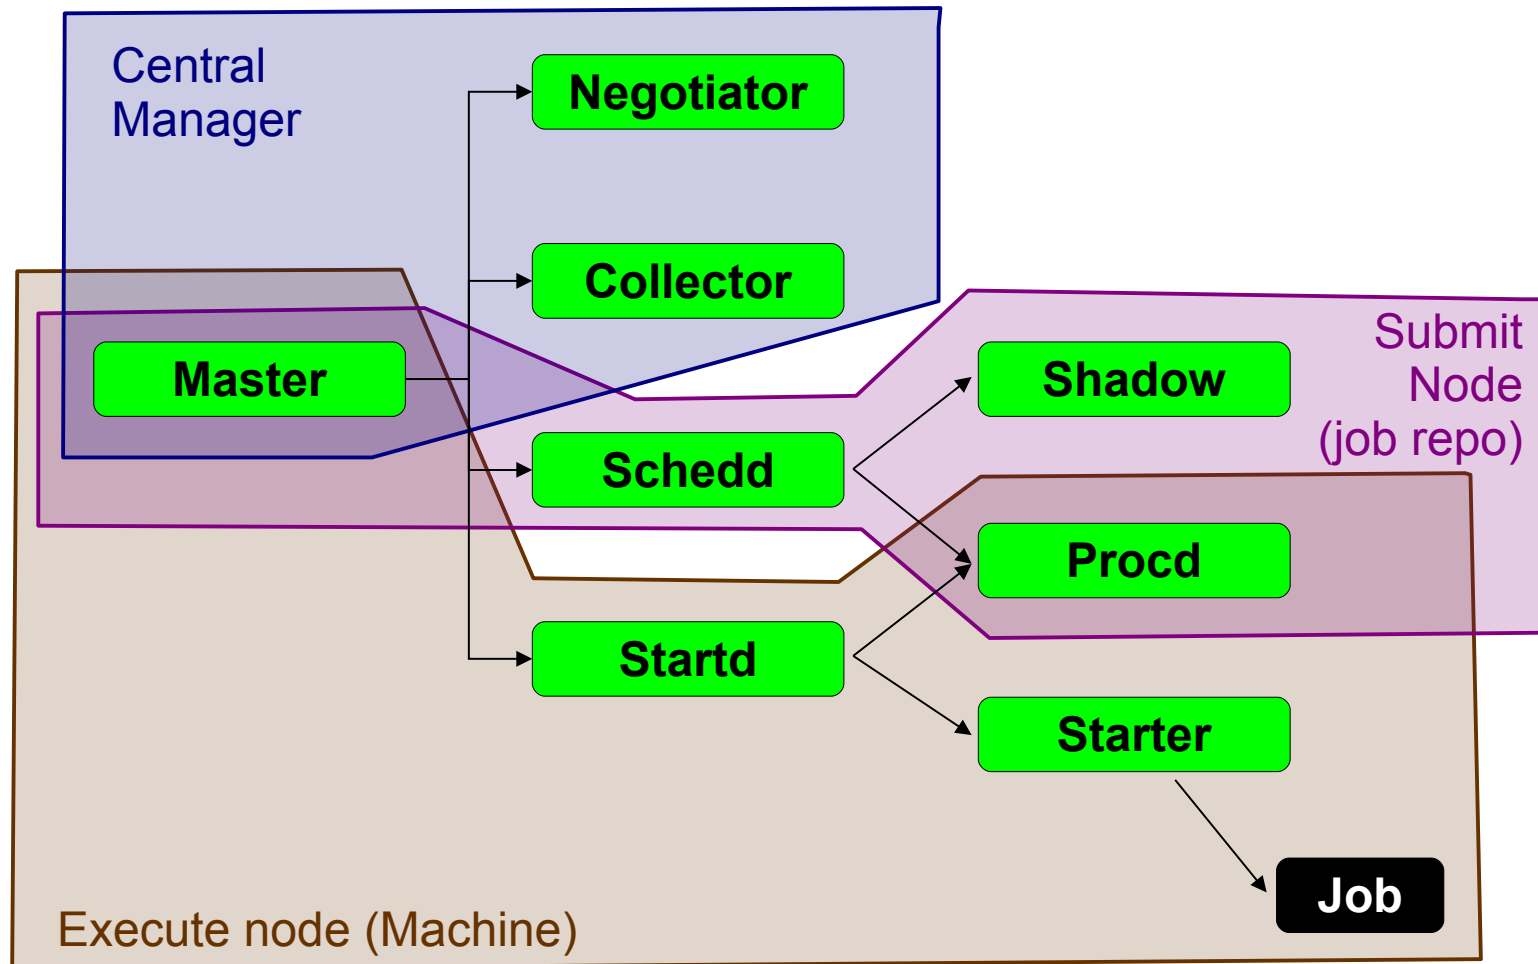
Buyer ~= Job

(Vanilla) Condor Daemons

Condor Daemons – Mix'n Match Components



Condor Daemons – Mix'n Match Components



condor_master

- You start it, it **starts up the other Condor daemons**
 - If a daemon exits unexpectedly, restarts daemon and emails administrator
 - If a daemon binary is updated (timestamp changed), restarts the daemon
- Provides access to many remote administration commands:
 - condor_reconfig, condor_restart, condor_off, condor_on, etc.
- Default server for many other commands:
 - condor_config_val, etc.

condor_procd

Submit node

Execute node

- Monitors all other processes on the node
 - Information then used by the other daemons
- Builds process tree
 - Tracks birth and death of processes
 - Monitors resource consumption (memory, CPU)

condor_schedd

Submit node

- **Represents jobs** to the Condor pool
- Maintains **persistent queue of jobs**
 - Queue is not strictly first-in-first-out (priority based)
 - Each machine running **condor_schedd** maintains its own independent queue
- Responsible for contacting available machines and spawning waiting jobs
 - When told to by condor_negotiator
- Services most user commands:
 - condor_submit, condor_rm, condor_q

condor_shadow

Submit node

- Spawned by condor_schedd
- **Represents a running job** on the submit machine
 - **Yes, one per running job**
- Handles file transfers
- Enforces Periodic_* expressions
 - Hold, release, remove, ...

condor_startd

Execute node

- **Represents a machine** willing to run jobs to the Condor pool
- Run on any machine you want to run jobs on
- **Enforces the wishes** of the machine owner (the owner's "policy")
- Starts, stops, suspends jobs
- Provides other administrative commands
 - for example, condor_vacate

condor_starter

Execute node

- Spawned by the condor_startd
- Handles all the details of **starting and managing the job**
 - Transfer job's binary to execute machine
 - Send back exit status
 - Etc.
- **One per running job**
 - The default configuration is willing to run one condor_starter per CPU

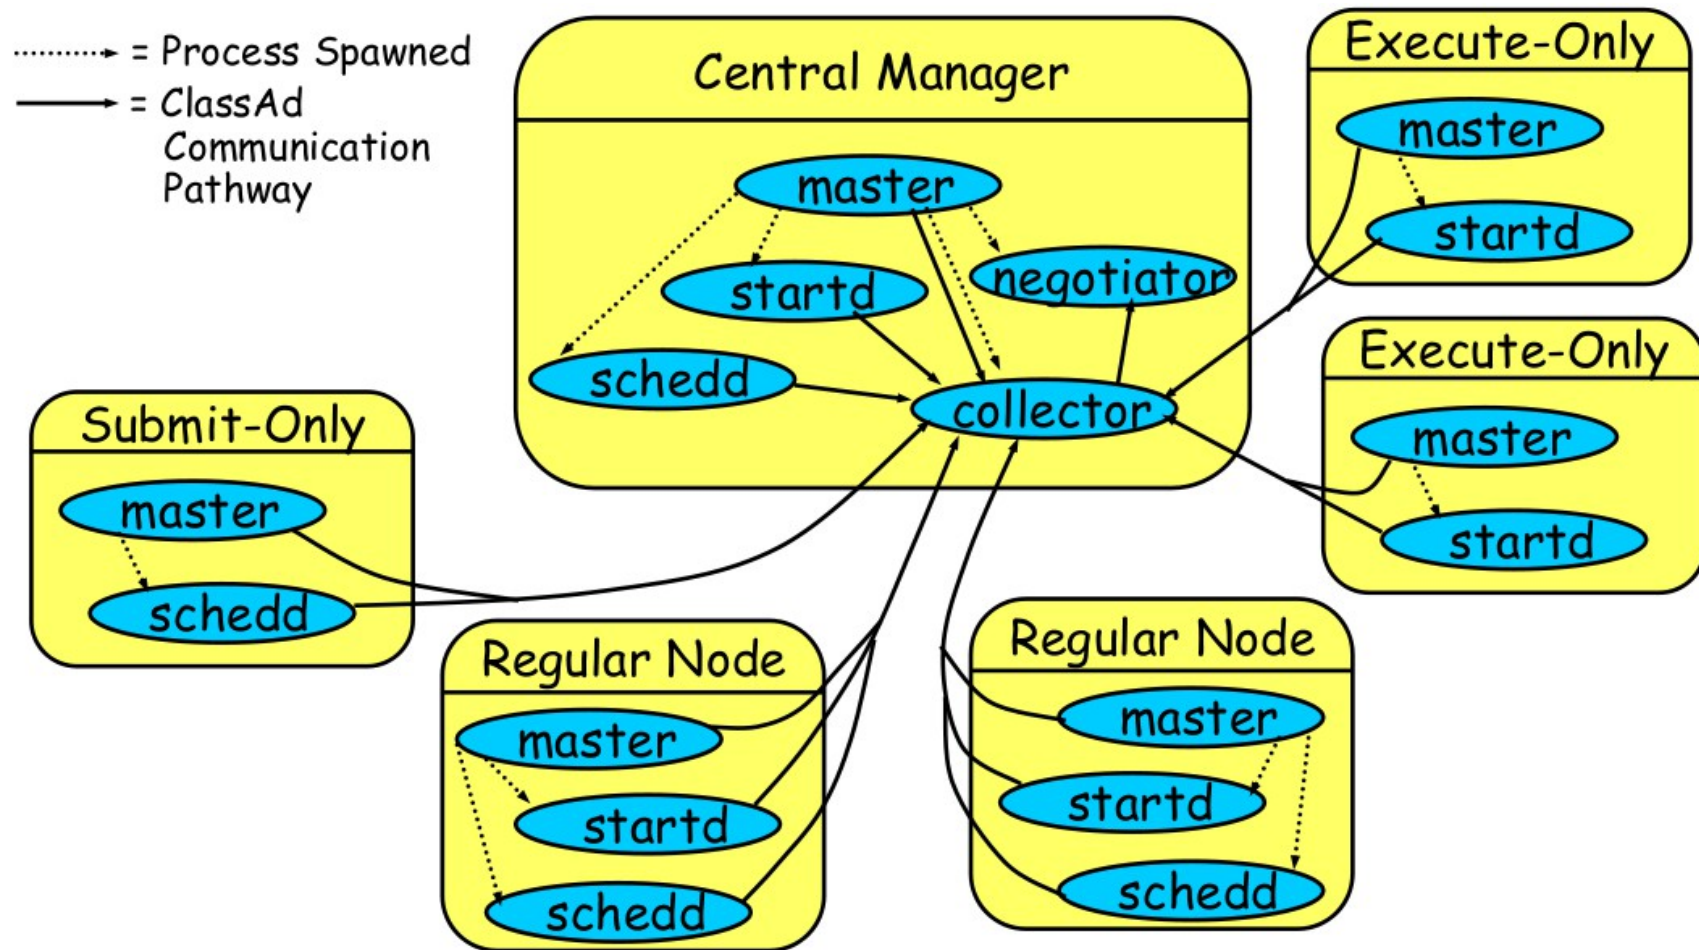
condor_collector

- **Collects information** from all other Condor daemons in the pool
- Each daemon sends a periodic update called a ClassAd to the collector
 - **Old ClassAds removed after a timeout** (~15 mins)
- Services queries for information:
 - Queries from other Condor daemons
 - Queries from users (condor_status)

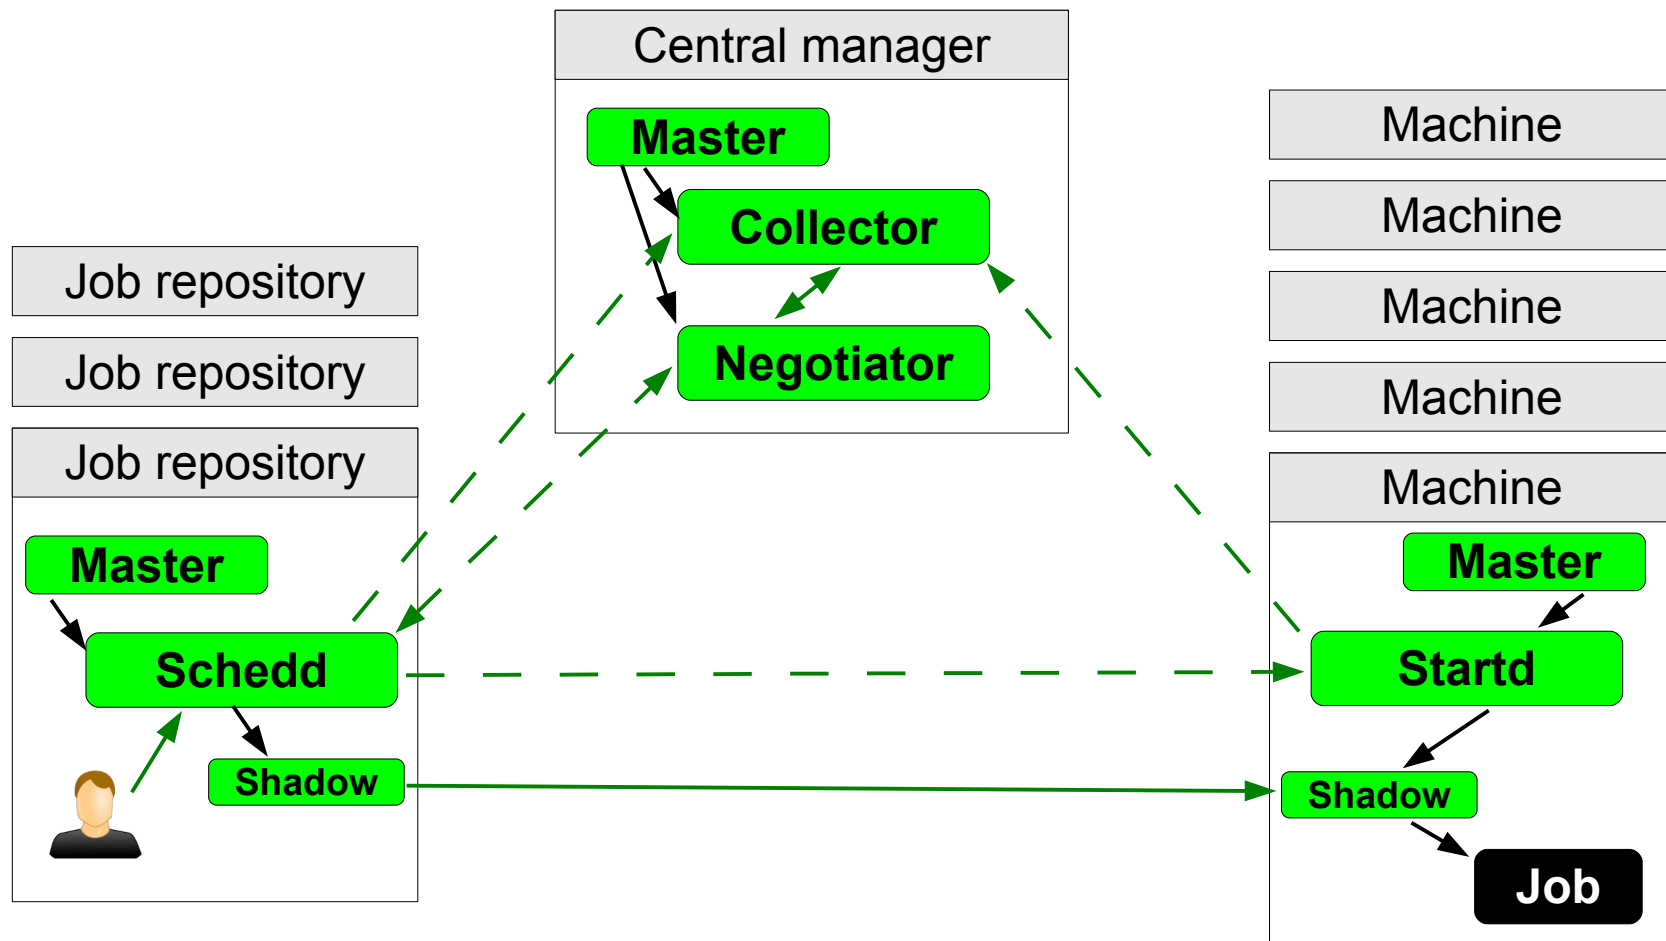
condor_negotiator

- **Performs matchmaking** in Condor
 - Pulls list of available machines from condor_collector, gets jobs from condor_schedds
 - Matches jobs with available machines
 - Both the job and the machine must satisfy each other's requirements (2-way matching)
- **Handles user priorities and accounting**

Sample Condor pool

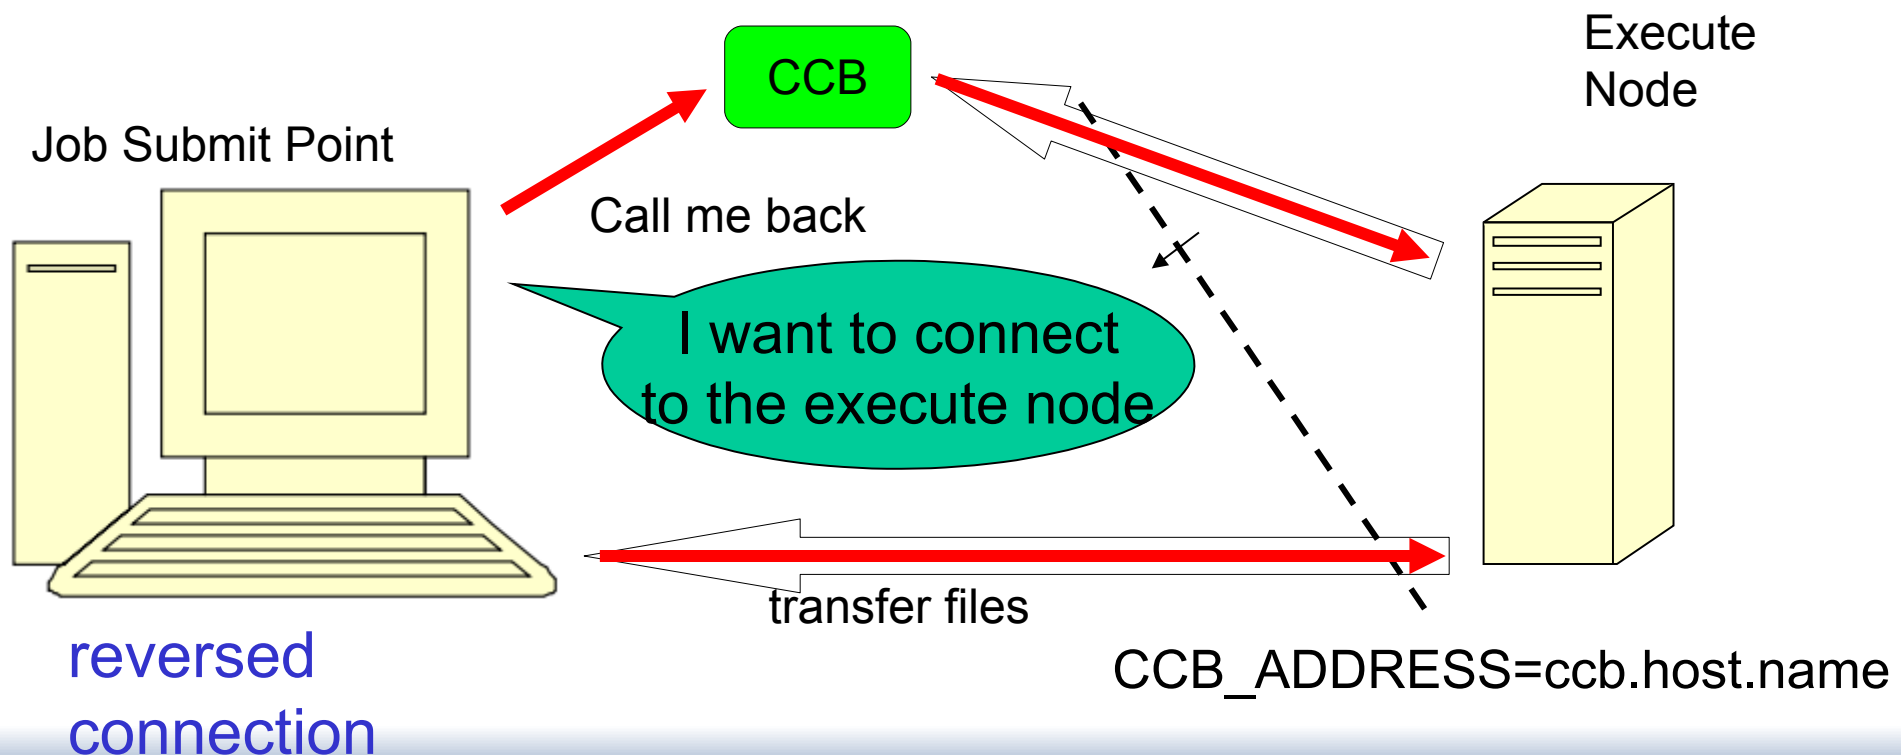


Sample Condor pool



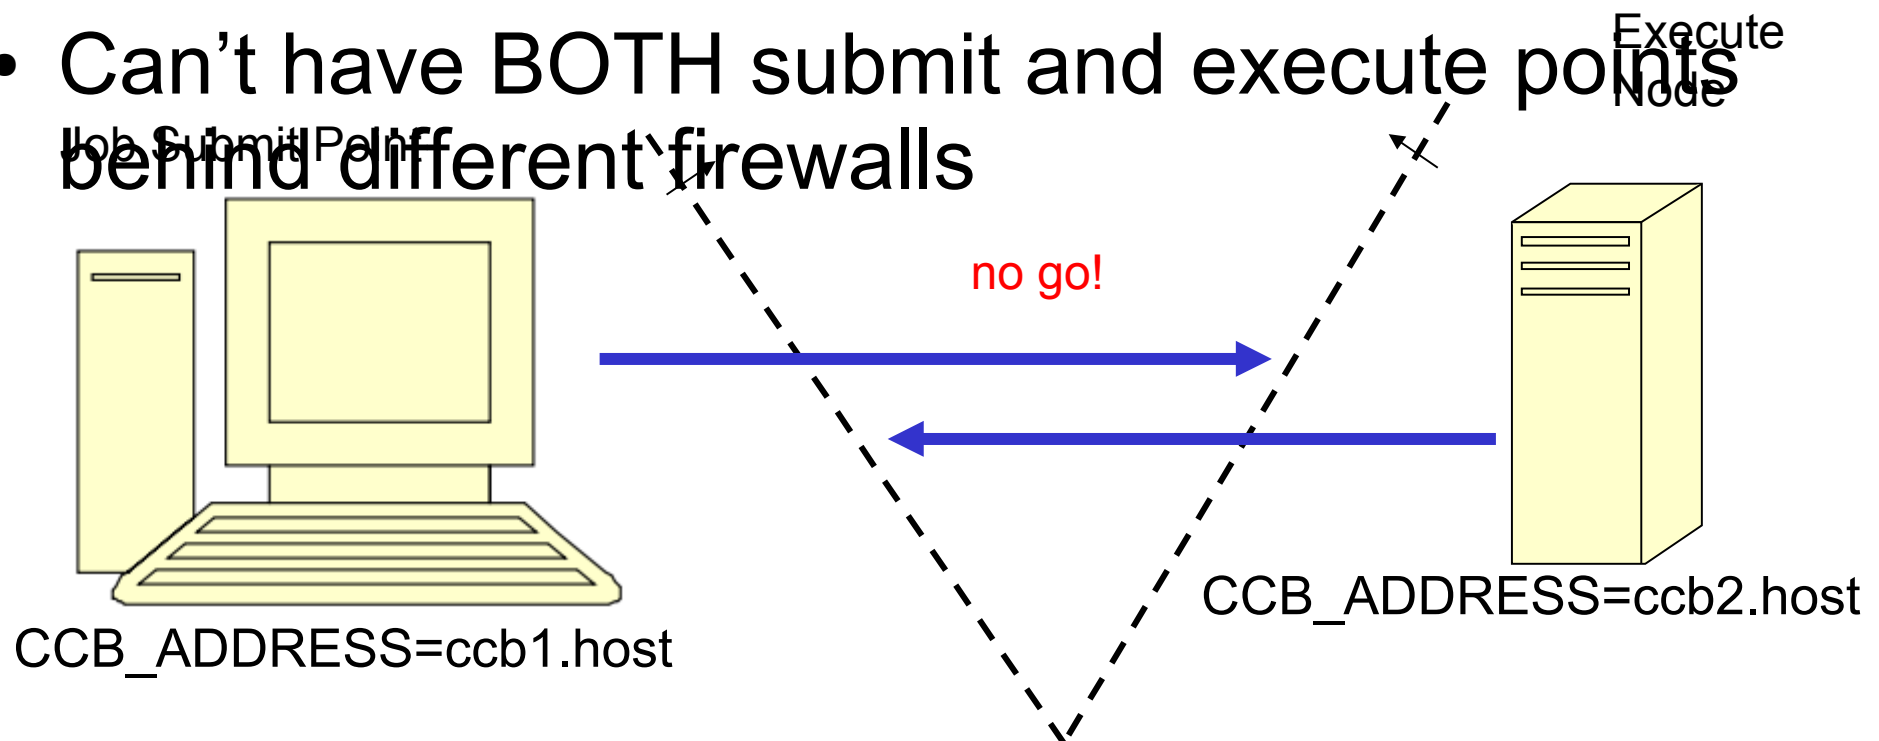
CCB: Condor Connection Broker

- Condor wants two-way p2p connectivity
- With CCB, one-way is good enough
 - Collector requests reversed connections for clients



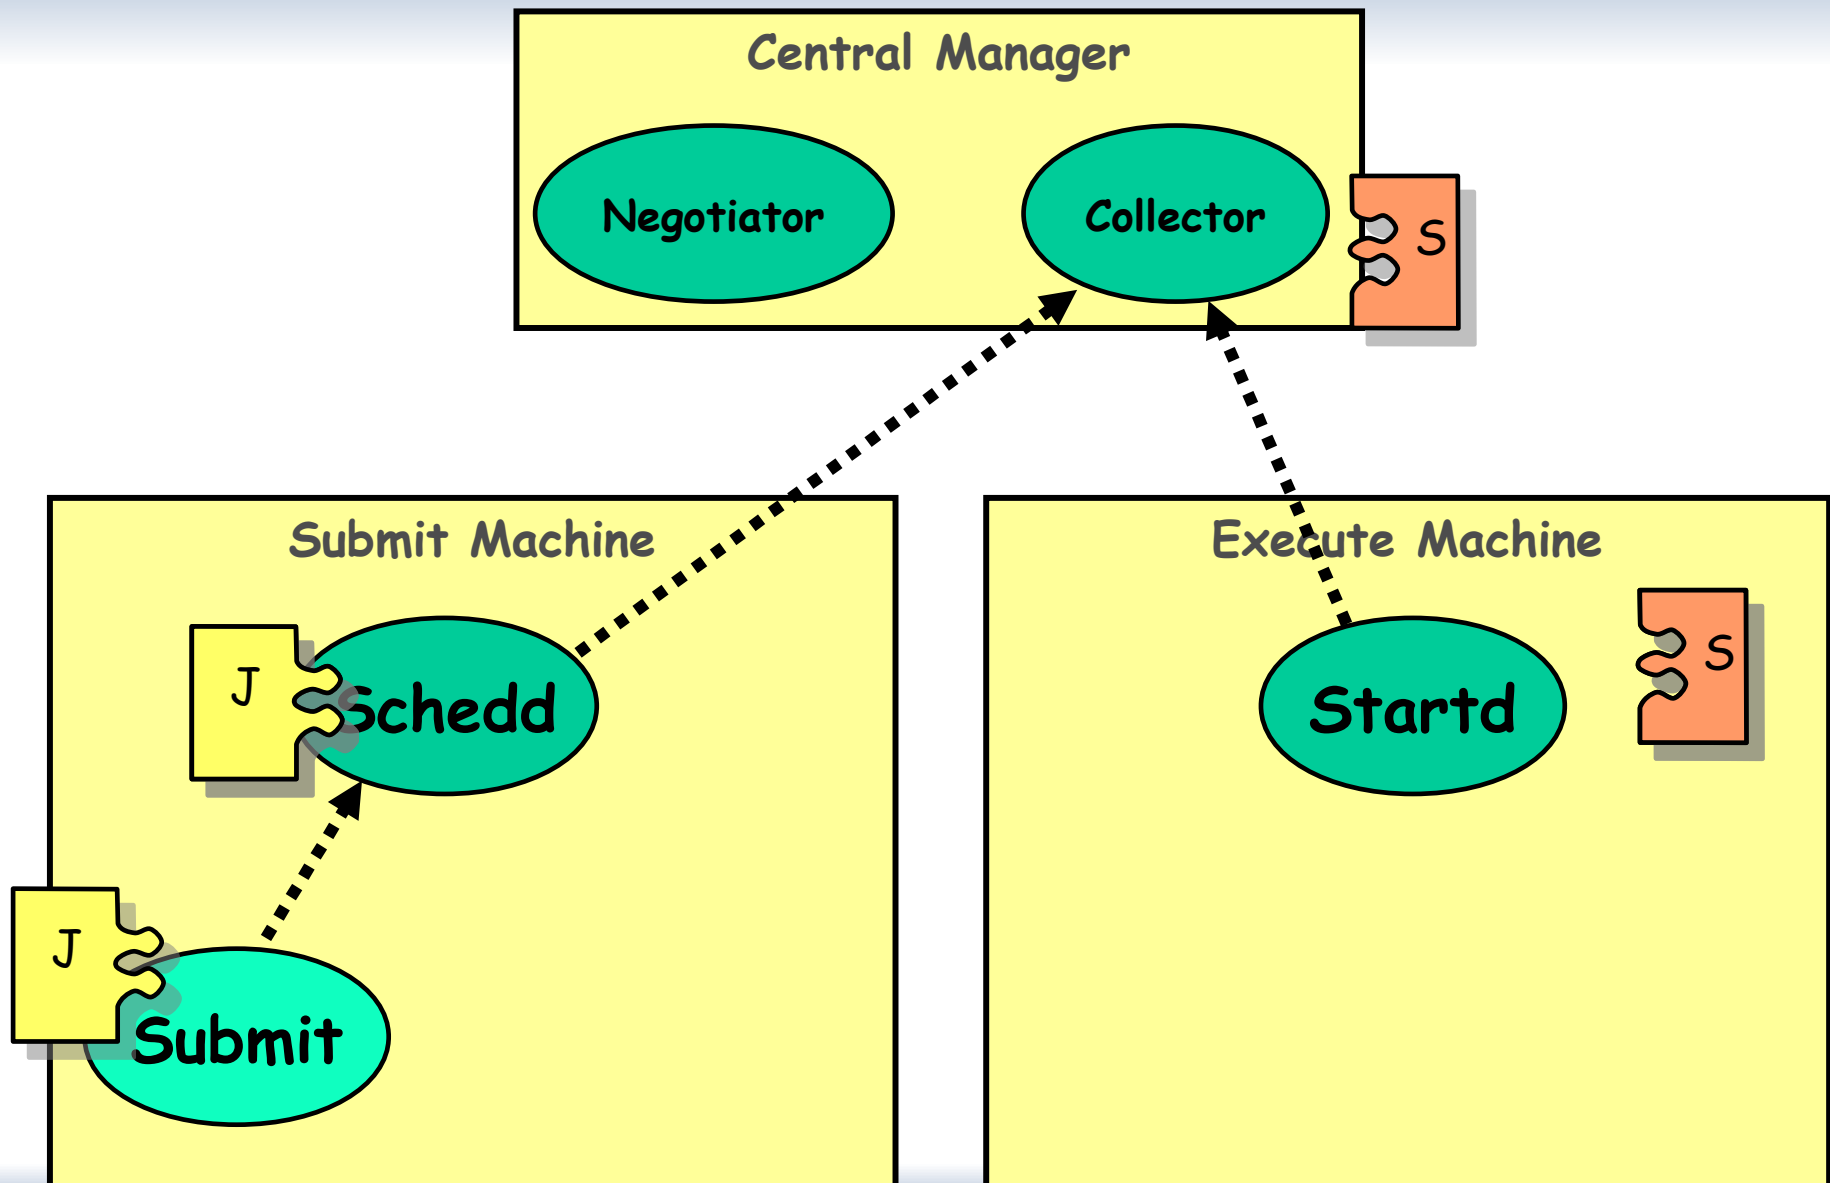
Limitations of CCB

- Collector (CCB Broker) needs to be accessible by everyone
- Requires outgoing connectivity
- Can't have BOTH submit and execute points behind different firewalls

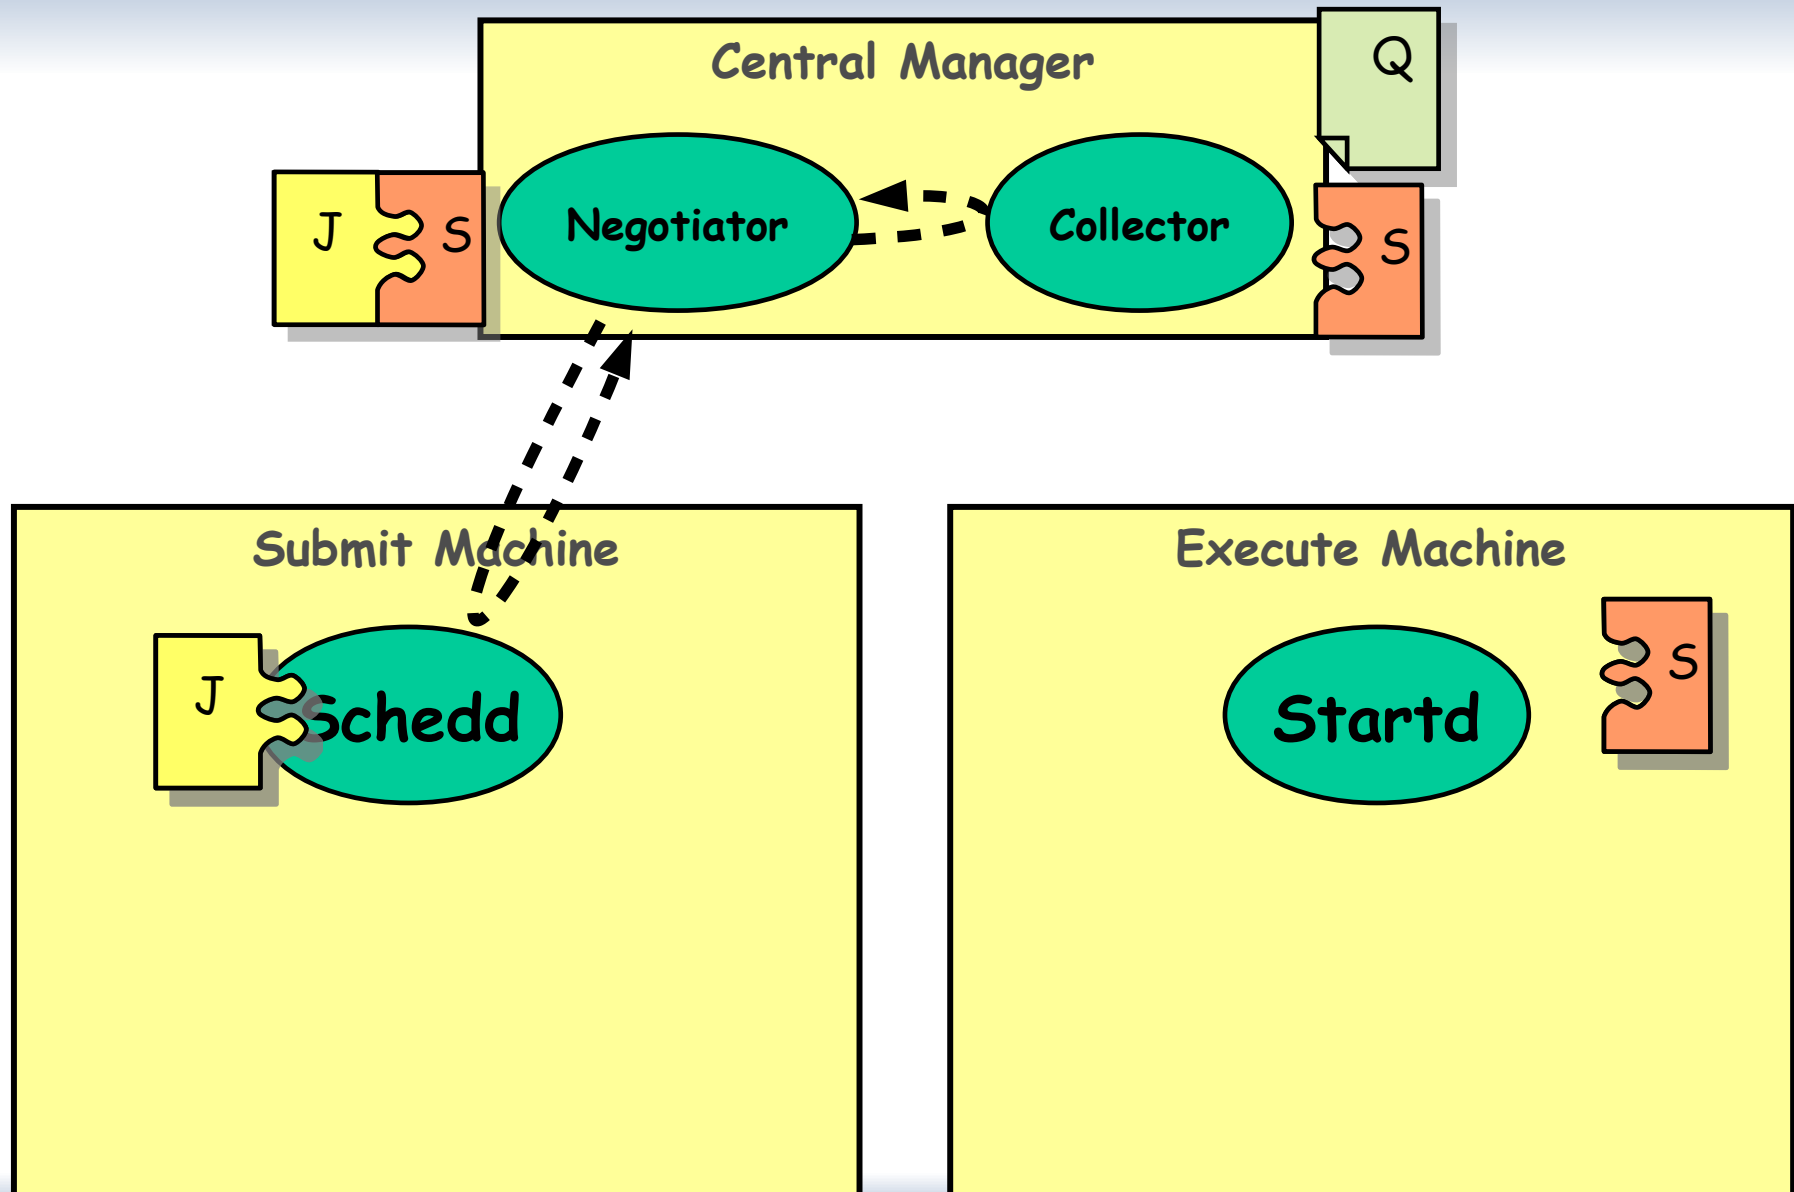


(Vanilla) Condor protocol

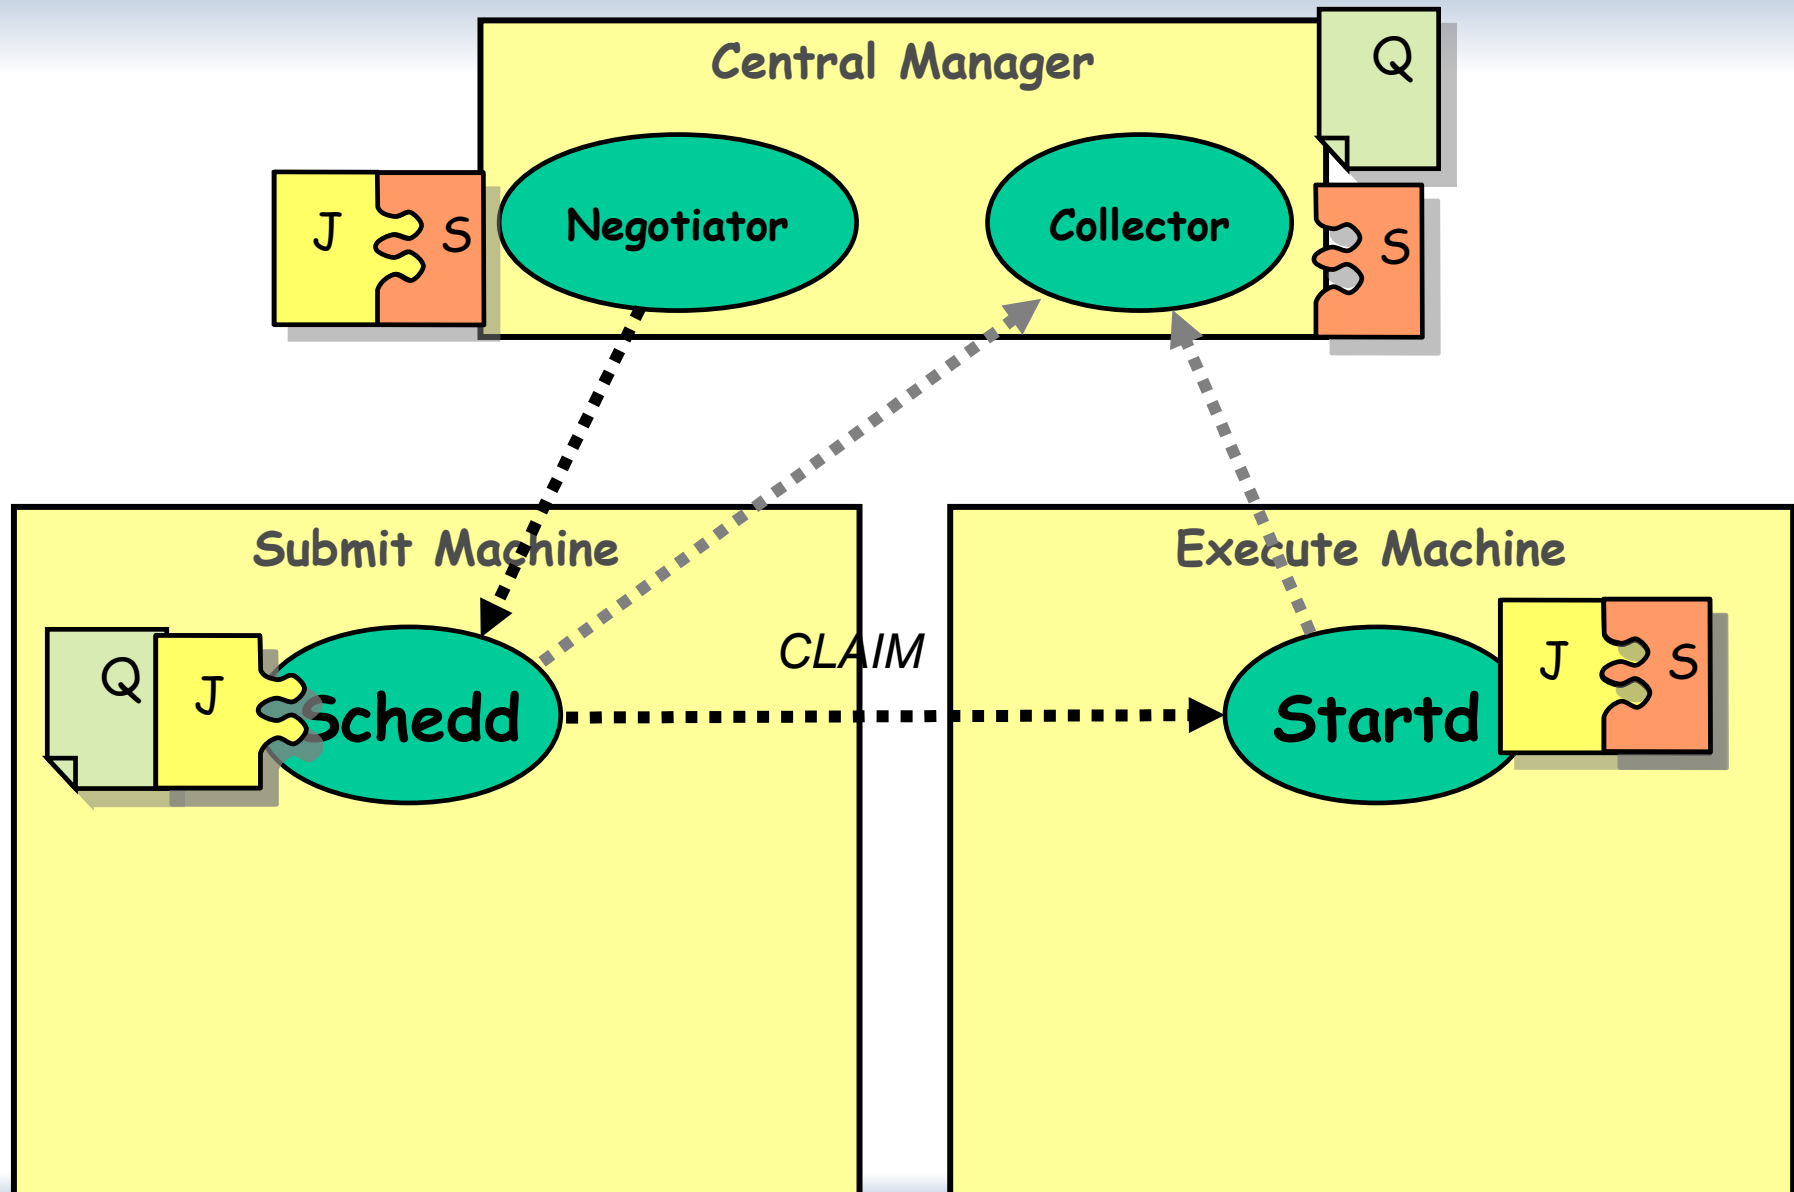
Claiming Protocol



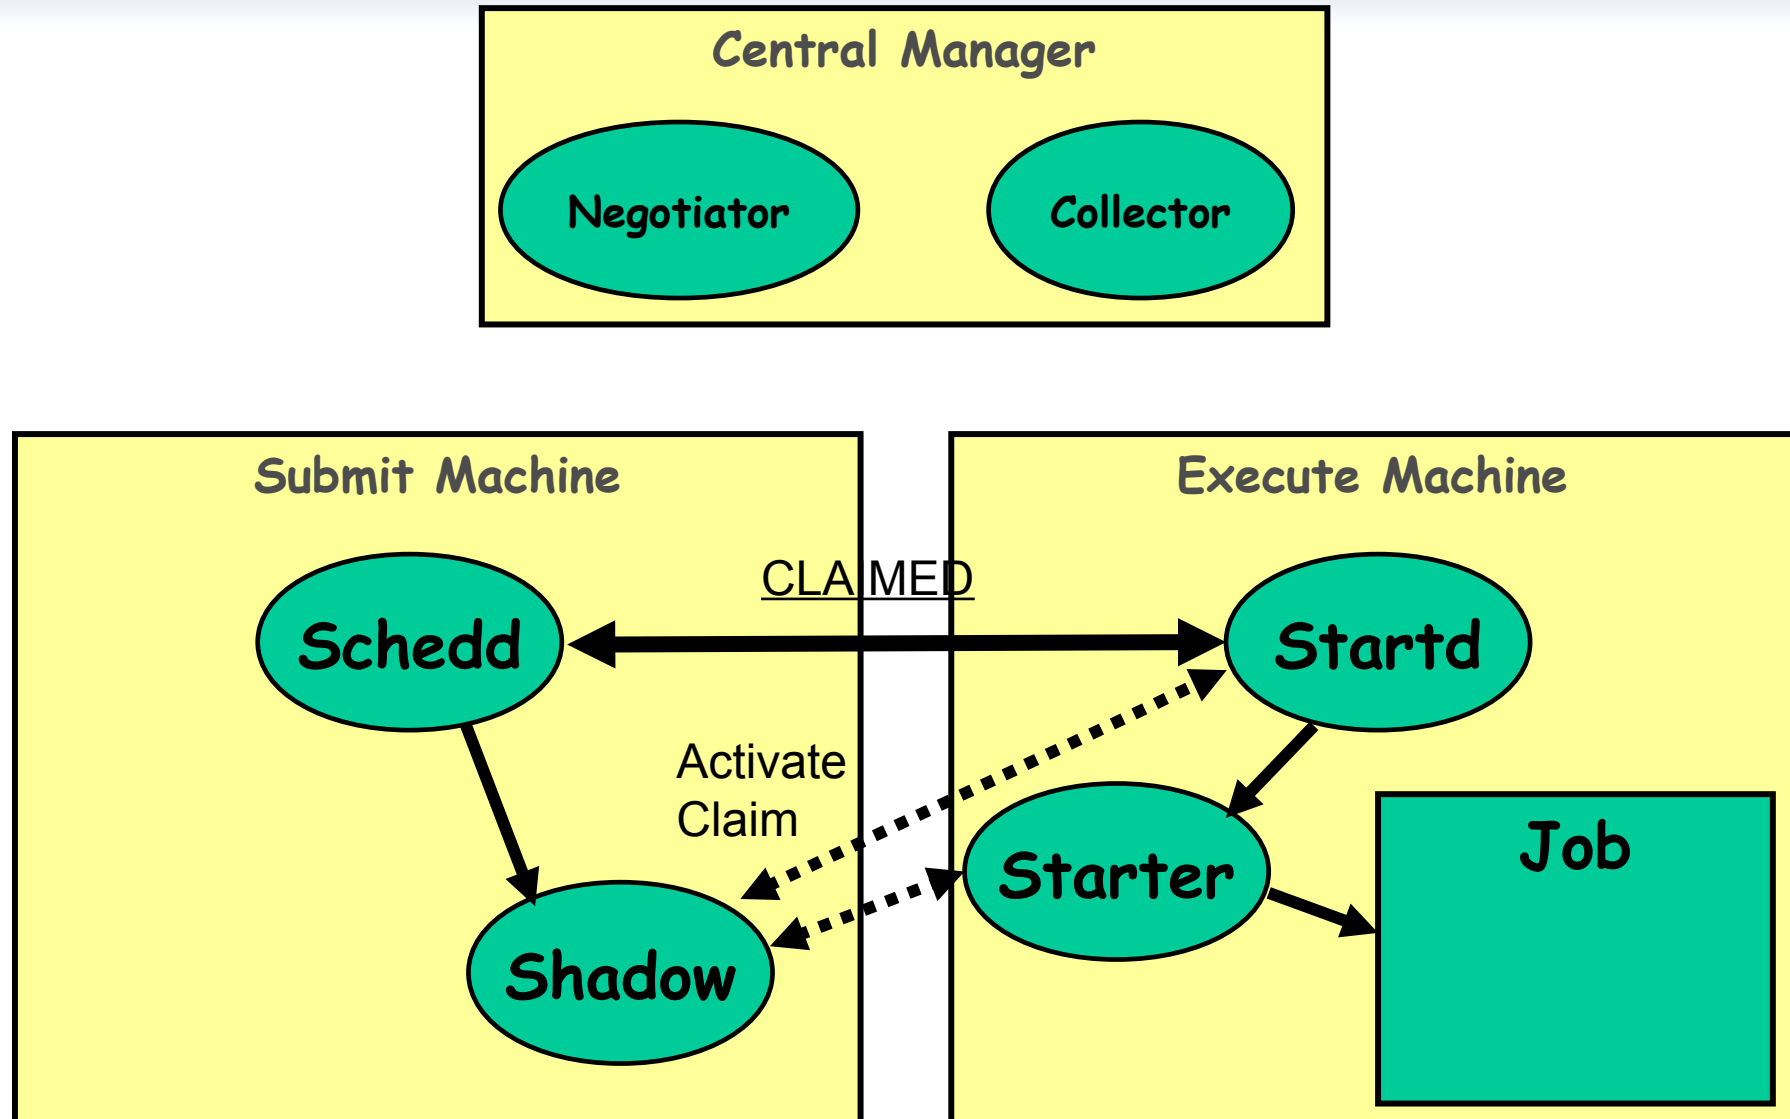
Claiming Protocol



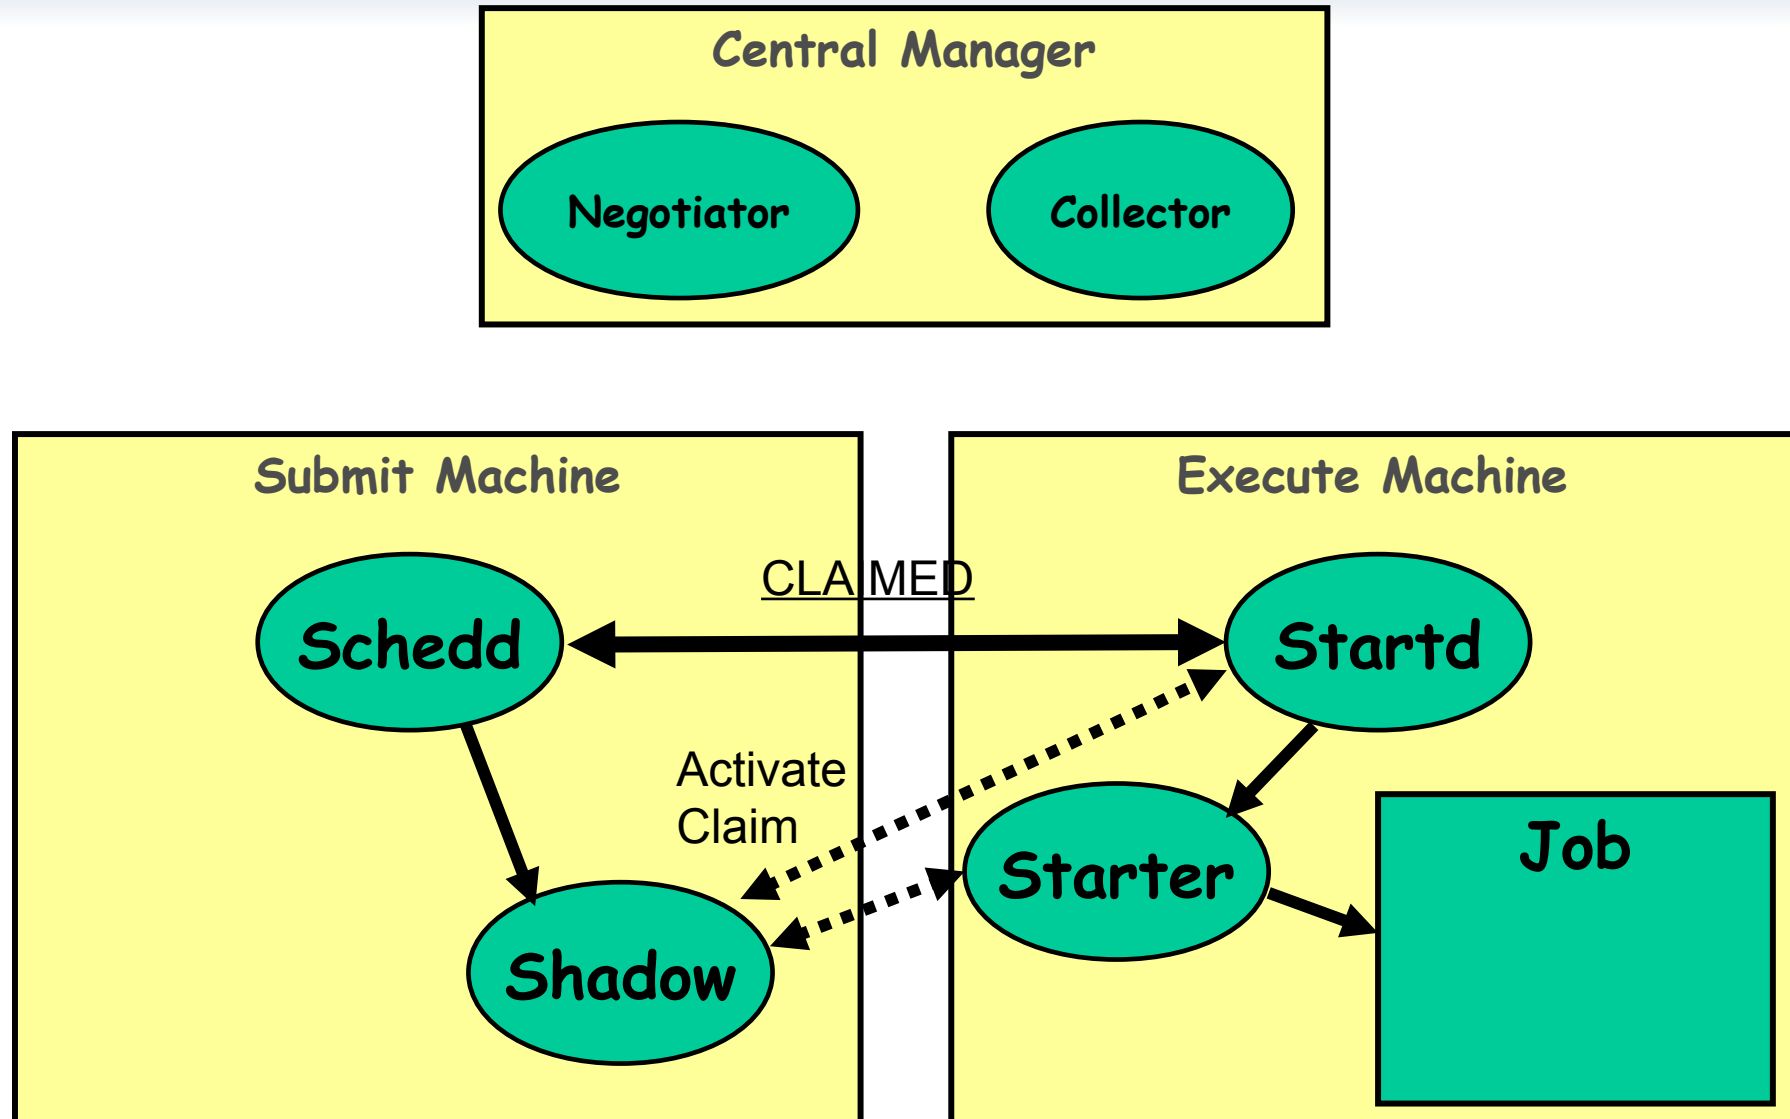
Claiming Protocol



Claim Activation



Repeat until Claim released



When is claim released?

- When relinquished by one of the following
 - lease on the claim is not renewed
 - Why? Machine powered off, disappeared, etc
 - schedd
 - Why? Out of jobs, shutting down, schedd didn't "like" the machine, etc
 - startd
 - Why? Policy re claim lifetime, prefers a different match (via Rank), non-dedicated desktop, etc
 - negotiator
 - Why? User priority inversion policy
 - explicitly via a command-line tool
 - E.g. condor_vacate

Defining Rank is dangerous!
(preemption)

The end

The Condor Project (Established '85)

- Research and Development in the Distributed High Throughput Computing field
- Team of ~35 faculty, full time staff and students
 - Face software engineering challenges in a distributed UNIX/Linux/NT environment
 - Are involved in national and international grid collaborations
 - Actively interact with academic and commercial entities and users
 - Maintain and support large distributed production environments
 - Educate and train students

The Condor Team



Pointers

- Condor Home Page
<http://www.cs.wisc.edu/condor/>
- Condor Manual
<http://www.cs.wisc.edu/condor/manual/v7.6/>
- Support
condor-user@cs.wisc.edu
condor-admin@cs.wisc.edu