Open Science Grid

Operations Training: The OSG Software Stack

Alain Roy

OSG Software Coordinator

Open Science Grid

Crash course: What you need to know to support the OSG Software Stack

Alain Roy

OSG Software Coordinator

Open Science Grid

Mistakes we made in creating OSG Software, and your life of pain in supporting it

Alain Roy

OSG Software Coordinator

Goals

- Help you understand OSG Software 3.0
- Help you deal with incoming support tickets

Open Science Grid

Rough Agenda

Time	Topic
9:00 – 10:00	Deep intro to OSG Software 3
10:00 – 11:15	Live demo and coffee break
11:15 – 12:00	Debugging common problems
12:00 – 1:00	Lunch
1:00 – 2:00	More debugging common problems
2:00 – 3:00	*Alain runs OSG Software Meeting*
3:00 – 3:15	Coffee break
3:15 – 5:00	Open discussion aka stump the expert
6:00	Dinner?

Part 1:
The Big Picture

Open Science Grid

I'm assuming:

- You know all the basics about OSG 1.2
- You know how to use Pacman

- If you want lessons on how to use Pacman, we'll have to do a separate training session.

Differences from OSG Software 1.2

- No more Pacman: RPM & yum
- Install software into standard locations
  - /usr/bin/grid-proxy-init
  - /usr/sbin/globus-gridftp-server
  - /etc/osg/config.d/
  - /var/log/globus-gatekeeper.log
- No more $VDT_LOCATION
  - Nor $OSG_LOCATION, $GLOBUS_LOCATION…
- Life is easiest if you're using standard tools & cluster management system

Requirements for using OSG Software 3.0

- RHEL, CentOS, or Scientific Linux
  - Version 5 right now
  - Version 6 on April 24th
- We rely on software from three sources:
  - OSG
  - OS
  - EPEL
- You install software with:
  - Yum
  - RPM

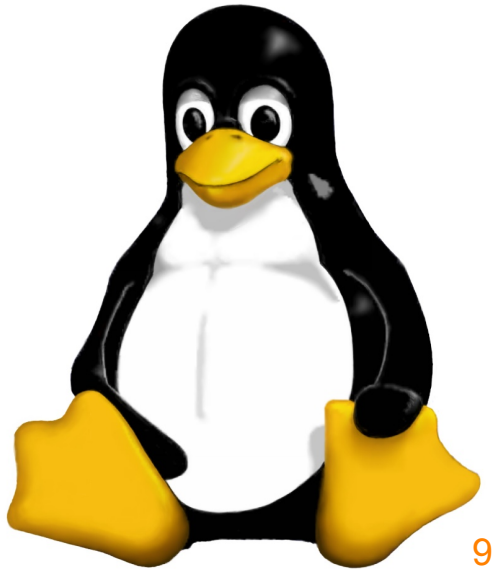

Open Science Grid

Software changes

Upgrades:

Software	New Version	Old Version
Globus	5.2	4.0
Bestman	2.2	2.0
VOMS	2.0	1.8
Condor	7.6	7.2
glexec	0.8	0.6
hdfs	0.2	---

Changes:

New	Old
lcmaps	PRIMA
configure-osg	osg-configure
---	vdt-control

Globus upgrade: Implications

- This is a really big upgrade!
  - From 4.0 to 5.2
  - And we shipped 5.2 *before* Globus did
- Big changes:
  - One jobmanager per user instead of per job
  - Globus gatekeeper is long-running daemon
- This all works pretty well…
  - But we've been fixing a steady stream of problems over the last few months
  - Expect more problems for a bit longer?

VOMS Upgrade: Implications

- VOMS Admin is fancier:
  - Remove need for VOMRS
  - Let's users sign AUP, expire access…
- New command-line client (`voms-proxy-init`) is *not* compatible with oldest VOMS servers:
  - We pushed hard to upgrade all VOMS servers in OSG
  - Fermilab:
    - Was last to upgrade (complex setup)
    - Is having many problems (see GOC Ticket)
    - Dan threatened downgrading
    - We're trying hard to resolve their problems ASAP

PRIMA -> lcmaps: implications

- It all just works, so don't worry…
  - …But the configuration is all different
  - …And the logging is all different
  - So it may end up generating some support
- Configuration for using gums is in:

`/etc/lcmaps.db`

No more vdt-control???

- It used to be:

```
# vdt-control --on
enabling init service apache... ok
enabling inetd service globus-gatekeeper... ok
...
```

- Now we use the standard interface:

```
# /sbin/service httpd start
Starting httpd:                          [   OK   ]

# /sbin/service globus-gatekeeper start
Started globus-gatekeeper                [   OK   ]
```

Migrating from OSG 1.2

- It's best to do a fresh install
  - Don't want libraries from old install messing with libraries from new install
- The old "config.ini" for osg-configure is now multiple files:

```
# ls /etc/osg/config.d
01-squid.ini        20-condor.ini   40-localsettings.ini
10-misc.ini         30-cemon.ini    40-network.ini
10-storage.ini      30-gip.ini      40-siteinfo.ini
15-managedfork.ini  30-gratia.ini
```

Migrating from OSG 1.2

- Option 1:
  - Install from RPMs
  - Edit relevant bits of all config.ini files
- Option 2:
  - Install from RPMs
  - Copy old config.ini to 99-config.ini
  - Edit 30-gip.ini to comment out GIP config

```
https://twiki.grid.iu.edu/bin/view/Documentation/
Release3/UpdateComputeElement
```

Open Science Grid

Part 2:
Let's explore a CE

Open Science Grid

A map of the CE

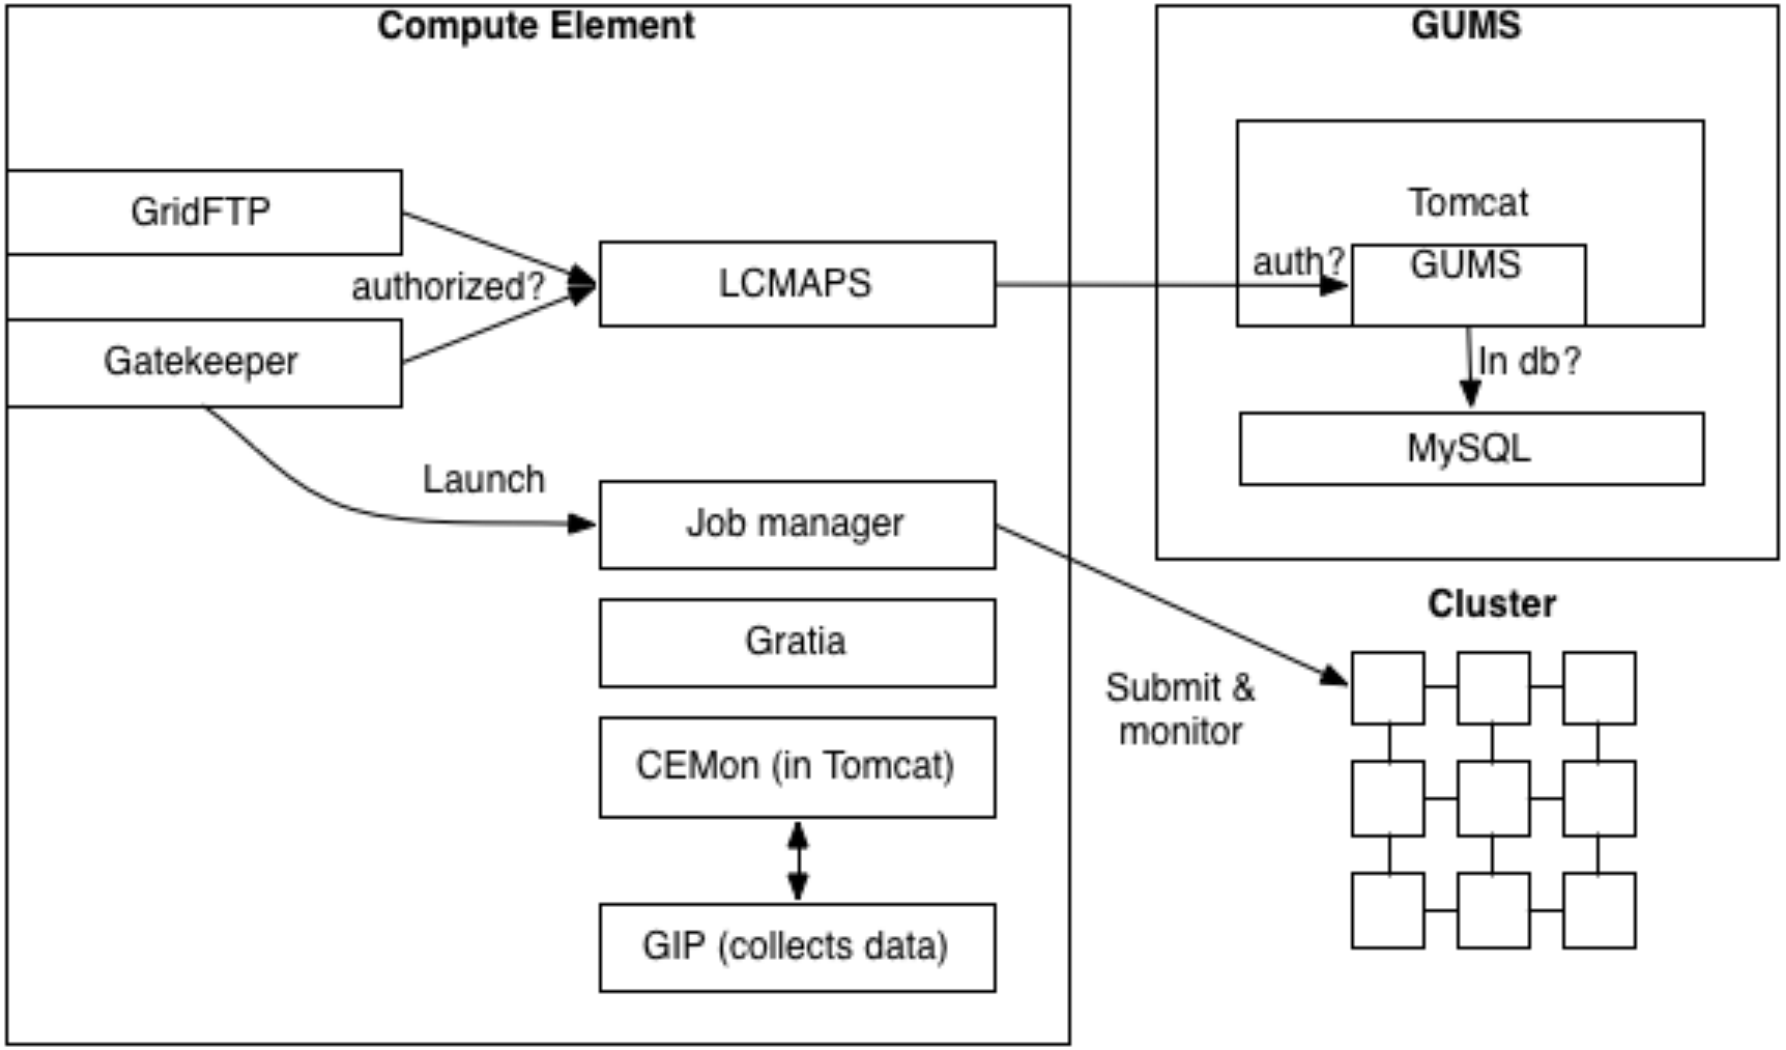

Open Science Grid

The gatekeeper

- Long-running daemon
  - Used to be invoked per request
- Authorizes user to submit jobs
- Creates job manager per user/DN
- Configuration in:
  - /etc/sysconfig/globus-gatekeeper
- Log in:
  - /var/log/globus-gatekeeper.log

The job manager

- One long-lived job manager per user
  - Handles job submission, removal, etc
  - The user's main point of contact
- Configuration:
  - /etc/globus/globus-gram-jobmanager.conf
  - /etc/globus/globus-fork.conf
  - /etc/globus/globus-condor.conf
  - /etc/globus/globus-pbs.conf
- Log: (not very useful)
  - /var/log/globus/gram_USER.log-date

Open Science Grid

GridFTP

- Long-lived daemon
- Configuration:
  - /etc/gridftp.conf
  - /etc/sysconfig/gridftp.conf.d
- Logs:
  - /var/log/gridftp-auth.log
  - /var/log/gridftp.log

Open Science Grid

lcmaps

- Lcmaps handles authorization
  - Can do gridmap (but we don't)
  - Used for GUMS
  - Used by Gatekeeper/GridFTP & glexec
- Configuration:
  - /etc/lcmaps.db
- Log:
  - /var/log/messages
  - Will be more useful with next update

Open Science Grid

CEMon & GIP

- Massive application in Tomcat
  - For our purposes, “fancy cron”
  - We could replace with “cron and a script”
    - We have, but it’s never been fully tested/supported
- CEMon invokes GIP
  - Every five minutes
  - GIP collects static & dynamic info
- CEMon pushes data to GOC
- Configuration:
  - /etc/osg/config.d/
- Log:
  - /var/log/glite-ce-monitor/glite-ce-monitor.log

Information Services

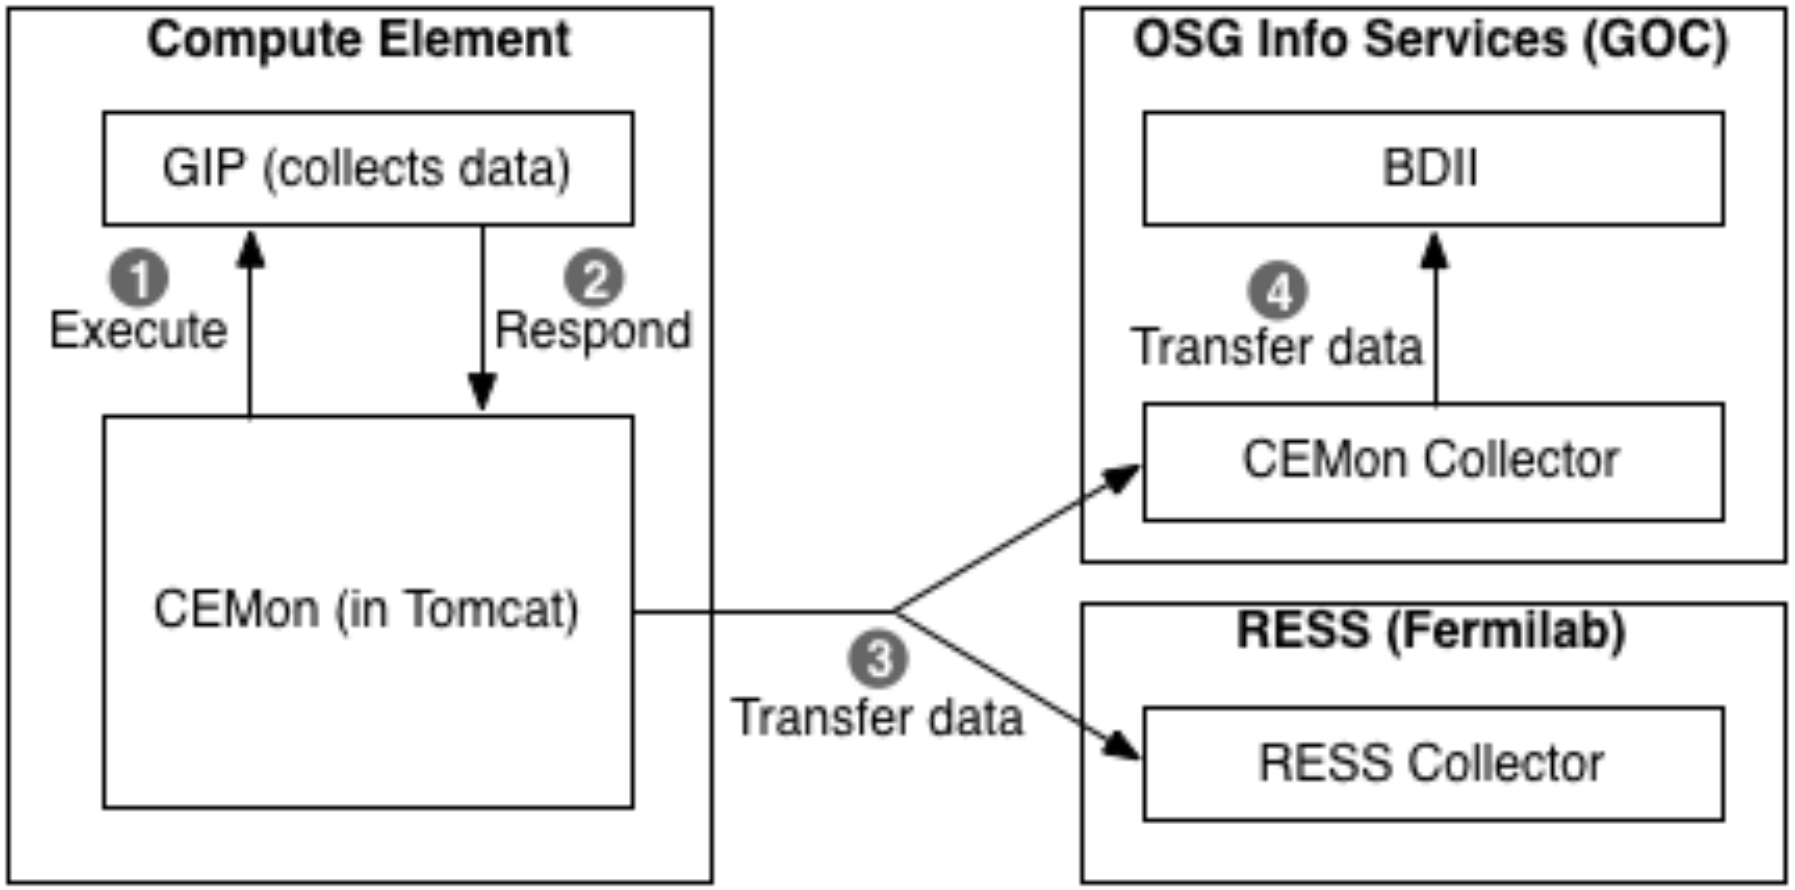

Gratia

- Gratia runs as cron job
  - But controlled from init script

    `/sbin/service gratia-probes-cron start`
- Periodically collects information from:
  - Batch system
  - GridFTP
  - …
- Configuration:
  - /etc/gratia/NAME/ProbeConfig
- Log:
  - /var/log/gratia/DATE.log

Open Science Grid

Gratia

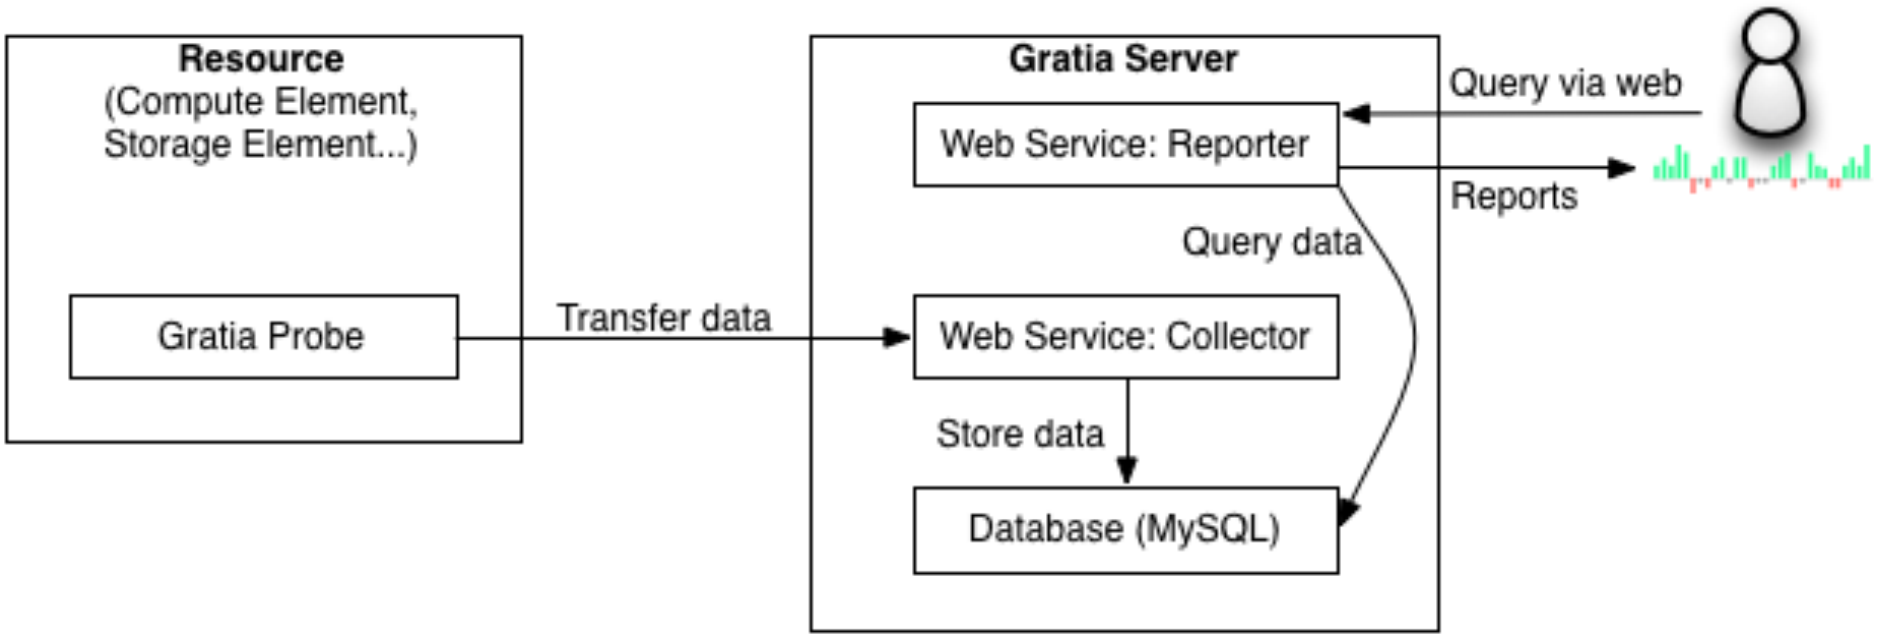

Fetch-crl

- Cron job to periodically update CA Certificate Revocation Lists
- But controlled from init:

```
/sbin/service fetch-crl-boot start
/sbin/service fetch-crl-cron start
```

Open Science Grid

Part 3: Let's do an example install

Install EPEL

```
# rpm -Uvh http://dl.fedoraproject.org/pub/epel/5/i386/
epel-release-5-4.noarch.rpm
Retrieving http://dl.fedoraproject.org/pub/epel/5/i386/
epel-release-5-4.noarch.rpm
Preparing...                ########################### [100%]
   1:epel-release           ########################### [100%]
```

Install Yum priorities

```
# yum install yum-priorities
...
================================================================
Package       Arch   Version          Repository  Size
================================================================
Installing:
yum-priorities noarch 1.1.16-14.el5   sl-base  14 k
...
Is this ok [y/N]: y
...
Installed:
  yum-priorities.noarch 0:1.1.16-14.el5
```

Install OSG repo

```
# rpm -Uvh http://repo.grid.iu.edu/osg-el5-release-latest.rpm
Retrieving http://repo.grid.iu.edu/osg-el5-release-latest.rpm
warning: /var/tmp/rpm-xfer.2wCD1t: Header V3 DSA signature: NOKEY, key ID 824b8603
Preparing...   ########################### [100%]
 1:osg-release ########################### [100%]
```

Install the CA certificates

You get to choose the variety of CA certificates. Here's the default:

```
# yum install osg-ca-certs
================================================================
Package        Arch    Version          Repository   Size
================================================================
Installing:
 osg-ca-certs noarch 1.26-3.osg.el5   osg          281 k
```

Install the Software (CE)

```
# yum install osg-ce-condor
================================================================
Package        Arch    Version           Repository   Size
================================================================
Installing:
osg-ce-condor  x86_64  3.0.0-26          osg          4.5 k

Installing for dependencies:
 CGSI-gSOAP    x86_64  1.3.4.2-3.osg     osg          56 k
 I2util        x86_64  1.1-1             osg          67 k
...
```

Some configuration

```
# ls -1 /etc/osg/config.d/
01-squid.ini
10-misc.ini
10-storage.ini
15-managedfork.ini
20-condor.ini
30-cemon.ini
30-gip.ini
30-gratia.ini
40-localsettings.ini
40-network.ini
40-siteinfo.ini
```

```
# less /etc/lcmaps.db
gumsclient =
  "lcmaps_gums_client.mod"
  "-resourcetype ce"
  "-actiontype execute-now"
...
  "--endpoint https://foo.osg.edu:
8443/gums/services/
GUMSXACMLAuthorizationServicePort"
```

Run osg-configure & start services

```
# osg-configure –v
Configuration verified successfully

# osg-configure –c
Running /usr/sbin/fetch-crl, this process make take
some time to fetch all the crl updates
/usr/sbin/osg-configure completed

# service globus-gatekeeper start
Started globus-gatekeeper[  OK  ]

# service globus-gridftp-server start
Started GridFTP Server[  OK  ]
```

It's ready to go!

```
% voms-proxy-init -voms Engage
Enter GRID pass phrase for this identity:
Your identity: /DC=org/DC=doegrids/OU=People/CN=Alain
Roy 424511
Creating temporary proxy ......................  Done
Contacting  osg-engage.renci.org:15001 [/DC=org/
DC=doegrids/OU=Services/CN=osg-engage.renci.org]
"Engage" Done
Creating proxy ................................ Done

Your proxy is valid until Mon Mar 19 03:07:41 2012

% globus-job-run fermicloud084.fnal.gov /bin/hostname
fermicloud084.fnal.gov
```

Open Science Grid

Q & A

Q: Can I use SL6?

A: For the worker node: yes

Everything else on April 24th (OSG 3.1)

Q: How do I update my system?

A: yum update

Can select specific packages if you wish

Open Science Grid

More Q & A

Q: Can I mix Pacman and RPM on one machine?

A: No

Q: Can I use Pacman on some machines and RPMs on others?

A: Yes

Coming in OSG 3.1: Add back missing functionality

- GUMS
  - In testing/evaluation right now
  - Release at end of March, if it goes well
- Gratia Service
  - In development right now
  - Release in April?

Open Science Grid

Part 4: Where does the software come from?

Principle of Community Packaging

The OSG Software Team should be a good community citizen when it comes to packaging: When possible, we should use packages from existing and/or broader communities; when that is not possible, we should make our own packaging but contribute them back to the broader communities.

Therefore, we should package software only when one of the following is true:

- The software is not already packaged; or
- The software is packaged but needs significant changes to be acceptable to our users. (Different version, extra patches, etc...)

Otherwise we should use the existing packaging provided by external developers or software repositories.

In short: leverage the community

Borrow

The VDT should borrow packages from the larger community whenever possible.

Contribute

The VDT should contribute packages to the larger community whenever possible

Accept Donations

The VDT should make it easy for people to donate packages to the VDT

OSG Software is an amalgamation of repos

- **Refer** to packages in OS/EPEL
- **Copy/patch** packages as needed
- **Create** packages as neede

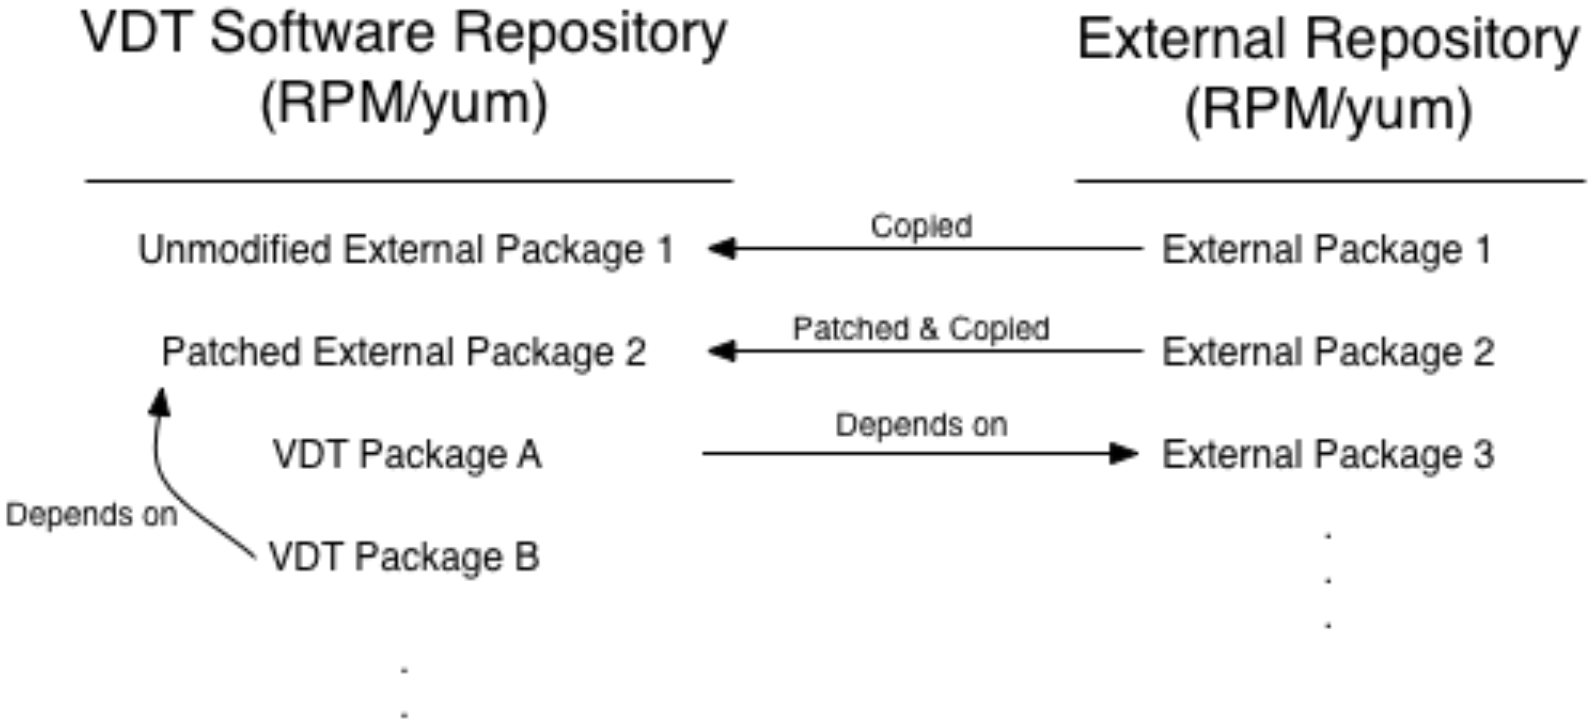

How we build software

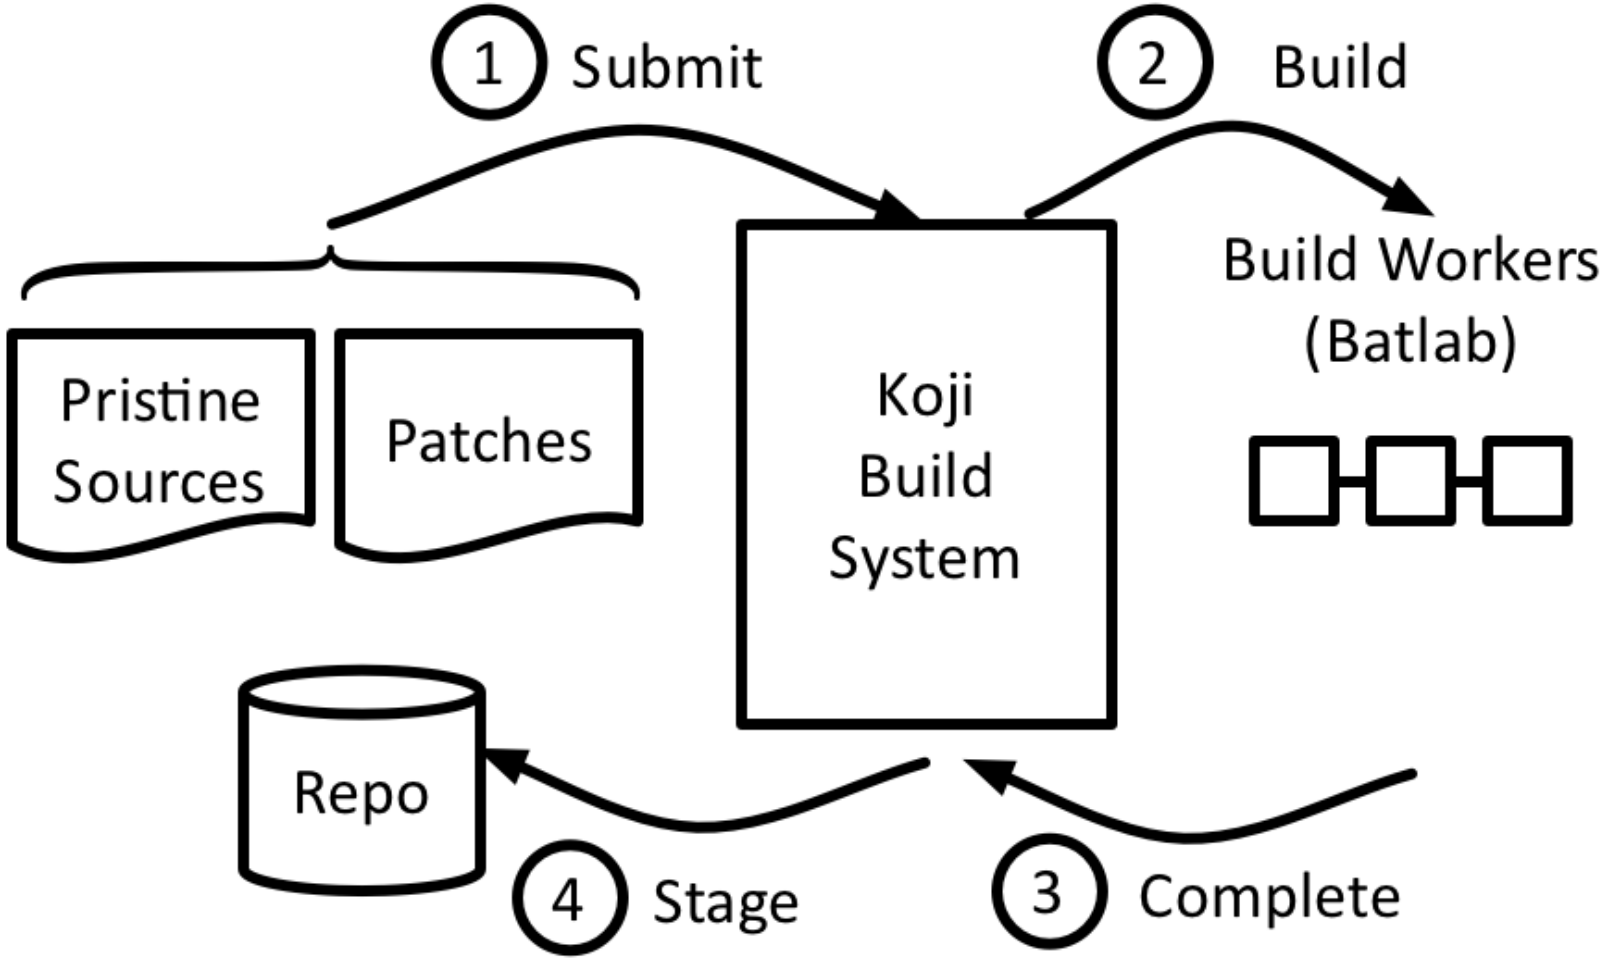

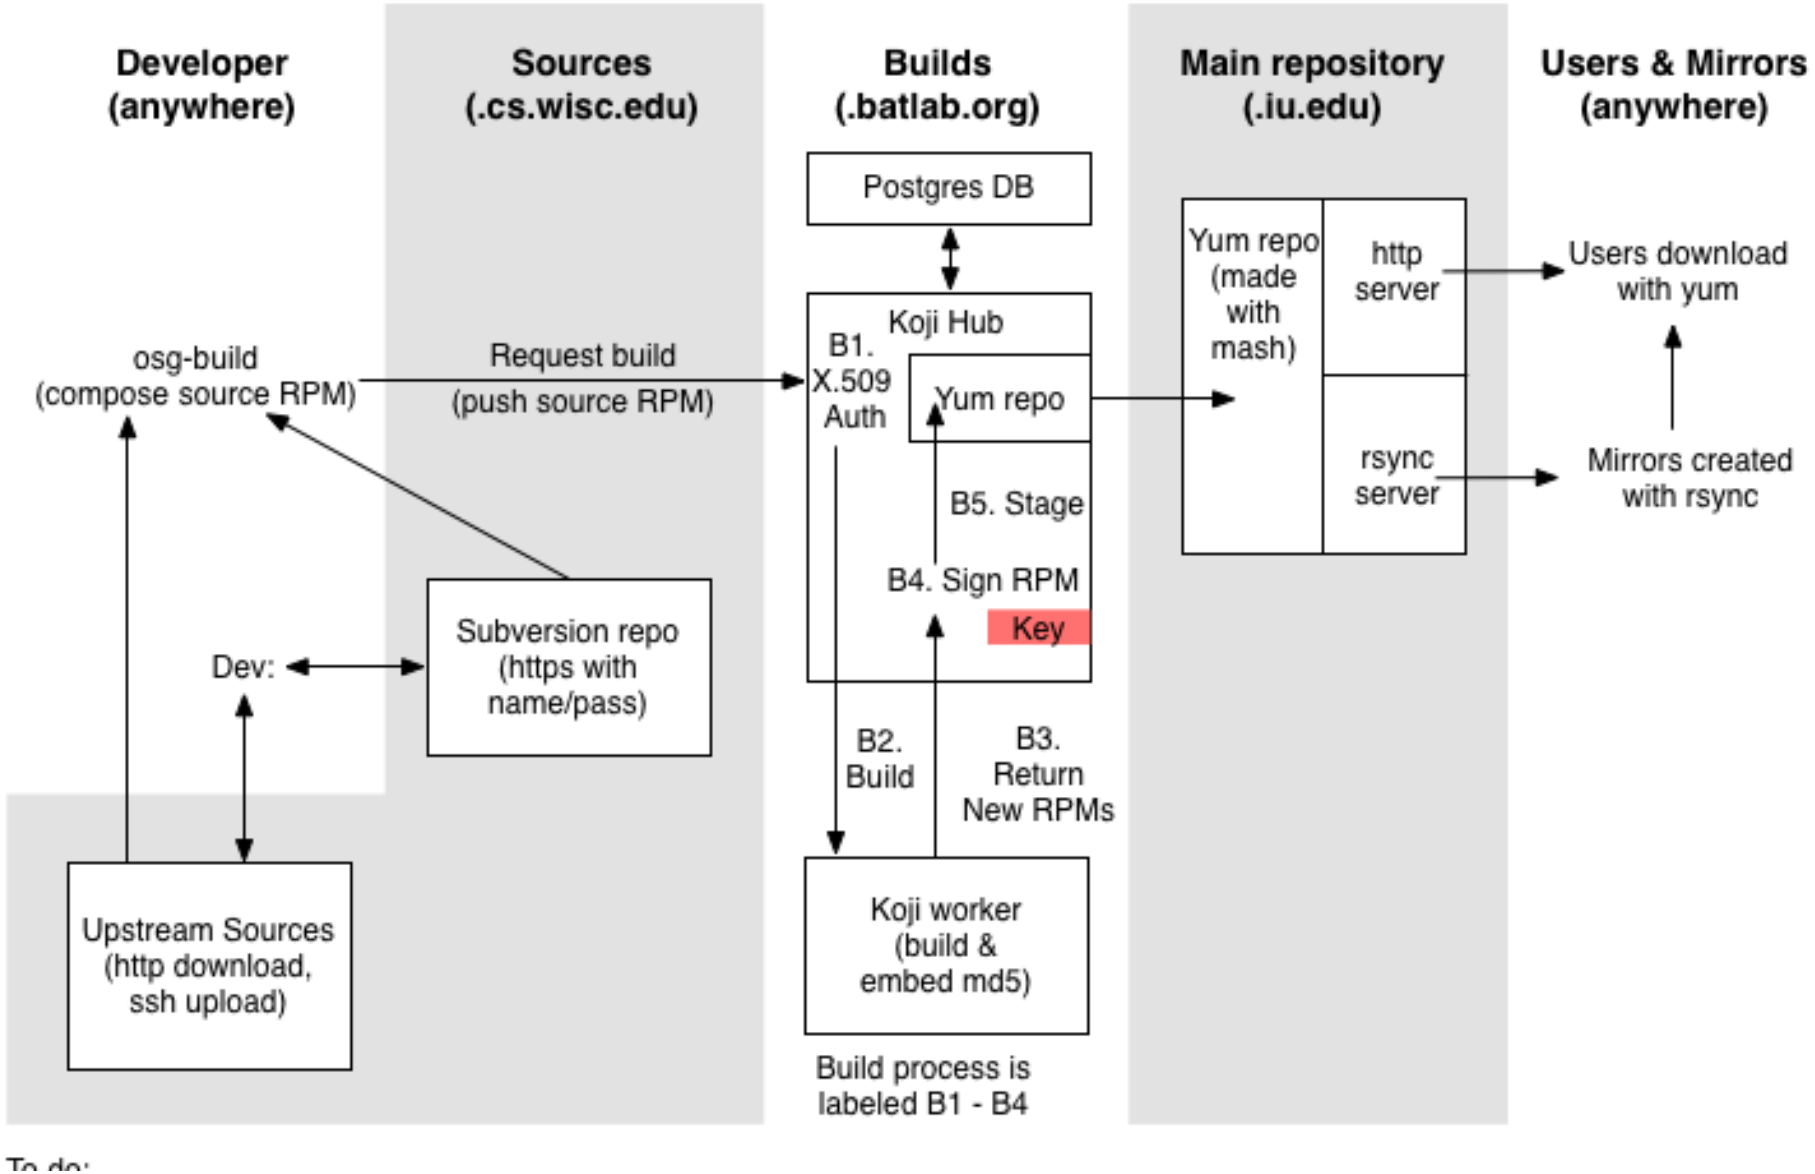
Developer (anywhere)
Sources (.cs.wisc.edu)
Builds (.batlab.org)
Main repository (.iu.edu)
Users & Mirrors (anywhere)
Postgres DB
Koji Hub
B1. X.509 Auth
Yum repo
B5. Stage
B4. Sign RPM
Key
Yum repo (made with mash)
http server
rsync server
Users download with yum
Mirrors created with rsync
osg-build (compose source RPM)
Request build (push source RPM)
Dev:
Subversion repo (https with name/pass)
Upstream Sources (http download, ssh upload)
B2. Build
B3. Return New RPMs
Koji worker (build & embed md5)
Build process is labeled B1 - B4

To do:

Open Science Grid

Questions? Feedback?