



TEXAS TECH UNIVERSITY™
High Performance Computing Center



Texas Tech University (TTU) – Big Tier 3 OSG Site Administrators & CMS Tier 3 workshop

10th August, 2010
ACCRE, Vanderbilt University

Poonam Mane
Graduate Assistant TTU, HPCC

Outline of this talk



- IT infrastructure at TTU
- Introduce people active in this area at TTU
- Describe range of resources available at TTU, and the fraction of those through the TTU high energy physics group for use in our Tier-3 CMS installation
- Identify areas of applicability and potentially useful improvement

Collaboration



- Large no of research groups + small no of support staff (HPCC + TTU IT department)
- Collaboration is the key
- Emphasize on IT infrastructure for research
- Effective even with minimum owned personal resources

TTU High Performance Computing Center



HPCC Staff

- Dr. Philip Smith (Sr. Director)
- Dr. James Abbott (Assoc. Director)
- Jerry Perez
- Srirangam Addepalli
- Dr. Alan Sill (Sr. Scientist)
- Jodi McMurray
- Dr. Per Andersen
- Dr. Ravi Vadapalli
- Huijun Zhu
- Poonam Mane

Purpose

- Facilitate high performance research computing
- Assist with PC to Linux migration and grid-based computing
- Provide consulting and assistance with use of HPC resources

TTU High Energy Physics



Faculty: Dr.Sungwon Lee (Dept. Chair), Dr.Richard Wigmans, Dr.Igor Volobouev, Dr. Nural Akchurin , Dr.Alan Sill (Adjunct Prof.)

Postdoc's: Efe Yazgan, Jordan Damgov

Students in TTU Tier-3 operations : Youn Roh, Chiyoun Jeong, Keng Kovitangoon, Terence Libeiro, Poonam Mane

Students in offline CMS CSP shift : Cemile Bardak

Personnel based at TTU, FNAL, and CERN. Physics analyses carried out using a mix of resources, primarily those provided at the LPC and via CRAB.

TTU T3 is available principally to run CRAB jobs



Resources and capabilities

- Linux clusters for parallel and serial computing
- Large-scale Lustre-based data storage
- High Speed Networking
- Oracle Grid Engine (Sun Grid Engine)
- Local campus grid
- Nationwide, regional and state-wide Globus grid access
- Operational for past decade
- 15% of the total grants to TTU are generated by research groups supported by HPCC

Infrastructure and Support



Infrastructure

- Experimental Sciences Building (Main Campus)
- Reese Center
- 12 miles apart, currently connected with 1Gb network that is being upgraded to 10Gb

Support

- TTU IT department
- HPCC



Antaeus (OSG and CMS Tier-3 resource)

- Production Cluster
- Mixed dual-dual core + dual-quad core Xeons
- 261.4 TB dedicated HEP (CMS) + 6TB user storage interfaced via Lustre


Special nodes

- **Hugin** – basic login node, Frontier Server , test queues for testing physicist
- **Munin** –basic login node, ITB storage element

Testwulf (ITB resource)

- Test Cluster
- 2 worker nodes

Antaeus Cluster



TTU Main Campus - HPCC Primary Linux Clusters



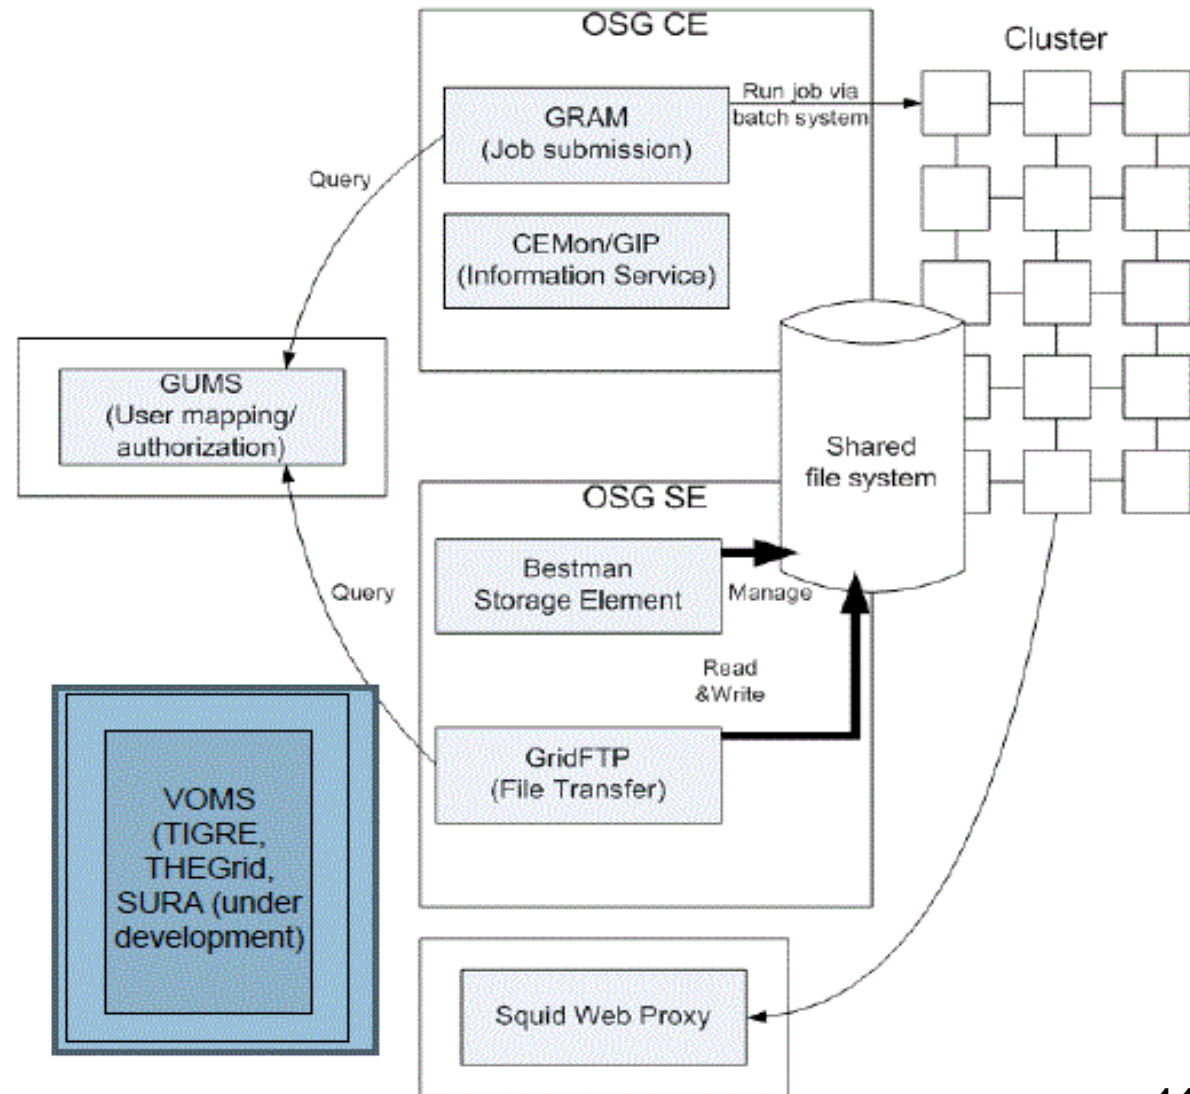
Grendel (primary parallel computing resource)

- Ranked 175 in the November 2009 Top 500 list, achieving 33.5 Teraflop/s max sustained and 40 Teraflop/s peak performance in LinPack test
- 420 dual slot quad core nodes with Intel(R) Xeon(R) CPU E5450 processors for a total of 3360 cores
- Each node has two Intel 5450 Quad Core 64 bit processors on a single board, as an SMP unit. Each node contains 16 GB of memory
- The core frequency is 3.0 GHz
- The core's are connected with DDR Infiniband, with a rating of 40.2 Tflop peak performance
- Interconnect: 4X DDR Non-Blocking InfiniBand
- A 80 TB Lustre file system runs over Infiniband
- HPCC website : <http://www.hpcc.ttu.edu/>

TTU Tier-3 site configuration



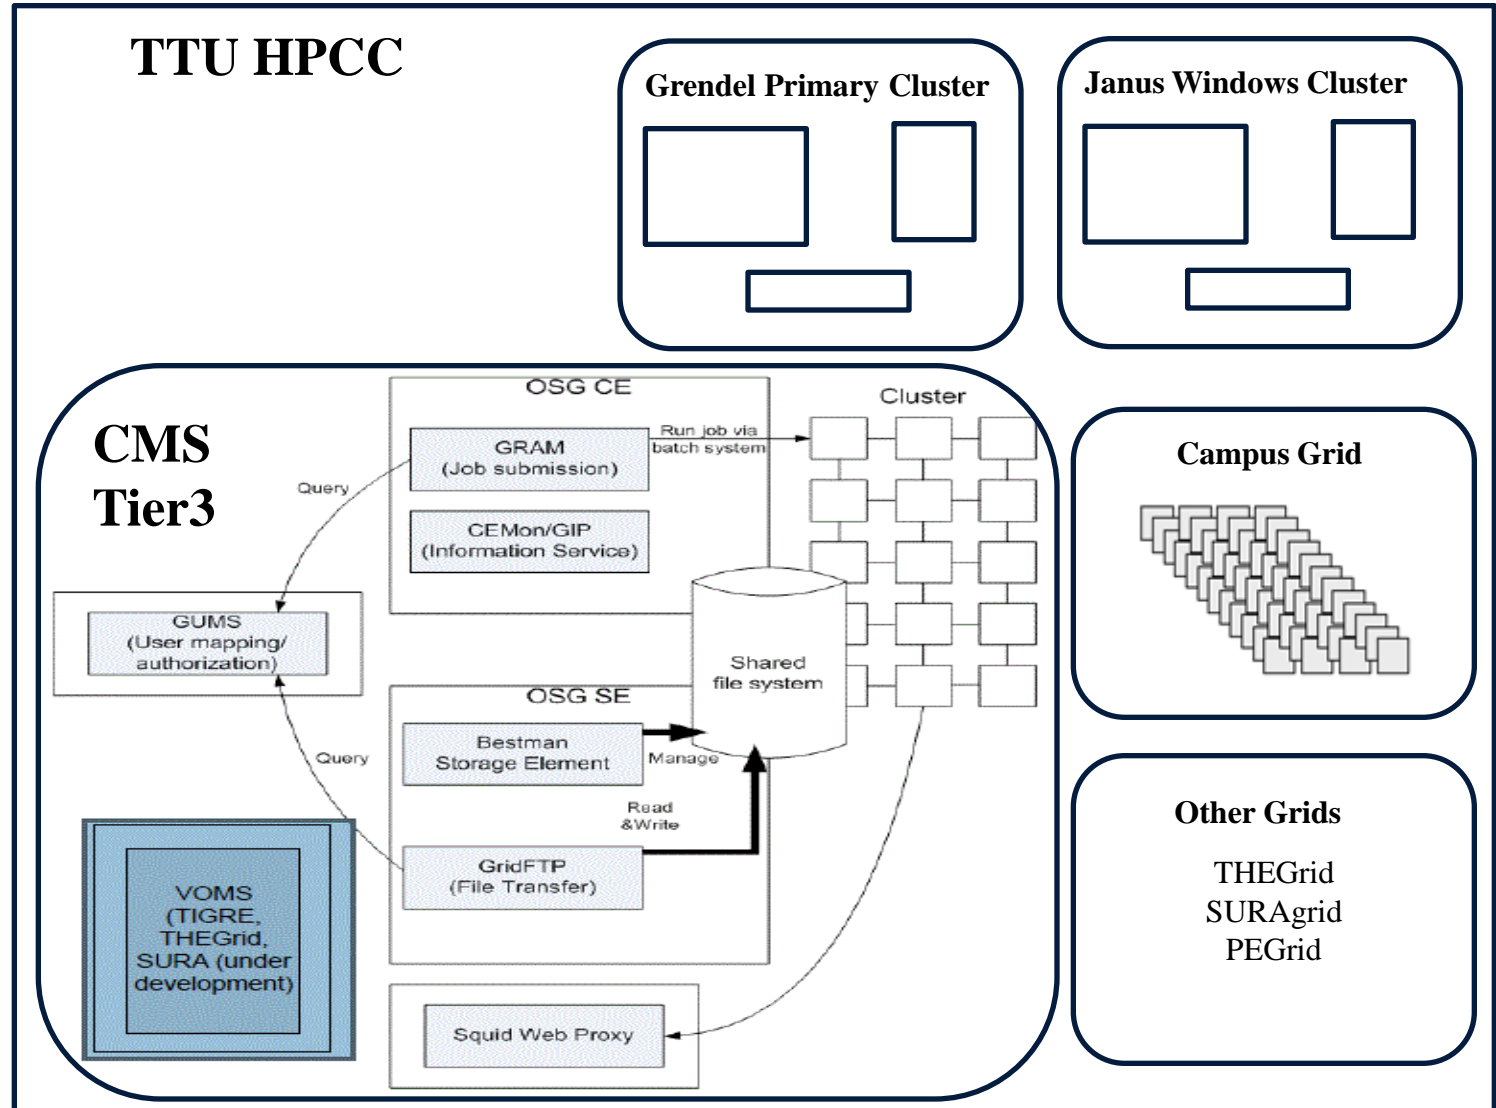
Basically the same as
an OSG medium site,
with additional
services for other VOs



TTU Tier-3 site configuration



Our T3 exists and is supported in the context of other clusters in the TTU High Performance Computing Center



TTU Tier-3 Current Status and Capabilities



CPU power:

- Antaeus - Community Cluster – HEP, Chemistry group , HPCC
- Theoretically total 240 cores,
- HEP – 142 cores (two are down and out of warranty)
 - 128 Xeon cores available to cms queue + 16 cores interactive login (HEP group)
- Chemistry – 64 cores, HPCC – 32 cores

Storage:

- 261.4 TB dedicated HEP disk space, 6 TB general use
- PhEDEx production T3_US_TTU instance
- SRMv2 interface (BeStMan-2.2.1.3.13) to cluster storage

Software & Batch system:

- Rocks 5.3 based on CentOS 5.4
- Oracle Grid Engine (previously known as Sun Grid Engine)



TTU Tier-3 installation methods

CE and worker node base grid software:

- OSG-1.2.11 on CE
- CMS software (CMSSW) is automatically installed
- Monitoring: SAM, Gratia, RSV, MonaLisa and Ganglia

GUMS (Grid Identity Mapping Service):

- OSG supplemented by TIGRE VO

VOMS (Virtual Organization Management Service):

- Used to support TIGRE, THEGrid and SURA, PEGrid
- Prima authentication

FroNtier and squid:

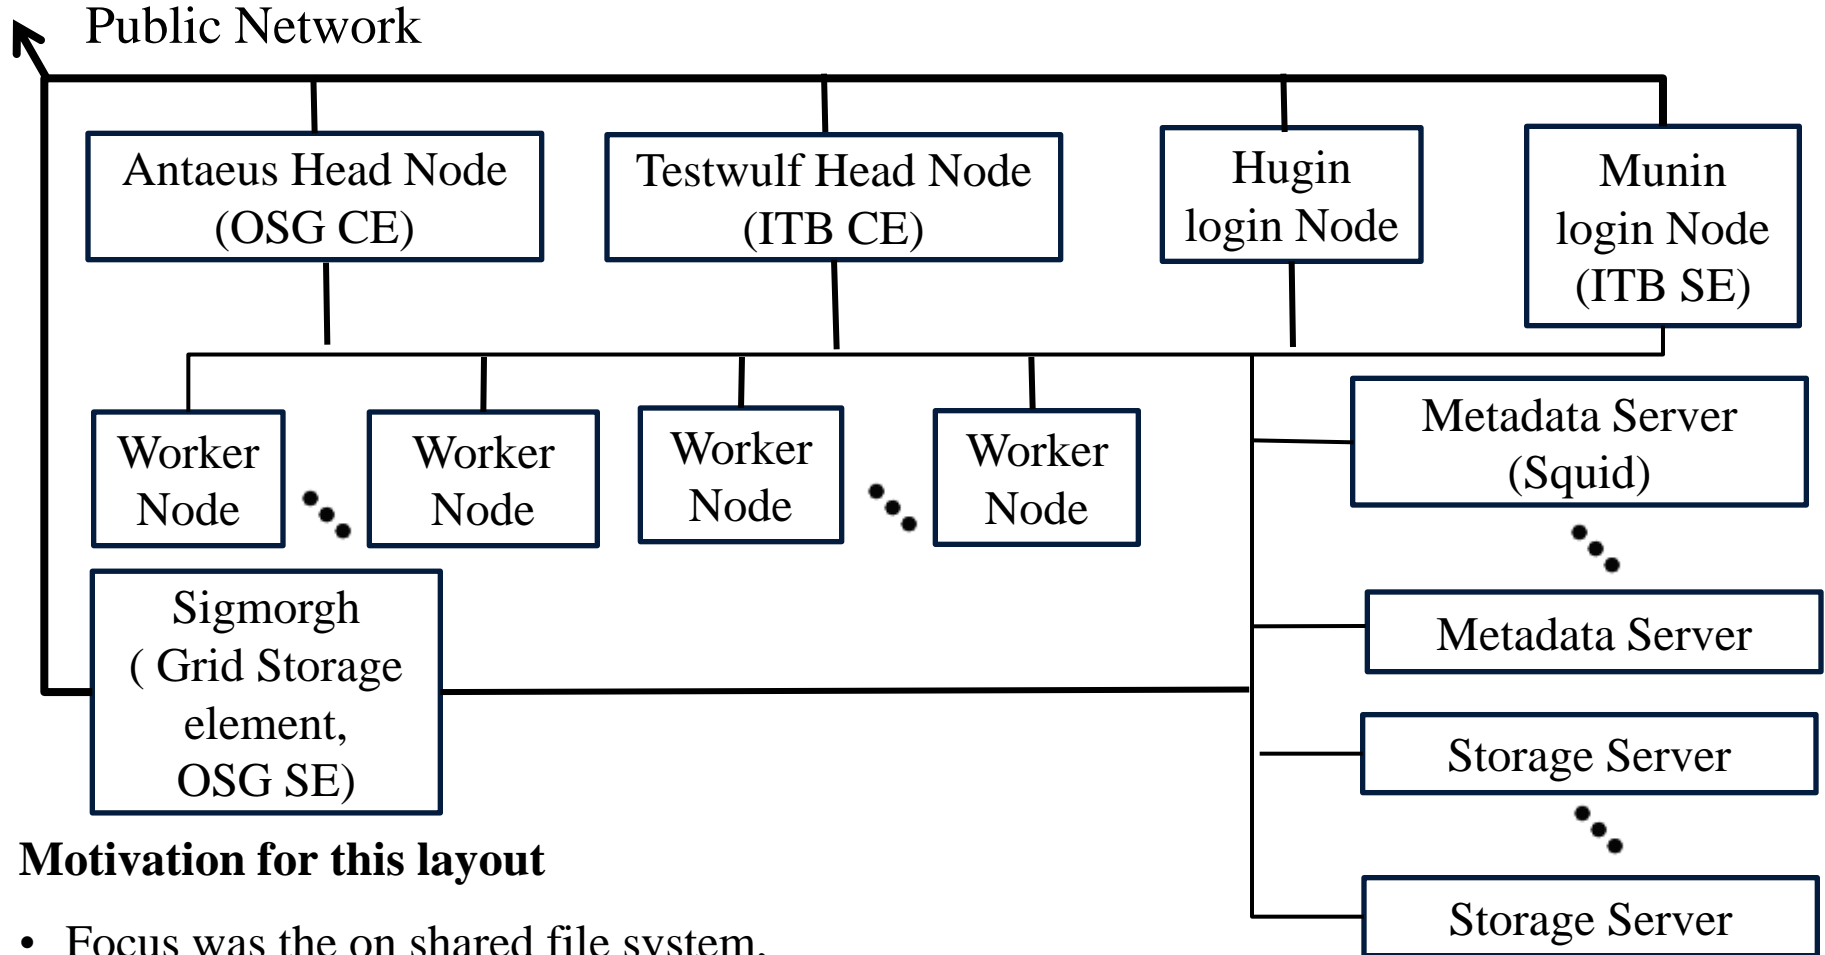
- CMS

Storage (BeStMan):

- BeStMan-2.2.1.3.13



Worker node and storage layout



Motivation for this layout

- Focus was the on shared file system,
- No finances for separate ITB
- Avoided problems caused by over loading of NFS mounts

CMS center



- First CMS center outside Fermi Lab
- 12 monitor, two dual core system
- Official offline CSP CMS shifts
- Count towards requirement of CMS authorship



Way ahead



- Integrate - HEP office desktops with shared file system
- Use the Grendel's idle time to submit jobs as it is 20 times faster
- Virtualization – multiple copies of GUMS, Frontier db
- REDDnet (Research and Education Data Depot network) initiation
- Improve the CMS CSP facility

Current problems and general topics for discussion



- Queue advertisement
 - Control advertisement of queues to include only those available to grid users by VO, so that CMS jobs are not submitted to the wrong queues
 - Assign queues to our user groups over the other CMS groups
 - Blacklist/whitelist parameters in the config.ini script helped
- Grid jobs in general & CMS jobs, do not clean up after execution. What are the best practices ?
- If all the jobs occupy all the queues, there no room for monitoring jobs. Hence, monitoring jobs should run at high priorities
- Pilot jobs have increased from UCSD
 - violation of grid certificate usage
 - certificate is assigned to each pilot job instead of assigning it to a person

