



TEXAS TECH UNIVERSITY
High Performance Computing Center



Texas Tech University (TTU) – Big Tier 3 OSG Site Administrators & CMS Tier 3 workshop

10th August, 2010
ACCRC, Vanderbilt University

Poonam Mane
Graduate Assistant TTU, HPCC

Outline and goals of this talk



- IT infrastructure
- Introduce people active in this area at TTU
- Describe range of resources available at TTU, and the fraction of those through the TTU high energy physics group for use in our Tier-3 CMS installation
- Seek to identify areas of applicability and potentially useful improvement

HPCC and IT infrastructure



- Research group + IT
- Collaboration is the key
- Emphasize on IT infrastructure for research
- Can be effective even with minimum owned personal resources
- IT department at TTU supports the user groups



HPCC Staff

- Philip Smith (Sr. Director)
- James Abbott (Assoc. Director)
- Jerry Perez
- Srirangam Addepalli
- Alan Sill (Sr. Scientist)
- Dr. Per Andersen
- Jodi McMurray
- Huijun Zhu
- Dr. Ravi Vadapalli
- Poonam Mane

Purpose

- Facilitate high performance research computing
- Assist with PC to Linux migration and grid-based computing
- Provide consulting and assistance with use of HPC resources

TTU High Energy Physics



Faculty: Nural Akchurin (Dept. Chair), Richard Wigmans, Igor Volobouev, Sungwon Lee, Alan Sill (Adjunct Prof.)

Postdoc's: Efe Yazgan, Jordan Damgov

Students in TTU Tier-3 operations : Youn Roh, Chiyoun Jeong, Keng Kovitangoon, Terence Libeiro, Poonam Mane

Students in offline CMS CSP shift : Cemile Bardak

Personnel based at TTU, FNAL, and CERN. Physics analyses carried out using a mix of resources, primarily those provided at the LPC and via CRAB.

TTU T3 is available principally to run CRAB jobs



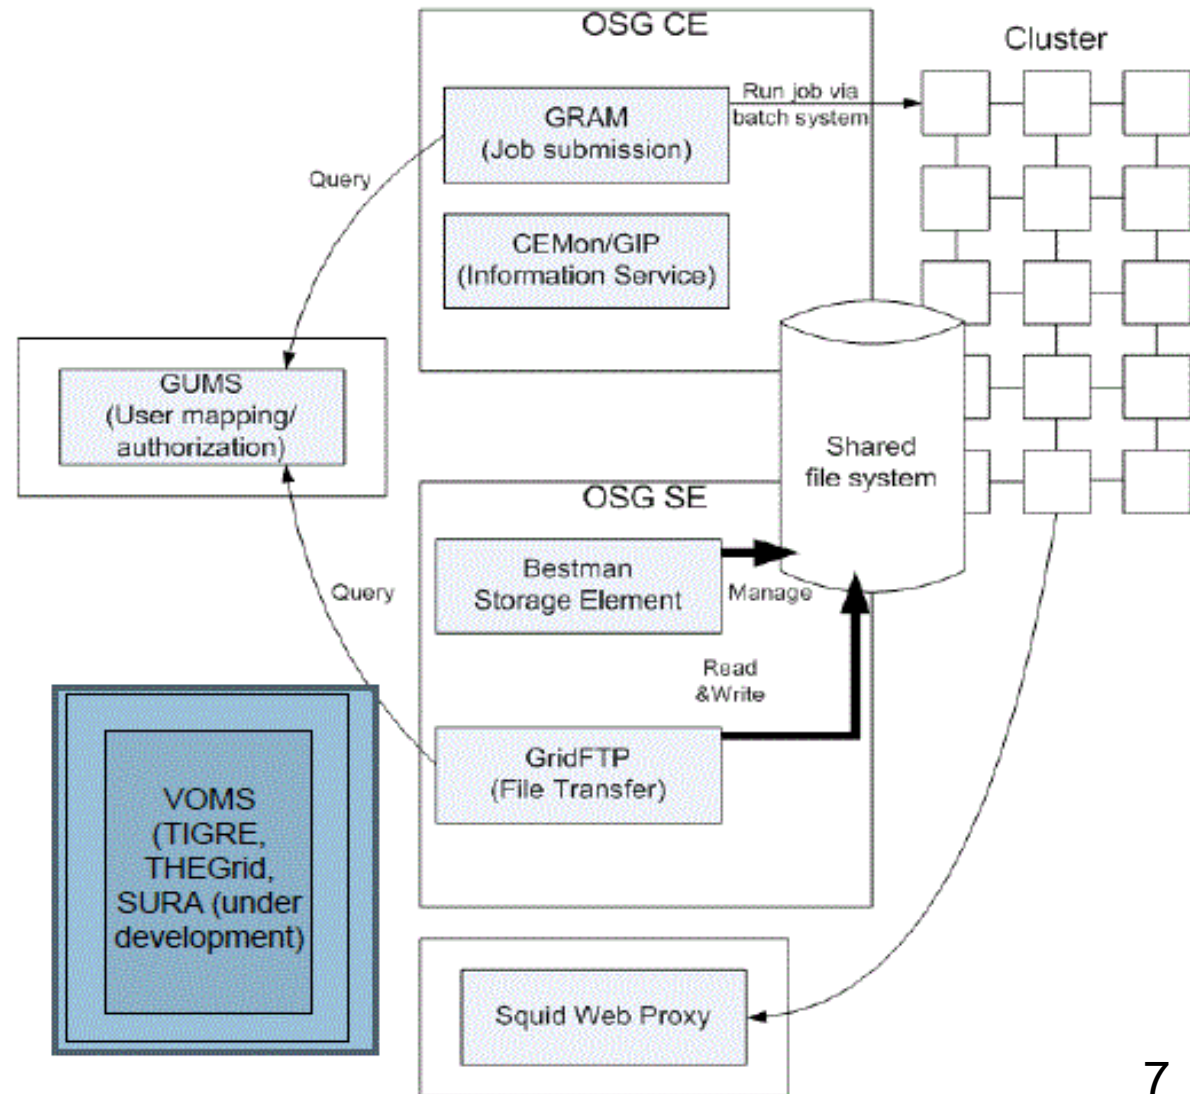
Resources and capabilities

- Linux clusters for parallel and serial computing
- Large-scale Lustre-based data storage
- High Speed Networking
- Oracle Grid Engine (Sun Grid Engine)
- Local campus grid
- Nationwide, regional and state-wide Globus grid access
- Operation for past decade
- 15% of the total grants to TTU are generated by research groups supported by HPCC

TTU Tier-3 site configuration



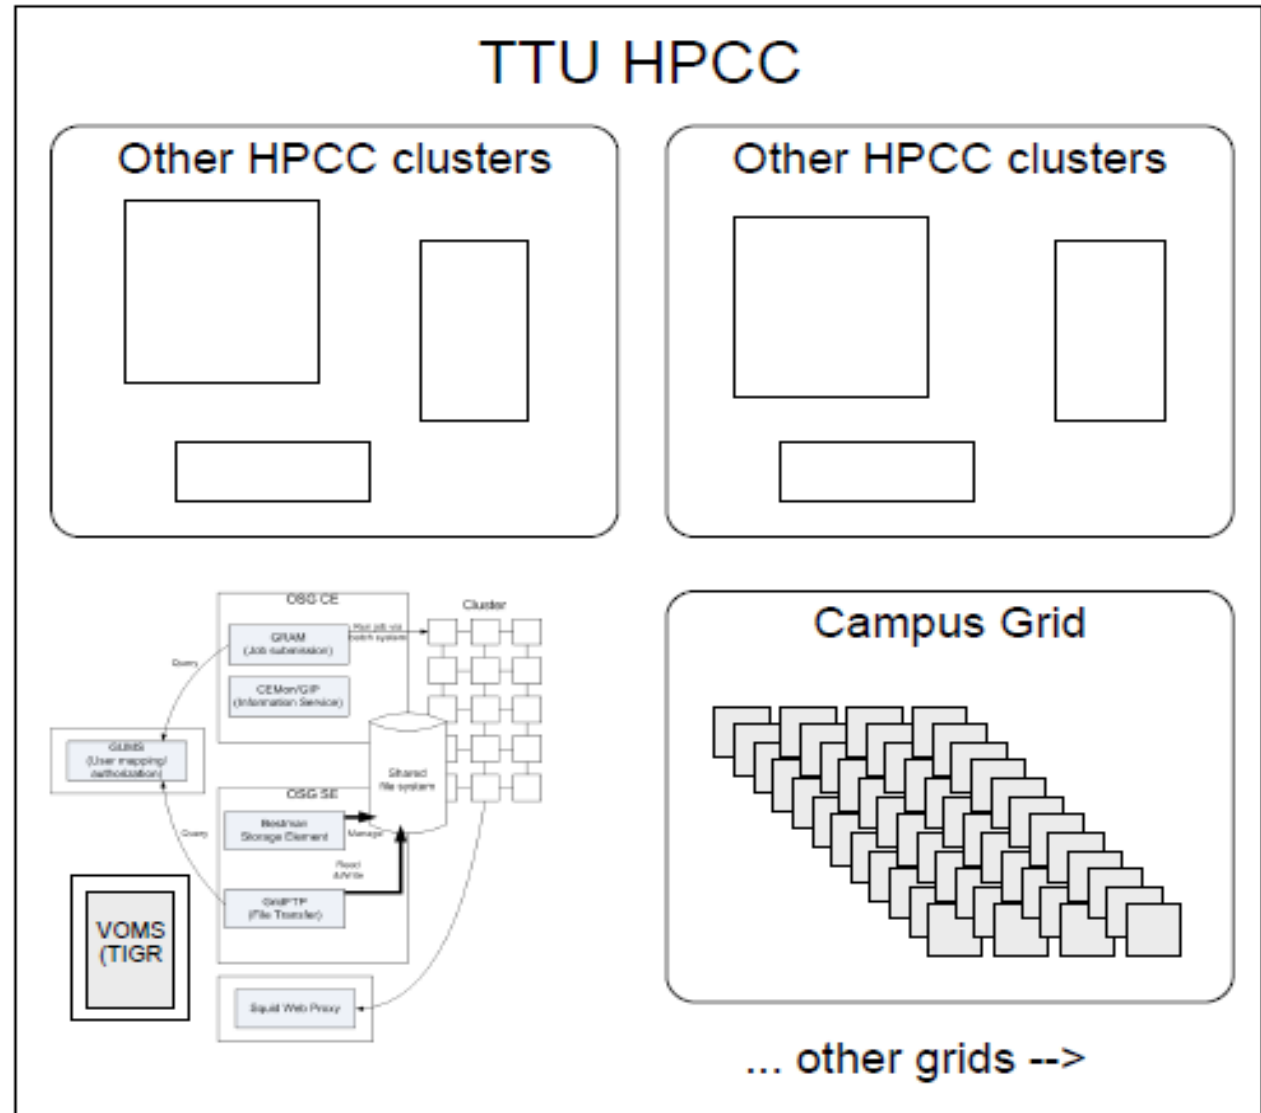
Basically the same as
an OSG medium site,
with additional
services for other VOs





TTU Tier-3 site configuration

Our T3 exists and is supported in the context of other clusters in the TTU High Performance Computing Center



TTU Tier-3 Current Status and Capabilities



CPU power:

- 128 Xeon core available to cms queue + 16 core interactive login (HEP group)
- Idle time access to remaining additional 96 cores

Storage:

- 261.4 TB dedicated HEP disk space, 6 TB general use
- PhEDEx production T3_US_TTU instance
- SRMv2 interface (BeStMan) to cluster storage

Batch system:

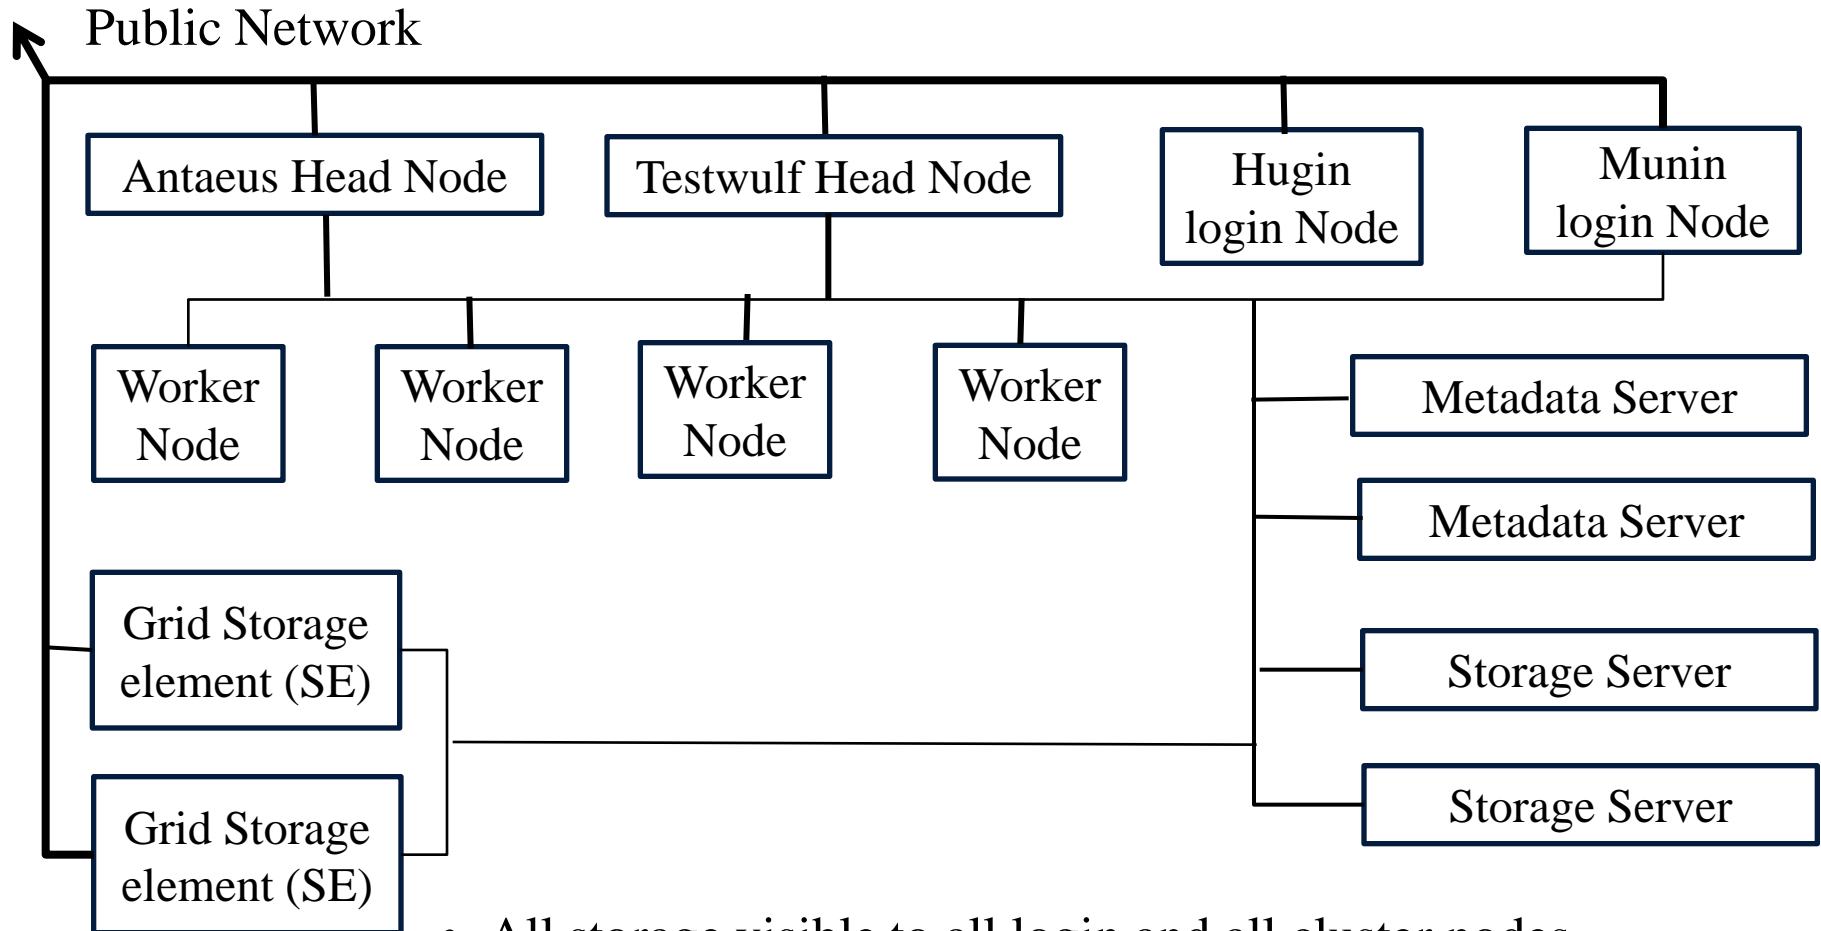
- SGE

Other resources:

- Idle-time access to campus cycle-scavenging grid and other clusters
- ITB resource (TTU-Testwulf), GUMS, and TIGRE VOMS server



Worker node and storage layout



- All storage visible to all login and all cluster nodes
- Public network access through grid storage elements
- Logical partition of storage servers adjustable via Lustre



Grendel (primary parallel computing resource)

- Ranked 175 in the November 2009 Top 500 list
- 420 dual slot quad core nodes with Intel(R) Xeon(R) CPU E5450 processors for a total of 3360 cores
- Each node has two Intel 5450 Quad Core 64 bit processors on a single board, as an SMP unit. Each node contains 16 GB of memory
- The core frequency is 3.0 GHz
- The core's are connected with DDR Infiniband, with a rating of 40.2 Tflop peak performance
- Interconnect: 4X DDR Non-Blocking InfiniBand
- A 80 TB Lustre file system runs over Infiniband



Antaeus (OSG and CMS Tier-3 resource)


- Mixed dual-dual core + dual-quad core Xeons, 240 cores total, 2 GB memory per core, Gb backplane. (Some owned by other researchers)
- 261.4 TB dedicated HEP (CMS) + 6TB user storage interfaced via Lustre

Hugin – special node , basic login node, Frontier Server

Munin – special node, basic login node, 1TB storage element

Network between Main Campus and Reese is upgraded to 10Gb

Antaeus Cluster



CMS center



- 12 monitor, two dual core system
- Periodic Shifts according to US times to support on global basis
- Official offline CSP CMS shifts
- Count towards requirement of CMS authorship



Way ahead



- Integrate - HEP office desktops with shared file system
- Use the Grendel's idle time to submit jobs as it is 20 times faster
- Virtualization – multiple copies of GUMS, Frontier db
- REDDnet (Research and Education Data Depot network) initiation
- Improve the CMS CPS facility

Current problems and general topics for discussion



- Queue advertisement
 - Control advertisement of queues to include only those available to grid users by VO, so that CMS jobs are not submitted to the wrong queues
 - Blacklist/whitelist parameters in the config.ini script helped
- Grid jobs in general & CMS jobs, do not clean up after execution. What are the best practices ?
- If all the jobs occupy all the queues, there no room for monitoring jobs. Hence, monitoring jobs should run at high priorities
- Pilot jobs have increased from UCSD
 - violation of grid certificate usage
 - certificate is assigned to each pilot job instead of assigning it to a person

