



August 9th 2011, OSG Site Admin Workshop
Jason Zurawski – Internet2 Research Liaison

BWCTL

Agenda

- Tutorial Agenda:
 - Network Performance Primer - Why Should We Care? (**30 Mins**)
 - Introduction to Measurement Tools (**20 Mins**)
 - Use of NTP for network measurements (**15 Mins**)
 - **Use of the BWCTL Server and Client (25 Mins)**
 - Use of the OWAMP Server and Client (**25 Mins**)
 - Use of the NDT Server and Client (**25 Mins**)
 - perfSONAR Topics (**30 Mins**)
 - Diagnostics vs Regular Monitoring (**20 Mins**)
 - Use Cases (**30 Mins**)
 - Exercises

BWCTL – What is it?


- BWCTL is:
 - A command line client application
 - A scheduling and policy daemon
 - Wraps the throughput testing tools [Iperf](#) and [Nuttcp](#).
- These tests are able to measure:
 - Maximum TCP bandwidth (with various tuning options available)
 - The delay, jitter, and datagram loss of a network when doing a UDP test

Problem Statement

- Users want to verify available bandwidth/throughput:
 - Between their site and a remote resource
 - Between two remote resources
 - Validate/Verify an SLA
- Methodology:
 - Verify available bandwidth from each endpoint to points in the middle
 - Determine problem area(s)
 - Re-run tests over time – requires access to tool instead of doing a ‘one off’ test

Throughput? Bandwidth?

- The term “throughput” is vague
 - Capacity: link speed
 - Narrow Link: link with the lowest capacity along a path
 - Capacity of the end-to-end path = capacity of the narrow link
 - Utilized bandwidth: current traffic load
 - Available bandwidth: capacity – utilized bandwidth
 - Tight Link: link with the least available bandwidth in a path
 - Achievable bandwidth: includes protocol and host issues



(Shaded portion shows background traffic)

INTERNET
2

Typical Solution

- Run “iperf” or similar tool on two endpoints and hosts on intermediate paths
 - Roadblocks:
 - Need software on all test systems
 - Need permissions on all systems involved (usually full shell accounts*)
 - Need to coordinate testing with others *
 - Need to run software on both sides with specified test parameters *
- Desirable features for an alternate method
 - ‘Daemon’ to run in the background
 - Protocol to exchange results/errors
 - Works with firewalls
 - Protect resources
- (* BWCTL was designed to help with these)

Implementation

- Applications
 - Daemon (bwctld)
 - Client (bwctl)
- Open Source License & Development
 - Modified BSD (<http://www.internet2.edu/membership/ip.html>)
 - Mailing lists for developer communication – come join us!
- Protocol Abstraction Library
 - Will support development of new clients
 - Add custom ‘hooks’ into the policy (e.g. add authentication via OpenID or similar)

Server Functionality (bwctld)

- bwctld on each test host
 - Accepts requests for “iperf” tests including time slot and parameters for test
 - Responds with a tentative reservation or a denied message
 - Reservations by a client must be confirmed with a “start session” message
 - Acts as the “Resource Broker”
 - Runs the test
 - Both “sides” of test get results

Client Functionality (bwctl)

- bwctl client application makes requests to both endpoints of a test
 - Communication can be “open”, “authenticated”, or “encrypted” (encrypted reserved for future use)
 - Requests include a request for a **time slot** as well as a full parameterization of the test
 - “Third party” requests – run a test on two distributed hosts
 - If no server is available on the localhost, client handles test endpoint
 - *Mostly* the same command line options as testers (e.g. iperf, nuttcp – read the help or man pages to be sure...)

TCP Measurements

- Measures TCP Achievable Bandwidth
 - Measurement includes the end system
 - Sometimes called “memory-to-memory” tests
 - Set expectations for well coded application
- Limits of what we can measure
 - TCP *hides* details
 - In hiding the details it can obscure what is causing errors
- Many things can limit TCP throughput
 - Loss
 - Congestion
 - Buffer Starvation
 - Out of order delivery

TCP – Quick Overview

- Data Packet
 - Contains some header overhead, and the broken up chunk of user data
- ACK Packet
 - Acknowledge the receipt of a data packet, “cumulative” in nature
- SACK Packet
 - Selective acknowledgement for a specific missing segment
- MSS
 - Maximum segment size (largest size of packets on a given network segment)
- Congestion Control
 - Process of self regulating flow speed due to loss in the network (e.g. making it fair)
- Slow Start
 - Avoid sending more data than the network is capable of consuming. Goal is to reach a loss (establishes window size by relying on acks)

TCP – Quick Overview

- Congestion Avoidance
 - Additive-increase/Multiplicative-decrease [AIMD] scheme to find a fair speed for a TCP flow by adjusting the sending window. Starts low ($2 \times \text{MSS}$) and increase
- Fast Retransmit
 - Retransmit a single segment after receiving duplicate ACKs for the prior numbered segment
- Fast Recovery
 - Re-send packets in a window after a timeout
- Bandwidth Delay Product
 - The amount of “in flight” data allowed for a TCP connection
 - $\text{BDP} = \text{bandwidth} * \text{round trip time}$
 - Example: 1Gb/s cross country, $\sim 100\text{ms}$
 - $1,000,000,000 \text{ b/s} * .1 \text{ s} = 100,000,000 \text{ bits}$
 - $100,000,000 / 8 = 12,500,000 \text{ bytes}$
 - $12,500,000 \text{ bytes} / (1024 * 1024) \sim 12\text{MB}$


TCP – Quick Overview

- Congestion Control Algorithms (selectable in the Linux Kernel)
 - RENO (Slow Start, Cong. Avoidance, Fast Retransmit, Fast Recovery)
 - Cubic (Optimized for LFN [Long Fat Networks] with large latency, Cubic growth pattern)
 - HTCP (still additive-increase/multiplicative-decrease [AIMD], more aggressive as loss decreases on high BDP paths)

TCP – Quick Overview

- General Operational Pattern
 - Sender buffers up data to send into segments (respect the MSS) and numbers each
 - The ‘window’ is established and packets are sent in order from the window
 - The flow of data and ACK packets will dictate the overall speed of TCP for the length of the transfer


TCP – Quick Overview (Typical Sawtooth)



TCP – Quick Overview

- General Operational Pattern – cont
 - TCP starts slow, until it can establish the available resources on the network.
 - The idea is to grow the window until a loss is observed
 - This is the signal to the algorithm that it must limit the window for the time being, it can slowly build it back up


TCP – Quick Overview (Slow Start)



TCP – Quick Overview

- General Operational Pattern – cont
 - Receiver will acknowledge packets as they arrive
 - ACK Each (old style)
 - Cumulative ACK (“I have seen everything up to this segment”)
 - Selective ACK (sent to combat a complete retransmit of the window)
 - TCP relies on loss to a certain extent – it will adjust it’s behavior after each loss
 - Congestive (e.g. reaching network limitation, or due to traffic)
 - Non-congestive (due to actual problems in the network)
 - Congestion avoidance stage follows slow start, window will remain a certain size and data rates will increase/decrease based on loss in the network
 - Congestion Control algorithms modify the behavior over time
 - Control how large the window may grow
 - Control how fast to recover from any loss

TCP – Quick Overview (Cong. Avoidance)



TCP Performance: Window Size

- Use TCP auto tuning if possible
 - Linux 2.6, Mac OS X 10.5, FreeBSD 7.x, and Windows Vista
 - Allow the OS to decide how large the window needs to be based on current resources and performance
- The `-w` option can be used to request a particular buffer size.
 - Use this if your OS doesn't have TCP auto tuning
 - This sets both send and receive buffer size.
 - The OS may need to be tweaked to allow buffers of sufficient size.
 - See <http://fasterdata.es.net/fasterdata/host-tuning/> for more details
- Parallel transfers may help as well, the `-P` option can be used for this
- To get full TCP performance the TCP window needs to be large enough to accommodate the Bandwidth Delay Product



TCP Parallel Streams

- Parallel streams can help in some situations
- TCP attempts to be “fair” and conservative
 - Sensitive to loss, but more streams hedge bet
 - Circumventing fairness mechanism
 - 1 bwctl stream vs. n background: bwctl gets $1/(n+1)$
 - X bwctl streams vs. n background: bwctl gets $x/(n+x)$
 - Example: 2 background, 1 bwctl stream: $1/3 = 33\%$
 - Example: 2 background, 8 bwctl streams: $8/10 = 80\%$
- How?
 - The `-P` option sets the number of streams/threads to use
 - There is a point of diminishing returns

TCP Performance: Read/Write Buffer Size

- TCP breaks the stream into pieces transparently
- Longer writes often improve performance
 - Let TCP “do it’s thing”
 - Fewer system calls
- How?
 - -l <size> (lower case ell)
 - Example -l 128K
- UDP doesn’t break up writes, don’t exceed Path MTU

UDP Measurements

- UDP provides greater transparency
- We can directly measure some things TCP hides
 - Loss
 - Jitter
 - Out of order delivery
- Use -b to specify target bandwidth
 - Default is 1M
 - Two sets of multipliers
 - k, m, g multipliers are 1000, 1000², 1000³
 - K, M, G multipliers are 1024, 1024², 1024³
 - Eg, -b 1m is 1,000,000 bits per second

Example

```
boote@nms-rthr2:~  
[boote@nms-rthr2 ~]$ bwctl -x -s bwctl.kans.net.internet2.edu  
bwctl: 19 seconds until test results available  
  
RECEIVER START  
3421251446.646488: iperf -B 2001:468:9:100::16:22 -P 1 -s -f b -m -p 5001 -t 10 -V  
-----  
Server listening on TCP port 5001  
Binding to local address 2001:468:9:100::16:22  
TCP window size: 87380 Byte (default)  
-----  
[ 14] local 2001:468:9:100::16:22 port 5001 connected with 2001:468:4:100::16:214 port 5001  
[ 14] 0.0-10.2 sec 1193058304 Bytes 939913512 bits/sec  
[ 14] MSS size 8928 bytes (MTU 8968 bytes, unknown interface)  
  
RECEIVER END  
  
SENDER START  
3421251448.787198: iperf -c 2001:468:9:100::16:22 -B 2001:468:4:100::16:214 -f b -m -p 5001 -t 10 -V  
-----  
Client connecting to 2001:468:9:100::16:22, TCP port 5001  
Binding to local address 2001:468:4:100::16:214  
TCP window size: 87380 Byte (default)  
-----  
[ 7] local 2001:468:4:100::16:214 port 5001 connected with 2001:468:9:100::16:22 port 5001  
[ 7] 0.0-10.0 sec 1193058304 Bytes 951107779 bits/sec  
[ 7] MSS size 8928 bytes (MTU 8968 bytes, unknown interface)  
  
SENDER END  
[boote@nms-rthr2 ~]$
```


BWCTL GUIs

performance **ps** toolkit

User Tools

- Local Performance Services
- Global Performance Services
- Java OWAMP Client
- Reverse Traceroute
- Reverse Ping
- PingER Web GUI

Service Graphs

- Throughput
- One-Way Latency
- Ping Latency
- SNMP Utilization
- Cacti Graphs

Toolkit Administration

- Administrative Information
- External BWCTL Limits
- External OWAMP Limits
- Enabled Services
- NTP
- Scheduled Tests
- Cacti SNMP Monitoring

Performance Toolkit

- Configuration Help
- Frequently Asked Questions
- About
- Credits

pS-Performance Node - Throughput Tests

https://desk172.internet2.edu/toolkit/gui/perfAdmin/serviceTest.cgi?url=http://localhost:8085/perfSONAR_PS/services/pSB&ev Google

Throughput Tests

Active Data Sets											
First Host	First Address	Second Host	Second Address	Protocol	Duration	Window Size	Bandwidth Limit	Bi-Directional	Line Graph	Scatter Graph	
bwctl.ucsc.edu	128.114.0.205	desk172.internet2.edu	207.75.164.172	TCP	20			Yes	-- Select --	-- Select --	
desk172.internet2.edu	207.75.164.172	infotech-sv-62.ggnnet.umn.edu	146.57.255.17	TCP	20			Yes	-- Select --	-- Select --	
desk172.internet2.edu	207.75.164.172	iperf.its.vanderbilt.edu	192.111.110.34	TCP	20			Yes	-- Select --	-- Select --	
desk172.internet2.edu	207.75.164.172	lab253.internet2.edu	207.75.164.253	TCP	20			Yes	-- Select --	-- Select --	
desk172.internet2.edu	207.75.164.172	ndt.ScrippsCollege.edu	134.173.151.207	TCP	20			Yes	-- Select --	-- Select --	
desk172.internet2.edu	207.75.164.172	perfsonar.its.iastate.edu	129.186.6.241	TCP	20			Yes	-- Select --	-- Select --	
desk172.internet2.edu	207.75.164.172	perfsonar.ndsu.NoDak.edu	134.129.90.1	TCP	20			Yes	-- Select --	-- Select --	


1 Week Average Bandwidth in Mbps

Host	In BW (Mbps)	Out BW (Mbps)
bwctl.ucsc.edu	~70	~45
desk172.internet2.edu	~75	~85
infotech-sv-62.ggnnet.umn.edu	~90	~85
iperf.its.vanderbilt.edu	~90	~70
lab253.internet2.edu	~90	~90
ndt.ScrippsCollege.edu	~85	~55
perfsonar.its.iastate.edu	~90	~85
perfsonar.ndsu.NoDak.edu	~90	~85

Non-Active Data Sets

First Host	First Address	Second Host	Second Address	Protocol	Duration	Window Size	Bandwidth Limit	Bi-Directional	Line Graph	Scatter Graph
------------	---------------	-------------	----------------	----------	----------	-------------	-----------------	----------------	------------	---------------

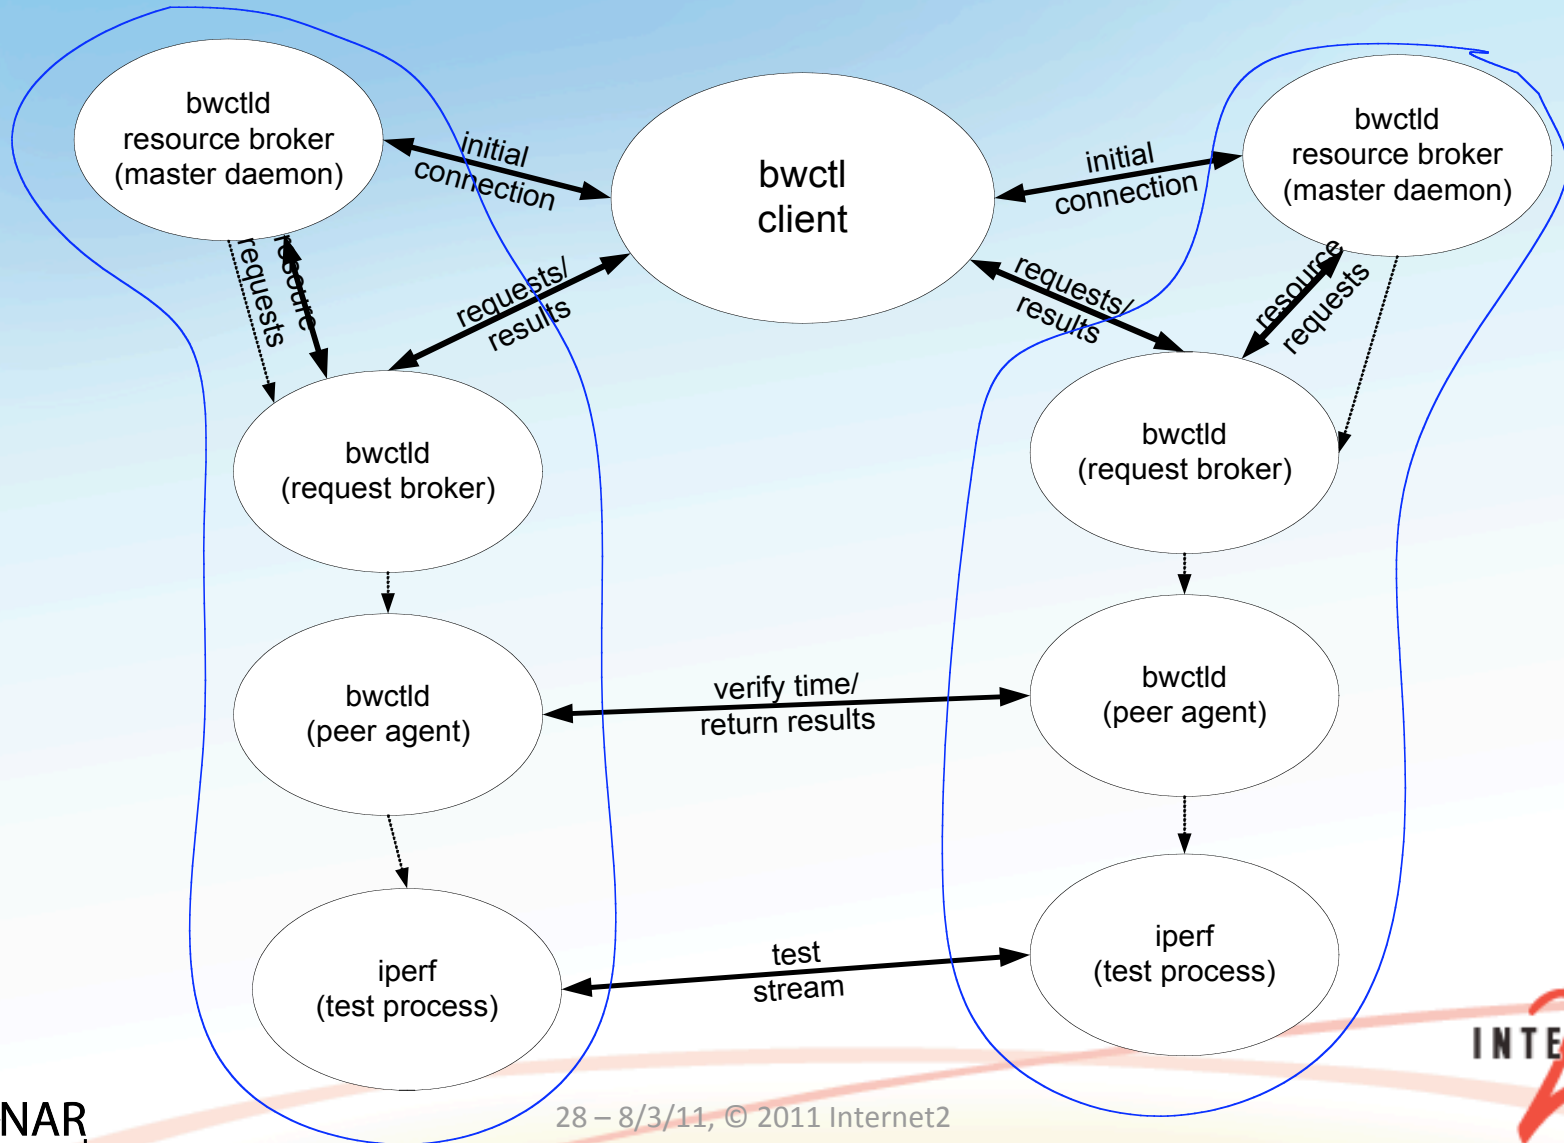
BWCTL GUIs



Resource Allocation

- Each connection is “classified” (authentication)
- Each classification is hierarchical and has an associated set of hierarchical limits:
 - Connection policy (allow_open_mode)
 - Bandwidth (allow_tcp,allow_udp,bandwidth)
 - Scheduling (duration,event_horizon,pending)
 - A time slot is simply a time-dependent resource that needs to be allocated just like any other resource. It therefore follows the resource allocation model.

3rd Party Testing



General Requirements

- Iperf version 2.0.x
- NTP (ntpd) synchronized clock on the local system
 - Used for scheduling
 - More important that errors are accurate than the clock itself
- Firewalls:
 - Lots of ports for communication and testing – see the web for specifics
- End hosts must be tuned!
 - <http://fasterdata.es.net/fasterdata/host-tuning>
 - http://www.psc.edu/networking/perf_tune.html

Supported Systems

- Source Code
 - All modern Unix distributions (Free BSD/Linux)
 - OS X
- Packages
 - Support for CentOS 5.5 (x86)
 - Packages have been shown to operate on similar systems (CentOS, Fedora, RHEL, and x86_64 architecture)

Security Considerations

- DoS source
 - Imagine a large number of compromised BWCTLD servers being used to direct traffic
- DoS target
 - Someone might attempt to affect statistics web pages to see how much impact they can have
- Resource consumption
 - Time slots
 - Network bandwidth

Policy Approaches

- Restrictive for UDP
 - Allow between peers
 - Limit bandwidth, and time of tests
- More liberal for TCP tests
 - Open for all (or peers)
 - Limit length of tests
- Protect AES keys!
 - If being used

Availability

- Currently available
 - <http://www.internet2.edu/performance/bwctl>
 - <http://software.internet2.edu>
- Mail lists:
 - <https://lists.internet2.edu/sympa/info/bwctl-users>
 - bwctl-users@internet2.edu
 - <https://lists.internet2.edu/sympa/info/bwctl-announce>
 - bwctl-announce@internet2.edu



BWCTL

August 9th 2011, OSG Site Admin Workshop

Jason Zurawski – Internet2 Research Liaison

For more information, visit <http://www.internet2.edu/workshops/npw>

Tester Applications

- Iperf is primary “tester”
 - Well known – widely used
- Problems integrating exec'd tool
 - Server initialization (port number allocation)
 - error conditions
 - No indication of partial progress (How full was the send buffer when the session was killed?)
- thrulay/nuttcp are available also

Testing with no “Local” Server

