August 9th 2011, OSG Site Admin Workshop

Jason Zurawski – Internet2 Research Liaison

Welcome & Performance Primer

Who are we, Who are you?

Agenda

- Welcome and Thanks
  - http://www.internet2.edu/workshops/npw/roster/neren.cfm
- Tutorial Agenda:
  - Network Performance Primer - Why Should We Care? (**30 Mins**)
  - Introduction to Measurement Tools (**20 Mins**)
  - Use of NTP for network measurements (**15 Mins**)
  - Use of the BWCTL Server and Client (**25 Mins**)
  - Use of the OWAMP Server and Client (**25 Mins**)
  - Use of the NDT Server and Client (**25 Mins**)
  - perfSONAR Topics (**30 Mins**)
  - Diagnostics vs Regular Monitoring (**20 Mins**)
  - Use Cases (**30 Mins**)
  - Exercises

Your Goals?

- What are your goals for this workshop?
  - Experiencing performance problems?
  - Responsible for the campus/lab network?
  - Learning about state of the art, e.g. 'What is perfSONAR'?
  - Developing or researching performance tools?
- Is there a Magic Bullet?
  - No, but we can give you access to strategies and tools that will help
  - Patience and diligence will get you to most goals
- This workshop is as much a learning experience for me as it is for you
  - What problem/problems need to be solved
  - What will make networking a less painful experience
  - How can we improve our goods/services

Problem: "The Network Is Broken"

- How can your users effectively report problems?
  - And how can you learn to take them seriously…
- How can users and the local administrators effectively solve multi-domain problems?
  - Eliminate the 'who you know' network to finding resources
  - Automate things when applicable
- Components:
  - Tools to use
  - Questions to ask
  - Methodology to follow
  - How to ask for (and receive) help

Motivation

- Proactive vs Reactive Positions
  - Do you want to find problems before the users do?
  - Can monitoring tools help in other aspects of operations?
    - Capacity Planning
    - Scheduling Maintenance
    - Traffic Engineering
- Automatic user response: “The Network is broken”
  - Is this justified behavior?
    - In actuality, there is a lot of “network” between the applications
    - What about those applications?
    - What about the host itself?
- Lets try to put this into an example …

Motivation – A Typical Scenario

- User and resource are geographically separated
  - Common case: Remote instrument + distributed users
- Both have access to high speed communication network
  - LAN infrastructure - 1Gbps Ethernet
  - WAN infrastructure – 10Gbps Optical Backbone

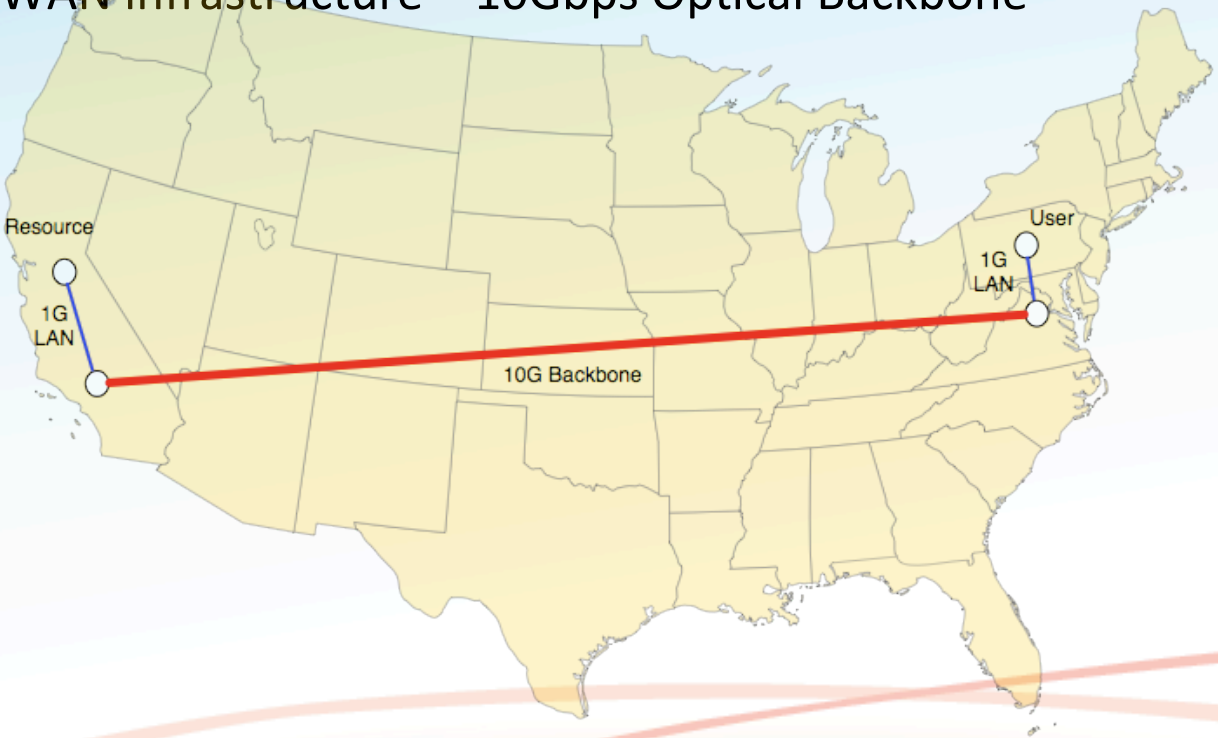

Motivation – A Typical Scenario

- User wants to access a file at the resource (e.g. ~600MB)
- Plans to use COTS tools (e.g. "scp", but could easily be something scientific like "GridFTP" or simple like a web browser)
- What are the expectations?
  - 1Gbps network (e.g. *bottleneck* speed on the LAN)
  - 600MB * 8 = 4,800 Mb file
  - User expects *line rate*, e.g. 4,800 Mb / 1000 Mbps = 4.8 Seconds
  - Audience Poll: Is this expectation too high?
- What are the realities?
  - Congestion and other network performance factors
  - Host performance
  - Protocol Performance
  - Application performance

Motivation – A Typical Scenario

- Real Example (New York USA to Los Angeles USA):

```
[zurawski@nms-rthr2 ~]$ scp zurawski@bwctl.losa.net.internet2.edu:pS-Performance
_Toolkit-3.1.1.iso .
pS-Performance_Toolkit-3.1.1.iso              2%   17MB   1.0MB/s   10:05 ETA_
```

- 1MB/s (8Mb/s) ??? 10 Minutes to transfer???
- Seems unreasonable given the investment in technology
  - Backbone network
  - High speed LAN
  - Capable hosts
- Performance realities as network speed decreases:
  - 100 Mbps Speed – 48 Seconds
  - 10 Mbps Speed – 8 Minutes
  - 1 Mbps Speed – 80 Minutes
- How could this happen? More importantly, why are there not more complaints?
- Audience Poll: Would you complain? If so, to whom?
- Brainstorming the above – where should we look to fix this?

Motivation – A Typical Scenario

- Expectation does not even come close to experience, time to debug. Where to start though?
  - Application
    - Have other users reported problems? Is this the most up to date version?
  - Protocol
    - Protocols typically can be tuned on an individual basis, consult your operating system.
  - Host
    - Are the hardware components (network card, system internals) and software (drivers, operating system) functioning as they should be?
  - LAN Networks
    - Consult with the local administrators on status and potential choke points
  - Backbone Network
    - Consult the administrators at remote locations on status and potential choke points (Caveat – do you [should you] know who they are?)

Motivation – A Typical Scenario

- Following through on the previous, what normally happens …
  - Application
    - This step is normally skipped, the application designer will *blame the network*
  - Protocol
    - These settings may not be explored. Shouldn't this be automatic (e.g. autotuning)?
  - Host
    - Checking and diagnostic steps normally stop after establishing connectivity. E.g. "can I ping the other side"
  - LAN Networks
    - Will assure "internal" performance, but LAN administrators will ignore most user complaints and shift blame to upstream sources. E.g. "our network is fine, there are no complaints"
  - Backbone Network
    - Will assure "internal" performance, but Backbone responsibilities normally stop at the demarcation point, blame is shifted to other networks up and down stream

* Denotes Problem Areas from Example

Why Worry About Network Performance?

- Most network design lends itself to the introduction of flaws:
  - Heterogeneous equipment
  - Cost factors heavily into design – e.g. *Get what you pay for*
  - Design heavily favors ***protection*** and ***availability*** over performance
- Communication protocols are not advancing as fast as networks
  - *TCP/IP* is the king of the protocol stack
    - Guarantees reliable transfers
    - Adjusts to failures in the network
    - Adjusts speed to be *fair* for all
- User Expectations
  - ***Big Science*** is prevalent globally
  - "The Network is Slow/Broken" – is this the response to almost any problem? Hardware? Software?
  - Empower users to be more informed/more helpful

"Big" Science

- A Few words on the LHC
  - 17 Mile Circumference "ring" in Switzerland/France
  - Collide opposing beams of particles (3.5 TeV each – 7TeV collision)
  - "Detectors" are present to gather data on the collision (ALICE, ATLAS, CMS, LHCb)
  - Data is stored at CERN (Tier0), and distributed world wide to other Tiers (1, 2, 3) for processing and analysis
    - Different types of data, Raw + several kinds of processed data to find areas of interest.
    - N.B. even the raw data doesn't capture anything – the machine would produce ***1PB*** (!) of data, ***per second*** (!!), if it was unfiltered
    - Typical processed data set (2011) = ***10 – 100 TB***.
  - Tier1s receive and distribute data to Tier2s, Tier2s do the same for Tier3s
  - Each Tier contains storage and processing software/hardware.
    - Goal is to get the data to the lowest tier ***within 4 hours*** (!)

“Big” Science – ATLAS Detector

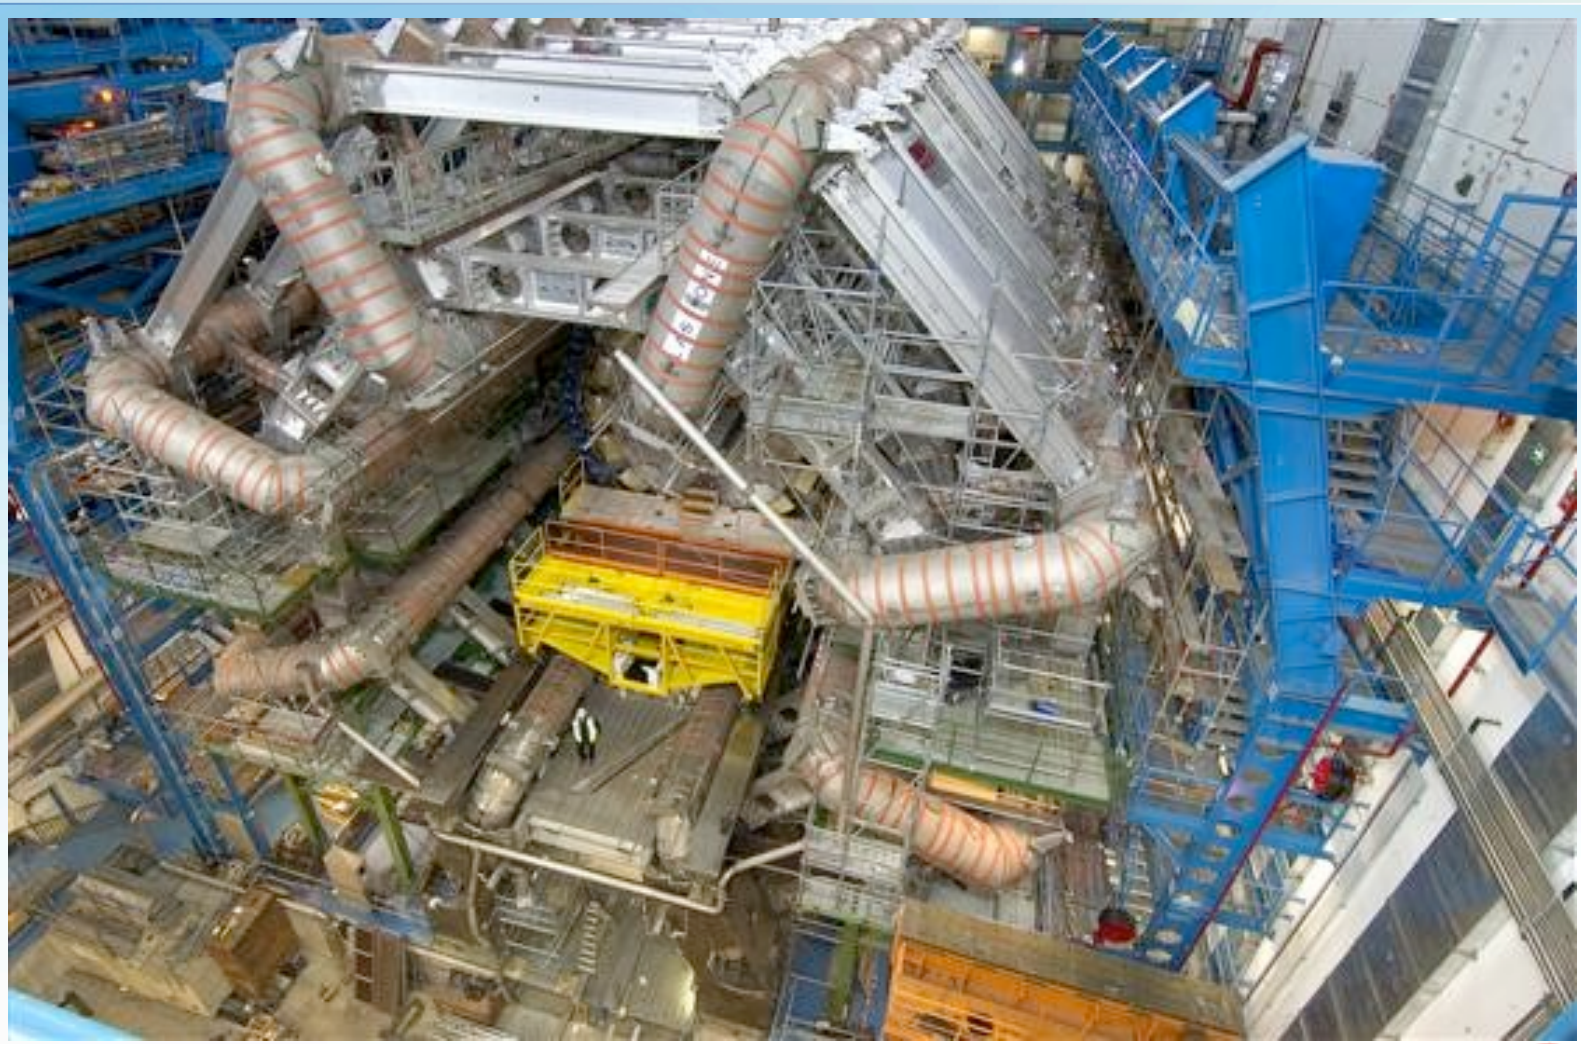

Why is Science Data Movement Different?

- Different Requirements
  - Campus network is not designed for large flows
    - ***Enterprise*** requirements
    - 100s of Mbits is common, any more is rare (or viewed as *strange*)
    - Firewalls
    - Network is designed to mitigate the risks since the common hardware (e.g. Desktops and Laptops) are un-trusted
  - Science is different
    - Network needs to be robust and stable (e.g. predictable performance)
    - 10s of Gbits of traffic (N.B. that its probably not sustained – but could be)
    - Sensitive to enterprise protections (e.g. firewalls, LAN design)
- ***Fixing*** is not easy
  - Design the base network for science, attach the enterprise on the side (expensive, time consuming, and good luck convincing your campus this is necessary…)
  - Mitigate the problems by moving your science equipment to the edge
    - Try to bypass that firewall at all costs
    - Get as close to the WAN connection as you can

Identifying Common Network Problems

- The above examples paint a broad picture: there is a problem, *somewhere*, that needs to be fixed
- What could be out there?
  - Architecture
  - Common Problems, e.g. “Soft Failures”
- Myths and Pitfalls
  - Getting trapped is easy
  - Following a bad lead is easy too

Identifying Common Network Problems

- Audience Question: Would you complain if you knew what you were getting was not correct?

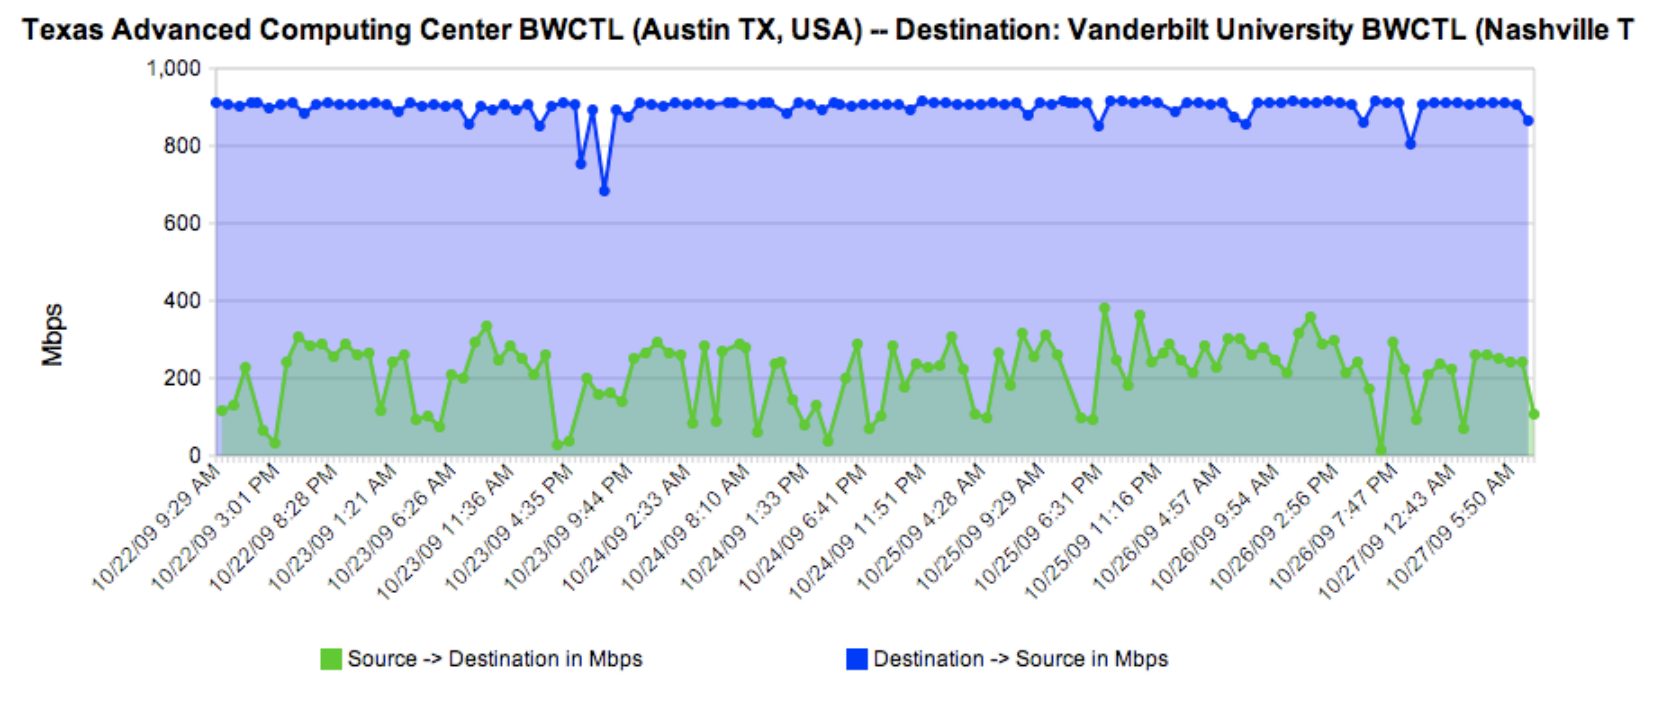

- N.B. Actual performance between Vanderbilt University and TACC – Should be about 1Gbps in both directions.

Identifying Common Network Problems

- Internet2/ESnet engineers will help members and customers debug problems if they are escalated to us
  - Goal is to solve the entire problem – end to end
  - Involves many parties (typical: End users as well as Campus, Regional, Backbone staff)
  - Slow process of locating and testing each segment in the path
  - Have tools to make our job easier (more on this later)
- Common themes and patterns for ***almost every*** debugging exercise emerge
  - Architecture (e.g. LAN design, Equipment Choice, Firewalls)
  - Configuration
  - "Soft Failures", e.g. something that doesn't severe connectivity, but makes the experience unpleasant

Architectural Considerations

- LAN vs WAN Design
  - Multiple Gbit flows [to the outside] should be close to the WAN connection
  - Eliminate the number of hops/devices/physical wires that may slow you down
  - Great performance on the LAN != Great performance on the WAN
- *You Get What you Pay For*
  - Cheap equipment will let you down
  - Network
    - Small Buffers, questionable performance (e.g. internal switching fabric can't keep up w/ LAN demand let alone WAN)
    - Lack of diagnostic tools (SNMP, etc.)
  - Storage
    - Disk throughput needs to be high enough to get everything on to the network
    - Plunking a load of disk into an incapable server is not great either
      - Bus performance
      - Network Card(s)

Architectural Considerations – cont.

- Firewalls
  - Designed to stop traffic
    - read this slowly a couple of times…
  - Small buffers
    - Concerned with protecting the network, not impacting your performance
  - Will be ***a lot*** slower than the original wire speed
  - A "***10G Firewall***" may handle 1 flow close to 10G, doubtful that it can handle a couple.
  - If *firewall-like* functionality is a must – consider using router filters instead

Configuration

- Host Configuration
  - Tune your hosts (especially compute/storage!)
  - Changes to several parameters can yield 4 – 10X improvement
  - Takes minutes to implement/test
  - Instructions: http://fasterdata.es.net/tuning.html
- Network Switch/Router Configuration
  - ***Out of the box*** configuration may include small buffers
  - Competing Goals: Video/Audio etc. needs small buffers to remain responsive. Science flows need large buffers to push more data into the network.
  - Read your manuals and test LAN host to a WAN host to verify (not LAN to LAN).

Host Configuration

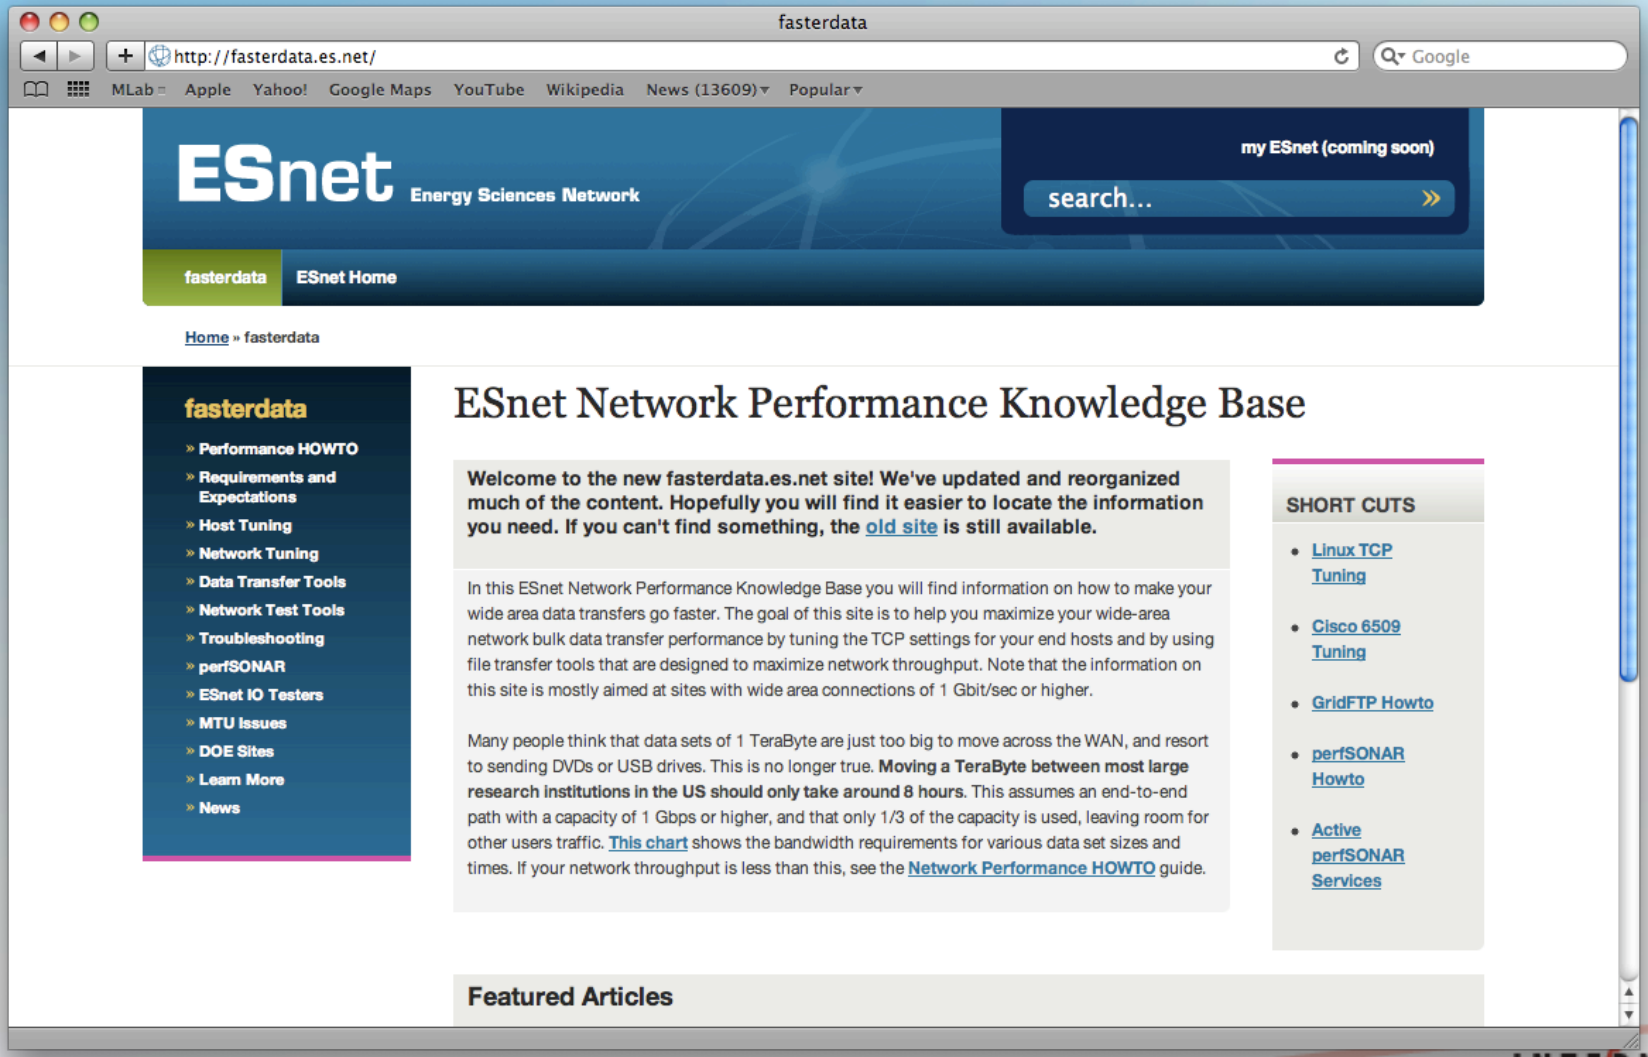

Configuration – cont.

- Host Configuration – spot when the settings were tweaked…

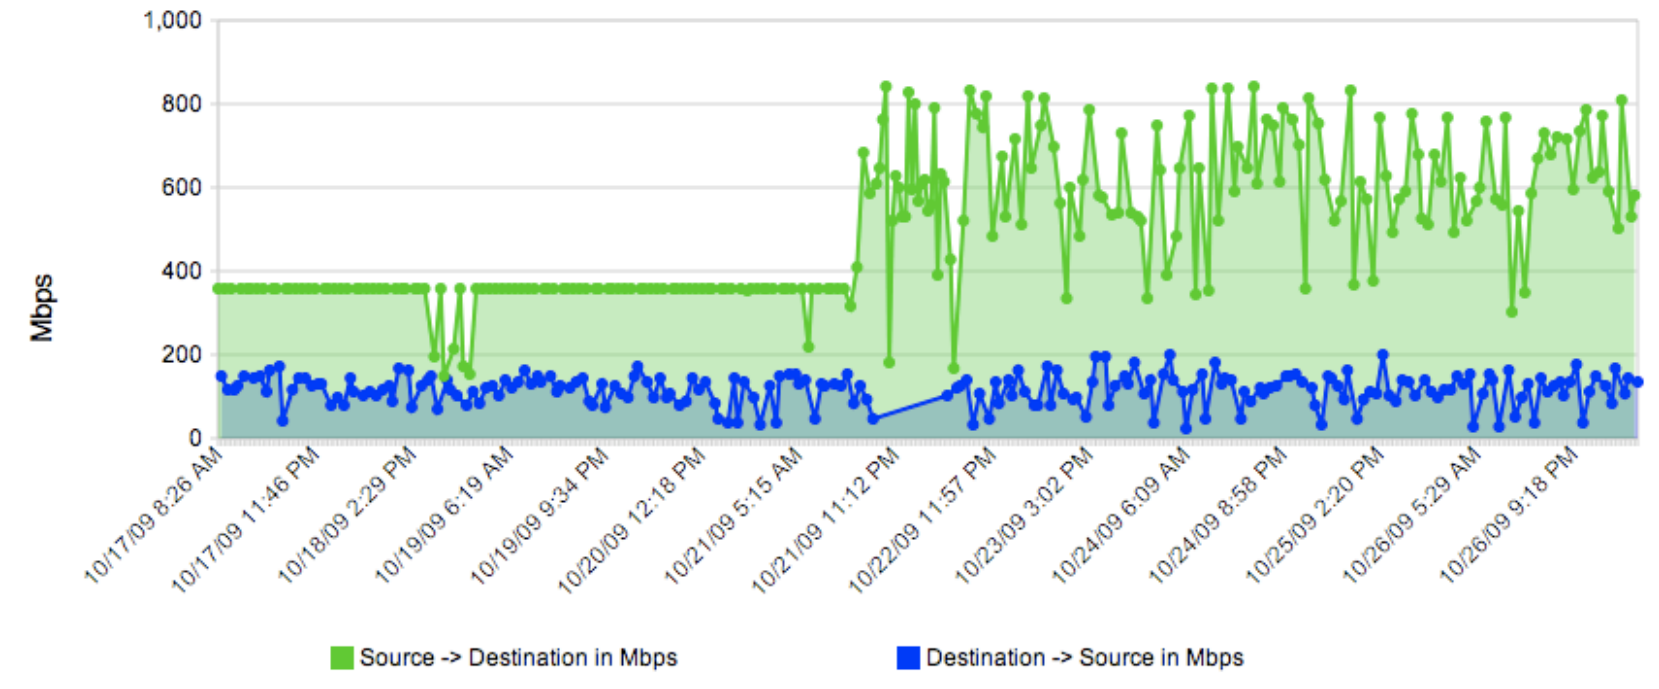

- N.B. Example Taken from REDDnet (UMich to TACC), using BWCTL measurement)

Soft Failures

- ***Soft Failures*** are any network problem that does not result in a loss of connectivity
  - Slows down a connection
  - Hard to diagnose and find
  - May go unnoticed by LAN users in some cases, but remote users may be the ones complaining
    - Caveat – How much time/energy do you put into listing to complaints of remote users?
- Common:
  - Dirty or Crimped Cables
  - Failing Optics/Interfaces
  - [Router] Process Switching, aka "*Punting*"
  - Router Configuration (Buffers/Queues)

Soft Failures – cont.

- Dirty or Crimped Cables and Failing Optics/Interfaces
  - Throw off very low levels of loss – may not notice on a LAN, will notice on the WAN
  - Will be detected with passive tools (e.g. SNMP monitoring)
  - Question: Would you fix it if you knew it was broken?
- [Router] Process Switching
  - "Punt" traffic to a slow path
  - Duplicate traffic onto multiple paths
- Router Configuration (Buffers/Queues)
  - Need to be large enough to handle science flows
  - Routing table overflow (e.g. system crawls to a halt when memory is exhausted)

Myths and Pitfalls

- "My LAN performance is great, WAN is probably the same"
  - TCP recovers from loss/congestion quickly on the LAN (low RTT)
  - TCP will cut speed in half for every loss/discard on the WAN – will take a long time to recover for a large RTT/
  - Small levels of loss on the LAN (ex. 1/1000 packets) will go unnoticed, will be very noticeable on the WAN.
- "Ping is not showing loss/latency differences"
  - ICMP May be blocked/ignored by some sites
  - Routers process ICMP differently than other packets (e.g. may show phantom delay)
  - ICMP may hide some (not all) loss.
  - Will not show asymmetric routing delays (e.g. taking a different path on send vs receive)
- Our goal is to dispel these and others by educating the proper way to verify a network – we have lots of tools at our disposal but using these in the appropriate order is necessary too

Topics of Discussion in this Workshop

- Diagnosis Methodology
- Partial Path Decomposition
- Systematic Troubleshooting
- On Demand vs Regular Testing

Topics of Discussion

- Diagnosis Methodology
  - Find a measurement server "near me"
    - Why is this important?
    - How hard is this to do?
  - Encourage user to participate in diagnosis procedures
  - Detect and report common faults in a manner that can be shared with admins/NOC
    - 'Proof' goes a long way
  - Provide a mechanism for admins to review test results
  - Provide feedback to user to ensure problems are resolved

Topics of Discussion – cont.

- Partial Path Decomposition
  - Networking is increasingly:
    - Cross domain
    - Large scale
    - Data intensive
  - Identification of the end-to-end path is key (must solve the problem end to end…)
  - Discover measurement nodes that are “near” this path
  - Provide proper authentication or receive limited authority to run tests
    - No more conference calls between 5 networks, in the middle of the night
  - Initiate tests between various nodes
  - Retrieve and store test data for further analysis

Topics of Discussion – cont.

- Systematic Troubleshooting
  - Having tools deployed (along the entire path) to enable adequate troubleshooting
  - Getting end-users involved in the testing
  - Combining output from multiple tools to understand problem
    - Correlating diverse data sets – only way to understand complex problems.
  - Ensuring that results are adequately documented for later review
- On Demand vs Regular Testing
  - On-Demand testing can help solve existing problems once they occur
  - Regular performance monitoring can quickly identify and locate problems before users complain
    - Alarms
    - Anomaly detection
  - Testing and measuring performance increases the value of the network to all participants

Our Goals

- To spread the word that today's networks really can, do, and will support demanding applications
  - Science
    - Physics
      - LHC, LIGO
    - Astronomy
      - LSST, SDSS, eVLBI
    - Biology and Climate
      - Genome Sequencing, Weather simulations, remote senors
  - Arts and Humanities
    - Distance learning, synchronized performance
  - Computational and Network Research
    - DYNES, GENI, MeasurementLab, etc.
- To increase the number of test points
  - Instrumenting the end to end path is key
  - Spread the knowledge and encourage adoption

Final Thoughts

- See a talk from the recent Joint Techs Conference:
  - http://www.internet2.edu/presentations/jt2010july/20100714-metzger-whatnext.pdf
- Take home points:
  - Close to $1 Billion USD spent on networking at all levels (Campus, Regional, Backbone) in the next 2 years due to ARRA Funding
  - Unprecedented access and capacity for many people
  - Ideal View:
    - Changes will be seamless
    - Completed on time
    - Bandwidth will solve all performance problems
  - Realistic View:
    - Network 'breaks' when it is touched (e.g. new equipment, configs)
    - Optimization will not be done in a global fashion (e.g. backbone fixes performance, but what about regional and campus?)
    - Bandwidth means nothing when you have a serious performance problem

Welcome & Performance Primer

August 9th 2011, OSG Site Admin Workshop

Jason Zurawski – Internet2 Research Liaison

For more information, visit http://www.internet2.edu/workshops/npw